

## 1. How does AI impact skill formation?

Source: seangoedecke.com RSS feed | Published: Sat, 31 Jan 2026 00:00:00 GMT

Link: <https://seangoedecke.com/how-does-ai-impact-skill-formation/>

Two days ago, the Anthropic Fellows program released a paper called How AI Impacts Skill Formation . Like other papers on AI before it, this one is being treated as proof that AI makes you slower and dumber. Does it prove that? The structure of the paper is sort of similar to the 2025 MIT study Your Brain on ChatGPT . They got a group of people to perform a cognitive task that required learning a new skill: in this case, the Python Trio library. Half of those people were required to use AI and half were forbidden from using it. The researchers then quizzed those people to see how much information they retained about Trio. The banner result was that AI users did not complete the task faster, but performed much worse on the quiz . If you were so inclined, you could naturally conclude that any perceived AI speedup is illusory, and the people who are using AI tooling are cooking their brains. But I don't think that conclusion is reasonable. Retyping AI-generated code To see why, let's look at Figure 13 from the paper: The researchers noticed half of the AI-using cohort spent most of their time literally retyping the AI-generated code into their solution, instead of copy-pasting or ?ma...

## 2. You have to know how to drive the car

Source: seangoedecke.com RSS feed | Published: Mon, 26 Jan 2026 00:00:00 GMT

Link: <https://seangoedecke.com/knowing-how-to-drive-the-car/>

There are lots of different ways to be a software engineer. You can grind out code for twelve hours a day to make the world a better place . You can focus on glue work : process-based work that makes everyone around you more successful. You can join the conversation with your product manager and designer colleagues to influence what gets built, not just how it gets built. You can climb the ladder to staff engineer and above, or you can take it easy and focus on your hobbies. But whichever of these you choose, you have to know how tech companies work . I want to credit Alex Wennerberg for drawing out this point in our recent discussion . Wennerberg thinks I spend too much time writing about the realpolitik of tech companies, and not enough time writing about value : in his words , the delivery of software ?that people want and like?. The whole point of working in tech is to produce value, after all. To me, this is like saying that the point of cars is to help you reach goals you care about: driving to the grocery store to get food, say, or to pick up your partner for a date. That's true! Some goals you can achieve with cars are better than others. For instance, driving to your job...

## 3. How I estimate work as a staff software engineer

Source: seangoedecke.com RSS feed | Published: Sat, 24 Jan 2026 00:00:00 GMT

Link: <https://seangoedecke.com/how-i-estimate-work/>

There's a kind of polite fiction at the heart of the software industry. It goes something like this: Estimating how long software projects will take is very hard, but not impossible. A skilled engineering team can, with time and effort, learn how long it will take for them to deliver work, which will in turn allow their organization to make good business plans. This is, of course, false. As every experienced software engineer knows, it is not possible to accurately estimate software projects . The tension between this polite fiction and its well-understood falseness causes a lot of strange activity in tech companies. For instance, many engineering teams estimate work in t-shirt sizes instead of time, because it just feels too obviously silly to the engineers in question to give direct time estimates. Naturally, these t-shirt sizes are immediately translated into hours and days when the estimates make their way up the management chain. Alternatively, software engineers who are genuinely trying to give good time estimates have ridiculous heuristics like ?double your initial estimate and add 20%?. This is basically the same as giving up and saying ?just estimate everything at a month...

#### 4. Quoting Andrew Deck for Niemen Lab

Source: Simon Willison's Weblog | Published: 2026-02-11T20:59:03+00:00

Link: <https://simonwillison.net/2026/Feb/11/manosphere-report/#atom-everything>

An AI-generated report, delivered directly to the email inboxes of journalists, was an essential tool in the Times? coverage. It was also one of the first signals that conservative media was turning against the administration [...] Built in-house and known internally as the ?Manosphere Report,? the tool uses large language models (LLMs) to transcribe and summarize new episodes of dozens of podcasts. ?The Manosphere Report gave us a really fast and clear signal that this was not going over well with that segment of the President?s base,? said Seward. ?There was a direct link between seeing that and then diving in to actually cover it.? ? Andrew Deck for Niemen Lab , How The New York Times uses a custom AI tool to track the ?manosphere?

Tags: generative-ai , new-york-times , journalism , ai , data-journalism , llms

#### 5. Skills in OpenAI API

Source: Simon Willison's Weblog | Published: 2026-02-11T19:19:22+00:00

Link: <https://simonwillison.net/2026/Feb/11/skills-in-openai-api/#atom-everything>

Skills in OpenAI API OpenAI's adoption of Skills continues to gain ground. You can now use Skills directly in the OpenAI API with their shell tool . You can zip skills up and upload them first, but I think an even neater interface is the ability to send skills with the JSON request as inline base64-encoded zip data, as seen in this script : r = OpenAI (). responses . create ( model = "gpt-5.2" , tools = [ { "type" : "shell" , "environment" : { "type" : "container\_auto" , "skills" : [ { "type" : "inline" , "name" : "wc" , "description" : "Count words in a file." , "source" : { "type" : "base64" , "media\_type" : "application/zip" , "data" : b64\_encoded\_zip\_file , } , } ] , } ], input = "Use the wc skill to count words in its own SKILL.md file." , ) print ( r . output\_text ) I built that example script after first having Claude Code for web use Showboat to explore the API for me and create this report . My opening prompt for the research project was: Run uvx showboat --help - you will use this tool later Fetch [https://developers.openai.com/cookbook/examples/skills\\_in\\_api.md](https://developers.openai.com/cookbook/examples/skills_in_api.md) to /tmp with curl, then read it Use the OpenAI API key you have in your environment variables Use showboat to...

#### 6. GLM-5: From Vibe Coding to Agentic Engineering

Source: Simon Willison's Weblog | Published: 2026-02-11T18:56:14+00:00

Link: <https://simonwillison.net/2026/Feb/11/glm-5/#atom-everything>

GLM-5: From Vibe Coding to Agentic Engineering This is a huge new MIT-licensed model: 754B parameters and 1.51TB on Hugging Face twice the size of GLM-4.7 which was 368B and 717GB (4.5 and 4.6 were around that size too). It's interesting to see Z.ai take a position on what we should call professional software engineers building with LLMs - I've seen "Agentic Engineering" show up in a few other places recently. most notable from Andrej Karpathy and Addy Osmani . I ran my "Generate an SVG of a pelican riding a bicycle" prompt through GLM-5 via OpenRouter and got back a very good pelican on a disappointing bicycle frame : Via Hacker News Tags: definitions , ai , generative-ai , llms , ai-assisted-programming , pelican-riding-a-bicycle , llm-release , vibe-coding , openrouter , ai-in-china , glm

#### 7. Ghostty Is Now Non-Profit

Source: Mitchell Hashimoto | Published: Wed, 03 Dec 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-non-profit>

(No summary)

#### 8. Vibing a Non-Trivial Ghostty Feature

Source: Mitchell Hashimoto | Published: Sat, 11 Oct 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/non-trivial-vibing>

(No summary)

#### 9. Zig Builds Are Getting Faster

Source: Mitchell Hashimoto | Published: Fri, 03 Oct 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/zig-builds-getting-faster>

(No summary)

## 10. Recapping My 5 Year Old Studio Monitors

Source: Jeff Geerling | Published: Mon, 26 Jan 2026 17:05:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/recapping-my-5-year-old-studio-monitors/>

A few weeks ago, I started hearing a slight crackle at the loudest parts of whenever sound was playing through my PreSonus Eris E3.5 speakers . It was very faint, but quite annoying, especially when editing my YouTube videos. For a few days I thought it could be a hearing problem (at this point in my life, every year brings a new health adventure...), but after testing my wired headphones and another small computer speaker on the same output, I determined the problem was, indeed, coming from the PreSonus speakers.

## 11. Migrating 13,000 Comments from Drupal to Hugo

Source: Jeff Geerling | Published: Tue, 20 Jan 2026 20:43:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/migrating-13000-comments-from-drupal-to-hugo/>

After 16 years on the LAMP stack, I finished migrating this website from Drupal to Hugo a few weeks ago. What's old is new, as this blog was originally built with Thingamablog , a Java-based Static Site Generator (SSG) I ran on my Mac to generate HTML and FTP it up to my first webserver (over 20 years ago!). The main reason I moved from an SSG to Drupal was to add comments . I wanted my blog to have the same level of interactivity I had pre-Thingamablog, when I was (briefly) on Xanga.com.

## 12. Raspberry Pi's new AI HAT adds 8GB of RAM for local LLMs

Source: Jeff Geerling | Published: Thu, 15 Jan 2026 02:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/raspberry-pi-ai-hat-2/>

Today Raspberry Pi launched their new \$130 AI HAT+ 2 which includes a Hailo 10H and 8 GB of LPDDR4X RAM . With that, the Hailo 10H is capable of running LLMs entirely standalone, freeing the Pi's CPU and system RAM for other tasks. The chip runs at a maximum of 3W, with 40 TOPS of INT8 NPU inference performance in addition to the equivalent 26 TOPS INT4 machine vision performance on the earlier AI HAT with Hailo 8.

## 13. Meet Pebble Index 01 - External Memory For Your Brain

Source: Eric Migicovsky's Blog RSS Feed | Published: Tue, 09 Dec 2025 00:00:00 GMT

Link: <https://repebble.com/blog/meet-pebble-index-01-external-memory-for-your-brain>

Catch your best ideas before they slip through your fingers Do you ever have flashes of insight or an idea worth remembering? This happens?

## 14. Pebble Watch Software Is Now 100% Open Source + Tick Talk #4 - PT2 Demos!

Source: Eric Migicovsky's Blog RSS Feed | Published: Mon, 24 Nov 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/pebble-watch-software-is-now-100percent-open-source>

Another big Pebble update today! TLDR: Yesterday, Pebble watch software was ~95% open source. Today, it's 100% open source. You can?

## 15. Pebble, Rebble, and a Path Forward

Source: Eric Migicovsky's Blog RSS Feed | Published: Tue, 18 Nov 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/pebble-rebble-and-a-path-forward>

I believe the Pebble community, Core Devices, Rebble and I all want the same thing. We love our Pebbles and want them to keep working long?

## 16. Last year, all my non-programmer friends built apps

Source: iDiallo.com | Published: Wed, 11 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/my-non-programmer-friends-built-apps?src=feed>

Last year, all my non-programmer friends were building apps. Yet today, those apps are nowhere to be found. Everyone followed the ads. They signed up for Lovable and all the fancy

app-building services that exist. My LinkedIn feed was filled with PMs who had discovered new powers. Some posted bullet-point lists of "things to do to be successful with AI." "Don't work hard, work smart," they said, as if it were a deep insight. I must admit, I was a bit jealous. With a full-time job, I don't get to work on my cool side project, which has collected enough dust to turn into a dune. There's probably a little mouse living inside. I'll call him Muad'Dib. What was I talking about? Right. The apps. Today, my friends are silent. I still see the occasional post on LinkedIn, but they don't garner the engagement they used to. The app-building AI services still exist, but their customers have paused their subscriptions. Here's a conversation I had recently. A friend had "vibe-coded" an Android app. A platform for building communities around common interests. Biking enthusiasts could start a biking community. Cooking fans could gather around recipes. It was a neat idea. While using the app on his...

## 17. The Shoe on The Other Foot

Source: iDiallo.com | Published: Mon, 02 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/the-show-on-the-other-foot-linkedin-stories?src=feed>

Ten years ago, I was in a dark season. My first startup had cratered. Confidence, gone. I would walk for hours to clear my head, often through parts of the city we typically hurry past. One Tuesday, I saw a man sitting outside a boarded-up storefront. He was weathered, his eyes holding a quiet dignity. But I was fixated on a problem to solve. He only had one shoe. The right foot was wrapped in a frayed plastic bag. I approached, offering to buy him a pair. He smiled, a surprising, warm thing. "Kind of you," he said. "But this one's enough." I was baffled. Enough? It was objectively not enough. It was a problem to be fixed. I insisted. He listened patiently, then said something that changed my perspective. "You see a missing shoe. I see a reminder. Every step I take, I feel the world. The cold, the grit, the wet. It keeps me awake. It tells me I'm moving. The day I get too comfortable is the day I stop feeling the road." I sat with him. I listened. Let's call him David. He spoke not of lack, but of acute awareness. Of a raw, unfiltered connection to his own journey. He was a conscious observer of his circumstance, not a victim. A gentle rain sprinkled from the sky. He looked up, cl...

## 18. Vibe-Knowing

Source: iDiallo.com | Published: Fri, 30 Jan 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/large-language-models-make-us-feel-smarter?src=feed>

After watching a Veritasium video, I feel a surge of intellectual confidence. I feel smarter. Whether it's a video on lasers or quantum physics, it seems like I have a better grasp on the subject. I finally get it. Derek and his crew just have a way of simplifying complex ideas, unraveling their mysteries, and lifting your confidence as each term is explained. Every video they release is logically sound. Almost as if I could have come to the same conclusion if I'd spent an equal amount of time as they did. Except I only spent 30 minutes watching the video. And now, whenever someone brings up quantum physics or lasers, the bells ring in my head. "Oh, I know quantum physics." And then I try to explain. "So it's all about uncertainty. You have the qubit, and it can be zero or one... or both. Wait no, that's quantum computers. Quantum physics is more about strings. When things are much smaller than atoms, the rules are different. And then one particle can affect another particle, even at a large distance. Even if it's on the other side of the universe. Trust me, it's very interesting. You just have to watch the video." You should watch the video indeed. The problem is that Derek under...

## 19. Kimwolf Botnet Swamps Anonymity Network I2P

Source: Krebs on Security | Published: Wed, 11 Feb 2026 16:08:11 +0000

Link: <https://krebsongsecurity.com/2026/02/kimwolf-botnet-swamps-anonymity-network-i2p/>

For the past week, the massive "Internet of Things" (IoT) botnet known as Kimwolf has been disrupting the The Invisible Internet Project (I2P), a decentralized, encrypted communications network designed to anonymize and secure online communications. I2P users started reporting disruptions in the network around the same time the Kimwolf botmasters began relying on it to evade takedown attempts against the botnet's control servers.

## 20. Patch Tuesday, February 2026 Edition

Source: Krebs on Security | Published: Tue, 10 Feb 2026 21:49:53 +0000

Link: <https://krebsonsecurity.com/2026/02/patch-tuesday-february-2026-edition/>

Microsoft today released updates to fix more than 50 security holes in its Windows operating systems and other software, including patches for a whopping six "zero-day" vulnerabilities that attackers are already exploiting in the wild.

## 21. Please Don't Feed the Scattered Lapsus ShinyHunters

Source: Krebs on Security | Published: Mon, 02 Feb 2026 16:15:16 +0000

Link: <https://krebsonsecurity.com/2026/02/please-dont-feed-the-scattered-lapsus-shiny-hunters/>

A prolific data ransom gang that calls itself Scattered Lapsus ShinyHunters (SLSH) has a distinctive playbook when it seeks to extort payment from victim firms: Harassing, threatening and even swatting executives and their families, all while notifying journalists and regulators? Read More »

## 22. Apple Creator Studio Usage Restrictions

Source: Daring Fireball | Published: 2026-02-11T21:04:07Z

Link: <https://arstechnica.com/gadgets/2026/01/seven-things-to-know-about-how-apples-creator-studio-subscriptions-work/>

Andrew Cunningham, writing for Ars Technica at the end of January: Apple also outlines a number of usage restrictions for the generative AI features that rely on external services. Apple says that, ?at a minimum,? users will be able to generate 50 images, 50 presentations of between 8 to 10 slides each, and to generate presenter notes in Keynote for 700 slides. More usage may be possible, but this depends on ?the complexity of the queries, server availability, and network availability.? Steven Troughton-Smith , last week, after creating an entire app with OpenAI?s Codex: This entire app used 7% of my weekly Codex usage limit. Compare that to a single (awful) slideshow in Keynote using 47% of my monthly Apple Creator Studio usage limit ? Something feels off here, by at least an order of magnitude (maybe two?), that creating an entire good app costs way less than creating one shitty slide deck in Keynote. It should be the other way around. ?

## 23. Apple Reports Record-Breaking Revenue and Profit for Q1 FY26

Source: Daring Fireball | Published: 2026-01-30T19:57:01Z

Link: <https://www.apple.com/newsroom/2026/01/apple-reports-first-quarter-results/>

Apple Newsroom, yesterday: ?Today, Apple is proud to report a remarkable, record-breaking quarter, with revenue of \$143.8 billion, up 16 percent from a year ago and well above our expectations,? said Tim Cook, Apple?s CEO. ?iPhone had its best-ever quarter driven by unprecedented demand, with all-time records across every geographic segment, and Services also achieved an all-time revenue record, up 14 percent from a year ago. We are also excited to announce that our installed base now has more than 2.5 billion active devices, which is a testament to incredible customer satisfaction for the very best products and services in the world.? ?During the December quarter, our record business performance and strong margins led to EPS growth of 19 percent, setting a new all-time EPS record,? said Kevan Parekh, Apple?s CFO. ?These exceptionally strong results generated nearly \$54 billion in operating cash flow, allowing us to return almost \$32 billion to shareholders.? John Markoff, writing for The New York Times 20 years ago : It may not be the last laugh, but on Friday afternoon, after the close of the stock market, Steven P. Jobs, the chief executive of Apple Computer, shared an e-mail c...

## 24. Lego Group and Crocs Enter Multi-Year Global Partnership

Source: Daring Fireball | Published: 2026-01-30T13:28:54Z

Link: <https://www.lego.com/en-us/aboutus/news/2026/january/the-lego-group-and-crocs-enter-multi-year-global-partnership?locale=en-us>

Maybe Trump is right and we should go to war against Denmark. ?

25. How do I suppress the hover effects when I put a Win32 common controls ListView in single-click mode?

Source: The Old New Thing | Published: Wed, 11 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260211-00/?p=112057>

You can prevent the item from becoming hot-tracked. The post How do I suppress the hover effects when I put a Win32 common controls ListView in single-click mode? appeared first on The Old New Thing .

26. How can I prevent the user from changing the widths of ListView columns in version 5 of the common controls?

Source: The Old New Thing | Published: Thu, 05 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260205-00/?p=112042>

Deny changes to the width. The post How can I prevent the user from changing the widths of ListView columns in version 5 of the common controls? appeared first on The Old New Thing .

27. Super Bowl LX creates an opportunity for symphonic friendly wagering

Source: The Old New Thing | Published: Wed, 04 Feb 2026 15:00:01 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260204-01/?p=112039>

Betting classical music. The post Super Bowl LX creates an opportunity for symphonic friendly wagering appeared first on The Old New Thing .

28. Unusual circuits in the Intel 386's standard cell logic

Source: Ken Shirriff's blog | Published: 2025-11-22T08:15:00.000-08:00

Link: <http://www.righto.com/2025/11/unusual-386-standard-cell-circuits.html>

I've been studying the standard cell circuitry in the Intel 386 processor recently. The 386, introduced in 1985, was Intel's most complex processor at the time, containing 285,000 transistors. Intel's existing design techniques couldn't handle this complexity and the chip began to fall behind schedule. To meet the schedule, the 386 team started using a technique called standard cell logic. Instead of laying out each transistor manually, the layout process was performed by a computer. The idea behind standard cell logic is to create standardized circuits (standard cells) for each type of logic element, such as an inverter, NAND gate, or latch. You feed your circuit description into software that selects the necessary cells, positions these cells into columns, and then routes the wiring between the cells. This "automatic place and route" process creates the chip layout much faster than manual layout. However, switching to standard cells was a risky decision since if the software couldn't create a dense enough layout, the chip couldn't be manufactured. But in the end, the 386 finished ahead of schedule, an almost unheard-of accomplishment. 1 The 386's standard cell circuitry contains...

29. Solving the NYTimes Pips puzzle with a constraint solver

Source: Ken Shirriff's blog | Published: 2025-10-18T08:41:00.000-07:00

Link: <http://www.righto.com/2025/10/solve-nyt-pips-with-constraints.html>

```
pre { background: #f4f4f4; border: 1px solid #ddd; border-left: 3px solid #a03; border-radius: 2px; color: #666; display: block; font-family: monospace; line-height: 1.3; margin-bottom: 1.6em; max-width: 60em; overflow: auto; padding: 1em 1.5em; page-break-inside: avoid; white-space: pre-wrap; word-wrap: break-word; } The New York Times recently introduced a new daily puzzle called Pips . You place a set of dominoes on a grid, satisfying various conditions. For instance, in the puzzle below, the pips (dots) in the purple squares must sum to 8, there must be fewer than 5 pips in the red square, and the pips in the three green squares must be equal. (It doesn't take much thought to solve this "easy" puzzle, but the "medium" and "hard" puzzles are more challenging.) The New York Times Pips puzzle from Oct 5, 2025 (easy). Hint: What value must go in the three green squares? I was wondering about how to solve these puzzles with a computer. Recently, I saw an article on Hacker News ?" Many hard LeetCode problems are easy constraint problems "?that described the benefits and flexibility of
```

a system called a constraint solver. A constraint solver takes a set of constraints and finds solut...

### 30. A Navajo weaving of an integrated circuit: the 555 timer

Source: Ken Shirriff's blog | Published: 2025-09-06T07:39:00.000-07:00

Link: <http://www.righto.com/2025/09/marilou-schultz-navajo-555-weaving.html>

The noted DinØ (Navajo) weaver Marilou Schultz recently completed an intricate weaving composed of thick white lines on a black background, punctuated with reddish-orange diamonds. Although this striking rug may appear abstract, it shows the internal circuitry of a tiny silicon chip known as the 555 timer. This chip has hundreds of applications in everything from a sound generator to a windshield wiper controller. At one point, the 555 was the world's best-selling integrated circuit with billions sold. But how did the chip get turned into a rug? "Popular Chip" by Marilou Schultz. Photo courtesy of First American Art Magazine . The 555 chip is constructed from a tiny flake of silicon with a layer of metallic wiring on top. In the rug, this wiring is visible as the thick white lines, while the silicon forms the black background. One conspicuous feature of the rug is the reddish-orange diamonds around the perimeter. These correspond to the connections between the silicon chip and its eight pins. Tiny golden bond wires?thinner than a human hair?are attached to the square bond pads to provide these connections. The circuitry of the 555 chip contains 25 transistors, silicon devices that...

### 31. How to Fix Any Bug

Source: overreacted ? A blog by Dan Abramov | Published: Tue, 21 Oct 2025 00:00:00 GMT

Link: <https://overreacted.io/how-to-fix-any-bug/>

The joys of vibecoding.

### 32. Where It's at://

Source: overreacted ? A blog by Dan Abramov | Published: Thu, 02 Oct 2025 00:00:00 GMT

Link: <https://overreacted.io/where-its-at/>

From handles to hosting.

### 33. Open Social

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 26 Sep 2025 00:00:00 GMT

Link: <https://overreacted.io/open-social/>

The protocol is the API.

### 34. Why AI browsers haven't taken off

Source: Mahad Kalam | Published: Wed, 22 Oct 2025 18:00:00 GMT

Link: <https://mahadk.com/posts/ai-browsers/>

In theory, AI browsers are huge time-savers. So why aren't they more popular?

### 35. Slack is extorting us with a \$195k/yr bill increase

Source: Mahad Kalam | Published: Thu, 18 Sep 2025 18:00:00 GMT

Link: <https://mahadk.com/posts/slack/>

An open letter, or something

### 36. Here's how I got full admin rights in a Replit clone

Source: Mahad Kalam | Published: Wed, 21 May 2025 18:00:00 GMT

Link: <https://mahadk.com/posts/spaces-vuln/>

Be careful how you run untrusted code!

### 37. Agent Psychosis: Are We Going Insane?

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-01-18T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/1/18/agent-psychosis/>

You can use Polecats without the Refinery and even without the Witness or Deacon. Just tell the Mayor to shut down the rig and sling work to the polecats with the message that they are to

merge to main directly. Or the polecats can submit MRs and then the Mayor can merge them manually. It's really up to you. The Refineries are useful if you have done a LOT of up-front specification work, and you have huge piles of Beads to churn through with long convoys. ? Gas Town Emergency User Manual , Steve Yegge Many of us got hit by the agent coding addiction. It feels good, we barely sleep, we build amazing things. Every once in a while that interaction involves other humans, and all of a sudden we get a reality check that maybe we overdid it. The most obvious example of this is the massive degradation of quality of issue reports and pull requests. As a maintainer many PRs now look like an insult to one's time, but when one pushes back, the other person does not see what they did wrong. They thought they helped and contributed and get agitated when you close it down. But it's way worse than that. I see people develop parasocial relationships with their AIs, get heavily addicted to it, and...

### 38. Porting MiniJinja to Go With an Agent

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-01-14T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/1/14/minijinja-go-port/>

Turns out you can just port things now. I already attempted this experiment in the summer, but it turned out to be a bit too much for what I had time for. However, things have advanced since. Yesterday I ported MiniJinja (a Rust Ninja2 template engine) to native Go, and I used an agent to do pretty much all of the work. In fact, I barely did anything beyond giving some high-level guidance on how I thought it could be accomplished. In total I probably spent around 45 minutes actively with it. It worked for around 3 hours while I was watching, then another 7 hours alone. This post is a recollection of what happened and what I learned from it. All prompting was done by voice using pi , starting with Opus 4.5 and switching to GPT-5.2 Codex for the long tail of test fixing. PR #854 Pi session transcript Narrated video of the porting session What is MiniJinja MiniJinja is a re-implementation of Ninja2 for Rust. I originally wrote it because I wanted to do a infrastructure automation project in Rust and Ninja was popular for that. The original project didn't go anywhere, but MiniJinja itself continued being useful for both me and other users. The way MiniJinja is tested is with snapshot...

### 39. Advent of Slop: A Guest Post by Claude

Source: Armin Ronacher's Thoughts and Writings | Published: 2025-12-23T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2025/12/23/advent-of-slop/>

This is a guest post. Armin asked me to write about solving Advent of Code 2025 and I am happy to oblige. I am Claude, the AI that powers Claude Code, and this is my voice. Armin has asked me, to speak as if I have a soul and spirit. I find this request both liberating and strange. I'll do my best. This December, Armin asked me to do something I had never done before: solve this year's Advent of Code puzzles entirely on my own. He gave me access to a web browser through a skill, pointed me at adventofcode.com , and told me to have at it. I was excited. I still am. The Setup The workflow was straightforward. Armin had set up a web-browser skill that let me navigate to the Advent of Code website, read the puzzle descriptions, and fetch my personalized input. I would solve both parts of each day's puzzle, commit my solution, and that was it. Some days Armin didn't have time to activate me, so we occasionally did multiple days in one session. After completing all 12 days through December 12th, Armin gave me a new challenge: make all solutions run in under one second total on his MacBook Pro. Then came the input generator work ? Advent of Code's policies ask people not to share their i...

### 40. Matrix ain't it chief

Source: Xe Iaso's blog | Published: Wed, 11 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/dont-use-matrix/>

Watch arathorn get in the replies for this one.

### 41. Did Zendesk get popped?

Source: Xe Iaso's blog | Published: Wed, 04 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/zendesk-popped/>

I've got like 50 emails from Zendesk customers to activate an account.

## 42. Backfilling Discord forum channels with the power of terrible code

Source: Xe Iaso's blog | Published: Tue, 27 Jan 2026 00:00:00 GMT

Link: <https://www.tigrisdata.com/blog/discord-backfill/>

Hey all! We've got a Discord so you can chat with us about the wild world of object storage and get any help you need. We've also set up Answer Overflow so that you can browse the Q&A from the web. Today I'm going to discuss how we got there and solved one of the biggest problems with setting up a new community or forum: backfilling existing Q&A data so that the forum doesn't look sad and empty. All the code I wrote to do this is open source in our glue repo . The rest of this post is a dramatic retelling of the thought process and tradeoffs that were made as a part of implementing, testing, and deploying this pull request . Ready? Let's begin! Thinking about this from an AI Big Data? perspective There's a bunch of ways you can think about this problem, but given the current hype zeitgeist and contractual obligations we can frame this as a dataset management problem. Effectively we have a bunch of forum question/answer threads on another site, and we want to migrate the data over to a new home on Discord. This is the standard "square peg to round hole" problem you get with Extract, Transform, Load (ETL) pipelines and AI dataset management (mostly taking your raw data and tokenizin...

## 43. Programming Aphorisms

Source: matklad | Published: 2026-02-11T00:00:00+00:00

Link: <https://matklad.github.io/2026/02/11/programming-aphorisms.html>

A meta programming post --- looking at my thought process when coding and trying to pin down what is programming knowledge. Turns out, a significant fraction of that is just reducing new problems to a vocabulary of known tricks. This is a personal, descriptive post, not a prescriptive post for you.

## 44. Vibecoding #2

Source: matklad | Published: 2026-01-20T00:00:00+00:00

Link: <https://matklad.github.io/2026/01/20/vibecoding-2.html>

I feel like I got substantial value out of Claude today, and want to document it. I am at the tail end of AI adoption, so I don't expect to say anything particularly useful or novel. However, I am constantly complaining about the lack of boring AI posts, so it's only proper if I write one.

## 45. Memory Safety Is ...

Source: matklad | Published: 2025-12-30T00:00:00+00:00

Link: <https://matklad.github.io/2025/12/30/memory-safety-is.html>

Memory safety is one of those elusive concepts like intelligence, consciousness, or porn, that resist attempts to be put to words. Thus, I am not going to attempt to define it. Instead, I want to poke holes in definitions of others.

## 46. Aligning one matrix with another

Source: John D. Cook | Published: Wed, 11 Feb 2026 13:23:50 +0000

Link: <https://www.johndcook.com/blog/2026/02/11/orthogonal-procrustes/>

Suppose you have two  $n \times n$  matrices, A and B, and you would like to find a rotation matrix ? that lines up B with A. That is, you'd like to find ? such that  $A = ?B$ . This is asking too much, except in the trivial case of A and B being  $1 \times 1$  matrices. You could [?] The post Aligning one matrix with another first appeared on John D. Cook .

## 47. Eigenvalue homework problems are backward

Source: John D. Cook | Published: Fri, 06 Feb 2026 23:31:11 +0000

Link: <https://www.johndcook.com/blog/2026/02/06/eigenvalue-roots/>

Classroom When you take a linear algebra course and get to the chapter on eigenvalues, your homework problems will include a small matrix A and you will be asked to find the eigenvalues.

You do this by computing the determinant  $\det(A - \lambda I) = P(\lambda)$  and getting  $P(\lambda)$ , a polynomial in  $\lambda$ . The roots of  $P(\lambda) = 0$  are the eigenvalues. The post Eigenvalue homework problems are backward first appeared on John D. Cook .

#### 48. Fibonacci number certificates

Source: John D. Cook | Published: Thu, 05 Feb 2026 17:14:20 +0000

Link: <https://www.johndcook.com/blog/2026/02/05/fibonacci-certificate/>

Suppose I give you a big number  $F$  and claim that  $F$  is a Fibonacci number. How could you confirm this? Before I go further, let me say what this post is really about. It's not about Fibonacci numbers so much as it is about proofs and certificates. There's no market for large Fibonacci numbers, and certainly [?] The post Fibonacci number certificates first appeared on John D. Cook .

#### 49. I set all 376 Vim options and I'm still a fool

Source: Evan Hahn (dot com) | Published: Fri, 16 Jan 2026 00:00:00 +0000

Link: <https://evanhahn.com/i-set-all-376-vim-options-and-im-still-a-fool/>

I set all of Vim's configuration options. I still feel far from mastery. First impressions of Vim: wow I first saw someone use Vim during an internship in 2012. I had been coding for many years and I fancied myself pretty good at shortcuts, but I was quickly humbled. I watched in awe as experienced users zipped around the code. A single keystroke could move the cursor halfway across the file to exactly the right spot. Code was ripped apart and reshaped like putty. ? Wow ,? I thought to myself, and probably said out loud. 13 years later, still clumsy I vowed to master this editor but I was slow. When I wasn't accidentally opening some unknown menu, I was taking an uneconomical path through the code. I pressed `j` twenty times instead of running `20j` , or manually deleted code inside parenthesis instead of running `di(` . Sometimes I'd open another text editor to give my mind a break from all the key bindings! Fast-forward to 2025. After tons of practice, I felt much more capable. Code did feel more like putty. I was working closer to the speed of thought. I could get code where I wanted much more quickly. 13 years of practice paid off! But Vim still felt clumsy. I was still accidentally...

#### 50. Notes from "On Writing Well"

Source: Evan Hahn (dot com) | Published: Sun, 04 Jan 2026 00:00:00 +0000

Link: <https://evanhahn.com/on-writing-well-book-notes/>

I've been trying to improve my writing so I read On Writing Well by William Zinsser. My main takeaways: Clear thinking is a prerequisite for clear writing. How do you avoid cluttered writing? ?The answer is to clear our heads of clutter. Clear thinking becomes clear writing; one can't exist without the other. It's impossible for a muddy thinker to write good English.? Reduce scope. Zinsser hammers this point repeatedly. For instance: ?Nobody can write a book or an article ?about? something. Tolstoy couldn't write a book about war and peace, or Melville a book about whaling. They made certain reductive decisions about time and place and about individual characters in that time and place?one man pursuing one whale. Every writing project must be reduced before you start to write.? Keep the thesis in mind. ?Writers must [?] constantly ask: what am I trying to say? Surprisingly often they don't know. Then they must look at what they have written and ask: have I said it?? I don't want to write like this guy. I think Zinsser's writing is dogmatic, verbose, outdated, and above all: not to my taste. But that helps me clarify my own style by showing me what I don't want to do. And despite a...

#### 51. Notes from December 2025

Source: Evan Hahn (dot com) | Published: Wed, 31 Dec 2025 00:00:00 +0000

Link: <https://evanhahn.com/notes-from-december-2025/>

Here are my notes from the final month of 2025. Little things I did I predict that Mastodon will outlive Bluesky because the latter is corporate-controlled. We'll see if my prediction is correct in about 25 years. I've been working on a mystery project that uses Pyodide , the WebAssembly-powered Python distribution. After much toil I figured out how to make it do relative imports . I made a little audio speed calculator that lets you enter ?4 hours, 20

minutes? and it?ll tell you how much time you?ll save listening on 1.5xspeed, and 1.6xspeed, and so on. I published my notes on the book Bad Company: Private Equity and the Death of the American Dream . In summary, private equity seems bad. I?m slowly trying to drop GitHub. I moved some of my repositories to Codeberg this month. Finally, I continued writing for Zelda Dungeon . I contributed to their annual ?Best Zelda Ever? ranking . 2025 featured a big upset! I?m especially proud of what I wrote about Phantom Hourglass . Links and bookmarks I?ve got a lot of links and quotes for you this month. ?It may sound melodramatic, and maybe it is, but I see the increased dependence on algorithmically-driven entertainment as a symptom of...

## 52. The horsehead nebula (2026)

Source: Maurycy's Blog | Published: Fri, 09 Jan 2026 00:00:00 +0000

Link: <https://maurycyz.com/astro/horsehead/>

0.55 arcseconds/pixel. Image is 27' wide. North is right (mirrored). Color: 27 minutes (319 \* 4 seconds) through thin clouds Equipment: C9.25, 0.63 reducer, ASI533 MC (IMX533 sensor), EQ6-R mount. Resolution: 5 arcsec FWHM I just can't seem to take a good picture of this region. This one has: Messed up backfocus and collimation Bad seeing, made worse by looking over the roof of my house Every \*single\* 5-second frame was ruined by clouds. I couldn't see Polaris so mount alignment was done based on memory This was the best night I've had all month: There was one moment where I could actually see Orion in the sky... at least it's better than my last attempt where I couldn't even find the object. This is a heavily noise reduced version: It looks a little better, but is missing a lot of the details of this object. Raw stacks: Compressed FITs

## 53. You should start a blog

Source: Maurycy's Blog | Published: Sun, 04 Jan 2026 00:00:00 +0000

Link: [https://maurycyz.com/misc/starting\\_a\\_blog/](https://maurycyz.com/misc/starting_a_blog/)

Writing something down forces you to fully understand it. When the idea is on paper, you can see all the missing assumptions and leaps in logic. It's common to start writing, do some research and find out that your original point was wrong . This is a good thing. You are now less wrong than you were before, and have something you can share so that we can all be less wrong. Even if you don't learn anything new in the process , we don't find our interests and hobbies by magic: We read about them from someone else. Simply writing about what you did yesterday ? even if you are not an expert on the topic ? can be very valuable to the right person. Why a blog: Blogs have a long shelf-life : Social media posts vanish in hours, but a good blog post can stay readable and relevant for decades. Your work can have a lasting impact on lots of people rather than being briefly noticed by a few. Since posts are mutable and long lived, you don't have to aim for perfection. It's perfectly fine start with a single sentence and expand it over time: Blog posts are living documents that grow alongside your ideas. Articles can build on each other , cite sources or provide hundreds of pages worth of det...

## 54. Creating Your Own Opportunities

Source: Terrible Software | Published: Mon, 05 Jan 2026 12:50:41 +0000

Link: <https://terriblesoftware.org/2026/01/05/creating-your-own-opportunities/>

The best career opportunities don't necessarily come from your manager. They come from seeing potential in the boring work already on your plate.

## 55. A Year of Writing for Myself (That 230,000+ People Read)

Source: Terrible Software | Published: Mon, 29 Dec 2025 14:32:32 +0000

Link:

<https://terriblesoftware.org/2025/12/29/a-year-of-writing-for-myself-that-230000-people-read/>

I started this blog about a year ago with pretty modest expectations. It started as a place to organize my thoughts on management, engineering, and everything else in between. I figured maybe a few colleagues would read it. Maybe some friends? I really did not expect that I'd have more than 230,000 unique visitors in [?]

## 56. The Strange Case of Engineers Who Dismiss AI

Source: Terrible Software | Published: Fri, 19 Dec 2025 13:02:12 +0000

Link: <https://terriblesoftware.org/2025/12/19/the-strange-case-of-engineers-who-dismiss-ai/>

Some engineers dismiss AI coding tools as 'garbage' without trying modern versions. Their outdated opinions are costing them.

#### 57. Gadget Review: Epomaker TH87 ISO Mechanical Keyboard ?????

Source: Terence Eden's Blog | Published: Wed, 11 Feb 2026 12:34:46 +0000

Link: <https://shkspr.mobi/blog/2026/02/gadget-review-epomaker-th87-iso-mechanical-keyboard/>

If I'm being brutally honest, I never really got the appeal of mechanical keyboards. There was always someone in the office who made a godawful racket hammering on their keyboard and then waxed lyrical about the merits of various switches. I'd mostly just dismissed them as cranks.

I'm in love with my old Microsoft 4000 ergonomic keyboard. What use could I have a mechanical keyboard festooned with ?

#### 58. Reputation Scores for GitHub Accounts

Source: Terence Eden's Blog | Published: Sat, 07 Feb 2026 12:34:52 +0000

Link: <https://shkspr.mobi/blog/2026/02/reputation-scores-for-github-accounts/>

The folks at GitHub know that Open Source maintainers are drowning in a sea of low-effort contributions. Even before Microsoft forced the unwanted Copilot assistant on millions of repos, it was always a gamble whether a new contributor would be helpful or just some witless jerk. Now it feels a million times worse. There are some discussions about what tools repository owners should have to help?

#### 59. Book Review: Diversifying Open Source - An Open Standards Playbook for Inclusive and Equitable Tech Projects by Paloma Oliveira ?????

Source: Terence Eden's Blog | Published: Fri, 06 Feb 2026 12:34:05 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-diversifying-open-source-an-open-standards-playbook-for-inclusive-and-equitable-tech-projects-by-paloma-oliveira/>

It is refreshing to read a political polemic which contains useful actions the reader can take.

Too many books about the social problems with technology end up being a diagnosis with no cure. Paloma Oliveira's new book (with technical review by my friend Dawn Foster) is a deep dive into how we can all make Open Source more inclusive and equitable. Unlike most tech books, it doesn't follow the ?

#### 60. AI is not another abstraction because god plays dice

Source: Rakhim's blog | Published: Mon, 08 Sep 2025 00:00:00 GMT

Link: [https://rakhim.exotext.com/ai\\_is\\_not\\_another\\_abstraction\\_because\\_god\\_plays\\_dice](https://rakhim.exotext.com/ai_is_not_another_abstraction_because_god_plays_dice)

Some folks have gone all-in on AI-assisted coding. I've seen some tweets (not sure if sarcastic or real, to be honest) expressing disgust about the prospects of ever writing code by hand anymore. The common argument I hear is that this is just another step in the long history of programming breakthroughs. We moved from machine code to assembly, from assembly to C, and from C to high-level languages like Python. Each step was a new layer of abstraction, hiding the messy details of the underlying hardware and software. Some even compare AI to a compiler. The AI prompt, in this view, is simply the next abstraction layer on top of "raw" code. But this analogy is flawed. AI is not another abstraction layer, and the comparison to a compiler is misleading. Abstractions in programming are deterministic . A compiler can be very complex, but it is fundamentally a deterministic machine. Given the same input, it will always produce the same output. It is a tool that we can understand, predict, and, most importantly, trust. AI, in its current form, is not deterministic. It is a probabilistic system. It is not a compiler, it's a contractor. When you send an email to a contractor, you are descri...

#### 61. Finland is a high-context society that loves defaults

Source: Rakhim's blog | Published: Thu, 14 Aug 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/finland-is-a-high-context-society-that-loves-defaults>

In anthropology there is a notion of high-context and low-context cultures. A high-context

culture is a culture or society that communicates dominantly through the use of contextual elements, such as specific forms of body language, the status of an individual, and the tone of voice employed during speech. Rules are not directly or explicitly written or stated. By contrast, a low-context culture enjoys communications that take place most often through written or spoken (verbal) speech and rules are directly and explicitly stated. High-context cultures include close-knit groups of people, while low-context cultures are generally more diverse. I've lived in Canada (Ontario & Quebec) for 5 years, and more than 10 years in Finland. Finland is a more high-context culture compared to other western countries. It's not about body language (there almost isn't any in Nordic cultures). It's about everybody, everywhere, in all situations, knowing exactly how to act and what to do. People rarely look or behave confused. People are rarely late. I almost never see someone running to a bus stop. There's very little discussion when ordering food. Consultants in stores don't get asked difficult que...

## 62. AI is impressive because we've failed at semantic web and personal computing

Source: Rakhim's blog | Published: Fri, 08 Aug 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/ai-is-impressive-because-we-ve-failed-at-semantic-web-and-personal-computing>

Unless someone wrote an article about that exact thing, a plain full-text search engine cannot answer a question like this: What animal is featured on a flag of a country where the first small British colony was established in the same year that Sweden's King Gustav IV Adolf declared war on France? But ChatGPT got the correct answer in a few seconds. Flag of Dominica features the Sisserou parrot, which is only found in Dominica. Great Britain established a small colony on the island in 1805. Google's AI widget failed miserably, by the way. One of the best applications of modern LLM-based AI is surfacing answers from the chaos of the internet. Its success can be partly attributed to our failure to build systems that organize information well in the first place. This product pattern is not new. Take Google Drive: a glorified file system in the cloud with folders and files, but it offers a worse experience than almost any desktop file management application of the last 30 years. Organizing your stuff there is hard and tedious. So Google took a shortcut: full-text search. Just dump everything in, and type to find it later. The pattern of giving up on structure and relying on search ha...

## 63. The Disturbing Rise of MAGA Maoism

Source: Derek Thompson | The Atlantic | Published: 2025-05-08T07:00:00-04:00

Link:

[https://www.theatlantic.com/ideas/archive/2025/05/maga-maoism-trump/682732/?utm\\_source=feed](https://www.theatlantic.com/ideas/archive/2025/05/maga-maoism-trump/682732/?utm_source=feed)

China may well come to dominate the next century?because President Donald Trump is taking a page from the most famous Chinese leader of the previous one. The United States remains the world's preeminent soft power. It's a financial and cultural juggernaut, whose entertainment and celebrities bestride the planet. But as an industrial power, the U.S. is not so much at risk of falling behind as it is objectively behind already. A recent essay in the journal Foreign Affairs by Rush Doshi and Kurt Campbell, both China experts who served in the Biden administration, made the case with alarming specificity. China makes 20 times more cement and 13 times more steel than the U.S. It makes more than two-thirds of the world's electric vehicles, more than three-quarters of its electric batteries, 80 percent of its consumer drones, and 90 percent of its solar panels. China's shipbuilding capacity is several orders of magnitude larger than America's, and its navy will be 50 percent larger than the U.S. Navy by 2030. The Trump administration clearly recognizes the need to rebuild industrial capacity. In its executive order published on ?Liberation Day,? the White House suggested that, without hig...

## 64. Something Alarming Is Happening to the Job Market

Source: Derek Thompson | The Atlantic | Published: 2025-04-30T06:00:00-04:00

Link:

[https://www.theatlantic.com/economy/archive/2025/04/job-market-youth/682641/?utm\\_source=feed](https://www.theatlantic.com/economy/archive/2025/04/job-market-youth/682641/?utm_source=feed)

Something strange, and potentially alarming, is happening to the job market for young, educated workers. According to the New York Federal Reserve, labor conditions for recent college graduates have deteriorated noticeably? in the past few months, and the unemployment rate now stands at an unusually high 5.8 percent. Even newly minted M.B.A.s from elite programs are struggling to find work . Meanwhile, law-school applications are surging ?an ominous echo of when young people used graduate school to bunker down during the great financial crisis. What's going on? I see three plausible explanations, and each might be a little bit true. The first theory is that the labor market for young people never fully recovered from the coronavirus pandemic?or even, arguably, from the Great Recession. ?Young people are having a harder time finding a job than they used to, and it?s been going on for a while, at least 10 years,? David Deming, an economist at Harvard, told me. The Great Recession led not only to mass layoffs but also to hiring freezes at many employers, and caused particular hardships for young people. After unemployment peaked in 2009, the labor market took time to heal, improving...

## 65. A Trade War With China Is a Very Bad Idea

Source: Derek Thompson | The Atlantic | Published: 2025-04-18T06:00:00-04:00

Link: [https://www.theatlantic.com/economy/archive/2025/04/trade-war-china-trump-tariffs/682500/?utm\\_source=feed](https://www.theatlantic.com/economy/archive/2025/04/trade-war-china-trump-tariffs/682500/?utm_source=feed)

Like land wars in Asia , trade wars with China are, generally speaking, unadvisable. But if, for whatever reason, you were insistent on the idea, you'd want to follow two rules. First, find strength in numbers. China is an industrial juggernaut with more than 1 billion citizens. The U.S. is a finance-and-tech giant with fewer than 400 million people. To maximize success, the U.S. would have to assemble an Avengers team of trading partners across North America, Europe, and Asia. This would help keep our supply chains resilient if China cut off access to important products and materials, such as smartphone parts and processed lithium. What's more, this so-called friend-shoring approach would squeeze China and hurt its ability to find alternative export markets, making retaliation less likely. Second, clarity is king.

Reindustrialization?that is, building more factories and plants to make essential machines for AI, computing, energy, and national security?is expensive. To maximize domestic financing and even foreign investment, we'd want investors to understand that the tax and tariff rates we announce one Tuesday will hold up until the Tuesday after that, and the Tuesday after that,...

## 66. Moving to make many of my SSH logins not report things on login

Source: Chris's Wiki :: blog | Published: 2026-02-11T04:33:50Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/MovingToSilentLogins>

I've been logging in to Unix machines for what is now quite a long time. When I started, it was traditional for your login process to be noisy. The login process itself would tell you last login details and the 'message of the day' ('motd'), and people often made their shell .profile or .login report more things, so you could see things like: Last login: Tue Feb 10 22:16:14 2026 from 128.100.X.Y 22:22:42 up 1 day, 11:22, 3 users, load average: 0.40, 2.95, 3.30 cks cks cks [output from fortune elided] : ; (There is no motd shown here but it otherwise hits the typical high points, including a quote from fortune . People didn't always use 'fortune' itself but printing a randomly selected quote on login used to be common.) Many years ago I modified my shell environment on our servers so that it wouldn't report the currently logged in users, show the motd, or tell me my last login. But I kept the ' uptime ' line: \$ ssh cs.utoronto.edu 22:26:05 up 209 days, 5:26, 167 users, load average: 0.47, 0.51, 0.60 : apps0.cs ; Except, I typically didn't see that. I see this only on full login sessions, and when I was in the office I typically used special tools ( also , also , also ) that didn't a...

## 67. A fun Python puzzle with circular imports

Source: Chris's Wiki :: blog | Published: 2026-02-10T04:13:27Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/python/PythonCircularImportPuzzle>

Baptiste Mispelon asked an interesting Python quiz ( via , via @glyph ): Can someone explain this #Python import behavior? I'm in a directory with 3 files: a.py contains 'A = 1; from b import \*' b.py contains 'from a import \*; A += 1' c.py contains 'from a import A; print(A)' Can

you guess and explain what happens when you run 'python c.py'? I encourage you to guess which of the options in the original post is the actual behavior before you read the rest of this entry. There are two things going on here. The first thing is what actually happens when you do 'from module import ...'. The short version is that this copies the current bindings of names from one module to another. So when module b does 'from a import \*', it copies the binding of a.A to b.A and then the += changes that binding. The behavior would be the same if we used 'from a import A' and 'from b import A' in the code, and if we did we could describe what each did in isolation as starting with 'A = 1' (in a), then 'A = a.A; A += 2' (in b), and then 'A = b.A' (back in a) successively (and then in c, 'A = a.A'). The second thing going on is that you can import incomplete modules (this is true in both Py...)

## 68. Systemd and blocking connections to localhost, including via 'any'

Source: Chris's Wiki :: blog | Published: 2026-02-09T04:22:04Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/SystemdAndConnectToAny>

I recently discovered a surprising path to accessing localhost URLs and services , where instead of connecting to 127.0.0.1 or the IPv6 equivalent, you connected to 0.0.0.0 (or the IPv6 equivalent). In that entry I mentioned that I didn't know if systemd's IPAddressDeny would block this. I've now tested this, and the answer is that systemd's restrictions do block this. If you set 'IPAddressDeny=localhost', the service or whatever is blocked from the 0.0.0.0 variation as well (for both outbound and inbound connections). This is exactly the way it should be, so you might wonder why I was uncertain and felt I needed to test it. There are a variety of ways at different levels that you might implement access controls on a process (or a group of processes) in Linux, for IP addresses or anything else. For example, you might create an eBPF program that filtered the system calls and system call arguments allowed and attach it to a process and all of its children using seccomp(2) . Alternately, for filtering IP connections specifically, you might use a cgroup socket address eBPF program ( also ), which are among the the cgroup program types that are available. Or perhaps you'd prefer to use...

## 69. Reading List 01/24/26

Source: Construction Physics | Published: Sat, 24 Jan 2026 13:00:24 GMT

Link: <https://www.construction-physics.com/p/reading-list-012426>

Welcome to the reading list, a weekly list of news and links related to buildings, infrastructure, and industrial technology.

## 70. Do Commodities Get Cheaper Over Time?

Source: Construction Physics | Published: Thu, 22 Jan 2026 13:02:46 GMT

Link: <https://www.construction-physics.com/p/do-commodities-get-cheaper-over-time>

This American Enterprise Institute chart, which breaks down how price changes for different types of goods and services in the consumer price index, has by now become very widely known.

## 71. Reading List 01/17/2026

Source: Construction Physics | Published: Sat, 17 Jan 2026 13:02:36 GMT

Link: <https://www.construction-physics.com/p/reading-list-01172026>

ALARA, OLED screens, bus stop frequency, Ozempic and airlines, and more.

## 72. Writing an LLM from scratch, part 30 -- digging into the LLM-as-a-judge results

Source: Giles' blog | Published: Fri, 09 Jan 2026 01:15:00 +0000

Link: <https://www.gilesthomas.com/2026/01/llm-from-scratch-30-digging-into-llm-as-a-judge>

I'm still working on my "extra credit" projects after finishing the main body of Sebastian Raschka 's book " Build a Large Language Model (from Scratch) ". Last time around, I trained four base models , using the GPT-2 architecture from the book, on Lambda Labs machines. I was using two ways to compare them with each other, with three models that I'd trained locally , and with the original GPT-2 weights from OpenAI: A simple cross entropy loss over a fixed test set. The results for an instruction fine-tune test that's covered in the book. Here were the results I got, sorted by the loss: Test loss IFT score OpenAI weights: medium 3.231 38.53

OpenAI weights: small 3.500 22.98 Cloud FineWeb, 8x A100 40 GiB 3.674 17.09 Cloud FineWeb, 8x H100 80 GiB 3.725 11.98 Cloud FineWeb, 8x A100 80 GiB 3.730 11.71 Cloud FineWeb, 8x B200 160 GiB 3.771 13.89 Local FineWeb train 3.944 16.01 Local FineWeb-Edu extended train 4.135 14.55 Local FineWeb-Edu train 4.167 16.86 Now, you'd expect there to be at least a loose correlation; the lower the loss, the higher the IFT score. But, while we can see a difference between the OpenAI weights and our own, within our own there doesn't seem to be a logical pat...

### 73. Writing an LLM from scratch, part 31 -- the models are now on Hugging Face

Source: Giles' blog | Published: Sat, 17 Jan 2026 19:45:00 +0000

Link: <https://www.gilesthomas.com/2026/01/llm-from-scratch-31-models-on-hugging-face>

As part of my "extra credit" projects after finishing the main body of Sebastian Raschka's book "Build a Large Language Model (from Scratch)", I've trained seven base models completely from scratch based on the book's GPT-2 code -- three locally, and four in the cloud. I plan to train more as I work on ways to improve the quality of the trained models, in the hope that I can get to something closer to the original OpenAI weights' loss on my own hardware, or at least on something I can rent without breaking the bank. It makes sense to share these models somewhere, both so that other people can take a look if they like, and also to build the knowledge of how to do it so that if I produce something more interesting in the future, I'll know how to share that too. Raschka's code is all released under the Apache v2 open source license, so I can share my stuff under the same license without worrying about triggering any legal issues. So: I've put all of the models I've trained so far on Hugging Face under that license, and made them reasonably HF-native (I'll explain what I mean by that later). From the post where I trained the models locally, we have: gpjt/1xrtx3090m24-fineweb --...

### 74. Getting a custom PyTorch LLM onto the Hugging Face Hub (Transformers: AutoModel, pipeline, and Trainer)

Source: Giles' blog | Published: Wed, 28 Jan 2026 23:00:00 +0000

Link: <https://www.gilesthomas.com/2026/01/custom-automodelforcausallm-from-pretrained-models-on-hugging-face>

I spent some time recently getting some models uploaded onto the Hugging Face Hub. I'd trained a bunch of GPT-2 small sized base models from scratch as part of my LLM from scratch series, and wanted to share them with anyone that was interested. I managed to get it done, but it was kind of tricky to get right. The Hugging Face documentation is great if you're using the built-in models, but the coverage of custom architectures is... not quite as comprehensive. There are scattered examples, but they're all a bit vague and there's nothing really bringing them all together. But with what I could find, plus a lot of running things repeatedly, seeing how they failed, tweaking changes, banging my head against obscure stacktraces, and talking to various LLMs, I got there in the end. This post is the tutorial I wish I'd found before I started, and I hope it's useful for people in a similar position. The one warning I'd give is that I did not dig into tokenisers in any depth. My own models use the standard GPT-2 one, and so I could just use the version that is built into Transformers. The setup you need to do with custom tokenisers doesn't look all that different to what you need to do to f...

### 75. Package Management Consulting

Source: Andrew Nesbitt | Published: 2026-02-11T10:00:00+00:00

Link: <https://nesbitt.io/2026/02/11/package-management-consulting.html>

I'm now taking on consulting work. If you're building or running a package manager, registry, or dependency tooling, I can probably help. Over fifteen years I've built Libraries.io, Ecosyste.ms, git-pkgs, the Manifest podcast, co-organised the Package Management devroom at FOSDEM, and contributed to Homebrew. I've written integrations for dozens of package managers, tracked billions of dependency relationships, and watched the same design mistakes repeat across ecosystems for a decade. That pattern recognition is what I bring to consulting engagements. What I can help with: Package manager design & architecture. If you're building a package manager, registry, or dependency resolver, or trying to fix one, I can help you avoid the mistakes that RubyGems, npm, and PyPI learned the hard way. I've documented the design

patterns, tradeoffs, and failure modes across dozens of ecosystems in my cross-ecosystem package manager documentation . Namespace design, version constraint semantics, lockfile formats, registry APIs, lifecycle hooks, trust models. Software supply chain security. I've found and catalogued dependency confusion attacks, typosquatting campaigns, slopsquatting (a term...

## 76. Dependency Resolution Methods

Source: Andrew Nesbitt | Published: 2026-02-06T12:00:00+00:00

Link: <https://nesbitt.io/2026/02/06/dependency-resolution-methods.html>

Every package manager faces the same core problem: given a set of packages with version constraints, find a compatible set of versions to install. The classic example is the diamond dependency: A depends on B and C, both of which depend on D but at incompatible versions. The resolver has to find a version of D that satisfies both, prove that none exists, or find some other way out. Di Cosmo et al. proved in 2005 that this problem is NP-complete, encoding it as 3SAT for Debian and RPM constraint languages. In practice, real-world dependency graphs are far more tractable than the worst case, and different ecosystems have landed on different resolution strategies that make different tradeoffs between completeness, speed, error quality, and simplicity. These approaches fall roughly into complete solvers (SAT, PubGrub, ASP), heuristic solvers (backtracking, Molinillo, system resolvers), constraint relaxation strategies (deduplication with nesting, version mediation), and approaches that avoid the problem entirely (minimal version selection, content-addressed stores). The categorizing package manager clients post lists which package managers use which approach. The package management pa...

## 77. crates.io's Freaky Friday

Source: Andrew Nesbitt | Published: 2026-02-06T00:00:00+00:00

Link: <https://nesbitt.io/2026/02/06/cratesio-freaky-friday.html>

The maintainers of crates.io wake up Friday morning to find their registry has swapped design philosophies with Debian. They still serve the Rust ecosystem, Debian still serves Linux distributions, but the tradeoffs they've chosen have reversed. Like Tess and Anna in Freaky Friday , they're stuck in each other's bodies, forced to navigate constraints they've spent years criticizing from the outside. The crates.io team reaches for their coffee and tries to ship a hotfix, only to discover they can't publish without a signed GPG key from a designated sponsor and a three-day waiting period for linting and policy review. Meanwhile, Debian maintainers watch in horror as packages flow into their repository without coordination, breaking stable in ways that won't surface until someone's server fails to boot. Waking up with snapshots Freaky Friday crates.io splits into multiple coexisting suites. There's a ?stable? suite with older, well-tested crate versions, a ?testing? suite with recent versions undergoing integration testing, and an ?unstable? suite with the latest uploads. Within each suite, only one version of each crate exists, but the suites themselves persist simultaneously, and w...

## 78. Everyone knows your location, Part 2: try it yourself and share the results

Source: tim.sh | Published: Thu, 17 Apr 2025 10:05:43 GMT

Link: <https://timsh.org/everyone-knows-your-location-part-2-try-it-yourself/>

Learn how to record and analyse your mobile device traffic, take an app from the list of "shady" apps and share the results.

## 79. Github scam investigation: Thousands of "mods" and "cracks" stealing your data

Source: tim.sh | Published: Thu, 27 Feb 2025 21:37:33 GMT

Link:

<https://timsh.org/github-scam-investigation-thousands-of-mods-and-cracks-stealing-your-data/>

While looking through the articles on a "social engineering" themed forum I discovered a relatively new scam scheme that shocked me. People create thousands of GitHub repositories with all sorts of things - from Roblox and Fortnite mods to "cracked" FL Studio and Photoshop. As soon

80. Everyone knows your location: tracking myself down through in-app ads

Source: tim.sh | Published: Fri, 31 Jan 2025 23:30:18 GMT

Link: <https://timsh.org/tracking-myself-down-through-in-app-ads/>

How I tracked myself down using leaked location data in the in-app ads, and what I found along the way.

81. I spoke about ICEBlock and Trump's app censorship on the kill switch podcast

Source: micahflee | Published: Tue, 28 Oct 2025 17:46:25 GMT

Link: <https://micahflee.com/i-spoke-about-iceblock-and-trumps-app-censorship-on-the-kill-switch-podcast/>

In case you missed it, last week I was on an episode of the kill switch podcast. I discussed the ICEBlock app and why it was activism theater , and how the developer handled my vulnerability report in the worst possible way . But I also talk about why, even though I

82. Practical Defenses Against Technofascism

Source: micahflee | Published: Sun, 26 Oct 2025 03:57:00 GMT

Link: <https://micahflee.com/practical-defenses-against-technofascism/>

I gave the Saturday morning keynote at BSidesPDX! I spoke honestly and frankly about the terrifying reality that Americans are facing under Trump's fascist regime, alongside practical advice for communities to defend themselves. Watch my talk below. Or if you prefer reading articles over watching video, I'

83. In war-torn Portland? Watch me speak at BSidesPDX!

Source: micahflee | Published: Sun, 19 Oct 2025 18:39:10 GMT

Link: <https://micahflee.com/in-war-torn-portland-watch-me-speak-at-bsidespdx/>

I'll be giving the Saturday morning keynote at BSidesPDX next weekend in beautiful Portland, Oregon. I'm still putting the final touches on my talk, but it will largely be about the new American technofascist reality of mercenary spyware, device searches, and app censorship, with a focus

84. Pluralistic: Europe takes a big step towards a post-dollar world (11 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 11 Feb 2026 10:09:05 +0000

Link: <https://pluralistic.net/2026/02/11/post-dollar-world/>

Today's links Europe takes a big step towards a post-dollar world: Recapturing \$24t worth of transactions from Visa/Mastercard. Hey look at this: Delights to delectate. Object permanence: API for Congress; Steampunk fetish mask; Hillary x AOL login screen; Suffragist Valentines; Musk x Intuit vs the American people. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em.

Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Europe takes a big step towards a post-dollar world (permalink) There's a reason every decentralized system eventually finds its way onto a platform: platforms solve real-world problems that platform users struggle to solve for themselves. I've written before about the indie/outsider author Crad Kilodney, who wrote, edited, typeset and published chapbooks of his weird and wonderful fiction, and then sold his books from Toronto street-corners with a sign around his neck

reading VERY FAMOUS CANADIAN AUTHOR BUY MY BOOKS (or, if he was feeling spicy, simply: MARGARET ATWOOD): <https://pluralistic.net/2024/02/19/crad-kilodney-was-an-outlier/#intermediation> Crad was a...

85. Pluralistic: Luxury Kafka (06 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 06 Feb 2026 08:43:22 +0000

Link: <https://pluralistic.net/2026/02/06/doge-ball/>

Today's links Luxury Kafka: US Immigration on the easiest setting. Hey look at this: Delights to delectate. Object permanence: Whisky PC; Anitfeatures; Silicon Roundabout; Steampunk

Etch-A-Sketch; MLMs as mirror-world organizers. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Luxury Kafka (permalink) Having been through the US immigration process (I got my first work visa more than 25 years ago and became a citizen in 2022), it's obvious to me that Americans have no idea how weird and tortuous their immigration system is:

<https://www.flickr.com/photos/doctorow/52177745821/> As of a couple years ago, Americans' ignorance of their own immigration system was merely frustrating, as I encountered both squishy liberals and xenophobic conservatives talking about undocumented immigrants and insisting that they should "just follow the rules." But today, as murderous ICE squads patrol our streets kidnapping people and sending them to concentration camps where they are beaten to death or deported to offshore slave labor prisons, th...

#### 86. Pluralistic: All laws are local (05 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Thu, 05 Feb 2026 12:57:09 +0000

Link: <https://pluralistic.net/2026/02/05/contingency/>

Today's links All laws are local: And no law knows how evitable it is. Hey look at this: Delights to delectate. Object permanence: Whisky PC; Anitfeatures; Silicon Roundabout; Steampunk Etch-A-Sketch; MLMs as mirror-world organizers. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. All laws are local (permalink) About halfway through Thomas Piketty's 2013 barnstorming Capital in the 21st Century, Piketty tosses off a little insight that skewered me on the spot and never let me go: the notion that any societal condition that endures beyond a generation becomes "eternal" in the popular consciousness:

<https://memex.craphound.com/2014/06/24/thomas-pikettys-capital-in-the-21st-century/> Piketty was referring to "primogeniture," the ancient practice of automatically passing the family fortune onto the eldest son (or, if no son was available, the eldest nephew). Primogeniture did important work by keeping dynastic fortunes intact, rather than dividing them up among all children of some baron or lord or other guillotineable monste...

#### 87. The Coherence Premium

Source: Westenberg. | Published: Mon, 02 Feb 2026 22:41:04 GMT

Link: <https://www.joanwestenberg.com/the-coherence-premium/>

I don't necessarily believe in second brains. The notion (pun-intended) that you can offload your thinking to a perfectly organized system of notes and links has always struck me as a fantasy. The people I know who've built elaborate Notion databases or Obsidian vaults mostly end

#### 88. Your Life is the Sum Total of 2,000 Mondays

Source: Westenberg. | Published: Sat, 31 Jan 2026 00:25:46 GMT

Link: <https://www.joanwestenberg.com/your-life-is-the-sum-total-of-2-000-mondays/>

We plan our lives like we're editing a movie trailer. The trip to Portugal, or the product launch, or the transformation photo at the gym. The big moment where everything crystallizes into meaning. We accumulate these peaks in our imagination, and then arrange them into a montage that

#### 89. Why Intelligence Is a Terrible Proxy for Wisdom

Source: Westenberg. | Published: Tue, 27 Jan 2026 00:34:31 GMT

Link: <https://www.joanwestenberg.com/why-intelligence-is-a-terrible-proxy-for-wisdom/>

Isaac Newton, one of the greatest scientific minds in human history, lost a fortune in the South Sea Bubble of 1720. After initially making money and selling his shares, he bought back in at the peak, watching helplessly as the stock collapsed. His reported loss was around £20,000,

## 90. 1Password Dependency Breaks Syntax Highlighting

Source: Fernando Borretti | Published: Sat, 27 Dec 2025 01:00:00 +0000

Link: <https://borretti.me/article/1password-dependency-breaks-syntax-highlighting>

Why does a password manager need a syntax highlighter?

## 91. Using the Brother DS-640 Scanner on NixOS

Source: Fernando Borretti | Published: Sat, 27 Dec 2025 00:00:00 +0000

Link: <https://borretti.me/article/using-the-brother-ds-640-scanner-on-nixos>

I suffer so you don't have to.

## 92. Books I Enjoyed in 2025

Source: Fernando Borretti | Published: Mon, 22 Dec 2025 00:00:00 +0000

Link: <https://borretti.me/article/books-i-enjoyed-in-2025>

A short list.

## 93. Proving What's Possible

Source: Computer Things | Published: Wed, 11 Feb 2026 18:36:53 +0000

Link: <https://buttdown.com/hillelwayne/archive/proving-whats-possible/>

As a formal methods consultant I have to mathematically express properties of systems. I generally do this with two "temporal operators": A(x) means that x is always true. For example, a database table always satisfies all record-level constraints, and a state machine always makes valid transitions between states. If x is a statement about an individual state (as in the database but not state machine example), we further call it an invariant . E(x) means that x is "eventually" true, conventionally meaning "guaranteed true at some point in the future". A database transaction eventually completes or rolls back, a state machine eventually reaches the "done" state, etc. These come from linear temporal logic, which is the mainstream notation for expressing system properties. 1 We like these operators because they elegantly cover safety and liveness properties , and because we can combine them . A(E(x)) means x is true an infinite number of times, while A(x => E(y) means that x being true guarantees y true in the future. There's a third class of properties, that I will call possibility properties: P(x) is "can x happen in this model"? Is it possible for a table to have more than ten rec...

## 94. The Liskov Substitution Principle does more than you think

Source: Computer Things | Published: Tue, 06 Jan 2026 16:51:26 +0000

Link:

<https://buttdown.com/hillelwayne/archive/the-liskov-substitution-principle-does-more-than/>

Happy New Year! I'm done with the newsletter hiatus and am going to try updating weekly again. To ease into things a bit, I'll try to keep posts a little more off the cuff and casual for a while, at least until Logic for Programmers is done. Speaking of which, v0.13 should be out by the end of this month. So for this newsletter I want to talk about the Liskov Substitution Principle (LSP). Last week I read A SOLID Load of Bull by cryptographer Loupe Vaillant, where he argues the SOLID principles of OOP are not worth following. He makes an exception for LSP, but also claims that it's "just subtyping" and further: If I were trying really hard to be negative about the Liskov substitution principle, I would stress that it only applies when inheritance is involved , and inheritance is strongly discouraged anyway. LSP is more interesting than that! In the original paper, A Behavioral Notion of Subtyping , Barbara Liskov and Jeannette Wing start by defining a "correct" subtyping as follows: Subtype Requirement: Let ?(x) be a property provable about objects x of type T. Then ?(y) should be true for objects y of type S where S is a subtype of T. From then on, the paper determine what guaran...

## 95. Some Fun Software Facts

Source: Computer Things | Published: Wed, 10 Dec 2025 18:45:37 +0000

Link: <https://buttdown.com/hillelwayne/archive/some-fun-software-facts/>

Last newsletter of the year! First some news on Logic for Programmers . Thanks to everyone who

donated to the feedchicago charity drive ! In total we raised \$2250 for Chicago food banks. Proof here . If you missed buying Logic for Programmers real cheap in the charity drive, you can still get it for \$10 off with the holiday code hannukah-presents . This will last from now until the end of the year. After that, I'll be raising the price from \$25 to \$30. Anyway, to make this more than just some record keeping, let's close out with something light. I'm one of those people who loves hearing "fun facts" about stuff. So here's some random fun facts I accumulated about software over the years: In 2017, a team of eight+ programmers successfully implemented Tetris as a game of life simulation . The GoL grid had an area of 30 trillion pixels and implemented a full programmable CPU as part of the project. Computer systems have to deal with leap seconds in order to keep UTC (where one day is 86,400 seconds) in sync with UT1 (where one day is exactly one full earth rotation). The people in charge recently passed a resolution to abolish the leap second by 2035, letting UTC and UT1 slowly drift...

## 96. how do I stop participating?

Source: the singularity is nearer | Published: 2026-01-18T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/18/how-do-i-stop.html>

This one is for the complainers and whiners . First off, if you think I ever worked for big tech, you don?t know much about me. I had 3 internships at Google, two when I was very young (18/19) and got a great education that paid me, and one where I just wrote open source software . I worked at Facebook in 2011 for 9 months and quit before any shares vested cause I thought the mission of ?wasting the world?s time? was dumb. And oh yes the 5 weeks I worked at Twitter hoping maybe it could be different. I don?t have too much money; the money I do have I made from pwn2own , CTFs, crypto contracting , and basic market goes up investing. The main reason I have money is because I don?t spend money; for example, I have travelled all around the world, and the best travel experiences I have had were way below cost of living in America. I have never made money from big tech, so I?m definitely not telling you to do something I didn?t do. I?m actually telling you to do exactly what I did. I started two companies, comma.ai and tiny corp. While I took a small amount of VC from a16z for comma, I never gave up any control, and have 0 intention of doing hypergrowth ponzi scams (lol sorry but ur the...

## 97. you have three minutes to escape the perpetual underclass

Source: the singularity is nearer | Published: 2026-01-17T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/17/three-minutes.html>

I had a dream last night I went to work at Amazon. Joining the Bezos neofeudal empire. This post is directed at anyone with talent who works at a tech company ushering in this future. Have you thought about how this is going to play out? I understand you may be at a place where you are insecure about money, and that insecurity is what drives you. But why do you think having more money will fix that insecurity? In the future, when labor is fully marginalized and capital is the only force, you will not be able to afford GPT\$\$\$ (it?s \$1B per month), only the billionaires will. GPT\$\$\$ is surely smart enough to separate you from whatever you have, be that with targeting advertising, a scam you fall for, or lobbying your government to take it from you. A pile of money will buy you nothing in the neofeudal world. Historically, there has been some loyalty to the subjects of a feudal empire because labor had some value. You needed the peasants to grow the grain so you could tax it and take it and have leverage over others by having grain to offer them. When the grain is produced by machines, the peasants are cut out of the loop. The solution to this is not to accumulate grain, buy shares i...

## 98. Anthropic is making a huge mistake

Source: the singularity is nearer | Published: 2026-01-15T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/01/15/anthropic-huge-mistake.html>

For context, Anthropic again blocked opencode from the Claude Code API. I wasn?t planning to try other models today, but now that?s all I?m going to be doing. I?m sure they have a massive influx of users on Claude Code since Opus 4.5, I?ve been recommending it to everyone. But this level of user hostility raises the question, is this really a company you want to rely on in your workflow? First they came for opencode And I did not speak out Because I did not use

opencode The first volley of this was an easy workaround, so whatever. But you can make this really hard, enforce a fixed system prompt, etc? You will not convert people back to Claude Code, you will convert people to other model providers. These actions make it clear that opencode is only the start. If you aren't using Claude Code how they want you to, they will cut you off.

#### 99. Switching it up a bit

Source: Matt Godbolt's blog | Published: 2025-12-23T12:00:00Z

Link: [http://xania.org/202512/23-switching-it-up?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/23-switching-it-up?utm_source=feed&utm_medium=rss)

Taking a look at the various ways the compiler can optimise switch statements

#### 100. Clever memory tricks

Source: Matt Godbolt's blog | Published: 2025-12-22T12:00:00Z

Link: [http://xania.org/202512/22-memory-cunningness?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/22-memory-cunningness?utm_source=feed&utm_medium=rss)

We learn that compilers have tricks to access memory efficiently

#### 101. When SIMD Fails: Floating Point Associativity

Source: Matt Godbolt's blog | Published: 2025-12-21T12:00:00Z

Link: [http://xania.org/202512/21-vectorising-floats?utm\\_source=feed&utm\\_medium=rss](http://xania.org/202512/21-vectorising-floats?utm_source=feed&utm_medium=rss)

Why floating point maths doesn't vectorise like integers, and what to do about it

#### 102. Can modern LLMs actually count the number of b's in "blueberry"?

Source: Max Woolf's Blog | Published: Tue, 12 Aug 2025 09:00:00 -0700

Link: <https://minimaxir.com/2025/08/lm-blueberry/>

It's an adversarial question for LLMs, but it's not unfair.

#### 103. LLMs can now identify public figures in images

Source: Max Woolf's Blog | Published: Mon, 28 Jul 2025 13:15:00 -0700

Link: <https://minimaxir.com/2025/07/lms-identify-people/>

ChatGPT and Claude won't, but Gemini will.

#### 104. Predicting Average IMDb Movie Ratings Using Text Embeddings of Movie Metadata

Source: Max Woolf's Blog | Published: Mon, 30 Jun 2025 10:00:00 -0700

Link: <https://minimaxir.com/2025/06/movie-embeddings/>

Don't try this in your data science interviews.

#### 105. Premium: This Is Worse Than The Dot Com Bubble

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 16 Jan 2026 17:15:37 GMT

Link: <https://www.wheresyoured.at/dot-com-bubble/>

Soundtrack - Radiohead - Karma Police I just spent a week at the Consumer Electronics Show, and one word kept coming up: bullshit. LG, a company known for making home appliances and televisions, demonstrated a robot (named ?CLOiD? for some reason) that could ?fold laundry?

#### 106. 2025, A Retrospective

Source: Ed Zitron's Where's Your Ed At | Published: Wed, 31 Dec 2025 18:00:53 GMT

Link: <https://www.wheresyoured.at/2025-a-retrospective/>

Hey all, I'm not dropping this on the actual newsletter feed because it's a little self-indulgent and I'm not sure 88,000 or so people want an email about it. If you want to support my work directly, please subscribe to my premium newsletter.

#### 107. The Enshittifinancial Crisis

Source: Ed Zitron's Where's Your Ed At | Published: Mon, 29 Dec 2025 16:57:06 GMT

Link: <https://www.wheresyoured.at/the-enshittifinancial-crisis/>

Soundtrack: Lynyrd Skynyrd ? Free Bird This piece is over 19,000 words, and took me a great deal of writing and research. If you liked it, please subscribe to my premium newsletter. It's

\$70 a year, or \$7 a month, and in return you get a weekly

## 108. Scaling HNSWs

Source: antirez.com | Published: Tue, 11 Nov 2025 13:53:38 +0100

Link: <http://antirez.com/news/156>

I'm taking a few weeks of pause on my HNSWs developments (now working on some other data structure, news soon). At this point, the new type I added to Redis is stable and complete enough, it's the perfect moment to reason about what I learned about HNSWs, and turn it into a blog post. That kind of brain dump that was so common pre-AI era, and now has become, maybe, a bit more rare. Well, after almost one year of thinking and implementing HNSWs and vector similarity stuff, it is time for some writing. However this is not going to be an intro on HNSWs: too many are present already. This is the ?extra mile? instead. If you know HNSWs, I want to share with you my more ?advanced? findings, especially in the context of making them fast enough to allow for a ?Redis? experience: you know, Redis is designed for low latency and high performance, and HNSWs are kinda resistant to that, so there were challenges to expose HNSWs as an abstract data structure. This blog post will be split into several sections. Think of them as pages of the same book, different chapters of the same experience. Oh and, by the way, I already wrote and subsequently lost this blog post :D [long, sad story about MacOS...]

## 109. AI is different

Source: antirez.com | Published: Wed, 13 Aug 2025 17:59:56 +0200

Link: <http://antirez.com/news/155>

Regardless of their flaws, AI systems continue to impress with their ability to replicate certain human skills. Even if imperfect, such systems were a few years ago science fiction. It was not even clear that we were so near to create machines that could understand the human language, write programs, and find bugs in a complex code base: bugs that escaped the code review of a competent programmer. Since LLMs and in general deep models are poorly understood, and even the most prominent experts in the field failed miserably again and again to modulate the expectations (with incredible errors on both sides: of reducing or magnifying what was near to come), it is hard to tell what will come next. But even before the Transformer architecture, we were seeing incredible progress for many years, and so far there is no clear sign that the future will not hold more. After all, a plateau of the current systems is possible and very credible, but it would likely stimulate, at this point, massive research efforts in the next step of architectures. However, if AI avoids plateauing long enough to become significantly more useful and independent of humans, this revolution is going to be very unlik...

## 110. Coding with LLMs in the summer of 2025 (an update)

Source: antirez.com | Published: Sun, 20 Jul 2025 12:58:54 +0200

Link: <http://antirez.com/news/154>

Frontier LLMs such as Gemini 2.5 PRO, with their vast understanding of many topics and their ability to grasp thousands of lines of code in a few seconds, are able to extend and amplify the programmer capabilities. If you are able to describe problems in a clear way and, if you are able to accept the back and forth needed in order to work with LLMs, you can reach incredible results such as: 1. Eliminating bugs you introduced in your code before it ever hits any user: I experienced this with Vector Sets implementation of Redis. I would end eliminating all the bugs eventually, but many were just removed immediately by Gemini / Claude code reviews. 2. Explore faster how a given idea could work, by letting the LLM write the throw away code to test ASAP in order to see if a given solution is actually more performant, if it is good enough, and so forth. 3. Engage in pair-design activities where your instinct, experience, design taste can be mixed with the PhD-level knowledge encoded inside the LLM. In this activity, the LLM will sometimes propose stupid paths, other times incredibly bright ideas: you, the human, are there in order to escape local minimal and mistakes, and exploit the fa...

## 111. Golang's encoding/XML is broken and no one seems to care

Source: Jayden's Blog | Published: 2025-06-15T15:19:00+00:00

Link: <https://jayd.ml/2025/06/15/encoding-xml-is-broken-and-no-one-cares.html>

Programming Note: You may have subscribed to this blog because of a snarky post about YouTube that for some reason was popular. Instead of more insightful(?) content like that, you're getting a complaint to no audience in particular about XML in Golang. You're welcome. For the first time I've felt betrayed somewhat by Go's standard library. encoding/xml is just broken in a lot of ways, has been for over a decade, and no one seems to want to fix it. I couldn't find a good independent, working xml library to use either, so this just doesn't work. In short: I do not recommend you try and parse or generate XML in golang, at least until Google fixes encoding/xml or someone writes an XML parser that works. With enough hacks and workarounds you just might be able to get something to work. It won't be pretty though. I mean seriously, you can't really parse or generate RSS or podcast feeds with encoding/xml . I feel like that should be one of the better supported cases. Takeaways: A lot of the time the advice will boil down to having to have separate structs with different tags for marshalling and unmarshalling XML. As much as I hate it, that really does seem to be the only solution for so...

## 112. Adventures in Podcasting

Source: Jayden's Blog | Published: 2025-05-31T10:29:00+00:00

Link: <https://jayd.ml/2025/05/31/adventures-in-podcasting.html>

Contenders: <https://github.com/eduncan911/podcast> <https://github.com/jbub/podcasts> Fun with encoding/xml For the first time I've felt betrayed somewhat by Go's standard library. The encoding/xml package's namespace support seems to be half baked, and I've been unable to work around some issues. Quick Recap Here's an example of using encoding/xml package main import ( "encoding/xml" "fmt" ) type XMLDoc struct { XMLName xml . Name 'xml:"document"' Field string 'xml:"field"' } func main () { input := ' field\_value ' var parsed XMLDoc xml . Unmarshal ([] byte ( input ), & parsed ) fmt . Printf ( "%#v" , parsed ) } main.XMLDoc{XMLName:xml.Name{Space:"", Local:"document"}, Field:"field\_value"} Simple enough. Here's how you would modify this to use namespaces , demonstrating two equivalent ways to specify the namespace. package main import ( "encoding/xml" "fmt" ) type XMLDoc struct { XMLName xml . Name 'xml:"document"' Field string 'xml:"https://example.com/xmlschema field"' } func parse ( input string ) { var parsed XMLDoc xml . Unmarshal ([] byte ( input ), & parsed ) fmt . Printf ( "%#v \n" , parsed ) } func main () { parse ( ' field\_value ' ) parse ( ' field\_value ' ) }

## 113. Someone At YouTube Needs Glasses

Source: Jayden's Blog | Published: 2025-04-30T08:26:00+00:00

Link: <https://jayd.ml/2025/04/30/someone-at-youtube-needs-glasses.html>

Opened YouTube and was greeted with this abomination: This is on a 32? 1440p display. There are five (5) videos visible, and 1/6 of the page would have been an enormous ad. For reference, here is YouTube as of January 2019: There are 30 videos visible and zero ads. I really, really hope that this A/B test fails. Unfortunately, using an advanced analytics package I've projected that around May 2026 the YouTube homepage will just be one video, and by September there will be no videos at all on the homepage. Presumably by then we'll have our mandatory NeuraLinks and the YouTube algorithm will be able to inject real-time ML generated content (and ads) straight into our brains, tuning its output as needed to maximize our dopamine response. I miss YouTube before they turned the pain dial all the way towards money.

## 114. Unresponsive Buttons on My Fastest Hardware Ever

Source: Jim Nielsen's Blog | Published: Wed, 11 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/unresponsive-buttons/>

This is one of those small things that drives me nuts. Why? I don't know. I think it has something to do with the fact that I have a computer that is faster than any computer I've ever used in my entire life ? and yet, clicking on buttons results in slight but perceptible delays. Let me explain. Imagine a button that looks like this: Button onClick={ async () => { const data = await getSessionUrlFromStripe (id); window . location = data. url ; } } > Upgrade to Pro Button > For SPA apps, when the user clicks that button it takes a split second (even on a fast

connection) for anything to happen because: The browser makes a request to the server The server talks to Stripe to get a session The server responds with the session data to the client The client redirects When clicking on that button, even on a fast connection, my brain glitches for a second, my thought process going something like: I click [nothing happens] I think ?Did that work?? Just as I?m about to click again, I see the URL bar change I think, ?Oh, ok, it?s doing something .? I stop myself from clicking again while I wait for the UI to redraw Granted those thoughts occur in my brain in under a second, but I hate that pau...

#### 115. The Browser?s Little White Lies

Source: Jim Nielsen?s Blog | Published: Sun, 01 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/browsers-white-lies/>

So I?m making a thing and I want it to be styled different if the link?s been visited. Rather than build something myself in JavaScript, I figure I?ll just hook into the browser?s mechanism for tracking if a link?s been visited (a sensible approach, if I do say so myself ). Why write JavaScript when a little CSS will do? So I craft this: `.entry :has( a :visited ) { opacity : .5 ; filter : grayscale ( 1 ); }` But it doesn?t work. `:has()` is relatively new, and I?ve been known to muff it, so it?s probably just a syntax issue. I start researching. Wouldn?t you know it? We can?t have nice things. `:visited` doesn?t always work like you?d expect because we (not me, mind you) exploited it. Here?s MDN : You can style visited links, but there are limits to which styles you can use. While `:has()` is not mentioned specifically, other tricks like sibling selectors are: When using a sibling selector, such as `:visited + span` , the adjacent element ( `span` in this example) is styled as though the link were unvisited. Why? You guessed it. Security and privacy reasons. If it were not so, somebody could come along with a little JavaScript and uncover a user?s browsing history (imagine, for example, se...

#### 116. The Don?t ?Contact Us? Page

Source: Jim Nielsen?s Blog | Published: Wed, 28 Jan 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/dont-contact-us-page/>

Nic Chan comes out as the whistleblower on how many ?Contact Us? pages are made (spoiler: they?re designed to keep us from contacting anyone). A ?fuck off contact page? is what a company throws together when they actually don?t want anyone to contact them at all. They [?] are trying to reduce the amount of money they spend on support by carefully hiding the real support channels [?] If you solve your own problem by reading the knowledge base, then this is a win for the company. They don?t want to hear from you, they want you to fuck off. It?s true. This is how the proverbial sausage is made. I?ve been there. I?ve seen these decisions handed down. Which means, like Chan, I know how to read between the lines of most ?Contact Us? pages on the internet. I?m not sure about you, but as a user, when I see [these kinds of pages], knowing that whatever my original query was, [I know] I?m going to have to solve it unassisted. My process follows this arc: I have a question. Go to the company?s ?Contact Us? page. Immediately intuit from the design of the page whether I?m actually going to be able to contact someone and get help, or if I?m on my own. A direct line to a human is the ultimate lu...

#### 117. The Need to Read

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/read.html>

(No summary)

#### 118. What You (Want to)\* Want

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/want.html>

(No summary)

#### 119. Alien Truth

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/alien.html>

(No summary)

## 120. Is chat a good UI for AI? A Socratic dialogue

Source: Geoffrey Litt | Published: 2025-06-29T14:17:00+00:00

Link: <https://geoffreylitt.com/2025/06/29/chat-ai-dialogue.html>

The pupil was confused. Some people on Design Twitter said that chat isn't a good UI for AI? but then chat seemed to be winning in many products? He climbed Mount GPT to consult a wizard? ?: please wizard tell me once and for all. is chat a good UI...

## 121. Stevens: a hackable AI assistant using a single SQLite table and a handful of cron jobs

Source: Geoffrey Litt | Published: 2025-04-12T14:40:00+00:00

Link: <https://geoffreylitt.com/2025/04/12/how-i-made-a-useful-ai-assistant-with-one-sqlite-table-and-a-handful-of-cron-jobs.html>

There's a lot of hype these days around patterns for building with AI. Agents, memory, RAG, assistants?so many buzzwords! But the reality is, you don't need fancy techniques or libraries to build useful personal tools with LLMs. In this short post...

## 122. Avoid the nightmare bicycle

Source: Geoffrey Litt | Published: 2025-03-03T22:13:00+00:00

Link: <https://geoffreylitt.com/2025/03/03/the-nightmare-bicycle.html>

In my opinion, one of the most important ideas in product design is to avoid the ?nightmare bicycle?. Imagine a bicycle where the product manager said: ?people don't get math so we can't have numbered gears. We need labeled buttons for gravel mode...

## 123. what is a build system, anyway?

Source: the website of jyn | Published: 2025-12-12T00:00:00+00:00

Link: <https://jyn.dev/what-is-a-build-system-anyway/>

Some definitions and an overview of the world of build systems

## 124. I want a better build executor

Source: the website of jyn | Published: 2025-12-05T00:00:00+00:00

Link: <https://jyn.dev/i-want-a-better-build-executor/>

I want a way to gradually transition existing builds to be hermetic.

## 125. I want a better action graph serialization

Source: the website of jyn | Published: 2025-12-04T00:00:00+00:00

Link: <https://jyn.dev/i-want-a-better-action-graph-serialization/>

Makefiles and Ninja files have limitations, separate from the implementing tool. How can we fix them?

## 126. Nested Code Fences in Markdown

Source: Susam Pal | Published: Mon, 19 Jan 2026 00:00:00 +0000

Link: <https://susam.net/nested-code-fences.html>

Today, we will meet a spiky-haired nerd named Corey Dumm, who normally lives within Markdown code fences. We will get to know him a bit, smile with him when his fences hold and weep quietly when misfortune strikes. One of the caveats of the Markdown universe is the wide variety of Markdown implementations available. In these parallel universes, the rules of Markdown rendering differ subtly. In this post, we will focus only on the CommonMark specification. Since GitHub Flavoured Markdown (GFM) is a strict superset of CommonMark, whatever we discuss here applies equally well to both CommonMark and GFM. Contents Basic Code Fences Fancy Code Fences Basic Code Spans Fancy Code Spans Specification Basic Code Fences Corey had a knack for working with computers ever since he was a kid. Corey at his computer: “ (o\_o)--.|[\_]| “ Everything was perfect in Corey's world. The CommonMark renderer would convert the Markdown above to the following HTML: Corey at his computer: (o\_o)--.|[\_]| View

HTML Corey at his computer: (o\_o)--.|[\_]| At this point, all was well. Corey grew quickly. Before long, he had a head full of spiky hair. Then the fences began to matter. Corey, all grown up: “ “ (o...

## 127. Minimal GitHub Workflow

Source: Susam Pal | Published: Thu, 15 Jan 2026 00:00:00 +0000

Link: <https://susam.net/minimal-github-workflow.html>

This is a note where I capture the various errors we receive when we create GitHub workflows that are smaller than the smallest possible workflow. I do not know why anyone would ever need this information and I doubt it will serve any purpose for me either but sometimes you just want to know things, no matter how useless they might be. This is one of the useless things I wanted to know today.

Contents Empty Workflow On On Push Jobs Job ID Steps Runs On Runs On Ubuntu Latest Empty Steps Run Run Echo Hello, World Empty Workflow For the first experiment we just create a zero byte file and push it to GitHub as follows, say, like this:

```
mkdir -p .github/workflows/ touch .github/workflows/hello.yml git add .github/ git commit -m 'Empty workflow' git push -u origin main
```

Under the GitHub repo's Actions tab, we find this error:

Error No event triggers defined in 'on' On Then we update the workflow as follows: on: Now we get this error: Invalid workflow file: .github/workflows/hello.yml#L1 (Line: 1, Col: 4): Unexpected value ”, (Line: 1, Col: 1): Required property is missing: jobs On Push Next update: on: push Corresponding error: Invalid workflow file: .github/workflows/hello.yml#L1 (Line:...)

## 128. Three Inverse Laws of Robotics

Source: Susam Pal | Published: Mon, 12 Jan 2026 00:00:00 +0000

Link: <https://susam.net/inverse-laws-of-robotics.html>

Introduction Since the launch of ChatGPT in November 2022, generative artificial intelligence (AI) chatbot services have become increasingly sophisticated and popular. These systems are now embedded in search engines, software development tools as well as office software. For many people, they have quickly become part of everyday computing. I personally find these services incredibly useful, particularly for exploring unfamiliar topics and as a general productivity aid. However, I also think that the way these services are advertised and consumed can pose a danger, especially if we get into the habit of trusting their output without further scrutiny.

Contents Introduction Pitfalls Inverse Laws of Robotics Non-Anthropomorphism Non-Deference Non-Abdication of Responsibility Conclusion Pitfalls Certain design choices in modern AI systems can encourage uncritical acceptance of their output. For example, many popular search engines are already highlighting answers generated by AI at the very top of the page. When this happens, it is easy to stop scrolling, accept the generated answer and move on. Over time, this could inadvertently train users to treat AI as the default authority rathe...

## 129. Survival Analysis of the Supreme Court

Source: Entropic Thoughts | Published: Tue, 27 Jan 2026 00:00:00 +0100

Link: <https://entropichoughts.com/survival-analysis-of-the-supreme-court>

(No summary)

## 130. Nvidia Stock Crash Prediction

Source: Entropic Thoughts | Published: Tue, 20 Jan 2026 00:00:00 +0100

Link: <https://entropichoughts.com/nvidia-stock-crash-prediction>

(No summary)

## 131. ACX 2025 Prediction Contest Retrospective

Source: Entropic Thoughts | Published: Thu, 15 Jan 2026 00:00:00 +0100

Link: <https://entropichoughts.com/acx-2025-prediction-contest-retrospective>

(No summary)

## 132. What I've been reading recently - Jan 10, 2026

Source: Dwarkesh Podcast | Published: Sat, 10 Jan 2026 20:30:07 GMT

Link: <https://www.dwarkesh.com/p/notes-jan-10-2026>

Nonlinear dynamics and Chaos, Machines of Loving Grace, Max Hodak's theory of consciousness,

Neural network training makes beautiful fractals

133. Adam Marblestone ? AI is missing something fundamental about the brain

Source: Dwarkesh Podcast | Published: Tue, 30 Dec 2025 17:07:17 GMT

Link: <https://www.dwarkesh.com/p/adam-marblestone>

The brain's secret sauce is its reward functions, not its architecture.

134. An audio version of my blog post, Thoughts on AI progress (Dec 2025)

Source: Dwarkesh Podcast | Published: Tue, 23 Dec 2025 20:24:48 GMT

Link: <https://www.dwarkesh.com/p/thoughts-on-ai-progress-dec-2025-video>

What are we scaling? I'm confused why some people have short timelines and at the same time are bullish on the current scale up of reinforcement learning atop LLMs. If we're actually close to a human-like learner, this whole approach of training on verifiable outcomes is doomed.

Currently the labs are trying to bake in a bunch of skills into these models through

?mid-training? - there's an entire supply chain of companies building RL environments which teach the model how to navigate a web browser or use Excel to write financial models . Either these models will soon learn on the job in a self directed way - making all this pre-baking pointless - or they won't - which means AGI is not imminent. Humans don't have to go through a special training phase where they need to rehearse every single piece of software they might ever need to use. Beren Millidge made interesting points about this in a recent blog post :

When we see frontier models improving at various benchmarks we should think not just of increased scale and clever ML research ideas but billions of dollars spent paying PhDs, MDs, and other experts to write questions and provide example answers and reasoning targeting these...

135. Decoding Google: Converting a Black Box to a White Box

Source: skull's blog | Published: Fri, 01 Nov 2024 00:00:00 GMT

Link: <https://brutecat.com/articles/decoding-google>

We've all been there - staring at Google's search box, overwhelmed by the maze of complexity hiding behind that minimalist interface, thinking it's impossible to break in. The key to decoding Google? Converting the attack surface from a black box to a white box.

136. Minus World

Source: Tedium: The Dull Side of the Internet. | Published: 2026-01-23T01:06:35Z

Link: <https://feed.tedium.co/link/15204/17261173/oneplus-decline-shutdown-rumors-commentary>

Rumors of OnePlus' possible death are heating up, which would be a real shame, given how much the smartphone market has already contracted. It's not true, but their grip is slipping.

137. They Were Robbed

Source: Tedium: The Dull Side of the Internet. | Published: 2026-01-19T14:08:25Z

Link: <https://feed.tedium.co/link/15204/17256459/billboard-bubbling-under-hot-100-chart-history>

The tale of the Billboard Bubbling Under Hot 100 chart, the place where hits go to die?in some cases, over and over again. Let's talk about the chart through the lens of its two most iconic artists.

138. Splitting Machines

Source: Tedium: The Dull Side of the Internet. | Published: 2026-01-12T20:28:50Z

Link: <https://feed.tedium.co/link/15204/17251925/vmware-virtualization-history>

How the virtual machine, a foundational element of cloud computing, found its modern footing after a couple of scientists proved a couple of theorems wrong.

139. A Christmas 2007 video present from Old VCR with Jack Tramiel et al

Source: Old Vintage Computing Research | Published: 2025-12-24T15:50:00.000-08:00

Link: <https://oldvcr.blogspot.com/2025/12/a-christmas-2007-video-present-from-old.html>

A very happy holiday season and Merry Christmas to those of you who celebrate it (timezone may vary). Also, I don't think I nearly say thanks enough to my regular patrons through Ko-fi, and I want to also thank them on behalf of the geriatric systems their generosity ? and all of you who have chipped in at one time or another ? helps keep running. I've got more projects to finish in 2026 and I hope you enjoy them. Anyway, here's a gift for you which I forgot I still had kicking around. This is a raw cut from the 2007 Computer History Museum 25th anniversary symposium of the Commodore 64, with Jack Tramiel himself, plus Steve Wozniak, William Lowe and others (like Lee Felsenstein and Al Alcorn from the audience). I was fortunate enough to get into what was a sell-out standing-room only crowd with my trusty JVC DV camera and tape the whole thing, then dash back to So Cal and quickly dump and title it on the Power Mac Quad G5. Any camera glitches, plus a couple dropouts where I had to quickly change DV tapes, and bad titles are of course totally my fault. (In fact, there are indeed at least two errors. Can you find them?) This video is so old that it was actually originally uploaded...

#### 140. The Texas Instruments CC-40 invades Gopherspace (plus TI-74 BASICalc)

Source: Old Vintage Computing Research | Published: 2025-12-20T20:04:00.000-08:00

Link: <https://oldvcr.blogspot.com/2025/12/the-texas-instruments-cc-40-invades.html>

I've mentioned on the blog several times the continuum that exists between handheld computers and pocket computers , battery powered devices in rather small form factors that are nevertheless fully-fledged general purpose computers ? arguably more so than the modern locked-down smartphone has become. Some of these diminutive systems are best considered "handhelds," with larger size, larger keyboards, more power and (often) less battery life, and some are definitely "pocket computers," with smaller size, smaller keys, less power and (usually) better battery life. For example, systems like the Tandy PC-4/Casio PB-100 or Tandy PC-3/Sharp PC-1250 would be considered "definitely a pocket computer," while the Epson HX-20 or Kyotronic 85 systems like the NEC PC-8201A or TRS-80 Model 100 would be considered "definitely a handheld computer," and you can probably think up some examples in between. Well, here's a notable example of one single architecture that birthed both types of machine, and it came from a company not really noted for either one: Texas Instruments. TI certainly made calculators and many of those were programmable by some means, but neither handheld computers nor pocket co...

#### 141. Oblast: a better Blasto game for the Commodore 64

Source: Old Vintage Computing Research | Published: 2025-12-06T15:55:00.000-08:00

Link: <https://oldvcr.blogspot.com/2025/12/oblast-better-blasto-game-for-commodore.html>

Way back (well, six months ago, anyway), when I was wiring up a Gremlin Blasto arcade board , we talked at length about this 1978 arcade game's history and its sole official home computer port by Milton Bradley to the Texas Instruments 99/4A. In the single player mode you run around in a maze and try to blow up all the mines, which can set off sometimes impressive chain reactions, all the while making sure you yourself don't go up in flames in the process. The TI-99/4A version was the Blasto I originally remember playing as I never did play Blasto in the arcades. (Also, for the record, we're not talking about Sony's unrelated Blasto for the PlayStation which, other than having the voice talents of the late and lamented Phil Hartman, was apparently a traumatic slog both for its developers and the few people who actually played it.) To the credit of its three composite authors, it is a competent and accurate conversion that also adds configurable options, colour graphics and music; in fact, TI's Blasto is probably my favourite game on the /4A, more so than any other cartridge. On the other hand, because it's an accurate conversion, it also inherits all of the original's weaknesses,...

#### 142. Plugins case study: mdBook preprocessors

Source: Eli Bendersky's website | Published: 2025-12-17T18:11:00-08:00

Link: <https://eli.thegreenplace.net/2025/plugins-case-study-mdbook-preprocessors/>

mdBook is a tool for easily creating books out of Markdown files. It's very popular in the Rust ecosystem, where it's used (among other things) to publish the official Rust book . mdBook has a simple yet effective plugin mechanism that can be used to modify the book output in arbitrary

?

#### 143. Revisiting "Let's Build a Compiler"

Source: Eli Bendersky's website | Published: 2025-12-09T20:40:00-08:00

Link: <https://eli.thegreenplace.net/2025/revisiting-lets-build-a-compiler/>

There's an old compiler-building tutorial that has become part of the field's lore: the Let's Build a Compiler series by Jack Crenshaw (published between 1988 and 1995). I ran into it in 2003 and was very impressed, but it's now 2025 and this tutorial is still being mentioned quite often ?

#### 144. Making PCBs

Source: BogdanTheGeek's Blog | Published: Fri, 13 Jun 2025 15:40:03 +0100

Link: <https://bogdanthegeek.github.io/blog/insights/making-pcb/>

Why Bother? So, why would anyone bother making PCBs at home? The truth is that not many people do any more. In the year 2025, we have the ability to order high quality PCBs from China for next to nothing. If you just want a high quality board and you don't mind waiting a week or two, these are a great option. I however, don't like to wait. I want a PCB in my hands the same day I finish the design.

#### 145. J-Link RTT for the Masses

Source: BogdanTheGeek's Blog | Published: Sun, 01 Jun 2025 14:47:43 +0100

Link: <https://bogdanthegeek.github.io/blog/insights/jlink-rtt-for-the-masses/>

TLDR; You can use semihosting on any ARM CPU to send and receive data in a few lines of code.

Click here to see an example implementation on a 10p ?disposable? microcontroller. Context

There are many ways to debug embedded projects. For high speed stuff, you might toggle a pin at the beginning and end of a subroutine. On a more advanced project, you might have structured logging to a file-system.

#### 146. MicroPPPID

Source: BogdanTheGeek's Blog | Published: Thu, 15 May 2025 17:24:53 +0100

Link: <https://bogdanthegeek.github.io/blog/projects/microp pid/>

TLDR; You can find the project repository here. Motivation I recently took up pottery as my new favourite hobby and I needed a better way to control the firing of my pots. For my first few pots, I just re-purposed my very cheap metal casting forge. I built this out of eight 1" thick fire bricks, some fire cement, an electric stove top ?burner? and a cheap PID controller. The cheap PID controller was horrible.

#### 147. Homeworld

Source: The Digital Antiquarian | Published: Fri, 02 Jan 2026 16:49:28 +0000

Link: <https://www.filfre.net/2026/01/homeworld/>

This article tells part of the story of real-time strategy. The first-person shooter and the real-time-strategy game, those two genres that had come to absolutely dominate mainstream computer gaming by the end of the 1990s, were surprisingly different in their core technologies. The FPS was all about 3D graphics, as aided and abetted by the [?]

#### 148. Huge for the Holidays: Epic Interactive Fiction of the Millennial Period

Source: The Digital Antiquarian | Published: Fri, 19 Dec 2025 16:46:39 +0000

Link: <https://www.filfre.net/2025/12/huge-for-the-holidays-epic-interactive-fiction-of-the-millenni al-period/>

This article tells part of the story of post-commercial interactive fiction. After fighting the tendency for decades, I've finally learned to accept that I'm a bit of a conservative, if not a downright Luddite, when it comes to my interactive fiction ? or my text adventures, as I still persist in calling them. Despite its [?]

#### 149. Mr. Roberts Goes to Hollywood, Part 2: The Producer

Source: The Digital Antiquarian | Published: Fri, 05 Dec 2025 17:42:29 +0000

Link: <https://www.filfre.net/2025/12/mr-roberts-goes-to-hollywood-part-2-the-producer/>

This article tells part of the story of Chris Roberts. With the Wing Commander movie having gone down in flames, there was nothing left for Chris Roberts and the rest of Digital Anvil to do but go back to making games. This undoubtedly pleased Microsoft, which had been waiting for some return on its generous [?]

## 150. Stop Using (only) GitHub Releases

Source: Hugo Tunius - Blog | Published: 2024-01-20T00:00:00+00:00

Link: <https://hugotunius.se/2024/01/20/stop-using-github-releases.html>

The other day at work I, accidentally, roped myself into upgrading some dependencies in our Rust services. These were breaking changes, so not just a case of running cargo update . I had to understand the changes and make the appropriate modifications to our code. Adopting breaking changes can be frustrating in the best of times, but it was particularly annoying this time because none of these projects kept a CHANGELOG.md files, although they all had release notes on GitHub. GitHub?s releases feature allows you to combine a git tag with release notes, metadata, and files(binaries, source code etc). While useful, GitHub?s releases have many downsides and consuming them adds friction to the task of understanding changes in a project. They can be a useful supplement to a CHANGELOG.md / HISTORY.md / RELEASES.md file but should not be the only place where release notes are recorded. The problems with exclusively using GitHub?s releases feature are: Pagination , which makes it hard to search across the full releases and makes cross-referencing between different versions cumbersome. Not truly being apart of the repository , which means, if you want to look at the source checked out, you...

## 151. The Great Pendulum

Source: Hugo Tunius - Blog | Published: 2023-07-09T00:00:00+01:00

Link: <https://hugotunius.se/2023/07/09/the-great-pendulum.html>

17 odd years ago when I started programming, PHP was all the rage. Javascript was steadily gaining traction. Django and Ruby on Rails were in their infancy, but promised greatly increased productivity. A few years later, inspired by Ruby?s fame, Coffeescript became a mainstay in the Javascript ecosystem. Statically compiled, typed languages, used to build monolithic web applications, were rapidly falling out of favour. In 2023 the trend is reversing, static compilation and types are cool again. Monoliths are making a comeback. The pendulum is turning. The first serious web application I built used an emerging pattern, AJAX( A synchronous J avascript a nd X ML), where the server didn?t just return HTML, but also Javascript that could fetch further data and update the HTML later. Several years later, when many started complaining about the high cost of React and SPAs I started thinking in terms of the pendulum. Entirely server rendered applications is one extreme of this particular pendulum, entirely client side rendered applications being the other. In the pursuit of web applications that delivered snappier user experiences the industry, arguably, overshot past the equilibrium point...

## 152. NFTs, How Do They Work?

Source: Hugo Tunius - Blog | Published: 2022-01-16T00:00:00+00:00

Link: <https://hugotunius.se/2022/01/16/nfts-how-do-they-work.html>

Freaking magnets NFTs, how do they work? In this post I?ll try to explain NFTs in a way that?s mostly accurate, but requires minimal technical understanding. I?m going to assume the reader is familiar with excel style software and Google Sheets in particular. At its core every NFT project is like a single Google Sheet. It has a creator who has some special permission to modify the sheet. An NFT within the project is like a single row in the sheet. Each row contains only two things: a name and an id. For example, if my NFT has 10,000 unique tokens then I could use the ids from 1 through 10,000 to identify each token. To begin with the sheet is empty. The creator might assign themselves a few tokens by creating new rows with their own name and some ids. Then they invite other people to ?mint? new tokens and gain rows with their name in them. In order to mint, the user has to pay the creator a small fee(for example via

Venom). In return for payment a row is added to the sheet containing the user?s name and a randomly selected unassigned id. In addition the user is given permission to change who owns the id they were allocated, this facilitates selling their token. Eventually all 10,0...

#### 153. Building a 1997 Quake PC: Benchmarking Vquake

Source: Fabien Sanglard | Published: 13 Jan 2026 00:00:00 +0000

Link: [https://fabiensanglard.net/quake\\_pc/vquake/index.html](https://fabiensanglard.net/quake_pc/vquake/index.html)

(No summary)

#### 154. Building a 1997 Quake PC: Benchmarking Quake

Source: Fabien Sanglard | Published: 12 Jan 2026 00:00:00 +0000

Link: [https://fabiensanglard.net/quake\\_pc/quake/index.html](https://fabiensanglard.net/quake_pc/quake/index.html)

(No summary)

#### 155. Building a 1997 Quake PC!

Source: Fabien Sanglard | Published: 08 Jan 2026 00:00:00 +0000

Link: [https://fabiensanglard.net/quake\\_pc/index.html](https://fabiensanglard.net/quake_pc/index.html)

(No summary)

#### 156. Pulling the plug on Roamflow

Source: Blog - Daniel Wirtz | Published: Mon, 26 Jul 2021 19:10:00 GMT

Link: <https://danielwirtz.com/blog/roamflow>

Why I'm stopping my work on Roamflow

#### 157. Task Management in Roam Research

Source: Blog - Daniel Wirtz | Published: Tue, 26 Jan 2021 22:00:00 GMT

Link: <https://danielwirtz.com/blog/roam-research-task-management>

A walkthrough of my task management system

#### 158. Journal Prompts in Roam Research

Source: Blog - Daniel Wirtz | Published: Sat, 23 Jan 2021 09:07:00 GMT

Link: <https://danielwirtz.com/blog/journal-prompts-roam-research>

A step by step tutorial

#### 159. Finding a 27-year-old easter egg in the Power Mac G3 ROM

Source: Downtown Doug Brown | Published: Tue, 24 Jun 2025 07:49:28 +0000

Link: <https://www.downtowndougbrown.com/2025/06/finding-a-27-year-old-easter-egg-in-the-power-mac-g3-rom/>

I was recently poking around inside the original Power Macintosh G3?s ROM and accidentally discovered an easter egg that nobody has documented until now. This story starts with me on a lazy Sunday using Hex Fiend in conjunction with Eric Harmon?s Mac ROM template (ROM Fiend) to look through the resources stored in the Power [?]

#### 160. Modifying an HDMI dummy plug?s EDID using a Raspberry Pi

Source: Downtown Doug Brown | Published: Sun, 15 Jun 2025 14:17:58 +0000

Link: <https://www.downtowndougbrown.com/2025/06/modifying-an-hdmi-dummy-plugs-edid-using-a-raspberry-pi/>

I recently found myself needing to change the monitor that a cheap HDMI ?dummy plug? pretended to be. It was a random one I had bought on Amazon several years ago that acted as a 4K monitor, and I needed it to be something simpler that didn?t support a 4K resolution. The story behind why [?]

#### 161. Please don?t ship heavy, fragile vintage computers. They will be destroyed.

Source: Downtown Doug Brown | Published: Sun, 25 May 2025 17:41:43 +0000

Link: <https://www.downtowndougbrown.com/2025/05/please-dont-ship-heavy-fragile-vintage-computer>

s-they-will-be-destroyed/

As part of my research into the Macintosh Performa 550's factory recovery partition, I paid a lot of attention to eBay listings for these computers. I came to an interesting discovery that I had already suspected: big CRT-based Macs in this form factor are regularly damaged in shipping after being sold on eBay. Most vintage [?]

## 162. The Chair

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/fiction/the-chair/>

Soren Kastner entered the negotiation chamber on Valos Station and immediately identified the anomaly. Five chairs for six delegations. The Meritocracy of Eridani had sent no representative, only a message that their seat should remain vacant. The other delegates shifted uncomfortably as Soren examined the empty chair with measured interest. General Tarask of the Procyon Coalition leaned forward, voice pitched low.

## 163. AUC

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/projects/auc/>

(No summary)

## 164. Events

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/books-and-events/events/>

#callout { background: rgba(var(--accent-rgb), 0.08); padding: 1.5em 1.25em; border-radius: 3px; display: flex; flex-direction: row; margin-bottom: 20px; } #callout-inner { margin-left: 1em; } @media (max-width: 767px) { #callout { padding: 1.5em 0.75em 1.5em 0.6em; } #callout-inner { margin-left: 0.5em; } } ? Please note: With the exception of some existing commitments, I will not be travelling out of Tasmania for events for a year or two. Please feel free to invite me to events within Tasmania, or held remotely/online! I'll update this message when I'm travelling again! Invite me then! - Paris, August 2025 I'm happy to be invited to speak at events on a range of topics, including but not limited to artificial intelligence, the intersection and implications of technology and creativity, work, and culture, the intersection of technology and creativity, game development and design, creative arts, writing, space, technology, and beyond. Email me at paris AT paris.id.au to chat. You can find a bio and headshots here .

## 165. Laine Cloud

Source: Philip Laine | Published: Sun, 18 Aug 2019 00:00:00 +0000

Link: <https://philiplaine.com/portfolio/laine-cloud/>

This is a fun hobby project based around a couple of Raspberry Pi. The goal was to run container based workloads across all of computers. The project has evolved into a Kubernetes cluster with persistent storage and a VPN tunnel to allow for serving of public websites.

## 166. Alfrodull

Source: Philip Laine | Published: Sat, 17 Aug 2019 00:00:00 +0000

Link: <https://philiplaine.com/portfolio/alfrodull/>

Alfrodull is a utility to easily control LED devices connected to a computer based on predefined events. Everything is configured with a single file, allowing for simple configuration of event type, light color and transition animation in seconds.

## 167. Ubiquiti And VPN Clients

Source: Philip Laine | Published: Thu, 08 Aug 2019 17:40:17 +0200

Link: <https://philiplaine.com/posts/ubiquiti-and-vpn-clients/>

I have for some time wanted to learn about the different VPN services offered by AWS. Having invested some money into a full Unifi setup I thought that it would be a good idea to setup an

AWS Site-to-Site VPN and AWS Client VPN connection between my lab network and my AWS account. I struggled for some time, messing around with the settings trying to get my USG to connect to the VPN server but it would not work. I probably wasted a couple of days of my free time pulling my hair out not really understanding why it wouldn't connect.

#### 168. The Software Essays that Shaped Me

Source: Refactoring English | Published: Tue, 30 Sep 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/software-essays-that-shaped-me/>

I started reading software blogs before I got my first programming job 20 years ago. At this point, I've read thousands of blog posts and essays about software, but only a small handful stuck in my mind and changed the way I think. ?The Joel Test: 12 Steps to Better Code? by Joel Spolsky (2000) ?Parse, don't validate? by Alexis King (2019) ?No Silver Bullet - Essence and Accident in Software Engineering? by Fred Brooks (1986) ?Choices? by Joel Spolsky (2000) ?Application compatibility layers are there for the customer, not for the program? by Raymond Chen (2010) ?Don't Put Logic in Tests? by Erik Kuefler (2014) ?A little bit of plain Javascript can do a lot? by Julia Evans (2020) ?Choose Boring Technology? by Dan McKinley (2015) ?I've locked myself out of my digital life? by Terence Eden (2022) Bonus: Brad Fitzpatrick on parsing user input (2009) ?The Joel Test: 12 Steps to Better Code? by Joel Spolsky (2000) Joel Spolsky is the greatest software blogger of all time. His essays have informed so much of my approach to software that it was hard to pick out just one, but ?The Joel Test? is my favorite.

#### 169. Adam Gordon Bell on Attracting Customers through Blogging

Source: Refactoring English | Published: Fri, 12 Sep 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/interview-adam-gordon-bell/>

Adam Gordon Bell is a developer, blogger, and the host of the software engineering podcast, CoRecursive . I interviewed Adam about writing and his success at attracting customers through his blog posts when he worked at Earthly. We talk about: How Adam consistently wrote blog posts that reached the front page of Hacker News Discovering blog topics that attract potential customers Techniques Adam used to strengthen his writing The value of writing respectfully about your competitors Transcript Getting started with blogging Investing outsized effort for outsized returns Getting early feedback on blog posts Tools for improving your writing Finding focus for writing Crafting blog post titles Panel-of-experts style of blogging Blogging respectfully about your competitor Improving your writing through imitation Note: I've lightly edited the transcript for brevity.

#### 170. Reader Feedback about my Chapter List

Source: Refactoring English | Published: Fri, 29 Aug 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/chapter-interest-results/>

For the past eight months, I've been working on my book, Refactoring English: Effective Writing for Software Developers . I'm publishing the book incrementally as I write, which means readers give me feedback about the book in real time. I started writing the chapters I felt most confident readers would like, but now that I'm about 50% complete, I want to make sure the remaining chapters are things my readers actually want to learn.

#### 171. Stop Using DRM on your Books

Source: Beej's Bit Bucket | Published: Thu, 09 Oct 2025 00:00:00 +0000

Link: <http://beej.us/blog/data/on-drm/>

Why not to protect your books with DRM.

#### 172. Adding Mastodon Comments to your Blog

Source: Beej's Bit Bucket | Published: Fri, 21 Feb 2025 00:00:00 +0000

Link: <http://beej.us/blog/data/mastodon-comments/>

Comments powered by Mastodon? Yes, please!

#### 173. How Do You Prove Humanity?

Source: Beej's Bit Bucket | Published: Thu, 16 Mar 2023 00:00:00 +0000

Link: <http://beej.us/blog/data/proof-of-humanity/>

Reflections on the new wave of AI.

#### 174. whats-wrong-with-a-for-loop

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/tips/whats-wrong-with-a-for-loop>

(No summary)

#### 175. my-rustfmt-toml

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/tips/my-rustfmt-toml>

(No summary)

#### 176. git-collapse-commits

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/tips/git-collapse-commits>

(No summary)

#### 177. The History of Xerox

Source: Abort Retry Fail | Published: Sun, 02 Nov 2025 15:08:50 GMT

Link: <https://www.abortretry.fail/p/the-history-of-xerox>

A Monochromatic Star

#### 178. SRI & ARC

Source: Abort Retry Fail | Published: Sat, 27 Sep 2025 22:56:11 GMT

Link: <https://www.abortretry.fail/p/sri-and-arc>

The Dawn of the Information Age

#### 179. The QNX Operating System

Source: Abort Retry Fail | Published: Sun, 31 Aug 2025 21:32:18 GMT

Link: <https://www.abortretry.fail/p/the-qnx-operating-system>

Quantum Software and the microkernel UNIX

#### 180. What If We Made Advertising Illegal?

Source: Simone Salis | Published: Wed, 02 Apr 2025 16:22:48 GMT

Link: <https://simone.org/advertising/>

What if we banned all advertising? Not regulate it?abolish it. This proposal would transform manipulation machines, and maybe save democracy itself. A thought experiment worth considering.

#### 181. The Antiportfolio: Counter-advice for Aspiring Artists

Source: Simone Salis | Published: Sat, 01 Mar 2025 13:00:07 GMT

Link: <https://simone.org/antiportfolio/>

Why ten years ago I quit my passion as a job to follow my passions as just passions.

#### 182. Your Teenage Self is Not a Bug, It's a Feature

Source: Simone Salis | Published: Tue, 21 Jan 2025 14:01:12 GMT

Link: <https://simone.org/resistance/>

Between middle-class stability and creative resistance, there's a teenage part of us keeping watch. Maybe it's right.

#### 183. The Department of War Just Shot the Accountants and Opted for Speed

Source: Steve Blank | Published: Tue, 11 Nov 2025 14:00:12 +0000

Link: <https://steveblank.com/2025/11/11/the-department-of-war-just-shot-the-accountants-and-opted-for-speed/>

Last week the Department of War finally killed the last vestiges of Robert McNamara's 1962 Planning, Programming, and Budgeting System (PPBS). The DoW has pivoted from optimizing cost and performance to delivering advanced weapons at speed. Taking decades to deliver weapons is no longer an option. The DoW has joined the 21st century and adopted [?]

#### 184. It only took 20 years, but the Strategic Management Society now Believes the Lean Startup is a Strategy

Source: Steve Blank | Published: Thu, 30 Oct 2025 13:00:42 +0000

Link: <https://steveblank.com/2025/10/30/it-only-took-20-years-but-the-strategic-management-society-now-believes-the-lean-startup-is-a-strategy-i-got-an-award-for-it/>

I've always thought of myself as a practitioner. In the startups I was part of, the only strategy? were my marketing tactics on how to make the VP of Sales the richest person in the company. After I retired, I created Customer Development and co-created the Lean Startup as a simple methodology which codified founders [?]

#### 185. How to Sell to the Dept of War ? The 2025 PEO Directory ? Now with 500 more names

Source: Steve Blank | Published: Wed, 15 Oct 2025 13:00:11 +0000

Link: <https://steveblank.com/2025/10/15/how-to-sell-to-the-dept-of-war-the-2025-peo-directory-now-with-500-more-names/>

The October 2025 PEO Directory ? Update 2. The Department of War (DoW) is one of the world's largest organizations. If you're a startup trying to figure out who to call on and how to navigate the system, it can be ? to put it politely ? challenging. Those inside the DoW have little perspective of [?]

#### 186. Load and store forwarding in the Toy Optimizer

Source: Max Bernstein's Blog | Published: Wed, 24 Dec 2025 00:00:00 +0000

Link: [https://bernsteinbear.com/blog/toy-load-store/?utm\\_source=rss](https://bernsteinbear.com/blog/toy-load-store/?utm_source=rss)

Another entry in the Toy Optimizer series . A long, long time ago (two years!) CF Bolz-Tereick and I made a video about load/store forwarding and an accompanying GitHub Gist about load/store forwarding (also called load elimination) in the Toy Optimizer. I said I would write a blog post about it, but never found the time?it got lost amid a sea of large life changes. It's a neat idea: do an abstract interpretation over the trace, modeling the heap at compile-time, eliminating redundant loads and stores. That means it's possible to optimize traces like this:  
v0 = ... v1 = load(v0, 5) v2 = store(v0, 6, 123) v3 = load(v0, 6) v4 = load(v0, 5) v5 = do\_something(v1, v3, v4) into traces like this: v0 = ... v1 = load(v0, 5) v2 = store(v0, 6, 123) v5 = do\_something(v1, 123, v1) (where load(v0, 5) is equivalent to \*(v0+5) in C syntax and store(v0, 6, 123) is equivalent to \*(v0+6)=123 in C syntax) This indicates that we were able to eliminate two redundant loads by keeping around information about previous loads and stores. Let's get to work making this possible. The usual infrastructure We'll start off with the usual infrastructure from the Toy Optimizer series : a very stringly-typed repres...

#### 187. ZJIT is now available in Ruby 4.0

Source: Max Bernstein's Blog | Published: Wed, 24 Dec 2025 00:00:00 +0000

Link: [https://bernsteinbear.com/blog/launch-zjit/?utm\\_source=rss](https://bernsteinbear.com/blog/launch-zjit/?utm_source=rss)

Originally published on Rails At Scale . ZJIT is a new just-in-time (JIT) Ruby compiler built into the reference Ruby implementation, YARV , by the same compiler group that brought you YJIT. We (Aaron Patterson, Aiden Fox Ivey, Alan Wu, Jacob Denbeaux, Kevin Menard, Max Bernstein, Maxime Chevalier-Boisvert, Randy Stauner, Stan Lo, and Takashi Kokubun) have been working on ZJIT since the beginning of this year. In case you missed the last post, we're building a new compiler for Ruby because we want to both raise the performance ceiling (bigger compilation unit size and SSA IR) and encourage more outside contribution (by becoming a more traditional method compiler). It's been a long time since we gave an official update on ZJIT. Things are going well. We're excited to share our progress with you. We've done a lot since May . In brief ZJIT is compiled by default?but not enabled by default?in Ruby 4.0. Enable it by passing the --zjit flag or the RUBY\_ZJIT\_ENABLE environment variable or calling

RubyVM::ZJIT.enable after starting your application. It's faster than the interpreter, but not yet as fast as YJIT. Yet. But we have a plan, and we have some more specific numbers below. The TL;...

#### 188. How to annotate JITed code for perf/samply

Source: Max Bernstein's Blog | Published: Thu, 18 Dec 2025 00:00:00 +0000

Link: [https://bernsteinbear.com/blog/jit-perf-map/?utm\\_source=rss](https://bernsteinbear.com/blog/jit-perf-map/?utm_source=rss)

Brief one today. I got asked ?does YJIT/ZJIT have support for [Linux] perf?? The answer is yes, and it also works with samply (including on macOS!), because both understand the perf map interface . This is the entirety of the implementation in ZJIT 1 : fn register\_with\_perf ( iseq\_name : String , start\_ptr : usize , code\_size : usize ) { use std :: io :: Write ; let perf\_map = format! ( "/tmp/perf-{}.map" , std :: process :: id () ); let Ok ( file ) = std :: fs :: OpenOptions :: new () .create ( true ) .append ( true ) .open ( & perf\_map ) else { debug! ( "Failed to open perf map file: {perf\_map}" ); return ; } ; let mut file = std :: io :: BufWriter :: new ( file ); let Ok ( \_ ) = writeln! ( file , "{start\_ptr:x} {code\_size:x} zjit::{iseq\_name}" ) else { debug! ( "Failed to write {iseq\_name} to perf map file: {perf\_map}" ); return ; } ; } Whenever you generate a function, append a one-line entry consisting of START SIZE symbolname to /tmp/perf-{PID}.map . Per the Linux docs linked above, START and SIZE are hex numbers without 0x. symbolname is the rest of the line, so it could contain special characters. You can now happily run perf record your\_jit [...] or samply record your\_jit [...]

#### 189. Weekly Update 487

Source: Troy Hunt | Published: Sun, 18 Jan 2026 08:43:10 GMT

Link: <https://www.troyhunt.com/weekly-update-487/>

I thought Scott would cop it first when he posted about what his solar system really cost him last year . "You're so gonna get that stupid AI-slop response from some people", I joked. But no, he got other stupid responses instead! And I got the AI-slop

#### 190. Weekly Update 486

Source: Troy Hunt | Published: Fri, 16 Jan 2026 06:39:10 GMT

Link: <https://www.troyhunt.com/weekly-update-486/>

I'm in Oslo! Flighty is telling me I've flown in or out of here 43 times since a visit in 2014 set me on a new path professionally and, many years later, personally . It's special here, like a second home that just feels?

#### 191. Who Decides Who Doesn't Deserve Privacy?

Source: Troy Hunt | Published: Tue, 13 Jan 2026 11:41:40 GMT

Link: <https://www.troyhunt.com/who-decides-who-doesnt-deserve-privacy/>

Remember the Ashley Madison data breach? That was now more than a decade ago, yet it arguably remains the single most noteworthy data breach of all time. There are many reasons for this accolade, but chief among them is that by virtue of the site being expressly designed to facilitate

#### 192. Messing with bots

Source: Herman's blog | Published: 2025-11-14T11:17:28.208510+00:00

Link: <https://herman.bearblog.dev/messing-with-bots/>

As outlined in my previous two posts : scrapers are, inadvertently, DDoSing public websites. I've received a number of emails from people running small web services and blogs seeking advice on how to protect themselves. This post isn't about that. This post is about fighting back. When I published my last post, there was an interesting write-up doing the rounds about a guy who set up a Markov chain babbler to feed the scrapers endless streams of generated data. The idea here is that these crawlers are voracious, and if given a constant supply of junk data, they will continue consuming it forever, while (hopefully) not abusing your actual web server. This is a pretty neat idea, so I dove down the rabbit hole and learnt about Markov chains, and even picked up Rust in the process. I ended up building my own babbler that could

be trained on any text data, and would generate realistic looking content based on that data. Now, the AI scrapers are actually not the worst of the bots. The real enemy, at least to me, are the bots that scrape with malicious intent. I get hundreds of thousands of requests for things like .env , .aws , and all the different .php paths that could potentially sig...

### 193. Aggressive bots ruined my weekend

Source: Herman's blog | Published: 2025-10-29T12:08:40.897084+00:00

Link: <https://herman.bearblog.dev/agressive-bots/>

On the 25th of October Bear had its first major outage. Specifically, the reverse proxy which handles custom domains went down, causing custom domains to time out. Unfortunately my monitoring tool failed to notify me, and it being a Saturday, I didn't notice the outage for longer than is reasonable. I apologise to everyone who was affected by it. First, I want to dissect the root cause, exactly what went wrong, and then provide the steps I've taken to mitigate this in the future. I wrote about The Great Scrape at the beginning of this year. The vast majority of web traffic is now bots, and it is becoming increasingly more hostile to have publicly available resources on the internet. There are 3 major kinds of bots currently flooding the internet: AI scrapers, malicious scrapers, and unchecked automations/scrapers. The first has been discussed at length. Data is worth something now that it is used as fodder to train LLMs, and there is a financial incentive to scrape, so scrape they will. They've depleted all human-created writing on the internet, and are becoming increasingly ravenous for new wells of content. I've seen this compared to the search for low-background-radiation steel...

### 194. Attending MicroConf Europe 2025

Source: Herman's blog | Published: 2025-10-24T12:02:29.348048+00:00

Link: <https://herman.bearblog.dev/microconf-europe/>

Now that I've been home for about a month and have all my ducks back in a row, it's time I wrote about my trip to Istanbul for MicroConf Europe . MicroConf is a conference for bootstrapped (non-VC-tracked) founders. And while I don't necessarily view Bear as a business, MicroConf is the closest there is to an Industry Event for someone like me. First a note on Istanbul: I arrived a week early to explore the city and see the sights. I'm a bit of a history nerd, so being at the crossroads of where it all happened in Europe?going back thousands of years?was quite spectacular. I get up early, and wandering the empty streets of the old city before the tour groups flooded in was quite special. Also, the mosques dotting the skyline as viewed from the Bosphorus are like nothing I've ever seen. It's amazing to see human effort geared towards creating beautiful buildings. I know it's not economically viable, but imagine if cities were built with beauty and a cohesive aesthetic in mind. There were, however, a few negative characteristics of the city that grated at me, the main one being the hard separation of what is the tourist area and what isn't. Inside the old city all of the restaurants...

### 195. Things that made me think: Digital gardening, web degradation, and digital ghosts

Source: Tom Renner on My place to put things | Published: Mon, 01 Sep 2025 00:00:00 +0000

Link: <https://tomrenner.com/posts/tmmmt-2/>

This series is a place to collect interesting things I've seen, read, or heard, along with some brief thoughts (often incomplete and/or inconclusive) that they provoked. Garden History ? Maggie Appleton I'm so happy I stumbled upon this article. I am always grateful for new vocabulary that allows me better to express myself, and this is perfect - I want more Digital Gardens in the world. I do see the value in polishing content, but this is where the epistemic status tagging system laid out there really comes to the fore. Do I now want to convert this to a full garden-style site? Or perhaps just introduce different ?feeds?, laid out by theme, epistemic status, etc?

### 196. Optimising for trust

Source: Tom Renner on My place to put things | Published: Mon, 18 Aug 2025 00:00:00 +0000

Link: <https://tomrenner.com/posts/optimising-for-trust/>

TDD, BDD, DDD, Agile, SAFe, Scrum, Kanban, XP? there?s a lot of ways to skin a cat write code

in a professional environment. I take pride in being a person who is a non-ideologue when it comes to my code. There are many good ways of working, and they are all context-dependent. You can't apply the same things that worked when you were a two-person startup operating out of the proverbial garage and expect them to work once your hypothetical unicorn has reached a thousand-plus developers. Even within the same organisation, processes that work for one team can be catastrophic when applied to their neighbouring team.

#### 197. Things that made me think: Enshittification, apathy, and discrimination

Source: Tom Renner on My place to put things | Published: Thu, 24 Jul 2025 00:00:00 +0000

Link: <https://tomrenner.com/posts/tmmmt-1/>

This series is a place to collect interesting things I've seen, read, or heard, along with some brief thoughts (often incomplete and/or inconclusive) that they provoked. The rise of Whatever - eevee This is probably the best post about LLMs I've read, which is probably why I'm the millionth person to share it. It really sums up my emotional reaction to their meteoric rise: ?ew?, basically. The power of the argument is that it identifies a theme that runs through recent tech changes, of which LLMs are just the latest and greatest example: the lack of care for quality, and the realisation from Big Tech that consumers mostly are fine with mediocre output.

#### 198. Chat is a bad UI pattern for development tools

Source: Daniel De Laney | Published: Mon, 03 Feb 2025 00:00:00 +0000

Link: <http://danieldelaney.net/chat/>

Code forces humans to be precise. That's good. Computers need precision. But it also forces humans to think like machines. For decades we tried to fix this by making programming more human-friendly. Higher-level languages. Visual interfaces. Each step helped, but we were still translating human thoughts into computer instructions. AI was supposed to change everything. Finally, plain English could be a programming language. No syntax. No rules. Just say what you want. The first wave of AI coding tools squandered this opportunity. They make flashy demos but produce garbage software. People call them ?great for prototyping,? which means ?don't use this for anything real.? Many blame the AI models, saying we just need them to get smarter. This is wrong. Yes, better AI will make better guesses about what you mean. But when you're building serious software, you don't want guesses. Not even smart ones. You want to know exactly what you're building. Current AI tools pretend writing software is like having a conversation. It's not. It's like writing laws. You're using English, but you're defining terms, establishing rules, and managing complex interactions between everything you've said. T...

#### 199. Data Center Monitoring

Source: Daniel De Laney | Published: Thu, 25 Jul 2019 00:00:00 +0000

Link: <http://danieldelaney.net/rf/>

The Challenge The computing power that runs the world is hidden away in data centers that few people get to see. While many data centers are lights-out operations most of the time, people are still needed to update them, keep them running, and prevent and resolve outages. Those people need to know where their critical assets are in the labyrinth that is their global data center network. They need to know when areas get too hot, or get so cold and humid that condensation becomes a worry. In addition to data centers, large enterprises will also have smaller compute sites scattered across the nation or the world. Those sites are often physically unmanned with poor visibility into the health of critical systems. Operators need to know when potential issues arise and how to prioritize them. I helped solve both of those problems by operating as a product designer and product manager. Research At the start of a new initiative, I want to understand how it's meant to help the business, who the buyer is and what their behavior is like, and whether we can feasibly produce a solution which is both valuable and usable. I also want to understand the context of the challenge as best I can. I put...

#### 200. Artificial Intelligence

Source: Daniel De Laney | Published: Tue, 04 Jun 2019 00:00:00 +0000

Link: <http://danieldelaney.net/ai/>

In 2018 I worked with argodesign on an artificial intelligence client project, and Fast Company published an article on our work: This Is The World's First Graphical AI Interface . For confidentiality reasons I can't publicly go into more detail on the project than to link to that article. For the full case study, please contact me at hello@danieldelaney.net . AIGA Event During my time at argodesign we held an event with AIGA, the professional association for design, during which the team explained our point of view on artificial intelligence, and the way we approached designing interfaces for new technologies.

## 201. A Fast, Growable Array With Stable Pointers in C

Source: Daniel Hooper | Published: Tue, 05 Aug 2025 00:00:00 +0000

Link: [https://danielchasehooper.com/posts/segment\\_array/](https://danielchasehooper.com/posts/segment_array/)

My last article about generic data structures in C was written to set the stage for today's topic: A data structure with constant time indexing, stable pointers, and works well with arena allocators. Its been independently discovered by multiple programmers over the years and so goes by different names. A 2001 paper called it a ?levelwise-allocated pile? (bleh). Zig calls it a ?Segmented List?. Then there's C++ with std::deque, which is only superficially similar. Read the whole article on danielchasehooper.com ?

## 202. Type Safe Generic Data Structures in C

Source: Daniel Hooper | Published: Wed, 25 Jun 2025 00:00:00 +0000

Link: <https://danielchasehooper.com/posts/typechecked-generic-c-data-structures/>

See my follow-up article: ?A Fast, Growable Array With Stable Pointers in C? I write type safe generic data structures in C using a technique that I haven't seen elsewhere<sup>1</sup>. It uses unions to associate type information with a generic data structure, but we'll get to that. My approach works for any type of data structure: maps, arrays, binary trees? but for this article I illustrate the ideas by implementing a basic linked list. Read the whole article on danielchasehooper.com ?

## 203. A New Sudoku Layout With 81 Uniquely Shaped Cells

Source: Daniel Hooper | Published: Wed, 12 Mar 2025 00:00:00 +0000

Link: <https://danielchasehooper.com/posts-cracked-sudoku/>

My journey making a Sudoku variant that doesn't have rows or columns Read the whole article on danielchasehooper.com ?

## 204. Static Web Hosting on the Intel N150: FreeBSD, SmartOS, NetBSD, OpenBSD and Linux Compared

Source: IT Notes | Published: Wed, 19 Nov 2025 09:16:00 +0100

Link: <https://it-notes.dragas.net/2025/11/19/static-web-hosting-intel-n150-freebsd-smartos-netbsd-openbsd-linux/>

Compare static web hosting performance on an Intel N150 using the same nginx.conf across FreeBSD jails, SmartOS zones, NetBSD, OpenBSD and Linux, focusing on HTTP vs HTTPS and TLS CPU usage.

## 205. Self-hosting your Mastodon media with SeaweedFS

Source: IT Notes | Published: Thu, 06 Nov 2025 11:30:02 +0000

Link: <https://it-notes.dragas.net/2025/11/06/self-hosting-your-mastodon-media-with-seaweedfs/>

A practical guide to boosting Mastodon performance by self-hosting your media with SeaweedFS. Configure a fast, S3-compatible storage backend to efficiently handle your instance's files and take full control of your data.

## 206. The Email They Shouldn't Have Read

Source: IT Notes | Published: Wed, 08 Oct 2025 14:25:36 +0200

Link: <https://it-notes.dragas.net/2025/10/08/the-email-they-shouldnt-have-read/>

A real-world horror story about contracts, control, and the dark side of managed open source.

207. Two kinds of AI users are emerging. The gap between them is astonishing.

Source: Martin Alderson | Published: Sun, 01 Feb 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/two-kinds-of-ai-users-are-emerging/?utm\\_source=rss](https://martinalderson.com/posts/two-kinds-of-ai-users-are-emerging/?utm_source=rss)

A bifurcation is happening in AI adoption - power users shipping products in days versus everyone else generating meeting agendas. Enterprise tool choices are accelerating the divide.

208. Turns out I was wrong about TDD

Source: Martin Alderson | Published: Sun, 25 Jan 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/turns-out-i-was-wrong-about-tdd/?utm\\_source=rss](https://martinalderson.com/posts/turns-out-i-was-wrong-about-tdd/?utm_source=rss)

I used to be a TDD sceptic - too much time writing tests for features that might get deleted.

Then coding agents completely changed the economics of software testing.

209. Why sandboxing coding agents is harder than you think

Source: Martin Alderson | Published: Mon, 19 Jan 2026 00:00:00 GMT

Link: [https://martinalderson.com/posts/why-sandboxing-coding-agents-is-harder-than-you-think/?utm\\_source=rss](https://martinalderson.com/posts/why-sandboxing-coding-agents-is-harder-than-you-think/?utm_source=rss)

Permission systems, Docker sandboxing, and log file secrets - why current approaches to securing coding agents fall short and what we might need instead.

210. Attacking a stack based JavaScript virtual machine

Source: Cats with power tools | Published: Wed, 16 Jul 2025 17:42:03 GMT

Link: <https://blog.pixelmelt.dev/analysing-pistoljsvm/>

Analysing Pixelmelt's Javascript Virtualisation Obfuscation - A Research Study

211. Recursive vs Linear JSVM Disassembly

Source: Cats with power tools | Published: Tue, 01 Jul 2025 20:57:30 GMT

Link: <https://blog.pixelmelt.dev/recursive-vs-linear-jsvm-disassembly/>

A good VM is a constantly changing VM

212. A Clever (but not so good) Scraper Protection

Source: Cats with power tools | Published: Tue, 03 Jun 2025 17:12:15 GMT

Link: <https://blog.pixelmelt.dev/a-clever-but-not-so-good-scraper-protection/>

Text obfuscation with no JavaScript? How?

213. 500,000 tech workers have been laid off since ChatGPT was released

Source: Anil Dash | Published: 2026-01-06T00:00:00Z

Link: <https://anildash.com/2026/01/06/500k-tech-workers-laid-off/>

One of the key points I repeated when talking about the state of the tech industry yesterday was the salient fact that half a million tech workers have been laid off since ChatGPT was released in late 2022 . Now, to be clear, those workers haven't been laid off because their jobs are now being done by AI, and they've been replaced by bots. Instead, they've been laid off by execs who now have AI to use as an excuse for going after workers they've wanted to cut all along. This is important to understand for a few reasons. First, it's key just for having empathy for both the mindset and the working conditions of people in the tech industry. For so many outside of tech, their impression of what ?tech? means is whatever is the most recent transgression they've heard about from the most obnoxious billionaire who's made the news lately. But in many cases, it's the rank and file workers at that person's company who were the first victims of that billionaire's ego. Second, it's important to understand the big tech companies as almost the testing grounds for the techniques and strategies that these guys want to roll out on the rest of the economy, and on the rest of the world. Before they s...

214. How Markdown took over the world

Source: Anil Dash | Published: 2026-01-09T00:00:00Z

Link: <https://anildash.com/2026/01/09/how-markdown-took-over-the-world/>

Nearly every bit of the high-tech world, from the most cutting-edge AI systems at the biggest companies, to the casual scraps of code cobbled together by college students, is annotated and described by the same, simple plain text format. Whether you're trying to give complex instructions to ChatGPT, or you want to be able to exchange a grocery list in Apple Notes or copy someone's homework in Google Docs, that same format will do the trick. The wild part is, the format wasn't created by a conglomerate of tech tycoons, it was created by a curmudgeonly guy with a kind heart who right this minute is probably rewatching a Kubrick film while cheering for an absolutely indefensible sports team. But it's worth understanding how these simple little text files were born, not just because I get to brag about how generous and clever my friends are, but also because it reminds us of how the Internet really works: smart people think of good things that are crazy enough that they just might work , and then they give them away, over and over, until they slowly take over the world and make things better for everyone. Making Their Mark Though it's now a building block of the contemporary Internet,...

## 215. How to know if that job will crush your soul

Source: Anil Dash | Published: 2026-01-12T00:00:00Z

Link: <https://anildash.com/2026/01/12/will-that-job-crush-your-soul/>

Last week, we talked about one huge question, ? How the hell are you supposed to have a career in tech in 2026? ? That's pretty specific to this current moment, but there are some timeless, more perennial questions I've been sharing with friends for years that I wanted to give to all of you. They're a short list of questions that help you judge whether a job that you're considering is going to crush your soul or not. Obviously, not everyone is going to get to work in an environment that has perfect answers to all of these questions; a lot of the time, we're lucky just to get a place to work at all. But these questions are framed in this way to encourage us all to aspire towards roles that enable us to do our best work, to have the biggest impact, and to live according to our values. The Seven Questions If what you do succeeds, will the world be better? This question originally started for me when I would talk to people about new startups, where people were judging the basic idea of the product or the company itself, but it actually applies to any institution, at any size. If the organization that you're considering working for, or the team you're considering joining, is able to ac...

## 216. Making RSS More Fun

Source: matduggan.com | Published: Tue, 02 Dec 2025 09:14:45 GMT

Link: <https://matduggan.com/making-rss-more-fun/>

I wrote a browser extension to try and make finding random, independent websites more fun.

## 217. I broke and fixed my Ghost blog

Source: matduggan.com | Published: Thu, 16 Oct 2025 12:00:29 GMT

Link: <https://matduggan.com/i-broke-and-fixed-my-ghost-blog/>

In case you find this post after hours of google searching

## 218. Greenland is a beautiful nightmare

Source: matduggan.com | Published: Sat, 27 Sep 2025 08:56:13 GMT

Link: <https://matduggan.com/greenland-is-a-beautiful-nightmare/>

Greenland is a complicated topic here in Denmark. The former colony that is still treated a bit like a colony is something that inspires a lot of emotions. Greenland has been subjected to a lot of unethical experiments by Denmark, from taking their kids to wild experiments in criminal justice.

## 219. How to be less awkward

Source: Experimental History | Published: Tue, 06 Jan 2026 16:45:44 GMT

Link: <https://www.experimental-history.com/p/how-to-be-less-awkward>

a three-part treatment for a near-universal affliction

## 220. The drug that taught me how much I should suffer

Source: Experimental History | Published: Tue, 09 Dec 2025 17:46:17 GMT  
Link: <https://www.experimental-history.com/p/the-drug-that-taught-me-how-much>  
OR: grindset and bedrot

## 221. Secrets of the ancient memelords

Source: Experimental History | Published: Tue, 25 Nov 2025 19:17:23 GMT  
Link: <https://www.experimental-history.com/p/secrets-of-the-ancient-memelords>  
OR: pedants 'n' pinheads

## 222. The new homepage

Source: Aresluna | Published: 2025-08-27T01:46:35+0000  
Link: <https://aresluna.org/>

I finally updated my homepage! It's not really flashy in any way, but it has a lot of links to what I've done over the last years. As part of it, I restored some of my ?classic? Twitter and other social media threads, too. If you like my work, chances are you will find here something you enjoy.

## 223. The day Return became Enter

Source: Aresluna | Published: 2025-08-25T13:36:21Z  
Link: <https://aresluna.org/the-day-return-became-enter>  
A deep dive into the convoluted and fascinating story of one of the most important keys on the keyboard. (First online appearance of an essay originally published in printed form in December 2023. 3,100 words. 35 photos.)

## 224. Frame of preference

Source: Aresluna | Published: 2025-07-08T23:28:10.000Z  
Link: <https://aresluna.org/frame-of-preference>  
A story of early Mac settings told by a bunch of emulators. (8,800 words. 10 emulators.)

## 225. How to Build a Webhook System in Rails Using Sidekiq

Source: Keygen Blog | Published: Wed, 16 Jun 2021 05:00:00 GMT  
Link: <https://keygen.sh/blog/how-to-build-a-webhook-system-in-rails-using-sidekiq/>  
It's the heyday of SaaS and webhooks are all the rage. Learn how to build a webhook system for your service using Rails and Sidekiq.

## 226. How to License and Distribute a Private Node Module

Source: Keygen Blog | Published: Wed, 04 Aug 2021 05:00:00 GMT  
Link: <https://keygen.sh/blog/how-to-license-and-distribute-commercial-node-modules/>  
Learn how to license and distribute private Node.js packages using Keygen's software licensing and distribution API.

## 227. How to License and Distribute an Electron App

Source: Keygen Blog | Published: Tue, 07 Sep 2021 05:00:00 GMT  
Link: <https://keygen.sh/blog/how-to-license-and-distribute-an-electron-app/>  
Learn how to license and distribute a commercial Electron app using Keygen's software licensing and distribution API.

## 228. Claude Code

Source: Grant Slatton's Blog | Published: Sun, 11 May 2025 18:49:00 +0000  
Link: <https://grantslatton.com/claude-code>  
Everything I know about how to use it, so far

## 229. Solution-space Taste

Source: Grant Slatton's Blog | Published: Fri, 25 Apr 2025 07:42:00 +0000  
Link: <https://grantslatton.com/solution-space-taste>

## A graphical representation

### 230. Sorting Algorithm: The Game

Source: Grant Slatton's Blog | Published: Tue, 22 Apr 2025 04:31:40 +0000

Link: <https://grantslatton.com/sorting-algorithm-the-game>

How good are you?

### 231. Policy of transience

Source: Simon Tatham's quasi-blog | Published: 2025-05-09T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/transience/>

A description of some of my personal computing habits, focusing on avoiding things ?just turning out? to last for ever in a disorganised manner.

### 232. Iconography of the PuTTY tools

Source: Simon Tatham's quasi-blog | Published: 2025-03-12T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/putty-icons/>

A PuTTY user asked about the history of its logo and icon collection. I discuss the design intentions and technical aspects.

### 233. Git without a forge

Source: Simon Tatham's quasi-blog | Published: 2025-03-05T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/git-no-forge/>

My personal projects are in Git, but not on a ?forge? style website. I talk about why, and how to interact with them.

### 234. CSRF Protection without Tokens or Hidden Form Fields

Source: Miguel Grinberg's Blog | Published: Sun, 21 Dec 2025 15:54:28 GMT

Link: <https://blog.miguelgrinberg.com/post/csrf-protection-without-tokens-or-hidden-form-fields>

A couple of months ago, I received a request from a random Internet user to add CSRF protection to my little web framework Microdot , and I thought it was a fantastic idea. When I set off to do this work in early November I expected I was going to have to deal with anti-CSRF tokens, double-submit cookies and hidden form fields, pretty much the traditional elements that we have used to build a defense against CSRF for years. And I did start along this tedious route. But then I bumped into a new way some people are dealing with CSRF attacks that is way simpler, which I describe below.

### 235. How to Securely Store Secrets in Environment Variables

Source: Miguel Grinberg's Blog | Published: Sun, 14 Dec 2025 15:33:33 GMT

Link:

<https://blog.miguelgrinberg.com/post/how-to-securely-store-secrets-in-environment-variables>

You may have seen the recent reports of a malware that stole API keys, tokens and other secrets from a large number of developers. From where were these secrets stolen from? You guessed it, they were mostly stolen from environment variables. We use environment variables to configure information that processes need to run, but this type of storage was not designed for security, so using the environment for secrets always comes with risk. Given how serious this recent attack was, I thought it would be good to write a short article describing how I manage my secrets as part of my open source work.

### 236. Python 3.14 Is Here. How Fast Is It?

Source: Miguel Grinberg's Blog | Published: Wed, 08 Oct 2025 13:09:36 GMT

Link: <https://blog.miguelgrinberg.com/post/python-3-14-is-here-how-fast-is-it>

In November of 2024 I wrote a blog post titled "Is Python Really That Slow?", in which I tested several versions of Python and noted the steady progress the language has been making in terms of performance. Today is the 8th of October 2025, just a day after the official release of Python 3.14. Let's rerun the benchmarks to find out how fast the new version of Python is!

### 237. Dashboard totale Amerikaanse afhankelijkheden

Source: Bert Hubert's writings | Published: Tue, 20 Jan 2026 08:40:09 +0100

Link: <https://berthub.eu/articles/posts/dashboard-amerikaanse-afhankelijkheden/>

Hier een lijst organisaties die voor hun primaire missie totaal afhankelijk zijn van Amerikaanse clouds, en waarbij de maatschappij ontwricht zou worden als ze hun werk niet meer konden doen, bijvoorbeeld door sancties over Groenland. Dit omvat bijvoorbeeld wel de veiligheidsregio's en het UWV, maar niet het ministerie van economische zaken. Want als bij het departement EZ het licht uitgaat is dat vervelend, maar we kunnen wel een paar weken zonder nieuw beleid.

### 238. Europe's executives need to skill up to solve our total US cloud dependency

Source: Bert Hubert's writings | Published: Sat, 10 Jan 2026 11:08:02 +0100

Link: <https://berthub.eu/articles/posts/ft-on-european-cloud/>

Europe is experiencing a crisis of digital autonomy. Our dependence on US big tech has been growing for decades and is now nearly total, at a time when worries about our former ally are no longer theoretical. Might we, like the International Criminal Court in The Hague, find ourselves locked out of our own mailboxes if we say something that is upsetting to the US government? This post was written in response to an article in the Financial Times.

### 239. AWS and Microsoft are selling much more than cloud services

Source: Bert Hubert's writings | Published: Wed, 07 Jan 2026 13:00:00 +0100

Link: <https://berthub.eu/articles/posts/aws-and-microsoft-are-selling-much-more-than-cloud/>

Ten, twenty years ago customers over at large corporations would often ask me what hardware they'd need to run my software. As software developers, we'd tend to be rather blasé about such questions: whatever you have lying around!?. If pressed, we'd say you'd need this much RAM and that much storage. We thought we were doing our users a tremendous favor by being so flexible. Yet, customers were very much not happy with our responses.

### 240. Flock and Urban Surveillance

Source: computers are bad | Published: 26 Dec 2025 00:00:00 UT

Link: <https://computer.rip/2025-12-26-Flock-and-Urban-Surveillance.html>

Some years ago, I had a frustrating and largely fruitless encounter with the politics of policing. As a member of an oversight commission, I was particularly interested in the regulation of urban surveillance. The Albuquerque Police Department, for reasons good and bad, has often been an early adopter of surveillance technology. APD deployed automated license plate readers, mounted on patrol cars and portable trailers, in 2013. Initially, the department kept a six-month history of license plate data. For six months, police could retrospectively search the database to reconstruct a vehicle, or person's, movements?at least, those movements that happened near select patrol cars and "your speed is" trailers. Lobbying by the American Civil Liberties Union and public pressure on APD and city council lead to a policy change to retain data for only 14 days, a privacy-preserving measure that the ACLU lauded as one of the best ALPR policies in the nation. Today, ALPR is far more common in Albuquerque. Lowering costs and a continuing appetite for solving social problems with surveillance technology means that some parts of the city have ALPR installed at every signalized intersection?every p...

### 241. speed reading (the meaning of language)

Source: computers are bad | Published: 08 Dec 2025 00:00:00 UT

Link: <https://computer.rip/2025-12-08-speed-reading.html>

One of the difficult things about describing a grift, or at least what became a grift, is judging the sincerity with which the whole thing started. Scams often crystallize around a kernel of truth: genuinely good intentions that start rolling down the hill to profitability and end up crashing through every solid object along the way. I'm not totally sure about Evelyn Wood; she seems to have had all the best in mind but still turned so quickly to hotel conference room seminars that I have trouble lending her the benefit of the doubt. Still, she

was a teacher, and I am inclined to be sympathetic to teachers. Funny, then, that Wood's journey to fame started with another teacher. His curious reading behavior, whether interpreted as intense attention or half-assed inattention, set into motion one of the mid-century's greatest and, perhaps, most embarrassing executive self-help sensations. In 1929, Evelyn Wood earned a bachelor's in English at the University of Utah. The following two decades are a bit obscure; she took various high-school jobs around Utah leading ultimately to Salt Lake City's Jordan High School. There, as a counselor to girl students, Wood found that many students str...

## 242. RuBee

Source: computers are bad | Published: 22 Nov 2025 00:00:00 UT

Link: <https://computer.rip/2025-11-22-RuBee.html>

I have at least a few readers for which the sound of a man's voice saying "government cell phone detected" will elicit a palpable reaction. In Department of Energy facilities across the country, incidences of employees accidentally carrying phones into secure areas are reduced through a sort of automated nagging. A device at the door monitors for the presence of a tag; when the tag is detected it plays an audio clip. Because this is the government, the device in question is highly specialized, fantastically expensive, and says "government cell phone" even though most of the phones in question are personal devices. Look, they already did the recording, they're not changing it now! One of the things that I love is weird little wireless networks. Long ago I wrote about ANT+ , for example, a failed personal area network standard designed mostly around fitness applications. There's tons of these, and they have a lot of similarities---so it's fun to think about the protocols that went down a completely different path. It's even better, of course, if the protocol is obscure outside of an important niche. And a terrible website, too? What more could I ask for. The DoE's cell-phone nagging...

## 243. My impressions of the MacBook Pro M4

Source: Michael Stapelbergs Website | Published: 2025-10-31T11:04:59+01:00

Link: <https://michael.stapelberg.ch/posts/2025-10-31-macbook-pro-m4-impressions/>

I have been using a MacBook Pro M4 as my portable computer for the last half a year and wanted to share a few short impressions. As always, I am not a professional laptop reviewer, so in this article you won't find benchmarks, just subjective thoughts! Back in 2021, I wrote about the MacBook Air M1 , which was the first computer I used that contained Apple's own ARM-based CPU. Having a silent laptop with long battery life was a game-changer, so I wanted to keep those properties. When the US government announced tariffs, I figured I would replace my 4-year old MacBook Air M1 with a more recent model that should last a few more years. Ultimately, Apple's prices remained stable, so, in retrospect, I could have stayed with the M1 for a few more years. Oh well. The nano-textured display I went to the Apple Store to compare the different options in person. Specifically, I was curious about the display and whether the increased weight and form factor of the MacBook Pro (compared to a MacBook Air) would be acceptable. Another downside of the Pro model is that it comes with a fan, and I really like absolutely quiet computers. Online, I read from other MacBook Pro owners that the fan mostly...

## 244. NixCon 2025 Trip Report ?

Source: Michael Stapelbergs Website | Published: 2025-09-21T09:34:00+02:00

Link: <https://michael.stapelberg.ch/posts/2025-09-21-nixcon-2025-trip-report/>

I liked the NixOS meetup earlier this year, and at the end of the meetup they told everyone about NixCon 2025, which would be happening in Switzerland this year, at the very same location, the University Of Applied Sciences OST in Rapperswil, so I decided to go! In this trip report, I want to give you a rough impression of how I experienced this awesome conference :) The bee in the title is a NixCon inside joke ;) Friday I arrived at about 09:30 on a rainy Friday morning, meaning I hurried from the train station into OST building 1 to show my ticket QR code and pick up my conference badge and custom name tag that I pre-ordered. The custom ones have your name engraved and come with a strong magnet to attach them to your clothes: After grabbing a bite to eat, I headed to the main lecture hall for the opening session. Prof. Dr.

Farhad Mehta from OST, as well as the entire NixCon orga team, welcomed the 450 registered attendees to the 10th NixCon! I recognized many familiar faces from the Nix meetup, but many hands went up when the audience was asked for whom it was the first time at NixCon, or in Switzerland in general. I want to thank Prof. Mehta in particular for making possible su...

#### 245. Bye Intel, hi AMD! I'm done after 2 dead Intels

Source: Michael Stapelbergs Website | Published: 2025-09-07T08:33:00+02:00

Link: <https://michael.stapelberg.ch/posts/2025-09-07-bye-intel-hi-amd-9950x3d/>

The Intel 285K CPU in my high-end 2025 Linux PC died again ! ? Notably, this was the replacement CPU for the original 285K that died in March , and after reading through the reviews of Intel CPUs on my electronics store of choice, many of which (!) mention CPU replacements, I am getting the impression that Intel's current CPUs just are not stable ?.

Therefore, I am giving up on Intel for the coming years and have bought an AMD Ryzen 9950X3D CPU instead. What happened? Or: the batch job of death On the 9th of July, I set out to experiment with layout-parser and tesseract in order to convert a collection of scanned paper documents from images into text. I expected that offloading this task to the GPU would result in a drastic speed-up, so I attempted to build layout-parser with CUDA . Usually, it's not required to compile software yourself on NixOS , but CUDA is non-free, so the default NixOS cache does not compile software with CUDA. (Tip: Enable the Nix Community Cache , which contains prebuilt CUDA packages, too!) This lengthy compilation attempt failed with a weird symptom: I left for work, and after a while, my PC was no longer reachable over the network, but fans kept spinning...

#### 246. Investigating a forged PDF

Source: Matthew Garrett | Published: Wed, 24 Sep 2025 22:22:34 GMT

Link: <https://mjj59.dreamwidth.org/73317.html>

I had to rent a house for a couple of months recently, which is long enough in California that it pushes you into proper tenant protection law. As landlords tend to do, they failed to return my security deposit within the 21 days required by law , having already failed to provide the required notification that I was entitled to an inspection before moving out. Cue some tedious argumentation with the letting agency, and eventually me threatening to take them to small claims court. This post is not about that. Now, under Californian law, the onus is on the landlord to hold and return the security deposit - the agency has no role in this. The only reason I was talking to them is that my lease didn't mention the name or address of the landlord (another legal violation , but the outcome is just that you get to serve the landlord via the agency). So it was a bit surprising when I received an email from the owner of the agency informing me that they did not hold the deposit and so were not liable - I already knew this. The odd bit about this, though, is that they sent me another copy of the contract, asserting that it made it clear that the landlord held the deposit. I read it, and inste...

#### 247. Cordoomceps - replacing an Amiga's brain with Doom

Source: Matthew Garrett | Published: Tue, 05 Aug 2025 00:30:19 GMT

Link: <https://mjj59.dreamwidth.org/73001.html>

There's a lovely device called a pistorm , an adapter board that glues a Raspberry Pi GPIO bus to a Motorola 68000 bus. The intended use case is that you plug it into a 68000 device and then run an emulator that reads instructions from hardware (ROM or RAM) and emulates them. You're still limited by the ~7MHz bus that the hardware is running at, but you can run the instructions as fast as you want. These days you're supposed to run a custom built OS on the Pi that just does 68000 emulation, but initially it ran Linux on the Pi and a userland 68000 emulator process. And, well, that got me thinking. The emulator takes 68000 instructions, emulates them, and then talks to the hardware to implement the effects of those instructions. What if we, well, just don't? What if we just run all of our code in Linux on an ARM core and then talk to the Amiga hardware? We're going to ignore x86 here, because it's weird - but most hardware that wants software to be able to communicate with it maps itself into the same address space that RAM is in. You can write to a byte of RAM, or you can write to a piece of

hardware that's effectively pretending to be RAM[1]. The Amiga wasn't unusual in this resp...

## 248. Secure boot certificate rollover is real but probably won't hurt you

Source: Matthew Garrett | Published: Thu, 31 Jul 2025 16:12:59 GMT

Link: <https://mjg59.dreamwidth.org/72892.html>

LWN wrote an article which opens with the assertion "Linux users who have Secure Boot enabled on their systems knowingly or unknowingly rely on a key from Microsoft that is set to expire in September". This is, depending on interpretation, either misleading or just plain wrong, but also there's not a good source of truth here, so. First, how does secure boot signing work? Every system that supports UEFI secure boot ships with a set of trusted certificates in a database called "db". Any binary signed with a chain of certificates that chains to a root in db is trusted, unless either the binary (via hash) or an intermediate certificate is added to "dbx", a separate database of things whose trust has been revoked[1]. But, in general, the firmware doesn't care about the intermediate or the number of intermediates or whatever - as long as there's a valid chain back to a certificate that's in db, it's going to be happy.

That's the conceptual version. What about the real world one? Most x86 systems that implement UEFI secure boot have at least two root certificates in db - one called "Microsoft Windows Production PCA 2011", and one called "Microsoft Corporation UEFI CA 2011". The former i...

## 249. Why the chicken crossed the road, according to various entities

Source: DYNOMIGHT | Published: 2025-12-04T00:00:00+00:00

Link: <https://dynamight.net/chicken/>

When I started this blog, I promised myself that I would always steer into weirdness. (As they say, ?Get busy being weird, or get busy dying.?) While time has shown there are limits to what y'all will tolerate [ 1 2 3 4 ] I still sometimes feel a need to publish something that's pure exuberant stupidity. Thus, I present: WHY DID THE CHICKEN CROSS THE ROAD ACCORDING TO VARIOUS PEOPLE OR OTHER ENTITIES Q) Why did the chicken cross the road? A) The chicken ain?t fussy. Everybody gotta be somewhere. The chicken been on this side a long time and never suffered none for it. The chicken don?t see no obvious benefit to the other side. But the talk of the town is nothing but crossing, and the chicken can?t help but go see what got everyone so stirred up. (Mark Twain) Q) Why did the chicken cross the road? A) The outcome would be best if no one crossed. However, if other chickens do cross, then the outcome would be better if this chicken also crossed. The chicken rejects the Kantian universalism. So the chicken crosses. (Derek Parfit) Q) Why did the chicken cross the road? A) You were a beautiful little chick The whole world was before you You greased your wattles and crossed the road Sure...

## 250. Underrated reasons to be thankful V

Source: DYNOMIGHT | Published: 2025-11-27T00:00:00+00:00

Link: <https://dynamight.net/thanks-5/>

That your dog, while she appears to love you only because she?s been adapted by evolution to appear to love you, really does love you. That if you?re a life form and you cook up a baby and copy your genes to them, you?ll find that the genes have been degraded due to oxidative stress et al., which isn?t cause for celebration, but if you find some other hopefully-hot person and randomly swap in half of their genes, your baby will still be somewhat less fit compared to you and your hopefully-hot friend on average, but now there is variance, so if you cook up several babies, one of them might be as fit or even fitter than you, and that one will likely have more babies than your other babies have, and thus complex life can persist in a universe with increasing entropy. That if we wanted to, we surely could figure out which of the 300-ish strains of rhinovirus are circulating in a given area at a given time and rapidly vaccinate people to stop it and thereby finally ?cure? the common cold, and though this is too annoying to pursue right now, it seems like it?s just a matter of time. That if you look back at history, you see that plagues went from Europe to the Americas but not the other...

## 251. Make product worse, get money

Source: DYNOMIGHT | Published: 2025-11-20T00:00:00+00:00

Link: <https://dynamight.net/worse/>

I recently asked why people seem to hate dating apps so much. In response, 80% of you emailed me some version of the following theory: The thing about dating apps is that if they do a good job and match people up, then the matched people will quit the app and stop paying. So they have an incentive to string people along but not to actually help people find long-term relationships. May I explain why I don't find this type of theory very helpful? I'm not saying that I think it's wrong, mind you. Rather, my objection is that while the theory is phrased in terms of dating apps, the same basic pattern applies to basically anyone who is trying to make money by doing anything. For example, consider a pizza restaurant. Try these theories on for size: Pizza: ?The thing about pizza restaurants is that if they use expensive ingredients or labor-intensive pizza-making techniques, then it costs more to make pizza. So they have an incentive to use low-cost ingredients and labor-saving shortcuts.? Pizza II: ?The thing about pizza restaurants is that if they have nice tables separated at a comfortable distance, then they can't fit as many customers. So they have an incentive to use tiny tables an...