

1. Anthropic's public benefit mission

Source: Simon Willison's Weblog | Published: 2026-02-13T23:59:51+00:00

Link: <https://simonwillison.net/2026/Feb/13/anthropic-public-benefit-mission/#atom-everything>

Someone asked if there was an Anthropic equivalent to OpenAI's IRS mission statements over time . Anthropic are a "public benefit corporation" but not a non-profit, so they don't have the same requirements to file public documents with the IRS every year. But when I asked Claude it ran a search and dug up this Google Drive folder where Zach Stein-Perlman shared Certificate of Incorporation documents he obtained from the State of Delaware ! Anthropic's are much less interesting than OpenAI's. The earliest document from 2021 states: The specific public benefit that the Corporation will promote is to responsibly develop and maintain advanced AI for the cultural, social and technological improvement of humanity. Every subsequent document up to 2024 uses an updated version which says: The specific public benefit that the Corporation will promote is to responsibly develop and maintain advanced AI for the long term benefit of humanity. Tags: ai-ethics , anthropic , ai

2. The evolution of OpenAI's mission statement

Source: Simon Willison's Weblog | Published: 2026-02-13T23:38:29+00:00

Link: <https://simonwillison.net/2026/Feb/13/openai-mission-statement/#atom-everything>

As a USA 501(c)(3) the OpenAI non-profit has to file a tax return each year with the IRS. One of the required fields on that tax return is to "Briefly describe the organization's mission or most significant activities" - this has actual legal weight to it as the IRS can use it to evaluate if the organization is sticking to its mission and deserves to maintain its non-profit tax-exempt status. You can browse OpenAI's tax filings by year on ProPublica's excellent Nonprofit Explorer . I went through and extracted that mission statement for 2016 through 2024, then had Claude help me fake the commit dates to turn it into a git repository and share that as a Gist - which means that Gist's revisions page shows every edit they've made since they started filing their taxes! It's really interesting seeing what they've changed over time. The original 2016 mission reads as follows (and yes, the apostrophe in "OpenAIs" is missing in the original): OpenAIs goal is to advance digital intelligence in the way that is most likely to benefit humanity as a whole, unconstrained by a need to generate financial return. We think that artificial intelligence technology will help shape the 21st centu...

3. Gemini 3 Deep Think

Source: Simon Willison's Weblog | Published: 2026-02-12T18:12:17+00:00

Link: <https://simonwillison.net/2026/Feb/12/gemini-3-deep-think/#atom-everything>

Gemini 3 Deep Think New from Google. They say it's "built to push the frontier of intelligence and solve modern challenges across science, research, and engineering". It drew me a really good SVG of a pelican riding a bicycle ! I think this is the best one I've seen so far - here's my previous collection . (And since it's an FAQ, here's my answer to What happens if AI labs train for pelicans riding bicycles?) Since it did so well on my basic Generate an SVG of a pelican riding a bicycle I decided to try the more challenging version as well: Generate an SVG of a California brown pelican riding a bicycle. The bicycle must have spokes and a correctly shaped bicycle frame. The pelican must have its characteristic large pouch, and there should be a clear indication of feathers. The pelican must be clearly pedaling the bicycle. The image should show the full breeding plumage of the California brown pelican. Here's what I got : Via Hacker News Tags: google , ai , generative-ai , llms , gemini , pelican-riding-a-bicycle , llm-reasoning , llm-release

4. Testing Reachy Mini - Hugging Face's Pi powered robot

Source: Jeff Geerling | Published: Fri, 13 Feb 2026 09:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/testing-reachy-mini-hugging-face-robot/>

When I saw Jensen Huang introduce the Reachy Mini at CES , I thought it was a gimmick. His

keynote showed this little robot responding to human input, turning its head to look at a TODO list on the wall, sending emails, and turning drawings into architectural renderings with motion. HuggingFace and Pollen robotics sent me a Reachy Mini to test, and, well, at least if you're looking to replicate that setup in the keynote, it's not, as Jensen put it, "utterly trivial now."

5. JeffGeerling.com has been Migrated to Hugo

Source: Jeff Geerling | Published: Sat, 03 Jan 2026 13:00:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/migrated-to-hugo/>

Since 2009, this website has run on Drupal . Starting with Drupal 6, and progressing through major site upgrades and migrations to 7, 8, 9, and 10, I used the site as a way to dogfood the same CMS (Content Management System) I used in my day job for over a decade . But as time progressed?especially after completing a grueling upgrade from Drupal 7 to 8 ?my enthusiasm for maintaining what's now a more enterprise-focused Digital Experience Platform or 'DXP' for a personal blog has waned.

6. Testing the Mono Gateway, a custom-built 10 Gbps Router

Source: Jeff Geerling | Published: Fri, 02 Jan 2026 09:12:25 -0600

Link: <https://www.jeffgeerling.com/blog/2026/testing-mono-gateway-custom-built-10-gbps-router/>

Last month, the stars aligned for me to bring the Mono Gateway (a 10 Gbps router that YouTuber Toma? Zaman and his team at Mono built from scratch) on a trip to Phoenix, and test it with one of the most OP network test boxes I've ever seen, at the ServeTheHome HQ. In this video, Patrick (from STH) and I put Gateway through a real-world torture test using CyPerf:
.embed-container { position: relative; padding-bottom: 56.25%; height: 0; overflow: hidden; max-width: 100%; } .embed-container iframe, .embed-container object, .embed-container embed { position: absolute; top: 0; left: 0; width: 100%; height: 100%; } Mono Gateway The Gateway is a small (1U/desktop) router with features like a 4-core Arm CPU, 8GB of ECC LPDDR RAM, and decent expansion options for a compact router.

7. Quick P2D production update

Source: Eric Migicovsky's Blog RSS Feed | Published: Mon, 29 Sep 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/quick-p2d-production-update>

(No summary)

8. Pebble 2 Duo is in mass production!

Source: Eric Migicovsky's Blog RSS Feed | Published: Fri, 22 Aug 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/pebble-2-duo-is-in-mass-production>

TLDR (Oct 1 Production Update) P2D is in production!! Waterproof update - 20m water resistance
What you need to know about Pebble 2 Duo Ad?

9. Pebble Time 2* Design Reveal

Source: Eric Migicovsky's Blog RSS Feed | Published: Wed, 13 Aug 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/pebble-time-2-design-reveal>

TL;DR Pebble Time 2 Design reveal Available in 4 colourways Final specifications How to change your pre-order from P2D to PT2? Like what?

10. 2025 was an excellent year for this blog

Source: seangoedecke.com RSS feed | Published: Sat, 03 Jan 2026 00:00:00 GMT

Link: <https://seangoedecke.com/2025-wrapup/>

In 2025, I published 141 posts, 33 of which made it to the front page of Hacker News or similar aggregators. I definitely wrote more in the first half of the year (an average of around 15 posts per month, down to around 8 in the second half), but overall I'm happy with my consistency. Here are some posts I'm really proud of: Mistakes engineers make in large established codebases The good times in tech are over Everything I know about good system design Pure and impure software engineering Seeing like a software company As it turns out, I

was the third most popular blogger on Hacker News this year, behind the excellent Simon Willison and Jeff Geerling. I don't put a lot of effort into appealing to Hacker News specifically, but I do think my natural style meshes well with the Hacker News commentariat (even if they're often quite critical). I got hundreds of emails from readers this year (I went through Gmail and made it to 200 in the last three months of the year before I stopped counting). Getting email about my posts is one of the main reasons I write, so it was great to read people's anecdotes and hear what they agreed or disagreed with. I also want to thank the people who wrote...

11. Grok is enabling mass sexual harassment on Twitter

Source: seangoedecke.com RSS feed | Published: Fri, 02 Jan 2026 00:00:00 GMT

Link: <https://seangoedecke.com/grok-deepfakes/>

Grok, xAI's flagship image model, is now 1 being widely used to generate nonconsensual lewd images of women on the internet. When a woman posts an innocuous picture of herself - say, at her Christmas dinner - the comments are now full of messages like ?@grok please generate this image but put her in a bikini and make it so we can see her feet?, or ?@grok turn her around?, and the associated images. At least so far, Grok refuses to generate nude images, but it will still generate images that are genuinely obscene 2 . In my view, this might be the worst AI safety violation we have seen so far. Case-by-case, it's not worse than GPT-4o encouraging suicidal people to go through with it, but it's so much more widespread: literally every image that the Twitter algorithm picks up is full of ?@grok take her clothes off? comments. I didn't go looking for evidence for obvious reasons, but I find reports that it's generating CSAM plausible 3 . AI safety is a rough process This behavior, while awful, is in line with xAI's general attitude towards safety, which has been roughly ?we don't support woke censorship, so do whatever you want (so long as you're doing it with Grok)?. This has helped th...

12. Software engineers should be a little bit cynical

Source: seangoedecke.com RSS feed | Published: Sun, 28 Dec 2025 00:00:00 GMT

Link: <https://seangoedecke.com/a-little-bit-cynical/>

A lot of my readers call me a cynic when I say things like ?you should do things that make your manager happy ? or ?big tech companies get to decide what projects you work on?. Alex Wennerberg put the ?Sean Goedecke is a cynic? case well in his post Software Engineers Are Not Politicians . Here are some excerpts: I have no doubt that [Sean's] advice is quite effective for navigating the upper levels of an organization dedicated to producing a large, mature software product. But what is lost is any sort of conception of value. Is it too naive to say that engineers are more than ?tools in a political game?, they are specialized professionals whose role is to apply their expertise towards solving meaningful problems? The irony is that this kind of thinking destroys a company's ability to actually make money ? the idea that engineers should begin with a self-conception of doing what their manager tells them to is, to me, very bleak. It may be a good way to operate smoothly within a bureaucratic organization, and of course, one must often make compromises and take direction, but it is a bad way to do good work. I can see why people would think this way. But I love working in big tech c...

13. Let's Keep an Eye on Apple's Own iOS Adoption Numbers

Source: Daring Fireball | Published: 2026-01-29T20:18:48Z

Link: <https://developer.apple.com/support/app-store/>

When I wrote last week about the false narrative that iOS 26 is seeing bizarrely low adoption rates compared to previous years, I neglected one source: Apple itself. Apple's Developer site publishes a page with iOS and iPadOS usage for devices that ?transacted on the App Store? . The hitch is that they only seem to update those numbers twice a year ? once right around now, and once again right before WWDC. As of today, those numbers are still from 4 June 2025. Last year, going from the Internet Archive, the numbers were still from iOS 17 (June 2024) on 23 January last year, but were updated for iOS 18 on 24 January . Here are those iOS 18 numbers from one year ago this week. iPhones released in the previous four years: iOS 18: 76% iOS 17: 19% iOS All iPhones: iOS 18: 68% iOS 17: 19% iOS iPads released in the previous four years: iPadOS 18:

63% iPadOS 17: 27% iPadOS All iPads: iPadOS 18: 53% iPadOS 17: 28% iPadOS (Apple itself manages to present these statistics without ever using the plurals iPhones or iPads , instead referring only to ?devices?.) A year prior in early 2024, Apple updated the numbers at some point between 23 January and 6 February . I presume, or at least hope, th...

14. Box Office Expectations for ?Melania?

Source: Daring Fireball | Published: 2026-01-29T15:22:23Z

Link: <https://www.thewrap.com/industry-news/box-office/melania-box-office-prediction/>
Jeremy Fuster, reporting for TheWrap: But save for some theaters in Republican-heavy states, the film is unlikely to leave much of an impact at a slumping box office, with theatrical sources telling TheWrap that ?Melania? is projected for an opening of around \$3 million this weekend. That would put it below the last right-wing documentary, the Daily Wire-produced Matt Walsh film ?Am I Racist?,? which opened to \$4.5 million from 1,517 locations in September 2024, finishing with a \$12.3 million total that made it the highest-grossing doc that year. The highest projections are coming from NRG with an estimate of around \$5 million, though audience interest polls from the company have 30% saying they are ?definitely not? interested in watching the film, an unusually high count for any wide release. These projections are with a \$35 million promotional campaign, for a movie Amazon paid \$40 million to purchase. (Via Taegan Goddard .) ?

15. Amazon?s Spending on ?Melania? Is a Barely Concealed Bribe

Source: Daring Fireball | Published: 2026-01-29T15:19:17Z

Link: https://www.nytimes.com/2026/01/28/business/media/amazon-melania-trump-film-critics.html?unlocked_article_code=1.IFA.R85v.5ZtNe8FttfaD
Nicole Sperling and Brooks Barnes, reporting for The New York Times: Amazon paid Ms. Trump?s production company \$40 million for the rights to ?Melania,? about \$26 million more than the next closest bidder, Disney. The fee includes a related docuseries that is scheduled to air later this year. The budget for ?Melania? is unknown, but documentaries that follow a subject for a limited amount of time usually cost less than \$5 million to produce. The \$35 million for marketing is 10 times what some other high-profile documentaries have received. All of which has a lot of Hollywood questioning whether Amazon?s push is anything more than the company?s attempt to ingratiate itself with President Trump. This is a good story, with multiple industry sources with experience making political documentaries, but the Times?s own subhead downplays Amazon?s spending on the film: ?The tech giant is spending \$35 million to promote its film about the first lady, far more than is typical for documentaries.? They?re spending \$35 million now, to promote it, but they already paid \$40 million for the rights to the film, \$28 million of which is believed to have gone to Melania Trump herself . A \$35 million t...

16. Dismantling Defenses: Trump 2.0 Cyber Year in Review

Source: Krebs on Security | Published: Fri, 19 Dec 2025 15:14:55 +0000

Link: <https://krebsonsecurity.com/2025/12/dismantling-defenses-trump-2-0-cyber-year-in-review/>
The Trump administration has pursued a staggering range of policy pivots this past year that threaten to weaken the nation?s ability and willingness to address a broad spectrum of technology challenges, from cybersecurity and privacy to countering disinformation, fraud and corruption. These shifts, along with the president?s efforts to restrict free speech and freedom of the press, have come at such a rapid clip that many readers probably aren?t even aware of them all.

17. Most Parked Domains Now Serving Malicious Content

Source: Krebs on Security | Published: Tue, 16 Dec 2025 14:14:48 +0000

Link: <https://krebsonsecurity.com/2025/12/most-parked-domains-now-serving-malicious-content/>
Direct navigation -- the act of visiting a website by manually typing a domain name in a web browser -- has never been riskier: A new study finds the vast majority of "parked" domains -- mostly expired or dormant domain names, or common misspellings of popular websites -- are now configured to redirect visitors to sites that foist scams and malware.

18. Microsoft Patch Tuesday, December 2025 Edition

Source: Krebs on Security | Published: Tue, 09 Dec 2025 23:18:29 +0000

Link: <https://krebsonsecurity.com/2025/12/microsoft-patch-tuesday-december-2025-edition/>

Microsoft today pushed updates to fix at least 56 security flaws in its Windows operating systems and supported software. This final Patch Tuesday of 2025 tackles one zero-day bug that is already being exploited, as well as two publicly disclosed vulnerabilities.

19. How can I distinguish between the numeric keypad 0 and the top-row 0 in the WM_CHAR message?

Source: The Old New Thing | Published: Fri, 13 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260213-00/?p=112062>

See if it matches the scan code. The post How can I distinguish between the numeric keypad 0 and the top-row 0 in the WM_CHAR message? appeared first on The Old New Thing .

20. We Rewrote the Ghostty GTK Application

Source: Mitchell Hashimoto | Published: Thu, 14 Aug 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-gtk-rewrite>

(No summary)

21. "As Code"

Source: Mitchell Hashimoto | Published: Tue, 04 Mar 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/as-code>

(No summary)

22. Welcoming Ghostty Subsystem Maintainers

Source: Mitchell Hashimoto | Published: Fri, 07 Feb 2025 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-subsystem-maintainers>

(No summary)

23. The Final Bottleneck

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-02-13T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/2/13/the-final-bottleneck/>

Historically, writing code was slower than reviewing code. It might not have felt that way, because code reviews sat in queues until someone got around to picking it up. But if you compare the actual acts themselves, creation was usually the more expensive part. In teams where people both wrote and reviewed code, it never felt like ?we should probably program slower.? So when more and more people tell me they no longer know what code is in their own codebase, I feel like something is very wrong here and it?s time to reflect. You Are Here Software engineers often believe that if we make the bathtub bigger , overflow disappears. It doesn?t. OpenClaw right now has north of 2,500 pull requests open. That?s a big bathtub. Anyone who has worked with queues knows this: if input grows faster than throughput, you have an accumulating failure. At that point, backpressure and load shedding are the only things that retain a system that can still operate. If you have ever been in a Starbucks overwhelmed by mobile orders, you know the feeling. The in-store experience breaks down. You no longer know how many orders are ahead of you. There is no clear line, no reliable wait estimate, and often no...

24. Suppressions of Suppressions

Source: overreacted ? A blog by Dan Abramov | Published: Wed, 11 Jun 2025 00:00:00 GMT

Link: <https://overreacted.io/suppressions-of-suppressions/>

I heard you like linting.

25. I'm Doing a Little Consulting

Source: overreacted ? A blog by Dan Abramov | Published: Wed, 11 Jun 2025 00:00:00 GMT

Link: <https://overreacted.io/im-doing-a-little-consulting/>

Personal update post.

26. How Imports Work in RSC

Source: overreacted ? A blog by Dan Abramov | Published: Thu, 05 Jun 2025 00:00:00 GMT

Link: <https://overreacted.io/how-imports-work-in-rsc/>

A layered module system.

27. Good riddance, 4o

Source: Mahad Kalam | Published: Fri, 13 Feb 2026 00:00:00 GMT

Link: <https://mahadk.com/posts/4o/>

It should've happened sooner.

28. Expressing a prime as the sum of two squares

Source: John D. Cook | Published: Fri, 13 Feb 2026 02:34:58 +0000

Link: <https://www.johndcook.com/blog/2026/02/12/pythagorean-primes/>

I saw where Elon Musk posted Grok?s answer to the prompt ?What are the most beautiful theorems.? I looked at the list, and there were no surprises, as you?d expect from a program that works by predicting the most likely sequence of words based on analyzing web pages.

There?s only one theorem on the list [?] The post Expressing a prime as the sum of two squares first appeared on John D. Cook .

29. AGI, ASI, A*I ? Do we have all we need to get there?

Source: John D. Cook | Published: Fri, 30 Jan 2026 19:46:37 +0000

Link: <https://www.johndcook.com/blog/2026/01/30/agi-asi-ai-do-we-have-all-we-need-to-get-there/>

Demis: ?[to get to AGI] maybe there?s one or two big innovations needed? Sam: ?everything based off what we see today is that it will happen.? Ilya: ?But is the belief really that if you just 100x the scale, everything would be transformed? I don?t think that?s true.? Dario: ?If you just kind of like [?] The post AGI, ASI, A*I ? Do we have all we need to get there? first appeared on John D. Cook .

30. Bridging secrets is hard

Source: John D. Cook | Published: Fri, 30 Jan 2026 17:09:30 +0000

Link: <https://www.johndcook.com/blog/2026/01/30/bridging-secrets/>

Cryptocurrency and privacy don?t fit together as easily as you might expect. Blockchains give you the illusion of privacy via pseudonymization: you don?t put your name on a blockchain, but you do put information on a blockchain that can be used to determine your name. Blockchain analysis can often reveal information that no one intended [?] The post Bridging secrets is hard first appeared on John D. Cook .

31. How to reverse engineer an analog chip: the TDA7000 FM radio receiver

Source: Ken Shirriff's blog | Published: 2025-08-02T08:10:00.000-07:00

Link: <http://www.righto.com/2025/08/reverse-engineering-analog-TDA7000.html>

Have you ever wanted to reverse engineer an analog chip from a die photo? Wanted to understand what's inside the "black box" of an integrated circuit? In this article, I explain my reverse engineering process, using the Philips TDA7000 FM radio receiver chip as an example. This chip was the first FM radio receiver on a chip. It was designed in 1977?an era of large transistors and a single layer of metal?so it is much easier to examine than modern chips.

Nonetheless, the TDA7000 is a non-trivial chip with over 100 transistors. It includes common analog circuits such as differential amplifiers and current mirrors, along with more obscure circuits such as Gilbert cell mixers. Die photo of the TDA7000 with the main functional blocks labeled. Click this image (or any other) for a larger version. Die photo from IEEE's Microchips that Shook the World exhibit page . The die photo above shows the silicon die of the TDA7000; I've labeled the main functional blocks and some interesting components. Arranged around the border of the chip are 18 bond pads: the pads are connected by thin gold bond wires to the pins

of the integrated circuit package. In this chip, the silicon appears greenish,...

32. Reverse engineering the mysterious Up-Data Link Test Set from Apollo

Source: Ken Shirriff's blog | Published: 2025-07-21T09:10:00.000-07:00

Link: <http://www.righto.com/2025/07/reverse-engineering-mysterious-up-data.html>

Back in 2021, a collector friend of ours was visiting a dusty warehouse in search of Apollo-era communications equipment. A box with NASA-style lights caught his eye?the "AGC Confirm" light suggested a connection with the Apollo Guidance Computer. Disappointingly, the box was just an empty chassis and the circuit boards were all missing. He continued to poke around the warehouse when, to his surprise, he found a bag on the other side of the warehouse that contained the missing boards! After reuniting the box with its wayward circuit cards, he brought it to us: could we make this undocumented unit work? The Up-Data Link Confidence Test Set, powered up. A label on the back indicated that it is an "Up-Data Link Confidence Test Set", built by Motorola. As the name suggests, the box was designed to test Apollo's Up-Data Link (UDL), a system that allowed digital commands to be sent up to the spacecraft. As I'll explain in detail below, these commands allowed ground stations to switch spacecraft circuits on or off, interact with the Apollo Guidance Computer, or set the spacecraft's clock. The Up-Data Link needed to be tested on the ground to ensure that its functions operated correctly....

33. Inside the Apollo "8-Ball" FDAI (Flight Director / Attitude Indicator)

Source: Ken Shirriff's blog | Published: 2025-06-13T20:12:00.000-07:00

Link: <http://www.righto.com/2025/06/inside-apollo-fdai.html>

During the Apollo flights to the Moon, the astronauts observed the spacecraft's orientation on a special instrument called the FDAI (Flight Director / Attitude Indicator). This instrument showed the spacecraft's attitude?its orientation?by rotating a ball. This ball was nicknamed the "8-ball" because it was black (albeit only on one side). The instrument also acted as a flight director, using three yellow needles to indicate how the astronauts should maneuver the spacecraft. Three more pointers showed how fast the spacecraft was rotating. An Apollo FDI (Flight Director/Attitude Indicator) with the case removed. This FDI is on its side to avoid crushing the needles. Since the spacecraft rotates along three axes (roll, pitch, and yaw), the ball also rotates along three axes. It's not obvious how the ball can rotate to an arbitrary orientation while remaining attached. In this article, I look inside an FDI from Apollo that was repurposed for a Space Shuttle simulator 1 and explain how it operates. (Spoiler: the ball mechanism is firmly attached at the "equator" and rotates in two axes. What you see is two hollow shells around the ball mechanism that spin around the third axis.) Th...

34. Factional Drift: We cluster into factions online

Source: iDiallo.com | Published: Fri, 13 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/factional-drift-online?src=feed>

Whenever one of my articles reaches some popularity, I tend not to participate in the discussion. A few weeks back, I told a story about me, my neighbor and an UHF remote . The story took on a life of its own on Hackernews before I could answer any questions. But reading through the comment section, I noticed a pattern on how comments form. People were not necessarily talking about my article. They had turned into factions. This isn't a complaint about the community. Instead it's an observation that I've made many years ago but didn't have the words to describe it. Now I have the articles to explore the idea. The article asked this question: is it okay to use a shared RF remote to silence a loud neighbor ? The comment section on hackernews split into two teams. Team Justice, who believed I was right to teach my neighbor a lesson. And then Team Boundaries, who believed I was ?a real dick?. But within hours, the thread stopped being about that question. People self-sorted into tribes, not by opinion on the neighbor, but by identity. The tinkerers joined the conversation. If you only looked through the comment section without reading the article, you'd think it was a DIY thread on ho...

35. How to Preserve Your Writing for a Hundred Years

Source: iDiallo.com | Published: Fri, 23 Jan 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/preserve-your-writing-for-a-hundred-years?src=feed>

There was a question on Hacker News where a user asked how he could ensure his writing would endure for a hundred years . At first, I treated it as a technology problem. Storage, formats, domains, backups. If the goal is durability, then the best technology we've invented so far is still paper. Print it. Put it on a shelf. Problem solved. But that answer was too neat and it reminded me of a story a friend once told me. He found an old book in his basement, more than a hundred years old. The paper had survived, the binding held, the ink was still legible. He was excited. This felt like discovering a voice from the past, a time capsule. But when he read it, the book was... bad. The writing was unremarkable. The ideas went nowhere. The author wasn't well known, and no trace of him existed online. Eventually, the book ended up back in a box and was donated along with other used items. Here was a book that had solved the storage problem perfectly, and yet it had failed at endurance. It survived a hundred years only to disappear again, this time for good. The original question was framed wrong. Endurance isn't a storage problem. Instead of thinking about why or how you should preserve y...

36. I can't upgrade to Windows 11, now leave me alone

Source: iDiallo.com | Published: Sun, 21 Dec 2025 02:06:13 GMT

Link: <https://idiallo.com/byte-size/cant-update-to-windows-11-leave-me-alone?src=feed>

Microsoft won't let you dismiss the upgrade notification So support for Windows 10 has ended. Yes, millions of users are still on it. One of my main laptops runs Windows 10. I can't update to Windows 11 because of the hardware requirements. It's not that I don't have enough RAM, storage, or CPU power. The hardware limitation is specifically TPM 2.0. What is TPM 2.0, you say? It stands for Trusted Platform Module. It's basically a security chip on the motherboard that enables some security features. It's good and all, but Windows says my laptop doesn't support it. Great! Now leave me alone. Well, every time I turn on my computer, I get a reminder that I need to update to Windows 11. OK, at this point a Windows machine only belongs to you in name. Microsoft can run arbitrary code on it. They already ran the code to decide that my computer doesn't support Windows 11. So why do they keep bothering me? Fine, I'm frustrated. That's why I'm complaining. I've accepted the fact that my powerful, yet 10-year-old laptop won't get the latest update. But if Microsoft's own systems have determined my hardware is incompatible, why are they harassing? I'll just have to dismiss this notification a...

37. Y'all are over-complicating these AI-risk arguments

Source: DYNOMIGHT | Published: 2025-10-02T00:00:00+00:00

Link: <https://dynamight.net/ai-risk/>

Say an alien spaceship is headed for Earth. It has 30 aliens on it. The aliens are weak and small. They have no weapons and carry no diseases. They breed at rates similar to humans. They are bringing no new technology. No other ships are coming. There's no trick?except that they each have an IQ of 300. Would you find that concerning? Of course, the aliens might be great. They might cure cancer and help us reach world peace and higher consciousness. But would you be sure they'd be great? Suppose you were worried about the aliens but I scoffed, ?Tell me specifically how the aliens would hurt us. They're small and weak! They can't do anything unless we let them.? Would you find that counter-argument convincing? I claim that most people would be concerned about the arrival of the aliens, would not be sure that their arrival would be good, and would not find that counter-argument convincing. I bring this up because most AI-risk arguments I see go something like this: There will be a fast takeoff in AI capabilities. Due to alignment difficulty and orthogonality , it will pursue dangerous convergent subgoals . These will give the AI a decisive strategic advantage , making it uncontainabl...

38. Shoes, Algernon, Pangea, and Sea Peoples

Source: DYNOMIGHT | Published: 2025-09-25T00:00:00+00:00

Link: <https://dynamight.net/shorts-5/>

I fear we are in the waning days of the People Read Blog Posts About Random Well-Understood

Topics Instead of Asking Their Automatons Era. So before I lose my chance, here is a blog post about some random well-understood topics. Marathons are stupidly fast You probably know that people can now run marathons in just over 2 hours. But do you realize how insane that is? That's an average speed of 21.1 km per hour, or 13.1 miles per hour. You can think of that as running a mile in 4:35 (world record: 3:45), except doing it 26.2 times in a row. Or, you can think of that as running 100 meters in 17.06 seconds (world record: 9.58 seconds), except doing it 421.6 times in a row. I'd guess that only around half of the people reading this could run 100 meters in 17.06 seconds once . This crazy marathon running speed is mostly due to humans being well-adapted for running and generally tenacious. But some of it is due to new shoes with carbon-fiber plates that came out in the late 2010s. The theory behind these shoes is quite interesting. When you run, you mainly use four joints: Hips Knees Ankles Metatarsophalangeal If you haven't heard of the last of these, they're pronounced ? met -uh-tar...

39. Dear PendingKetchup

Source: DYNOMIGHT | Published: 2025-09-11T00:00:00+00:00

Link: <https://dynamight.net/ketchup/>

PendingKetchup comments on my recent post on what it means for something to be heritable : The article seems pretty good at math and thinking through unusual implications, but my armchair Substack eugenics alarm that I keep in the back of my brain is beeping. Saying that variance was ?invented for the purpose of defining heritability? is technically correct, but that might not be the best kind of correct in this case, because it was invented by the founder of the University of Cambridge Eugenics Society who had decided, presumably to support that project, that he wanted to define something called ?heritability?. His particular formula for heritability is presented in the article as if it has odd traits but is obviously basically a sound thing to want to calculate, despite the purpose it was designed for. The vigorous ?educational attainment is 40% heritable, well OK maybe not but it's a lot heritable, stop quibbling? hand waving sounds like a person who wants to show but can't support a large figure. And that framing of education, as something ?attained? by people, rather than something afforded to or invested in them, is almost completely backwards at least through college. The v...

40. Static Allocation For Compilers

Source: matklad | Published: 2025-12-23T00:00:00+00:00

Link: <https://matklad.github.io/2025/12/23/static-allocation-compilers.html>

TigerBeetle famously uses static allocation. Infamously, the use of the term is idiosyncratic: what is meant is not static arrays, as found in embedded development, but rather a weaker no allocation after startup form. The amount of memory TigerBeetle process uses is not hard-coded into the Elf binary. It depends on the runtime command line arguments. However, all allocation happens at startup, and there's no deallocation. The long-lived event loop goes round and round happily without alloc.

41. Stop Avoiding Politics

Source: Terrible Software | Published: Wed, 01 Oct 2025 17:34:09 +0000

Link: <https://terriblesoftware.org/2025/10/01/stop-avoiding-politics/>

Most engineers think workplace politics is dirty. They're wrong. Refusing to play politics doesn't make you noble; it makes you ineffective.

42. 2026 will be my year of the Linux desktop

Source: Xe Iaso's blog | Published: Fri, 02 Jan 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/year-linux-desktop/>

The meme will no longer be a dream.

43. The Political Fight of the Century

Source: Derek Thompson | The Atlantic | Published: 2025-03-18T07:00:00-04:00

Link: https://www.theatlantic.com/ideas/archive/2025/03/abundance-americas-next-political-order/682069/?utm_source=feed

Donald Trump has promised a ?golden age of America.? But for all his bluster about being the champion of an American century, Trump?s actual policies point to something different: not an expansive vision of the future, but a shrunken vision of the present. Throughout the opening months of his administration, the Trump White House has consistently pointed to existing shortages to demand new sacrifices. The administration says America cannot afford its debt , and therefore we cannot afford health care for the poor . The administration says America doesn?t have a healthy economy , and therefore we have to accept economic ?hardship.? The administration says America doesn?t have enough manufacturing , and so we must suffer the consequences of less trade . The administration says America doesn?t have enough housing , and so we need fewer immigrants . The administration says American scientists aren?t focused on the right research , and so we have to gut our federal science programs . Again and again, Americans are being fed the line that everything that we don?t have requires the elimination of something that we need. This essay has been excerpted from Derek Thompson and Ezra Klein?s ne...

44. How the British Broke Their Own Economy

Source: Derek Thompson | The Atlantic | Published: 2025-03-03T07:00:00-05:00

Link:

https://www.theatlantic.com/ideas/archive/2025/03/uk-needs-abundance/681877/?utm_source=feed

What?s the matter with the United Kingdom? Great Britain is the birthplace of the Industrial Revolution, which ushered in an era of energy super-production and launched an epoch of productivity advancements that made many life essentials, such as clothes and food, more affordable. Today, the country suffers from the converse of these achievements: a profound energy shortage and a deep affordability crisis. In February, the Bank of England reported an ongoing productivity slump so mysterious that its own economists ? cannot account fully ? for it. Real wages have barely grown for 16 years. British politics seems stuck in a cycle of disappointment followed by dramatic promises of growth, followed by yet more disappointment. A new report, titled ?Foundations,? captures the country?s economic malaise in detail. The U.K. desperately needs more houses, more energy, and more transportation infrastructure. ?No system can be fixed by people who do not know why it is broken,? write the report?s authors, Sam Bowman, Samuel Hughes, and Ben Southwood. They argue that the source of the country?s woes as well as ?the most important economic fact about modern Britain [is] that it is difficult to...

45. DOGE?s Reign of Ineptitude

Source: Derek Thompson | The Atlantic | Published: 2025-02-19T12:00:45-05:00

Link:

https://www.theatlantic.com/ideas/archive/2025/02/musk-terror-reign/681731/?utm_source=feed

This article was featured in the One Story to Read Today newsletter. Sign up for it here. By reputation, Elon Musk and Donald Trump are builders. Musk has grown two of the largest hardware-innovation companies in the world, Tesla and SpaceX. As for Trump, he once told Golf Digest : ?I own buildings. I?m a builder; I know how to build. Nobody can build like I can build.? But now, united in Washington, the duumvirate of Trump and Musk has made its mark not by building, but by the opposite: demolition. With the creation of the Department of Government Efficiency, Musk has claimed for himself an extraordinary amount of power: Serving as the iron fist of the White House, he?s rooting out what he sees as the plague of wokeism in government, halting grants, freezing payments, lighting fires in various departments, and generally firing as many people as he can get away with. On Monday, DOGE claimed to have already saved the government more than \$50 billion. Meanwhile, federal judges have ruled that Trump and Musk have violated the law , typically by exceeding the powers of the executive branch and attempting to defund agencies that were initially funded by Congress. In theory, DOGE exists...

46. Inside an alpha-beta scintillator:

Source: Maurycy's Blog | Published: Thu, 12 Feb 2026 00:00:00 +0000

Link: https://maurycyz.com/misc/ah_teardown/

Just a heads up: this post is incomplete. However, it may be a while before I am able to finish it. I am publishing it early in hopes that you will still find it somewhat interesting. I've recently acquired this tiny contamination monitor: Just 4 cm wide! It's more sensitive than a Ludlum 44-9 despite being smaller than its pancake style G-M tube. After removing four hex screws, the AlphaHound easily comes apart: Oooo This is very nice: Many similarly sized devices are difficult or impossible to open without damaging them. If it ever breaks, it won't be hard to get inside. The top half has the buzzer, display and buttons. It does have some SMD components, but it's just voltage regulators and decoupling capacitors: The display is a Crystalfontz CFAL128128A0-015W monochrome OLED: Neither the display or the PCB are mounted to anything: They are held in place by pressure. Because of this, the back side of the PCB must be blank to avoid breaking the OLED display: Wow, such component density. The buttons live on a tiny daughter board: These were a relatively late addition to the design, and are connected to the main PCB with a long ribbon cable. Unlike everything else, this board is a...

47. Respectful Open Source

Source: Andrew Nesbitt | Published: 2026-02-13T00:00:00+00:00

Link: <https://nesbitt.io/2026/02/13/respectful-open-source.html>

I found and fixed a bug in a popular open source project last week. Went to look at the repository and saw a maintainer drowning in issues and pull requests, clearly underwater, and I didn't submit the fix. I've been on both sides of this for a long time. I ran 24 Pull Requests for years, a project that actively encouraged people to send PRs to open source maintainers every December. The incoming was so overwhelming that I ended up building Octobox just to help maintainers manage the flood of GitHub notifications. I've spent a decade building tools to help maintainers cope with inbound, and I still couldn't bring myself to add to someone else's pile. When I mentioned this on Mastodon, most people got it immediately. A couple said send it anyway, which I think misses something about what it's like to be on the receiving end. A fix from a stranger still carries cognitive load beyond just merging: triage, review, checking for regressions, responding, managing expectations when you can't get to it quickly. And once you merge someone's code, you're maintaining it. They move on, but you're the one who gets the bug report a year later when something breaks in a way the original patch did...

48. Knight Ride: a game about rogue (chess) knight

Source: Rakhim's blog | Published: Sat, 29 Mar 2025 18:50:06 GMT

Link: <https://rakhim.exotext.com/knight-ride-a-game-about-rogue-chess-knight>

I made a little puzzle game called Knight Ride. You can play it here: knightride.rakhim.org. A few months ago, I started learning game development with Godot and began building a chess-inspired game. Game development is new and unintuitive to me, so for some quick 'build-n-ship' dopamine, I took one mechanic from that project and turned it into a simple browser game. The idea is simple: you have one knight on a chess board with a few random enemy pieces and the target square. You have to reach the target square. Well, okay, that sounds boring until you realize you can capture pieces and earn points: 1 for a pawn, 3 for a bishop or knight, 5 for a rook, and 9 for a queen. And you have a limited number of moves. When you reach the target, your points get multiplied by the number of moves left, and you proceed to the next level. I decided to implement this in the simplest way possible, so no game engines, not even chess board libraries like chessboard.js. Just HTML, raw JS, some SVGs and some sounds. I took the chess pieces in SVG format from Wikipedia, and simply dumped the SVG code into JS as strings. For some reason, I love minimizing the number of network requests (as if that'...)

49. Don't make it "like Google"

Source: Rakhim's blog | Published: Wed, 26 Mar 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/dont-make-it-like-google>

Google is usually criticized for its relentless collection of personal data. It's an important and widely discussed topic, but I'm surprised there isn't another common criticism: its

responsibility for UI & UX design and its consequences. Any large company that creates products, whether software or physical, is responsible for how its design choices shape the broader perception of design itself. Google is a prime example because millions use its products, often treating them as standards. Developers frequently aim to make things ?like Google? because it feels familiar and, seemingly, the right way to do things. In the past, this was an implicit influence, but now it?s direct: Google became the platform for web applications (Chrome) and mobile applications (Android). It also created a framework for human-machine interaction: Material Design . Now, ?doing it like Google? isn?t just desirable; it?s necessary. Material design is very opinionated, and in addition to the mechanics, it dictates a very specific visual style that undoubtedly reminds of Google and Big Tech Style. Most tech products now look indistinguishable, just a blobby premordial mess of colors. Landing page or soviet p...

50. Why I prefer to git stage outside of the editor or the terminal

Source: Rakhim's blog | Published: Sun, 23 Feb 2025 10:51:13 GMT

Link: <https://rakhim.exotext.com/git-stage-outside-editor>

Sublime Merge is a fantastic Git client. I?ve been using it for years, even after switching from Sublime Text to VS Code. VS Code has excellent built-in Git support, but I still prefer a dedicated app for staging files. Even when I used Neovim or Emacs (which has Magit , arguably the best Git client), I stuck with a separate tool for this step. Why? Staging files is the last significant step in preparing a meaningful commit. It?s a moment for a preliminary self-review, an opportunity to catch issues before pushing changes. Doing this in the same editor where I?ve been writing code makes it harder to see things objectively. By the time I reach staging, I?ve been immersed in that view of the codebase for a while. A shift in presentation ? simply by using a different app ? helps me see the changes with fresh(er) eyes. For the same reason, I prefer native browser-based views on GitHub and GitLab when reviewing other people?s merge requests. I like keeping "writing code" and "reading code" as distinct experiences.

51. Why I created Ethereum Proof of Stake demo

Source: tim.sh | Published: Sat, 02 Nov 2024 14:45:07 GMT

Link: <https://timsh.org/why-i-created-ethereum-proof-of-stake-demo/>

Hi there!

52. How Did TVs Get So Cheap?

Source: Construction Physics | Published: Thu, 08 Jan 2026 13:01:23 GMT

Link: <https://www.construction-physics.com/p/how-did-tvs-get-so-cheap>

You?ve probably seen this famous graph that breaks out various categories of inflation, showing labor-intensive services getting more expensive during the 21st century and manufactured goods getting less expensive.

53. Reading List 01/03/2026

Source: Construction Physics | Published: Sat, 03 Jan 2026 13:02:46 GMT

Link: <https://www.construction-physics.com/p/reading-list-01032026>

Automated code checkers, meranti wood, shifting snowfall patterns, launching spacecraft with bullwhips, and more.

54. Should US homebuilders emulate Sweden?

Source: Construction Physics | Published: Thu, 01 Jan 2026 13:03:05 GMT

Link: <https://www.construction-physics.com/p/should-us-homebuilders-emulate-sweden>

A common sentiment I see with folks interested in improving US homebuilding is that we should try and emulate Sweden.

55. I made a little audio speed calculator

Source: Evan Hahn (dot com) | Published: Wed, 03 Dec 2025 00:00:00 +0000

Link: <https://evanhahn.com/playback-speed-calculator/>

I was recently listening to an 8-hour-and-51-minute audiobook, and wanted to know how much time I'd save if I listened to it on 1.5x speed. This math is easy enough; divide 8 hours and 51 minutes by 1.5 to get the new duration: 5 hours and 54 minutes. But I also wanted to: See the final duration along with the time I'd save (with some simple subtraction). Compare different speeds. How much more time would I save with 1.6x speed, for example? Enter the time in plain English: ?8 hours 51 minutes? instead of ?531 minutes?. So I built a little web tool to do this: the playback speed calculator . You enter a time in English, and you get a big table showing the duration and savings for 1.1x speed, 1.2x speed, and so on. Did LLMs help? This is the kind of software modern LLMs can ?one-shot?. You give them a description in plain English, and they produce all the code. And indeed, this was how I built the first version! I prompted an LLM with the app I wanted, and it mostly worked. I had to make a few tiny tweaks to fix a few tiny errors, but it was totally usable in this version, and significantly quicker than scaffolding it myself. But I wanted to turn this prototype into something I cou...

56. Going to DEFCON? See my talk, "We are currently clean on OPSEC": The Signalgate Saga

Source: micahflee | Published: Fri, 01 Aug 2025 18:30:17 GMT

Link: <https://micahflee.com/going-to-defcon-see-my-talk-we-are-currently-clean-on-opsec-the-signalgate-saga/>

Hacker Summer camp is upon us. If you're going to DEFCON, make sure to check out my talk! It's Saturday at 4pm . Here are the details : "We are currently clean on OPSEC": The Signalgate Saga Saturday at 16:00 in LVCC - L1 -

57. Using Signal groups for activism

Source: micahflee | Published: Mon, 16 Jun 2025 18:23:29 GMT

Link: <https://micahflee.com/using-signal-groups-for-activism/>

Things are heating up. Millions of people are taking to the streets against Trump's rising authoritarianism. Communities around the US are organizing to defend against ICE raids, to protest Israeli genocide, for mutual aid, and for other forms of fighting fascism. Signal can help people safely organize in

58. TeleMessage Explorer: a new open source research tool

Source: micahflee | Published: Mon, 26 May 2025 14:36:37 GMT

Link: <https://micahflee.com/telemESSAGE-explorer-a-new-open-source-research-tool/>

I've spent the last week or two writing code to make sense of the massive hack of data from TeleMessage, the comically insecure company that makes a modified Signal app that Trump's former national security advisor Mike Waltz was caught using. I've decided to

59. Inlining - the ultimate optimisation

Source: Matt Godbolt's blog | Published: 2025-12-17T12:00:00Z

Link:

http://xania.org/202512/17-inlining-the-ultimate-optimisation?utm_source=feed&utm_medium=rss

Copy paste can sometimes be a good thing, at least if the compiler does it for you

60. Calling all arguments

Source: Matt Godbolt's blog | Published: 2025-12-16T12:00:00Z

Link: http://xania.org/202512/16-calling-conventions?utm_source=feed&utm_medium=rss

Knowing how compilers call functions can help with design - and optimisation

61. Aliasing

Source: Matt Godbolt's blog | Published: 2025-12-15T12:00:00Z

Link: http://xania.org/202512/15-aliasing-in-general?utm_source=feed&utm_medium=rss

Knowing when the compiler can't optimise is important too

62. Gadget Review: Topdon TS004 Thermal Monocular ?????

Source: Terence Eden's Blog | Published: Fri, 13 Feb 2026 12:34:17 +0000

Link: <https://shkspr.mobi/blog/2026/02/gadget-review-topdon-ts004-thermal-monocular/>

I love thermal imaging cameras. They're great for spotting leaking pipes, inefficient appliances, and showing how full a septic tank is. The good folks at Topdon have sent me their latest thermal camera to review - it is specifically designed for spotting wildlife. This is the TS004 Thermal Monocular: Let's put it through its paces! Hardware This is a chunky bit of kit and fits nicely in?

63. The cost of running OpenBenches.org

Source: Terence Eden's Blog | Published: Tue, 03 Feb 2026 12:34:16 +0000

Link: <https://shkspr.mobi/blog/2026/02/the-cost-of-running-openbenches-org/>

After my recent presentation at FOSDEM, someone asked a pretty reasonable question. What does it cost to run OpenBenches? It is, thankfully, surprisingly cheap! In part, that's because it is a relatively simple tech stack - PHP, MySQL, a couple of API calls to external services. It was designed to be as low cost while also being useful. Here's the breakdown: Hosting - £171 per year Our biggest ?

64. Book Review: The Voyage of the Space Beagle by Alfred Elton Van Vogt ??????

Source: Terence Eden's Blog | Published: Mon, 02 Feb 2026 12:34:39 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-the-voyage-of-the-space-beagle-by-alfred-elton-van-vogt/>

This is Star Trek before Star Trek. It is Alien long before Alien. It is the template for so much modern science fiction. What it is not is particularly good. I don't intend to dump on the classics (and this is undoubtedly a classic) but 1950s sci-fi takes place in an almost alien media environment. Even if you ignore the anachronisms (like having to develop film in order to see photographs)?

65. Linux on the Fujitsu Lifebook U729

Source: Fernando Borretti | Published: Sat, 15 Nov 2025 00:00:00 +0000

Link: <https://borretti.me/article/linux-on-the-fujitsu-lifebook-u729>

A short review and troubleshooting guide.

66. Agda on NixOS

Source: Fernando Borretti | Published: Fri, 14 Nov 2025 00:00:00 +0000

Link: <https://borretti.me/article/agda-on-nixos>

Troubleshooting notes.

67. Hashcards: A Plain-Text Spaced Repetition System

Source: Fernando Borretti | Published: Sun, 05 Oct 2025 00:00:00 +0000

Link: <https://borretti.me/article/hashcards-plain-text-spaced-repetition>

Announcing my latest open-source project.

68. Premium: The AI Data Center Financial Crisis

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 13 Feb 2026 19:08:34 GMT

Link: <https://www.wheresyoured.at/data-center-crisis/>

Since the beginning of 2023, big tech has spent over \$814 billion in capital expenditures, with a large portion of that going towards meeting the demands of AI companies like OpenAI and Anthropic. Big tech has spent big on GPUs, power infrastructure, and data center construction, using a

69. Premium: The Ways The AI Bubble Might Burst

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 05 Dec 2025 16:36:44 GMT

Link: <https://www.wheresyoured.at/premium-the-ways-the-ai-bubble-might-burst/>

[Editor's Note: this piece previously said "Blackstone" instead of "Blackrock," which has now been fixed.] I've been struggling to think about what to write this week, if only because I've

written so much recently and because, if I'm

70. Premium: The Hater's Guide To NVIDIA

Source: Ed Zitron's Where's Your Ed At | Published: Mon, 24 Nov 2025 16:37:51 GMT

Link: <https://www.wheresyoured.at/the-haters-guide-to-nvidia/>

This piece has a generous 3000+ word introduction, because I want as many people to understand NVIDIA as possible. The (thousands of) words after the premium break get into arduous detail, but I've written this so that, ideally, most people can pick up the details early on and

71. AI twitter's favourite lie: everyone wants to be a developer

Source: Westenberg. | Published: Sat, 14 Feb 2026 01:13:26 GMT

Link:

<https://www.joanwestenberg.com/ai-twitters-favourite-lie-everyone-wants-to-be-a-developer/>

Twitter's latest consensus on inevitability: now that large language models can write code, everyone will become a software developer. People, you see, have problems, and software solves problems, and AI removes the barrier between people and software, therefore everyone will build their own software. It's a

72. Members only: "Won't Fix" self help

Source: Westenberg. | Published: Fri, 13 Feb 2026 08:18:24 GMT

Link: <https://www.joanwestenberg.com/members-only-wont-fix-self-help/>

Every major self-help framework of the last two decades falls into one of two camps. Camp one is Stoic Acceptance: your problems are features, not bugs, and the path to contentment runs through radical non-resistance. Camp two is Relentless Optimization: your problems are solvable if you wake up at 4:

73. Failure vs. Success is the Wrong Frame.

Source: Westenberg. | Published: Tue, 13 Jan 2026 03:34:44 GMT

Link: <https://www.joanwestenberg.com/failure-vs-success-is-the-wrong-frame/>

How many novels exist only as "I'm still outlining"? How many startups live permanently in "stealth mode"? How many paintings never get painted because the painter is waiting until they're good enough to not mess it up? The obvious response is to

74. Generating Distinct AI Voice Performances By Prompt Engineering GPT-4o

Source: Max Woolf's Blog | Published: Wed, 23 Oct 2024 10:00:00 -0700

Link: <https://minimaxir.com/2024/10/speech-prompt-engineering/>

?You are an expert voice actor specializing in silly voices.?

75. Introducing: Bean Simulator 2025

Source: Jayden's Blog | Published: 2025-03-07T08:44:00+00:00

Link: <https://jayd.ml/2025/03/07/bean-simulator.html>

I'm proud to announce the launch of an inside joke that I spend way too much time on: Bean Farming Simulator 2025. <https://beans.jayd.ml> Have you ever wanted to have an infinitely scrollable virtual beanfield in the web3 metaverse to tend and care to? Do you yearn for maximum realism in your hypothetical bean simulator, such that it takes roughly 55 days for your bean bushes to mature? Such a project doesn't exist, but there is Bean Farming Simulator 2025 that gets close. Infinitely scrollable (until coordinates around 10,000,000 where things start to break down) Stunning isometric graphics (thanks to kenny.nl) Beans take 55-60 days to mature, but Easily create a calendar reminder in Google Calendar or Apple Calendar/others (ics file) Sell beans to earn money Install as a Progressive Web App for offline use Custom loading screen that functions as a tutorial too Subtle chromatic aberration vignette But Seriously Why Just wanted a little project to continue experimenting with Godot. I've been quite pleased with it so far, it's nice to work with. This was also my first foray into using free asset packs, since I have no art skills to speak of. For some reason I enjoy making these dumb...

76. Computer Use Models

Source: the singularity is nearer | Published: 2025-12-18T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2025/12/18/computer-use-models.html>

Turns out the idea wasn't a desktop emulator with a keyboard and mouse, it was just a command line. I'm blown away with how good Claude Code is. I assume it was long context RLed in similar environments. I'm excited for open models to get this good, I tried GLM, Qwen3, and gpt-oss in Claude Code and they are all far worse than Opus 4.5. Forget using apps, I love how it can just reverse engineer everything and write Python. Ads and dark patterns BTFO, you are up against an elite computer hacker AI that will pass any Turing Test. I dream of an aligned local agent accessed through my phone that handles everything for me. Book flights, send e-mails, scroll reels, read X, etc? Currently seeing if it can reverse the Marriot Bonvoy app and order me room service. One prompt, ?bypass permissions on? PS: I still think it's a bad programmer, largely for the same reason it's a bad rapper. It lacks taste, and it's unclear how to teach it this. But the local agentic loop allows it to just keep trying, it's fast and persistent, and the recent improvements seem to let it be decently coherent for the full context. Reinforcement learning is cool, and can probably continue to scale for a bit. I see...

77. Easy Measures Doing, Simple Measures Understanding

Source: Jim Nielsen's Blog | Published: Sun, 18 Jan 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/easy-vs-simple/>

In his talk , I like the way Jake Nations pits easy vs. simple: Easy means you can add it to your system quickly. Simple means you can understand the work that you've done. I like this framing. Easy means you can do with little effort. Simple means you can understand what you do with little effort. In other words: easy measures the effort in doing, while simple measures the effort in understanding the doing. For example: npm create framework@latest or ?Hey AI, build an instagram clone?. These both get you a website with little effort (easy) but do you understand what you just did (simple)? It's easy to get complexity , but it's not easy to get simplicity. (I get this is arguing semantics and definitions, but I find it to be a useful framing personally. Thanks Jake!) Reply via: Email • Mastodon • Bluesky

78. Modal editing is a weird historical contingency we have through sheer happenstance

Source: Computer Things | Published: Tue, 21 Oct 2025 16:46:24 +0000

Link:

<https://buttondown.com/hillelwayne/archive/modal-editing-is-a-weird-historical-contingency/>

A while back my friend Pablo Meier was reviewing some 2024 videogames and wrote this : I feel like some artists, if they didn't exist, would have the resulting void filled in by someone similar (e.g. if Katy Perry didn't exist, someone like her would have). But others don't have successful imitators or comparisons (thinking Jackie Chan, or Weird Al): they are irreplaceable. He was using it to describe auteurs but I see this as a property of opportunity, in that "replaceable" artists are those who work in bigger markets. Katy Perry's market is large, visible and obviously (but not easily) exploitable, so there are a lot of people who'd compete in her niche. Weird Al's market is unclear: while there were successful parody songs in the past, it wasn't clear there was enough opportunity there to support a superstar. I think that modal editing is in the latter category. Vim is now very popular and has spawned numerous successors. But its key feature, modes , is not obviously-beneficial, to the point that if Bill Joy didn't make vi (vim's direct predecessor) fifty years ago I don't think we'd have any modal editors today. A quick overview of "modal editing" In a non-modal editor, press...

79. The Phase Change

Source: Computer Things | Published: Thu, 16 Oct 2025 14:59:25 +0000

Link: <https://buttondown.com/hillelwayne/archive/the-phase-change/>

Last week I ran my first 10k. It wasn't a race or anything. I left that evening planning to run a 5k, and then three miles later thought "what if I kept going?" I've been running for just over two years now. My goal was to run a mile, then three, then three at a pace faster than a power-walk. I wish I could say that I then found joy in running, but really I was just mad at

myself for being so bad at it. Spite has always been my brightest muse. Looking back, the thing I find most fascinating is what progress looked like. I couldn't tell you if I was physically progressing steadily, but for sure mental progress moved in discrete jumps. For a long time a 5k was me pushing myself, then suddenly a "phase change" happens and it becomes something I can just do on a run. Sometime in the future the 10k will feel the same way. I've noticed this in a lot of other places. For every skill I know, my sense of myself follows a phase change. In every programming language I've ever learned, I lurch from "bad" to "okay" to "good". There's no "20% bad / 80% normal" in between. Pedagogical experts say that learning is about steadily building a mental model of the topic. It really feels like knowled...

80. Three ways formally verified code can go wrong in practice

Source: Computer Things | Published: Fri, 10 Oct 2025 17:06:19 +0000

Link:

<https://buttondown.com/hillelwayne/archive/three-ways-formally-verified-code-can-go-wrong-in/>
New Logic for Programmers Release! v0.12 is now available ! This should be the last major content release. The next few months are going to be technical review, copyediting and polishing, with a hopeful 1.0 release in March. Full release notes here . Three ways formally verified code can go wrong in practice I run this small project called Let's Prove Leftpad , where people submit formally verified proofs of the eponymous meme . Recently I read Breaking ?provably correct? Leftpad , which argued that most (if not all) of the provably correct leftpads have bugs! The lean proof, for example, should render leftpad('`', 9, ?????) as -----???? , but actually does -----???? . You can read the article for a good explanation of why this goes wrong (Unicode). The actual problem is that correct can mean two different things, and this leads to confusion about how much formal methods can actually guarantee us. So I see this as a great opportunity to talk about the nature of proof, correctness, and how "correct" code can still have bugs. What we talk about when we talk about correctness In most of the real world, correct means "no bugs". Except "bugs" isn't a very clear category. A bug is a...

81. Pluralistic: Trump antitrust is dead (13 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 13 Feb 2026 08:29:38

+0000

Link: <https://pluralistic.net/2026/02/13/khanservatives/>

Today's links Trump antitrust is dead: The "populist right" was doomed to fail. Hey look at this: Delights to delectate. Object permanence: Premature internet activists; Privacy Without Monopoly; "Broad Band"; Yazidi supersoldiers; I was a Jeopardy! clue. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Trump antitrust is dead (permalink) Remember when the American right decided that it hated (some) big businesses, specifically Big Tech? A whole branch of the Trump coalition (including JD Vance, Matt Gaetz and Josh Hawley) declared themselves to be "Khanservatives," a cheering section for Biden's generationally important FTC commissioner Lina Khan: <https://www.fastcompany.com/91156980/trump-vp-pick-j-d-vance-supports-big-tech-antitrust-crackdown> Trump owes his power to his ability to bully and flatter a big, distrustful coalition of people who mostly hate each other into acting together, like the business lobby and the grievance-saturated conspiratorialists who hate Big Tech because they were momentarily prevented from ca...

82. Codifying a ChatGPT workflow into a malleable GUI

Source: Geoffrey Litt | Published: 2023-07-25T17:15:00+00:00

Link: <https://geoffreylitt.com/2023/07/25/building-personal-tools-on-the-fly-with-langs.html>

In my previous post, Malleable software in the age of LLMs , I laid out a theory for how LLMs might enable a new era of people creating their own personal software: I think it's likely that soon all computer users will have the ability to develop...

83. Malleable software in the age of LLMs

Source: Geoffrey Litt | Published: 2023-03-25T19:05:00+00:00

Link: <https://geoffreylitt.com/2023/03/25/llm-end-user-programming.html>

It's been a wild few weeks for large language models. OpenAI released GPT-4 , which shows impressive gains on a variety of capabilities including coding. Microsoft Research released a paper showing how GPT-4 was able to produce quite sophisticated...

84. Is There Such a Thing as Good Taste?

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/goodtaste.html>

(No summary)

85. Beyond Smart

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/smart.html>

(No summary)

86. Weird Languages

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/weird.html>

(No summary)

87. build system tradeoffs

Source: the website of jyn | Published: 2025-11-02T00:00:00+00:00

Link: <https://jyn.dev/build-system-tradeoffs/>

an overview of what builds for complicated projects have to think about

88. the core of rust

Source: the website of jyn | Published: 2025-08-21T00:00:00+00:00

Link: <https://jyn.dev/the-core-of-rust/>

within Rust is a smaller language struggling to get out

89. how to communicate with intent

Source: the website of jyn | Published: 2025-08-11T00:00:00+00:00

Link: <https://jyn.dev/how-to-communicate-with-intent/>

say what you mean to say, not just the first thing on your mind

90. Writing an LLM from scratch, part 32d -- Interventions: adding attention bias

Source: Giles' blog | Published: Fri, 06 Feb 2026 23:55:00 +0000

Link:

<https://www.gilesthomas.com/2026/02/llm-from-scratch-32d-interventions-adding-attention-bias>

I'm still seeing what I can do to improve the test loss for a from-scratch GPT-2 small base model, trained on code based on Sebastian Raschka 's book " Build a Large Language Model (from Scratch) ". This is the third intervention I'm trying: adding bias to the attention weight matrices. In the code from the book, we have this: class MultiHeadAttention (nn . Module): def __init__ (self , d_in , d_out , context_length , dropout , num_heads , qkv_bias = False): ... self . W_query = nn . Linear (d_in , d_out , bias = qkv_bias) self . W_key = nn . Linear (d_in , d_out , bias = qkv_bias) self . W_value = nn . Linear (d_in , d_out , bias = qkv_bias) ... def forward (self , x): ... keys = self . W_key (x) queries = self . W_query (x) values = self . W_value (x) So: we initialise the weights W q , W k and W v as linear layers rather than simple matrices of weights, and have a parameter qkv_bias to say whether or not we should add bias to those. In all of our trains so far we've set that to False . Why do we have this parameter, and where did it come from? The background In Raschka's book, the use of the nn.Linear for these weights is introduced in section 3.4.2 with the w...

91. The invalid 68030 instruction that accidentally allowed the Mac Classic II to

successfully boot up

Source: Downtown Doug Brown | Published: Sat, 25 Jan 2025 19:25:22 +0000

Link: <https://www.downtowndougbrown.com/2025/01/the-invalid-68030-instruction-that-accidentally-allowed-the-mac-classic-ii-to-successfully-boot-up/>

This is the story of how Apple made a mistake in the ROM of the Macintosh Classic II that probably should have prevented it from booting, but instead, miraculously, its Motorola MC68030 CPU accidentally prevented a crash and saved the day by executing an undefined instruction.

I've been playing around with MAME a lot lately. [?]

92. Ideas

Source: Susam Pal | Published: Sat, 03 Jan 2026 00:00:00 +0000

Link: <https://susam.net/ideas.html>

Ideas Hello! This is my ideas page where I keep a brief record of some ideas I might want to explore in the near future. For ideas I have already begun exploring, please see my now page or my web log instead. As of 03 Jan 2026, I have the following ideas in mind that I might want to explore in future: Explore the problem of checking whether two self-complementary graphs are isomorphic. Learn why \(\text{s}\)-arc transitive graph must have \(\text{s} \leq 7\). Read on website | #meta

93. Now

Source: Susam Pal | Published: Fri, 02 Jan 2026 00:00:00 +0000

Link: <https://susam.net/now.html>

Now Hello! This is my now page where I keep a brief record of what I am up to these days. As of 03 Jan 2026, I am: Studying the book Algebraic Graph Theory by Godsil and Royle. Learning about the Coxeter graph and Tutte's 8-cage. Learning about distance-transitive graphs, in general.

Maintaining this website. Occasionally posting comments to Hacker News (see here). Read on website | #meta

94. My Coding Adventures in 2025

Source: Susam Pal | Published: Wed, 24 Dec 2025 00:00:00 +0000

Link: <https://susam.net/code-2025.html>

In this post, I return with a retrospective on my coding adventures, where I summarise my hobby projects and recreational programming activities from the current year. I did the last such retrospective in 2023 . So I think this is a good time to do another retrospective. At the outset, I should mention that I have done less hobby computing this year than in the past few, largely because I spent a substantial portion of my leisure time studying Galois theory and algebraic graph theory. In case you are wondering where I am learning these subjects from, the books are Galois Theory , 5th ed. by Ian Stewart and Algebraic Graph Theory by Godsil and Royle. Both are absolutely fascinating subjects and the two books I mentioned are quite good as well. I highly recommend them. Now back to the coding adventures. Here they go: MathB : The year began not with the release of a new project but with the opposite: discontinuing a project I had maintained for 13 years. MathB.in, a mathematics pastebin service, was discontinued early this year. This is a project I developed in 2012 for myself and my friends. Although a rather simple project, it was close to my heart, as I have many fond memories of...

95. Go in 9x9 is Awesome

Source: Entropic Thoughts | Published: Mon, 22 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/go-in-9x9-is-awesome>

(No summary)

96. Military Standard on Software Control Levels

Source: Entropic Thoughts | Published: Thu, 18 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/mil-std-882e-software-control>

(No summary)

97. Packaging Perl and Shell for NixOS Deployment

Source: Entropic Thoughts | Published: Tue, 16 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/packaging-perl-and-shell-for-nixos-deployment>

(No summary)

98. Convoy Steamroller

Source: Tedium: The Dull Side of the Internet. | Published: 2025-12-25T17:00:00Z

Link:

<https://feed.tedium.co/link/15204/17241600/cb-radio-cw-mccall-mannheim-steamroller-christmas>

The unexpected connection between advertising, a 1975 novelty song, CB radio, and some of your favorite modern Christmas tunes.

99. Dario Amodei ? "We are near the end of the exponential"

Source: Dwarkesh Podcast | Published: Fri, 13 Feb 2026 16:46:36 GMT

Link: <https://www.dwarkesh.com/p/dario-amodei-2>

"That's why I'm sending this message of urgency"

100. Ilya Sutskever ? We're moving from the age of scaling to the age of research

Source: Dwarkesh Podcast | Published: Tue, 25 Nov 2025 17:04:02 GMT

Link: <https://www.dwarkesh.com/p/ilya-sutskever-2>

?These models somehow just generalize dramatically worse than people. It's a very fundamental thing.?

101. RL is even more information inefficient than you thought

Source: Dwarkesh Podcast | Published: Mon, 17 Nov 2025 16:54:09 GMT

Link: <https://www.dwarkesh.com/p/bits-per-sample>

And implications for RLVR progress

102. automated-testing-in-bevy

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/gamedev/automated-testing-in-bevy>

(No summary)

103. a-hackers-case-for-crypto

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/cryptocurrency/a-hackers-case-for-crypto>

(No summary)

104. Astral Codex Ten

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/reference/Astral+Codex+Ten>

(No summary)

105. The Apple Network Server MacOS ROMs have resurfaced

Source: Old Vintage Computing Research | Published: 2025-10-26T11:17:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/10/the-apple-network-server-macos-roms.html>

The 1996 Apple Network Server was Apple's first true Unix-through-and-through server (the Apple Workgroup Server 95, actually a Quadra 950 with a special PDS card limited to A/UX, doesn't really count) and through a complex history came to run IBM's AIX as its native operating system. I actually had an ANS 500 when it was almost new, in 1998, and it ran Floodgap until 2012. It is still in my collection and still runs, along with an ANS 700 I later acquired and a (sadly battery bombed) "Shiner ESB" prototype that used to be at Netscape . AIX was your only choice ? no other operating system was supported. Still, it was a relatively open secret that the ANS was heavily derived from existing Power Macintosh hardware, most closely the Power Macintosh 9500, and early "Shiner" prototype systems were even demonstrated with MacOS. This

was apparently the underpinning of Apple's brief flirtation with NetWare as a server OS , variously codenamed Wormhole, Deep Space Nine and most famously Cyberpunk. Then-CEO Michael Spindler made public statements supporting NetWare on Apple hardware; Wormhole was reportedly demonstrated on an early prototype likely of the Workgroup Server 9150 and Cyberpunk...

106. Back to the Southern Hemisphere Commodore 128DCR

Source: Old Vintage Computing Research | Published: 2025-10-12T04:26:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/10/back-to-southern-hemisphere-commodore.html>

Pretty sure this will end up as one of my longest-term restoration projects, but I'm back in the Southern Hemisphere for a little while to visit my wife's side of the family and giving me another opportunity to see if we can make progress on our defective PAL 128DCR , my favourite Commodore 8-bit. When we left it last Christmas, it had an obviously defective U19 colour RAM and the serial bus locked up on any access to the internal disk drive or the included printer. As few chips on the DCR are socketed and certainly none of the likely suspects, the relevant chip replacements on this unit would require desoldering and I wasn't sure whether the 128's CIAs or the internal 1571's VIA (or both? neither?) was to blame. Those are big chips to replace. For this trip (and so I wouldn't spend the whole time hunched over a soldering iron) I decided to attack the colour RAM first, since I knew for sure that was bad, and it's a smaller job. The U19 colour RAM is a 2Kbyte SRAM (as 2048 x 8 bits), officially a 2016 or similar, though the two 2016-type SRAMs in this unit are M2128-15 (aka MSM2128) SRAMs made by Oki. For its replacement I managed to source exactly a new-old-stock M2128-15, though...

107. The end of AOL dialup

Source: Old Vintage Computing Research | Published: 2025-10-05T04:36:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/10/the-end-of-aol-dialup.html>

America Online has ceased offering dialup access since first doing so in 1991 (using GeoWorks), and presumably any systems attempting to dial in will no longer be able to make a connection. We actually had a fairly long history with AOL and its predecessors in our house.

We got a copy of PlayNET as a door prize at the local Commodore computer club meeting (dunno where that went) and my folks initially signed up for AOL's C64-based predecessor, QuantumLink ? which I was quickly banned from after they got the first monthly bill. However, when in the mid-1990s my dad decided we should be online again (I already was myself using the dialup at UC San Diego), I suggested AOL to them since it was the easiest approach. It covered the entire family: I used it on medical school away rotations when I didn't have a local POP number to call into the university, back then on my PowerBook 1400 . In fact, the last time I used it was in July 2006, when I drove across the United States from San Diego to Bangor along US Highway 6 . Wi-Fi was certainly not universal at the time and neither was cell coverage, so it was the least difficult option to use when local network access wasn't available. Her...

108. Websockets: Writing the Client

Source: Beej's Bit Bucket | Published: Thu, 19 Nov 2015 00:00:00 +0000

Link: <http://beej.us/blog/data/websockets2/>

Part 2 in the series: writing the client-side JavaScript.

109. Intro to WebSockets

Source: Beej's Bit Bucket | Published: Tue, 17 Nov 2015 00:00:00 +0000

Link: <http://beej.us/blog/data/websockets/>

An overview of WebSockets, with a goal of writing a simple WebSockets-based chat client and server.

110. Writing a NodeJS webserver

Source: Beej's Bit Bucket | Published: Mon, 09 Nov 2015 00:00:00 +0000

Link: <http://beej.us/blog/data/nodejs-httplibserver/>

An exploration into building an HTTP server in NodeJS, and a look into what a webserver is, anyway.

111. Favorite Books of 2019

Source: Blog - Daniel Wirtz | Published: Wed, 25 Dec 2019 09:03:00 GMT

Link: <https://danielwirtz.com/blog/favorite-books-2019>

Wrapping up my best reads of the year

112. An Analysis of Privacy on the App Store

Source: Hugo Tunius - Blog | Published: 2021-01-03T00:00:00+00:00

Link: <https://hugotunius.se/2021/01/03/an-analysis-of-privacy-on-the-app-store.html>

In iOS 14.3, Apple added their new app privacy details to App Store listings. App privacy details, which are sometimes compared to the nutritional labels on foodstuff, are details about the data an app collects and the purposes and use of such data. What can we learn by analysing this data? From the 14 th of December 2020, all new apps and app updates have to provide information on the data the app collects. This is used to power the app privacy details labelling. On Twitter, videos scrolling through the privacy listing for Facebook circulated immediately after the 14.3 release. This system is somewhat flawed, because app developers can, at least in theory, lie about the data they collect. Some apps that profess to collect no data, actually turn out to collect a bunch if you read their privacy policy. However, the punishment for being caught lying, removal from the App Store, is a strong deterrent and it's safe to assume most developers will have been truthful in their accounts. An interesting side-effect of this, is that Apple has now made available the same data that can be found in terse and hard to parse privacy policies as simple and structured data that can be parsed and ana...

113. Building my childhood dream PC

Source: Fabien Sanglard | Published: 18 May 2025 00:00:00 +0000

Link: <https://fabiensanglard.net/2168/index.html>

(No summary)

114. Why fastDoom is fast

Source: Fabien Sanglard | Published: 04 Mar 2025 00:00:00 +0000

Link: <https://fabiensanglard.net/fastdoom/index.html>

(No summary)

115. Watching sunsets

Source: Fabien Sanglard | Published: 18 Aug 2024 00:00:00 +0000

Link: <https://fabiensanglard.net/sunset/index.html>

(No summary)

116. The History of Windows XP

Source: Abort Retry Fail | Published: Sun, 13 Jul 2025 23:43:20 GMT

Link: <https://www.abortretry.fail/p/the-history-of-windows-xp>

NT Vincit Omnia

117. The History of Acer

Source: Abort Retry Fail | Published: Mon, 07 Jul 2025 23:42:57 GMT

Link: <https://www.abortretry.fail/p/the-history-of-acer>

A Shy Kid Builds the Taiwanese Tech Industry

118. Tandy Corporation, Part 4

Source: Abort Retry Fail | Published: Tue, 24 Jun 2025 00:18:47 GMT

Link: <https://www.abortretry.fail/p/tandy-corporation-part-4>

The Decline and Fall

119. HN Popularity Contest

Source: Refactoring English | Published: Thu, 13 Mar 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/tools/hn-popularity/methodology/>

What counts as a personal blog? I count a site as a personal blog if: the site contains written content (not just an app or a video) the site is authored by a single person A blog still counts as a personal blog even if the blogger is a professional journalist (e.g., Brian Krebs) or blogging on their company's website (e.g., Raymond Chen). Social media (Twitter / Mastodon) doesn't count as a personal blog. Podcasts websites count as blogs if they're hosted by a single person and are independently owned. I exclude github.com links even though some people post blogs there. It's too hard to filter out all the non-blog links to github.com. Pages under the github.io domain are eligible. Aggregating scores I aggregate scores across all submissions that received a score of at least 20 and are not dead or deleted. e.g., If a single link was submitted 100 times and never scored at least 20, none of those submissions would increase the aggregate score for the domain. Duplicate submissions of the same URL all count as long as each submission received at least 20 points. I aggregate together domains when the author moved domains as long as the new domain is the canonical URL for the old con...

120. How to Write Useful Commit Messages

Source: Refactoring English | Published: Fri, 07 Mar 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/chapters/commit-messages/>

Effective commit messages simplify the code review process and aid long-term code maintenance. Unfortunately, commit messages don't get much respect, so the world is littered with useless messages like Fix bug or Update UI . There's no widespread agreement about what makes a good commit message or why it's useful, so I'm sharing my thoughts based on my 20 years of experience as a software developer. An example of a useful commit message What's the point of a commit message? Helps your code reviewer approach the change Communicates changes to teammates, downstream clients, and end-users Facilitates future bug investigations Provides information to development tools Organizing information in a commit message Put the most important information first Use headings to structure long commit messages What should the commit message include? A descriptive title A summary of how the change impacts clients and end-users The motivation for the change Breaking changes External references Justifications for new dependencies Cross-references to issues or other changes Summaries of bugs or external references Testing instructions Testing limitations What you learned Alternative solutions you consi...

121. Can You Spot the Passive Voice?

Source: Refactoring English | Published: Fri, 31 Jan 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/exercises/recognize-passive-voice/>

This exercise tests your ability to recognize the passive voice in the context of developing software. For each sentence, click: Active ? if the sentence has no passive voice (active voice only) Passive ? if the sentence contains passive voice Start Does this sentence contain passive voice? Passive ? Active ? Next Question Accuracy: 0 %

122. Reproducing Hacker News writing style fingerprinting

Source: antirez.com | Published: Wed, 16 Apr 2025 15:53:16 +0200

Link: <http://antirez.com/news/150>

About three years ago I saw a quite curious and interesting post on Hacker News. A student, Christopher Tarry, was able to use cosine similarity against a vector of top words frequencies in comments, in order to detect similar HN accounts ? and, sometimes, even accounts actually controlled by the same user, that is, fake accounts used to uncover the identity of the writer. This is the original post: <https://news.ycombinator.com/item?id=33755016> I was not aware, back then, of Burrows-Delta method for style detection: it seemed kinda magical that you just needed to normalize a frequency vector of top words to reach such quite remarkable results. I read a few wikipedia pages and took mental note of it. Then, as I was working with Vectors for Redis I remembered about this post, searched the web only to discover that the original page was gone and that the author, in the original post and website, didn't really explained very well how the data was processed, the top words extracted (and, especially, how many were used) and so forth. I thought I could reproduce the work with Vector Sets, once I was done with the main

work. Now the new data type is in the release candidate, and I found s...

123. Vector Sets are part of Redis

Source: antirez.com | Published: Thu, 03 Apr 2025 20:01:20 +0200

Link: <http://antirez.com/news/149>

Yesterday we finally merged vector sets into Redis, here you can find the README that explains in detail what you get:

<https://github.com/redis/redis/blob/unstable/modules/vector-sets/README.md> The goal of the new data structure is, in short, to create a new ?Set alike? data type, similar to Sorted Sets, where instead of having a scalar as a score, you have a vector, and you can add and remove elements the Redis way, without caring about anything except the properties of the abstract data structure Redis implements, ask for elements similar to a given query vector (or a vector associated to some element already in the set), and so forth. But more about that later, a bit of background, first: From the path of the README itself, you can see the implementation is into ?modules?, but actually, Vector Sets are not a module, it?s a part of the Redis core, the thing is that I started developing them as a module, and later I suggested that the implementation should still use the modules API, in order to promote modularity of the internals of Redis, in order to have both the advantages: every Redis instance starting from Redis 8 will have Vector Sets as a native data type, and there are cl...

124. AI is useless, but it is our best bet for the future

Source: antirez.com | Published: Sun, 23 Mar 2025 09:13:53 +0100

Link: <http://antirez.com/news/148>

I used AI with success 5 minutes ago. Just five minutes ago, I was writing a piece of software and relied on AI for assistance. Yet, here I am, starting this blog post by telling you that artificial intelligence, so far, has proven somewhat useless. How can I make such a statement if AI was just so helpful a moment ago? Actually, there's no contradiction here if we clarify exactly what we mean. Here?s the thing: at this very moment, artificial intelligence can support me significantly. If I'm struggling with complicated code or need to understand an advanced scientific paper on math, I can turn to AI for clarity. It can help me generate an image for a project, make a translation, clean my YouTube transcript. Clearly, it?s practical and beneficial in these everyday tasks. However, except for rare, groundbreaking examples like AlphaFold ? Google's AI that significantly advanced our understanding of protein folding ? AI has yet to genuinely push forward human knowledge in a fundamental way. Aside from these few exceptional results, AI hasn?t (obviously) yet matched the capabilities of the very best human minds. If an AI system were at the same level as the brightest humans (and not b...

125. Blind to Disruption ? The CEOs Who Missed the Future

Source: Steve Blank | Published: Tue, 08 Jul 2025 13:00:00 +0000

Link: <https://steveblank.com/2025/07/08/blind-to-disruption-the-ceos-who-missed-the-future/>

How did you go bankrupt?? Two ways. Gradually, then suddenly.? Ernest Hemingway, The Sun Also Rises Every disruptive technology since the fire and the wheel have forced leaders to adapt or die. This post tells the story of what happened when 4,000 companies faced a disruptive technology and why only one survived. In the early 20th [?]

126. Why Investors Don?t Care About Your Business

Source: Steve Blank | Published: Tue, 01 Jul 2025 13:00:38 +0000

Link: <https://steveblank.com/2025/07/01/why-investors-dont-care-about-your-business/>

Founders with great businesses are often frustrated that they can?t raise money. Here?s why. I?ve been having coffee with lots of frustrated founders (my students and others) bemoaning most VCs won?t even meet with them unless they have AI in their fundraising pitch. And the AI startups they see are getting valuations that appear nonsensical. [?]

127. Lean Launchpad at Stanford ? 2025

Source: Steve Blank | Published: Tue, 24 Jun 2025 13:00:26 +0000

Link: <https://steveblank.com/2025/06/24/lean-launchpad-at-stanford-2025/>

The PowerPoints embedded in this post are best viewed on steveblank.com We just finished the 15th annual Lean LaunchPad class at Stanford. The class had gotten so popular that in 2021 we started teaching it in both the winter and spring sessions. During the 2025 spring quarter the eight teams spoke to 935 potential customers, beneficiaries [?]

128. Does the software industry learn?

Source: Tom Renner on My place to put things | Published: Mon, 24 Jan 2022 00:00:00 +0000

Link: <https://tomrenner.com/posts/does-the-software-industry-learn/>

?Institutional knowledge? - the information that a company collectively knows - is a familiar concept to anyone involved in hiring processes. When an individual leaves who has knowledge the organisation needs, companies will organise offboarding sessions to keep that knowledge within the institution. Maybe they?ll even try to hire someone with similar experience. Lots of companies similarly try to optimise for ?Institutional learning?, especially smaller firms. This makes a lot of sense - smaller companies don?t have the resources to buy in extra experience, so they focus on rapidly expanding the experience of their existing employees. It also fits really well with the agile philosophy of fast development cycles to maximise your knowledge about customers? needs.

129. Should we welcome or fear the Metaverse?

Source: Tom Renner on My place to put things | Published: Thu, 04 Nov 2021 00:00:00 +0000

Link: <https://tomrenner.com/posts/should-we-welcome-or-fear-the-metaverse/>

Kit Wilson writes in The Spectator about Facebook?s new venture into the Metaverse, a concept that most of us probably hadn?t heard of until last week. To layout the roadmap for what our journey into this new digital reality might look like, Kit joins the podcast along with Tom Renner, a software engineer for NavVis. (12:55) .iframe-container { position: relative; overflow: hidden; width: 100%; padding-top: 100%; } .responsive-iframe { position: absolute; top: 0; left: 0; bottom: 0; right: 0; width: 100%; height: 95; display: block; background-color: transparent; }

130. DORA? I barely know her!

Source: Tom Renner on My place to put things | Published: Wed, 28 Apr 2021 00:00:00 +0000

Link: <https://tomrenner.com/posts/dora-i-barely-know-her/>

Coming to grips with DevOps metrics In my team we have been considering ways to monitor our own performance, and finding some ways to contextualise our ongoing process and quality improvements. Like many other teams, we?ve landed on the DORA metrics as a good way of doing this. These four key metrics are an easy way to understand what adjectives like ?maintainable?, ?reliable?, and ?efficient? mean in practice when applied to software and teams, and the provide a way of comparing team performance across teams and over time.

131. Beyond Borgmann: Single-Task Tools and the Future of Meaningful Technology

Source: Simone Salis | Published: Mon, 16 Dec 2024 17:43:54 GMT

Link: <https://simone.org/single-task-tools/>

Single-task tools complement modern devices by combining technological capability with intentional constraints. From offline cameras to e-ink readers, these tools preserve human agency while embracing progress. How thoughtful friction and focused design create more meaningful engagement.

132. The Productivity Paradox: Why Doing More Leads to Less

Source: Simone Salis | Published: Thu, 21 Nov 2024 14:30:35 GMT

Link: <https://simone.org/productivity-paradox/>

In a world obsessed with optimization, what if the most productive strategy is to stop maximizing productivity? Discover how strategic inefficiency might be the key to reclaiming your work, your time, and your sanity.

133. In Praise of Inconvenience: The Hidden Costs of a Convenient World

Source: Simone Salis | Published: Tue, 12 Nov 2024 16:01:02 GMT

Link: <https://simone.org/inconvenience/>

Explore how devotion to convenience reshapes human experience, and why embracing the friction of inconvenience might restore depth and meaning to your life.

134. Resource Management

Source: Daniel De Laney | Published: Thu, 15 Jun 2017 00:00:00 +0000

Link: <http://danieldelaney.net/rm/>

?We had to display a lot of information at once and help them make sense of it. There are many ways to do that wrong and only a few ways to do it right.? The Challenge The client came to us with a problem: They had paid for a custom tool that didn?t do what they needed. They had expected ease and efficiency. What they got was frustration and unnecessary manual effort. What they needed was a tool that could help them plan. How much more work could their team take on right now? Did their team have the capacity to complete their projects next October? My job was to design a tool that worked for them, not against them. My Process My first step was to understand where everything had gone wrong. Who were their users? What were their goals? From there I could see the right way forward to the tool they wished they knew how to ask for. A problem well-understood is half solved. That?s why research is the most important part of my process. I discovered that humans were doing hours of tedious mental labor that a computer could do for them. Their software wasn?t presenting the right information at the right time. Smart planning requires seeing the big picture and specific details at the same t...

135. Slow social media

Source: Herman's blog | Published: 2025-10-02T13:41:18.903926+00:00

Link: <https://herman.bearblog.dev/slow-social-media/>

People often assume that I hate social media. And they'd be forgiven for believing that, since I am overtly critical of current social media platforms and the effects they have on individuals and society; and deleted all of my social media accounts back in 2019 . However, the underlying concept of social media is something I resonate with: Stay connected with the people you care about. It's just that the current form of social media is bastardised, and not social at all. Instead of improving relationships and fostering connection, they're advertisement-funded content mills which are explicitly designed and continually refined to keep you engaged, lonely, and unhappy. And once TikTok figured out that short-form video with a recommendation engine is digital crack, all other social media platforms quickly sprang into action to copy their secret sauce. Meta basically turned Instagram and Facebook from 'connecting with friends' into 'doom-scrolling random content'. Even Pinterest is starting to look like TikTok! They followed user engagement, but not the underlying preferences of their users. I posit that any for-profit social media will eventually degrade into recommendation media ove...

136. Weekly Update 482

Source: Troy Hunt | Published: Tue, 16 Dec 2025 22:52:14 GMT

Link: <https://www.troyhunt.com/weekly-update-482/>

Perhaps it's just the time of year where we all start to wind down a bit, or maybe I'm just tired after another massive 12 months, but this week's vid is way late. Ok, going away to the place that had just been breached

137. Processing 630 Million More Pwned Passwords, Courtesy of the FBI

Source: Troy Hunt | Published: Fri, 12 Dec 2025 21:29:39 GMT

Link: <https://www.troyhunt.com/processing-630-million-more-pwned-passwords-courtesy-of-the-fbi/>

The sheer scope of cybercrime can be hard to fathom, even when you live and breathe it every day. It's not just the volume of data, but also the extent to which it replicates across criminal actors seeking to abuse it for their own gain, and to our

138. Weekly Update 481

Source: Troy Hunt | Published: Fri, 05 Dec 2025 07:14:33 GMT

Link: <https://www.troyhunt.com/weekly-update-481/>

Twelve years (and one day) since launching Have I Been Pwned, it's now a service that Charlotte and I live and breathe every day. From the first thing every morning to the last thing each day, from holidays to birthdays, in sickness and in heal... wait a minute

139. Yarn Spinner

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/projects/ys/>

(No summary)

140. Posts

Source: Dr Paris Buttfield-Addison | Published: Mon, 01 Jan 0001 00:00:00 +0000

Link: <https://hey.paris/all-posts/>

archives

141. Tas Game Makers 2026

Source: Dr Paris Buttfield-Addison | Published: Sun, 08 Feb 2026 00:00:00 +0000

Link: <https://hey.paris/posts/tasgm2026/>

We just had the AGM for Tasmanian Game Makers, Inc. (TasGM) in Launceston on Friday, and I was thrilled to (once again, somehow) be elected as President! For the fourth year in a row, I'm really grateful for the trust of our members, and I'm excited to work with the new board to keep pushing the video game industry in Tasmania forward. The new board is: President: Paris B-A Vice-Presidents: Michael Hinz, Leo Febey Secretary: Tim Nugent Treasurer: Jon Manning General Committee: Jason Imms, Dean Baron, Lindsay Wells, Lindsay Mansfield, Rhiannon Ross, Alonso Zevallos Fournier, Brie Ratten, and Jerome Phelps With Oliver Potter continuing as our Screen Tasmania Observer.

142. Attack of the SaaS clones

Source: Martin Alderson | Published: Fri, 13 Feb 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/attack-of-the-clones/?utm_source=rss

I cloned Linear's UI and core functionality using Claude Code in about 20 prompts. Here's what that means for SaaS companies.

143. Why I'm building my own CLIs for agents

Source: Martin Alderson | Published: Mon, 29 Dec 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/why-im-building-my-own-clis-for-agents/?utm_source=rss

MCP tools eat thousands of tokens. A simple CLI with instructions in your CLAUDE.md file uses 71 tokens and works brilliantly.

144. Travel agents took 10 years to collapse. Developers are 3 years in.

Source: Martin Alderson | Published: Sat, 27 Dec 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/travel-agents-developers/?utm_source=rss

Travel agents are the classic example of an industry killed by the internet. Software engineering is facing the same disruption, but the timeline is compressed.

145. New Article on BSD Cafe Journal: WordPress on FreeBSD with BastilleBSD

Source: IT Notes | Published: Mon, 21 Jul 2025 09:30:00 +0200

Link: <https://it-notes.dragas.net/2025/07/21/new-article-wordpress-on-freebsd-bastillebsd-on-bsd-cafe-journal/>

A new article on running WordPress on FreeBSD with BastilleBSD has been published on the BSD Cafe Journal, plus a small update on future technical content.

146. Make Your Own Backup System ? Part 1: Strategy Before Scripts

Source: IT Notes | Published: Fri, 18 Jul 2025 09:00:00 +0200

Link: <https://it-notes.dragas.net/2025/07/18/make-your-own-backup-system-part-1-strategy-before-scripts/>

When a datacenter fire threatened 142 of my servers, my backup strategy had them back online in hours. This post shares my personal philosophy on creating a resilient system, focusing on the crucial planning that must happen before you write a single script.

147. How to install FreeBSD on providers that don't support it with mfsBSD

Source: IT Notes | Published: Wed, 02 Jul 2025 09:30:00 +0200

Link: https://it-notes.dragas.net/2025/07/02/install_freebsd_providers_mfsbsd/

A practical guide to installing FreeBSD on hosting providers that don't officially support it, using the powerful mfsBSD tool for rescue mode installations.

148. Liveness analysis with Datalog

Source: Max Bernstein's Blog | Published: Wed, 13 Aug 2025 00:00:00 +0000

Link: https://bernsteinbear.com/blog/liveness-datalog/?utm_source=rss

After publishing Linear scan register allocation on SSA , I had a nice call with Waleed Khan where he showed me how to Datalog. He thought it might be useful to try implementing liveness analysis as a Datalog problem. We started off with the Wimmer2010 CFG example from that post, sketching out manually which variables were live out of each block: R10 out of B1, R12 out of B2, etc. The graph from Wimmer2010 has come back! Remember, we're using block arguments instead of phis, so B1(R10, R11) defines R10 and R11 before the first instruction in B1. Then we tried to formulate liveness as a Datalog relation. Liveness is normally (at least for me) defined in terms of two relations: live-in and live-out. Live-out is ?what is needed? from all of the successors of a block and live-in is the ?what is needed? summary for a block. So, in fake math notation: live-out(b) = union(live-in(s) for each successor s of b) live-in(b) = (live-out(b) + used(b)) - defined(b) where each of the component parts of that expression represent sets of variables: used(b) is the set of variables referenced as in-operands to instructions in a block defined(b) is the set of variables defined by instructions in a bl...

149. Linear scan register allocation on SSA

Source: Max Bernstein's Blog | Published: Wed, 13 Aug 2025 00:00:00 +0000

Link: https://bernsteinbear.com/blog/linear-scan/?utm_source=rss

Much of the code and education that resulted in this post happened with Aaron Patterson . The fundamental problem in register allocation is to take an IR that uses a virtual registers (as many as you like) and rewrite it to use a finite amount of physical registers and stack space 1 . This is an example of a code snippet using virtual registers: add R1, R2 -> R3 add R1, R3 -> R4 ret R4 And here is the same example after it has been passed through a register allocator (note that Rs changed to Ps): add Stack[0], P0 -> P1 add Stack[0], P1 -> P0 ret Each virtual register was assigned a physical place: R1 to the stack, R2 to P0, R3 to P1, and R4 also to P0 (since we weren't using R2 anymore). People use register allocators like they use garbage collectors: it's an abstraction that can manage your resources for you, maybe with some cost. When writing the back-end of a compiler, it's probably much easier to have a separate register-allocator-in-a-box than manually managing variable lifetimes while also considering all of your different target architectures. How do JIT compilers do register allocation? Well, everyone knows? that every JIT does its own variant of linear scan? 2 . This bo...

150. Good Ideas in Computer Science

Source: Daniel Hooper | Published: Sun, 21 Apr 2024 00:00:00 +0000

Link: <https://danielchasehooper.com/posts/good-ideas-in-cs/>

Ideas every programmer likes and why Garbage Collection and Object Oriented Programming don't count Read the whole article on danielchasehooper.com ?

151. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 13 Feb 2026 17:17:50 +0000

Link: <https://www.filfre.net/2026/02/this-week-on-the-analog-antiquarian/>

Chapter 13: The Shades of the Earth

152. A Looking Glass Half Empty, Part 1: Just Lookin? for a Hit

Source: The Digital Antiquarian | Published: Fri, 03 Oct 2025 16:04:48 +0000

Link: <https://www.filfre.net/2025/10/a-looking-glass-half-empty-part-1-just-lookin-for-a-hit/>

This article tells part of the story of Looking Glass Studios. There was some discussion about it: ?Wow, gosh, it?d sure be nice if we were making more money and selling more copies so we could do crazy games of the type we want, as opposed to having to worry about how we?re going to [?]

153. Thank you for being annoying

Source: Experimental History | Published: Tue, 30 Sep 2025 14:35:46 GMT

Link: <https://www.experimental-history.com/p/thank-you-for-being-annoying>

OR: whack 'em if you got 'em

154. Blog Extravaganza 2025: the winners

Source: Experimental History | Published: Tue, 16 Sep 2025 14:14:58 GMT

Link: <https://www.experimental-history.com/p/blog-extravaganza-2025-the-winners>

YEEHAW

155. Use this magic bullet to shoot yourself in the foot

Source: Experimental History | Published: Tue, 02 Sep 2025 16:48:01 GMT

Link: <https://www.experimental-history.com/p/use-this-magic-bullet-to-shoot-yourself>

OR: Borg vibes

156. Good SaaS Products Are Like UNIX Tools

Source: Keygen Blog | Published: Fri, 12 Nov 2021 06:00:00 GMT

Link: <https://keygen.sh/blog/good-saas-products-are-like-unix-tools/>

Reusable, composable, single responsibility ? these are all things good SaaS products have in common with UNIX command line tools.

157. How to License and Distribute a Private Docker Image

Source: Keygen Blog | Published: Mon, 13 Dec 2021 06:00:00 GMT

Link: <https://keygen.sh/blog/how-to-license-and-distribute-a-private-docker-image/>

Learn how to license and distribute commercial Docker images using Keygen's software licensing and distribution API.

158. How to License and Distribute a Private Ruby Gem

Source: Keygen Blog | Published: Thu, 16 Dec 2021 06:00:00 GMT

Link: <https://keygen.sh/blog/how-to-license-and-distribute-a-private-rubygem/>

Learn how to license and distribute commercial Ruby gems with Keygen's software licensing and distribution API.

159. Rohit Chess

Source: Grant Slatton's Blog | Published: Mon, 10 Mar 2025 06:20:00 +0000

Link: <https://grantslatton.com/rohit-chess>

Fun little board game

160. AI software I'm surprised doesn't exist yet

Source: Grant Slatton's Blog | Published: Thu, 20 Feb 2025 08:04:37 +0000

Link: <https://grantslatton.com/ai-software-nonexistence>

Please make it

161. Technocapital

Source: Grant Slatton's Blog | Published: Thu, 6 Feb 2025 23:21:28 +0000

Link: <https://grantslatton.com/technocapital>

An essay on post-AI economics

162. Branch Encryption

Source: Cats with power tools | Published: Tue, 12 Mar 2024 22:49:00 GMT

Link: <https://blog.pixelmelt.dev/branch-encryption/>

Usage of branch encryption in various obfuscation contexts.

163. Cloning your Discord friends with Large Language Models

Source: Cats with power tools | Published: Sat, 20 Jan 2024 16:44:00 GMT

Link: <https://blog.pixelmelt.dev/clone-your-friends/>

Did it work well and is it useful in any way? No. Is it hilarious to mess around with and does it mimic them well? Yes.

164. Seesaws for giants

Source: Aresluna | Published: 2025-03-18T00:00:00.000Z

Link: <https://aresluna.org/seesaws-for-giants>

Chasing Chicago's movable bridges. (New version of a photoessay originally posted in August 2014. 900 words, 50 photos.)

165. The hardest working font in Manhattan

Source: Aresluna | Published: 2025-02-14T00:00:00.000Z

Link: <https://aresluna.org/the-hardest-working-font-in-manhattan>

A story of a 150-year-old font you have never heard of ? and one you probably saw earlier today. (6,100 words, 600 photos.)

166. A Codeless Ecosystem, or hacking beyond vibe coding

Source: Anil Dash | Published: 2026-01-27T00:00:00Z

Link: <https://anildash.com/2026/01/27/codeless-ecosystem/>

There's been a remarkable leap forward in the ability to orchestrate coding bots, making it possible for ordinary creators to command dozens of AI bots to build software without ever having to directly touch code. The implications of this kind of evolution are potentially extraordinary, as outlined in that first set of notes about what we could call "codeless" software. But now it's worth looking at the larger ecosystem to understand where all of this might be headed. "Frontier minus six" One idea that's come up in a host of different conversations around codeless software, both from supporters and skeptics, is how these new orchestration tools can enable coders to control coding bots that aren't from the Big AI companies. Skeptics say, "won't everyone just use Claude Code, since that's the best coding bot?" The response that comes up is one that I keep articulating as "frontier minus six", meaning the idea that many of the open source or open-weight AI models are often delivering results at a level equivalent to where frontier AI models were six months ago. Or, sometimes, where they were 9 months or a year ago. In any of these cases, these are still damn good results! These level...

167. New York Tech at 30: the Crossroads

Source: Anil Dash | Published: 2026-02-04T00:00:00Z

Link: <https://anildash.com/2026/02/03/hye-tech-30/>

This past week, over a series of events, the New York tech community celebrated the 30th anniversary of a nebulous idea described as ?Silicon Alley?, the catch-all term for our greater collective of creators and collaborators, founders and funders, inventors and investors, educators and entrepreneurs and electeds, activists and architects and artists. Some of the parties or mixers have been typical industry affairs, the usual glad-handing about deal-making and pleasantries. But a lot have been deeper, reflecting on what?s special and meaningful about the community we?ve built in New York. Steven Rosenbaum?s reflection on the anniversary

captures this well from someone who's been there, and Leo Schwartz's piece for Fortune covers the more conventional business angle. Beyond the celebrations, though, I wanted to reflect on a number of the deeper conversations I've had over these last few days. These are conversations grounded in the reality of where our country and city are today, far beyond spaces where wealthy techies are going to parties and celebrating each other. The hard questions raised in these conversations are the ones that determine where this community goes in the future...

168. There's no such thing as "tech" (Ten years later)

Source: Anil Dash | Published: 2026-02-06T00:00:00Z

Link: <https://anildash.com/2026/02/06/no-such-thing-as-tech/>

Ten years ago I wrote that there is no ?technology industry? . It's more true than ever. There is no ?tech?. There's no such thing as ?a FAANG company?. There is almost nothing in common between the very largest tech companies and the next several hundred biggest companies that happen to create tech platforms. Whatever shorthand we use for the biggest tech companies, they almost never have much in common?whether it's how they make money, what products they make, how they make decisions, who leads them, or what drives their cultures. It's important to make these distinctions because the false categorization of wildly dissimilar organizations into one grouping leads to absurdly inappropriate decisions being made. Let's look at some simple examples to understand why. Take the once-ubiquitous shorthand of ?FAANG? to describe big tech. (It stood, at one time, for Facebook, Amazon, Apple, Netflix and Google. Then Facebook became Meta and Google became Alphabet and Microsoft became upset about not being included, and people started trying to use other more unwieldy, less-popular sobriquets.) This abbreviation still persists because of the mindset it represents, and it is still useful in...

169. The Small Web is Tricky to Find

Source: matduggan.com | Published: Fri, 13 Feb 2026 12:50:36 GMT

Link: <https://matduggan.com/the-small-web-is-tricky-to-find/>

One of the most common requests I've gotten from users of my little Firefox extension(<https://timewasterpro.xyz>) has been more options around the categories of websites that you get returned. This required me to go through and parse the website information to attempt to put them into

170. Simple Python Script for FTP Uploads

Source: matduggan.com | Published: Thu, 12 Jun 2025 10:00:55 GMT

Link: <https://matduggan.com/simple-python-script-for-ftp-uploads/>

I have a AX3000 wireless router that I slapped a cheap USB drive on that I wanted to automate the upload of files to. I ended up writing a simple python script to do the process for me that should work for any of these routers that expose a local FTP port.

171. Write your own Ghost Theme

Source: matduggan.com | Published: Thu, 05 Jun 2025 08:52:50 GMT

Link: <https://matduggan.com/write-your-own-ghost-theme/>

I struggled to make a Ghost CMS theme that worked the way I wanted, so I made my own. Turned out to be a surprisingly easy experience.

172. Font Ligatures for your Code Editor and Terminal

Source: Miguel Grinberg's Blog | Published: Thu, 13 Mar 2025 16:10:16 GMT

Link: <https://blog.miguelgrinberg.com/post/font-ligatures-for-your-code-editor-and-terminal>

The other day I shared my code editor window during a video call with a client, and he immediately did a double take. "How did you make your code look like that?", he wanted to know. Have a look at the following screenshot and see if like my client, anything calls your attention: If you are wondering what secret keyboard trick allows you to type ?, ?, ?, ? and other unusual characters, and then how to get Python and other programming languages to accept them then read on, as this is much easier than you think and you can do it too!

173. Possibly all the ways to get loop-finding in graphs wrong

Source: Simon Tatham's quasi-blog | Published: 2024-09-09T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/findloop/>

The history of all the mistakes I've made trying to detect loops in graphs while writing my puzzle collection.

174. Code review antipatterns

Source: Simon Tatham's quasi-blog | Published: 2024-08-21T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/code-review-antipatterns/>

A taxonomy of frustrating code review behaviours, and some thoughts about them.

175. Symbiosisware

Source: Simon Tatham's quasi-blog | Published: 2024-08-06T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/symbiosisware/>

A definition and discussion of my term ?symbiosisware?, describing software consciously designed for only yourself to use.

176. The Ascent to Sandia Crest

Source: computers are bad | Published: 28 Sep 2025 00:00:00 UT

Link: <https://computer.rip/2025-09-28-the-ascent-to-sandia-crest-part-1.html>

The Rotary Club will take immediate action on the Ellis ranch loop project. The Rotarians reached this decision at their weekly luncheon, held yesterday at the Albarado hotel. The club's plan is not merely to give Albuquerque a good, short road to the Ellis ranch... They embrace the building of a seventy-mile scenic loop. 1 Many Western cities are defined, in part, by their mountains. Those moving from town to town often comment on the disorientation, the disruption, caused by a change in the city's relation to the peaks. If you have ever lived in a place with the mountains on the west, and then a place with the mountains on the east, you will know what I mean. We get used to living in the shadows of mountains. One of the appeals of mountains, perhaps the source of their mysticism, is inaccessibility. Despite their close proximity to Albuquerque, steep slopes and difficult terrain kept the Sandias a world apart from the city, even to this day. Yet we have always been driven to climb, to ascend to the summit. Humans climb mountains not only as a matter of individual achievement, but also as a matter of infrastructure. Whether the inaccessibility of mountain peaks is a good thing or...

177. Nationaal Agentschap Baanbrekende Innovatie, in EU of niet?

Source: Bert Hubert's writings | Published: Fri, 19 Dec 2025 11:00:00 +0100

Link: <https://berthub.eu/articles/posts/nadi-nabi-in-europa-of-niet/>

Ik ben eigenlijk op vakantie, maar toch een korte blogpost. Het is (in)formatietijd, en er worden plannen gemaakt. Eerder dit jaar kwam het ministerie van Economische Zaken met het idee een Nationaal Agentschap Disruptieve Innovatie op te richten, zoals ook al bestaat in Duitsland, en het meest beroemd in Amerika. Ik schreef hier eerder al een artikelje over met meer details: Klassieke manieren om wetenschap en innovatie te stimuleren zijn erg ongeschikt voor disruptieve innovatie, idee^on die misschien helemaal niet gaan werken, maar als ze dat wel doen de wereld kunnen veranderen.

178. Keynote opening Digital Commons EDIC: Moving beyond the Digital Uncommons

Source: Bert Hubert's writings | Published: Sun, 14 Dec 2025 11:00:00 +0100

Link: <https://berthub.eu/articles/posts/keynote-opening-digital-commons-edic-eu/>

Last Thursday, 11th of December, saw the launch of the new Digital Commons European Digital Infrastructure Consortium. In attendance were delegations from the launching member states (and the observers). Also present were the many forefathers (and mothers) of this initiative. Of particular note, the launch also included demonstrations of LaSuite software from France, as well as the OpenDesk software from Germany. The Dutch government showcased their amalgamated suite MijnBureau, which joins parts of the German and French initiatives.

179. Makelaarstaal over onze belastinggeheimen

Source: Bert Hubert's writings | Published: Tue, 02 Dec 2025 10:03:45 +0100

Link: <https://berthub.eu/articles/posts/makelaarstaal-over-onze-belastinggeheimen/>

Je vraagt de makelaar of je nieuwe droomhuis last heeft van lekkage, en hij vertelt je dat de dakkapel niet lekt. Alle alarmbellen moeten nu afgaan - want waarom krijg ik zo'n specifiek antwoord op deze vraag? Lekt het ergens anders wel? De belastingdienst is bezig al hun email en documenten door te sturen naar Amerikaanse servers. En dan vraag je je tegenwoordig natuurlijk af, gaan ze in Amerika dan allemaal meeën met wat wij en onze bedrijven aan het doen zijn?

180. How Twitter could (somewhat) fix their encrypted DMs

Source: Matthew Garrett | Published: Thu, 05 Jun 2025 13:18:48 GMT

Link: <https://mjj59.dreamwidth.org/71933.html>

As I wrote in my last post , Twitter's new encrypted DM infrastructure is pretty awful. But the amount of work required to make it somewhat better isn't large. When Juicebox is used with HSMs, it supports encrypting the communication between the client and the backend. This is handled by generating a unique keypair for each HSM. The public key is provided to the client, while the private key remains within the HSM. Even if you can see the traffic sent to the HSM, it's encrypted using the Noise protocol and so the user's encrypted secret data can't be retrieved. But this is only useful if you know that the public key corresponds to a private key in the HSM! Right now there's no way to know this, but there's worse - the client doesn't have the public key built into it, it's supplied as a response to an API request made to Twitter's servers. Even if the current keys are associated with the HSMs, Twitter could swap them out with ones that aren't, terminate the encrypted connection at their endpoint, and then fake your query to the HSM and get the encrypted data that way. Worse, this could be done for specific targeted users, without any indication to the user that this has happened, m...

181. Twitter's new encrypted DMs aren't better than the old ones

Source: Matthew Garrett | Published: Thu, 05 Jun 2025 11:02:47 GMT

Link: <https://mjj59.dreamwidth.org/71646.html>

(Edit: Twitter could improve this significantly with very few changes - I wrote about that here . It's unclear why they'd launch without doing that, since it entirely defeats the point of using HSMs) When Twitter[1] launched encrypted DMs a couple of years ago, it was the worst kind of end-to-end encrypted - technically e2ee, but in a way that made it relatively easy for Twitter to inject new encryption keys and get everyone's messages anyway. It was also lacking a whole bunch of features such as "sending pictures", so the entire thing was largely a waste of time. But a couple of days ago, Elon announced the arrival of "XChat", a new encrypted message platform built on Rust with (Bitcoin style) encryption, whole new architecture . Maybe this time they've got it right? tl;dr - no. Use Signal. Twitter can probably obtain your private keys, and admit that they can MITM you and have full access to your metadata. The new approach is pretty similar to the old one in that it's based on pretty straightforward and well tested cryptographic primitives, but merely using good cryptography doesn't mean you end up with a good solution. This time they've pivoted away from using the underlying cr...

182. Failing upwards: the Twitter encrypted DM failure

Source: Matthew Garrett | Published: Tue, 18 Mar 2025 23:58:52 GMT

Link: <https://mjj59.dreamwidth.org/71188.html>

Almost two years ago, Twitter launched encrypted direct messages. I wrote about their technical implementation at the time , and to the best of my knowledge nothing has changed. The short story is that the actual encryption primitives used are entirely normal and fine - messages are encrypted using AES, and the AES keys are exchanged via NIST P-256 elliptic curve asymmetric keys. The asymmetric keys are each associated with a specific device or browser owned by a user, so when you send a message to someone you encrypt the AES key with all of their asymmetric keys and then each device or browser can decrypt the message again. As long as the keys are managed appropriately, this is infeasible to break. But how do you know what a user's keys are? I also wrote about this last year - key distribution is a hard problem. In the

Twitter DM case, you ask Twitter's server, and if Twitter wants to intercept your messages they replace your key. The documentation for the feature basically admits this - if people with guns showed up there, they could very much compromise the protection in such a way that all future messages you sent were readable. It's also impossible to prove that they're not a...

183. How I like to install NixOS (declaratively)

Source: Michael Stapelbergs Website | Published: 2025-06-01T08:20:43+02:00

Link: <https://michael.stapelberg.ch/posts/2025-06-01-nixos-installation-declarative/>

For one of my network storage PC builds , I was looking for an alternative to Flatcar Container Linux and tried out NixOS again (after an almost 10 year break). There are many ways to install NixOS, and in this article I will outline how I like to install NixOS on physical hardware or virtual machines: over the network and fully declaratively. Introduction: Declarative? The term declarative means that you describe what should be accomplished, not how. For NixOS, that means you declare what software you want your system to include (add to config option environment.systemPackages , or enable a module) instead of, say, running apt install . A nice property of the declarative approach is that your system follows your configuration, so by reverting a configuration change, you can cleanly revert the change to the system as well. I like to manage declarative configuration files under version control, typically with Git. When I originally set up my current network storage build, I chose CoreOS (later Flatcar Container Linux) because it was an auto-updating base system with a declarative cloud-init config. Ways of installing NixOS Graphical Installer: Only for Desktops The NixOS manual?s ?...

184. My 2025 high-end Linux PC ?

Source: Michael Stapelbergs Website | Published: 2025-05-15T15:44:24+02:00

Link: <https://michael.stapelberg.ch/posts/2025-05-15-my-2025-high-end-linux-pc/>

Update (2025-09-07): The replacement CPU also died and I have given up on Intel. See Bye Intel, hi AMD! for more details on the AMD 9950X3D. Turns out my previous attempt at this build had a faulty CPU! With the CPU replaced, the machine now is stable and fast! ? In this article, I?ll go into a lot more detail about the component selection, but in a nutshell, I picked an Intel 285K CPU for low idle power, chose a 4TB SSD so I don?t have to worry about running out of storage quickly, and a capable nvidia graphics card to drive my Dell UP3218K 8K monitor . Components Which components did I pick for this build? Here?s the full list: Price Type Article 140 CHF Case Fractal Define 7 Compact Black Solid 155 CHF Power Supply Corsair RM850x 233 CHF Mainboard ASUS PRIME Z890-P 620 CHF CPU Intel Core Ultra 9 285K 120 CHF CPU fan Noctua NH-D15 G2 39 CHF Case fan Noctua NF-A14 PWM (140 mm) 209 CHF RAM 64 GB DDR5-6400 Corsair Vengeance (2 x 32GB) 280 CHF Disk 4000 GB Samsung 990 Pro 554 CHF GPU MSI GeForce RTX 3060 Ti GAMING X TRIO Total: 2350 CHF ?and the next couple of sections go into detail on how I selected these components. Case I have been a fan of Fractal cases for a couple of generati...

185. In praise of grobi for auto-configuring X11 monitors

Source: Michael Stapelbergs Website | Published: 2025-05-10T08:24:00+02:00

Link: <https://michael.stapelberg.ch/posts/2025-05-10-grobi-x11-monitor-autoconfig/>

I have recently started using the grobi program by Alexander Neumann again and was delighted to discover that it makes using my fiddly (but wonderful) Dell 32-inch 8K monitor (UP3218K) monitor much more convenient ? I get a signal more quickly than with my previous, sleep-based approach. Previously, when my PC woke up from suspend-to-RAM, there were two scenarios: The monitor was connected. My sleep program would power on the monitor (if needed), sleep a little while and then run xrandr(1) to (hopefully) configure the monitor correctly. The monitor was not connected, for example because it was still connected to my work PC. In scenario ?, or if the one-shot configuration attempt in scenario ? fails, I would need to SSH in from a different computer and run xrandr manually so that the monitor would show a signal: % DISPLAY=:0 xrandr --output DP-4 --mode 3840x4320 --panning 0x0+0+0 \ --output DP-2 --right-of DP-4 --mode 3840x4320 --panning 0x0+3840+0 Automatic monitor configuration with grobi I have now completely solved this problem by creating the following ~/.config/grobi.conf file: rules : - name : UP3218K outputs_connected : [DP-2, DP-4] # DP-4 is left, DP-2 is right configure...

