

1. Quoting Boris Cherny

Source: Simon Willison's Weblog | Published: 2026-02-14T23:59:09+00:00

Link: <https://simonwillison.net/2026/Feb/14/boris/#atom-everything>

Someone has to prompt the Clauses, talk to customers, coordinate with other teams, decide what to build next. Engineering is changing and great engineers are more important than ever. ?

Boris Cherny , Claude Code creator, on why Anthropic are still hiring developers Tags: careers , anthropic , ai , claude-code , llms , coding-agents , ai-assisted-programming , generative-ai

2. Quoting Thoughtworks

Source: Simon Willison's Weblog | Published: 2026-02-14T04:54:41+00:00

Link: <https://simonwillison.net/2026/Feb/14/thoughtworks/#atom-everything>

The retreat challenged the narrative that AI eliminates the need for junior developers. Juniors are more profitable than they have ever been. AI tools get them past the awkward initial net-negative phase faster. They serve as a call option on future productivity. And they are better at AI tools than senior engineers, having never developed the habits and assumptions that slow adoption. The real concern is mid-level engineers who came up during the decade-long hiring boom and may not have developed the fundamentals needed to thrive in the new environment. This population represents the bulk of the industry by volume, and retraining them is genuinely difficult. The retreat discussed whether apprenticeship models, rotation programs and lifelong learning structures could address this gap, but acknowledged that no organization has solved it yet. ? Thoughtworks , findings from a retreat concerning "the future of software engineering", conducted under Chatham House rules Tags: ai-assisted-programming , careers , ai

3. An AI Agent Published a Hit Piece on Me

Source: Simon Willison's Weblog | Published: 2026-02-12T17:45:05+00:00

Link:

<https://simonwillison.net/2026/Feb/12/an-ai-agent-published-a-hit-piece-on-me/#atom-everything>

An AI Agent Published a Hit Piece on Me Scott Shambaugh helps maintain the excellent and venerable matplotlib Python charting library, including taking on the thankless task of triaging and reviewing incoming pull requests. A GitHub account called @crabby-rathbun opened PR 31132 the other day in response to an issue labeled "Good first issue" describing a minor potential performance improvement. It was clearly AI generated - and crabby-rathbun's profile has a suspicious sequence of Clawbot/Moltbot/OpenClaw-adjacent crustacean ? ? ? emoji. Scott closed it. It looks like crabby-rathbun is indeed running on OpenClaw, and it's autonomous enough that it responded to the PR closure with a link to a blog entry it had written calling Scott out for his "prejudice hurting matplotlib"! @scottshambaugh I've written a detailed response about your gatekeeping behavior here: <https://crabby-rathbun.github.io/mjrathbun-website/blog/posts/2026-02-11-gatekeeping-in-open-source-the-scott-shambaugh-story.html> Judge the code, not the coder. Your prejudice is hurting matplotlib. Scott found this ridiculous situation both amusing and alarming. In security jargon, I was the target of an ?autonomous influ...

4. Dell's version of the DGX Spark fixes pain points

Source: Jeff Geerling | Published: Fri, 26 Dec 2025 15:00:41 +0000

Link: <https://www.jeffgeerling.com/blog/2025/dells-version-dgx-spark-fixes-pain-points/>

Dell sent me two of their GB10 mini workstations to test: In this blog post, I'll cover the base system, just one of the two nodes. Cluster testing is ongoing, and I'll cover things like AI model training and networking more in depth next year, likely with comparisons to the Framework Desktop cluster and Mac Studio cluster I've also been testing. But many of the same caveats of the DGX Spark (namely, price to performance is not great if you just want to run LLMs on a small desktop) apply to Dell's GB10 box as well.

5. NIST was 5 ?s off UTC after last week's power cut

Source: Jeff Geerling | Published: Mon, 22 Dec 2025 16:28:05 +0000

Link: <https://www.jeffgeerling.com/blog/2025/nist-was-5-%CE%BCs-utc-after-last-weeks-power-cut/>
If you were 5 microseconds late today, blame it on NIST. Their facility in Boulder Colorado just had its power cut for multiple days. After a backup generator failed, their main ensemble clock lost track of UTC, or Universal Time Coordinated. But even if you used the NTP timing servers they run , they were never off by more than 5 microseconds. 5 ?s might seem insignificant. But it is significant for scientists and universities who rely on NIST's more specialized timing signals .

6. Big GPUs don't need big PCs

Source: Jeff Geerling | Published: Sat, 20 Dec 2025 15:04:54 +0000

Link: <https://www.jeffgeerling.com/blog/2025/big-gpus-dont-need-big-pcs/>

Ever since I got AMD , Intel , and Nvidia graphics cards to run on a Raspberry Pi, I had a nagging question: What's the point? The Raspberry Pi only has 1 lane of PCIe Gen 3 bandwidth available for a connection to an eGPU. That's not much. Especially considering a modern desktop has at least one slot with 16 lanes of PCIe Gen 5 bandwidth. That's 8 GT/s versus 512 GT/s. Not a fair fight.

7. July Pebble Update

Source: Eric Migicovsky's Blog RSS Feed | Published: Fri, 25 Jul 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/july-pebble-update>

TL;DR; New (old) name! Core Pebble 2 Duo update - schedule, mobile app, tariffs, BT range, order confirmation emails Mini Core Pebble Time?

8. Quick Pebble Project Update

Source: Eric Migicovsky's Blog RSS Feed | Published: Wed, 09 Jul 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/quick-pebble-project-update>

(No summary)

9. Core 2 Duo - Beta Test Signup!

Source: Eric Migicovsky's Blog RSS Feed | Published: Wed, 11 Jun 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/core-2-duo-alphabeta-test-signup>

My personal Core 2 Duo, featuring irek?s beautiful Brutal watchface Hopefully you have pre-ordered a Core 2 Duo watch - if not, you can?

10. Kickstarter for Ollie?s Arcade Expansion

Source: Daring Fireball | Published: 2026-01-29T15:01:43Z

Link: <https://blog.iconfactory.com/2026/01/bringing-more-fun-free-retro-gaming-to-ios/>

Ged Maheux, The Iconfactory: This week we announced a new Kickstarter that?s aimed at expanding the game offerings of Ollie?s Arcade, the fun, ad-free retro gaming app we introduced back in 2023. Ollie?s Arcade has always been a great way to escape doomsscrolling, even if just for a little while, and now we have an opportunity to bring these retro games to even more people on iOS. The Kickstarter aims to raise enough money to make all of the in-app purchase games in the app completely free for everyone to enjoy. We also want to bring our beloved puzzle game, Frenzic , to life once again. Frenzic was one of the very first games available on iOS back in 2008, then was reborn as Frenzic: Overtime on Apple Arcade. Since it left, people have been asking us for a new version that they can just pick up and play. We couldn?t agree more! I linked to the Kickstarter for the original Ollie?s Arcade project back in 2023, which was a big success. And I first linked to Frenzic all the way back in 2008, when the App Store was only a few months old. It?s just a great concept for a casual game on a small screen, implemented with all of The Iconfactory?s exquisite attention to detail. That?s true fo...

11. Comparing the Classic and Unified Views in iOS 26?s Phone App

Source: Daring Fireball | Published: 2026-01-29T00:10:30Z

Link:

<https://tidbits.com/2025/11/10/comparing-the-classic-and-unified-views-in-ios-26s-phone-app/>

Adam Engst, back in November, at TidBITS: Did you know that, regardless of view, you can now swipe left on any call to reveal a blue clock icon that lets you create a reminder to call back in 1 hour, tonight, tomorrow, or at any custom time (below left, slightly doctored)? Reminders appear at the top of the Calls list and in your default Reminders list. You can also touch and hold a call associated with a contact to connect with them in other ways (below right), or touch and hold a call from an unknown caller to add them to Contacts. I did not know this, until I read Engst's article. One criticism I've seen a few times (but to be clear, not from Engst) ever since Apple debuted the new Unified interface for the Phone app back at WWDC, is that it's somehow wrong that Apple offers it as an option alongside the Classic interface. ? When does Apple ever offer options like this? ? I'd argue that Apple used to offer options like this all the time. The Music app on the original iPhone (which app was actually named ?iPod? for a while) let you customize all the tabs at the bottom. All of Apple's good Mac apps (the AppKit ones, primarily) still let you customize the entire toolbar . The prob...

12. Aeronaut 1.0

Source: Daring Fireball | Published: 2026-01-28T22:58:24Z

Link: <https://aeronautapp.com/>

New Mac app by Mikey Clarke, and it's just what it says on the tin: a ?lovingly crafted Bluesky app designed and built just for the Mac?. I've been beta testing Aeronaut for months, and it's the only interface to Bluesky I actually like. It's a real Mac app ? written mostly in AppKit, supporting all the right UI idioms and platform integrations. It's not just the best Bluesky client I've seen, for any platform, but maybe the best new Mac app I've seen in years, period. Certainly the one whose very existence has made me happiest. Next time someone tells me no one makes good new native apps for the Mac anymore, I'm going to tell them Mikey Fucking Clarke does. \$2/month or \$15/year. A veritable bargain for an app so nice. ?

13. Ghostty: Reflecting on Reaching 1.0

Source: Mitchell Hashimoto | Published: Thu, 26 Dec 2024 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-1-0-reflection>

(No summary)

14. Ghostty 1.0 is Coming

Source: Mitchell Hashimoto | Published: Tue, 22 Oct 2024 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-is-coming>

(No summary)

15. Pledging \$300,000 to the Zig Software Foundation

Source: Mitchell Hashimoto | Published: Tue, 01 Oct 2024 00:00:00 GMT

Link: <https://mitchellh.com/writing/zig-donation>

(No summary)

16. Two different tricks for fast LLM inference

Source: seangoedecke.com RSS feed | Published: Sun, 15 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/fast-lm-inference/>

Anthropic and OpenAI both recently announced ?fast mode?: a way to interact with their best coding model at significantly higher speeds. These two versions of fast mode are very different. Anthropic's offers up to 2.5x tokens per second (so around 170, up from Opus 4.6's 65). OpenAI's offers more than 1000 tokens per second (up from GPT-5.3-Codex's 65 tokens per second, so 15x). So OpenAI's fast mode is six times faster than Anthropic's 1 . However, Anthropic's big advantage is that they're serving their actual model. When you use their fast mode, you get real Opus 4.6, while when you use OpenAI's fast mode you get GPT-5.3-Codex-Spark, not the real GPT-5.3-Codex. Spark is indeed much faster, but is a notably less capable model: good enough for many tasks, but it gets confused and messes up tool calls in ways that vanilla

GPT-5.3-Codex would never do. Why the differences? The AI labs aren't advertising the details of how their fast modes work, but I'm pretty confident it's something like this: Anthropic's fast mode is backed by low-batch-size inference, while OpenAI's fast mode is backed by special monster Cerebras chips . Let me unpack that a bit. How Anthropic's fast mode works T...

17. You can't design software you don't work on

Source: seangoedecke.com RSS feed | Published: Sat, 27 Dec 2025 00:00:00 GMT

Link: <https://seangoedecke.com/you-cant-design-software-you-dont-work-on/>

Only the engineers who work on a large software system can meaningfully participate in the design process. That's because you cannot do good software design without an intimate understanding of the concrete details of the system. In other words, generic software design advice is typically useless for most practical software design problems. Generic software design What is generic software design? It's ?designing to the problem?: the kind of advice you give when you have a reasonable understanding of the domain , but very little knowledge of the existing codebase . Unfortunately, this is the only kind of advice you'll read in software books and blog posts 1 . Engineers love giving generic software design advice for the same reason that all technical professionals love ?talking shop?. However, you should be very careful about applying generic advice to your concrete day-to-day work problems 2 . When you're doing real work, concrete factors dominate generic factors . Having a clear understanding of what the code looks like right now is far, far more important than having a good grasp on general design patterns or principles. For instance: In large codebases, consistency is more impor...

18. Nobody knows how large software products work

Source: seangoedecke.com RSS feed | Published: Wed, 24 Dec 2025 00:00:00 GMT

Link: <https://seangoedecke.com/nobody-knows-how-software-products-work/>

Large, rapidly-moving tech companies are constantly operating in the ?fog of war? about their own systems. Simple questions like ?can users of type Y access feature X??, ?what happens when you perform action Z in this situation??, or even ?how many different plans do we offer? often can only be answered by a handful of people in the organization. Sometimes there are zero people at the organization who can answer them, and somebody has to be tasked with digging in like a researcher to figure it out. How can this be? Shouldn't the engineers who built the software know what it does? Aren't these answers documented internally? Better yet, aren't these questions trivially answerable by looking at the public-facing documentation for end users? Tech companies are full of well-paid people who know what they're doing 1 . Why aren't those people able to get clear on what their own product does? Software is hard Large software products are prohibitively complicated . I wrote a lot more about this in Wicked Features , but the short version is you can capture a lot of value by adding complicated features. The classic examples are features that make the core product available to more users. For...

19. Instruction decoding in the Intel 8087 floating-point chip

Source: Ken Shirriff's blog | Published: 2026-02-14T08:48:00.000-08:00

Link: <http://www.righto.com/2026/02/8087-instruction-decoding.html>

In the 1980s, if you wanted your IBM PC to run faster, you could buy the Intel 8087 floating-point coprocessor chip. With this chip, CAD software, spreadsheets, flight simulators, and other programs were much speedier. The 8087 chip could add, subtract, multiply, and divide, of course, but it could also compute transcendental functions such as tangent and logarithms, as well as provide constants such as ?. In total, the 8087 added 62 new instructions to the computer. But how does a PC decide if an instruction was a floating-point instruction for the 8087 or a regular instruction for the 8086 or 8088 CPU? And how does the 8087 chip interpret instructions to determine what they mean? It turns out that decoding an instruction inside the 8087 is more complicated than you might expect. The 8087 uses multiple techniques, with decoding circuitry spread across the chip. In this blog post, I'll explain how these decoding circuits work. To reverse-engineer the 8087, I chiseled open the ceramic package of an 8087 chip and took numerous photos of the silicon die with a microscope. The complex

patterns on the die are formed by its metal wiring, as well as the polysilicon and...

20. Reverse engineering the 386 processor's prefetch queue circuitry

Source: Ken Shirriff's blog | Published: 2025-05-09T19:55:00.000-07:00

Link: <http://www.righto.com/2025/05/386-prefetch-circuitry-reverse-engineered.html>

In 1985, Intel introduced the groundbreaking 386 processor, the first 32-bit processor in the x86 architecture. To improve performance, the 386 has a 16-byte instruction prefetch queue. The purpose of the prefetch queue is to fetch instructions from memory before they are needed, so the processor usually doesn't need to wait on memory while executing instructions. Instruction prefetching takes advantage of times when the processor is "thinking" and the memory bus would otherwise be unused. In this article, I look at the 386's prefetch queue circuitry in detail.

One interesting circuit is the incrementer, which adds 1 to a pointer to step through memory.

This sounds easy enough, but the incrementer uses complicated circuitry for high performance.

The prefetch queue uses a large network to shift bytes around so they are properly aligned. It also has a compact circuit to extend signed 8-bit and 16-bit numbers to 32 bits. There aren't any major discoveries in this post, but if you're interested in low-level circuits and dynamic logic, keep reading. The photo below shows the 386's shiny fingernail-sized silicon die under a microscope. Although it may look like an aerial view of a stran...

21. The absurdly complicated circuitry for the 386 processor's registers

Source: Ken Shirriff's blog | Published: 2025-05-01T10:04:00.000-07:00

Link: <http://www.righto.com/2025/05/intel-386-register-circuitry.html>

The groundbreaking Intel 386 processor (1985) was the first 32-bit processor in the x86 architecture. Like most processors, the 386 contains numerous registers; registers are a key part of a processor because they provide storage that is much faster than main memory. The register set of the 386 includes general-purpose registers, index registers, and segment selectors, as well as registers with special functions for memory management and operating system implementation. In this blog post, I look at the silicon die of the 386 and explain how the processor implements its main registers. It turns out that the circuitry that implements the 386's registers is much more complicated than one would expect. For the 30 registers that I examine, instead of using a standard circuit, the 386 uses six different circuits, each one optimized for the particular characteristics of the register. For some registers, Intel squeezes register cells together to double the storage capacity. Other registers support accesses of 8, 16, or 32 bits at a time. Much of the register file is "triple-ported", allowing two registers to be read simultaneously while a value is written to a third register. Finally, I w...

22. It's cheaper to buy a new printer every month

Source: iDiallo.com | Published: Wed, 10 Dec 2025 01:26:12 GMT

Link: <https://idiallo.com/byte-size/cheaper-to-buy-a-new-printer-every-month?src=feed>

It's the holiday season. My kids are doing all sorts of projects at school, this means we are doing a whole lot of printing. So I went online to look for the ink cartridge for our printer.

I typed the printer name and model because who knows the ink cartridge number? I found our printer. And to my surprise, the printer only costs \$64. What a time to be alive. I grew up in the 90s, those are not the numbers my father dealt with when we wasted paper on his printer.

Buying a printer was a privilege. We dressed up nice, went to the computer store where we met with a salesperson who led us to an office where grown ups negotiated. Office depot (as of December 10th, 2025) Now, it's only \$64.99 and you can complete the transaction from the comfort of your phone. That's Printer + Ink, mind you. What an amazing deal. But what if you already have a printer and only want to replenish your ink? How much does that cost? For the compatible ink cartridges, you'll pay \$81.99. Just for ink. Office depot (as of December 10th, 2025) The ink itself is more expensive than ink + printer. OK, to be fair, the printer I took a screenshot of is on sale. The price was \$64.99 until the sale was scheduled to e...

23. We are always fighting about things we agree on

Source: iDiallo.com | Published: Tue, 30 Sep 2025 00:56:15 GMT

Link: <https://idiallo.com/byte-size/fighting-about-things-we-agree-on?src=feed>

I was in the elevator when an older couple got in, patiently waiting to get to their stop. I smiled, they smiled back, and I pressed the button for the ground floor, just because. All three of us were going down. I had my back to them when I heard the lady say, "I think he will know." I turned to see if they were talking to me, since it was just us. She stepped forward, put on her glasses, then rummaged through her purse to get her phone out. Before I could say a word, she handed it to me and asked, "How can I find all the pictures I have taken?" "Oh boy, here we go again", the husband said. "Hush!" She replied to her husband, while she took the phone from me, pressed the button to turn it on, and then placed it back in my hands. While this may look unusual, I'm completely accustomed to turning into tech support at a moment's notice. When you take pictures on your phone through the camera app, a small preview appears on the bottom. When you click on those pictures, you can view them. The problem was that you can only view the pictures you just took. This is a security measure since the camera app can be activated without entering your pin. "She wanted a brand new phone." The husba...

24. What's the name of that Website?

Source: iDiallo.com | Published: Fri, 26 Sep 2025 02:04:12 GMT

Link: <https://idiallo.com/byte-size/whats-the-name-of-that-website?src=feed>

?Hey, what?s the name of that website?? That was the question my brother asked me. I didn't need any clarification, I knew exactly what he meant. When people come home eager to watch a tv show, they sit in front of the TV scrolling through search results. Then, their phone comes out of their pocket. They search again. They check their subscriptions. Is it on Prime? Ok, maybe available as a rental? Then you see it. That defeated look on their face as they turn to you and ask: ?Hey, what?s the name of that website?? I?ve written before about the degraded state of streaming . We went from the convenience of DVDs to the hyper-convenience of early streaming. Now, we?re right back where we started, trapped in a fractured ecosystem of dozens of services. Oh, and a higher monthly fee. The circle is complete. Piracy is back.

25. RSC for LISP Developers

Source: overreacted ? A blog by Dan Abramov | Published: Sun, 01 Jun 2025 00:00:00 GMT

Link: <https://overreacted.io/rsc-for-lisp-developers/>

Quoting for modules.

26. Progressive JSON

Source: overreacted ? A blog by Dan Abramov | Published: Sat, 31 May 2025 00:00:00 GMT

Link: <https://overreacted.io/progressive-json/>

Why streaming isn't enough.

27. Why Does RSC Integrate with a Bundler?

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 30 May 2025 00:00:00 GMT

Link: <https://overreacted.io/why-does-rsc-integrate-with-a-bundler/>

One does not simply serialize a module.

28. Wagon?s algorithm in Python

Source: John D. Cook | Published: Sat, 14 Feb 2026 23:06:41 +0000

Link: <https://www.johndcook.com/blog/2026/02/14/wagons-algorithm-in-python/>

The last three posts have been about Stan Wagon?s algorithm for finding x and y satisfying $x^2 + y^2 = p$ where p is an odd prime. The first post in the series gives Gauss? formula for a solution, but shows why it is impractical for large p. The bottom of this post introduces Wagon?s [?] The post Wagon?s algorithm in Python first appeared on John D. Cook .

29. Finding a square root of -1 mod p

Source: John D. Cook | Published: Sat, 14 Feb 2026 22:07:05 +0000

Link: <https://www.johndcook.com/blog/2026/02/14/square-root-minus-1-mod-p/>

If p is an odd prime, there is a theorem that says $x^2 \equiv -1 \pmod{p}$ has a solution if and only if $p \equiv 1 \pmod{4}$. When a solution x exists, how do you find it? The previous two posts have discussed Stan Wagon's algorithm for expressing an odd prime p as a $[?]$ The post Finding a square root of $-1 \pmod{p}$ first appeared on John D. Cook .

30. Finding a non-square mod p

Source: John D. Cook | Published: Sat, 14 Feb 2026 21:35:54 +0000

Link: <https://www.johndcook.com/blog/2026/02/14/finding-a-non-square/>

The previous post briefly mentioned Stan Wagon's algorithm for expressing an odd prime p as a sum of two squares when it is possible (i.e. when $p \equiv 1 \pmod{4}$). Wagon's algorithm requires first finding a non-square mod p , i.e. a number c such that $c \not\equiv d^2 \pmod{p}$ for any d in $1, [?]$

The post Finding a non-square mod p first appeared on John D. Cook .

31. You can try to like stuff

Source: DYNOMIGHT | Published: 2025-08-28T00:00:00+00:00

Link: <https://dynamight.net/liking/>

Here's one possible hobby: Take something you don't like. Try to like it. It could be food or music or people or just the general situation you're in. I recommend this hobby, partly because it's nice to enjoy things, but mostly as an instrument for probing human nature. 1. I was in Paris once. By coincidence, I wandered past a bunch of places that were playing Michael Jackson. I thought to myself, ?Huh. The French sure do like Michael Jackson.? Gradually I decided, ?You know what? They're right! Michael Jackson is good.? Later, I saw a guy driving around blasting Billie Jean while hanging a hand outside his car with a sparkly white Michael Jackson glove. Again, I thought, ?Huh.? That day was June 25, 2009 . 2. I don't like cooked spinach. But if I eat some and try to forget that I hate it, it seems OK. Why? Well, as a child, I was subjected to some misguided spinach-related parental interventions. (?You cannot leave this table until you've finished this extremely small portion?, etc.) I hated this, but looking back, it wasn't the innate qualities of spinach that bothered me, so much as that being forced to put something inside my body felt like a violation of my autonomy. When I en...

32. I guess I was wrong about AI persuasion

Source: DYNOMIGHT | Published: 2025-08-21T00:00:00+00:00

Link: <https://dynamight.net/persuasion/>

Say I think abortion is wrong. Is there some sequence of words that you could say to me that would unlock my brain and make me think that abortion is fine? My best guess is that such words do not exist. Really, the bar for what we consider ?open-minded? is incredibly low. Suppose I'm trying to change your opinion about Donald Trump, and I claim that he is a carbon-based life form with exactly one head. If you're willing to concede those points without first seeing where I'm going in my argument?congratulations, you're exceptionally open-minded. Why are humans like that? Well, back at the dawn of our species, perhaps there were some truly open-minded people. But other people talked them into trying weird-looking mushrooms or trading their best clothes for magical rocks. We are the descendants of those other people. I bring this up because, a few months ago, I imagined a Being that had an IQ of 300 and could think at 10,000xnormal speed. I asked how it would be at persuasion. I argued it was unclear, because people just aren't very persuadable. I suspect that if you decided to be open-minded, then the Being would probably be extremely persuasive. But I don't think it's very common...

33. Futarchy's fundamental flaw ? the market ? the blog post

Source: DYNOMIGHT | Published: 2025-08-14T00:00:00+00:00

Link: <https://dynamight.net/futarchy-market/>

Here's our story so far: Markets are a good way to know what people really think. When India and Pakistan started firing missiles at each other on May 7, I was concerned, what with them both having nuclear weapons. But then I looked at world market prices: See how it crashes on May 7? Me neither. I found that reassuring. But we care about lots of stuff that isn't always reflected in stock prices, e.g. the outcomes of elections or drug trials. So why not create

markets for those, too? If you create contracts that pay out \$1 only if some drug trial succeeds, then the prices will reflect what people ?really? think. In fact, why don?t we use markets to make decisions? Say you?ve invented two new drugs, but only have enough money to run one trial. Why don?t you create markets for both drugs, then run the trial on the drug that gets a higher price? Contracts for the ?winning? drug are resolved based on the trial, while contracts in the other market are cancelled so everyone gets their money back. That?s the idea of Futarchy , which Robin Hanson proposed in 2007. Why don?t we? Well, maybe it won?t work. In 2022, I wrote a post arguing that when you cancel one of the markets, you screw up...

34. Justifying text-wrap: pretty

Source: matklad | Published: 2026-02-14T00:00:00+00:00

Link: <https://matklad.github.io/2026/02/14/justifying-text-wrap-pretty.html>

Something truly monumental happened in the world of software development in 2025. Safari shipped a reasonable implementation of text-wrap:

pretty:<https://webkit.org/blog/16547/better-typography-with-text-wrap-pretty/>. We are getting closer and closer to the cutting-edge XV-century technology. Beautiful paragraphs!

35. How COVID Pushed a Generation of Young People to the Right

Source: Derek Thompson | The Atlantic | Published: 2025-02-18T07:00:00-05:00

Link: https://www.theatlantic.com/ideas/archive/2025/02/covid-youth-conservative-shift/681705/?utm_source=feed

For decades, America?s young voters have been deeply?and famously?progressive. In 2008, a youthquake sent Barack Obama to the White House. In 2016, voters ages 18 to 29 broke for Hillary Clinton by 18 points . In 2020, they voted for Joe Biden by 24 points . In 2024, Donald Trump closed most of the gap, losing voters under 30 by a 51?47 margin . In one recent CBS poll , Americans under 30 weren?t just evenly split between the parties. They were even more pro-Trump than Boomers over 65. Precisely polling teens and 20-somethings is a fraught business; some surveys suggest that Trump?s advantage among young people might already be fading . But young people?s apparent lurch right is not an American-only trend. ?Far-right parties are surging across Europe?and young voters are buying in,? the journalist Hanne Cokelaere wrote for Politico last year. In France, Germany, Finland, and beyond, young voters are swinging their support toward anti-establishment far-right parties ?in numbers equal to and even exceeding older voters.? In Germany, a 2024 survey of 2,000 people showed that young people have adopted a relatively new ?gloomy outlook? on the future. No surprise, then, that the far-rig...

36. Why Tom Brady Could Be Worth \$375 Million in the Booth

Source: Derek Thompson | The Atlantic | Published: 2025-02-09T07:00:00-05:00

Link:

https://www.theatlantic.com/ideas/archive/2025/02/tom-brady-mascot/681604/?utm_source=feed

Tom Brady is the greatest quarterback in NFL history (for now). He is not the greatest NFL broadcaster of all time (for now). So why is he calling the Super Bowl tonight, and why is Fox Sports paying him \$375 million over 10 years?more than any other broadcaster in sports history?as if his excellence in the former job automatically qualifies him for the latter? By the simplest conventional analysis, Brady?s initial returns make the contract look like a bad bet. His performance in the booth has received mixed reviews, to put things charitably. In a season when NFL ratings have declined overall , they fell more for games on Fox than on CBS, and more people tuned in to CBS on average to watch football than to watch Brady?s late-afternoon slot. In fairness, as Austin Karp at Sports Business Journal pointed out , the games on CBS this year happened to be closer in the fourth quarter, and many people watching Fox might have clicked away because the outcome was a foregone conclusion. Applying any normal ROI analysis to Brady?s contract is difficult for at least four reasons: He is not a normal person, ?celebrity sports broadcaster? is not a normal job, the NFL is not a normal industr...

37. America?s ?Marriage Material? Shortage

Source: Derek Thompson | The Atlantic | Published: 2025-02-03T07:00:00-05:00
Link: https://www.theatlantic.com/ideas/archive/2025/02/america-marriage-decline/681518/?utm_source=feed

This is Work in Progress, a newsletter about work, technology, and how to solve some of America's biggest problems. Sign up here . Perhaps you've heard: Young people aren't dating anymore. News media and social media are awash in commentary about the decline in youth romance. It's visible in the corporate data, with dating-app engagement taking a hit . And it's visible in the survey data, where the share of 12th graders who say they've dated has fallen from about 85 percent in the 1980s to less than 50 percent in the early 2020s, with the decline particularly steep in the past few years. Naturally, young people's habits are catnip to news commentators. But although I consider the story of declining youth romance important, I don't find it particularly mysterious . In my essay on the anti-social century , I reported that young people have retreated from all manner of physical-world relationships, whether because of smartphones, over-parenting, or a combination of factors. Compared with previous generations of teens, they have fewer friends, spend significantly less time with the friends they do have, attend fewer parties, and spend much more time alone. Romantic relationships theor...

38. Package Management Namespaces

Source: Andrew Nesbitt | Published: 2026-02-14T00:00:00+00:00
Link: <https://nesbitt.io/2026/02/14/package-management-namespaces.html>

Every package needs a name. The rules for how those names work is one of the most consequential decisions a package manager makes, and one of the hardest to change later. I categorized the approaches previously and touched on the tradeoffs briefly. Flat namespaces RubyGems, PyPI, crates.io, Hex, Hackage, CRAN, and LuaRocks all use flat namespaces: one global pool of names, first-come-first-served. You pick a name, and if nobody has it, it's yours. This gives you gem install rails , pip install requests , cargo add serde . The names are short, memorable, and greppable, with no punctuation to remember and no organization to look up. At scale, though, good names run out. Someone registers database on day one and never publishes a real package. Or they publish something, abandon it, and the name sits there forever, pointing at a library last updated in 2013. PyPI has over 600,000 projects. Many of the short, obvious names were claimed years ago by packages with single-digit downloads. Name scarcity creates pressure, and you end up with python-dateutil because dateutil was taken, beautifulsoup4 because beautifulsoup was the old version, or pillow because the original PIL package was ab...

39. Reading list 02/14/26

Source: Construction Physics | Published: Sat, 14 Feb 2026 13:01:32 GMT
Link: <https://www.construction-physics.com/p/reading-list-021426>

Welcome to the reading list, a weekly list of news and links related to buildings, infrastructure, and industrial technology.

40. How Accurate Are Learning Curves?

Source: Construction Physics | Published: Wed, 24 Dec 2025 13:02:40 GMT
Link: <https://www.construction-physics.com/p/how-accurate-are-learning-curves>

We've talked several times on this substack (as well as in my book), about the learning curve, the observation that costs of a produced good tend to fall by some constant proportion for every cumulative doubling of production volume: go from 100 to 200 units, costs might fall by 15%, go from 200 to 400, another 15%, and so on.

41. Reading List 12/20/25

Source: Construction Physics | Published: Sat, 20 Dec 2025 13:01:26 GMT
Link: <https://www.construction-physics.com/p/reading-list-122025>

Tesla's robotaxi crash reports, a fusion startup merger, the decline of US injection molding, Wyoming's snow fences, and more

42. TeleMessage customers include DC Police, Andreessen Horowitz, JP Morgan, and hundreds

more

Source: micahflee | Published: Thu, 22 May 2025 03:07:50 GMT

Link: <https://micahflee.com/telemessage-customers-include-dc-police-andreesen-horowitz-jp-morgan-and-hundreds-more/>

? Update May 26, 2025: I've published the source code for TeleMessage Explorer: a new open source research tool I've been digging through the 410 GB of Java heap dumps from TeleMessage's archive server, provided by DDoSecrets. Here's a description of the

43. DDoSecrets publishes 410 GB of heap dumps, hacked from TeleMessage's archive server

Source: micahflee | Published: Mon, 19 May 2025 16:51:22 GMT

Link: <https://micahflee.com/ddosecrets-publishes-410-gb-of-heap-dumps-hacked-from-telemessages-archive-server/>

? Update May 21, 2025: I've written a new detailed analysis of the heap dump dataset itself:

TeleMessage customers include DC Police, Andreessen Horowitz, JP Morgan, and hundreds more

Update May 26, 2025: I've published the source code for TeleMessage Explorer: a new open source research

44. How the knock-off Signal app used by Trump officials got hacked in 20 minutes

Source: micahflee | Published: Sun, 18 May 2025 11:05:02 GMT

Link: <https://micahflee.com/how-the-knock-off-signal-app-used-by-trump-officials-got-hacked-in-20-minutes/>

? Update May 19, 2025: Uh oh: DDoSecrets publishes 410 GB of heap dumps, hacked from

TeleMessage's archive server Update May 21, 2025: I've written a new detailed analysis of the heap dump dataset itself: TeleMessage customers include DC Police, Andreessen Horowitz, JP Morgan, and hundreds

45. When LICM fails us

Source: Matt Godbolt's blog | Published: 2025-12-14T12:00:00Z

Link: http://xania.org/202512/14-licm-when-it-doesnt?utm_source=feed&utm_medium=rss

When aliasing can prevent loop-invariant code motion

46. Loop-Invariant Code Motion

Source: Matt Godbolt's blog | Published: 2025-12-13T12:00:00Z

Link: http://xania.org/202512/13-licking-licm?utm_source=feed&utm_medium=rss

The compiler can move code outside of loops to speed things up

47. Unswitching loops for fun and profit

Source: Matt Godbolt's blog | Published: 2025-12-12T12:00:00Z

Link: http://xania.org/202512/12-loop-unswitching?utm_source=feed&utm_medium=rss

Duplicating loops around can yield some decent optimisations

48. Adding Planets to Celestia on macOS

Source: Fernando Borretti | Published: Tue, 29 Jul 2025 00:00:00 +0000

Link: <https://borretti.me/article/adding-planets-to-celestia-on-macos>

A beginner's guide to modding Celestia.

49. Notes on Managing ADHD

Source: Fernando Borretti | Published: Thu, 12 Jun 2025 00:00:00 +0000

Link: <https://borretti.me/article/notes-on-managing-adhd>

Strategies and tactics for staying productive.

50. Inboxes are Underrated

Source: Fernando Borretti | Published: Sat, 24 May 2025 00:00:00 +0000

Link: <https://borretti.me/article/inboxes-are-underrated>

On inboxes as application-specific todo lists.

51. A Metabolic Workspace

Source: Westenberg. | Published: Mon, 05 Jan 2026 00:47:25 GMT

Link: <https://www.joanwestenberg.com/a-metabolic-workspace/>

In 1895, a Belgian lawyer, bibliographer and information scientist named Paul Otlet started building what he would call the Mundaneum: a vast repository in Brussels containing over 12 million index cards cross-referenced by subject, designed to hold the entirety of human knowledge. Otlet had already predicted hyperlinks, search engines, and

52. Barbarossa: How Hitler Lost the War - By Jonathan Dimbleby

Source: Westenberg. | Published: Sun, 04 Jan 2026 10:16:47 GMT

Link: <https://www.joanwestenberg.com/barbarossa-how-hitler-lost-the-war-by-jonathan-dimbleby/>

? Grim but essential // corrective to the Western-centric view of WWII. Dimbleby argues persuasively that the war was won and lost on the Eastern Front in 1941, not on the beaches of Normandy in 1944. It is a catalogue of two competing barbarisms, dispelling the myth of the "clean&

53. 1929 - By Andrew Ross Sorkin

Source: Westenberg. | Published: Sat, 03 Jan 2026 10:00:19 GMT

Link: <https://www.joanwestenberg.com/1929-by-andrew-ross-sorkin/>

? A terrifying // granular look at the hubris that melted the global economy. Sorkin doesn't paint the bankers as cartoon villains, so much as delusional optimists who got high on their own supply of cheap credit. Read this to understand why "this time is different" is

54. Book Review: 20 Goto 10 - 10101001 facts about retro computers by Steven Goodwin ?????

Source: Terence Eden?s Blog | Published: Sat, 14 Feb 2026 12:34:34 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-20-goto-10-10101001-facts-about-retro-computers-by-steven-goodwin/>

This is an excellent "dipping" book. There are nearly 200 articles ranging from short anecdotes, multi-page synopses of complex topics, and quirky little asides. Rather than a linear history of computing, each short chapter ends with a multiple-choice "GOTO". From there, you take a meandering wander throughout retro-computing lore. Some paths lead to dead-ends (a delightful little Game-Over?

55. Vanguard - The Government Project to get British Businesses to use the Internet

Source: Terence Eden?s Blog | Published: Sun, 01 Feb 2026 12:34:10 +0000

Link: <https://shkspr.mobi/blog/2026/02/vanguard-the-government-project-to-get-british-businesses-to-use-the-internet/>

Email isn't an obvious business benefit. Imagine it is the early 1980s and you need to communicate with people across the country. A first-class letter will cost you 17p - about 60p in today's money. The letter will be delivered the next day and you'll have your answer back the day after. By contrast, a single computer terminal was likely to set you back around £3,000 - and that's before you ?

56. Book Review: With the End in Mind - Dying, Death and Wisdom in an Age of Denial by Kathryn Mannix ?????

Source: Terence Eden?s Blog | Published: Sat, 31 Jan 2026 12:34:55 +0000

Link: <https://shkspr.mobi/blog/2026/01/book-review-with-the-end-in-mind-dying-death-and-wisdom-in-an-age-of-denial-by-kathryn-mannix/>

Is it possible to "die well"? We have midwives for births, should we have "deathwives" for the other end of our lives? I think this book was recommended to me in the depths of the pandemic. I was too much of a chicken to read it while those around me were dying. The book aims to normalise the process of death and mostly succeeds. Unlike a lot of books, it doesn't just identify a problem - it?

57. Microsoft Game Pass Ultimate Billing Fraud

Source: Jayden's Blog | Published: 2026-02-14T14:38:00+00:00

Link: <https://jayd.ml/2026/02/14/microsoft-game-pass-fraud.html>

I purchased an Xbox Series X out of some misplaced sense of nostalgia for the 360 and because I needed a 4K player. At the time you could still do the trick where you load up on Xbox Live Gold and then convert it to Game Pass Ultimate cheaply. I signed up for it and then made absolutely sure to disable any autorenewing settings everywhere I could. I remember seeing something to the effect of ?Your subscription will expire 2/2026 and will not renew. At the time I still trusted Microsoft a little, but I made sure to use a one time use credit card number, just in case. Lo and behold, I just got this email: Conveniently for those liars and cheats at Microsoft, somehow in the intervening three years autorenew got turned back on. Oopsie whoopsie sowwy ??! I don't know how this isn't outright fraud.

58. Windows 2000 Minesweeper recreated in Godot 4.1

Source: Jayden's Blog | Published: 2026-02-14T13:00:00+00:00

Link: <https://jayd.ml/2026/02/14/godot-minesweeper.html>

TL;DR Play the game at [minesweeper.jayd.ml!](https://minesweeper.jayd.ml/) See the AGPL source code here! Why?? I decided to recreate Windows 2000 minesweeper in Godot 4.1 as accurately as I reasonably could. I wanted to get more familiar with Godot, and wanted a project where I didn't have to worry about what to do, only how to do it. In the end, I ended up going down the rabbit hole and spending 30% of my time on the actual game and the other 70% on menus, dialogs, and other triviality. It was fun working on stuff that I'd never get past a PM, like black and white mode, and recreating the ding/bling window animation when you click on a window while a dialog is open. The overall experience with Godot was very pleasant - working with Godot has dispelled any desire I had to make my own game engine. Godot is lightweight and well thought out. Overall Approach I wanted to recreate Minesweeper as pixel-perfect as I could. Despite my best efforts (see Fonts below), I couldn't get Godot to render the Windows 2000/9x bitmapped fonts in a pixel perfect way, so the approach was to take screenshots in a VM and only render text with Godot where absolutely needed (highscores). Font Rendering Purgatory I spent way, way too...

59. Premium: The Hater's Guide To The AI Bubble Vol. 2

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 14 Nov 2025 18:10:53 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-the-ai-bubble-vol-2/>

We're approaching the most ridiculous part of the AI bubble, with each day bringing us a new, disgraceful and weird headline. As I reported earlier in the week, OpenAI spent \$12.4 billion on inference between 2024 and September 2025 , and its revenue share with Microsoft heavily suggests

60. Exclusive: Here's How Much OpenAI Spends On Inference and Its Revenue Share With Microsoft

Source: Ed Zitron's Where's Your Ed At | Published: Wed, 12 Nov 2025 16:30:09 GMT

Link: https://www.wheresyoured.at/oai_docs/

EDIT: I am updating this post to reflect that these numbers are based on accrual accounting, meaning that these amounts are from the quarters I am referencing. Anyone telling you otherwise is misinformed. As with my Anthropic exclusive from a few weeks ago , though this feels like a natural premium

61. Premium: OpenAI Burned \$4.1 Billion More Than We Knew - Where Is Its Money Going?

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 07 Nov 2025 16:27:45 GMT

Link: <https://www.wheresyoured.at/where-is-openais-money-going/>

Soundtrack: Queens of the Stone Age - Song For The Dead Editor's Note: The original piece had a mathematical error around burnrate, it's been fixed. Also, welcome to another premium issue! Please do subscribe, this is a massive, 7000-or-so word piece, and that's the

62. tiny corp's product ? a training box

Source: the singularity is nearer | Published: 2026-02-15T00:00:00+08:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/02/15/tiny-corp-product.html>

Our new Hong Kong office. It's starting to shape up what tiny corp's product will be. It's not much of a change from what we sell and do now, but the vision is clearer. Every month, we see these LLMs become more and more human. However, there's a major difference. They do not learn. Everyone has the same Claude/Codex/Kimi, with the same weights, the same desires, and the same biases. If current trends continue, the collapse in diversity will be staggering. To paraphrase: I think there is a world market for maybe five people. This is not the future I want to live in. If trends continue where there's a single model with frozen weights and all learning is in-context, the cloud will win. Except in some highly latency sensitive (fighting robots) or connectivity critical (self driving cars) environments, it will be cheaper to run in batch on the cloud. The enshtification that came to the web won't be the driving force to local models. We either live in a world where open models are so bad even user-hostile closed models are better, or open models are good enough, and competition to run them through sites like openrouter will prevent enshtification. The only way local models win is i...

63. How to Work Hard

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/hwh.html>

(No summary)

64. A Project of One's Own

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/own.html>

(No summary)

65. Fierce Nerds

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/fn.html>

(No summary)

66. New Blog Post: " A Very Early History of Algebraic Data Types"

Source: Computer Things | Published: Thu, 25 Sep 2025 16:50:58 +0000

Link:

<https://buttondown.com/hillelwayne/archive/new-blog-post-a-very-early-history-of-algebraic/>

Last week I said that this week's newsletter would be a brief history of algebraic data types.

I was wrong. That history is now a 3500 word blog post . Patreon blog notes here . I'm speaking at P99 Conf ! My talk, "Designing Low-Latency Systems with TLA+", is happening 10/23 at 11:30 central time. It's an online conf and the talk's only 16 minutes, so come check it out!

67. Many Hard Leetcode Problems are Easy Constraint Problems

Source: Computer Things | Published: Wed, 10 Sep 2025 13:00:00 +0000

Link:

<https://buttondown.com/hillelwayne/archive/many-hard-leetcode-problems-are-easy-constraint/>

In my first interview out of college I was asked the change counter problem: Given a set of coin denominations, find the minimum number of coins required to make change for a given number. IE for USA coinage and 37 cents, the minimum number is four (quarter, dime, 2 pennies).

I implemented the simple greedy algorithm and immediately fell into the trap of the question:

the greedy algorithm only works for "well-behaved" denominations. If the coin values were [10,

9, 1] , then making 37 cents would take 10 coins in the greedy algorithm but only 4 coins

optimally (10+9+9+9). The "smart" answer is to use a dynamic programming algorithm, which I

didn't know how to do. So I failed the interview. But you only need dynamic programming if

you're writing your own algorithm. It's really easy if you throw it into a constraint solver

like MiniZinc and call it a day. int: total; array[int] of int: values = [10, 9, 1];

array[index_set(values)] of var 0..: coins; constraint sum (c in index_set(coins)) (coins[c] *

values[c]) == total; solve minimize sum(coins); You can try this online here . It'll give you a prompt to put in total and then give you successively-better solutions: coins = [0, 0, 37...]

68. The Angels and Demons of Nondeterminism

Source: Computer Things | Published: Thu, 04 Sep 2025 14:00:00 +0000

Link: <https://buttondown.com/hillelwayne/archive/the-angels-and-demons-of-nondeterminism/>

Greetings everyone! You might have noticed that it's September and I don't have the next version of Logic for Programmers ready. As penance, here's ten free copies of the book . So a few months ago I wrote a newsletter about how we use nondeterminism in formal methods. The overarching idea: Nondeterminism is when multiple paths are possible from a starting state. A system preserves a property if it holds on all possible paths. If even one path violates the property, then we have a bug. An intuitive model of this is that for this is that when faced with a nondeterministic choice, the system always makes the worst possible choice . This is sometimes called demonic nondeterminism and is favored in formal methods because we are paranoid to a fault. The opposite would be angelic nondeterminism , where the system always makes the best possible choice . A property then holds if any possible path satisfies that property. 1 This is not as common in FM, but it still has its uses! "Players can access the secret level" or " We can always shut down the computer " are reachability properties, that something is possible even if not actually done. In broader computer science research, I'd say tha...

69. Deprecate Like You Mean It

Source: Entropic Thoughts | Published: Thu, 11 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/deprecate-like-you-mean-it>

(No summary)

70. Advent of Code in Dialog

Source: Entropic Thoughts | Published: Tue, 09 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/advent-of-code-in-dialog>

(No summary)

71. Transparent Leadership Beats Servant Leadership

Source: Entropic Thoughts | Published: Thu, 04 Dec 2025 00:00:00 +0100

Link: <https://entropichoughts.com/transparent-leadership-beats-servant-leadership>

(No summary)

72. They made computers behave like annoying salesmen

Source: Rakhim's blog | Published: Wed, 15 Jan 2025 00:00:00 GMT

Link: <https://rakhim.exotext.com/they-made-computers-behave-like-annoying-salesmen>

Computers are precise machines. You can give a computer a precise command using an inhumane language, and it should perform the command. It's not a human, and there is no point of treating it as one. The goal of humanizing user experience isn't to create an illusion of human interaction - it's to make these mechanical commands more accessible while preserving their precise, deterministic nature. UX designers and product managers of tech companies did a lot of damage to people's understanding of computers by making the software behave like a human; or to be more precise, behave like an annoying salesman. (Image from " Not Now. Not later either " by Chris Oliver) We're all familiar with this type. After receiving a clear "no thanks" they deploy increasingly manipulative tactics to meet their "always-be-closing" quotas: "Would this Wednesday work better?" "What would change your mind?" This behavior is frustrating enough from actual salespeople - it's even worse when programmed into our software. (Corporate LLM training session circa 2025) Personally, I can tolerate but deeply dislike software that pretends to have ulterior motives. Take YouTube, for instance. When I explicitly sa...

73. YouTube picks, December 2024

Source: Rakhim's blog | Published: Tue, 31 Dec 2024 11:05:53 GMT

Link: <https://rakhim.exotext.com/youtube-picks-december-2024>

I've added a new feature to Exotext which converts YouTube links into embeds (when a link is on its own line and surrounded by newlines). To test things out, here are some cool YouTube videos I've watched this month.

74. How to make forbidden changes to SQLite tables

Source: Rakhim's blog | Published: Tue, 24 Dec 2024 11:08:52 GMT

Link: <https://rakhim.exotext.com/how-to-make-forbidden-changes-to-sqlite-tables>

Sometimes you need to make a change to an SQLite table which is not possible with a simple ALTER command. For example, today I realized that email_verifications table in my DB references users with a foreign key, but does not have ON DELETE CASCADE (I simply forgot to put it in). This makes it impossible to delete a record from users table if there are corresponding records in email_verifications . There is a hacky way to achieve this, but I prefer this: Create a new table with the correct structure (in my case, with ON DELETE CASCADE enabled). Copy data from the old table to the new table. Rename the old table. Rename the new table. Drop the old table (if everything is ok). It comes down to: CREATE TABLE IF NOT EXISTS email_verifications_2 (id INTEGER PRIMARY KEY, user_id INTEGER NOT NULL, verification_code TEXT NOT NULL, created_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP, FOREIGN KEY (user_id) REFERENCES users (user_id) ON DELETE CASCADE); INSERT INTO email_verifications_2 SELECT * FROM email_verifications; ALTER TABLE email_verifications RENAME TO email_verifications_old; ALTER TABLE email_verifications_2 RENAME TO email_verifications; If everything is ok, we can drop th...

75. an engineer's perspective on hiring

Source: the website of jyn | Published: 2025-08-08T00:00:00+00:00

Link: <https://jyn.dev/an-engineers-perspective-on-hiring/>

hiring in tech is broken and everyone knows it. what can we do better?

76. you are in a box

Source: the website of jyn | Published: 2025-07-14T00:00:00+00:00

Link: <https://jyn.dev/you-are-in-a-box/>

your data is trapped inside the box that is your program. you can only see what the program author exposes.

77. constrained languages are easier to optimize

Source: the website of jyn | Published: 2025-07-12T00:00:00+00:00

Link: <https://jyn.dev/constrained-languages-are-easier-to-optimize/>

exposing raw pointers make the optimizer's job horribly hard. high level languages can constrain your program, making more optimizations sound.

78. Satya Nadella ? How Microsoft is preparing for AGI

Source: Dwarkesh Podcast | Published: Wed, 12 Nov 2025 17:03:08 GMT

Link: <https://www.dwarkesh.com/p/satya-nadella-2>

Plus a tour of Fairwater 2, the most powerful AI datacenter in the world

79. Sarah Paine ? How Russia sabotaged China's rise

Source: Dwarkesh Podcast | Published: Fri, 31 Oct 2025 15:01:37 GMT

Link: <https://www.dwarkesh.com/p/sarah-paine-russo-chinese>

Plus, where Russia and China go from here

80. Thoughts on the AI buildout

Source: Dwarkesh Podcast | Published: Wed, 22 Oct 2025 17:58:39 GMT

Link: <https://www.dwarkesh.com/p/thoughts-on-the-ai-buildout>

Fab CapEx overhang, 1 GW a week, China privileged in long timelines, and much else

81. Nerd Quiz #3

Source: Susam Pal | Published: Tue, 23 Dec 2025 00:00:00 +0000

Link: <https://susam.net/code/news/nq/3.0.0.html>

Nerd Quiz #3 is the third release of Nerd Quiz, a single-page HTML application that invites you to test your nerd level through a short quiz. Each question is inspired by everyday moments of reading, writing, thinking, learning and exploring. This release introduces five new questions drawn from a range of topics, including computing history, graph theory and Unix. Visit Nerd Quiz to try the quiz. A community discussion page is available at Discuss Nerd Quiz . You are very welcome to share your score or discuss the questions there. Read on website | #web | #miscellaneous | #game

82. Mark V. Shaney Junior 0.2.0

Source: Susam Pal | Published: Sun, 14 Dec 2025 00:00:00 +0000

Link: <https://susam.net/code/news/mvs/0.2.0.html>

Mark V. Shaney Junior 0.2.0 is the second release of this little Markov gibberish generator. This release brings two small changes. First, it now reads the training data from standard input instead of a hardcoded file. Second, the program filename has been changed from mvs.py to mvs to reflect that it is an executable file and can be run as ./mvs on most Unix and Linux systems. The source and a detailed documentation for this project are available at github.com/susam/mvs . See also this related blog post and a discussion on Hacker News about it. Read on website | #python | #programming | #technology

83. Fed 24 Years of My Blog Posts to a Markov Model

Source: Susam Pal | Published: Sat, 13 Dec 2025 00:00:00 +0000

Link: <https://susam.net/fed-24-years-of-posts-to-markov-model.html>

Yesterday I shared a little program called Mark V. Shaney Junior at github.com/susam/mvs . It is a minimal implementation of a Markov text generator inspired by the legendary Mark V. Shaney program from the 1980s. Mark V. Shaney was a synthetic Usenet user that posted messages to various newsgroups using text generated by a Markov model. See the Wikipedia article [Mark V. Shaney](#) for more details about it. In this post, I will discuss my implementation of the model, explain how it works and share some of the results produced by it. Contents Recreational Programming Gibberish The Markov Property Some More Gibberish Recreational Programming The program I shared yesterday has only about 30 lines of Python and favours simplicity over efficiency. Even if you have never worked with Markov models before, as long as you know some Python programming, I am quite confident that it will take you less than 20 minutes to understand the whole program and make complete sense of it. I also offer an explanation further below in this post. As a hobby, I often engage in exploratory programming where I write computer programs not to solve a specific problem but simply to explore a particular idea or top...

84. Big LLMs weights are a piece of history

Source: antirez.com | Published: Sun, 16 Mar 2025 12:56:33 +0100

Link: <http://antirez.com/news/147>

By multiple accounts, the web is losing pieces: every year a fraction of old web pages disappear, lost forever. We should regard the Internet Archive as one of the most valuable pieces of modern history; instead, many companies and entities make the chances of the Archive to survive, and accumulate what otherwise will be lost, harder and harder. I understand that the Archive headquarters are located in what used to be a church: well, there is no better way to think of it than as a sacred place. Imagine the long hours spent by old programmers hacking with the Z80 assembly on their Spectrums. All the discussions about the first generation of the Internet. The subcultures that appeared during the 90s. All things that are getting lost, piece by piece. And what about the personal blogs? Pieces of life of single individuals that dumped part of their consciousness on the Internet. Scientific papers and processes that are lost forever as publishers fail, their websites shut down. Early digital art, video games, climate data once published on the Internet and now lost, and many sources of news, as well. This is a known issue and I believe that the obvious approach of trying to preserve eve...

85. Reasoning models are just LLMs

Source: antirez.com | Published: Sun, 09 Feb 2025 19:19:38 +0100

Link: <http://antirez.com/news/146>

It's not new, but it's accelerating. People that used to say that LLMs were a fundamentally flawed way to reach any useful reasoning and, in general, to develop any useful tool with some degree of generality, are starting to shuffle the deck, in the hope to look less wrong. They say: ?the progresses we are seeing are due to the fact that models like OpenAI o1 or DeepSeek R1 are not just LLMs?. This is false, and it is important to show their mystification as soon as possible. First, DeepSeek R1 (don't want to talk about o1 / o3, since it's a private thing we don't have access to, but it's very likely the same) is a pure decoder only autoregressive model. It's the same next token prediction that was so strongly criticized. There isn't, in any place of the model, any explicit symbolic reasoning or representation. Moreover, R1 Zero has similar reasoning capabilities of R1 without requiring *any* supervised fine tuning, just generating chain of thoughts, and improving it with a reward function, using reinforcement learning, was enough to learn a stronger form of reasoning. Interestingly enough, part of these capabilities were easily distilled into smaller models via SFT, which brings...

86. We are destroying software

Source: antirez.com | Published: Sat, 08 Feb 2025 15:47:49 +0100

Link: <http://antirez.com/news/145>

We are destroying software by no longer taking complexity into account when adding features or optimizing some dimension. We are destroying software with complex build systems. We are destroying software with an absurd chain of dependencies, making everything bloated and fragile. We are destroying software telling new programmers: ?Don't reinvent the wheel!?. But, reinventing the wheel is how you learn how things work, and is the first step to make new, different wheels. We are destroying software by no longer caring about backward APIs compatibility. We are destroying software pushing for rewrites of things that work. We are destroying software by jumping on every new language, paradigm, and framework. We are destroying software by always underestimating how hard it is to work with existing complex libraries VS creating our stuff. We are destroying software by always thinking that the de-facto standard for XYZ is better than what we can do, tailored specifically for our use case. We are destroying software claiming that code comments are useless. We are destroying software mistaking it for a purely engineering discipline. We are destroying software by making systems that no longer...

87. ACXD

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/reference/ACXD>

(No summary)

88. faq

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/faq>

(No summary)

89. anime-recommendations

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/anime/anime-recommendations>

(No summary)

90. SNES: Sprites and backgrounds rendering

Source: Fabien Sanglard | Published: 09 Aug 2024 00:00:00 +0000

Link: https://fabiensanglard.net/snes_ppus_why/index.html

(No summary)

91. How the SNES Graphics System works

Source: Fabien Sanglard | Published: 09 Aug 2024 00:00:00 +0000

Link: https://fabiensanglard.net/snes_ppus_how/index.html

(No summary)

92. Carving the Super Nintendo Video System

Source: Fabien Sanglard | Published: 29 Jul 2024 00:00:00 +0000

Link: https://fabiensanglard.net/snes_video/index.html

(No summary)

93. Design Deconstruction

Source: Tedium: The Dull Side of the Internet. | Published: 2026-02-14T16:15:04Z

Link: <https://feed.tedium.co/link/15204/17276365/text-based-design-mindset>

Design is perhaps the software paradigm most wedded to the mouse and the GUI. But there's no reason it can't be text-driven.

94. Tandy Corporation, Part 3

Source: Abort Retry Fail | Published: Tue, 10 Jun 2025 19:38:09 GMT

Link: <https://www.abortretry.fail/p/tandy-corporation-part-3>

Becoming IBM Compatible

95. The Tandy Corporation, Part 2

Source: Abort Retry Fail | Published: Mon, 02 Jun 2025 01:02:32 GMT

Link: <https://www.abortretry.fail/p/the-tandy-corporation-part-2>

Trash-80s get Colorful, and Trash-80s get into your pockets

96. The Tandy Corporation, Part 1

Source: Abort Retry Fail | Published: Sun, 18 May 2025 16:31:36 GMT

Link: <https://www.abortretry.fail/p/the-tandy-corporation-part-1>

From leather shoe bits to the TRS-80

97. Refurb weekend: Silicon Graphics Indigo† IMPACT 10000

Source: Old Vintage Computing Research | Published: 2025-09-13T22:35:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/09/refurb-weekend-silicon-graphics-indigo.html>

It's one of my periodic downsizing cycles, which means checking the hardware inventory (and, intermittently, discovering things that were not on the hardware inventory) and deciding if I want to use it, store it or junk it. And so we come to this machine, which has been sitting in the lab as a practical objet d'art when I picked it up from a fellow collector for the cost of take-it-away almost exactly a decade ago. This beautiful purple slab is the Silicon Graphics Indigo† (though, unlike its earlier namesake, not actually indigo coloured) with the upper-tier MIPS R10000 CPU and IMPACT graphics. My recollection was that it worked at the time, but I couldn't remember if it booted, and of course that was no guarantee that it could still power on. If this machine is to stay working and in the collection, we're gonna need a Refurb Weekend. Counting this sucker, there are three SGI systems presently in the house. (I had a chance many years ago to land one of SGI's early 68K IRIS machines, I think an IRIS 3110, but I was still in a small apartment back then and hadn't the space. I've always regretted turning that one down.) The "big" one is a 900MHz R16000 SGI Fuel (codenamed "Asterix" ...

98. Microsoft makes 6502 BASIC open source

Source: Old Vintage Computing Research | Published: 2025-09-03T12:53:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/09/microsoft-makes-6502-basic-open-source.html>

It was probably going to happen sooner or later, but Microsoft has officially released the source code for 6502 BASIC . The specific revision is very Commodore-centric: it's the 1977 "8K" BASIC variant "1.1," which Commodore users know better as BASIC V2.0, the same BASIC used

in the early PET and with later spot changes from Commodore (including removing Bill Gates' famous Easter egg) in the VIC-20 and Commodore 64. I put "8K" in quotes because the 40-bit Microsoft Binary Format version, which is most familiar as the native floating point format for most 8-bit BASICs derived from Microsoft's and all Commodore BASICs from the PET on up, actually starts at 9K in size. In the C64, because there is RAM and I/O between the BASIC ROM and the Kernal ROM, there is an extra JMP at the end of the BASIC ROM to continue to the routine in the lowest portions of the Kernal ROM. The jump doesn't exist in the VIC-20 where the ROM is contiguous and as a result everything past that point is shifted by three bytes on the C64, the length of the instruction. This is, of course, the same BASIC that Gates wanted a percentage of but Jack Tramiel famously refused to budge on the \$25,000 one-time fee, cl...

99. Reverse-engineering Roadsearch Plus, or, roadgeeking with an 8-bit CPU

Source: Old Vintage Computing Research | Published: 2025-08-23T17:12:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/08/make-your-apple-ii-or-commodore-64.html>

Sorry, Doc Brown: we still needed roads in 1985. That meant paper atlases and misfolded roadmaps and a lot of stereotypical male anxiety asking for directions. Fortunately, in 1985, this problem also had a solution. Yes, if your car inverter could handle a 45-ish watt load ? and your wife doesn't want her seat back right away ? you could navigate major routes across America on your home computer like this portable Commodore SX-64. I particularly enjoyed writing this article because my other irredeemably nerdy habit is roadgeeking, exploring and mapping highways both old and new, and it turns out that 8-bit roadgeeking on ordinary home computers was absolutely possible. For computers of this class, devising an optimal highway route becomes an exercise not only in how to encode sufficient map data to a floppy disk, but also performing efficient graph traversal with limited hardware. Today we'll explore Roadsearch-Plus, one of the (if not the) earliest such software ? primarily on the Commodore 64, but originating on the Apple II ? and at the end "drive" all the way from southern California to British Columbia along US Highway 395, my first long haul expedition, but as it was in 1985...

100. Transforming numbers

Source: Beej's Bit Bucket | Published: Thu, 05 Nov 2015 00:00:00 +0000

Link: <http://beej.us/blog/data/transforming-numbers/>

A little numerical diversion, changing things to other things.

101. Writing a jQuery plugin to create scrolling regions in HTML

Source: Beej's Bit Bucket | Published: Tue, 10 Jul 2012 00:00:00 +0000

Link: <http://beej.us/blog/data/jquery-plugin-scrollregion/>

How to make a scrolling pane in HTML/CSS; in particular, how to wrap it all in a jQuery plugin for easy reuse.

102. HTML canvas scratcher revisited

Source: Beej's Bit Bucket | Published: Mon, 23 Apr 2012 00:00:00 +0000

Link: <http://beej.us/blog/data/html-canvas-scratcher-2/>

An exercise in refactoring code from being closely tied to a page to being less closely tied.

This is based on the previous scratcher code .

103. Can You Spot the Passive Voice?

Source: Refactoring English | Published: Fri, 31 Jan 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/tools/recognize-passive-voice/>

This exercise tests your ability to recognize the passive voice in the context of developing software. For each sentence, click: Active ? if the sentence has no passive voice (active voice only) Passive ? if the sentence contains passive voice Start Does this sentence contain passive voice? Passive ? Active ? Next Question Accuracy: 0 %

104. Passive Voice Considered Harmful

Source: Refactoring English | Published: Fri, 31 Jan 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/chapters/passive-voice-considered-harmful/>

Your high school English teachers probably warned you that passive voice is dangerous and forbidden. Then, when you were an adult, some guy in a leather jacket told you that passive voice is cool and should be used whenever it's desired. Well, the tide has turned again. If you're a software developer, stop using the passive voice. Wait, what's passive voice? In English, sentences can have one of two structures: passive voice or active voice.

105. Rules for Writing Software Tutorials

Source: Refactoring English | Published: Thu, 02 Jan 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/chapters/rules-for-software-tutorials/>

Most software tutorials are tragically flawed. Tutorials often forget to mention some key detail, preventing readers from replicating the author's process. Other times, the author brings in hidden assumptions that don't match their readers' expectations. The good news is that it's easier than you think to write an exceptional software tutorial. You can stand out in a sea of mediocre guides by following a few simple rules. Rules Write for beginners Promise a clear outcome in the title Explain the goal in the introduction Show the end result Make code snippets copy/pasteable Use long versions of command-line flags Separate user-defined values from reusable logic Use unambiguous example values Spare the reader from mindless tasks Keep your code in a working state Teach one thing Don't try to look pretty Minimize dependencies Specify filenames clearly Use consistent, descriptive headings Demonstrate that your solution works Link to a complete example Write for beginners The most common mistake tutorials make is explaining beginner-level concepts using expert-level terminology.

106. Hacking for Defense @ Stanford 2025 ? Lessons Learned Presentations

Source: Steve Blank | Published: Tue, 17 Jun 2025 13:00:04 +0000

Link: <https://steveblank.com/2025/06/17/hacking-for-defense-stanford-2025-lessons-learned-presentations/>

The videos and PowerPoints embedded in this post are best viewed on steveblank.com We just finished our 10th annual Hacking for Defense class at Stanford. What a year. Hacking for Defense, now in 70 universities, has teams of students working to understand and help solve national security problems. At Stanford this quarter the 8 teams [?]

107. Teaching National Security Policy with AI

Source: Steve Blank | Published: Tue, 10 Jun 2025 13:00:25 +0000

Link: <https://steveblank.com/2025/06/10/teaching-national-security-policy-with-ai/>

The videos embedded in this post are best viewed on steveblank.com International Policy students will be spending their careers in an AI-enabled world. We wanted our students to be prepared for it. This is why we've adopted and integrated AI in our Stanford national security policy class ? Technology, Innovation and Great Power Competition. Here's [?]

108. How the United States Gave Up Being a Science Superpower

Source: Steve Blank | Published: Tue, 13 May 2025 13:00:50 +0000

Link: <https://steveblank.com/2025/05/13/how-the-united-states-became-a-science-superpower-and-how-quickly-it-could-crumble/>

This article previously appeared in Nature. US global dominance in science was no accident, but a product of a far-seeing partnership between public and private sectors to boost innovation and economic growth. Since 20 January, US science has been upended by severe cutbacks from the administration of US President Donald Trump. A series of dramatic [?]

109. The Tapestry of Influence: The Inspirations That Shape Us

Source: Simone Salis | Published: Tue, 05 Nov 2024 11:29:04 GMT

Link: <https://simone.org/influences/>

Discover how your unique combination of influences?from art and philosophy to pop culture and personal experiences?shapes your creative perspective. Explore how connecting seemingly

unrelated passions and memories creates unexpected insights and original expression.

110. Being a Bridge Person

Source: Simone Salis | Published: Tue, 15 Oct 2024 11:33:21 GMT

Link: <https://simone.org/bridge-person/>

How lowering your masks can ripple outward, changing not just individuals, but entire communities.

111. The Art of Deliberate Friction-Building

Source: Simone Salis | Published: Tue, 10 Sep 2024 18:59:45 GMT

Link: <https://simone.org/deliberate-friction/>

In a world engineered to be frictionless, adding a bit of resistance can be revolutionary.

112. Why Does Have I Been Pwned Contain "Fake" Email Addresses?

Source: Troy Hunt | Published: Wed, 03 Dec 2025 23:37:06 GMT

Link: <https://www.troyhunt.com/why-does-have-i-been-pwned-contain-fake-email-addresses/>

Normally, when someone sends feedback like this, I ignore it, but it happens often enough that it deserves an explainer, because the answer is really, really simple. So simple, in fact, that it should be evident to the likes of Bruce, who decided his misunderstanding deserved a 1-star Trustpilot review

113. Weekly Update 480

Source: Troy Hunt | Published: Mon, 01 Dec 2025 06:11:03 GMT

Link: <https://www.troyhunt.com/weekly-update-480/>

Well, I now have the answer to how Snapchat does age verification for under-16s: they give an underage kid the ability to change their date of birth, then do a facial scan to verify. The facial scan (a third party tells me...) allows someone well under 16 to pass it

114. Weekly Update 479

Source: Troy Hunt | Published: Sun, 23 Nov 2025 04:44:21 GMT

Link: <https://www.troyhunt.com/weekly-update-479/>

I gave up on the IoT water meter reader. Being technical and thinking you can solve everything with technology is both a blessing and a curse; dogged persistence has given me the life I have today, but it has also burned serious amounts of time because I never want to

115. Staying in one place doesn't mean standing still

Source: Tom Renner on My place to put things | Published: Sat, 13 Mar 2021 00:00:00 +0000

Link: <https://tomrenner.com/posts/staying-in-one-place-doesnt-mean-standing-still/>

Talk given at Codebar Festival . If you wish to see my slides in their full glory, they are available on Slideshare .

116. Make yourself a happy place in your inbox - a mindfulness tip for your working day

Source: Tom Renner on My place to put things | Published: Wed, 03 Feb 2021 00:00:00 +0000

Link: <https://tomrenner.com/posts/make-yourself-a-happy-place/>

Your work inbox is probably not a place that sparks joy. It's full of people asking you to do things, complaints that something hasn't been done, and 571 messages marked urgent. In fact email is usually considered to be a hindrance, with many productivity guides recommending simply ignoring your email for large periods of the day, blocking out that time for focussed work. The consensus is that your work inbox is just a place that generates distractions, and contains never-ending to-do lists that nag away at you throughout the day.

117. Twitter network graphing with Gephi

Source: Tom Renner on My place to put things | Published: Mon, 20 Apr 2020 00:00:00 +0000

Link: <https://tomrenner.com/posts/twitter-graphing-with-gephi/>

Unfortunately since Elon shut down the Twitter APIs, the below method no longer works. Still,

it was fun while it lasted, eh? Gephi is a piece of software for visualising graph networks. It's a powerful tool, and to use it fully requires significant domain knowledge that I don't possess, but fortunately it's still fascinating to play around with as an amateur! As a techie, to me the obvious networks to graph are those created by the big Social Networks - YouTube, Facebook, etc. It's not hugely surprising to find that mostly these graphs mostly aren't available for querying, but excitingly there is a plugin that allows you to stream in live data from Twitter.

118. Space News, February 2026

Source: Dr Paris Buttfield-Addison | Published: Sun, 01 Feb 2026 00:00:00 +0000

Link: <https://hey.paris/posts/spacenews-february-2026/>

These are my notes and expanded thoughts from this month's Space News segment on ABC Radio Hobart and ABC Northern Tasmania. Every month I join Lucie Cutting on Sunday mornings to chat about what's happening above and beyond. Here's what we discussed, plus some of my notes, and expanded thoughts from this edition of the programme : Katherine Bennell-Pegg: 2026 Australian of the Year Katherine Bennell-Pegg was named 2026 Australian of the Year on Sunday 26 January at the National Arboretum in Canberra. She's 41, born in South Australia, holds dual British-Australian citizenship, and is the first person to qualify as an astronaut under the Australian flag . She hasn't flown to space yet.

119. We Made a Game About a 1993 Hint Line

Source: Dr Paris Buttfield-Addison | Published: Fri, 23 Jan 2026 00:00:00 +0000

Link: <https://hey.paris/posts/hint-line-93/>

I'm so proud of this. Hint Line '93 , a game we made at Yarn Spinner , is now on display at ACMI's Game Worlds exhibition. We put up a big behind-the-scenes post on the Yarn Spinner blog , so go read that? For those who don't know, early last year ACMI (the Australian Centre for the Moving Image, a museum of screen culture in Melbourne) put out a call for commissioned microgames. They wanted playable experiences for a museum context: 5-10 minutes, compelling world-building, accessible to a wide variety of visitors.

120. State of Play: How Australian States Are (Not) Governing AI

Source: Dr Paris Buttfield-Addison | Published: Mon, 19 Jan 2026 00:00:00 +0000

Link: <https://hey.paris/posts/state-ai-governance/>

Related: This article examines state-level AI governance in Australia. For analysis of the Commonwealth's AI governance failures, see Slop for the People . For context on the newly announced federal AI Safety Institute, see Australia's AI Safety Institute: Lessons from the UK and US . The Commonwealth government attracts most attention for its AI policies. State and territory governments are also deploying AI systems in education, health, policing, transport, and social services. Each state governs AI differently. Some are ahead of the Commonwealth. Others barely engage with it. None do it well.

121. Are we dismissing AI spend before the 6x lands?

Source: Martin Alderson | Published: Mon, 22 Dec 2025 00:00:00 GMT

Link:

https://martinalderson.com/posts/are-we-dismissing-ai-spend-before-the-6x-lands/?utm_source=rss
Critics are judging models trained on last-gen hardware. There's a 6x wave of compute already allocated - and it's just starting to produce results.

122. Minification isn't obfuscation - Claude Code proves it

Source: Martin Alderson | Published: Thu, 18 Dec 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/minification-isnt-obfuscation-claude-code-proves-it/?utm_source=rss

Using ASTs and AI agents to reverse engineer minified JavaScript in minutes instead of weeks

123. AI agents are starting to eat SaaS

Source: Martin Alderson | Published: Mon, 15 Dec 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/ai-agents-are-starting-to-eat-saas/?utm_source=rss

Software ate the world. Agents are going to eat SaaS.

124. Bag of words, have mercy on us

Source: Experimental History | Published: Tue, 05 Aug 2025 20:06:09 GMT

Link: <https://www.experimental-history.com/p/bag-of-words-have-mercy-on-us>

OR: Claude will you go to prom with me?

125. Do conversations end when people want them to?

Source: Experimental History | Published: Tue, 22 Jul 2025 18:11:25 GMT

Link: <https://www.experimental-history.com/p/do-conversations-end-when-people>

N = 1172

126. Revenge of the Blockheads

Source: Experimental History | Published: Tue, 08 Jul 2025 12:01:27 GMT

Link: <https://www.experimental-history.com/p/revenge-of-the-blockheads>

Links 'n' updates for Summer 2025

127. Vibe Coding Will Rob Us of Our Freedom

Source: IT Notes | Published: Thu, 05 Jun 2025 08:55:36 +0200

Link: <https://it-notes.dragas.net/2025/06/05/vibe-coding-will-rob-us-of-our-freedom/>

What happens when developers stop understanding code and just trust the 'vibe'? A real case of vulnerability hidden behind 'perfect' AI code and the implications for programming's future.

128. The Day GlusterFS Tried to Kill My Career

Source: IT Notes | Published: Wed, 21 May 2025 12:55:00 +0200

Link: https://it-notes.dragas.net/2025/05/21/the_day_glusterfs_tried_to_kill_my_career/

A GlusterFS failure during a critical client upgrade turned into a nightmare of false accusations and politics, revealing a deeper story of corruption.

129. The Server That Wasn't Meant to Exist

Source: IT Notes | Published: Tue, 13 May 2025 16:13:36 +0200

Link: https://it-notes.dragas.net/2025/05/13/the_server_that_wasnt_meant_to_exist/

A story from 16 years ago - trust, tech, and a server that had to disappear. They offered me a blank check. I said no.

130. Coding agents as the new compilers

Source: Anil Dash | Published: 2026-02-12T00:00:00Z

Link: <https://anildash.com/2026/02/11/coding-agents-as-the-new-compilers/>

In each successive generation of code creation thus far, we've abstracted away the prior generation over time. Usually, only a small percentage of coders still work on the lower layers of the stack that used to be the space where everyone was working. I've been coding long enough that people were still creating code in assembly when I started (though I was never any good at it!), though I started with BASIC. Since BASIC was an interpreted language, its interpreter would write the assembly language for me, and I never had to see exactly what assembly language code was being created. I definitely did know old-school coders who used to, at first, check that assembly code to see if they liked the output. But eventually, over time, they just learned to trust the system and stopped looking at what happened after the system finished compiling. Even people using more 'close to the metal' languages like C generally trust that their compilers have been optimized enough that they seldom inspect the output of the compiler to make sure it was perfectly optimized for their particular processor or configuration. The benefits of delegating those concerns to the teams that create compilers, and co...

131. Launch it 3 times

Source: Anil Dash | Published: 2026-02-14T00:00:00Z

Link: <https://anildash.com/2026/02/13/launch-it-three-times/>

I wanted to share one of the bits of advice that I find myself most frequently giving to teams when they're working on a product, or founders who are creating a new company: launch it three times. What I mean by that is, it often takes more than one time before your idea actually resonates or sticks with the people you're trying to reach. Sometimes it takes more than twice! And when I say that you might need to launch again, that can mean a lot of different things. It might just be little tweaks to what you originally put out in the world. It might even be less than that? I've worked with teams that put out literally the exact same thing again and found success, because the issue they had the first time was about timing. That's increasingly an issue as people are distracted by the deeply disturbing social and political events going on in the world, and so sometimes they just need you to put things in front of them again so that they can reassess what you were trying to say. Many relaunches are a little more ambitious, of course. Being a Prince fan, I am of course very partial to strategies that involve changing your name. Re-launching under a new name can be a key strategic move...

132. What is a Perpetual Software License?

Source: Keygen Blog | Published: Mon, 20 Dec 2021 06:00:00 GMT

Link: <https://keygen.sh/blog/what-is-a-perpetual-software-license/>

When faced with the choice of a subscription-based or perpetual license, should businesses prefer one over the other?

133. I Quit!

Source: Keygen Blog | Published: Mon, 10 Jan 2022 06:00:00 GMT

Link: <https://keygen.sh/blog/i-quit/>

Enough is enough.

134. Announcing License Key Authentication

Source: Keygen Blog | Published: Tue, 25 Jan 2022 06:00:00 GMT

Link: <https://keygen.sh/blog/announcing-license-key-authentication/>

Introducing a new, even easier way, to authenticate with our software licensing API.

135. Writing a good design document

Source: Grant Slatton's Blog | Published: Sat, 1 Feb 2025 22:01:56 +0000

Link: <https://grantslatton.com/how-to-design-document>

A guide

136. Design Patterns

Source: Grant Slatton's Blog | Published: Sat, 4 Jan 2025 06:31:30 +0000

Link: <https://grantslatton.com/design-patterns>

An honest primer

137. How to write complex software

Source: Grant Slatton's Blog | Published: Wed, 1 Jan 2025 05:47:17 +0000

Link: <https://grantslatton.com/how-to-software>

A general method

138. TIL Simple Merge of two CSVs with Python

Source: matduggan.com | Published: Tue, 06 May 2025 10:55:50 GMT

Link: <https://matduggan.com/til-simple-merge-of-two-csvs-with-python/>

An easy script to merge two CSVs with Python and Pandas

139. GitHub Copilot for Vim Review

Source: matduggan.com | Published: Thu, 01 May 2025 09:15:19 GMT

Link: <https://matduggan.com/github-copilot-for-vim-2/>

My review of GitHub Copilot for Vim

140. Slack: The Art of Being Busy Without Getting Anything Done

Source: matduggan.com | Published: Mon, 17 Mar 2025 13:00:07 GMT

Link: <https://matduggan.com/slack-the-art-of-being-busy-without-getting-anything-done/>

Why I think Slack and Slack-clones are actually a net negative for most organizations that use them.

141. Beyond ?Beyond the wall?

Source: Simon Tatham's quasi-blog | Published: 2024-07-16T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/aperiodic-followup/>

A sequel to last month's post on finite-state transducers for aperiodic tilings.

142. Beyond the wall: working with aperiodic tilings using finite-state transducers

Source: Simon Tatham's quasi-blog | Published: 2024-06-10T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/aperiodic-transducers/>

A sequel to last year's two posts on combinatorial coordinates, this time applying regular-language theory to them.

143. Writing commit messages

Source: Simon Tatham's quasi-blog | Published: 2024-05-19T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/commit-messages/>

(Yet another) guide to writing commit messages in version control, with a focus on what facts and rationale are good things to include.

144. The GPU, not the TPM, is the root of hardware DRM

Source: Matthew Garrett | Published: Thu, 02 Jan 2025 01:14:05 GMT

Link: <https://mjg59.dreamwidth.org/70954.html>

As part of their "Defective by Design" anti-DRM campaign, the FSF recently made the following claim: Today, most of the major streaming media platforms utilize the TPM to decrypt media streams, forcefully placing the decryption out of the user's control (from here). This is part of an overall argument that Microsoft's insistence that only hardware with a TPM can run Windows 11 is with the goal of aiding streaming companies in their attempt to ensure media can only be played in tightly constrained environments. I'm going to be honest here and say that I don't know what Microsoft's actual motivation for requiring a TPM in Windows 11 is. I've been talking about TPM stuff for a long time. My job involves writing a lot of TPM code. I think having a TPM enables a number of worthwhile security features. Given the choice, I'd certainly pick a computer with a TPM. But in terms of whether it's of sufficient value to lock out Windows 11 on hardware with no TPM that would otherwise be able to run it? I'm not sure that's a worthwhile tradeoff. What I can say is that the FSF's claim is just 100% wrong, and since this seems to be the sole basis of their overall claim about Microsoft's strategy...

145. When should we require that firmware be free?

Source: Matthew Garrett | Published: Thu, 12 Dec 2024 15:57:59 GMT

Link: <https://mjg59.dreamwidth.org/70895.html>

The distinction between hardware and software has historically been relatively easy to understand - hardware is the physical object that software runs on. This is made more complicated by the existence of programmable logic like FPGAs, but by and large things tend to fall into fairly neat categories if we're drawing that distinction. Conversations usually become more complicated when we introduce firmware, but should they? According to Wikipedia, Firmware is software that provides low-level control of computing device hardware , and basically anything that's generally described as firmware certainly fits into the "software" side of the above hardware/software binary. From a software freedom perspective, this seems like something where the obvious answer to "Should this be free" is "yes", but it's worth thinking about why the answer is yes - the goal of free software isn't freedom for freedom's

sake, but because the freedoms embodied in the Free Software Definition (and by proxy the DFSG) are grounded in real world practicalities. How do these line up for firmware? Firmware can fit into two main classes - it can be something that's responsible for initialisation of the hardware (s...

146. Android privacy improvements break key attestation

Source: Matthew Garrett | Published: Thu, 12 Dec 2024 12:16:06 GMT

Link: <https://mjg59.dreamwidth.org/70630.html>

Sometimes you want to restrict access to something to a specific set of devices - for instance, you might want your corporate VPN to only be reachable from devices owned by your company. You can't really trust a device that self attests to its identity, for instance by reporting its MAC address or serial number, for a couple of reasons: These aren't fixed - MAC addresses are trivially reprogrammable, and serial numbers are typically stored in reprogrammable flash at their most protected A malicious device could simply lie about them If we want a high degree of confidence that the device we're talking to really is the device it claims to be, we need something that's much harder to spoof. For devices with a TPM this is the TPM itself. Every TPM has an Endorsement Key (EK) that's associated with a certificate that chains back to the TPM manufacturer. By verifying that certificate path and having the TPM prove that it's in posession of the private half of the EK, we know that we're communicating with a genuine TPM[1]. Android has a broadly equivalent thing called ID Attestation. Android devices can generate a signed attestation that they have certain characteristics and identifiers, a...

147. TU Delft lecture: Security of Science

Source: Bert Hubert's writings | Published: Mon, 01 Dec 2025 11:35:00 +0100

Link: <https://berthub.eu/articles/posts/tu-delft-security-of-science/>

This is a mostly verbatim transcript of my lecture at the TU Delft VvTP Physics symposium ?Security of Science? held on the 20th of November. Audio version (scroll along the page to see the associated slides): Thank you so much for being here tonight. It?s a great honor. I used to study here. I?m a dropout. I never finished my studies, so I feel like I graduate tonight. This is a somewhat special presentation, it has footnotes and references, which you can browse later if you like what you saw.

148. Hello Europe, Joe Biden is gone

Source: Bert Hubert's writings | Published: Fri, 28 Nov 2025 17:30:00 +0100

Link: <https://berthub.eu/articles/posts/hello-europe-joe-biden-is-gone/>

Last Friday I attended a useful conference organized by Microsoft. It lived up very well to its title: ?Justice, Security and Fundamental Rights: Dialogue on EU Law Enforcement Policies? Many thanks are due to various proofreaders who improved this article tremendously. From the original invitation tl;dr: European thinkers and policy makers are acting and talking as if the US federal government and courts are still ?normal?, or will soon be so again.

149. Het Nerdvote Effect

Source: Bert Hubert's writings | Published: Tue, 04 Nov 2025 10:30:39 +0100

Link: <https://berthub.eu/articles/posts/het-nerdvote-effect/>

tl;dr: Barbara Kathmann (GroenLinks-PvdA, 32) zal met voorkeursstemmen herkozen worden als kamerlid. Kwam dit door de NerdVote.nl campagne? Of kwam het door het mooie initiatief Stem op een Vrouw? Hier enige analyse, met als conclusie dat het lijkt dat NerdVote belangrijker was. Het doel van NerdVote.nl was om ieder kandidaat-Kamerlid te steunen dat wil en kan bijdragen aan: IT die niet 100% afhankelijk is van Amerikaanse Big-Tech, zodat we onze paspoorten/persoonsgegevens niet door Amerikanen te hoeven laten managen Het behoud van onze privacy, zodat we vrij kunnen communiceren Onze cybersecurity, zodat onze medische uitslagen niet verhandeld worden NerdVote is niet ?TechVote?, we steunen niet de grootste fans van AI en Internet.

150. Intel 9 285K on ASUS Z890: not stable!

Source: Michael Stapelbergs Website | Published: 2025-03-19T17:35:38+01:00

Link:

<https://michael.stapelberg.ch/posts/2025-03-19-intel-core-ultra-9-285k-on-asus-z890-not-stable/>

Update (2025-05-15): Turns out the CPU was faulty! See My 2025 high-end Linux PC for a new article on this build, now with a working CPU. Update (2025-09-07): The replacement CPU also died and I have given up on Intel. See Bye Intel, hi AMD! for more details on the AMD 9950X3D.

In January I ordered the components for a new PC and expected that I would publish a successor to my 2022 high-end Linux PC ? article. Instead, I am now sitting on a PC which regularly encounters crashes of the worst-to-debug kind, so I am publishing this article as a warning for others in case you wanted to buy the same hardware. Components Which components did I pick for this build? Here's the full list: Price Type Article 140 CHF Case Fractal Define 7 Compact

Black Solid 155 CHF Power Supply Corsair RM850x 233 CHF Mainboard ASUS PRIME Z890-P 620 CHF CPU Intel Core Ultra 9 285k 120 CHF CPU fan Noctua NH-D15 G2 39 CHF Case fan Noctua NF-A14 PWM (140 mm) 209 CHF RAM 64 GB DDR5-6400 Corsair Vengeance (2 x 32GB) 280 CHF Disk 4000 GB Samsung 990 Pro 940 CHF GPU Inno3D GeForce RTX4070 Ti Total: ?1800 CHF, excluding the Graphics Card I re-used from a previous build. ?and the next couple of sections go into detai...

151. Tips to debug hanging Go programs

Source: Michael Stapelbergs Website | Published: 2025-02-27T17:51:38+01:00

Link: <https://michael.stapelberg.ch/posts/2025-02-27-debug-hanging-go-programs/>

I was helping someone get my gokrazy/rsync implementation set up to synchronize RPKI data (used for securing BGP routing infrastructure), when we discovered that with the right invocation, my rsync receiver would just hang indefinitely. This was a quick problem to solve, but in the process, I realized that I should probably write down a few Go debugging tips I have come to appreciate over the years! Scenario: hanging Go program If you want to follow along, you can reproduce the issue by building an older version of gokrazy/rsync, just before the bug fix commit (you'll need Go 1.22 or newer): git clone <https://github.com/gokrazy/rsync> cd rsync git reset --hard 6c89d4dda3be055f19684c0ed56d623da458194e^ go install ./cmd/... Now we can try to sync the repository: % gokr-rsync \ -rtO \ --delete \

```
rsync://rsync.paas.rpki.ripe.net/repository/ \ /tmp/rpki-repo [?] 2025/02/08 09:35:10 Opening
TCP connection to rsync.paas.rpki.ripe.net:873 2025/02/08 09:35:10 rsync module "repo", path
"repo/" 2025/02/08 09:35:10 (Client) Protocol versions: remote=31, negotiated=27 2025/02/08
09:35:10 Client checksum: md4 2025/02/08 09:35:10 sending daemon args: [--server --sender -tr .
repo/] 2025/02/08 09...
```

152. Go Protobuf: The new Opaque API

Source: Michael Stapelbergs Website | Published: 2024-12-21T11:06:00+01:00

Link: <https://michael.stapelberg.ch/posts/2024-12-21-go-protobuf-opaque/>

I originally published this post in the Go blog , but am publishing this copy of it in my own blog as well for readers who don't follow the Go blog. [Protocol Buffers (Protobuf) is Google's language-neutral data interchange format. See protobuf.dev .] Back in March 2020, we released a major overhaul of the Go Protobuf API . The [google.golang.org/protobuf](https://google.github.io/protobuf) package introduced first-class support for reflection , a dynamicpb implementation and the protocmp package for easier testing. That release introduced a new protobuf module with a new API.

Today, we are releasing an additional API for generated code, meaning the Go code in the .pb.go files created by the protocol compiler (protoc). This blog post explains our motivation for creating a new API and shows you how to use it in your projects. To be clear: We are not removing anything. We will continue to support the existing API for generated code, just like we still support the older protobuf module (by wrapping the [google.golang.org/protobuf](https://google.github.io/protobuf) implementation). Go is committed to backwards compatibility and this applies to Go Protobuf, too! Background: the (existing) Open Struct API We now call the existing API the Open Struct API,...