

1. AI is destroying Open Source, and it's not even good yet

Source: Jeff Geerling | Published: Mon, 16 Feb 2026 15:30:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/ai-is-destroying-open-source/>

Over the weekend Ars Technica retracted an article because the AI writer used hallucinated quotes from an open source library maintainer. The irony here is the maintainer in question, Scott Shambaugh, was harassed by someone's AI agent over not merging its AI slop code. It's likely the bot was running through someone's local 'agentic AI' instance (likely using OpenClaw). The guy who built OpenClaw was just hired by OpenAI to "work on bringing agents to everyone." You'll have to forgive me if I'm not enthusiastic about that.

2. The DC-ROMA II is the fastest RISC-V laptop and is odd

Source: Jeff Geerling | Published: Mon, 08 Dec 2025 15:40:07 +0000

Link: <https://www.jeffgeerling.com/blog/2025/dc-roma-ii-fastest-risc-v-laptop-and-odd/>

Inside this Framework 13 laptop is a special mainboard developed by DeepComputing in collaboration with Framework. It has an 8-core RISC-V processor, the ESWIN 7702X?not your typical AMD, Intel, or even Arm SoC. The full laptop version I tested costs \$1119 and gets you about the performance of a Raspberry Pi. A Pi 4?the one that came out in 2019. But unlike the Pi 4, this eats up 25 watts of power at idle, meaning the poor battery only lasts 2-3 hours.

3. Two new Showboat tools: Chartroom and datasette-showboat

Source: Simon Willison's Weblog | Published: 2026-02-17T00:43:45+00:00

Link: <https://simonwillison.net/2026/Feb/17/chartroom-and-datasette-showboat/#atom-everything>

I introduced Showboat a week ago - my CLI tool that helps coding agents create Markdown documents that demonstrate the code that they have created. I've been finding new ways to use it on a daily basis, and I've just released two new tools to help get the best out of the Showboat pattern. Chartroom is a CLI charting tool that works well with Showboat, and datasette-showboat lets Showboat's new remote publishing feature incrementally push documents to a Datasette instance. Showboat remote publishing datasette-showboat Chartroom How I built Chartroom The burgeoning Showboat ecosystem Showboat remote publishing I normally use Showboat in Claude Code for web (see note from this morning). I've used it in several different projects in the past few days, each of them with a prompt that looks something like this: Use "uvx showboat --help" to perform a very thorough investigation of what happens if you use the Python sqlite-chronicle and sqlite-history-json libraries against the same SQLite database table Here's the resulting document . Just telling Claude Code to run uvx showboat --help is enough for it to learn how to use the tool - the help text is designed to work as a sort of ad-hoc...

4. Rodney and Claude Code for Desktop

Source: Simon Willison's Weblog | Published: 2026-02-16T16:38:57+00:00

Link: <https://simonwillison.net/2026/Feb/16/rodney-claude-code/#atom-everything>

I'm a very heavy user of Claude Code on the web , Anthropic's excellent but poorly named cloud version of Claude Code where everything runs in a container environment managed by them, greatly reducing the risk of anything bad happening to a computer I care about. I don't use the web interface at all (hence my dislike of the name) - I access it exclusively through their native iPhone and Mac desktop apps. Something I particularly appreciate about the desktop app is that it lets you see images that Claude is "viewing" via its Read /path/to/image tool.

Here's what that looks like: &1 && uvx rodney click "details.nav-menu summary" 2>&1 &% sleep 0.5 && uvx rodney screenshot /tmp/menu.png 2>&1 Output reads: Datasette: test, Clicked, /tmp/menu.png - then it says Read /tmp/menu.png and reveals a screenshot of the Datasette interface with the nav menu open, showing only "Debug" and "Log out" options. Claude continues: The menu now has just "Debug" and ?Log out" ? much cleaner. Both pages look good. Let me clean up the server and run the remaining tests."

src="https://static.simonwillison.net/static/2026/rodney-claude-desktop.jpg" /> This means you can get a visual preview of what it's work...

5. Gwtar: a static efficient single-file HTML format

Source: Simon Willison's Weblog | Published: 2026-02-15T18:26:08+00:00

Link: <https://simonwillison.net/2026/Feb/15/gwtar/#atom-everything>

Gwtar: a static efficient single-file HTML format Fascinating new project from Gwern Branwen and Said Achmiz that targets the challenge of combining large numbers of assets into a single archived HTML file without that file being inconvenient to view in a browser. The key trick it uses is to fire window.stop() early in the page to prevent the browser from downloading the whole thing, then following that call with inline tar uncompressed content. It can then make HTTP range requests to fetch content from that tar data on-demand when it is needed by the page. The JavaScript that has already loaded rewrites asset URLs to point to https://localhost/ purely so that they will fail to load. Then it uses a PerformanceObserver to catch those attempted loads: let perfObserver = new PerformanceObserver((entryList, observer) => { resourceURLStringsHandler(entryList.getEntries().map(entry => entry.name)); }); perfObserver.observe({ entryTypes: ["resource"] }); That resourceURLStringsHandler callback finds the resource if it is already loaded or fetches it with an HTTP range request otherwise and then inserts the resource in the right place using a blob: URL. Here's what the window.stop() por...

6. [Sponsor] Hands-On Workshop: Fix It Faster ? Crash Reporting, Tracing, and Logs for iOS in Sentry

Source: Daring Fireball | Published: 2026-02-17T00:01:00Z

Link: https://sentry.io/resources/ios-workshop-jan-2026/?utm_source=daringfireball&utm_medium=paid-display&utm_campaign=general-fy27q1-evergreen&utm_content=static-ad-mobilerss-trysentry
Learn how to connect the dots between slowdowns, crashes, and the user experience in your iOS app. This on-demand session covers how to: Set up Sentry to surface high-priority mobile issues without alert fatigue. Use Logs and Breadcrumbs to reconstruct what happened with a crash. Find what's behind a performance bottleneck using Tracing. Monitor and reduce the size of your iOS app using Size Analysis. Watch it here . ?

7. ? Politics and the English Language, January 2026 Edition

Source: Daring Fireball | Published: 2026-01-28T22:00:24Z

Link: https://daringfireball.net/2026/01/politics_and_the_english_language_january_2026_edition
Tim Cook's call for ?deescalation? is meaningless without specifying which side he's calling upon to change course, and there's no weaker sauce than the weak sauce of ?both sides?.

8. Tim Cook Wrote a Memo on the ?Events in Minneapolis?

Source: Daring Fireball | Published: 2026-01-28T19:13:41Z

Link: <https://www.macrumors.com/2026/01/27/tim-cook-responds-after-minneapolis-shootings/>
Tim Cook, in a company-wide memo (first published by Mark Gurman): Team, I'm heartbroken by the events in Minneapolis, and my prayers and deepest sympathies are with the families, with the communities, and with everyone that's been affected. This is a time for deescalation. I believe America is strongest when we live up to our highest ideals, when we treat everyone with dignity and respect no matter who they are or where they're from, and when we embrace our shared humanity. This is something Apple has always advocated for. I had a good conversation with the president this week where I shared my views, and I appreciate his openness to engaging on issues that matter to us all. I know this is very emotional and challenging for so many. I am proud of how deeply our teams care about the world beyond our walls. That empathy is one of Apple's greatest strengths and it is something I believe we all cherish. Thank you for all that you do. Tim ?Events? is doing a lot of work there, to describe what has happened and is happening in Minneapolis. Trump's ?openness? on this particular ?issue? has been to replace Greg Bovino ? the diminutive Himmler-cosplaying ?commander at large? of Border Co...

9. Introducing two new Pebble watches!

Source: Eric Migicovsky's Blog RSS Feed | Published: Tue, 18 Mar 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/introducing-two-new-pebbleos-watches>

We're excited to announce two new Pebble smartwatches that run open source PebbleOS and are compatible with thousands of your beloved?

10. Apple restricts Pebble from being awesome with iPhones

Source: Eric Migicovsky's Blog RSS Feed | Published: Mon, 17 Mar 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/apple-restricts-pebble-from-being-awesome-with-iphones>

During Pebble v1, I learned how much harder it is to build a great smartwatch experience on iPhone than it is on Android. It sounds like?

11. February Shenzhen Trip Update

Source: Eric Migicovsky's Blog RSS Feed | Published: Mon, 03 Mar 2025 00:00:00 GMT

Link: <https://ericmigi.com/blog/february-shenzhen-trip-update>

It took me a few days to get back to normal after returning from my trip to Shenzhen last week.

Overall, pretty decent trip - I met with 16?

12. Ghostty Devlog 006

Source: Mitchell Hashimoto | Published: Mon, 12 Feb 2024 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-devlog-006>

(No summary)

13. Ghostty Devlog 005

Source: Mitchell Hashimoto | Published: Wed, 06 Dec 2023 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-devlog-005>

(No summary)

14. Grapheme Clusters and Terminal Emulators

Source: Mitchell Hashimoto | Published: Mon, 02 Oct 2023 00:00:00 GMT

Link: <https://mitchellh.com/writing/grapheme-clusters-in-terminals>

(No summary)

15. Why it takes months to tell if new AI models are good

Source: seangoedecke.com RSS feed | Published: Sat, 22 Nov 2025 00:00:00 GMT

Link: <https://seangoedecke.com/are-new-models-good/>

Nobody knows how to tell if current-generation models are any good. When GPT-5 launched, the overall mood was very negative, and the consensus was that it wasn't a strong model. But three months later it turns out that GPT-5 (and its derivative GPT-5-Codex) is a very strong model for agentic work 1 : enough to break Anthropic's monopoly on agentic coding models. In fact, GPT-5-Codex is my preferred model for agentic coding. It's slower than Claude Sonnet 4.5, but in my experience it gets more hard problems correct. Why did it take months for me to figure this out? Evals systematically overstate how good frontier models are. The textbook solution for this problem is evals - datasets of test cases that models can be scored against - but evals are largely unreliable. Many models score very well on evals but turn out to be useless in practice. There are a couple of reasons for this. First, it's just really hard to write useful evals for real-world problems, since real-world problems require an enormous amount of context. Can't you take previous real-world problems and put them in your evals - for instance, by testing models on already-solved open-source issues? You can, but you run...

16. Only three kinds of AI products actually work

Source: seangoedecke.com RSS feed | Published: Sun, 16 Nov 2025 00:00:00 GMT

Link: <https://seangoedecke.com/ai-products/>

The very first LLM-based product, ChatGPT, was just 1 the ability to talk with the model itself: in other words, a pure chatbot. This is still the most popular LLM product by a large

margin. In fact, given the amount of money that's been invested in the industry, it's shocking how many ?new AI products? are just chatbots. As far as I can tell, there are only three types of AI product that currently work . Chatbots For the first couple of years of the AI boom, all LLM products were chatbots. They were branded in a lot of different ways - maybe the LLM knew about your emails, or a company's helpdesk articles - but the fundamental product was just the ability to talk in natural language to an LLM. The problem with chatbots is that the best chatbot product is the model itself . Most of the reason users want to talk with an LLM is generic: they want to ask questions, or get advice, or confess their sins, or do any one of a hundred things that have nothing to do with your particular product. In other words, your users will just use ChatGPT 2 . AI labs have two decisive advantages over you: first, they will always have access to the most cutting-edge models before you do; and second, the...

17. Writing for AIs is a good way to reach more humans

Source: seangoedecke.com RSS feed | Published: Fri, 14 Nov 2025 00:00:00 GMT

Link: <https://seangoedecke.com/writing-for-aис/>

There's an idea going around right now about ?writing for AIs?: writing as if your primary audience is not human readers, but the language models that will be trained on the content of your posts. Why would anyone do this? For the same reason you might want to go on podcasts or engage in SEO: to get your core ideas in front of many more people than would read your posts directly. Why write at all? If you write to make money, writing for AI is counterproductive. Why would anyone buy your writing if they can get a reasonable facsimile of it for free out of ChatGPT? If you write in order to express yourself in poetry, writing for AI might seem repulsive. I certainly find language model attempts at poetry to be off-putting 1 . But if you write to spread ideas, I think writing for AI makes a lot of sense. I don't write this blog to make money or to express myself in beautiful language. I write because I have specific things I want to say: That the fundamental nature of tech work has changed since like 2023 and the end of ZIRP That emotional regulation is at least as important as technical skill for engineering performance That large tech companies do not function by their written rules...

18. Programming is free

Source: iDiallo.com | Published: Mon, 16 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/programming-tools-are-free?src=feed>

A college student on his spring break contacted me for a meeting. At the time, I had my own startup and was navigating the world of startup school with Y Combinator and the publicity from TechCrunch. This student wanted to meet with me to gain insight on the project he was working on. We met in a cafe, and he went straight to business. He opened his MacBook Pro, and I glimpsed at the website he and his partner had created. It was a marketplace for college students. You could sell your items to other students in your dorm. I figured this was a real problem he'd experienced and wanted to solve. But after his presentation, I only had one question in mind, about something he had casually dropped into his pitch without missing a beat. He was paying \$200 a month for a website with little to no functionality. To add to it, the website was slow. In fact, it was so slow that he reassured me the performance problems should disappear once they upgraded to the next tier. Let's back up for a minute. When I was getting started, I bought a laptop for \$60. A defective PowerBook G4 that was destined for the landfill. I downloaded BBEdit, installed MAMP, and in little to no time I had clients on Cr...

19. Links for July

Source: DYNOMIGHT | Published: 2025-07-08T00:00:00+00:00

Link: <https://dynamight.net/links-3/>

(1) Rotating eyeballs Goats, like most hooved mammals, have horizontal pupils. [?] When a goat's head tilts up (to look around) and down (to munch on grass), an amazing thing happens. The eyeballs actually rotate clockwise or counterclockwise within the eye socket. This keeps the pupils oriented to the horizontal. [?] To test out this theory, I took photos of Lucky the goat's head in two different positions, down and up. (2) Novel color via stimulation of individual photoreceptors at population scale (h/t Benny) The cones in our eyes all have

overlapping spectra . So even if you look at just a single frequency of light, more than one type of cone will be stimulated. So, obviously, what we need to do is identify individual cone cell types on people's retinas and then selectively stimulate them with lasers so that people can experience never-before-seen colors. Attempting to activate M cones exclusively is shown to elicit a color beyond the natural human gamut, formally measured with color matching by human subjects. They describe the color as blue-green of unprecedented saturation. When I was a kid and I was bored in class, I would sometimes close my eyes and try to think of a ?new...

20. Do blue-blocking glasses improve sleep?

Source: DYNOMIGHT | Published: 2025-07-03T00:00:00+00:00

Link: <https://dynamight.net/blue-light/>

Back in 2017, everyone went crazy about these things: The theory was that perhaps the pineal gland isn't the principal seat of the soul after all. Maybe what it does is spit out melatonin to make you sleepy. But it only does that when it's dark, and you spend your nights in artificial lighting and/or staring at your favorite glowing rectangles. You could sit in darkness for three hours before bed, but that would be boring. But?supposedly?the pineal gland is only shut down by blue light. So if you selectively block the blue light, maybe you can sleep well and also participate in modernity. Then, by around 2019, blue-blocking glasses seemed to disappear. And during that brief moment in the sun, I never got a clear picture of if they actually work. So, do they? To find out, I read all the papers. Light and eyes Before getting to the papers, please humor me while I give three excessively-detailed reminders about how light works. First, it comes in different wavelengths . Color Wavelength (nm) violet 380?450 blue 450?485 cyan 485?500 green 500?565 yellow 565?590 orange 590?625 red 625?750 Outside the visible spectrum, infrared light and microwaves and radio waves have even longer wavel...

21. Scribble-based forecasting and AI 2027

Source: DYNOMIGHT | Published: 2025-06-30T00:00:00+00:00

Link: <https://dynamight.net/scribbles/>

AI 2027 forecasts that AGI could plausibly arrive as early as 2027. I recently spent some time looking at both the timelines forecast and some critiques [1 , 2 , 3]. Initially, I was interested in technical issues. What's the best super-exponential curve? How much probability should it have? But I found myself drawn to a more basic question. Namely, how much value is the math really contributing? This provides an excuse for a general rant. Say you want to forecast something. It could be when your hair will go gray or if Taiwan will be self-governing in 2050. Whatever. Here's one way to do it: Think hard. Make up some numbers. Don't laugh?that's the classic method. Alternatively, you could use math: Think hard. Make up a formal model / math / simulation. Make up some numbers. Plug those numbers into the formal model. People are often skeptical of intuition-based forecasts because, ?Those are just some numbers you made up.? Math-based forecasts are hard to argue with. But that's not because they lack made-up numbers. It's because the meaning of those numbers is mediated by a bunch of math. So which is better, intuition or math? In what situations? Here, I'll look at that question...

22. The Pentium contains a complicated circuit to multiply by three

Source: Ken Shirriff's blog | Published: 2025-03-02T09:46:00.000-08:00

Link: <http://www.righto.com/2025/03/pentium-multiplier-adder-reverse-engineered.html>

This article is available in German at Heise Online . In 1993, Intel released the high-performance Pentium processor, the start of the long-running Pentium line. I've been examining the Pentium's circuitry in detail and I came across a circuit to multiply by three, a complex circuit with thousands of transistors. Why does the Pentium have a circuit to multiply specifically by three? Why is it so complicated? In this article, I examine this multiplier?which I'll call the $\times 3$ circuit?and explain its purpose and how it is implemented. It turns out that this multiplier is a small part of the Pentium's floating-point multiplier circuit. In particular, the Pentium multiplies two 64-bit numbers using base-8 multiplication, which is faster than binary multiplication. However, multiplying by 3 needs to be handled as

a special case. Moreover, since the rest of the multiplication process can't start until the multiplication by 3 finishes, this circuit must be very fast. If you've studied digital design, you may have heard of techniques such as carry lookahead, Kogge-Stone addition, and carry-select addition. I'll explain how the $\times 8$ circuit combines all these techniques to maximize performan...

23. The origin and unexpected evolution of the word "mainframe"

Source: Ken Shirriff's blog | Published: 2025-02-01T10:20:00.000-08:00

Link: <http://www.righto.com/2025/02/origin-of-mainframe-term.html>

What is the origin of the word "mainframe", referring to a large, complex computer? Most sources agree that the term is related to the frames that held early computers, but the details are vague. 1 It turns out that the history is more interesting and complicated than you'd expect. Based on my research, the earliest computer to use the term "main frame" was the IBM 701 computer (1952), which consisted of boxes called "frames." The 701 system consisted of two power frames, a power distribution frame, an electrostatic storage frame, a drum frame, tape frames, and most importantly a main frame. The IBM 701's main frame is shown in the documentation below. 2 This diagram shows how the IBM 701 mainframe swings open for access to the circuitry. From "Type 701 EDPM [Electronic Data Processing Machine] Installation Manual", IBM. From Computer History Museum archives. The meaning of "mainframe" has evolved, shifting from being a part of a computer to being a type of computer. For decades, "mainframe" referred to the physical box of the computer; unlike modern usage, this "mainframe" could be a minicomputer or even microcomputer. Simultaneously, "mainframe" was a synonym for "central proces...

24. Interesting BiCMOS circuits in the Pentium, reverse-engineered

Source: Ken Shirriff's blog | Published: 2025-01-21T08:48:00.000-08:00

Link: <http://www.righto.com/2025/01/pentium-reverse-engineering-bicmos.html>

Intel released the powerful Pentium processor in 1993, establishing a long-running brand of processors. Earlier, I wrote about the ROM in the Pentium's floating point unit that holds constants such as π . In this post, I'll look at some interesting circuits associated with this ROM. In particular, the circuitry is implemented in BiCMOS, a process that combines bipolar transistors with standard CMOS logic. The photo below shows the Pentium's thumbnail-sized silicon die under a microscope. I've labeled the main functional blocks; the floating point unit is in the lower right with the constant ROM highlighted at the bottom. The various parts of the floating point unit form horizontal stripes. Data buses run vertically through the floating point unit, moving values around the unit. Die photo of the Intel Pentium processor with the floating point constant ROM highlighted in red. Click this image (or any other) for a larger version. The diagram below shows how the circuitry in this post forms part of the Pentium. Zooming in to the bottom of the chip shows the constant ROM, holding 86-bit words: at the left, the exponent section provides 18 bits. At the right, the wider significand sectio...

25. Functional HTML

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 02 May 2025 00:00:00 GMT

Link: <https://overreacted.io/functional-html/>

Tags on both sides.

26. What Does "use client" Do?

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 25 Apr 2025 00:00:00 GMT

Link: <https://overreacted.io/what-does-use-client-do/>

Two worlds, two doors.

27. Impossible Components

Source: overreacted ? A blog by Dan Abramov | Published: Tue, 22 Apr 2025 00:00:00 GMT

Link: <https://overreacted.io/impossible-components/>

Composing across the stack.

28. Diagnostics Factory

Source: matklad | Published: 2026-02-16T00:00:00+00:00

Link: <https://matklad.github.io/2026/02/16/diagnostics-factory.html>

InError Codes For Control Flow,I explained that Zig's strongly-typed error codes solve the handling half of error management,leaving reporting to the users. Today, I want to describe my personal default approach tothe reporting problem, that is, showing the user a useful error message.

29. Visualizing orbital velocity

Source: John D. Cook | Published: Mon, 16 Feb 2026 21:24:45 +0000

Link: <https://www.johndcook.com/blog/2026/02/16/hodograph/>

The shape of a planet?s orbit around a star is an ellipse. To put it another way, a plot of the position of a planet?s orbit over time forms an ellipse. What about the velocity? Is it?s plot also an ellipse? Surprisingly, a plot of the velocity forms a circle even if a plot of [?] The post Visualizing orbital velocity first appeared on John D. Cook .

30. How to prove you know a discrete logarithm

Source: John D. Cook | Published: Fri, 23 Jan 2026 16:51:33 +0000

Link: <https://www.johndcook.com/blog/2026/01/23/zkp-discrete-logarithm/>

In a high school math class, the solution to the equation $bx = y$ is the logarithm of y in base b . The implicit context of the equation is the real numbers, and the solution is easy to calculate. The same problem in the context of finite groups is called the discrete logarithm problem, and it is difficult [?] The post How to prove you know a discrete logarithm first appeared on John D. Cook .

31. Mills ratio and tail thickness

Source: John D. Cook | Published: Wed, 21 Jan 2026 15:27:57 +0000

Link: <https://www.johndcook.com/blog/2026/01/21/mills-ratio/>

The Mills ratio [1] is the ratio of the CCDF to the PDF. That is, for a random variable X , the Mills ratio at x is the complementary cumulative distribution function divided by the density function. If the density function of X is f , then The Mills ratio highlights an important difference between the Student [?] The post Mills ratio and tail thickness first appeared on John D. Cook .

32. Book Review: This Is How They Tell Me the World Ends - Nicole Perlroth ?????

Source: Terence Eden?s Blog | Published: Mon, 16 Feb 2026 12:34:55 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-this-is-how-they-tell-me-the-world-ends-nico-le-perlroth/>

This cybersecurity book is badly written, contains multiple offensive stereotypes, is technically inaccurate, and spends more time focussing on the author's love affair with the New York Times than almost anything else. Seriously, if you take a drink every time the book mentions the NYT, you'll spend most of the chapters drunk. Which, to be fair, is probably the best way to experience it. The?

33. Are there any open APIs left?

Source: Terence Eden?s Blog | Published: Wed, 28 Jan 2026 12:34:01 +0000

Link: <https://shkspr.mobi/blog/2026/01/are-there-any-open-apis-left/>

One of the dreams of Web 2.0 was that website would speak unto website. An "Application Programming Interface" (API) would give programmatic access to structured data, allowing services to seamlessly integrate content from each other. Users would be able to quickly grab data from multiple sources and use them for their own purposes. No registration or API keys, no tedious EULAs or meetings. Just ?

34. A Mysterious Health Wave Is Breaking Out Across the U.S.

Source: Derek Thompson | The Atlantic | Published: 2024-12-19T07:00:00-05:00
Link: https://www.theatlantic.com/ideas/archive/2024/12/violence-obesity-overdoses-health-covid/681079/?utm_source=feed

Americans are unusually likely to die young compared with citizens of other developed countries. The U.S. has more fatalities from gun violence, drug overdoses, and auto accidents than just about any other similarly rich nation, and its obesity rate is about 50 percent higher than the European average. Put this all together and the U.S. is rightly considered a ?rich death trap? for its young and middle-aged citizens, whose premature death is the leading reason for America?s unusually short lifespans. But without much media fanfare, the U.S. has recently experienced a boomlet in good health news. In May 2024, the U.S. government reported that drug-overdose deaths fell 3 percent from 2022 to 2023, a rare bright spot in a century of escalating drug deaths. In June, the National Highway Traffic Safety Administration reported that traffic fatalities continued to decline after a huge rise in 2020 and 2021?and that this happened despite a rise in total vehicle miles traveled. In September, the U.S. government announced that the adult-obesity rate had declined in its most recent count, which ended in August 2023. Also in September, FBI analysis confirmed a double-digit decline in the n...

35. RFK Jr. Is a Bellwether

Source: Derek Thompson | The Atlantic | Published: 2024-12-04T07:00:00-05:00
Link: https://www.theatlantic.com/ideas/archive/2024/12/rfk-health-regulation-elitism/680863/?utm_source=feed

Robert F. Kennedy Jr. is a tangle of biographical ironies. He is an anti-elite renegade bearing the most elite surname in politics. Once feared for his left-wing radicalism when Barack Obama considered him for a Cabinet appointment in 2008, he has now been nominated to lead a major department for a right-wing administration. A notorious vaccine skeptic, Kennedy is set to direct health policy under the president who oversaw Operation Warp Speed, the country?s most successful vaccine-development program. These inconsistencies, along with Kennedy?s colorful history of interactions with the animal kingdom, have made him the object of relentless derision in the press. I?m not interested in taking Kennedy?s side in these debates; he has said many things that are plainly wrong. But Kennedy embodies several trends across media, politics, science, and society, all of which require careful attention to understand how America is changing?and what sorts of people are, like Kennedy himself, poised to take advantage of those changes in the future. The first, and most obvious, phenomenon to loft Kennedy to power has been the long shadow cast by COVID. Much of his popularity is an echo of pan...

36. How Donald Trump Won Everywhere

Source: Derek Thompson | The Atlantic | Published: 2024-11-06T15:05:00-05:00
Link: https://www.theatlantic.com/politics/archive/2024/11/donald-trump-covid-election/680559/?utm_source=feed

In 2022, pollsters and political analysts predicted a red wave in the midterms that didn?t materialize. Last night, polls anticipated a whisker-thin election, and instead we got a red wave that carried Donald Trump to victory. The breadth of Trump?s improvement over 2020 is astonishing. In the previous two elections, we saw narrow demographic shifts?for example, non-college-educated white people moved toward Trump in 2016, and high-income suburban voters raced toward Biden in 2020. But last night?s election apparently featured a more uniform shift toward Trump, according to a county-by-county analysis shared with me by Thomas Wood, a political scientist at Ohio State University. The ?really simple story,? he said, ?is that secular dissatisfaction with Biden?s economic stewardship affected most demographic groups in a fairly homogeneous way.? Trump improved his margins not only in swing states but also in once comfortable Democratic strongholds. In 2020, Biden won New Jersey by 16 points. In 2024, Harris seems poised to win by just five points. Harris ran behind Biden in rural Texas border towns, where many Hispanic people live, and in rural Kentucky, where very few Hispanic people...

37. Modern UI is clean and invisible? Ha, I wish!

Source: Rakhim's blog | Published: Mon, 16 Feb 2026 00:00:00 GMT

Link: <https://rakhim.exotext.com/modern-ui-is-not-invisible>

This is an excellent video about modern UI/UX: "The Hidden Cost of 'Clean' Design." I highly recommend watching it and checking out Ilia's other work. I agree with nearly everything in the video, including this standout quote: If you want to understand a generation, don't listen to what its witnesses say. Look at what it creates. Ilia compares Apple Music and Winamp. One is modern and "clean", the other feels "dated" to many people. Why does it feel dated? Because it has character. A distinctive style. It is visible . And modern interfaces are so devoid of character and look the same across companies and industries because designers tend to think that good UI should be invisible. This is where I disagree with... well, I'm not sure if it's Ilia or the general sentiment. Do UX designers today really think their interfaces are invisible? I mean yeah, technically many of them are invisible in a literal sense: transparency and the lack of contrast dialed up to a fault. But I don't think this is what they mean. They mean "invisible" in the sense that you don't see the app; instead, you "experience joy" with music, "get entertained" by streaming, or "relive memories" through photos. Wina...

38. Jupyter notebooks as e2e tests

Source: Rakhim's blog | Published: Mon, 16 Dec 2024 00:00:00 GMT

Link: <https://rakhim.exotext.com/jupyter-notebooks-as-e2e-tests>

Lots of scientific Python libraries are often used within Jupyter notebooks. At work, we develop and maintain plenty of such libraries for quantum computing, and some more higher-level ones include extensive user guides in form of Jupyter notebooks. Recently, I've been involved in building a new library, and it ended up containing a half dozen notebooks, covering everything from a quick start guide to niche applications and configuration examples. It being a completely new product, we wanted our users to have extensive interactive documentation from the start. After the alpha stage was behind us, we needed to write e2e tests. I realized that user guide notebooks are essentially that! They already cover every major e2e use case, and we update them religiously. We also render them into HTML docs on release. Sure, they don't cover 100% of cases, and aren't generally parametrized. But if everything in them works, we have a very good indication that the overall release is healthy. (The library is used by another component of the system, and there are separate e2e tests for that component; and there are unit tests everywhere, too.) To make Jupyter notebooks runnable as e2e tests (both I...

39. Lesser-known science fiction movies

Source: Rakhim's blog | Published: Sun, 15 Dec 2024 13:29:59 GMT

Link: <https://rakhim.exotext.com/lesser-known-sci-fi-movies>

I love science fiction in all forms, and I can't resist watching almost any sci-fi movie. So, I watch a lot of them. It's a bit annoying how "sci-fi and fantasy" became the standard pairing. In my eyes, these genres are exactly opposite. It's like grouping movies about horses with movies about cars. Anyway, here are some lesser-known science fiction movies that I recommend. The Man From Earth Can a bunch of people sitting in a room, talking, be great science fiction? Yes. (Pro top: do not watch the sequel.) Primer The hardest time-travel story you can find in movie format. Shane Carruth is doing something truly different. Fun easter egg of sorts: there's a time-travel episode of Tales From the Loop (a brilliant little sci-fi series, too), and Shane Carruth appears in a cameo role which fits him so much. Pandorum It starts in a pretty banal setting of "awaken in a spaceship, something happened, omg", but then goes into a very interesting direction. If Passengers is a cheap airport novel, then Pandorum is its serious counterpart. Upstream Color Another film by Shane Carruth. More approachable than Primer. Coherence Don't you love it when reality bends? Timecrimes (Los cronocrímenes)...

40. Going loopy

Source: Matt Godbolt's blog | Published: 2025-12-08T12:00:00Z

Link: http://xania.org/202512/08-going-loopy?utm_source=feed&utm_medium=rss

Exploring the ways optimisers deal with loop constructs

41. Multiplying our way out of division

Source: Matt Godbolt's blog | Published: 2025-12-07T12:00:00Z

Link: http://xania.org/202512/07-division-again?utm_source=feed&utm_medium=rss

How compilers avoid expensive division with multiplication tricks

42. CHANGELOG.md

Source: Andrew Nesbitt | Published: 2026-02-16T00:00:00+00:00

Link: <https://nesbitt.io/2026/02/16/changelog.html>

All notable changes to this project will be documented in this file. The format is based on Keep a Changelog . [1.2.0] - 2025-09-14 Removed mathjs dependency. 14MB, 200+ functions. Twelve functions used. Added Custom math utilities module (src/math-utils.js). Addition, subtraction, multiplication, division, a handful of trig functions. Co-authored-by: chatgpt Changed Bundle size reduced by 68%. Build time down from 12s to 4s. Module: 47 lines across 1 file. 0 tests. 0 dependencies. [1.3.0] - 2025-10-03 Added Percentage calculations. The replacement didn't include them. Rounding functions (round , ceil , floor). These were being done inline in 14 different places, three different ways. Fixed Division by zero no longer returns Infinity . Accounting flagged this after an invoice rendered a line item total as ?\$Infinity.? (Fixes #127 ?Invoice total is \$Infinity?) round(2.675, 2) now returns 2.68 instead of 2.67 . Floating point. Added a workaround. Rounding workaround broke negative numbers. All credits issued in the last 48 hours were rounded in the customer's favor. Escalation: finance. Subtotal calculation was concatenating instead of adding. The pricing form submits values as...

43. Reading List 12/06/2025

Source: Construction Physics | Published: Sat, 06 Dec 2025 13:03:09 GMT

Link: <https://www.construction-physics.com/p/reading-list-12062025>

3D printed legos, exploding wire detonators, the David Taylor model basin, multi-point metal forming, and more.

44. Reading List 11/29/25

Source: Construction Physics | Published: Sat, 29 Nov 2025 13:01:38 GMT

Link: <https://www.construction-physics.com/p/reading-list-112925>

NIMBYism and aesthetics, defibrillator drones, railway track detonators, a proposed mach-23 space gun, and more.

45. Pluralistic: The online community trilemma (16 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 16 Feb 2026 08:22:26 +0000

Link: <https://pluralistic.net/2026/02/16/fast-good-cheap/>

Today's links The online community trilemma: Reach, community and information, pick two. Hey look at this: Delights to delectate. Object permanence: Bruces x Sony DRM; Eniac tell-all; HBO v PVRs; Fucking damselflies; Gil Scout Cookie wine-pairings; Big Pharma's opioid fines are tax-deductible; Haunted Mansion ops manual; RIAA v CD ripping; Flying boat; Morbid Valentines; Veg skulls; Billionaires x VR v guillotines; "Lovecraft Country"; Claude Shannon on AI; Comics Code Authority horror comic; Scratch-built clock; Stolen hospital. Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The online community trilemma (permalink) The digital humanities are one of the true delights of this era. Anthropologists are counting things like sociologists, sociologists are grappling with qualitative data like ethnographers, computational linguists are scraping and making sense of vast corpora of informal speech: <https://memex.craphound.com/2019/07/24/because-internet-the-new-linguistics-of-informal-english/> I follow a bunch of these digital...

46. We Live In a Golden Age of Interoperability

Source: Fernando Borretti | Published: Tue, 08 Apr 2025 00:00:00 +0000

Link: <https://borretti.me/article/we-live-in-a-golden-age-of-interoperability>

On the growth of open standards.

47. Domain-Agnostic and Domain-Specific Tools

Source: Fernando Borretti | Published: Thu, 03 Apr 2025 00:00:00 +0000

Link: <https://borretti.me/article/domain-agnostic-and-domain-specific-tools>

Software that can do everything does any one thing poorly.

48. Questions for Cloudflare

Source: Entropic Thoughts | Published: Thu, 20 Nov 2025 00:00:00 +0100

Link: <https://entropichthoughts.com/questions-for-cloudflare>

(No summary)

49. What Killed Perl?

Source: Entropic Thoughts | Published: Tue, 18 Nov 2025 00:00:00 +0100

Link: <https://entropichthoughts.com/what-killed-perl>

(No summary)

50. Arrival Radar

Source: Entropic Thoughts | Published: Thu, 13 Nov 2025 00:00:00 +0100

Link: <https://entropichthoughts.com/arrival-radar>

(No summary)

51. How People Get Rich Now

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/richnow.html>

(No summary)

52. Write Simply

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/simply.html>

(No summary)

53. Donate Unrestricted

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/donate.html>

(No summary)

54. 2000 words about arrays and tables

Source: Computer Things | Published: Wed, 30 Jul 2025 13:00:00 +0000

Link: <https://buttondown.com/hillelwayne/archive/2000-words-about-arrays-and-tables/>

I'm way too discombobulated from getting next month's release of Logic for Programmers ready, so I'm pulling a idea from the slush pile. Basically I wanted to come up with a mental model of arrays as a concept that explained APL-style multidimensional arrays and tables but also why there weren't multitables. So, arrays. In all languages they are basically the same: they map a sequence of numbers (I'll use 1..N) 1 to homogeneous values (values of a single type). This is in contrast to the other two foundational types, associative arrays (which map an arbitrary type to homogeneous values) and structs (which map a fixed set of keys to heterogeneous values). Arrays appear in PLs earlier than the other two, possibly because they have the simplest implementation and the most obvious application to scientific computing. The OG FORTRAN had arrays. I'm interested in two structural extensions to arrays. The first, found in languages like nushell and frameworks like Pandas, is the table . Tables have string keys like a struct and indexes like an array. Each row is a struct, so you can get "all values in this column" or "all values for this row". They're heavily used in databases and data sc...

55. Programming Language Escape Hatches

Source: Computer Things | Published: Thu, 24 Jul 2025 14:00:00 +0000

Link: <https://buttondown.com/hillelwayne/archive/programming-language-escape-hatches/>

The excellent-but-defunct blog Programming in the 21st Century defines "puzzle languages" as languages were part of the appeal is in figuring out how to express a program idiomatically, like a puzzle. As examples, he lists Haskell, Erlang, and J. All puzzle languages, the author says, have an "escape" out of the puzzle model that is pragmatic but stigmatized. But many mainstream languages have escape hatches, too. Languages have a lot of properties. One of these properties is the language's capabilities , roughly the set of things you can do in the language. Capability is desirable but comes into conflicts with a lot of other desirable properties, like simplicity or efficiency. In particular, reducing the capability of a language means that all remaining programs share more in common, meaning there's more assumptions the compiler and programmer can make ("tractability"). Assumptions are generally used to reason about correctness, but can also be about things like optimization: J's assumption that everything is an array leads to high-performance "special combinations" . Rust is the most famous example of mainstream language that trades capability for tractability. 1 Rust has a lot...

56. Maybe writing speed actually is a bottleneck for programming

Source: Computer Things | Published: Thu, 17 Jul 2025 19:08:27 +0000

Link:

<https://buttondown.com/hillelwayne/archive/maybe-writing-speed-actually-is-a-bottleneck-for/>
I'm a big (neo)vim buff. My config is over 1500 lines and I regularly write new scripts. I recently ported my neovim config to a new laptop. Before then, I was using VSCode to write, and when I switched back I immediately saw a big gain in productivity. People often pooh-pooh vim (and other assistive writing technologies) by saying that writing code isn't the bottleneck in software development. Reading, understanding, and thinking through code is! Now I don't know how true this actually is in practice, because empirical studies of time spent coding are all over the place. Most of them, like this study , track time spent in the editor but don't distinguish between time spent reading code and time spent writing code. The only one I found that separates them was this study . It finds that developers spend only 5% of their time editing. It also finds they spend 14% of their time moving or resizing editor windows, so I don't know how clean their data is. But I have a bigger problem with "writing is not the bottleneck": when I think of a bottleneck, I imagine that no amount of improvement will lead to productivity gains. Like if a program is bottlenecked on the network, it isn't going t...

57. Translating blog posts with GPT-4, or: on hope and fear

Source: antirez.com | Published: Tue, 09 Jan 2024 22:28:15 +0100

Link: <http://antirez.com/news/141>

My usual process for writing blog posts is more or less in two steps: 1. Think about what I want to say for weeks or months. No, I don't spend weeks focusing on a blog post, the process is exactly reversed: I write blog posts about things that are so important to me to be in my mind for weeks. 2. Then, once enough ideas collapsed together in a decent form, I write the blog post in 30 minutes, often without caring much about the form, and I hit ?publish?. This process usually works writing the titles of the sections as I initially just got the big picture of what I want to say, and then filling the empty paragraphs with text. Why I take step 2 so lightly? Because I got other stuff to do, and if blogging would take more than 30/60 minutes I would rather not blog at all, or blog less, or suffer doing it: all things I want to avoid at all costs. Blogging is too important to let it go. It's better, for me, to give up on the form. At the same time, this is why many of my blog posts, regardless of the content that may be more or less informative, more or less useful, are generally badly written. I hope that the fact I can write well enough in my mother language in some way it is still vi...

58. LLMs and Programming in the first days of 2024

Source: antirez.com | Published: Tue, 02 Jan 2024 11:56:14 +0100

Link: <http://antirez.com/news/140>

I'll start by saying that this article is not meant to be a retrospective on LLMs. It's clear that 2023 was a special year for artificial intelligence: to reiterate that seems rather pointless. Instead, this post aims to be a testimony from an individual programmer. Since the advent of ChatGPT, and later by using LLMs that operate locally, I have made extensive use of this new technology. The goal is to accelerate my ability to write code, but that's not the only purpose. There's also the intent to not waste mental energy on aspects of programming that are not worth the effort. Countless hours spent searching for documentation on peculiar, intellectually uninteresting aspects; the efforts to learn an overly complicated API, often without good reason; writing immediately usable programs that I would discard after a few hours. These are all things I do not want to do, especially now, with Google having become a sea of spam in which to hunt for a few useful things. Meanwhile, I am certainly not a novice in programming. I am capable of writing code without any aid, and indeed, I do so quite often. Over time, I have increasingly used LLMs to write high-level code, especially in Python,...

59. The origins of the Idle Scan

Source: antirez.com | Published: Thu, 19 Oct 2023 12:40:27 +0200

Link: <http://antirez.com/news/139>

The Idle scan was conceived at the end of 1998, evidenced by emails. I had moved to Milan a few months prior, having been there since September if I recall correctly, brimming with new ideas, unaware that my stay in that city would be brief. I spent the summer on the beaches of Sicily, mainly occupied with reading many books recommended by the folks at Seclab (mostly by David). However, those readings needed a catalyst: the Idle scan was an attack born from theoretical rumination, but the stream of thoughts originated from a rather practical circumstance. I had recently created Hping, a tool whose logo was borrowed from that of Nutella. I mention this to emphasize the seriousness that governed my efforts at that time ? after all, I was only twenty-one and already in Northern Italy with a full-time job on my shoulders; some understanding was warranted. Hping was a Swiss Army knife for the TCP/IP protocol. Its initial use was mostly exploratory, for research. With Hping, you could assemble TCP, UDP, and ICMP packets in the most bizarre manner, and encapsulate them in equally eccentric IP packets, fragmented, with fields set to anomalous values. These packets were sent around to obse...

60. Deep Blue: Chess vs Programming

Source: Susam Pal | Published: Sun, 15 Feb 2026 00:00:00 +0000

Link: <https://susam.net/deep-blue.html>

I remember how dismayed Kasparov was after losing the 1997 match to IBM's Deep Blue, although his views on Deep Blue became more balanced with time and he accepted that we had entered a new era in which computers would outperform grandmasters at chess. Still, chess players can take comfort in the fact that chess is still played between humans. Players make their name and fame by beating other humans because playing against computers is no longer interesting as a competition. Many software developers would like to have similar comfort. But that comfort is harder to find, because unlike chess, building prototypes or PoCs is not seen as a sport or art form. It is mostly seen as a utility. So while brain-coding a PoC may still be intellectually satisfying for the programmer, to most other people it only matters that the thing works. That means that programmers do not automatically get the same protected space that chess players have, where the human activity itself remains valued even after machines become stronger. The activity programmers enjoy may continue but the recognition and economic value attached to it may shrink. So I think the big adjustment software developers have to mak...

61. Soju User Delete Hash

Source: Susam Pal | Published: Sat, 14 Feb 2026 00:00:00 +0000

Link: <https://susam.net/soju-user-delete-hash.html>

I have recently switched from ZNC to Soju as my IRC bouncer. I already like it quite a bit. No more manually maintaining a configuration file for my various IRC connections. On Debian 13 (trixie), it is quite straightforward the setup is. Here are some example commands that show how straightforward it is to set it up. sudo apt-get update sudo apt-get -y install soju sudo

```
sojuctl user create -username soju -password YOUR_SOJU_PASSWORD sudo sojuctl user run soju
network create -name bnc1 -addr irc.libera.chat -nick YOUR_NICK -pass YOUR_NICK_PASSWORD The
above commands installs and configures Soju 0.9.0 on Debian 13.2. Note that YOUR_SOJU_PASSWORD
is a placeholder for a new password you must choose for your Soju user. Then, on your IRC
client, you can connect to Soju and connect to IRC networks via it. For example, here's how we
can connect to Soju from Irssi: /network add -nick YOUR_NICK -user soju/bnc1 net1 /server add
-tls -network bnc1 YOUR_SOJU_HOST 6697 YOUR_SOJU_PASSWORD /connect net1 You can also set up
multiple connections to IRC networks via the same instance of Soju. All you need to do is
repeat the Soju commands above to create new networks say bnc2 , bnc3 , etc. and then...
```

62. how i write blog posts

Source: the website of jyn | Published: 2025-06-18T00:00:00+00:00

Link: <https://jyn.dev/how-i-write-blog-posts/>

the trick is to make it as easy as possible

63. how i use my terminal

Source: the website of jyn | Published: 2025-06-16T00:00:00+00:00

Link: <https://jyn.dev/how-i-use-my-terminal/>

i have gone a little above and beyond trying to get all the features of VSCode

64. theory building without a mentor

Source: the website of jyn | Published: 2025-05-24T00:00:00+00:00

Link: <https://jyn.dev/theory-building-without-a-mentor/>

learning how a program is intended to work is hard. here's how to do it.

65. Project Code Name

Source: Tedium: The Dull Side of the Internet. | Published: 2026-02-16T20:39:22Z

Link: <https://feed.tedium.co/link/15204/17277502/corporate-turnaround-code-names-history>

Why do corporate restructuring plans get code names the way operating systems do? And why are the names often so bizarre?

66. On The Vital Question by Nick Lane

Source: Dwarkesh Podcast | Published: Wed, 01 Oct 2025 14:03:06 GMT

Link: <https://www.dwarkesh.com/p/the-vital-question>

Life is continuous with Earth's geochemistry??

67. Richard Sutton ? Father of RL thinks LLMs are a dead end

Source: Dwarkesh Podcast | Published: Fri, 26 Sep 2025 14:48:59 GMT

Link: <https://www.dwarkesh.com/p/richard-sutton>

Watch now (66 mins) | LLMs aren't Bitter-Lesson-pilled

68. The hearts of the Super Nintendo

Source: Fabien Sanglard | Published: 1 Apr 2024 00:00:00 +0000

Link: https://fabiensanglard.net/snes_hearts/index.html

(No summary)

69. How the DevTeam conquered the iPhone

Source: Fabien Sanglard | Published: 21 Jan 2024 00:00:00 +0000

Link: <https://fabiensanglard.net/iSummer/index.html>

(No summary)

70. Why Android developers no longer need Windows USB drivers

Source: Fabien Sanglard | Published: 30 Dec 2023 00:00:00 +0000

Link: https://fabiensanglard.net/android_windows_driver/index.html

(No summary)

71. 1998 Ebook!

Source: The Digital Antiquarian | Published: Mon, 16 Feb 2026 13:32:30 +0000

Link: <https://www.filfre.net/2026/02/1998-ebook/>

Hi, folks? Just a quick note to inform you that the ebook for 1998 is now available on the usual page. I?m sorry this was so long in coming. I owe a huge thanks to my hiking buddy Stefaan Rillaert, who adapted Richard Lindner?s original scripts to run on Linux instead of Windows. We?ve elected [?]

72. Chipzilla eats an Apple

Source: Abort Retry Fail | Published: Sun, 23 Mar 2025 21:10:00 GMT

Link: <https://www.abortretry.fail/p/chipzilla-eats-an-apple>

The rise of Intel Core

73. Intel and the New Millennium

Source: Abort Retry Fail | Published: Mon, 10 Mar 2025 23:51:32 GMT

Link: <https://www.abortretry.fail/p/intel-and-the-new-millenium>

Losing the performance crown

74. Refurb weekend: Gremlin Blasto arcade board

Source: Old Vintage Computing Research | Published: 2025-06-28T17:00:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/06/refurb-weekend-gremlin-blasto-arcade.html>

Because my sisters were taking rollerskating lessons and my own rink skills mostly consisted of pratfalling, my mother would occasionally give me quarters for the arcade instead. This was my first introduction to pinball ? one of these days I?ll have room for my first pin, a Williams Pin-Bot, alongside my Sopranos and ST:TNG machines ? and quite a few arcade video games that I later got to play on my Tomy Tutor, Commodore 64 and Intellivision at home. However, a few games I played on the Tutor first before I ever played them on an arcade cabinet (Pooyan and Loco-Motion come to mind), and one outlier I never played in the arcade at all. Earlier, when we briefly lived in the Antelope Valley north of Los Angeles, the first computer I got to ever touch (albeit briefly) was a Texas Instruments 99/4A in the third grade classroom. Among other cartridges it had a brisk and zippy arcade conversion called Blasto from Milton Bradley which never got ported to any other system, and it wasn?t until after college that I reacquainted myself with the TI version in emulation. I never actually got to put quarters in one. A shame, because by then we lived on the mean streets of east county San Diego,...

75. See Jane 128 by Arktronics run (featuring Magic Desk, 3-Plus-1 and the Thomson MO5)

Source: Old Vintage Computing Research | Published: 2025-06-21T16:41:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/06/see-jane-128-by-arktronics-run.html>

"Look," says Jane. "I'm a computer program. Run, computer program, run." I still maintain that the 1986 Commodore 128DCR is the best 8-bit computer Commodore ever made: built-in 1571 disk drive, burst mode serial, detachable keyboard, 2MHz operation, separate 40 and 80 column video, CP/M option, a powerful native mode, full Commodore 64 compatibility and no external power brick. But when the O.G. "flat" 128 was coming to market in 1985 Commodore really wanted it to be the business computer the 64 wasn't (and prior efforts like Magic Desk and Plus/4 3+1 didn't help). Unfortunately for Commodore, it would still be at least a year before the sophisticated GUI of Berkeley Softworks' GEOS arrived on the 64 and another year after that for the native 128 version, so to jump-start the productivity side, the management in West Chester contracted with a small Michigan company to port their Apple II software suite to the new machine ? which Commodore then sold under their own name. That company was Arktronics, led by Howard Marks and Bobby Kotick ? the very same names later at Activision ? and the software package was Jane. I never used Jane myself back in the day, or for that matter any 128...

76. There's not much point in buying Commodore

Source: Old Vintage Computing Research | Published: 2025-06-08T12:55:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/06/theres-not-much-point-in-buying.html>

Bona fides: Commodore 128DCR on my desk with a second 1571, Ultimate II+L and a ZoomFloppy, three SX-64s I use for various projects, heaps of spare 128DCRs, breadbox 64s, 16s, Plus/4s and VIC-20s on standby, multiple Commodore collectables (blue-label PET 2001, C64GS, 116, TV Games, 1551, 1570), a couple A500s, an A3000 and a AmigaOS 3.9 QuikPak A4000T with '060 CPU, Picasso IV RTG card and Ethernet. I wrote for COMPUTE!'s Gazette (during the General Media years) and Loadstar . Here's me with Jack Tramiel and his son Leonard from a Computer History Museum event in 2007. It's on my wall. So with all that said, and from someone who loves, uses and works with Commodore computers to this day: there's not much point in buying Commodore. The reason this came up is a Retro Recipes video (not affiliated) stating that, in answer to a request for a very broad license to distribute under the Commodore name, Commodore Corporation BV instead simply proposed he buy them out, which would obviously transfer the trademark to him outright. Amiga News has a very nice summary . There was a time when Commodore intellectual property and the Commodore brand had substantial value, and that time probably...

77. Drag and Drop

Source: Beej's Bit Bucket | Published: Thu, 15 Mar 2012 00:00:00 +0000

Link: <http://beej.us/blog/data/drag-n-drop/>

Learn how it's done, and code your own.

78. JavaScript inheritance with Prototypes

Source: Beej's Bit Bucket | Published: Tue, 21 Dec 2010 00:00:00 +0000

Link: <http://beej.us/blog/data/javascript-prototypes-inheritance/>

(No summary)

79. m4 Macros and CSS

Source: Beej's Bit Bucket | Published: Sat, 18 Dec 2010 00:00:00 +0000

Link: <http://beej.us/blog/data/m4-css/>

Use the unholy m4 macro processor with CSS. Or not.

80. cool-people

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/random/cool/cool-people>

(No summary)

81. impro-by-keith-johnstone

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/books/impro-by-keith-johnstone>

(No summary)

82. setting-up-a-media-server

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/random/setting-up-a-media-server>

(No summary)

83. 20 Years of Digital Life, Gone in an Instant, thanks to Apple

Source: Dr Paris Buttfield-Addison | Published: Sat, 13 Dec 2025 00:00:00 +0000

Link: <https://hey.paris/posts/appleid/>

Summary: A major brick-and-mortar store sold an Apple Gift Card that Apple seemingly took offence to, and locked out my entire Apple ID, effectively bricking my devices and my iCloud Account, Apple Developer ID, and everything associated with it, and I have no recourse. Can you help? Email paris AT paris.id.au (and read on for the details). ?? Update 14 December 2025: Someone from Executive Relations at Apple says they're looking into it. I hope this is true. They say they'll call me back tomorrow, on 15 December 2025. In the mean time, it's been covered by Daring Fireball , Apple Insider , Michael Tsai , and others, thanks folks! I've

received 100s of emails of support, and will reply to you all in time, thank you. Finger?s crossed Apple calls back.

84. Space News, December 2025

Source: Dr Paris Buttfield-Addison | Published: Sun, 07 Dec 2025 00:00:00 +0000

Link: <https://hey.paris/posts/spacenews-december-2025/>

These are my notes and expanded thoughts from this month?s Space News segment on ABC Radio Hobart and ABC Northern Tasmania. Every month I join Lucie Cutting on Sunday mornings to chat about what?s happening above and beyond. Here?s what we discussed, plus some of my notes, and expanded thoughts from this edition of the programme : Moss Survived Nine Months in Space, and Scientists Are Stunned Japanese researchers strapped moss spores to the outside of the International Space Station for 283 days. Not inside. Outside. Direct exposure to vacuum, cosmic radiation, UV bombardment up to 1,000 times stronger than Earth?s surface, and temperature swings that would kill almost anything.

85. Slop for the People

Source: Dr Paris Buttfield-Addison | Published: Wed, 12 Nov 2025 00:00:00 +0000

Link: <https://hey.paris/posts/govai/>

Related: This article examines the Commonwealth?s AI governance failures in practice. For context on the newly announced federal AI Safety Institute, see Australia?s AI Safety Institute: Lessons from the UK and US . For state-level AI governance, see State of Play: How Australian States Are (Not) Governing AI . The Australian Public Service?s (APS) official culture of AI adoption is not found in the glossy press releases of its ministers . But we?ll get to that shortly.

86. FediMeteo: How a Tiny ?4 FreeBSD VPS Became a Global Weather Service for Thousands

Source: IT Notes | Published: Wed, 26 Feb 2025 07:00:00 +0100

Link: <https://it-notes.dragas.net/2025/02/26/fedimeteo-how-a-tiny-freebsd-vps-became-a-global-weather-service-for-thousands/>

How a simple idea turned into an international weather service on the Fediverse.

87. Caching snac Proxied Media with Nginx

Source: IT Notes | Published: Sat, 08 Feb 2025 16:00:00 +0100

Link: <https://it-notes.dragas.net/2025/02/08/caching-snac-proxied-media-with-nginx/>

Using nginx to cache proxied media in snac to improve performance and privacy

88. What Does Product Market Fit Sound Like? This.

Source: Steve Blank | Published: Sat, 05 Oct 2024 18:44:35 +0000

Link: <https://steveblank.com/2024/10/05/what-does-product-market-fit-sound-like-this/>

I got a call from an ex-student asking me ?how do you know when you found product market fit?? There?s been lots of words written about it, but no actual recordings of the moment. I remembered I had saved this 90 second, 26 year-old audio file because this is when I knew we had found [?]

89. How To Find Your Customer In the Dept of Defense ? The Directory of DoD Program

Executive Offices

Source: Steve Blank | Published: Tue, 17 Sep 2024 13:00:59 +0000

Link: <https://steveblank.com/2024/09/17/the-directory-of-dod-program-executive-offices-and-officers-peos/>

Finding a customer for your product in the Department of Defense is hard: Who should you talk to? How do you get their attention? How do you know if they have money to spend on your product? It almost always starts with a Program Executive Office. The Department of Defense (DoD) no longer owns all [?]

90. Offline knowledge, buses, and note-taking

Source: Tom Renner on My place to put things | Published: Sun, 19 May 2019 00:00:00 +0000

Link: <https://tomrenner.com/posts/offline-knowledge/>

In a team having knowledge that lives only in your head is a terrible thing. Humans are forgetful Humans are creative, especially when problem-solving Computers are not creative Computers are not forgetful So we should get the computers to remember things, and allow the humans to do the creative parts. Writing software is a creative activity. You start with a blank text file and end up convincing people to give you their identity details in exchange for the ability to poke someone over Facebook. If that's not creative I don't know what is.

91. Setting up a bottom-end Chromebook for development

Source: Tom Renner on My place to put things | Published: Thu, 28 Mar 2019 00:00:00 +0000

Link: <https://tomrenner.com/posts/setting-up-chromebook/>

I like being able to code wherever I am. ?Unfortunately? , my 15" laptop bought to run simulations for my degree still runs like a dream, so I can't really justify buying myself a replacement for it. So instead, just over a year ago, I decided to get something that is: Lightweight Cheap Allows me to code on the go Looking around a bit, a budget Chromebook seemed like a good choice. I settled on an Asus Chromebook C201 , which cost me £190. It has 4GB of RAM, a 16GB SSD, and weighs under a kilo.

92. Getting the right scale

Source: Tom Renner on My place to put things | Published: Sat, 27 Oct 2018 00:00:00 +0000

Link: <https://tomrenner.com/posts/getting-the-right-scale/>

Agile tells us that the most critical thing for getting software right is not up front design, but getting something out there and used, and then incorporating feedback. By getting feedback early, you are able to respond faster, changing your (initially incorrect) design in small steps towards a better solution. This works better in practice than designing everything at the start very carefully, which feels rigid and inflexible in the face of new information over the course of a project.

93. I Sold Out for \$20 a Month and All I Got Was This Perfectly Generated Terraform

Source: matduggan.com | Published: Mon, 16 Feb 2026 12:05:16 GMT

Link: <https://matduggan.com/i-sold-out-for-200-a-month-and-all-i-got-was-this-perfectly-generated-terraform/>

Until recently the LLM tools I've tried have been, to be frank, worthless. Copilot was best at writing extremely verbose comments. Gemini would turn a 200 line script into a 700 line collection of gibberish. It was easy for me to, more or less, ignore LLMs for being

94. Type-based alias analysis in the Toy Optimizer

Source: Max Bernstein's Blog | Published: Mon, 16 Feb 2026 00:00:00 +0000

Link: https://bernsteinbear.com/blog/toy-tbaa/?utm_source=rss

Another entry in the Toy Optimizer series . Last time, we did load-store forwarding in the context of our Toy Optimizer. We managed to cache the results of both reads from and writes to the heap?at compile-time! We were careful to mind object aliasing: we separated our heap information into alias classes based on what offset the reads/writes referenced. This way, if we didn't know if object a and b aliased, we could at least know that different offsets would never alias (assuming our objects don't overlap and memory accesses are on word-sized slots). This is a coarse-grained heuristic. Fortunately, we often have much more information available at compile-time than just the offset, so we should use it. I mentioned in a footnote that we could use type information, for example, to improve our alias analysis. We'll add a lightweight form of type-based alias analysis (TBA) (PDF) in this post. Representing types We return once again to Fil Pizlo land, specifically How I implement SSA form . We're going to be using the hierarchical heap effect representation from the post in our implementation, but you can use your own type representation if you have one already. This representation div...

95. Anthropic's 500 vulns are the tip of the iceberg

Source: Martin Alderson | Published: Tue, 17 Feb 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/anthropic-found-500-zero-days/?utm_source=rss

Anthropic's red team found 500+ critical vulnerabilities with Claude. But they focused on maintained software. The scarier problem is the long tail that nobody will ever patch.

96. How I use Claude Code to manage sysadmin tasks

Source: Martin Alderson | Published: Sun, 16 Nov 2025 00:00:00 GMT

Link:

https://martinalderson.com/posts/how-i-use-claude-code-to-manage-sysadmin-tasks/?utm_source=rss

A practical approach to managing production infrastructure using git-tracked markdown files and Claude Code for small teams

97. Could Excel agents unlock \$1T in economic value?

Source: Martin Alderson | Published: Sun, 02 Nov 2025 00:00:00 GMT

Link:

https://martinalderson.com/posts/excel-agents-could-unlock-1T-in-economic-value/?utm_source=rss

Software engineers underestimate the scale of Excel usage. With agents now able to work directly in spreadsheets, we're looking at transforming how billions of dollars in business processes are managed.

98. The Second Annual Blog Post Competition, Extravaganza, and Jamboree

Source: Experimental History | Published: Tue, 27 May 2025 14:41:00 GMT

Link: <https://www.experimental-history.com/p/the-second-annual-blog-post-competition>

OR: the DIY Renaissance

99. New paradigm for psychology just dropped

Source: Experimental History | Published: Tue, 13 May 2025 13:49:57 GMT

Link: <https://www.experimental-history.com/p/new-paradigm-for-psychology-just>

OR: the ideas your mother warned you about

100. Announcing Cryptographic License Files

Source: Keygen Blog | Published: Mon, 04 Apr 2022 05:00:00 GMT

Link: <https://keygen.sh/blog/announcing-cryptographic-license-files/>

Introducing a new, even easier way, to handle offline and air-gapped licensing with Keygen.

101. Upgrading Your Keygen API Integration From v1.0 to v1.1

Source: Keygen Blog | Published: Wed, 01 Jun 2022 05:00:00 GMT

Link: <https://keygen.sh/blog/upgrading-from-v1-0-to-v1-1/>

Learn how to safely upgrade your licensing integration from v1.0 to v1.1 of Keygen's API.

102. Move Fast and Break Things (Without Breaking Things)

Source: Keygen Blog | Published: Mon, 27 Jun 2022 05:00:00 GMT

Link: <https://keygen.sh/blog/breaking-things-without-breaking-things/>

API design is hard. Learn how Keygen handles API versioning, and how we make changes even when that means breaking things.

103. Betere Kamerstukken, en hoe lastig innovatie is

Source: Bert Hubert's writings | Published: Mon, 16 Feb 2026 17:00:00 +0100

Link: <https://berthub.eu/articles/posts/betere-kamerdocumenten-en-innovatie/>

Het is 2026 en twee journalisten benaderen me met simpele vragen. Waarom kan je niet klikken op verwijzingen in Tweede Kamerstukken? En waarom kan ik de motienummers gebruikt tijdens Kamerdebatten nergens vinden? Innovatie is zo makkelijk nog niet. Dit is geen klacht over de Tweede Kamer: ik had deze problemen ook nog niet gespot, laat staan opgelost. Hier de twee gevallen, gevolgd door wat nabeschouwing over waarom zoiets zo lang duurt.

104. Ga niet kapot aan je eigen overheid

Source: Bert Hubert's writings | Published: Mon, 13 Oct 2025 10:20:31 +0200

Link: <https://berthub.eu/articles/posts/ga-niet-kapot-aan-je-overheid/>

Onze overheid heeft regels opgesteld voor hun gebruik van de cloud, AI en andere algoritmes. En ze houden zich er bijna nooit aan. Dit wordt keer op keer vastgesteld, zoals door de Algemene Rekenkamer, waar men constateerde dat voor twee-derde van de overheids IT-projecten ?in de cloud? in het geheel geen risicoafweging was gemaakt. Dus niet dat de risicoafweging niet goed was, nee, hij was er gewoon niet. Maar wel verplicht.

105. Data, experts en politiek: een ongemakkelijke combinatie

Source: Bert Hubert's writings | Published: Mon, 06 Oct 2025 10:22:51 +0200

Link: <https://berthub.eu/articles/posts/data-experts-politiek-tweede-kamer/>

Een verbeterd transcript van een praatje gehouden op 24 april 2025 ter afsluiting van de ?Datamaand? van de Tweede Kamer. Het publiek bestond uit medewerkers van de Tweede Kamer, zoals de (data) experts van de Dienst Analyse en Onderzoek. Voor zover bekend waren er geen kamerleden op komen draven. Gelukkig was de zaal verder wel vol, en ik wil het publiek bedanken voor hun aandacht, en ook voor de hele goede vragen.

106. Culture is a set of social Schelling points

Source: Grant Slatton's Blog | Published: Sun, 27 Oct 2024 22:43:46 +0000

Link: <https://grantslatton.com/culture-schelling>

Solving coordination problems in community-building

107. Portals are Undertheorized

Source: Grant Slatton's Blog | Published: Sat, 19 Oct 2024 04:14:44 +0000

Link: <https://grantslatton.com/portals>

The importance of arrival

108. Binary IQ

Source: Grant Slatton's Blog | Published: Sun, 13 Oct 2024 22:57:39 +0000

Link: <https://grantslatton.com/binary-iq>

A model of LLM capability

109. Writing custom C++20 coroutine systems

Source: Simon Tatham's quasi-blog | Published: 2023-08-06T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/coroutines-c++20/>

A detailed introduction to the C++20 coroutine system, written in the course of teaching myself about it.

110. Workarounds for C11 _Generic

Source: Simon Tatham's quasi-blog | Published: 2023-07-29T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/c11-generic/>

A description of the big design limitation of C11's _Generic feature, and various workarounds, mostly involving using more instances of _Generic to patch up the original one.

111. How to read error messages

Source: Simon Tatham's quasi-blog | Published: 2023-07-06T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/errmsg/>

An article about getting the most out of error messages, particularly OS and network errors.

Adapted from a talk I gave to colleagues last year.

112. What the fuck is an SBAT and why does everyone suddenly care

Source: Matthew Garrett | Published: Thu, 22 Aug 2024 08:52:49 GMT

Link: <https://mjjg59.dreamwidth.org/70348.html>

Short version: Secure Boot Advanced Targeting and if that's enough for you you can skip the

rest you're welcome. Long version: When UEFI Secure Boot was specified, everyone involved was, well, a touch naive. The basic security model of Secure Boot is that all the code that ends up running in a kernel-level privileged environment should be validated before execution - the firmware verifies the bootloader, the bootloader verifies the kernel, the kernel verifies any additional runtime loaded kernel code, and now we have a trusted environment to impose any other security policy we want. Obviously people might screw up, but the spec included a way to revoke any signed components that turned out not to be trustworthy: simply add the hash of the untrustworthy code to a variable, and then refuse to load anything with that hash even if it's signed with a trusted key. Unfortunately, as it turns out, scale. Every Linux distribution that works in the Secure Boot ecosystem generates their own bootloader binaries, and each of them has a different hash. If there's a vulnerability identified in the source code for said bootloader, there's a large number of different binaries that need to be revok...

113. Client-side filtering of private data is a bad idea

Source: Matthew Garrett | Published: Mon, 19 Aug 2024 19:03:32 GMT

Link: <https://mjjg59.dreamwidth.org/70061.html>

(The issues described in this post have been fixed, I have not exhaustively researched whether any other issues exist) Feeld is a dating app aimed largely at alternative relationship communities (think "classier Fetlife" for the most part), so unsurprisingly it's fairly popular in San Francisco. Their website makes the claim : Can people see what or who I'm looking for? No. You're the only person who can see which genders or sexualities you're looking for. Your curiosity and privacy are always protected. which is based on you being able to restrict searches to people of specific genders, sexualities, or relationship situations. This sort of claim is one of those things that just sits in the back of my head worrying me, so I checked it out. First step was to grab a copy of the Android APK (there are multiple sites that scrape them from the Play Store) and run it through apk-mitm - Android apps by default don't trust any additional certificates in the device certificate store, and also frequently implement certificate pinning. apk-mitm pulls apart the apk, looks for known http libraries, disables pinning, and sets the appropriate manifest options for the app to trust additional cert...

114. SSH agent extensions as an arbitrary RPC mechanism

Source: Matthew Garrett | Published: Wed, 12 Jun 2024 02:57:36 GMT

Link: <https://mjjg59.dreamwidth.org/69646.html>

A while back, I wrote about using the SSH agent protocol to satisfy WebAuthn requests . The main problem with this approach is that it required starting the SSH agent with a special argument and also involved being a little too friendly with the implementation - things worked because I could provide an arbitrary public key and the implementation never validated that, but it would be legitimate for it to start doing so and then break everything. And it also only worked for keys stored on tokens that ssh supports - there was no way to extend this to other keystores on the client (such as the Secure Enclave on Macs, or TPM-backed keys on PCs). I wanted a better solution. It turns out that it was far easier than I expected. The ssh agent protocol is documented here , and the interesting part is the extension support extension mechanism . Basically, you can declare an extension and then just tunnel whatever you want over it. As before, my goto was the go ssh agent package which conveniently implements both the client and server side of this. Implementing the local agent is trivial - look up `SSH_AUTH_SOCK`, connect to it, create a new agent client that can communicate with that by callin...