

1. Frigate with Hailo for object detection on a Raspberry Pi

Source: Jeff Geerling | Published: Wed, 18 Feb 2026 14:50:00 -0600

Link: <https://www.jeffgeerling.com/blog/2026/frigate-with-hailo-for-object-detection-on-a-raspberry-pi/>

I run Frigate to record security cameras and detect people, cars, and animals when in view. My current Frigate server runs on a Raspberry Pi CM4 and a Coral TPU plugged in via USB. Raspberry Pi offers multiple AI HAT+'s for the Raspberry Pi 5 with built-in Hailo-8 or Hailo-8L AI coprocessors, and they're useful for low-power inference (like for image object detection) on the Pi. Hailo coprocessors can be used with other SBCs and computers too, if you buy an M.2 version .

2. LadybirdBrowser/ladybird: Abandon Swift adoption

Source: Simon Willison's Weblog | Published: 2026-02-19T01:25:33+00:00

Link: <https://simonwillison.net/2026/Feb/19/ladybird/#atom-everything>

LadybirdBrowser/ladybird: Abandon Swift adoption Back in August 2024 the Ladybird browser project announced an intention to adopt Swift as their memory-safe language of choice. As of this commit it looks like they've changed their mind: Everywhere: Abandon Swift adoption After making no progress on this for a very long time, let's acknowledge it's not going anywhere and remove it from the codebase. Via Hacker News Tags: ladybird , swift

3. Typing without having to type

Source: Simon Willison's Weblog | Published: 2026-02-18T18:56:56+00:00

Link: <https://simonwillison.net/2026/Feb/18/typing/#atom-everything>

25+ years into my career as a programmer I think I may finally be coming around to preferring type hints or even strong typing. I resisted those in the past because they slowed down the rate at which I could iterate on code, especially in the REPL environments that were key to my productivity. But if a coding agent is doing all that typing for me, the benefits of explicitly defining all of those types are suddenly much more attractive. Tags: ai-assisted-programming , programming , programming-languages , static-typing

4. The A.I. Disruption We?ve Been Waiting for Has Arrived

Source: Simon Willison's Weblog | Published: 2026-02-18T17:07:31+00:00

Link: <https://simonwillison.net/2026/Feb/18/the-ai-disruption/#atom-everything>

The A.I. Disruption We?ve Been Waiting for Has Arrived New opinion piece from Paul Ford in the New York Times. Unsurprisingly for a piece by Paul it's packed with quoteworthy snippets, but a few stood out for me in particular. Paul describes the November moment that so many other programmers have observed, and highlights Claude Code's ability to revive old side projects: [Claude Code] was always a helpful coding assistant, but in November it suddenly got much better, and ever since I?ve been knocking off side projects that had sat in folders for a decade or longer. It?s fun to see old ideas come to life, so I keep a steady flow. Maybe it adds up to a half-hour a day of my time, and an hour of Claude?s. November was, for me and many others in tech, a great surprise. Before, A.I. coding tools were often useful, but halting and clumsy. Now, the bot can run for a full hour and make whole, designed websites and apps that may be flawed, but credible. I spent an entire session of therapy talking about it. And as the former CEO of a respected consultancy firm (Postlight) he's well positioned to evaluate the potential impact: When you watch a large language model slice through some horribl...

5. Ghostty Devlog 003

Source: Mitchell Hashimoto | Published: Thu, 24 Aug 2023 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-devlog-003>

(No summary)

6. Ghostty Devlog 002

Source: Mitchell Hashimoto | Published: Sat, 05 Aug 2023 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-devlog-002>

(No summary)

7. Ghostty Devlog 001

Source: Mitchell Hashimoto | Published: Thu, 13 Jul 2023 00:00:00 GMT

Link: <https://mitchellh.com/writing/ghostty-devlog-001>

(No summary)

8. Why do AI models use so many em-dashes?

Source: seangoedecke.com RSS feed | Published: Thu, 30 Oct 2025 00:00:00 GMT

Link: <https://seangoedecke.com/em-dashes/>

If you asked most people to name a defining feature of AI-generated writing, they'd probably say the em-dash ? like this. Language models use em-dashes so much that real humans who like em-dashes have stopped using them out of fear of being confused with AI. It's also surprisingly hard to prompt models to avoid em-dashes: take this thread from the OpenAI forums where users share their unsuccessful attempts. Given all that, it's kind of weird that we don't really know why language models use the em-dash so much . Explanations I don't find convincing One common explanation is that normal English text contains a lot of em-dashes, so it's just learned behavior from the training data. I find this fairly unconvincing, for the reason that everyone thinks AI uses a lot of em-dashes . If em-dashes were as common in AI prose as human prose, they would be as unremarkable as the use of other punctuation marks. Another explanation I'm not convinced by is that AI models like em-dashes because they're so versatile. When the model is trying to predict the next token, an em-dash keeps its options open: it could either continue on the same point or make a brand new point. Since models are just try...

9. Mistakes I see engineers making in their code reviews

Source: seangoedecke.com RSS feed | Published: Sat, 25 Oct 2025 00:00:00 GMT

Link: <https://seangoedecke.com/good-code-reviews/>

In the last two years, code review has gotten much more important. Code is now easy to generate using LLMs, but it's still just as hard to review 1 . Many software engineers now spend as much (or more) time reviewing the output of their own AI tools than their colleagues? code. I think a lot of engineers don't do code review correctly. Of course, there are lots of different ways to do code review, so this is largely a statement of my engineering taste . Don't just review the diff The biggest mistake I see is doing a review that focuses solely on the diff 2 . Most of the highest-impact code review comments have very little to do with the diff at all, but instead come from your understanding of the rest of the system. For instance, one of the most straightforwardly useful comments is ?you don't have to add this method here, since it already exists in this other place?. The diff itself won't help you produce a comment like this. You have to already be familiar with other parts of the codebase that the diff author doesn't know about. Likewise, comments like ?this code should probably live in this other file? are very helpful for maintaining the long-term quality of a codebase. The car...

10. Should LLMs just treat text content as an image?

Source: seangoedecke.com RSS feed | Published: Tue, 21 Oct 2025 00:00:00 GMT

Link: <https://seangoedecke.com/text-tokens-as-image-tokens/>

Several days ago, DeepSeek released a new OCR paper . OCR, or ?optical character recognition?, is the process of converting an image of text - say, a scanned page of a book - into actual text content. Better OCR is obviously relevant to AI because it unlocks more text data to train language models on 1 . But there's a more subtle reason why really good OCR might have deep implications for AI models. Optical compression According to the DeepSeek paper, you can pull out 10 text tokens from a single image token with near-100% accuracy. In other words, a model's internal representation of an image is ten times as efficient as its internal representation of text. Does this mean that models shouldn't consume text at all? When I paste a few paragraphs

into ChatGPT, would it be more efficient to convert that into an image of text before sending it to the model? Can we supply 10x or 20x more data to a model at inference time by supplying it as an image of text instead of text itself? This is called ?optical compression?. It reminds me of a funny idea from June of this year to save money on OpenAI transcriptions: before uploading the audio, run it through ffmpeg to speed it up by 2x. The mo...

11. February Pebble Production and Software Updates

Source: Eric Migicovsky's Blog RSS Feed | Published: Wed, 18 Feb 2026 00:00:00 GMT

Link: <https://repebble.com/blog/february-pebble-production-and-software-updates>

Mega update on Pebble Time 2, Pebble Round 2 and Index 01 Things are busy in Pebbleland! We're getting close to shipping 3 new hardware?

12. Beeper is joining Automattic

Source: Eric Migicovsky's Blog RSS Feed | Published: Tue, 09 Apr 2024 00:00:00 GMT

Link: <https://ericmigi.com/blog/beeper-is-joining-automattic>

(No summary)

13. Beeper And The Power Of Protocols - Podcast

Source: Eric Migicovsky's Blog RSS Feed | Published: Wed, 17 Jan 2024 00:00:00 GMT

Link: <https://ericmigi.com/blog/beeper-and-the-power-of-protocols-podcast>

(No summary)

14. Paul Ford: ?The A.I. Disruption Has Arrived, and It Sure Is Fun?

Source: Daring Fireball | Published: 2026-02-18T21:15:10Z

Link: https://www.nytimes.com/2026/02/18/opinion/ai-software.html?unlocked_article_code=1.NFA.djaw.TBIAp8kE_N-i

Paul Ford, in an op-ed for The New York Times (gift link): All of the people I love hate this stuff, and all the people I hate love it. And yet, likely because of the same personality flaws that drew me to technology in the first place, I am annoyingly excited. ?

15. Meta?s Response to Reuters Report on ?Romance AI Chatbots? for Teenagers

Source: Daring Fireball | Published: 2026-01-28T16:02:25Z

Link: <https://x.com/andymstone/status/2016242174447223049>

Andy Stone, VP of communications at Meta, responding, in a series of tweets on Twitter/X, to Jeff Horwitz?s report at Reuters yesterday, linked here last night, which claimed that ?Zuckerberg blocked curbs on sex-talking chatbots for minors?: Never let the facts get in the way of a good story, eh, @Reuters , @JeffHorwitz ! The documents you cite in the story itself contradict this headline. The headline says ?Zuckerberg blocked curbs on sex-talking chatbots for minors? But the story cites a document that says ?Zuckerberg believed that AI companions should be blocked from engaging in sexually ?explicit? conversations? w young people. Huh?! After my post last night, a friend of mine, with a career of experience working in a large company, sent me this: A word of caution. ?Scumbag middle manager says CEO said? is not the same as ?CEO said.? I could believe Zuck shitcanned parental controls, but I am certain there are thousands of snakes inside that company who would lie about it to get what they want. That?s a good and fair point, and I think it?s what Stone is trying to emphasize above. The New Mexico lawsuit filing doesn?t contain evidence that Zuckerberg nixed parental controls fo...

16. Court Filing Claims Zuckerberg Blocked Curbs at Meta on Sex-Talking Chatbots for Minors

Source: Daring Fireball | Published: 2026-01-28T00:17:02Z

Link: <https://www.reuters.com/legal/government/meta-ceo-zuckerberg-blocked-curbs-sex-talking-chatbots-minors-court-filing-2026-01-27/>

Jeff Horwitz, reporting for Reuters: Meta Chief Executive Mark Zuckerberg approved allowing minors to access AI chatbot companions that safety staffers warned were capable of sexual interactions, according to internal Meta documents filed in a New Mexico state court case and

made public Monday. The lawsuit ? brought by the state?s attorney general, Raul Torrez, and scheduled for trial next month ? alleges that Meta ?failed to stem the tide of damaging sexual material and sexual propositions delivered to children? on Facebook and Instagram. [...] Messages between two employees from March of 2024 state that Zuckerberg had rejected creating parental controls for the chatbots, and that staffers were working on ?Romance AI chatbots? that would be allowed for users under the age of 18. We ?pushed hard for parental controls to turn GenAI off ? but GenAI leadership pushed back stating Mark decision,? one employee wrote in that exchange. Horwitz was with The Wall Street Journal for a long time; his is a byline worth paying attention to . ?

17. Thinking Improves Thinking

Source: iDiallo.com | Published: Wed, 18 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/taking-our-mind-for-granted?src=feed>

How did we do it before ChatGPT? How did we write full sentences, connect ideas into a coherent arc, solve problems that had no obvious answer? We thought. That's it. We simply sat with discomfort long enough for something to emerge. I find this fascinating. You have a problem, so you sit down and think until you find a solution. Sometimes you're not even sitting down. You go for a walk, and your mind quietly wrestles with the idea while your feet carry you nowhere in particular. A solution emerges not because you forced it, but because you thought it through. What happened in that moment is remarkable: new information was created from the collision of existing ideas inside your head. No prompt. No query. Just you. I remember the hours I used to spend debugging a particularly stubborn problem at work. I would stare at the screen, type a few keystrokes, then delete them. I'd meet with our lead engineer and we would talk in circles. At home, I would lie in bed still turning the problem over. And then one night, somewhere around 3 a.m., I dreamt I was running the compiler, making a small change, watching it build, and suddenly it worked. I woke up knowing the answer before I had even...

18. Could Write>Process>Memory be made faster by avoiding the intermediate buffer?

Source: The Old New Thing | Published: Wed, 18 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260218-00/?p=112069>

I guess it could, but why bother? The post Could Write>Process>Memory be made faster by avoiding the intermediate buffer? appeared first on The Old New Thing .

19. Please take my weird moral puzzles quiz

Source: DYNOMIGHT | Published: 2025-06-17T00:00:00+00:00

Link: <https://dynamight.net/puzzles/>

For reasons, I ask that you take a short moral puzzles survey. I?ll provide 12 scenarios. For each of them, I?ll ask (1) What percentage of current Western adults you believe would agree, and (2) If you personally agree. Please don?t overthink. I?m not trying to trap you or make some kind of tricky point, I swear. You can go here to take the survey. Or, if you want to see what you?d be getting into, here are the puzzles, ordered roughly by increasing weirdness. Chickens Since male ?layer? chickens serve no economic purpose, each year seven billion are killed immediately after hatching, typically by grinding or asphyxiation. We now have the technology to prevent male chicks from being born by detecting their sex as eggs. This raises the cost per egg by around \$0.01. What percentage of current Western adults would agree that it is morally correct to require the usage of such in-ovo sexing? Hydrocephalus Suppose a woman wishes to have a baby and becomes pregnant. Near the end of the second term, the baby is diagnosed with hydrocephalus , which is correlated with intellectual disability and reduced lifespan. The mother wishes to abort the baby so she can have another without this cond...

20. Futarchy?s fundamental flaw

Source: DYNOMIGHT | Published: 2025-06-12T00:00:00+00:00

Link: <https://dynamight.net/futarchy/>

Say you?re Robyn Denholm , chair of Tesla?s board. And say you?re thinking about firing Elon Musk. One way to make up your mind would be to have people bet on Tesla?s stock price six

months from now in a market where all bets get cancelled unless Musk is fired . Also, run a second market where bets are cancelled unless Musk stays CEO . If people bet on higher stock prices in Musk-fired world, maybe you should fire him. That's basically Futarchy : Use conditional prediction markets to make decisions. People often argue about fancy aspects of Futarchy. Are stock prices all you care about? Could Musk use his wealth to bias the market? What if Denholm makes different bets in the two markets, and then fires Musk (or not) to make sure she wins? Are human values and beliefs somehow inseparable? My objection is more basic: It doesn't work. You can't use conditional predictions markets to make decisions like this, because conditional prediction markets reveal probabilistic relationships, not causal relationships. The whole concept is faulty. There are solutions?ways to force markets to give you causal relationships. But those solutions are painful and I get the shakes when I see everyone a...

21. Optimizing tea: An N=4 experiment

Source: DYNOMIGHT | Published: 2025-06-05T00:00:00+00:00

Link: <https://dynamight.net/tea/>

Tea is a little-known beverage, consumed for flavor or sometimes for conjectured effects as a stimulant. It's made by submerging the leaves of *C. Sinensis* in hot water. But how hot should the water be? To resolve this, I brewed the same tea at four different temperatures, brought them all to a uniform serving temperature, and then had four subjects rate them along four dimensions. Subjects Subject A is an experienced tea drinker, exclusively of black tea w/ lots of milk and sugar. Subject B is also an experienced tea drinker, mostly of black tea w/ lots of milk and sugar. In recent years, Subject B has been pressured by Subject D to try other teas. Subject B likes fancy black tea and claims to like fancy oolong, but will not drink green tea. Subject C is similar to Subject A. Subject D likes all kinds of tea, derives a large fraction of their joy in life from tea, and is world's preeminent existential angst + science blogger. Tea and brewing For a tea that was as ?normal? as possible, I used pyramidal bags of PG Tips tea (Lipton Teas and Infusions, Trafford Park Rd., Trafford Park, Stretford, Manchester M17 1NH, UK). I brewed it according to the instructions on the box, by submerg...

22. Intel's \$475 million error: the silicon behind the Pentium division bug

Source: Ken Shirriff's blog | Published: 2024-12-28T10:54:00.000-08:00

Link: <http://www.righto.com/2024/12/this-die-photo-of-pentium-shows.html>

In 1993, Intel released the high-performance Pentium processor, the start of the long-running Pentium line. The Pentium had many improvements over the previous processor, the Intel 486, including a faster floating-point division algorithm. A year later, Professor Nicely, a number theory professor, was researching reciprocals of twin prime numbers when he noticed a problem: his Pentium sometimes generated the wrong result when performing floating-point division. Intel considered this "an extremely minor technical problem", but much to Intel's surprise, the bug became a large media story. After weeks of criticism, mockery, and bad publicity, Intel agreed to replace everyone's faulty Pentium chips, costing the company \$475 million. In this article, I discuss the Pentium's division algorithm, show exactly where the bug is on the Pentium chip, take a close look at the circuitry, and explain what went wrong. In brief, the division algorithm uses a lookup table. In 1994, Intel stated that the cause of the bug was that five entries were omitted from the table due to an error in a script. However, my analysis shows that 16 entries were omitted due to a mathematical mistake in the definitio...

23. Prime gaps and Gapcoin

Source: John D. Cook | Published: Mon, 19 Jan 2026 01:20:19 +0000

Link: <https://www.johndcook.com/blog/2026/01/18/prime-gaps-gapcoin/>

The previous post looked at tightly clustered primes. This post looks at the opposite, large gaps between primes. Riecoin is a cryptocurrency that uses finding prime clusters as its proof of work task. Gapcoin uses finding prime gaps as its proof of work task. There's some nuance to defining prime gaps. It's trivial to produce [?] The post Prime gaps and Gapcoin first appeared on John D. Cook .

24. Prime clusters and Riecoin

Source: John D. Cook | Published: Sun, 18 Jan 2026 20:23:09 +0000

Link: <https://www.johndcook.com/blog/2026/01/18/prime-clusters-riecoin/>

Prime clusters are sets of primes that appear as close together as is generally possible. There is one pair of consecutive prime numbers, 2 and 3, but there cannot be any more: in any larger pair of consecutive numbers, one of the pair will be even. But there are a lot of twin primes, perhaps infinitely [?] The post Prime clusters and Riecoin first appeared on John D. Cook .

25. Efficiently testing multiple primes at once

Source: John D. Cook | Published: Fri, 16 Jan 2026 15:52:05 +0000

Link: <https://www.johndcook.com/blog/2026/01/16/testing-multiple-primes/>

The previous post looked at a technique for inverting multiple integers mod m at the same time, using fewer compute cycles than inverting each integer individually. This post will do something analogous for prime chains, revisiting a post from a few days ago about testing prime chains. A prime chain is a sequence of primes in [?] The post Efficiently testing multiple primes at once first appeared on John D. Cook .

26. A Chain Reaction

Source: overreacted ? A blog by Dan Abramov | Published: Mon, 11 Dec 2023 00:00:00 GMT

Link: <https://overreacted.io/a-chain-reaction/>

The limits of my language mean the limits of my world.

27. npm audit: Broken by Design

Source: overreacted ? A blog by Dan Abramov | Published: Wed, 07 Jul 2021 00:00:00 GMT

Link: <https://overreacted.io/npm-audit-broken-by-design/>

Found 99 vulnerabilities (84 moderately irrelevant, 15 highly irrelevant)

28. Before You memo()

Source: overreacted ? A blog by Dan Abramov | Published: Tue, 23 Feb 2021 00:00:00 GMT

Link: <https://overreacted.io/before-you-memo/>

Rendering optimizations that come naturally.

29. Anubis v1.25.0: Necron

Source: Xe Iaso's blog | Published: Wed, 18 Feb 2026 00:00:00 GMT

Link: <https://github.com/TecharoHQ/anubis/releases/tag/v1.25.0>

Hey all, I'm sure you've all been aware that things have been slowing down a little with Anubis development, and I want to apologize for that. A lot has been going on in my life lately (my blog will have a post out on Friday with more information), and as a result I haven't really had the energy to work on Anubis in publicly visible ways. There are things going on behind the scenes, but nothing is really shippable yet, sorry! I've also been feeling some burnout in the wake of perennial waves of anger directed towards me. I'm handling it, I'll be fine, I've just had a lot going on in my life and it's been rough. I've been missing the sense of wanderlust and discovery that comes with the artistic way I playfully develop software. I suspect that some of the stresses I've been through (setting up a complicated surgery in a country whose language you aren't fluent in is kind of an experience) have been sapping my energy. I'd gonna try to mess with things on my break, but realistically I'm probably just gonna be either watching Stargate SG-1 or doing unreasonable amounts of ocean fishing in Final Fantasy 14. Normally I'd love to keep the details about my medical state fairly private, bu...

30. Book Review: All Systems Red - The Murderbot Diaries by Martha Wells ?????

Source: Terence Eden's Blog | Published: Wed, 18 Feb 2026 12:34:05 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-all-systems-red-the-murderbot-diaries-by-martha-wells/>

Everyone raves about this series, so I thought I'd grab the first book. It's basically fine, I

guess. It is moderately amusing having the Muderbot be an awkward teenage boy who just wants to watch videos and cringes when people stare at him. But it is a bit one-note. Similarly, evil corporations hiding details from exo-planet surveyors is a trope which has been a thousand times before. This?

31. The Urban Family Exodus Is a Warning for Progressives

Source: Derek Thompson | The Atlantic | Published: 2024-08-05T06:00:00-04:00

Link: https://www.theatlantic.com/ideas/archive/2024/08/the-urban-family-exodus-is-a-warning-for-progressives/679350/?utm_source=feed

This is Work in Progress, a newsletter about work, technology, and how to solve some of America's biggest problems. Sign up here . Children?and the millions of private decisions to have or not have them?are in the news these days, for regrettable reasons. Ohio Senator J. D. Vance, the Republican vice-presidential candidate, has made a habit of excoriating progressives who don?t have a record of procreation. In November 2020, he implied that childless Democratic leaders are ? sociopathic .? In an interview with the Fox News host Tucker Carlson in 2021, he lamented that the country was being run by ? a bunch of childless cat ladies .? Later that year, in an address in Southern California, he said he wanted ?to take aim at the left, specifically the childless left ? because I think the rejection of the American family is perhaps the most pernicious and most evil thing that the left has done in this country.? Vance?s commentary is rude and revisionist. Childless adults aren?t psychotic, and many childless people are desperate to bear children. Suggesting that their unsuccessful reproductive efforts amount to sociopathy is cruel. More substantively, in 2022, it was progressive Democra...

32. What Package Registries Could Borrow from OCI

Source: Andrew Nesbitt | Published: 2026-02-18T00:00:00+00:00

Link: <https://nesbitt.io/2026/02/18/what-package-registries-could-borrow-from-oci.html>

Every package manager ships code as an archive, and every one of them has a slightly different way to do it. npm wraps tarballs in a package/ directory prefix. RubyGems nests gzipped files inside an uncompressed tar. Alpine concatenates three gzip streams and calls it a package. Python cycled through four distribution formats in twenty years. RPM used cpio as its payload format for nearly three decades before finally dropping it in 2025. Meanwhile, the container world converged on a single format: OCI, the Open Container Initiative spec. And over the past few years, OCI registries have quietly started storing things that aren?t containers at all: Helm charts, Homebrew bottles, WebAssembly modules, AI models. The format was designed for container images, but the underlying primitives turn out to be general enough that it?s worth asking whether every package manager could use OCI for distribution. What OCI actually is OCI defines three specifications: a Runtime Spec (how to run containers), an Image Spec (how to describe container contents), and a Distribution Spec (how to push and pull from registries). At the storage level, an OCI registry deals in two primitives: manifests and bl...

33. Here we go

Source: Rakhim's blog | Published: Thu, 31 Oct 2024 17:18:12 GMT

Link: <https://rakhim.exotext.com/here-we-go>

I just wanted to write again. My old blog at rakhim.org is built with Hugo, and after installing a new version of Hugo things broke. After 25 minutes of attempts to fix it I thought "yeah, that's it", and decided (as you do) to just build a simple blogging platform. Static site generators are great. Hugo though has too many features I never use. But my overall issue with static site generators is UX. Unlike, apparently, many others in tech circles, I don't enjoy writing in my normal code editor, nor do I find joy in publishing via git, and not having simple image uploads. Sure, there are fully fledged blog engines like Ghost and Wordpress (with its drama). I've used both for a long time in the past in various projects, and was generally ok with either. Today, ghost is trying really hard to become a "publishing" thing, with its magazine-like themes, paid memberships, and newsletters. Wordpress is just scary, to be honest (same vibe as Java or Kubernetes). There are very nice, small platforms like Bear , Pika ,

Mataroa . Each would actually be pretty much perfect for me. But I always wanted to build my own platform, with some specific features and UX nuances. So, yeah. I'm starting...

34. The mandatory smartphone

Source: Rakhim's blog | Published: Tue, 18 Aug 2020 00:00:00 GMT

Link: <https://rakhim.exotext.com/the-mandatory-smartphone>

I hate to use a smartphone. Being able to view maps, look up info on the go and listen to music is great, but I always felt that smartphones are like plastic utensils. Very useful on the go, but temporary in nature. A compromise. Using a smartphone for more than a few minutes in a row feels wrong. Wrong because either I'm trying to do something I should've done before or should do later, at my normal-sized computing device. Or, more often, because I should not do this thing at all. Using a smartphone makes me physically and mentally uncomfortable, and noticing it makes things even worse. Like eating from a PVC container with a single-use plastic fork while sitting in a full-equipped kitchen. Half of the reason I dislike using any smartphone is the operating system. Both major OSes are an insult to human bodies and ergonomics, mental models and years of research of human-computer interaction. (Have you ever though how "iOS and Android" sounds and feels so similar to American politics? Two parties, an illusion of choice, both crappy in their own unique ways.) It's really unfortunate that the promised personal computer turned out to be a half-baked proprietary attention-grabbing devi...

35. It's okay to walk uphill

Source: Rakhim's blog | Published: Thu, 23 Jul 2020 12:14:00 GMT

Link: <https://rakhim.exotext.com/its-okay-to-walk-uphill>

At first I cycled up the hill with all my force trying to impress the world. My athletic sweat was the only thing that showed the effort. My face was trying to do the opposite, to convince the unseen viewer that the ascent is effortless, easy even. Then I started noticing it. Who is the viewer? Am I in a movie? Sometimes, there aren't even people around. This is madness. I can convince myself that the unseen viewer is a good abstraction, a tool to force better action. Is that naive religion? Then I cycled up the hill with all my force trying to impress myself. Clearly, this is the next level. Who cares about others, I need this for myself. I can become better, stronger, faster without external pressure. Is that true religion? Later still, I cycled up the hill with all my force, yet with no motivation. I killed the uphill ego twice now. Clearly, this is the final level. "For you" and "for them" are different, but the problem is not the object of this anchoring, but the anchoring itself. Eliminate the "for"! I can struggle up the hill for nothing. Is that true spirituality? Finally, this morning I stopped cycling and walked up the hill. Two birds landed in the tree. The wind was war...

36. Stream of Consciousness Driven Development

Source: Computer Things | Published: Wed, 18 Feb 2026 16:33:08 +0000

Link: <https://buttondown.com/hillelwayne/archive/stream-of-consciousness-driven-development/>

This is something I just tried out last week but it seems to have enough potential to be worth showing unpolished. I was pairing with a client on writing a spec. I saw a problem with the spec, a convoluted way of fixing the spec. Instead of trying to verbally explain it, I started by creating a new markdown file: NameOfProblem.md Then I started typing. First the problem summary, then a detailed description, then the solution and why it worked. When my partner asked questions, I incorporated his question and our discussion of it into the flow. If we hit a dead end with the solution, we marked it out as a dead end. Eventually the file looked something like this: Current state of spec Problems caused by this Elaboration of problems What we tried that didn't work Proposed Solution Theory behind proposed solution How the solution works Expected changes Other problems this helps solve Problems this does *not* help with Only once this was done, my partner fully understood the chain of thought, and we agreed it represented the right approach, did we start making changes to the spec. How is this better than just making the change? The change was conceptually complex. A rough analogy: imagi...

37. Solving LinkedIn Queens with SMT

Source: Computer Things | Published: Thu, 12 Jun 2025 15:43:25 +0000

Link: <https://buttondown.com/hillelwayne/archive/solving-linkedin-queens-with-smt/>

No newsletter next week I'll be speaking at Systems Distributed . My talk isn't close to done yet, which is why this newsletter is both late and short. Solving LinkedIn Queens in SMT The article Modern SAT solvers: fast, neat and underused claims that SAT solvers 1 are "criminally underused by the industry". A while back on the newsletter I asked "why": how come they're so powerful and yet nobody uses them? Many experts responded saying the reason is that encoding SAT kinda sucked and they rather prefer using tools that compile to SAT. I was reminded of this when I read Ryan Berger's post on solving ?LinkedIn Queens? as a SAT problem. A quick overview of Queens. You're presented with an NxN grid divided into N regions, and have to place N queens so that there is exactly one queen in each row, column, and region. While queens can be on the same diagonal, they cannot be adjacently diagonal. (Important note: LinkedIn ?Queens? is a variation on the puzzle game Star Battle , which is the same except the number of stars you place in each row/column/region varies per puzzle, and is usually two. This is also why 'queens' don't capture like chess queens.) Ryan solved this by writing Queens...

38. AI is a gamechanger for TLA+ users

Source: Computer Things | Published: Thu, 05 Jun 2025 14:59:11 +0000

Link: <https://buttondown.com/hillelwayne/archive/ai-is-a-gamechanger-for-tla-users/>

New Logic for Programmers Release v0.10 is now available ! This is a minor release, mostly focused on logic-based refactoring, with new material on set types and testing refactors are correct. See the full release notes at the changelog page . Due to conference pressure v0.11 will also likely be a minor release. AI is a gamechanger for TLA+ users TLA+ is a specification language to model and debug distributed systems. While very powerful, it's also hard for programmers to learn, and there's always questions of connecting specifications with actual code. That's why The Coming AI Revolution in Distributed Systems caught my interest. In the post, Cheng Huang claims that Azure successfully used LLMs to examine an existing codebase, derive a TLA+ spec, and find a production bug in that spec. "After a decade of manually crafting TLA+ specifications", he wrote, "I must acknowledge that this AI-generated specification rivals human work". This inspired me to experiment with LLMs in TLA+ myself. My goals are a little less ambitious than Cheng's: I wanted to see how LLMs could help junior specifiers write TLA+, rather than handling the entire spec automatically. Details on what did and didn'...

39. The case for gatekeeping, or: why medieval guilds had it figured out

Source: Westenberg. | Published: Wed, 18 Feb 2026 02:21:09 GMT

Link: <https://www.joanwestenberg.com/the-case-for-gatekeeping-or-why-medieval-guilds-had-it-figured-out/>

Every open source maintainer I've talked to in the last six months has the same complaint: the absolute flood of mass-produced, AI-generated, mass-submitted slop requests have turned their repositories into a slush pile. The contributions look like contributions, they have commit messages, they reference issues and they follow

40. The Airbnbs

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/airbnbs.html>

(No summary)

41. How to Think for Yourself

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/think.html>

(No summary)

42. Early Work

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/early.html>

(No summary)

43. A Few Rambling Observations on Care

Source: Jim Nielsen's Blog | Published: Wed, 18 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/observations-on-care/>

In this new AI world, ?taste? is the thing everyone claims is the new supreme skill. But I think ?care? is the one I want to see in the products I buy. Can you measure care? Does scale drive out care? If a product conversation is reduced to being arbitrated exclusively by numbers, is care lost? The more I think about it, care seems antithetical to the reductive nature of quantification ? ?one death is a tragedy, one million is a statistic?. Care considers useful, constructive systematic forces ? rules, processes, etc. ? but does not take them as law. Individual context and sensitivity are the primary considerations. That?s why the professional answer to so many questions is: ?it depends?. ?This is the law for everyone, everywhere, always? is not a system I want to live in. Businesses exist to make money, so one would assume a business will always act in a way that maximizes the amount of money that can be made. That?s where numbers take you. They let you measure who is gaining or losing the most quantifiable amount in any given transaction. But there?s an unmeasurable, unquantifiable principle lurking behind all those numbers: it can be good for business to leave money on the tabl...

44. Weather Forecasts Are Punctuated

Source: Entropic Thoughts | Published: Tue, 28 Oct 2025 00:00:00 +0100

Link: <https://entropichoughts.com/weather-forecasts-are-punctuated>

(No summary)

45. Markdown?s Moment

Source: Tedium: The Dull Side of the Internet. | Published: 2026-02-18T04:01:53Z

Link: <https://feed.tedium.co/link/15204/17278321/markdown-growing-influence-cloudflare-ai>

For some reason, a bunch of big companies are really leaning into Markdown right now. AI may be the reason, but I kind of love the possible side benefits.

46. Terrible Horrible No Good Very Bad Python

Source: the website of jyn | Published: 2025-03-28T00:00:00+00:00

Link: <https://jyn.dev/terrible-horrible-no-good-very-bad-python/>

new and exciting ways to write buggy code

47. Scheduling Do-Not-Disturb in GNOME

Source: the website of jyn | Published: 2025-02-22T00:00:00+00:00

Link: <https://jyn.dev/do-not-disturb-in-gnome/>

Do Not Disturb GNOME has a little button that lets you turn on Do-Not-Disturb for notifications: Unfortunately, it has no way of scheduling DnD . Good news, though! It does support turning on DnD through the CLI: gsettings set org.gnome.desktop.notifications show-banners false . I put that in a script named toggle-dnd in my dotfiles: \$ cat bin/toggle-dnd #!/bin/sh case \${ 1 :- } in # "show-banners" is reversed from what you would expect "do not disturb" to mean true) new = false ;; false) new = true ;; *) if ["\$" \$(gsettings get org.gnome.desktop.notifications show-banners) " = true] then new = false else new = true fi ;; esac gsettings set org.gnome.desktop.notifications show-banners \$ new scheduling I tried putting that in cron 1 , had a sneaking suspicion it wouldn't work, set it to run every minute, and saw this very unhelpful line of logging: \$ journalctl -- unit cron -- since ' 5m ago ' Feb 22 12:00:01 pop-os CRON[1623131]: (CRON) info (No MTA installed, discarding output) Ok, fine. Let's pipe the output to the system log 2 , since clearly cron can't handle that itself. * * * * bash -lc 'org.gnome.desktop.notifications show-banners false 2>&1 | logger -t toggle-dnd...

48. building your own rustc_driver

Source: the website of jyn | Published: 2024-10-24T00:00:00+00:00

Link: <https://jyn.dev/rustc-driver/>

what happens when you run cargo clippy ?

49. Programming and Writing

Source: antirez.com | Published: Fri, 14 May 2021 11:47:18 +0200

Link: <http://antirez.com/news/135>

One year ago I paused my programming life and started writing a novel, with the illusion that my new activity was deeply different than the previous one. A river of words later, written but more often rewritten, I'm pretty sure of the contrary: programming big systems and writing novels have many common traits and similar processes. The most obvious parallel between the two activities is that in both of them you write something. Code is not prose written in a natural language, yet it has a set of fixed rules (a grammar), certain forms that most programmers will understand as natural and others that, while formally correct, will sound hard to grasp. There is, however, a much deeper connection between the two activities: a good program and a good novel are both the sum of local and global elements that work well. Good code must be composed of well written and readable single statements, but overall the different parts of the program must be orthogonal, designed in a coherent way, and have clean interactions. A good novel must also succeed in the same two scales of the micro and the macro. Sentences must be well written, but the overall structure and relationship between the parts is...

50. The open source paradox

Source: antirez.com | Published: Sat, 03 Oct 2020 11:11:59 +0200

Link: <http://antirez.com/news/134>

A new idea is insinuating in social networks and programming communities. It's the proportionality between the money people give you for coding something, and the level of demand for quality they can claim to have about your work. As somebody said, the best code is written when you are supposed to do something else [1]. Like a writer will do her best when writing that novel that, maybe, nobody will pay a single cent for, and not when doing copywriting work for a well known company, programmers are likely to spend more energies in their open source side projects than during office hours, while writing another piece of a project they feel stupid, boring, pointless. And, if the company is big enough, chances are it will be cancelled in six months anyway or retired one year after the big launch. Open source is different, it's an artifact, it's a transposition in code of what you really want to do, of what you feel software should be, or just of all your fun and joy, or even anger you are feeling while coding. And you want it to rock, to be perfect, and you can't sleep at night if there is a fucking heisenbug. So if a user of your software is addressing you because some part of your co...

51. The end of the Redis adventure

Source: antirez.com | Published: Tue, 30 Jun 2020 15:00:16 +0200

Link: <http://antirez.com/news/133>

When I started the Redis project more than ten years ago I was in one of the most exciting moments of my career. My co-founder and I had successfully launched two of the major web 2.0 services of the Italian web. In order to make them scalable we had to invent many new concepts, that were already known in the field most of the times, but we didn't know, nor we cared to check. Problem? Let's figure out a solution. We wanted to solve problems but we wanted, even more, to have fun. This was the playful environment where Redis was born. But now Redis is, incredibly, one of the main parts of so many things. And year after year my work changed from building this thing to making sure that it was also as useful as possible, as reliable as possible. And in recent years, what I do every day changed so much that most of my attention is spent in checking what other developers tell me about the Redis code, how to improve it, the changes it requires to be more correct or faster or more secure. However I never wanted to be a software maintainer. I write code in order to express myself, and I consider what I code an

artifact, rather than just something useful to get things done. I would say that...

52. From ZNC to Soju

Source: Susam Pal | Published: Thu, 12 Feb 2026 00:00:00 +0000

Link: <https://susam.net/from-znc-to-soju.html>

I have recently switched from ZNC to Soju as my IRC bouncer and I am already quite pleased with it. I usually run my bouncer on a Debian machine, where Soju is well packaged and runs smoothly right after installation. By contrast, the ZNC package included with Debian 13 (Trixie) and earlier fails to start after installation because of a missing configuration file. As a result, I was forced to maintain my own configuration file along with a necessary PEM bundle, copy them to the Debian system and carefully set the correct file permissions before I could run ZNC successfully. None of this is necessary with Soju, since installing it from the Debian package repository automatically sets up the configuration and certificate files. I no longer have to manage any configuration or certificate files myself. Setup It is quite straightforward to install and set up Soju on Debian. The following two commands install Soju: sudo apt-get update sudo apt-get -y install soju Then setting up an IRC connection involves another two commands: sudo sojuctl user create -username soju -password YOUR_SOJU_PASSWORD sudo sojuctl user run soju network create -name bnc1 -addr irc.libera.chat -nick YOUR_NICK -p...

53. Forty years of programming

Source: Fabien Sanglard | Published: 08 Oct 2023 00:00:00 +0000

Link: <https://fabiensanglard.net/40/index.html>

(No summary)

54. Exploring Command-line space time

Source: Fabien Sanglard | Published: 26 Sep 2023 00:00:00 +0000

Link: <https://fabiensanglard.net/st/index.html>

(No summary)

55. Ode to the M1

Source: Fabien Sanglard | Published: 12 Aug 2023 00:00:00 +0000

Link: <https://fabiensanglard.net/m1/index.html>

(No summary)

56. What went wrong with wireless USB

Source: Old Vintage Computing Research | Published: 2025-05-03T18:54:00.000-07:00

Link: <https://oldvcr.blogspot.com/2025/05/what-went-wrong-with-wireless-usb.html>

(Hat tip to the late Bill Strauss and The Capitol Steps' Lirty Dies.) A few moons ago I was thinking of ways to take my Palm OS Fossil Wrist PDA smartwatch mobile . It has no on-board networking libraries but can be coerced into doing PPP over its serial port (via USB) by using the libraries from my Palm m505. Of course, that then requires it be constantly connected to a USB port, which is rather inconvenient for a wristwatch. But what if the USB connection could be made wirelessly ? For a few years, real honest-to-goodness wireless USB devices were actually a thing. Competing standards led to market fracture and the technologies fizzled out relatively quickly in the marketplace, but like the parallel universe of FireWire hubs there was another parallel world of wireless USB devices, at least for a few years. As it happens, we now have a couple of them here, so it's worth exploring what wireless USB was and what happened to it, how the competing standards worked (and how well), and if it would have helped. You can probably blame Wi-Fi for this: while early patents for Wi-Fi existed as far back as 1991, after the introduction of 802.11 in 1997 and Apple's use of 802.11b for the iBo...

57. the-good-and-bad-of-cpp-as-a-rust-dev

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/pltd/the-good-and-bad-of-cpp-as-a-rust-dev>

(No summary)

58. how-i-think-of-the-expression-problem

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/pltd/how-i-think-of-the-expression-problem>

(No summary)

59. openai-structured-outputs-are-really-useful

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/ai/openai-structured-outputs-are-really-useful>

(No summary)

60. The Government Wants Your Selfie to Use Instagram

Source: Dr Paris Buttfield-Addison | Published: Sat, 18 Oct 2025 00:00:00 +0000

Link: <https://hey.paris/posts/socialmediaban/>

Australia has a new plan to keep kids safe online. On the surface, it sounds simple, maybe even sensible. The government is banning anyone under sixteen from having a social media account. From December, platforms like Instagram and TikTok will have to kick the kids off in an effort to stop social harm. It's a nice idea. But when you pull back the curtain, you find a privacy disaster waiting to happen. It's a plan that won't work, and it will end up hurting the very people it's supposed to protect.

61. Space News, October 2025

Source: Dr Paris Buttfield-Addison | Published: Sun, 05 Oct 2025 00:00:00 +0000

Link: <https://hey.paris/posts/spacenews-october-2025/>

These are my notes and expanded thoughts from this month's Space News segment on ABC Radio Hobart and ABC Northern Tasmania. Every month I join Lucie Cutting on Sunday mornings to chat about what's happening above and beyond. Here's what we discussed, plus some of my notes, and expanded thoughts from this edition of the programme : The World's Space Community Came to Sydney The 76th International Astronautical Congress wrapped up in Sydney this month, bringing thousands and thousands of attendees from most countries. It's the largest global gathering of space professionals, and this year's theme was ?Sustainable Space: Resilient Earth.?

62. Space News, September 2025

Source: Dr Paris Buttfield-Addison | Published: Sun, 07 Sep 2025 00:00:00 +0000

Link: <https://hey.paris/posts/spacenews-september-2025/>

These are my notes and expanded thoughts from this month's Space News segment on ABC Radio Hobart and ABC Northern Tasmania. Every month I join Lucie Cutting on Sunday mornings to chat about what's happening above and beyond. Here's what we discussed, plus some of my notes, from this edition of the programme : Space Command: Politics Over Strategy President Trump announced he's moving US Space Command from Colorado Springs to Huntsville, Alabama ?reversing the Biden administration's decision and bringing domestic politics into military space strategy. This isn't just bureaucratic reshuffling, it's a \$1 billion economic decision affecting 1,400 jobs and America's space defence capabilities.

63. Minimax

Source: Beej's Bit Bucket | Published: Sun, 19 Feb 2012 00:00:00 +0000

Link: <http://beej.us/blog/data/minimax/>

Lookahead AI, just like what powers the WOPR

64. Adobe AIR

Source: Beej's Bit Bucket | Published: Wed, 17 Feb 2010 00:00:00 +0000

Link: <http://beej.us/blog/data/adobe-air/>

Very very short getting-started document.

65. Image Interlacing

Source: Beej's Bit Bucket | Published: Sat, 13 Feb 2010 00:00:00 +0000

Link: <http://beej.us/blog/data/image-interlacing/>

(No summary)

66. You Only Think They Work For You

Source: Steve Blank | Published: Wed, 18 Feb 2026 14:00:39 +0000

Link: <https://steveblank.com/2026/02/18/you-only-think-they-work-for-you/>

When I was a new VP of Marketing I got a painful lesson of who my PR (Public Relations) agency actually worked for. Later I realized that it was true for all of my external vendors. And much later I realized what I really should have been asking them to do. The lessons still apply [?]

67. "Full-stack" and why I don't like it

Source: Tom Renner on My place to put things | Published: Mon, 10 Oct 2016 00:00:00 +0000

Link: <https://tomrenner.com/posts/fullstack/>

Last week I went to a jobs event, recruiting for my company. I was there on my own, and recruiting is a pretty new experience for me, so I was kinda excited about it. The attendees were a mix of graduates, bootcampers, and a few more senior developers, but mostly the crowd were looking for their first or second job. I enjoy going out and talking to people about their experiences with software ? I think breadth of knowledge is really valuable ? so I found it interesting seeing what other people are looking for in a job. I?ve only ever had one; maybe I?ve aimed for all the wrong things!

68. Notes from MCP Dev Summit Europe: Where the Protocol Is Headed

Source: Martin Alderson | Published: Thu, 02 Oct 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/notes-from-mcp-europe/?utm_source=rss

Insights from MCP Dev Summit Europe on agentic discovery, client compatibility challenges, and the emerging field of agentic experience design

69. How I make CI/CD (much) faster and cheaper

Source: Martin Alderson | Published: Sun, 28 Sep 2025 00:00:00 GMT

Link: https://martinalderson.com/posts/how-i-make-cicd-much-faster-and-cheaper/?utm_source=rss

Why GitHub Actions runners are slow and how bare metal servers can make your CI/CD 2-10x faster while costing 10x less

70. Google AI Studio API has been unreliable for the past 2 weeks

Source: Martin Alderson | Published: Wed, 24 Sep 2025 00:00:00 GMT

Link:

https://martinalderson.com/posts/google-ai-studio-api-unreliable-for-two-weeks/?utm_source=rss

Google's Gemini AI Studio API has been suffering from severe reliability issues with little transparency about the problems on their status page.

71. That One Time Keygen Went Down for 5 Hours (Twice)

Source: Keygen Blog | Published: Wed, 21 Feb 2024 06:00:00 GMT

Link: <https://keygen.sh/blog/that-one-time-keygen-went-down-for-5-hours-twice/>

A postmortem of the dashboard and API outage that occurred on February 5th and 6th, 2024.

72. Announcing Multi-User Licenses

Source: Keygen Blog | Published: Mon, 29 Apr 2024 05:00:00 GMT

Link: <https://keygen.sh/blog/announcing-multi-user-licenses/>

Introducing support for multi-user licenses in our software licensing API.

73. Your 14-Day Free Trial Ain't Gonna Cut It

Source: Keygen Blog | Published: Fri, 03 May 2024 05:00:00 GMT

Link: <https://keygen.sh/blog/your-14-day-free-trial-aint-gonna-cut-it/>

Ramblings about common free trial lengths and time-to-convert.

74. How did we end up threatening our kids' lives with AI?

Source: Anil Dash | Published: 2026-02-18T00:00:00Z

Link: <https://anildash.com/2026/02/18/threatening-kids-with-AI/>

I have to begin by warning you about the content in this piece; while I won't be dwelling on any specifics, this will necessarily be a broad discussion about some of the most disturbing topics imaginable. I resent that I have to give you that warning, but I'm forced to because of the choices that the Big AI companies have made that affect children. I don't say this lightly. But this is the point we must reckon with if we are having an honest conversation about contemporary technology. Let me get the worst of it out of the way right up front, and then we can move on to understanding how this happened. ChatGPT has repeatedly produced output that encouraged and incited children to end their own lives. Grok's AI generates sexualized imagery of children, which the company makes available commercially to paid subscribers. It used to be that encouraging children to self-harm, or producing sexualized imagery of children, were universally agreed upon as being amongst the worst things one could do in society. These were among the rare truly non-partisan, unifying moral agreements that transcended all social and cultural barriers. And now, some of the world's biggest and most powerful compa...

75. Sports vs Games

Source: Grant Slatton's Blog | Published: Thu, 26 Sep 2024 03:32:34 +0000

Link: <https://grantslatton.com/sports-vs-games>

An aesthetic distinction

76. Nobody Cares

Source: Grant Slatton's Blog | Published: Wed, 25 Sep 2024 04:54:18 +0000

Link: <https://grantslatton.com/nobody-cares>

A rant about caring

77. Lightweight property-based testing at Row Zero

Source: Grant Slatton's Blog | Published: Wed, 4 Sep 2024 04:31:00 +0000

Link: <https://grantslatton.com/rowzero-property-testing>

How we verify correctness

78. History of PuTTY's development

Source: Simon Tatham's quasi-blog | Published: 2020-03-24T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/putty-history/>

An overview of PuTTY's development from its genesis in 1996 to the state of affairs in 2020.

Written in response to an email query.

79. Writing a soluble-grid generator for Mines

Source: Simon Tatham's quasi-blog | Published: 2019-08-26T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/mines-solver/>

A description of the technique for generating guaranteed-soluble Minesweeper grids in my puzzle game Mines . Written in response to an email query.

80. Writing a solver for Net

Source: Simon Tatham's quasi-blog | Published: 2019-03-27T00:00:00+00:00

Link: <https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/net-solver/>

A description of the automated solver in my puzzle game Net . Written in response to an email query.

81. Je kunt je digitale autonomie niet kopen in Amerika

Source: Bert Hubert's writings | Published: Sat, 27 Sep 2025 14:58:59 +0200

Link: <https://berthub.eu/articles/posts/je-kunt-je-digitale-autonomie-niet-kopen-in-amerika/>

Gelukkig staat digitale autonomie nu overal stevig op de agenda. Het idee dat we in Nederland

en Europa al ons computerwerk over moeten laten aan Amerikaanse big tech en hyperscalers is nu op z'n minst controversieel. Moet dat nou? Want het is toch raar dat we alleen nog onze belastingdienst draaiend kunnen houden als Amerika met alles mee kan kijken, en ook ieder moment de stekker uit de dienstverlening kan trekken.

82. Johan Remkes: Het kon minder

Source: Bert Hubert's writings | Published: Sun, 21 Sep 2025 12:55:26 +0200

Link: <https://berthub.eu/articles/posts/johan-remkes-het-kon-minder/>

Net uitgelezen, het handzame boekje ?Het kon minder, de lessen van Johan Remkes? (bol.com e-book). Geschreven door Thijs Broer en Peter Kee die samen toch al heel wat van de politieke wereld gezien en besproken hebben. Het boek is geen biografie, verwacht geen stapels voetnoten en analyses. Maar dat belooft de intro ook al niet. ??Een boek over mij, dat hoeft niet zo nodig,? zei Remkes. Pas toen we hem verzekерden dat het geen biografie moest worden, maar dat we wilden beschrijven hoe hij het had gebracht tot veelgevraagd probleemoplosser in dit stuurloze land [?] en wat anderen daar misschien van konden leren, pas toen bromde hij: ?Laten we dan maar eens praten??.

83. Microstacks or megadependencies over at Webdevcon 2025

Source: Bert Hubert's writings | Published: Sat, 20 Sep 2025 17:20:23 +0200

Link: <https://berthub.eu/articles/posts/microstacks-or-megadependencies-webdev/>

tl;dr: A meandering talk that reminds us of the risks of shipping unknown build time dependencies (as from npm), and also highlights how many services now have runtime third party service dependencies. These are very serious choices about the future of a project, and we can wonder how well we are making these choices. Is anyone thinking about the long term costs and impact? The talk then covers ?microstacks?, whereby a project relies in a minimal number of build dependencies, and no runtime service dependencies.

84. Making SSH host certificates more usable

Source: Matthew Garrett | Published: Tue, 19 Dec 2023 19:48:14 GMT

Link: <https://mjj59.dreamwidth.org/68721.html>

Earlier this year, after Github accidentally committed their private RSA SSH host key to a public repository , I wrote about how better support for SSH host certificates would allow this sort of situation to be handled in a user-transparent way without any negative impact on security. I was hoping that someone would read this and be inspired to fix the problem but sadly that didn't happen so I've actually written some code myself. The core part of this is straightforward - if a server presents you with a certificate associated with a host key, then make the trust in that host be whoever signed the certificate rather than just trusting the host key. This means that if someone needs to replace the host key for any reason (such as, for example, them having published the private half), you can replace the host key with a new key and a new certificate, and as long as the new certificate is signed by the same key that the previous certificate was, you'll trust the new key and key rotation can be carried out without any user errors. Hurrah! So obviously I wrote that bit and then thought about the failure modes and it turns out there's an obvious one - if an attacker obtained both the pri...

85. Why does Gnome fingerprint unlock not unlock the keyring?

Source: Matthew Garrett | Published: Tue, 05 Dec 2023 06:32:10 GMT

Link: <https://mjj59.dreamwidth.org/68537.html>

There's a decent number of laptops with fingerprint readers that are supported by Linux, and Gnome has some nice integration to make use of that for authentication purposes. But if you log in with a fingerprint, the moment you start any app that wants to access stored passwords you'll get a prompt asking you to type in your password, which feels like it somewhat defeats the point. Mac users don't have this problem - authenticate with TouchID and all your passwords are available after login. Why the difference? Fingerprint detection can be done in two primary ways. The first is that a fingerprint reader is effectively just a scanner - it passes a graphical representation of the fingerprint back to the OS and the OS decides whether or not it

matches an enrolled finger. The second is for the fingerprint reader to make that determination itself, either storing a set of trusted fingerprints in its own storage or supporting being passed a set of encrypted images to compare against. Fprint supports both of these, but note that in both cases all that we get at the end of the day is a statement of "The fingerprint matched" or "The fingerprint didn't match" - we can't associate anything else...

86. Why ACPI?

Source: Matthew Garrett | Published: Wed, 01 Nov 2023 06:30:06 GMT

Link: <https://mjjg59.dreamwidth.org/68350.html>

"Why does ACPI exist" -- the greatest thread in the history of forums, locked by a moderator after 12,239 pages of heated debate, wait no let me start again. Why does ACPI exist? In the beforetimes power management on x86 was done by jumping to an opaque BIOS entry point and hoping it would do the right thing. It frequently didn't. We called this Advanced Power Management (Advanced because before this power management involved custom drivers for every machine and everyone agreed that this was a bad idea), and it involved the firmware having to save and restore the state of every piece of hardware in the system. This meant that assumptions about hardware configuration were baked into the firmware - failed to program your graphics card exactly the way the BIOS expected? Hurrah! It's only saved and restored a subset of the state that you configured and now potential data corruption for you. The developers of ACPI made the reasonable decision that, well, maybe since the OS was the one setting state in the first place, the OS should restore it. So far so good. But some state is fundamentally device specific, at a level that the OS generally ignores. How should this state be managed?

O...