RSS Digest
Generated at: 2026-02-20 09:38:42

1. Beeper ? Moving Forward
Source: Eric Migicovsky’s Blog RSS Feed | Published: Thu, 21 Dec 2023 00:00:00 GMT
Link: https://ericmigi.com/blog/beeper-moving-forward
(No summary)

2. Introducing Beeper Mini ? get blue bubbles on Android ?
Source: Eric Migicovsky’s Blog RSS Feed | Published: Tue, 05 Dec 2023 00:00:00 GMT
Link: https://ericmigi.com/blog/introducing-beeper-mini-get-blue-bubbles-on-android
(No summary)

3. The better, smarter future of messaging | The Vergecast
Source: Eric Migicovsky’s Blog RSS Feed | Published: Mon, 16 Oct 2023 00:00:00 GMT
Link: https://ericmigi.com/blog/the-better-smarter-future-of-messaging-or-the-vergecast
(No summary)

4. Gemini 3.1 Pro
Source: Simon Willison’s Weblog | Published: 2026-02-19T17:58:37+00:00
Link: https://simonwillison.net/2026/Feb/19/gemini-31-pro/#atom-everything
Gemini 3.1 Pro The first in the Gemini 3.1 series, priced the same as Gemini 3 Pro ($2/million input, $12/million output under 200,000 tokens, $4/$18 for 200,000 to 1,000,000). That’s less than half the price of Claude Opus 4.6 with very similar benchmark scores to that model. They boast about its improved SVG animation performance compared to Gemini 3 Pro in the announcement! I tried "Generate an SVG of a pelican riding a bicycle" in Google AI Studio and it thought for 323.9 seconds (thinking trace here) before producing this one: It’s good to see the legs clearly depicted on both sides of the frame (should satisfy Elon), the fish in the basket is a nice touch and I appreciated this comment in the SVG code : I’ve added the two new model IDs gemini-3.1-pro-preview and gemini-3.1-pro-preview-customtools to my llm-gemini plugin for LLM . That "custom tools" one is described here - apparently it may provide better tool performance than the default model in some situations. The model appears to be incredibly slow right now - it took 104s to respond to a simple "hi" and a few of my other tests met "Error: This model is currently experiencing high demand. Spikes in demand are usually...

5. Experimenting with sponsorship for my blog and newsletter
Source: Simon Willison’s Weblog | Published: 2026-02-19T05:44:29+00:00
Link: https://simonwillison.net/2026/Feb/19/sponsorship/#atom-everything
I’ve long been resistant to the idea of accepting sponsorship for my blog. I value my credibility as an independent voice, and I don’t want to risk compromising that reputation. Then I learned about Troy Hunt’s approach to sponsorship , which he first wrote about in 2016 . Troy runs with a simple text row in the page banner - no JavaScript, no cookies, unobtrusive while providing value to the sponsor. I can live with that! Accepting sponsorship in this way helps me maintain my independence while offsetting the opportunity cost of not taking a full-time job. To start with I’m selling sponsorship by the week. Sponsors get that unobtrusive banner across my blog and also their sponsored message at the top of my newsletter . I will not write content in exchange for sponsorship . I hope the sponsors I work with understand that my credibility as an independent voice is a key reason I have an audience, and compromising that trust would be bad for everyone. Freeman & Forrest helped me set up and sell my first slots. Thanks also to Theo Browne for helping me think through my approach. Tags: newsletter , blogging , troy-hunt

6. SWE-bench February 2026 leaderboard update
Source: Simon Willison’s Weblog | Published: 2026-02-19T04:48:47+00:00
Link: https://simonwillison.net/2026/Feb/19/swe-bench/#atom-everything

SWE-bench February 2026 leaderboard update SWE-bench is one of the benchmarks that the labs love to list in their model releases. The official leaderboard is infrequently updated but they just did a full run of it against the current generation of models, which is notable because it’s always good to see benchmark results like this that weren’t self-reported by the labs. The fresh results are for their "Bash Only" benchmark, which runs their mini-swe-bench agent (~9,000 lines of Python, here are the prompts they use) against the SWE-bench dataset of coding problems - 2,294 real-world examples pulled from 12 open source repos: django/django (850), sympy/sympy (386), scikit-learn/scikit-learn (229), sphinx-doc/sphinx (187), matplotlib/matplotlib (184), pytest-dev/pytest (119), pydata/xarray (110), astropy/astropy (95), pylint-dev/pylint (57), psf/requests (44), mwaskom/seaborn (22), pallets/flask (11). Correction : The Bash only benchmark runs against SWE-bench Verified, not original SWE-bench. Verified is a manually curated subset of 500 samples described here , funded by OpenAI. Here’s SWE-bench Verified on Hugging Face - since it’s just 2.1MB of Parquet it’s easy to browse using D...

7. IMAX and Apple Collaborate to Screen F1 Races Live in Theaters
Source: Daring Fireball | Published: 2026-02-19T23:29:06Z
Link: https://www.motorsport.com/f1/news/f1-to-screen-live-in-imax-theatres-in-2026-as-apple-tv-unveils-new-us-viewing-experience/10798974/
Lydia Mee, reporting for Motorsport: IMAX has announced that a select number of races will be shown live in IMAX locations across the United States in 2026. The new fan viewing experience is part of a collaboration with Apple TV, which has taken over the broadcasting rights for the championship in the US on a multi-year deal from 2026. ?F1 is a rapidly growing force in sports and culture in the US, and by bringing F1 on Apple TV live to IMAX theatres nationwide, we?re delivering the energy and excitement to even more screens in a truly immersive way,? said Oliver Schusser, Apple?s vice president of music, sports, and Beats. You know what would add even more screens in an immersive way? If Vision Pro users had access to the same live screenings on virtual IMAX screens. ?

8. One More Spitball Idea for Apple?s March 4 Media Event ?Experience?: Immersive F1 on Vision Pro?
Source: Daring Fireball | Published: 2026-02-19T04:41:38Z
Link: https://www.formula1.com/en/latest/article/official-grand-prix-start-times-for-2026-f1-season-confirmed.2UgPfArqH76tzlOYh21jSG
A reader pointed out that the 2026 Formula 1 season starts in Australia on March 8. You will recall from October that Apple TV is now the exclusive broadcast partner for F1 in the U.S. Apple is already dabbling with live immersive sports broadcasting for VisionOS with a limited slate of Lakers games this season . If they have something planned for streaming F1 races live on Vision Pro, with some level of immersion, March 4 would be a pretty good date to demo that experience to the media. It doesn?t even have to be live race coverage. Technically that?s probably impossible for this season. It would just be a sign of confidence and interest in the platform long-term merely to see some sort of immersive component to F1 on Apple TV, even if it?s not live. Like ?ride the track? to experience the turns and elevation changes. Could just be a total coincidence that the Formula 1 season is starting the weekend after this event. But it seems worth noting. ?

9. ?The Secret Fear of the Morally Depraved?
Source: Daring Fireball | Published: 2026-01-27T23:36:54Z
Link: https://www.theatlantic.com/ideas/2026/01/the-neighbors-defending-minnesota-from-ice/685769/?gift=Je3D9AQS-C17lUTOnl2W8L893jn-xkg4gA0ahaD_Ltw
Adam Serwer, reporting from the streets of Minneapolis for The Atlantic, ?Minnesota Proved MAGA Wrong? (gift link): The secret fear of the morally depraved is that virtue is actually common, and that they?re the ones who are alone. In Minnesota, all of the ideological cornerstones of MAGA have been proved false at once. Minnesotans, not the armed thugs of ICE and the Border Patrol, are brave. Minnesotans have shown that their community is socially cohesive ? because

of its diversity and not in spite of it. Minnesotans have found and loved one another in a world atomized by social media, where empty men have tried to fill their lonely soul with lies about their own inherent superiority. Minnesotans have preserved everything worthwhile about ?Western civilization,? while armed brutes try to tear it down by force. ?

10. What have we learned about building agentic AI tools?
Source: seangoedecke.com RSS feed | Published: Sun, 19 Oct 2025 00:00:00 GMT
Link: https://seangoedecke.com/ideas-in-agentic-ai-tooling/
In the middle of 2025, agentic coding finally became a thing: first with the release of Claude Sonnet 4, the first ?smart enough to be useful? agentic model, and then with OpenAI?s GPT-5-Codex, which is for my money the best-in-class agentic model. ?Agent mode? is now the main way to interact with your preferred AI coding tool (whether that?s Claude Code, Codex, Copilot, or Cursor) 1 . Obviously, much of this improvement is the result of better models. Sonnet 4 and Codex simply get lost less often and make fewer mistakes than their predecessors. But we?ve also seen a ton of improvements to the agentic harness : the code that wraps the LLM in a loop with tools. For an interesting time capsule, you can read my 2023 post Building LLM-driven agents , where I wrote about my attempts to build an agentic coding system on top of GPT-3 (!). This was before tool calls existed - I had to prompt the model to include structured tool call content in its output and then parse that out. In hindsight, I was right about a few important things: LLM agentic coding could actually work If you build a good agentic coding tool, it?s only going to get more useful as better models come out You should tune...

11. We are in the "gentleman scientist" era of AI research
Source: seangoedecke.com RSS feed | Published: Sat, 18 Oct 2025 00:00:00 GMT
Link: https://seangoedecke.com/ai-and-informal-science/
Many scientific discoveries used to be made by amateurs. William Herschel , who discovered Uranus, was a composer and an organist. Antoine Lavoisier , who laid the foundation for modern chemistry, was a politician. In one sense, this is a truism. The job of ?professional scientist? only really appeared in the 19th century, so all discoveries before then logically had to have come from amateurs, since only amateur scientists existed. But it also reflects that any field of knowledge gets more complicated over time . In the early days of a scientific field, discoveries are simple: ?air has weight?, ?white light can be dispersed through a prism into different colors?, ?the mass of a burnt object is identical to its original mass?, and so on. The way you come up with those discoveries is also simple: observing mercury in a tall glass tube, holding a prism up to a light source, weighing a sealed jar before and after incinerating it, and so on. The 2025 Nobel prize in physics was just awarded ?for the discovery of macroscopic quantum mechanical tunnelling and energy quantisation in an electric circuit?. The press release gallantly tries to make this discovery understandable to the layman...

12. How I provide technical clarity to non-technical leaders
Source: seangoedecke.com RSS feed | Published: Sun, 12 Oct 2025 00:00:00 GMT
Link: https://seangoedecke.com/clarity/
My mission as a staff engineer is to provide technical clarity to the organization. Of course, I do other stuff too. I run projects, I ship code, I review PRs, and so on. But the most important thing I do - what I?m for - is to provide technical clarity. What is technical clarity? In an organization, technical clarity is when non-technical decision makers have a good-enough practical understanding of what changes they can make to their software systems. The people in charge of your software organization 1 have to make a lot of decisions about software. Even if they?re not setting the overall strategy, they?re still probably deciding which kinds of users get which features, which updates are most important to roll out, whether projects should be delayed or rushed, and so on. These people may have been technical once. They may even have fine technical minds now. But they?re still ?non-technical? in the sense I mean, because they simply don?t have the time or the context to build an accurate mental model of the system. Instead, they rely on a vague mental model, supplemented by advice from engineers they trust. To the extent that their vague mental model is accurate and the advice

th...

13. My Approach to Building Large Technical Projects
Source: Mitchell Hashimoto | Published: Thu, 01 Jun 2023 00:00:00 GMT
Link: https://mitchellh.com/writing/building-large-technical-projects
(No summary)

14. Integrating Zig and SwiftUI
Source: Mitchell Hashimoto | Published: Sat, 27 May 2023 00:00:00 GMT
Link: https://mitchellh.com/writing/zig-and-swiftui
(No summary)

15. Prompt Engineering is for Transactional Prompting
Source: Mitchell Hashimoto | Published: Mon, 24 Apr 2023 00:00:00 GMT
Link: https://mitchellh.com/writing/prompt-engineering-transactional-prompting
(No summary)

16. Exploring the signals the dialog manager uses for dismissing a dialog
Source: The Old New Thing | Published: Thu, 19 Feb 2026 15:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20260219-00/?p=112072
Summarizing the flow. The post Exploring the signals the dialog manager uses for dismissing a dialog appeared first on The Old New Thing .

17. My advice on (internet) writing, for what it?s worth
Source: DYNOMIGHT | Published: 2025-05-29T00:00:00+00:00
Link: https://dynomight.net/writing-advice/
A lot of writing advice seems to amount to: I start by having verbal intelligence that?s six standard deviations above the population mean. I find that this is really helpful! Also, here are some tips on spelling and how I cope with the never-ending adoration. I think this leaves space for advice from people with less godlike levels of natural talent e.g. your friend dynomight. The advice Here it is: Make something you would actually like. Actually, let me bold the important words: Make something you would actually like. Why you ? Why make something you would like? To be clear, I?m not suggesting you write ?for yourself?. I assume that your terminal goal is to make something other people like. But try this experiment: Go write a few thousand words and give them to someone who loves you. Now, go through paragraph-by-paragraph and try to predict what was going through their head while reading. It?s impossible. I tell you, it cannot be done! Personally, I think this is because nobody really understands anyone else. (I recently discovered that my mother secretly hates tomatoes.) If you try to make something that other people would like, rather than yourself, you?ll likely just end up...

18. DumPy: NumPy except it?s OK if you?re dum
Source: DYNOMIGHT | Published: 2025-05-22T00:00:00+00:00
Link: https://dynomight.net/dumpy/
What I want from an array language is: Don?t make me think. Run fast on GPUs. Really, do not make me think. Do not . I say NumPy misses on three of these. So I?d like to propose a ?fix? that?I claim?eliminates 90% of unnecessary thinking, with no loss of power. It would also fix all the things based on NumPy, for example every machine learning library. I know that sounds grandiose. Quite possibly you?re thinking that good-old dynomight has finally lost it. So I warn you now: My solution is utterly non-clever. If anything is clever here, it?s my single-minded rejection of cleverness. To motivate the fix, let me give my story for how NumPy went wrong. It started as a nice little library for array operations and linear algebra. When everything has two or fewer dimensions, it?s great. But at some point, someone showed up with some higher-dimensional arrays. If loops were fast in Python, NumPy would have said, ?Hello person with ?3 dimensions, please call my ?2 dimensional functions in a loop so I can stay nice and simple, xox, NumPy.? But since loops are slow, NumPy instead took all the complexity that

would usually be addressed with loops and pushed it down into individual functions....

19. The Heat Mirage: My least-favorite internet maneuver
Source: DYNOMIGHT | Published: 2025-05-19T00:00:00+00:00
Link: https://dynomight.net/heat/
Alice and Bob are driving through the desert. Alice : Looks dry. Bob : That?s wrong, what we see ahead is caused by the sun heating up the road. While the speed of light is constant in vacuum, when light moves through matter, the atoms emit new light that destructively interferes with the old light, in effect causing a ?delay?. This happens more with more atoms, meaning light travels slower through denser media. Alice : OK, but? Bob : So when the sun heats up the road, this creates a layer of thin warm air with denser cooler air higher up. As light passes through these layers, it refracts upwards towards the denser cooler layer. If conditions are right, this can bend the light back up towards your eyes. Does that make sense? Alice : Yeah, but? Bob : Now you might ask, ?Why does this look like water ?? The situations seem similar at first, with two fluids of different densities. But think about it: With an ocean, the dense water is below the thin air. While, in front of you, the dense cool air is above the thin warm air. The situations are actually reversed! Alice : Please just? Bob : With an ocean, there?s a sharp increase in density where the air meets the water. The Fresnel equa...

20. The WET Codebase
Source: overreacted ? A blog by Dan Abramov | Published: Mon, 13 Jul 2020 00:00:00 GMT
Link: https://overreacted.io/the-wet-codebase/
Come waste your time with me.

21. Goodbye, Clean Code
Source: overreacted ? A blog by Dan Abramov | Published: Sat, 11 Jan 2020 00:00:00 GMT
Link: https://overreacted.io/goodbye-clean-code/
Let clean code guide you. Then let it go.

22. My Decade in Review
Source: overreacted ? A blog by Dan Abramov | Published: Wed, 01 Jan 2020 00:00:00 GMT
Link: https://overreacted.io/my-decade-in-review/
A personal reflection.

23. Tighter bounds in the prime number theorem
Source: John D. Cook | Published: Fri, 16 Jan 2026 15:50:28 +0000
Link: https://www.johndcook.com/blog/2026/01/16/prime-number-theorem-bounds/
The most elementary form of the prime number theorem says that ?(x), the number of prime numbers less than x, is asymptotically equal to x / log(x). That?s true, but a more accurate result says ?(x) is asymptotically equal to li(x) where Five years ago I wrote about a result that was new at the [?] The post Tighter bounds in the prime number theorem first appeared on John D. Cook .

24. Efficiently computing multiple modular inverses at once
Source: John D. Cook | Published: Wed, 14 Jan 2026 15:06:03 +0000
Link: https://www.johndcook.com/blog/2026/01/14/montgomerys-trick/
Suppose you have a large prime number M and you need to find the inverse of several numbers mod M. Montgomery?s trick is a way to combine the computation of the inverses to take less time than computing the inverses individually. Peter Montgomery (1947?2020) came up with this trick in 1985. We will illustrate Montgomery?s trick by [?] The post Efficiently computing multiple modular inverses at once first appeared on John D. Cook .

25. The middle binomial coefficient
Source: John D. Cook | Published: Mon, 12 Jan 2026 12:41:29 +0000
Link: https://www.johndcook.com/blog/2026/01/12/the-middle-binomial-coefficient/

The previous post contained an interesting observation: Is it true more generally that for large n? Sorta, but the approximation gets better if we add a correction factor. If we square both sides of the approximation and move the factorials to one side, the question becomes whether Now the task becomes to estimate the middle coefficient [?] The post The middle binomial coefficient first appeared on John D. Cook .

26. AI is a NAND Maximiser
Source: Terence Eden?s Blog | Published: Thu, 19 Feb 2026 12:34:33 +0000
Link: https://shkspr.mobi/blog/2026/02/ai-is-a-nand-maximiser/
PC Gamer is reporting that the current demand by AI companies for computer chips is having a disastrous effect on the rest of the industry. In an interview, the CEO of Phison said: If NVIDIA Vera Rubin ships tens of millions of units, each requiring 20+TB SSDs, it will consume approximately 20% of last year’s global NAND production capacity ?HaYaO NAND is a type of microchip. Rather than b?

27. Go Modules for Package Management Tooling
Source: Andrew Nesbitt | Published: 2026-02-19T00:00:00+00:00
Link: https://nesbitt.io/2026/02/19/go-modules-for-package-management-tooling.html
I?ve been working on a reusable layer for building ecosystem-agnostic package and supply chain tools in Go: fourteen modules under git-pkgs covering manifest parsing, registry clients, license normalization, platform translation, vulnerability feeds, and more. These are rebuilds of libraries I?ve written and used in Ruby for years, some going back to Libraries.io and more recently for Ecosyste.ms , which I wrote about previously . I built the Go versions for git-pkgs , a tool for exploring the dependency history of your repositories that compiles to a single binary with no runtime dependencies, which matters for a git subcommand that needs to just work on any machine. When I went looking for Go equivalents of my Ruby libraries, most were either abandoned, incomplete, or only covered a single ecosystem, so I rebuilt them. Identification purl Package URL (now ECMA-427) is the standard format for identifying packages across ecosystems. This handles parsing, generation, and type-specific configuration for around 40 ecosystems, including registry URL generation and the reverse: parsing a registry URL back into a PURL. p , _ := purl . Parse ("pkg:npm/%40babel/core@7.24.0") p . FullNa...

28. Pluralistic: Six Years of Pluralistic (19 Feb 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Thu, 19 Feb 2026 14:08:05 +0000
Link: https://pluralistic.net/2026/02/19/now-we-are-six/
Today’s links Six years of Pluralistic: Time flies when you’re writing the web. Hey look at this: Delights to delectate. Object permanence: MBA phrenology; Sony’s DRM CEO is out; Midwestern Tahrir; Reverse Centaurs and AI. Upcoming appearances: Where to find me. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. Six years of Pluralistic (permalink) Six years ago today, after 19 years with Boing Boing, during which time I wrote tens of thousands of blog posts, I started a new, solo blog, with the semi-ironic name "Pluralistic." I didn’t know what Pluralistic was going to be, but I wasn’t writing Boing Boing anymore, and I knew I wanted to keep writing the web in some fashion. Six years and more than 1,500 posts later, I am so satisfied with how Pluralistic is going. I spent a couple of decades processing everything that seemed interesting or significant through a blog, which created a massive database (and mnemonically available collection of partially developed thoughts) that I’m now reprocessing as a series of essays that make sense of today in light of everythi...

29. Programming vs. Coding vs. Software Engineering
Source: Rakhim’s blog | Published: Wed, 27 Nov 2019 00:00:00 GMT
Link: https://rakhim.exotext.com/programming-vs-coding-vs-software-engineering
Programming, coding, and software engineering are often used interchangeably. Here is how I usually think of them. Programming Programming is solving explicit problems in a verifiable

manner. It is similar to mathematical proofs. "Find an element in a sorted sequence" is a problem. A binary search algorithm is a solution. You can prove it works, and even analyze whether it’s optimal in terms of time or memory. A better word for programming might be computing ?a branch of mathematics. At this level, code isn’t the focus; algorithms are. Only essential complexity matters. Accidental complexity doesn’t exist here. The goal of programming is to create an algorithm?or prove that one cannot exist. Coding Coding is expressing a programming solution in a formal language. It?s like writing or speaking. Binary search can be coded imperatively in C or JavaScript, or declaratively in Haskell. The implementation details differ, but the algorithm remains the same. Coding is both easier and harder than programming. Easier, because it assumes the problem is already solved. Harder, because translating an abstract solution into a real system often introduces vague, hard-to-verify constraints (e.g.,...

30. Products aren’t for people yet
Source: Rakhim’s blog | Published: Sat, 04 Feb 2017 00:00:00 GMT
Link: https://rakhim.exotext.com/products-arent-for-people-yet
Remember how your parents would try to use Windows 95 or something like Norton Commander. They?d copy an app shortcut to a floppy disk and be amazed how much stuff they were able to put inside. All the games, and lots of space left! And you?d think they don?t understand anything at all, they are just clicking pretty much randomly, hoping this magic machine would at some point understand them and do the right thing. That was the time when programmers were building products for programmers . It wasn?t awesome, but it was sincere. Nobody was pretending that software was built for regular people, and the consensus was: in order to use a computer you have to learn something. Or know someone who can help. Today we?re living in times when programmers are building products for people believe that they?re building products for people . This isn?t awesome and it?s not sincere. Today all of us regularly feel like our parents with Windows 95. You have to learn or, more commonly, just remember how to use a certain website or an app. And I?m not talking about small pieces of software, I?m mostly talking about huge, global products: Google?s user interfaces are chaos and madness, Facebook is mad...

31. Producing a video lesson
Source: Rakhim’s blog | Published: Wed, 11 Jan 2017 00:00:00 GMT
Link: https://rakhim.exotext.com/producing-a-video-lesson
A few years ago, I’ve made a bunch of videos on YouTube about foundational topics in programming. There are lots of hand-drawn cartoons and some rudimentary animation. People keep asking how I made those, so here we go. Inspiration Structure and Interpretation of Computer Programs (SICP). If you?re only going to read one book on programming, it should be this one. Interactive articles and talks by Bret Victor . An amazing person ? at the very least, watch his talk Inventing on Principle . Mindstorms by Seymour Papert. About math anxiety, poor teaching methods, and innovative approaches to education. Every teacher ? especially math or CS teachers ? should read it. A bunch of YouTube channels with great visualizations: Ted-ED , Vox , Minute Physics , Numberphile , Computerphile . Scriptwriting First, I define the topic and write the script. I have an initial course outline, but the order sometimes shifts to maintain logical flow. I start with all major topics written on sticky notes, laid out on the desk: past topics, current candidates, and future ones. This helps visualize what concepts were already introduced, which ones I need to reinforce, and which to hint at. The upcoming les...

32. Is the Future ?AWS for Everything??
Source: Construction Physics | Published: Thu, 19 Feb 2026 13:01:06 GMT
Link: https://www.construction-physics.com/p/is-the-future-aws-for-everything
A theme running through my book is the idea that efficiency improvements, and the various methods for making products cheaper over time, have historically been dependent on some degree of repetition, on running your production process over and over again.

33. What does "Undecidable" mean, anyway

Source: Computer Things | Published: Wed, 28 May 2025 19:34:02 +0000
Link: https://buttondown.com/hillelwayne/archive/what-does-undecidable-mean-anyway/
Systems Distributed I'll be speaking at Systems Distributed next month! The talk is brand new and will aim to showcase some of the formal methods mental models that would be useful in mainstream software development. It has added some extra stress on my schedule, though, so expect the next two monthly releases of Logic for Programmers to be mostly minor changes. What does "Undecidable" mean, anyway Last week I read Against Curry-Howard Mysticism , which is a solid article I recommend reading. But this newsletter is actually about one comment : I like to see posts like this because I often feel like I can?t tell the difference between BS and a point I?m missing. Can we get one for questions like ?Isn?t XYZ (Undecidable|NP-Complete|PSPACE-Complete)?? I've already written one of these for NP-complete , so let's do one for "undecidable". Step one is to pull a technical definition from the book Automata and Computability : A property P of strings is said to be decidable if ... there is a total Turing machine that accepts input strings that have property P and rejects those that do not. (pg 220) Step two is to translate the technical computer science definition into more conventional pr...

34. Finding hard 24 puzzles with planner programming
Source: Computer Things | Published: Tue, 20 May 2025 18:21:01 +0000
Link: https://buttondown.com/hillelwayne/archive/finding-hard-24-puzzles-with-planner-programming/
Planner programming is a programming technique where you solve problems by providing a goal and actions, and letting the planner find actions that reach the goal. In a previous edition of Logic for Programmers , I demonstrated how this worked by solving the 24 puzzle with planning. For reasons discussed here I replaced that example with something more practical (orchestrating deployments), but left the code online for posterity. Recently I saw a family member try and fail to vibe code a tool that would find all valid 24 puzzles, and realized I could adapt the puzzle solver to also be a puzzle generator. First I'll explain the puzzle rules, then the original solver, then the generator. 1 For a much longer intro to planning, see here . The rules of 24 You're given four numbers and have to find some elementary equation (+-*/ +groupings) that uses all four numbers and results in 24. Each number must be used exactly once, but do not need to be used in the starting puzzle order. Some examples: [6, 6, 6, 6] -> 6+6+6+6=24 [1, 1, 6, 6] -> (6+6)*(1+1)=24 [4, 4, 4, 5] -> 4*(5+4/4)=24 Some setups are impossible, like [1, 1, 1, 1] . Others are possible only with non-elementary operations, lik...

35. Modeling Awkward Social Situations with TLA+
Source: Computer Things | Published: Wed, 14 May 2025 16:02:21 +0000
Link: https://buttondown.com/hillelwayne/archive/modeling-awkward-social-situations-with-tla/
You're walking down the street and need to pass someone going the opposite way. You take a step left, but they're thinking the same thing and take a step to their right , aka your left. You're still blocking each other. Then you take a step to the right, and they take a step to their left, and you're back to where you started. I've heard this called "walkwarding" Let's model this in TLA+ . TLA+ is a formal methods tool for finding bugs in complex software designs, most often involving concurrency. Two people trying to get past each other just also happens to be a concurrent system. A gentler introduction to TLA+'s capabilities is here , an in-depth guide teaching the language is here . The spec ---- MODULE walkward ---- EXTENDS Integers VARIABLES pos vars == > Double equals defines a new operator, single equals is an equality check. > is a sequence, aka array. you == "you" me == "me" People == {you, me} MaxPlace == 4 left == 0 right == 1 I've gotten into the habit of assigning string "symbols" to operators so that the compiler complains if I misspelled something. left and right are numbers so we can shift position with right - pos . direction == [you |-> 1, me |-> -1] goal == [you...

36. AI is the Best Thing to Happen to Art
Source: the singularity is nearer | Published: 2026-02-19T00:00:00+08:00
Link: https://geohot.github.io//blog/jekyll/update/2026/02/19/ai-art.html

I watched this video about how AI has already ruined music. Her mom sent her a song and she told her mom it was AI. She played the song and it sounded like slop. It had inspired lyrics like: From quiet roots, a garden grows She?s got that light, and now it shows Yes, she rises, and she glows Oh, she rises, now she knows Pure slop. Compare it to: I?m in the cut acting crazy I?m in the whip doing eighty Only God can judge me And only she can save me Note that ?cut? and ?whip? are not exactly words, but products of a culture . Ulysses is particularly hard to read because you don?t know 1910?s Irish pop culture. I can?t believe Marvel movies were popular. The first Iron Man was good, but by the time we got to Spider-Man: No Way Home it was practically a clip show with triple inside references and cringy fourth wall breaking humor. How it got a 8.1 on IMDB is beyond me, and just reduced my trust for IMDB. Marvel movies are also the easiest things to make with AI. Little story and long term coherence, no ?progression of the genre?, tons of eye candy special effects. It?s shocking with how large the budgets for them were that they couldn?t pay a story guy a little. I felt similarly about...

37. Modeling a Wealth Tax
Source: Paul Graham: Essays | Published: -
Link: http://www.paulgraham.com/wtax.html
(No summary)

38. The Four Quadrants of Conformism
Source: Paul Graham: Essays | Published: -
Link: http://www.paulgraham.com/conformism.html
(No summary)

39. Orthodox Privilege
Source: Paul Graham: Essays | Published: -
Link: http://www.paulgraham.com/orth.html
(No summary)

40. Redis 6.0.0 GA is out!
Source: antirez.com | Published: Thu, 30 Apr 2020 15:33:35 +0200
Link: http://antirez.com/news/132
Finally Redis 6.0.0 stable is out. This time it was a relatively short cycle between the release of the first release candidate and the final release of a stable version. It took about four months, that is not a small amount of time, but is not a lot compared to our past records :) So the big news are the ones announced before, but with some notable changes. The old stuff are: SSL, ACLs, RESP3, Client side caching, Threaded I/O, Diskless replication on replicas, Cluster support in Redis-benchmark and improved redis-cli cluster support, Disque in beta as a module of Redis, and the Redis Cluster Proxy (now at https://github.com/RedisLabs/redis-cluster-proxy). So what changed between RC1 and today, other than stability? 1. Client side caching was redesigned in certain aspects, especially the caching slot approach was discarded in favor of just using key names. After analyzing the alternatives, with the help of other Redis core team members, in the end this approach looks better. Other than that, finally the feature was completed with the things I had in the backlog for the feature, especially the ?broadcasting mode?, that I believe will be one of the most popular usage modes of the f...

41. Redis 6 RC1 is out today
Source: antirez.com | Published: Thu, 19 Dec 2019 17:27:00 +0100
Link: http://antirez.com/news/131
So it happened again, a new Redis version reached the release candidate status, and in a few months it will hit the shelves of most supermarkets. I guess this is the most ?enterprise? Redis version to date, and it?s funny since I took quite some time in order to understand what ?enterprise? ever meant. I think it?s word I genuinely dislike, yet it has some meaning. Redis is now everywhere, and it is still considerably able to ?scale down?: you can still download

it, compile it in 30 seconds, and run it without any configuration to start hacking. But being everywhere also means being in environments where things like encryption and ACLs are a must, so Redis, inevitably, and more than thanks to me, I would say, in spite of my extreme drive for simplicity, adapted. But what?s interesting is that, even additions may be done in very opinionated ways. Redis ACLs hardly resemble something you saw in other systems, and SSL support was written in a few iterations in order to finally pick the idea that was the most sounding, from the point of view of letting the core as clean as possible. I?m quite happy with the result. Redis 6 does not bring just ACLs and SSL, it is the largest release of...

42. Client side caching in Redis 6
Source: antirez.com | Published: Thu, 04 Jul 2019 19:10:34 +0200
Link: http://antirez.com/news/130
[Note: this post no longer describes the client side implementation in the final implementation of Redis 6, that changed significantly, see https://redis.io/topics/client-side-caching] The New York Redis day was over, I get up at the hotel at 5:30, still pretty in sync with the Italian time zone and immediately went walking on the streets of Manhattan, completely in love with the landscape and the wonderful feeling of being just a number among millions of other numbers. Yet I was thinking at the Redis 6 release with the feeling that, what was probably the most important feature at all, the new version of the Redis protocol (RESP3), was going to have a very slow adoption curve, and for good reasons: wise people avoid switching tools without very good reasons. After all why I wanted to improve the protocol so badly? For two reasons mainly, to provide clients with more semantical replies, and in order to open to new features that were hard to implement with the old protocol; one feature in particular was the most important to me: client side caching. Rewind back to about one year ago. I arrived at Redis Conf 2018, in San Francisco, with the firm idea that client side caching was the...

43. the rust project has a burnout problem
Source: the website of jyn | Published: 2024-01-16T00:00:00+00:00
Link: https://jyn.dev/the-rust-project-has-a-burnout-problem/
the number of people who have left the rust project due to burnout is shockingly high. the number of people in the project who are close to burnout is also shockingly high. this post is about myself, but it’s not just about myself. i’m not going to name names because either you know what i’m talking about, in which case you know at least five people matching this description, or you don’t, in which case sorry but you’re not the target audience. consider, though, that the project has been around for 15 years, and compare that to the average time a maintainer has been active ... what does this look like (i apologize in advance if this story does not match your experience; hopefully the suggestions on what to do about burnout will still be helpful to you.) the pattern usually goes something like this: you want to work on rust. you go to look at the issue tracker. you find something you care about, since the easy/mentored issues are taken. it’s hard to find a mentor because all the experienced people are overworked and burned out, so you end up doing a lot of the work independently. guess what you’ve already learned at this point: work in this project doesn’t happen unless you persona...

44. wonder
Source: the website of jyn | Published: 2023-12-20T00:00:00+00:00
Link: https://jyn.dev/wonder/
a sense of wonder is a wonderful thing

45. How to maintain an Open Source project
Source: the website of jyn | Published: 2023-12-04T00:00:00+00:00
Link: https://jyn.dev/how-to-maintain-an-open-source-project/
How to maintain a project without burning yourself out

46. mDNS Primer
Source: Fabien Sanglard | Published: 11 Aug 2023 00:00:00 +0000

Link: https://fabiensanglard.net/mdns/index.html
(No summary)

47. Commander Keen: Adaptive Tile Scrolling
Source: Fabien Sanglard | Published: 27 Jul 2023 00:00:00 +0000
Link: https://fabiensanglard.net/ega/index.html
(No summary)

48. 10NES
Source: Fabien Sanglard | Published: 18 Jul 2023 00:00:00 +0000
Link: https://fabiensanglard.net/10nes/index.html
(No summary)

49. HTML5 Canvas Part II: Pixel Manipulation
Source: Beej's Bit Bucket | Published: Tue, 09 Feb 2010 00:00:00 +0000
Link: http://beej.us/blog/data/html5s-canvas-2-pixel/
(No summary)

50. HTML5 Canvas Introduction
Source: Beej's Bit Bucket | Published: Sun, 07 Feb 2010 00:00:00 +0000
Link: http://beej.us/blog/data/html5-canvas/
(No summary)

51. Function Pointers and Cooperative Multitasking
Source: Beej's Bit Bucket | Published: Fri, 29 Jan 2010 00:00:00 +0000
Link: http://beej.us/blog/data/function-pointers-cooperative-multitasking/
Just like Grandma used to multitask.

52. CommBank's AI boyfriend
Source: Dr Paris Buttfield-Addison | Published: Tue, 12 Aug 2025 00:00:00 +0000
Link: https://hey.paris/posts/cba/
CBA has been naughty, and too reliant on their AI boyfriend . Here?s what happened: A CBA customer contacted bank requesting contact details for Secretlab (a company that is not us, and makes chairs) CBA staff member queried ChatGPT (possibly via their own personal, unauthenticated access to ChatGPT) to obtain phone number for Secretlab (chairs) CBA staff disclosed the retrieved phone number to the requesting customer The retrieved phone number is a number belonging to one of the directors of Secret Lab (a company that does not make chairs, and is us), a customer of CBA, and is used for our CBA account and our Director?s CBA account Therefore, CBA disclosed customer personal information to another, unrelated customer, and trusted a third-party LLM (ChatGPT), accessed seemingly unauthenticated on the consumer ChatGPT platform, as a source for data to provide to another customer During investigation, CBA staff replicated the same ChatGPT query process, seemingly on a personal phone, again unauthenticated, and on the consumer ChatGPT platform: Signs indicate this might be routine practice amongst CBA staff Oh, and they gave out the phone number for us, Secret Lab (not chairs), to som...

53. Space News, August 2025
Source: Dr Paris Buttfield-Addison | Published: Sun, 03 Aug 2025 00:00:00 +0000
Link: https://hey.paris/posts/spacenews-august-2025/
These are my notes and expanded thoughts from this month?s Space News segment on ABC Radio Hobart and ABC Northern Tasmania. Every month I join Lucie Cutting on Sunday mornings to chat about what?s happening above and beyond. Here?s what we discussed, plus some of my notes, from this edition of the programme : The most valuable 14 seconds in Australian aerospace history happened last month, while the most expensive corporate rivalry ever fought in orbit just turned 40.

54. We’re all adults here
Source: Dr Paris Buttfield-Addison | Published: Fri, 25 Jul 2025 00:00:00 +0000
Link: https://hey.paris/posts/werealladultshere/
The recent delisting of over 17,000 games from the independent platform Itch.io wasn?t a content moderation decision made by the platform itself. It was an act of financial coercion. Instigated by an Australian activist group and carried out by payment processors like Mastercard, it sets a dangerous precedent for all creative fields online. The stated goal of the campaign was to combat media that glorifies sexual violence. An aim that sounds laudable, making it difficult to argue against in public. But like calls to age-gate social media, such goals can act as a Trojan horse. A socially acceptable pretext for gaining more information and control over people. And the outcome here reveals a profound and damaging overreach.

55. error-driven-development
Source: Chad Nauseam Home | Published: -
Link: https://chadnauseam.com/coding/tips/error-driven-development
(No summary)

56. the-game-design-of-math-academy
Source: Chad Nauseam Home | Published: -
Link: https://chadnauseam.com/coding/gamedev/the-game-design-of-math-academy
(No summary)

57. calculator-app
Source: Chad Nauseam Home | Published: -
Link: https://chadnauseam.com/coding/random/calculator-app
(No summary)

58. What happens when coding agents stop feeling like dialup?
Source: Martin Alderson | Published: Fri, 19 Sep 2025 00:00:00 GMT
Link: https://martinalderson.com/posts/what-happens-when-coding-agents-stop-feeling-like-dialup/?utm_source=rss
From magical to frustrating in months. Why AI coding agents feel like dial-up internet and what ultra-fast inference could unlock for developer productivity.

59. Solving Claude Code’s API Blindness with Static Analysis Tools
Source: Martin Alderson | Published: Mon, 01 Sep 2025 00:00:00 GMT
Link: https://martinalderson.com/posts/claude-code-static-analysis/?utm_source=rss
How to give AI coding assistants complete visibility into APIs and third-party libraries using static analysis instead of basic text search.

60. Are OpenAI and Anthropic Really Losing Money on Inference?
Source: Martin Alderson | Published: Wed, 27 Aug 2025 00:00:00 GMT
Link: https://martinalderson.com/posts/are-openai-and-anthropic-really-losing-money-on-inference/?utm_source=rss
Deconstructing the real costs of running AI inference at scale. My napkin math suggests the economics might be far more profitable than commonly claimed.

61. Keygen is now Fair Source
Source: Keygen Blog | Published: Tue, 06 Aug 2024 05:00:00 GMT
Link: https://keygen.sh/blog/keygen-is-now-fair-source/
Keygen relicenses from ELv2 to the Fair Core License.

62. Licensing is packaging
Source: Keygen Blog | Published: Mon, 02 Sep 2024 05:00:00 GMT

Link: https://keygen.sh/blog/licensing-is-packaging/
On the first- and second-order effects of software licensing.

63. Keygen joins the Open Source Pledge
Source: Keygen Blog | Published: Mon, 09 Sep 2024 05:00:00 GMT
Link: https://keygen.sh/blog/keygen-joins-osspledge/
Keygen is joining the Open Source Pledge, committing to donate at least $2,000 per-developer annually to support Open Source.

64. Porting PuTTY to Windows on Arm
Source: Simon Tatham's quasi-blog | Published: 2019-03-05T00:00:00+00:00
Link: https://www.chiark.greenend.org.uk/~sgtatham/quasiblog/putty-woa/
What I had to do to make PuTTY compile for Windows on Arm as well as on x86, plus a deep dive into the NEON-accelerated cryptography.

65. Hoe herken je een (toekomstig) invloedrijk kamerlid?
Source: Bert Hubert's writings | Published: Sat, 30 Aug 2025 18:45:36 +0200
Link: https://berthub.eu/articles/posts/invloedrijk-kamerlid/
Eerder schreef ik in de Volkskrant dat het niet makkelijk is om als kamerlid iets te bereiken: Het inhoudelijk gemotiveerde Kamerlid loopt grote kans om geknipt en geschoren te worden door het kabinet. Als Kamerleden in een debat inhoudelijke vragen stellen aan een minister staan ze er vrijwel alleen voor. Maar de minister heeft op de achtergrond een zaal vol beleidsmedewerkers die live een antwoord op alles kunnen componeren, inclusief antwoorden waarmee het Kamerlid geheel het bos in wordt gestuurd.

66. De Digitale Bermudadriehoek in Den Haag
Source: Bert Hubert's writings | Published: Tue, 26 Aug 2025 08:01:23 +0200
Link: https://berthub.eu/articles/posts/de-digitale-bermudadriehoek-in-den-haag/
Als er kamerleden voor de commissie Digitale Zaken gevonden moeten worden steekt bijna niemand in Den Haag zijn of haar hand op. Want iedereen weet, het kost je je carrłre als je te digitaal bent. Je partij straft je af met een lullige positie op de lijst voor de volgende verkiezingen. Dit gebeurde harde werkers Hind Dekker en Lisa van Ginneken al eerder. Inmiddels is er NerdVote.nl wat aandacht vraagt voor dit probleem!

67. Nederlands Agentschap voor Disruptieve Innovatie
Source: Bert Hubert's writings | Published: Mon, 25 Aug 2025 09:30:06 +0200
Link: https://berthub.eu/articles/posts/nederlands-agentschap-disruptieve-innovatie/
De hele korte versie: klassieke manieren om wetenschap en innovatie te stimuleren zijn erg ongeschikt voor disruptieve innovatie, ideeºn die misschien helemaal niet gaan werken, maar als ze dat wel doen de wereld kunnen veranderen. En dat soort innovatie hebben we keihard nodig. De Duitse overheid heeft hiervoor een succesvol agentschap geschapen waar dit wel lukt, maar dat ging niet zonder slag of stoot. Nederland overweegt nu ook een ?Agentschap voor Disruptieve Innovatie? op te richten.

68. Rust Macros: Zero to Hero
Source: Grant Slatton's Blog | Published: Sun, 1 Sep 2024 17:45:17 +0000
Link: https://grantslatton.com/rust-macros
A comprehensive guide on Rust macros

69. Algorithms we develop software by
Source: Grant Slatton's Blog | Published: Sat, 17 Aug 2024 17:33:46 +0000
Link: https://grantslatton.com/software-pathfinding
Pathfinding applied to the software solution domain

70. Status among whom?

Source: Grant Slatton’s Blog | Published: Thu, 15 Aug 2024 16:33:36 +0000
Link: https://grantslatton.com/status-among-whom
An essay about status relativism

71. Defending abuse does not defend free software
Source: Matthew Garrett | Published: Thu, 12 Oct 2023 16:32:49 GMT
Link: https://mjg59.dreamwidth.org/68004.html
The Free Software Foundation Europe and the Software Freedom Conservancy recently released a statement that they would no longer work with Eben Moglen , chairman of the Software Freedom Law Center. Eben was the general counsel for the Free Software Foundation for over 20 years, and was centrally involved in the development of version 3 of the GNU General Public License. He’s devoted a great deal of his life to furthering free software. But, as described in the joint statement, he’s also acted abusively towards other members of the free software community. He’s acted poorly towards his own staff . In a professional context, he’s used graphically violent rhetoric to describe people he dislikes. He’s screamed abuse at people attempting to do their job. And, sadly, none of this comes as a surprise to me. As I wrote in 2017 , after it became clear that Eben’s opinions diverged sufficiently from the FSF’s that he could no longer act as general counsel, he responded by threatening an FSF board member at an FSF-run event (various members of the board were willing to tolerate this, which is what led to me quitting the board). There’s over a decade’s evidence of Eben engaging in abusive beh...