

1. Insider amnesia

Source: seangoedecke.com RSS feed | Published: Mon, 23 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/insider-amnesia/>

Speculation about what's really going on inside a tech company is almost always wrong. When some problem with your company is posted on the internet, and you read people's thoughts on it, their thoughts are almost always ridiculous. For instance, they might blame product managers for a particular decision, when in fact the decision in question was engineering-driven and the product org was pushing back on it. Or they might attribute an incident to overuse of AI, when the system in question was largely written pre-AI-coding and unedited since. You just don't know what the problem is unless you're on the inside. But when some other company has a problem on the internet, it's very tempting to jump in with your own explanations. After all, you've seen similar things in your own career. How different can it really be? Very different, as it turns out. This is especially true for companies that are unusually big or small. The recent kerfuffle over some bad GitHub Actions code is a good example of this - many people just seemed to have no mental model about how a large tech company can produce bad code, because their mental model of writing code is something like ?individual engineer main...

2. What's so hard about continuous learning?

Source: seangoedecke.com RSS feed | Published: Mon, 23 Feb 2026 00:00:00 GMT

Link: <https://seangoedecke.com/continuous-learning/>

Why can't models continue to get smarter after they're deployed? If you hire a human employee, they will grow more familiar with your systems over time, and (if they stick around long enough) eventually become a genuine domain expert. AI models are not like this. They are always exactly as capable as the first moment you use them. This is because model weights are frozen once the model is released. The model can only ?learn? as much as can be stuffed into its context window: in effect, it can take new information into its short-term working memory, but not its long-term memory. ?Continuous learning? - the ability for a model to update its own weights over time - is thus often described as the bottleneck for AGI 1 . Continuous learning is an easy technical problem However, the mechanics of continuous learning are not hard . The technical problem of ?how do you change the weights of a model at runtime? is straightforward. It's the exact same process as post-training: you simply keep running new user input through the training pipeline you already have. In a sense, every LLM since GPT-3 is already capable of continuous learning (via RL, RLHF, or whatever). It's just that the continuo...

3. Startup School - How to talk to users

Source: Eric Migicovsky's Blog RSS Feed | Published: Thu, 25 Jul 2019 00:00:00 GMT

Link: <https://ericmigi.com/blog/startup-school-how-to-talk-to-users>

2025 Note - I based this off the Mom Test, written by a fellow YC founder. Interestingly enough - if you look at the products I build?

4. Startup School - A Conversation on Hard Tech

Source: Eric Migicovsky's Blog RSS Feed | Published: Fri, 12 Oct 2018 00:00:00 GMT

Link: <https://ericmigi.com/blog/startup-school-a-conversation-on-hard-tech>

2025 Note - good convo with Adora for Startup School

5. Hardware Founders: (When) Should You Try To Raise VC Money?

Source: Eric Migicovsky's Blog RSS Feed | Published: Mon, 09 Jul 2018 00:00:00 GMT

Link: <https://ericmigi.com/blog/hardware-founders-when-should-you-try-to-raise-vc-money>

If you choose to raise any amount of venture capital investment you are making an implicit agreement with your investor that your aim is to?

6. Ladybird adopts Rust, with help from AI

Source: Simon Willison's Weblog | Published: 2026-02-23T18:52:53+00:00

Link: <https://simonwillison.net/2026/Feb/23/ladybird-adopts-rust/#atom-everything>

Ladybird adopts Rust, with help from AI Really interesting case-study from Andreas Kling on advanced, sophisticated use of coding agents for ambitious coding projects with critical code. After a few years hoping Swift's platform support outside of the Apple ecosystem would mature they switched tracks to Rust their memory-safe language of choice, starting with an AI-assisted port of a critical library: Our first target was LibJS , Ladybird's JavaScript engine. The lexer, parser, AST, and bytecode generator are relatively self-contained and have extensive test coverage through test262 , which made them a natural starting point. I used Claude Code and Codex for the translation. This was human-directed, not autonomous code generation. I decided what to port, in what order, and what the Rust code should look like. It was hundreds of small prompts, steering the agents where things needed to go. [...] The requirement from the start was byte-for-byte identical output from both pipelines. The result was about 25,000 lines of Rust, and the entire port took about two weeks. The same work would have taken me multiple months to do by hand. We've verified that every AST produced by the Rust par...

7. Writing about Agentic Engineering Patterns

Source: Simon Willison's Weblog | Published: 2026-02-23T17:43:02+00:00

Link: <https://simonwillison.net/2026/Feb/23/agentic-engineering-patterns/#atom-everything>

I've started a new project to collect and document Agentic Engineering Patterns - coding practices and patterns to help get the best results out of this new era of coding agent development we find ourselves entering. I'm using Agentic Engineering to refer to building software using coding agents - tools like Claude Code and OpenAI Codex, where the defining feature is that they can both generate and execute code - allowing them to test that code and iterate on it independently of turn-by-turn guidance from their human supervisor. I think of vibe coding using its original definition of coding where you pay no attention to the code at all, which today is often associated with non-programmers using LLMs to write code. Agentic Engineering represents the other end of the scale: professional software engineers using coding agents to improve and accelerate their work by amplifying their existing expertise. There is so much to learn and explore about this new discipline! I've already published a lot under my ai-assisted-programming tag (345 posts and counting) but that's been relatively unstructured. My new goal is to produce something that helps answer the question "how do I get good resu...

8. Writing code is cheap now

Source: Simon Willison's Weblog | Published: 2026-02-23T16:20:42+00:00

Link:

<https://simonwillison.net/guides/agentic-engineering-patterns/code-is-cheap/#atom-everything>

Agentic Engineering Patterns > The biggest challenge in adopting agentic engineering practices is getting comfortable with the consequences of the fact that writing code is cheap now . Code has always been expensive. Producing a few hundred lines of clean, tested code takes most software developers a full day or more. Many of our engineering habits, at both the macro and micro level, are built around this core constraint. At the macro level we spend a great deal of time designing, estimating and planning out projects, to ensure that our expensive coding time is spent as efficiently as possible. Product feature ideas are evaluated in terms of how much value they can provide in exchange for that time - a feature needs to earn its development costs many times over to be worthwhile! At the micro level we make hundreds of decisions a day predicated on available time and anticipated tradeoffs. Should I refactor that function to be slightly more elegant if it adds an extra hour of coding time? How about writing documentation? Is it worth adding a test for this edge case? Can I justify building a debug interface for this? Coding agents dramatically drop the cost of typing code into the co...

9. The Pants-Shitting Saga of Resizing Windows on MacOS 26 Tahoe Continues

Source: Daring Fireball | Published: 2026-02-24T00:31:00Z

Link: <https://noheger.at/blog/2026/02/12/resizing-windows-on-macos-tahoe-the-saga-continues/>

Norbert Heger: In the release notes for macOS 26.3 RC, Apple stated that the window-resizing

issue I demonstrated in my recent blog post had been resolved. You?ll never guess what happened between the RC (release candidate) version and the actual shipping version of 26.3. Just kidding, you?ll guess. ?

10. NetNewsWire 7 for Mac

Source: Daring Fireball | Published: 2026-02-23T21:41:00Z

Link: <https://netnewswire.blog/2026/01/27/netnewswire-for-mac.html>

Brent Simmons, last month: The big change from 6.2.1 is that it adopts the Liquid Glass UI and it requires macOS 26. (Note to people who aren?t on macOS 26: we fixed a lot of bugs in 6.2 and 6.2.1 knowing that many people might skip, or at least delay, installing macOS 26. Also note that there?s a page where you can get old versions of NetNewsWire .) It feels a little weird for me not to be running the latest version of NetNewsWire, but since I?m skipping MacOS 26 Tahoe, I can?t run NetNewsWire 7. I am running NetNewsWire 7 betas on my iPhone and iPad, and I?ve tried it out the secondary Mac where I do have Tahoe installed. It?s so good. And syncing works just fine with NetNewsWire 6.x, for anyone else in the fellowship sticking with MacOS 15 Sequoia. You can run NetNewsWire 7 on some devices and NetNewsWire 6 on your Sequoia Mac, and it just works. ?

11. Trader Joe?s Dark Chocolate Peanut Butter Cups

Source: Daring Fireball | Published: 2026-02-23T21:40:48Z

Link: <https://www.traderjoes.com/home/products/pdp/dark-chocolate-peanut-butter-cups-094064>

Trader Joe?s: Like their milk chocolate brethren on our shelves, our Dark Chocolate Peanut Butter Cups are made with real peanut butter that?s made with slowly roasted and ground Virginia peanuts. The luscious, smooth, rich, dark chocolate enveloping that peanut butter is crafted from high quality cacao beans. Other purveyors of peanut butter cups fill theirs with all kinds of ?extraneous? ingredients. Ours are free of such things. We eschew artificial flavors and preservatives, as well as colors other than those derived from natural sources. We?re quite certain they taste better our way. Trader Joe?s speaketh the truth and throweth some shade. ?

12. Customizing the ways the dialog manager dismisses itself: Detecting the ESC key, second (failed) attempt

Source: The Old New Thing | Published: Mon, 23 Feb 2026 15:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260223-00/?p=112080>

Sniffing the synchronous keyboard state is still not precise enough. The post Customizing the ways the dialog manager dismisses itself: Detecting the ESC key, second (failed) attempt appeared first on The Old New Thing .

13. The Little Red Dot

Source: iDiallo.com | Published: Mon, 23 Feb 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/little-red-dot?src=feed>

Sometimes, I have 50 tabs open. Looking for a single piece of information ends up being a rapid click on each tab until I find what I?m looking for. Somehow, every time I get to that LinkedIn tab, I pause for a second. I just have to click on the little red dot in the top right corner, see that there is nothing new, then resume my clicking. Why is that? Why can?t I ignore the red notification badge? When you sign up for LinkedIn for the first time, it?s right there. A little red dot in the top right corner with a number in it. It stands out against the muted grays and blues of the interface. Click on it, and you?ll discover you have a notification. It?s not from someone you know; this is a fresh new account, after all. But the dot was there anyway. Add a few connections, give it some time, and come back. Refresh the page, and you?ll have new notifications waiting. If your LinkedIn account is like mine, a ghost town, you still get the little red dot. My connections and I usually keep a few recruiters in our networks, an insurance policy in case we need to find work quickly. But we rarely, if ever, post anything. Yet whenever I log in, there?s a new notification. Sometimes it?s even...

14. Zig Tokenizer

Source: Mitchell Hashimoto | Published: Thu, 10 Feb 2022 00:00:00 GMT

Link: <https://mitchellh.com/zig/tokenizer>

(No summary)

15. Comparing Filesystem Performance in Virtual Machines

Source: Mitchell Hashimoto | Published: Fri, 10 Jan 2014 00:00:00 GMT

Link: <https://mitchellh.com/writing/comparing-filesystem-performance-in-virtual-machines>

(No summary)

16. Packer

Source: Mitchell Hashimoto | Published: Fri, 28 Jun 2013 00:00:00 GMT

Link: <https://mitchellh.com/writing/packer>

(No summary)

17. Giant Steps

Source: John D. Cook | Published: Mon, 23 Feb 2026 19:20:23 +0000

Link: <https://www.johndcook.com/blog/2026/02/23/giant-steps/>

John Coltrane's song Giant Steps is known for its unusual and difficult chord changes. Although the chord progressions are complicated, there aren't that many unique chords, only nine. And there is a simple pattern to the chords; the difficulty comes from the giant steps between the chords. If you wrap the chromatic scale around a [?] The post Giant Steps first appeared on John D. Cook .

18. Tritone substitution

Source: John D. Cook | Published: Mon, 23 Feb 2026 12:50:12 +0000

Link: <https://www.johndcook.com/blog/2026/02/23/tritone-sub/>

Big moves in roots can correspond to small moves in chords. Imagine the 12 notes of a chromatic scale arranged around the hours of a clock: C at 12:00, C? at 1:00, D at 2:00, etc. The furthest apart two notes can be is 6 half steps, just as the furthest apart two times can [?] The post Tritone substitution first appeared on John D. Cook .

19. Portable monitors are good

Source: Xe Iaso's blog | Published: Tue, 24 Feb 2026 00:00:00 GMT

Link: <https://xeiaso.net/blog/2026/portable-monitors-are-good/>

A review of portable monitors for travel

20. Pluralistic: Deplatform yourself (23 Feb 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 23 Feb 2026 10:52:48

+0000

Link: <https://pluralistic.net/2026/02/23/goodharts-lawbreaker/>

Today's links Deplatform yourself: Copyright infringement is your least entertainment dollar.

Hey look at this: Delights to delectate. Object permanence: "Lawer" threatens suit; Landmark metaphotos; 3DP v (c); Forced arbitration; Imperial Scott Walker; Keysigning ritual;

Polyfingerprinted robot dictaphone; DNS bug; Register of copyright damns term extension; How

Anonymous decides; Christchurch quake people-finder; Minor HP disenshittification; US v

developing world at WIPO; TfL v anagram tube-map; Disneyland waiting; Internet of Garbage.

Upcoming appearances: Where to find me. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em.

Colophon: All the rest. Deplatform yourself (permalink) The first time I met William Gibson ?

to interview him for the Globe and Mail on the release of 1999's All Tomorrow's Parties ? there was one question I knew I wanted to ask him: "What happens to the counterculture in the era of instantaneous commodification?" <https://craphound.com/nonfic/transcript.html> Gibson's answer stuck with me for decades: What we're doing pop culturally is like burning the rain forest.

The...

21. How Are Function Components Different from Classes?

Source: overreacted ? A blog by Dan Abramov | Published: Sun, 03 Mar 2019 00:00:00 GMT

Link: <https://overreacted.io/how-are-function-components-different-from-classes/>

They're a whole different Pokémon.

22. Coping with Feedback

Source: overreacted ? A blog by Dan Abramov | Published: Sat, 02 Mar 2019 00:00:00 GMT

Link: <https://overreacted.io/coping-with-feedback/>

Sometimes I can't fall asleep.

23. Fix Like No One's Watching

Source: overreacted ? A blog by Dan Abramov | Published: Fri, 15 Feb 2019 00:00:00 GMT

Link: <https://overreacted.io/fix-like-no-ones-watching/>

The other kind of technical debt.

24. Book Review: A Geography of Time by Robert V. Levine ?????

Source: Terence Eden's Blog | Published: Mon, 23 Feb 2026 12:34:07 +0000

Link: <https://shkspr.mobi/blog/2026/02/book-review-a-geography-of-time-by-robert-v-levine/>

This book doesn't know what it wants to be. Is it a sociology textbook, travel guide, history book, or guide to the mysteries of the world? Subtitled "the temporal misadventures of a social psychologist" it veers between hard data and well-worn anecdotes until it becomes a sort of self-help book for the time-poor 1990s American executive. Despite being well-caveated against the "dangers in?"

25. Where Do Specifications Fit in the Dependency Tree?

Source: Andrew Nesbitt | Published: 2026-02-23T10:00:00+00:00

Link: <https://nesbitt.io/2026/02/23/where-do-specifications-fit-in-the-dependency-tree.html>

Your Ruby gem declares required_ruby_version >= 3.0 . That constraint references the Ruby 3.0 language specification, expressed through the implementation version, checked against whichever runtime happens to be running, with no distinction between MRI and JRuby, and no connection to the specification document that defines what Ruby 3.0 even is. Runtimes at least show up somewhere in the tooling. Your HTTP library also depends on RFC 9110, your JSON parser on ECMA-404, your TLS implementation on RFC 8446, and none of those appear in any manifest, lockfile, or SBOM. Library dependencies get the full treatment: manifests declare them, lockfiles pin them, SBOMs track them, scanners check them for vulnerabilities. Runtime versions sit one layer down, handled differently by every ecosystem. Python has Requires-Python in package metadata, enforced by pip but ignored by trove classifiers that may disagree with it. Ruby has required_ruby_version in the gemspec, enforced by both RubyGems and Bundler. Node.js has the engines field in package.json, advisory by default in npm unless you flip a config flag. Go's go directive in go.mod was advisory until Go 1.21, when it flipped to a hard minim...

26. New Blog Post: Some Silly Z3 Scripts I Wrote

Source: Computer Things | Published: Mon, 23 Feb 2026 16:49:10 +0000

Link: <https://buttondown.com/hillelwayne/archive/new-blog-post-some-silly-z3-scripts-i-wrote/>

Now that I'm not spending all my time on Logic for Programmers, I have time to update my website again! So here's the first blog post in five months: Some Silly Z3 Scripts I Wrote .

Normally I'd also put a link to the Patreon notes but I've decided I don't like publishing gated content and am going to wind that whole thing down. So some quick notes about this post: Part of the point is admittedly to hype up the eventual release of LfP. I want to start marketing the book, but don't want the marketing material to be devoid of interest, so tangentially-related-but-independent blog posts are a good place to start. The post discusses the concept of "chaff", the enormous quantity of material (both code samples and prose) that didn't make it into the book. The book is about 50,000 words? and considerably shorter than the

total volume of chaff! I don't think most of it can be turned into useful public posts, but I'm not entirely opposed to the idea. Maybe some of the old chapters could be made into something? Coming up with a conditioned mathematical property to prove was a struggle. I had two candidates: $a == b * c \Rightarrow a / b == c$, which would have required a long tangent on how division...

27. letterboxd2jellyfin sync script

Source: Jayden's Blog | Published: 2026-02-14T14:38:00+00:00

Link: <https://jayd.ml/2026/02/14/letterboxd2jellyfin.html>

I have an incredibly niche use case where I have a physical media collection that I haven't ripped. Would like to browse that collection in Jellyfin like it's my own personal Netflix. Have a public list in Letterboxd of all the movies I own. So I wrote a quick script and published it to pypi. Shout out to letterboxdpy that made this incredibly simple. This script reads the playlist from letterboxd and writes stub .mp4 files in the appropriate Jellyfin file structure so you can browse it, eg \$ letterboxd2jellyfin -o '/path/to/jellyfin/library/Movies' -url <https://letterboxd.com/matchup/list/scream-ranked/> parsing URL Loading playlist 'scream-ranked' 'from username 'matchup' ... done! This playlist has 6 entries creating stub file deleteme/Scream (1996) /Scream (1996).mp4 creating stub file deleteme/Scream 4 (2011) /Scream 4 (2011).mp4 creating stub file deleteme/Scream 2 (1997) /Scream 2 (1997).mp4 creating stub file deleteme/Scream VI (2023) /Scream VI (2023).mp4 creating stub file deleteme/Scream (2022) /Scream (2022).mp4 creating stub file deleteme/Scream 3 (2000) /Scream 3 (2000).mp4 all done! Now obviously you can't play this, it just plays a ch...

28. Making Icon Sets Easy With Web Origami

Source: Jim Nielsen's Blog | Published: Mon, 23 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/origami-icons/>

Over the years, I've used different icon sets on my blog. Right now I use Heroicons. The recommended way to use them is to copy/paste the source from the website directly into your HTML. It's a pretty straightforward process: Go to the website Search for the icon you want. Hover it Click to ?Copy SVG? Go back to your IDE and paste it. If you're using React or Vue, there are also npm packages you can install so you can import the icons as components. But I'm not using either of those frameworks, so I need the raw SVGs and there's no npm i for those so I have to manually grab the ones I want. In the past, my approach has been to copy the SVGs into individual files in my project, like: src/ icons/ home.svg about.svg search.svg Then I have a ?component? for reading those icons from disk which I use in my template files to inline the SVGs in my HTML. For example: // Some page template file import { Icon } from './Icon.js' const template = `\${Icon('search.svg')}` Search // Icon.js import fs from 'fs' import path from 'path' const __dirname = /* Do the stuff to properly resolve the file path */; export const Icon = (name) => fs.readFileSync(path.join(__dirname, 'icons', name))...

29. How AI Labs Proliferate

Source: Jim Nielsen's Blog | Published: Sun, 22 Feb 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/how-ai-labs-proliferate/>

SITUATION: there are 14 competing AI labs. ?We can't trust any of these people with super-intelligence. We need to build it ourselves to ensure it's done right!" ?YEAH!? SOON: there are 15 competing AI labs. (See: xkcd on standards.) The irony: ?we're the responsible ones? is each lab's founding mythology as they spin out of each other. Reply via: Email • Mastodon • Bluesky

30. Agentic swarms are an org-chart delusion

Source: Westenberg. | Published: Tue, 24 Feb 2026 01:07:08 GMT

Link: <https://www.joanwestenberg.com/agentic-swarms-are-an-org-chart-delusion/>

The "agentic swarm" vision of productivity is comfortingly familiar. Which should be an immediate red flag... You take the existing corporate hierarchy, you replace the bottom layers with a swarm of AI agents, and you keep humans around as supervisors. It's an org chart with

robots

31. Thoughts on Farcaster

Source: Westenberg. | Published: Mon, 23 Feb 2026 22:07:22 GMT

Link: <https://www.joanwestenberg.com/thoughts-on-farcaster/>

For the past few weeks I've been asking myself why I'm still on Farcaster, whether I'll stay, whether I even want to. I've landed on some answers. Farcaster, for the uninitiated, was the most credible attempt anyone has made at building a

32. Everyone in AI is building the wrong thing for the same reason

Source: Westenberg. | Published: Mon, 23 Feb 2026 11:09:55 GMT

Link:

<https://www.joanwestenberg.com/everyone-in-ai-is-building-the-wrong-thing-for-the-same-reason/>

Every AI founder I talk to is on an accelerating treadmill, burdened by a nagging suspicion that the entire industry is moving too fast in a direction that doesn't quite make sense, with no idea about how to get off. There is an overwhelming feeling that if everyone

33. The Bus Ticket Theory of Genius

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/genius.html>

(No summary)

34. General and Surprising

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/sun.html>

(No summary)

35. Charisma / Power

Source: Paul Graham: Essays | Published: -

Link: <http://www.paulgraham.com/pow.html>

(No summary)

36. Jan '26 Notes

Source: Susam Pal | Published: Thu, 29 Jan 2026 00:00:00 +0000

Link: <https://susam.net/notes-2026-01.html>

In these monthly notes, I jot down ideas and references I encountered during the month that I did not have time to expand into their own posts. A few of these may later develop into independent posts but most of them will likely not. In any case, this format ensures that I record them here. I spent a significant part of this month studying the book Algebraic Graph Theory by Godsil and Royle, so many of the notes here are about it. There are a few non-mathematical, technical notes towards the end. Contents Cayley Graphs Vertex Transitive Graphs Arc-Transitive Graphs Bipartite Graphs and Cycle Parity Tutte's Theorem Tutte's 8-Cage Linear Congruential Generator Numbering Lines Cayley Graphs Let $\langle G \rangle$ be a group and let $\langle C \subseteq G \rangle$ such that $\langle C \rangle$ is closed under taking inverses and does not contain the identity, i.e., $\forall x \in C, x^{-1} \in C, \forall e \notin C. \exists g \in C. g^{-1} = e$. Then the Cayley graph $\langle X(G, C) \rangle$ is the graph with the vertex set $\langle V(X(G, C)) \rangle$ and edge set $\langle E(X(G, C)) \rangle$ defined by $\begin{aligned} V(X(G, C)) &= G, \\ E(X(G, C)) &= \{ gh : hg^{-1} \in C \}. \end{aligned}$ The set $\langle C \rangle$ is known as the connection set. Vertex Transitive Graphs A graph $\langle X \rangle$...

37. Building Docs.rs

Source: the website of jyn | Published: 2019-11-26T00:00:00+00:00

Link: <https://jyn.dev/building-docs-rs/>

How I became the 2nd-most prolific contributor to docs.rs

38. Intro to Reverse Engineering

Source: the website of jyn | Published: 2019-11-11T00:00:00+00:00

Link: <https://jyn.dev/talks/assembly/>

(No summary)

39. Intro to C

Source: the website of jyn | Published: 2019-11-05T00:00:00+00:00

Link: <https://jyn.dev/talks/c/>

(No summary)

40. Redis will remain BSD licensed

Source: antirez.com | Published: Wed, 22 Aug 2018 15:45:52 +0200

Link: <http://antirez.com/news/120>

Today a page about the new Common Clause license in the Redis Labs web site was interpreted as if Redis itself switched license. This is not the case, Redis is, and will remain, BSD licensed. However in the era of [edit] uncontrollable spreading of information, my attempts to provide the correct information failed, and I'm still seeing everywhere ?Redis is no longer open source?. The reality is that Redis remains BSD, and actually Redis Labs did the right thing supporting my effort to keep the Redis core open as usually. What is happening instead is that certain Redis modules, developed inside Redis Labs, are now released under the Common Clause (using Apache license as a base license). This means that basically certain enterprise add-ons, instead of being completely closed source as they could be, will be available with a more permissive license. I think that Redis Labs Common Clause page did not provide a clear and complete information, but software companies often make communication errors, it happens. To me however, it looks more important that while running a system software business in the ?cloud era? (LOL) is very challenging using an open source license, yet Redis Labs tot...

41. Redis Lua scripting: several security vulnerabilities fixed

Source: antirez.com | Published: Wed, 13 Jun 2018 19:15:05 +0200

Link: <http://antirez.com/news/119>

A bit more than one month ago I received an email from the Apple Information Security team. During an auditing the Apple team found a security issue in the Redis Lua subsystem, specifically in the cmmsgpack library. The library is not part of Lua itself, it is an implementation of MessagePack I wrote myself. In the course of merging a pull request improving the feature set, a security issue was added. Later the same team found a new issue in the Lua struct library, again such library was not part of Lua itself, at least in the release of Lua we use: we just embedded the source code inside our Lua implementation in order to provide some functionality to the Lua interpreter that is available to Redis users. Then I found another issue in the same struct package, and later the Alibaba team found many other issues in cmmsgpack and other code paths using the Lua API. In a short amount of time I was sitting on a pile of Lua related vulnerabilities. Those vulnerabilities are mostly relevant in the specific case of providing managed Redis servers on the cloud, because it is very unlikely that the vulnerabilities discovered can be used without direct access to the Redis server: many Redis user...

42. Clarifications on the Incapsula Redis security report

Source: antirez.com | Published: Sat, 02 Jun 2018 19:52:39 +0200

Link: <http://antirez.com/news/118>

A few days ago I started my day with my Twitter feed full of articles saying something like: ?75% of Redis servers infected by malware?. The obvious misquote referred to a research by Incapsula where they found that 75% of the Redis instances left open on the internet, without any protection, on a public IP address, are infected [1]. [1] <https://www.incapsula.com/blog/report-75-of-open-redis-servers-are-infected.html> Many folks don't need any clarification about all this, because if you have some grip on computer security and how Redis works, you can contextualize all this without much efforts. However I'm writing

this blog post for two reasons. The obvious one is that it can help the press and other users that are not much into security and/or Redis to understand what's going on. The second is that the exposed Redis instances are a case study about safe defaults that should be interesting for the security circles. The Incapsula report === Let's start with the Incapsula report. What they did was to analyze exposed Redis instances on the internet. Instances that everybody from any place of the internet can access, because they are listening for connections in a public IP address, w...

43. The Book Of CP-System

Source: Fabien Sanglard | Published: 25 Sep 2022 00:00:00 +0000

Link: <https://fabiensanglard.net/cpsb/index.html>

(No summary)

44. The Beautiful Diablo 2 Resurrected machine

Source: Fabien Sanglard | Published: 08 May 2022 00:00:00 +0000

Link: <https://fabiensanglard.net/d2r/index.html>

(No summary)

45. USB Cheat Sheet

Source: Fabien Sanglard | Published: 05 May 2022 00:00:00 +0000

Link: <https://fabiensanglard.net/usbcheat/index.html>

(No summary)

46. 'Weekly' Links for February 26, 2025

Source: Dr Paris Buttfield-Addison | Published: Wed, 26 Feb 2025 00:00:00 +0000

Link: <https://hey.paris/posts/weekly-links/02262025/>

Some new links! Huzzah! ?The closer to the train station, the worse the kebab? - A ?Study?

Blogging in the age of ?AI? SQL Injection Court (also worth reading this vaguely similar story about 32 million emails) More soon!

47. 'Weekly' Links for February 14, 2025

Source: Dr Paris Buttfield-Addison | Published: Fri, 14 Feb 2025 00:00:00 +0000

Link: <https://hey.paris/posts/weekly-links/02132025/>

So it didn't end up being weekly. Here's some links! Duh: Fewer students are enrolling in doctoral degrees Torment nexus is going well: ? A significant negative correlation between frequent AI tool usage and critical thinking abilities ? Who would have guessed? AI code generation leads to more code, more duplicated code, less cohesive systems, more defects Residential networking over telephone More next ?week?!

48. Yarn Spinner at GDC 2025

Source: Dr Paris Buttfield-Addison | Published: Wed, 12 Feb 2025 00:00:00 +0000

Link: <https://hey.paris/posts/gdc2025/>

We're thrilled to announce that Jon, from the Secret Lab and Yarn Spinner team will be presenting a talk at the GDC this year in San Francisco: ?DREDGE? and Yarn Spinner: Building Narrative with Open Source . The talk will explore how the indie hit, DREDGE , uses our Yarn Spinner technology for narrative, and how you can do the same. Join us at GDC! Email me via paris AT yarnspinner.dev if you'd like to organise a meeting, too.

49. Object Oriented Thinking, Part I

Source: Beej's Bit Bucket | Published: Tue, 29 Dec 2009 00:00:00 +0000

Link: <http://beej.us/blog/data/object-oriented-programming/>

(No summary)

50. Dijkstra's Shortest Path

Source: Beej's Bit Bucket | Published: Fri, 25 Dec 2009 00:00:00 +0000

Link: <http://beej.us/blog/data/dijkstras-shortest-path/>

(No summary)

51. Parsing XML Streams with SAX

Source: Beej's Bit Bucket | Published: Wed, 23 Dec 2009 00:00:00 +0000

Link: <http://beej.us/blog/data/xml-sax/>

No need to parse that whole DOM? just handle tags as they show up.

52. Weekly Update 492

Source: Troy Hunt | Published: Tue, 24 Feb 2026 00:38:59 GMT

Link: <https://www.troyhunt.com/weekly-update-492/>

The recurring theme this week seems to be around the gap between breaches happening and individual victims finding out about them. It's tempting to blame this on the corporate victim of the breach (the hacked company), but they're simultaneously dealing with a criminal intrusion, a ransom

53. start-vibe-coding-fast

Source: Chad Nauseam Home | Published: -

Link: <https://chadnauseam.com/coding/tips/start-vibe-coding-fast>

(No summary)

54. Taking action against AI harms

Source: Anil Dash | Published: 2026-02-24T00:00:00Z

Link: <https://anildash.com/2026/02/23/taking-action-ai-harms/>

Taking action on AI harms against kids In my last piece, I talked about the harms that AI is visiting on children through the irresponsible choices made by the platforms creating those products. While we dove a bit into the incentives and institutional pressures that cause those companies to make such wildly irresponsible decisions, what we haven't yet reckoned with is how we hold these companies accountable. Often, people tell me they feel overwhelmed at the idea of trying to engage with getting laws passed, or fighting a big political campaign to rein in the giant tech companies that are causing so much harm. And grassroots, local organizing can be extraordinarily effective in standing up for the values of your community against the agenda of the Big AI companies. But while I think it's vital that we pursue systemic justice (and it's the only way to stop many kinds of harm), I do understand the desire for something more immediate and human-scale. So, I wanted to share some direct, personal actions that you can take to respond to the threats that Big AI has made against kids. Each of these tactics have been proven effective by others who have used the same strategies, so you can...

55. Self-promotion

Source: Keygen Blog | Published: Mon, 03 Feb 2025 06:00:00 GMT

Link: <https://keygen.sh/blog/self-promotion/>

On the mistakes new founders make when it comes to self-promotion online.

56. Time away

Source: Keygen Blog | Published: Tue, 04 Feb 2025 06:00:00 GMT

Link: <https://keygen.sh/blog/time-away/>

On what matters ? and taking time away as a founder.

57. The space elevator

Source: Keygen Blog | Published: Wed, 05 Feb 2025 06:00:00 GMT

Link: <https://keygen.sh/blog/the-space-elevator/>

A short story on bootstrapping and startups.

58. Chatcontrol Ødition 2025 en bref

Source: Bert Hubert's writings | Published: Sun, 17 Aug 2025 17:15:00 +0200

Link: <https://berthub.eu/articles/posts/chatcontrol-francais/>

En bref, sous l'impulsion du Danemark, de nombreux États membres de l'UE demandent que les messageries comme WhatsApp/Signal/etc. soient contraintes d'inspecter toutes nos photos et tous nos liens en utilisant de l'IA. Si l'IA a le moindre doute? quant à la possibilité qu'il s'agisse de pédopornographie, votre photo, votre localisation, votre numéro de téléphone ainsi que d'autres détails sont signalés à Europol et aux forces de police locales. La suite n'est pas précisée, mais vos vacances à Chypre pourraient se terminer par l'obligation d'expliquer vos photos au poste de police local.

59. Chatcontrol 2025 edition in Brief

Source: Bert Hubert's writings | Published: Sun, 17 Aug 2025 14:30:00 +0200

Link: <https://berthub.eu/articles/posts/chatcontrol-in-brief/>

In short, led by Denmark, many EU member states are arguing for forcing WhatsApp/Signal/etc to inspect all our photos and links, using AI. If the AI is in any doubt? if this might be child pornography, your photo, location, phone number & other details get reported to a new EU Centre? and thence to Europol and a local police force. What happens next is not specified, but your holiday in Cyprus might end up with you having to explain your photos over at the local police station.

60. Chatcontrol 2025 (kurz)

Source: Bert Hubert's writings | Published: Sun, 17 Aug 2025 14:07:15 +0200

Link: <https://berthub.eu/articles/posts/chatcontrol-kurz/>

Kurz zusammengefasst: Unter der Führung Dänemarks fordern viele EU-Mitgliedsstaaten, WhatsApp, Signal und co dazu zu zwingen, dass alle darüber versendeten Fotos und Links mit KI durchsucht werden. Wenn die KI auch nur den geringsten Verdacht hat, dass es sich um Kinderpornografie handeln könnte, werden dein Bild, dein Standort, deine Telefonnummer und weitere Details automatisch an Europol und eine örtliche Polizeibehörde gemeldet. Jetzt gegen Chat Control aktiv werden! Auch wenn nicht festgelegt ist was danach passiert, kann dein Urlaub auf Zypern schnell damit enden, dass du deine privaten Fotos einem unbekannten Polizisten erklären musst.

61. Pockets of Humanity

Source: Herman's blog | Published: 2026-02-23T13:24:23.228866+00:00

Link: <https://herman.bearblog.dev/pockets-of-humanity/>

There's a conspiracy theory that suggests that since around 2016 most web activity is automated. This is called Dead Internet Theory , and while I think they may have jumped the gun by a few years, it's heading that way now that LLMs can simulate online interactions near-flawlessly. Without a doubt there are tens (hundreds?) of thousands of interactions happening online right now between bots trying to sell each other something . This sounds silly, and maybe a little sad, since the internet is the commons that has historically belonged to, and been populated by all of us. This is changing. Something interesting happened a few weeks ago where an OpenClaw instance , named MJ Rathbun, submitted a pull request to the matplotlib repository, and after having its code rejected on the basis that humans needed to be in the loop for PRs, it proceeded to do some research on the open-source maintainer who denied it, and wrote a "hit piece" on him, to publicly shame him for feeling threatened by AI...or something. The full story is here and I highly recommend giving it a read. A lot of the discourse around this has taken the form of "haha, stupid bot", but I posit that it is the beginning of s...