

1. Patch Tuesday, May 2026 Edition

Source: Krebs on Security | Published: Tue, 12 May 2026 21:46:45 +0000

Link: <https://krebsonsecurity.com/2026/05/patch-tuesday-may-2026-edition/>

Artificial intelligence platforms may be just as susceptible to social engineering as human beings, but they are proving remarkably good at finding security vulnerabilities in human-made computer code. That reality is on full display this month with some of the more widely-used software makers -- including Apple, Google, Microsoft, Mozilla and Oracle -- fixing near record volumes of security bugs, and/or quickening the tempo of their patch releases.

2. Bambu Lab is abusing the open source social contract

Source: Jeff Geerling | Published: Tue, 12 May 2026 09:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/bambu-lab-abusing-open-source-social-contract/>

Last year I said I'd probably never recommend another Bambu Lab printer again . I still use my P1S, but after Bambu Lab started pushing their always-connected cloud solution as the new default: I blocked the printer from the Internet via my OPNsense Firewall I stopped updating the firmware I locked the printer into Developer mode I deleted Bambu Studio and started using OrcaSlicer I had to do that to keep it under my control, instead of Bambu's.

3. datasette 1.0a29

Source: Simon Willison's Weblog | Published: 2026-05-12T23:41:06+00:00

Link: <https://simonwillison.net/2026/May/12/datasette/#atom-everything>

Release: datasette 1.0a29 New TokenRestrictions.abbreviated(datasette) utility method for creating "_" dictionaries. #2695 Table headers and column options are now visible even if a table contains zero rows. #2701 Fixed bug with display of column actions dialog on Mobile Safari. #2708 Fixed bug where tests could crash with a segfault due to a race condition between Datasette.close() and Datasette.close(). #2709 That segfault bug was gnarly . I added a mechanism to Datasette recently that would automatically close connections at the end of each test, but it turned out that introduced a race condition where an in-flight query could sometimes be executing in a thread against a connection while it was being closed. I ended up solving that by having Codex CLI (with GPT-5.5 xhigh) create a minimal Dockerfile that recreated the bug. Tags: projects , datasette

4. Quoting Mo Bitar

Source: Simon Willison's Weblog | Published: 2026-05-12T22:59:58+00:00

Link: <https://simonwillison.net/2026/May/12/mo-bitar/#atom-everything>

Now, if your CEO has never heard the phrase Ralph Loop, oh man, you are less than 30 days away from your next promotion. I'm not even exaggerating. Walk into his office, close the door, and say, hey chief, been experimenting with something. It's called Ralph Loops. And I think it could change literally everything. And he's gonna say, what's a Ralph loop? And you will say, give me \$18,000 worth of API credits and I'll show you. Now you won't actually do anything, because you can't do anything. Because nobody can, because nobody knows what they're doing. But by the time he figures that out, you'll have a new title, and equity bump. [...] Talk about automation constantly. Nothing arouses the slumbering capitalists than the mention of automation. Drop names too, bro. Like talk about specific team members you can automate out of existence. Be like, yo, I automated Gary, bro. Tag Gary in the message. Tag him in Slack in a very public channel. Be like, yo, I just automated @Gary. His function has been Ralph Looped. And tag your CEO in the same message. You think you're getting laid off after that? ? Mo Bitar , The Unethical Guide to Surviving AI Layoffs, TikTok Tags: ai-ethics , tiktok ,...

5. Quoting Mitchell Hashimoto

Source: Simon Willison's Weblog | Published: 2026-05-12T22:21:51+00:00

Link: <https://simonwillison.net/2026/May/12/mitchell-hashimoto/#atom-everything>

The thing about 90% of TDMs [Technical Decision Makers] is that they're motivated primarily by NOT GETTING FIRED. These aren't people who browser Lobsters or push to GH on the weekend. These are people that work 9 to 5, get paid, go home, and NEVER THINK ABOUT WORK AGAIN. So to achieve all that, they follow secular trends supported by analysts and broad public sentiment. Oh, Gartner said that "AI strategy" is most important? McKinsey said "context" needs to be managed? Well, "Context Engine for AI Apps" is going to be defensible. Buy it. ? Mitchell Hashimoto , in a conversation about the design of the Redis homepage Tags: marketing , mitchell-hashimoto , redis

6. Kagi Snaps

Source: Daring Fireball | Published: 2026-05-12T21:06:21Z

Link: <https://help.kagi.com/kagi/features/snaps.html>

Kagi's documentation: Typing @r headphones will search for ?headphones? but limit the results to reddit.com (r is the short code for Reddit). This allows you to quickly find relevant content on a specific site using Kagi's powerful index. It is effectively the same as doing headphones site:old.reddit.com . Its relative, Bangs feature, invoked by using ?!r headphones?, would redirect the user to Reddit's internal search. I learned about the snaps feature from a Kagi blog post a few months ago, and I've been loving it ever since. From that post: You can also use Snaps to quickly search within Kagi's knowledgeable. For example, the @help snap searches the Kagi help docs, handy for when you want to quickly look into a feature. User tip: ?I recommend combining Snaps with the ?!m feeling lucky? bang: ! with no short code. Like searching @gh curl ! to go to the curl repo. I've never actually looked any of these up. I just guessed at the ones I most want to use and they all worked on the first try. @nyt returns results from The New York Times; @wsj is for The Wall Street Journal. Take a guess what @df does. And you can add your own custom bangs/snaps in Kagi's settings. It's easy. In fa...

7. Seriously, Give Kagi a Try

Source: Daring Fireball | Published: 2026-05-12T20:31:49Z

Link: https://daringfireball.net/2025/04/try_switching_to_kagi

Quoting from a post I wrote a year ago: Like, even if I use the magic &udm=14 parameter with Google search, to get ?disenshittified? results from Google , I find I get better results from Kagi. When I know there's one right answer (say, a specific article I remember reading and want to find again), Kagi is more likely than Google to list it first. If it's a years-old article, Kagi is way more likely than Google to find it at all. For me, Google (and, alas, DuckDuckGo too) have largely stopped working reliably for finding not-recent stuff on the web. Not true with Kagi. I used DuckDuckGo for years as my default search, and for those years, I found it largely on par with Google. But it felt like every once in a while ? maybe, say, once or twice a month ? DuckDuckGo would come up dry in its results. DuckDuckGo pioneered a trick they call Bangs . Include !g to any search terms, and instead of performing the search itself, DuckDuckGo will redirect that search to Google. They have a whole bunch of these Bangs ? ?!a? for Amazon search, ?!nf? for Netflix. There are literally thousands of them (which of course they allow you to search for). The only one I ever really used though was !g , f...

8. Search Ads as a Vector for Travel Scams

Source: Daring Fireball | Published: 2026-05-12T20:22:20Z

Link: <https://www.wsj.com/lifestyle/travel/the-simple-travel-scam-that-cost-a-seasoned-traveler-over-12-000-7d317f20?st=WDTpv5>

Dawn Gilbertson, writing for The Wall Street Journal (gift link): Calder says that he tried to rebook at the given link a few times but that it wouldn't work. He became worried new flight options were dwindling, so he googled the airline's customer-service number. (There was a link to customer-service contacts way down in the email that he initially overlooked.) The rest of the story is sadly familiar to the Better Business Bureau, Federal Trade Commission, airlines and consumer advocates. It's called an impostor scam. This can occur when a company impersonates an airline's customer-service number or site, often by buying a sponsored ad on a search platform. The company is hoping that panicked consumers trying, say, to rebook a flight

will click on the first link they see, bringing them to unscrupulous parties that try to charge exorbitant fees. I've written before about such tactics, and they are only becoming more sophisticated with AI. "Scammers thrive on that sense of urgency," says John Breyault, a vice president at the National Consumers League whose coverage area includes fraud. The person who answered Calder's call identified himself as a Lufthansa representative and aske...

9. Initialize and print 128-bit integers in C

Source: John D. Cook | Published: Tue, 12 May 2026 12:20:20 +0000

Link: <https://www.johndcook.com/blog/2026/05/12/c-128-bit-int/>

If you look very closely at my previous post, you'll notice that I initialize a 128-bit integer with a 64-bit value. The 128-bit unsigned integer represents the internal state of a random number generator. Why not initialize it to a 128-bit value? I was trying to keep the code simple. A surprising feature of C [?] The post Initialize and print 128-bit integers in C first appeared on John D. Cook .

10. Hacking the lehmer64 RNG

Source: John D. Cook | Published: Tue, 12 May 2026 11:07:31 +0000

Link: <https://www.johndcook.com/blog/2026/05/12/hacking-the-lehmer64-rng/>

A couple days ago I wrote about hacking the Mersenne Twister. I explained how to recover the random number generator's internal state from a stream of 640 outputs. This post will do something similar with the lehmer64 random number generator. This generator is very simple to implement. Daniel Lemire found it to be "the fastest [?]" The post Hacking the lehmer64 RNG first appeared on John D. Cook .

11. Learning Software Architecture

Source: matklad | Published: 2026-05-12T00:00:00+00:00

Link: <https://matklad.github.io/2026/05/12/software-architecture.html>

In reply to an email asking about learning software design skills as a researcher physicist:

12. Pluralistic: A fascist paradigm (12 May 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 12 May 2026 07:22:17 +0000

Link: <https://pluralistic.net/2026/05/12/donella-meadows/>

Today's links A fascist paradigm: The change that changed everything. Hey look at this: Delights to delectate. Object permanence: Openstreetmap x Isle of Wight; Found grocery lists; Mayor wants to pray away potholes; Designing a D120; "Too Like the Lightning." Upcoming appearances: Barcelona, Berlin, Hay-on-Wye, London, NYC, Edinburgh. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. A fascist paradigm (permalink) Yesterday, I attended a workshop on systems thinking and political change, which included a presentation on the work of Donella Meadows, whose Thinking in Systems is a canonical work on the subject: https://en.wikipedia.org/wiki/Thinking_In_Systems:_A_Primer "Systems thinking" is an analytical framework that treats the world as a mesh of interconnected, nonlinear components and relationships that can't be easily understood or steered. A complex system isn't merely "complicated." A mechanical watch is complicated, in that it has many parts that work together in ways that require training and specialized knowledge to understand. But it isn't "complex" b...

13. In praise of the Linux kernel netconsole (in the right circumstances)

Source: Chris's Wiki :: blog | Published: 2026-05-13T00:30:34Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/NetconsoleForDesktopPanics>

The Linux kernel's netconsole is a kernel module that will "log kernel printk messages over UDP" to a remote system, which makes it another form of kernel (message) console . These days it can be activated either on boot or after boot, and in the past I've had mixed views of it . However, I recently had a nice experience with netconsole that's left me more well inclined to

it in specific situations. A while back, my home desktop started locking up every once in a while . Several years ago my home desktop had a somewhat similar problem that was due to hardware issues , but the lockups this time were different, in that the machine would lock up for a bit and then reboot on its own. Local logs showed nothing, but I happen to have another machine sitting around so I thought I might as well try netconsole again. These days netconsole can be enabled on the fly: `modprobe netconsole cd /sys/kernel/config/netconsole mkdir heedra cd heedra echo em0 >dev_name echo 192.168.X.Y >remote_ip echo 1 >enabled` (This other machine is called heedra for obscure reasons.) On the other machine I ran a simple script to capture output inside a screen session: `#!/bin/sh while ;; do nc --recv-only -u -l 6666...`

14. A code (reformatting) conundrum in Python, and heuristics

Source: Chris's Wiki :: blog | Published: 2026-05-12T02:27:54Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/python/CodeFormattingBlockHeuristics>

Suppose that you are a Python code reformatter, and someone hands you the following snippet of Python code to act on: if something: blah blah blah [...] final-line some-statement [... more statements ...] Here's the question: should you reindent 'some-statement' so that it's part of the ' if ' block? One answer is that you absolutely should not. The current code is valid Python code, and you are a reformatter for style, not to correct (presumed) errors. Since this is valid code, you should re-flow line wrapping and so on within blocks, but not change what block valid code is part of. Another answer is that maybe the person writing this code made a mistake. Style wise, it's common to add a blank line between the end of an indented block and following code; the lack of a blank line suggests that a mistake was made. So maybe you should reindent 'some-statement' to where it properly should be, especially if you have a style rule that says that there should be blank lines in this sort of situation. (Of course, you could also opt to add the blank line that your style guide says should be there and not change what block a statement goes in. But we're in heuristics territory here.) If you...

15. Not a Security Issue

Source: Andrew Nesbitt | Published: 2026-05-12T10:00:00+00:00

Link: <https://nesbitt.io/2026/05/12/not-a-security-issue.html>

Daniel Stenberg wrote yesterday about Mythos finding a real curl vulnerability, which sent me poking around the curl tree with a couple of AI-assisted scanners to see what else turned up in a codebase that has already had every fuzzer and auditor on earth pointed at it. The findings were better than expected, and one reason stood out: the scanners had read docs/VULN-DISCLOSURE-POLICY.md and applied it. A whole class of results came back labelled in effect ?real bug, worth fixing, but not a security issue per the project?s own policy?: server-triggered NULL dereferences, small leaks, things that only fire if you can already control the command line. The tool had found them, checked them against curl?s published list of what doesn?t count, and demoted them before I had to. One example from the demoted pile: the scanner found that tool_formparse.c walks a linked list of -F form parts recursively with no depth limit, built a 150k-line config file to prove it, and got an ASan stack-overflow trace out the other end. It then wrote, in its own summary, ?trigger requires user to run curl with attacker-supplied config or args, so excluded by policy,? filed it under quality bugs, and moved o...

16. Position or Perish: The Narrative Blueprint

Source: Westenberg. | Published: Tue, 12 May 2026 10:15:17 GMT

Link: <https://www.joanwestenberg.com/position-or-perish-the-narrative-blueprint/>

Avis was losing \$3.2 million a year; and they'd been unprofitable for thirteen straight. In 1962, they sat at number two in American car rental, well behind Hertz, with no plausible path to catching up. Robert Townsend, the new president, hired Doyle Dane Bernbach and asked

17. Where Are All The Data Centers?

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 12 May 2026 16:17:30 GMT

Link: <https://www.wheresyoured.at/where-are-all-the-data-centers/>

If you liked this piece, please subscribe to my premium newsletter. It's \$70 a year, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA , Anthropic and OpenAI&

18. Regatta Starting Stations ? Chi-squared Continued

Source: Entropic Thoughts | Published: Tue, 12 May 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/regatta-starting-stations>

In the Henley Royal Regatta two teams at a time propel their boats up a river and compete to be first to go a distance. Teams get assigned to their starting stations ? Berkshire or Buckinghamshire ? at random. From there, it is a straight shot up the river, with the lane from each starting station being seemingly identical. I didn't know any of this, but a reader reached out some time ago because they had noticed something odd about this, and they wanted to borrow me as a sounding board. Here's the odd thing: the team that starts from the Berkshire station has won 53.5 % of the 7555 races in the historic data this reader looked at. This is highly unexpected. If teams are assigned at random, and the starting stations are practically equal, then the starting station of the winning team should be a coin flip. If we flip 7555 coins, we would never have as many as 53.5 % come up heads. (Continue reading the full article on the web.)

19. Building Software Requires Digestion

Source: Jim Nielsen's Blog | Published: Tue, 12 May 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/software-requires-digestion/>

Here's Scott Jenson in his insightful piece ?The Ma of a New Machine? : the chatbot interface [makes us] feel like deep cognitive work is happening. But the interface is fundamentally reactive. It spits complex text at you, you skim it quickly, and you immediately type a reaction to keep the momentum going. My hypothesis is that the very structure of the chatbot interface (type, read, type again) actively discourages reflection. When you are moving too fast, you get stuck in a groove. You literally need to take a break, step back, and basically step out of this groove so you can view the problem from a new angle. We've all walked away from a tough problem only to have the solution arrive unbidden into our thoughts later in the day. In my decades+ experience designing and developing software, I can't count the number of times I've stepped away from a problem at the computer only to return and find the problem magically resolved in my brain. But the human-computer interaction of prompting doesn't encourage the use of that skill in our subconscious. In fact, I think it actively discourages it (our tools shape us). Scott talks about this Japanese concept called ?Ma? which is about del...

20. Shame them, shun them, ban them, beat them!

Source: Experimental History | Published: Tue, 12 May 2026 13:23:30 GMT

Link: <https://www.experimental-history.com/p/shame-them-shun-them-ban-them-beat>

OR: crock pots and lightning bolts

21. Hoe werken de nieuwe (mRNA) Corona Vaccins

Source: Bert Hubert's writings | Published: Sat, 21 Nov 2020 19:05:02 +0100

Link: <https://berthub.eu/articles/posts/hoe-werken-de-nieuwe-corona-vaccins/>

Er zijn nu twee vaccins die meer dan 90% effectief lijken te zijn tegen Corona. Dit is heel hard gegaan, maar dat is goed te verklaren. Universitair Hoofddocent en kinderarts Patricia Bruijning heeft het hier heel duidelijk uitgelegd. In het kort, we hebben zo snel vaccins want: Het SARS-CoV-2 virus is een relatief makkelijk doelwit om een vaccin tegen te ontwikkelen. (het is geen griep) De coronavaccins worden ontwikkeld met nieuwe en veel snellere technologie Geld was geen probleem Toelatingsinstanties voor vaccins waren vanaf het begin nauw betrokken.

22. Patricia Bruijning: Hoe kan het dat we zo snel een coronavaccin hebben?

Source: Bert Hubert's writings | Published: Fri, 20 Nov 2020 20:32:29 +0100

Link: <https://berthub.eu/articles/posts/patricia-bruijning-over-corona-vaccins/>

Ik heb hier een Twitter-draadje van Patricia Bruijning omgezet naar een webpagina, waardoor het verhaal wat makkelijker te lezen is voor mensen zonder Twitter. Dr Bruijning is Universitair Hoofddocent epidemiologie bij UMC Utrecht maar is ook kinderarts. Ze weet dus goed waar ze het over heeft en ook hoe de praktijk werkt. En ik ben erg blij dat ze de tijd heeft genomen ons goed voor te lichten! Hier haar uitleg hoe het komt dat we zo snel een coronavaccin krijgen, en waarom het veilig zal zijn:

23. Corona-Update november: Twee vaccins en sneltests!

Source: Bert Hubert's writings | Published: Mon, 16 Nov 2020 12:42:49 +0100

Link: <https://berthub.eu/articles/posts/corona-update-november/>

Ja, daar zijn we weer! Voor het eerst [sinds

juli](<https://berthub.eu/articles/posts/corona-update-eind-juli/>) weer een Corona Update. Er is inmiddels weer genoeg te melden: vaccins, tests, cijfers, medicijnen en vitamine D. Vaccins Om te beginnen, wow. En nog een keer wow. SINDS VANDAAG ER ZIJN TWEE WERKENDE VACCINS! En nog gekker, ze zijn in februari 2020 al ontworpen, toen het RNA (de broncode) van het virus pas net bekend was. Vergeleken met zelfs maar een paar jaar geleden is dit pure science fiction.