

## 1. Warelay -> OpenClaw

Source: Simon Willison's Weblog | Published: 2026-05-16T20:23:30+00:00

Link: <https://simonwillison.net/2026/May/16/openclaw-names/#atom-everything>

In preparation for a lightning talk I'm giving at PyCon US this afternoon I decided to figure out how many names OpenClaw has actually had since that first commit back in November. Thanks to this `first_line_history.py` tool ( code here ) the answer, according to the Git history of the OpenClaw README, is: Warelay ? CLAWDIS ? CLAWDBOT ? Clawdbot ? Moltbot ?? OpenClaw Or in detail (the output from the tool): 2025-11-24T11:23:15+01:00 16dfc1a # Warelay ? WhatsApp Relay CLI (Twilio) 2025-11-24T11:41:37+01:00 d4153da # ? Warelay ? WhatsApp Relay CLI (Twilio) 2025-11-24T17:47:57+01:00 343ef9b # ? warelay ? WhatsApp Relay CLI (Twilio) 2025-11-25T04:44:10+01:00 14b3c6f # ? warelay ? WhatsApp Relay CLI 2025-11-25T12:48:40+01:00 4814021 # ? warelay ? Send, receive, and auto-reply on WhatsApp?Twilio-backed or QR-linked. 2025-11-25T13:50:18+01:00 d51a3e9 # warelay ? - Send, receive, and auto-reply on WhatsApp via Twilio or QR-linked WhatsApp Web; webhook setup in one command 2025-11-25T13:51:13+01:00 4d2a8a8 # ? warelay ? Send, receive, and auto-reply on WhatsApp?Twilio-backed or QR-linked. 2025-11-25T14:52:43+01:00 1ef7f4d # ? warelay ? Send, receive, and auto-reply on WhatsApp. 2025-12-03T15...

## 2. Quoting Julia Evans

Source: Simon Willison's Weblog | Published: 2026-05-16T16:45:37+00:00

Link: <https://simonwillison.net/2026/May/16/julia-evans/#atom-everything>

[...] in the last 10 years I've learned to really love and respect CSS as a technology. So I decided years ago that I wanted to react to 'CSS is hard' by getting better at CSS and taking it seriously as a technology, instead of devaluing it. Doing that changed everything for me: I learned that so many of my frustrations ('centering is impossible?') had been addressed in CSS a long time ago, and that also what 'centering?' means is not always straightforward and it makes sense that there are many ways to do it. CSS is hard because it's solving a hard problem! ? Julia Evans , Moving away from Tailwind, and learning to structure my CSS Tags: css , julia-evans

## 3. datasette-llm-limits 0.1a0

Source: Simon Willison's Weblog | Published: 2026-05-15T00:42:09+00:00

Link: <https://simonwillison.net/2026/May/15/datasette-llm-limits/#atom-everything>

Release: datasette-llm-limits 0.1a0 This plugin works in conjunction with datasette-llm and datasette-llm-accountant to let you configure a per-user (or global) spending limit for LLM usage inside of Datasette. Configuration looks something like this: plugins : datasette-llm-limits : limits : per-user-daily : scope : actor window : rolling-24h amount\_usd : 1.00 Tags: llm , datasette

## 4. How I use LLMs as a staff engineer in 2026

Source: seangoedecke.com RSS feed | Published: Sun, 17 May 2026 00:00:00 GMT

Link: <https://seangoedecke.com/how-i-use-llms-in-2026/>

A bit over a year ago I wrote How I use LLMs as a staff engineer . Here's a brief summary of what I used AI for last year: Smart autocomplete with Copilot Short tactical changes in areas I don't know well (always reviewed by a SME) Writing lots of use-once-and-throwaway research code Asking lots of questions to learn about new topics (e.g. the Unity game engine) Last-resort bugfixes, just in case it can figure it out immediately Big-picture proofreading for long-form English communication Here are some tasks I explicitly didn't use AI for last year: Writing whole PRs for me in areas I'm familiar with Writing ADRs or other technical communications Research in large codebases and finding out how things are done February 2025 was a long time ago. Back then the best model was the first reasoning model, OpenAI's o1. Agents sort of worked, but would often get stuck or thrown off by compaction. What's changed since then?

Agents are good now The biggest change is that I now use LLMs to produce entire PRs in areas I'm familiar with . A year ago I would very occasionally ask an agent to make changes to a single file if it was a simple change I couldn't be bothered typing out. Sometimes I wo...

#### 5. DeepSeek-V4-Flash means LLM steering is interesting again

Source: seangoedecke.com RSS feed | Published: Sat, 16 May 2026 00:00:00 GMT

Link: <https://seangoedecke.com/steering-vectors/>

Ever since Golden Gate Claude I've been fascinated with 'steering': the idea that you can guide LLM outputs by directly manipulating the activations of the model mid-flight. DeepSeek V4 Flash I was inspired to write this post by antirez's recent project DwarfStar 4 , which is a version of llama.cpp that's been stripped down to run only DeepSeek-V4-Flash. What's so special about this model? It might be what many engineers have been waiting for: a local model good enough to compete with at least the low end of frontier model agentic coding. Since steering requires a local model, it's now practical for many engineers to try it out for the first time. And indeed, antirez has baked steering into DwarfStar 4 as a first-class citizen. Right now it's very rudimentary (basically just the toy 'verbosity' example you can replicate via prompting), but the initial release was only eight days ago . I plan to follow this project closely. How steering works The basic idea behind steering is extracting a concept (like 'respond tersely') from the model's internal brain state, then reaching in during inference and boosting the numerical activations that form that concept. One way you might do this i...

#### 6. Reddit Is Blocking Some Users From Accessing Its Website From Mobile Devices

Source: Daring Fireball | Published: 2026-05-16T21:22:38Z

Link: <https://arstechnica.com/information-technology/2026/05/why-reddit-blocked-my-daily-visit-to-its-mobile-website/>

Nate Anderson, writing at Ars Technica: But I was surprised this weekend to suddenly find myself cut off; Reddit simply would not let me visit the site on my mobile phone. Instead, a new overlay popped up, saying, 'Get the app to keep using Reddit.' There was no way to skip, bypass, or close the overlay. It did not provide any instructions or alternatives for continuing to use the mobile web version. What it did offer was a large button I could press to get the app. If I did so, the overlay told me, I would be able to 'search better' and 'personalize your feed' - two things I don't care to do. [...] I reached out to the company to ask what was going on. According to a spokesperson, 'We recently started running a test for a small subset of frequent logged-out mobile users that prompts them to download the app after visiting the site. These users are already familiar with Reddit and we've seen that the experience is much better for them in the app. The app offers a more personalized experience and users can more easily find communities that match their interests.' Yes, they're doing this for the users' benefit. Sure. ?

#### 7. Santa Clara County Sues Meta Over Alleged Scam Ads

Source: Daring Fireball | Published: 2026-05-16T21:17:15Z

Link: <https://sanjosespotlight.com/santa-clara-county-sues-meta-over-alleged-scam-ads/>

Brandon Pho, reporting for San Jose Spotlight: The lawsuit filed Monday alleges that instead of cracking down on deceptive ads designed to trick users out of their money, Meta has hamstrung its own fraud prevention teams and helped fake companies bypass its filters to enable the tech powerhouse to enjoy an estimated \$7 billion in ad revenue from the scams every year. [...] The county lawsuit seeks attorney fees and a ruling barring Meta from further alleged violations of false advertising and unfair competition laws. Much of the lawsuit's allegations stem from a 2025 Reuters investigation suggesting Meta was at one point involved in one-third of all successful Internet scams in the U.S. The company has vowed to fight the lawsuit. 'This claim relies on Reuters reporting that distorts our motives and ignores the full range of actions we take to combat scams every day,' a spokesperson for the company told San JosØ Spotlight. 'We aggressively fight scams on and off our platforms because they're not good for us or the people and businesses that rely on our services.' Reuters's Jeff Horwitz and Engen Tham were awarded the Pulitzer Prize for Beat Reporting for their reporting on this sto...

#### 8. ? AI Is Technology, Not a Product

Source: Daring Fireball | Published: 2026-05-16T20:32:51Z

Link: [https://daringfireball.net/2026/05/ai\\_is\\_technology\\_not\\_a\\_product](https://daringfireball.net/2026/05/ai_is_technology_not_a_product)

It's not even a feature. It's just technology.

#### 9. Getting C code navigation even for Debian (or Ubuntu) packages

Source: Chris's Wiki :: blog | Published: 2026-05-16T02:06:07Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/CCodeIntelligenceForDebs>

Every so often, I want (or need) to make modifications to programs in an Ubuntu package, and often the programs are written in C (and these days I'm using dgit to manipulate the package ). One of my challenges when I do this is that I generally don't start out knowing where and how to change the code to do what I want; instead, I have to navigate around an unfamiliar code base and work out enough of its structure to find the specific bit of code I need to change. These days, the dominant way to get smart code navigation and other code knowledge things is through LSP servers and clients . A variety of modern and semi-modern languages have LSP servers that you can immediately use in your editor of choice and then navigate around random code bases with handy features like 'find definition' and 'find references' (for example, Go, Python, and Rust). Unfortunately, C isn't such a language. In the general case, understanding C code requires knowing how it's compiled, and that means you often have to tell C LSP servers this information. Well, specifically you have to tell this stuff to clangd , the dominant LSP server for C and C++. (There's also ccls , which may work out part of this inf...

#### 10. Make ZIP files smaller with ZIP Shrinker

Source: Evan Hahn (dot com) | Published: Sat, 16 May 2026 00:00:00 +0000

Link: <https://evanhahn.com/make-zip-files-smaller-with-zip-shrinker/>

I built ZIP Shrinker, a little browser tool to shrink ZIP files. It also works with formats that are secretly ZIPs underneath, like APK, EPUB, JAR, and many more. Try it out! How does it work? At a high level, this tool (1) re-compresses every file in the ZIP archive with higher compression (2) removes all metadata (3) removes entries for directories. Re-compressing ZIP files are typically compressed with an algorithm called Deflate . There are a few tools that can re-compress Deflate data and make it smaller, usually by spending more time on the computation. I took one of these tools, libdeflate , and applied it to each compressed entry in the ZIP. I chose libdeflate because of its performance; alternatives like Zopfli can achieve marginally smaller results but take much longer. I created libdeflate.js , a WebAssembly wrapper for libdeflate, as part of this work. (I always relish my time working with WASM!) Removing metadata and directories Each entry in a ZIP file can contain additional metadata like comments. These aren't typically used, and if they're there, my shrinker removes them. This usually doesn't save too many bytes, but it doesn't hurt. Removing directories is a sligh...

#### 11. Reading List 05/16/26

Source: Construction Physics | Published: Sat, 16 May 2026 12:01:26 GMT

Link: <https://www.construction-physics.com/p/reading-list-051626>

Tokyo's cheap housing and expensive land, the House response to the Senate housing bill, an IED near an Alabama dam, Fervo's IPO, and more.

#### 12. Pluralistic: Making sense of Trump's unscheduled sudden midair disassembly of the American empire (16 May 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 16 May 2026 08:35:41 +0000

Link: <https://pluralistic.net/2026/05/16/technopoly/>

Today's links Making sense of Trump's unscheduled sudden midair disassembly of the American empire: Don't mistake "powerful" for "durable." Hey look at this: Delights to delectate. Object permanence: Copyrighted law; Viral videos v cops; Crooked banker v tiny bat; "Infested"; Can the means of computation be seized? Upcoming appearances: Berlin, Hay-on-Wye, London, LA, Menlo

Park, NYC, Edinburgh. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Making sense of Trump's unscheduled sudden midair disassembly of the American empire (permalink) For generations, the American empire was the most powerful force on earth, and so we tended to assume that it was the most durable force on earth ? surely anything so powerful must also be eternal? But power and durability aren't the same thing, as Le Guin reminded us with her oft-quoted maxim that "We live in capitalism, its power seems inescapable ? but then, so did the divine right of kings": <https://www.ursulaklequin.com/nbf-medal> Monarchs may be powerful, but that power is derived from a manifestly incorrect belief in spe...

### 13. The Applicability of Spaced Repetition

Source: Fernando Borretti | Published: Sun, 17 May 2026 00:00:00 +0000

Link: <https://borretti.me/article/the-applicability-of-spaced-repetition>

On factual vs. conceptual knowledge.

### 14. Five Minutes of Prime Time

Source: Susam Pal | Published: 2026-05-16T00:00:00Z

Link: <https://susam.net/five-minutes-of-prime-time.html>

Let me share a very silly story from roughly 18 years ago! In 2008, I joined RSA , the network security company named after the initials of the inventors of the RSA algorithm, Rivest, Shamir and Adleman, who were also the founders of RSA, the company. There was a bit of a nerd culture in the workplace where topics like prime numbers, combinatorics, probability theory, etc. were discussed fervently. A prime-number employee number was considered a lucky charm. I had a rather nice large five-digit prime number as my employee number, which I remember being quite pleased about. There were internal forums for almost all kinds of topics. A few I remember fondly were a mathematics forum where colleagues would challenge each other with mathematical puzzles and a similar physics forum where, for some reason, crafting contrived paradoxes using special and general relativity and putting them up for debate was a common activity. The participants would analyse each paradox to determine if it truly was one or if it could be resolved into something that was no longer a paradox. I loved hanging out on those forums and made many friends there. The mathematics forum, in particular, gave me plenty of...

### 15. The mistake of conflating intelligence and power

Source: Dwarkesh Podcast | Published: Sat, 16 May 2026 19:04:16 GMT

Link: <https://www.dwarkesh.com/p/the-mistake-of-conflating-intelligence>

f this is your definition of intelligence is the ability to achieve your goals across a wide variety of domains, then Stalin was the most intelligent person who ever lived.

### 16. Notes on pretraining parallelisms and failed training runs.

Source: Dwarkesh Podcast | Published: Sat, 16 May 2026 19:01:16 GMT

Link: <https://www.dwarkesh.com/p/notes-on-pretraining-parallelisms>

Deeply researched interviews

### 17. RLVR might be disproportionately bad at science

Source: Dwarkesh Podcast | Published: Sat, 16 May 2026 19:00:59 GMT

Link: <https://www.dwarkesh.com/p/rlvr-might-be-disproportionately>

the verification loop for theories can be on the order of decades and centuries, and even then we know today as the better theory can often actually make worse predictions

### 18. The Tomy Tutor and the state of 1983 home computers

Source: Old Vintage Computing Research | Published: 2026-05-16T18:09:46.083-07:00

Link: <https://oldvcr.blogspot.com/2026/05/the-tomy-tutor-and-state-of-1983-home.html>

The Tomy Tutor was my first computer, in late 1983. I was seven and we got it at Federated . I've acquired several more since then, but this is the actual one I used and it still works perfectly. Using a design modeled on the doomed Texas Instruments 99/8, one of several

unreleased successors to the TI 99/4A, the Tomy Tutor and its overseas siblings, the Japanese Pyuuta (????) series, promised an easy kid-friendly introduction to computers with a durable case, nice graphics and sound, games on cartridge, and two, count 'em, two internal dialects of BASIC (one on early systems). It had 16K of RAM, though this was entirely tied up in the 9918A video display processor with only 256 bytes of RAM directly addressable by its 2.7MHz TMS 9995 CPU, and of Tomy's promised peripherals only game controllers and a tape deck were ever offered. Still, despite the bowdlerized operating system and bupkis contemporary expansion options, the Tutor was nevertheless one of the first true 16-bit home computers, and as part of the 1983 low-end home system cavalcade, an inexpensive choice as well. Another thing the Tomy doesn't have much of is history: the story of its development is rather murky, and mo...

#### 19. SQLAlchemy 2 In Practice - Chapter 8: SQLAlchemy and the Web

Source: Miguel Grinberg's Blog | Published: Sat, 16 May 2026 13:35:14 GMT

Link: <https://blog.miguelgrinberg.com/post/sqlalchemy-2-in-practice---chapter-8-sqlalchemy-and-the-web>

This is the eighth and final chapter of my SQLAlchemy 2 in Practice book. If you'd like to support my work, I encourage you to buy this book, either directly from my store or on Amazon . Thank you! Whether you are building a traditional web application, or a web API that works alongside a web front end or smartphone app, SQLAlchemy is one of the best choices to add database support to a Python web server. In this chapter two example integrations with Flask and FastAPI will be demonstrated. These are two of the most popular Python web frameworks and should serve as examples even if you use another web framework.

#### 20. Corona-update Eind Juli

Source: Bert Hubert's writings | Published: Sat, 25 Jul 2020 14:42:06 +0200

Link: <https://berthub.eu/articles/posts/corona-update-eind-juli/>

Toch weer een kleine Corona-update, vrij snel na de vorige van 1 juli. Want poeh, gaat het nou ineens weer mis in Nederland? Of valt het mee? Moeten we allemaal aan de maskers en komt het dan goed? En gaan de vaccins en medicijnen ons redden voor de winter? In tegenstelling tot vele roepstoeters heb ik daar ook allemaal geen zekere antwoorden op, maar hier toch een set updates over de meest recente ontwikkelingen.

#### 21. Multi-Signal GNSS & the Curious Case of Galileo E24: Part 1

Source: Bert Hubert's writings | Published: Sat, 18 Jul 2020 09:07:23 +0200

Link: <https://berthub.eu/articles/posts/the-curious-case-of-galileo-e24/>

GPS, GLONASS, Galileo, BeiDou (Global Navigational Satellite Systems, GNSS) are miracles of engineering. And within any complex system hide surprises and unsuspected complications. In this post I go over how we resolved an odd observation: that everyone (including major institutions) reported that one of the Galileo satellites (E24) appeared to be off by 1.5 meters, while actual positioning equipment (like phones) showed no problems at all. Through the explanation, we'll casually learn a thing or two about the cool concept of multi-frequency GNSS.

#### 22. Review of the Uputronics GPS/RTC Raspberry Pi Expansion Board

Source: Bert Hubert's writings | Published: Thu, 09 Jul 2020 13:38:00 +0200

Link: <https://berthub.eu/articles/posts/uputronics-gps-rtc-expansion-board/>

Hi everyone, Disclaimer: this is a review of a free sample I got sent by Uputronics. I recently spotted on Twitter that Uputronics had launched a GPS/RTS Raspberry Pi Expansion Board, and that they were looking for worthy projects to test it. I thought our Galileo/GPS/BeiDou/GLONASS monitoring project galmon.eu might be worthy and Uputronics agreed. Relevant for our purposes, this is an actual ?timing receiver?. Any GNSS chip will get you a location, but some of them go the extra mile to also deliver a very stable clock signal.