

### 1. Quoting SpaceX S-1

Source: Simon Willison's Weblog | Published: 2026-05-20T22:26:36+00:00

Link: <https://simonwillison.net/2026/May/20/spacex-s1/#atom-everything>

We have the ability to use compute resources to support our proprietary AI applications (such as Grok 5, which is currently being trained at COLOSSUS II), while also providing access to select compute capacity to third-party customers. For example, in May 2026, we entered into Cloud Services Agreements with Anthropic PBC (?Anthropic?), an AI research and development public benefit corporation, with respect to access to compute capacity across COLOSSUS and COLOSSUS II . Pursuant to these agreements, the customer has agreed to pay us \$1.25 billion per month through May 2029, with capacity ramping in May and June 2026 at a reduced fee. The agreements may be terminated by either party upon 90 days? notice. ? SpaceX S-1 , highlights mine Tags: anthropic , grok , generative-ai , ai , llms

### 2. How fast is 10 tokens per second really?

Source: Simon Willison's Weblog | Published: 2026-05-20T17:57:45+00:00

Link: <https://simonwillison.net/2026/May/20/tokens-per-second/#atom-everything>

How fast is 10 tokens per second really? Neat little HTML app by Mike Veerman ( source code here ) which simulates LLM token output speeds from 5/second to 800/second. Useful if you see a model advertised as "30 tokens/second" and want to get a feel for what that actually looks like. Via Hacker News Tags: ai , generative-ai , llms

### 3. Google I/O, Gemini Spark, Antigravity

Source: Simon Willison's Weblog | Published: 2026-05-20T15:32:17+00:00

Link: <https://simonwillison.net/2026/May/20/google-io/#atom-everything>

It's hard to find much to write about Google I/O this year because I have a policy of not writing about anything that I can't try out myself, and a lot of the big announcements are "coming soon". I actually prefer to write about things that are in general availability, because I've had instances in the past where the previews didn't match what was released to the general public later on. Aside from Gemini 3.5 Flash the most interesting announcement looks to be Google's upcoming OpenClaw competitor Gemini Spark , described as "your personal AI agent" which can "connect natively with your favorite Google apps like Gmail, Calendar, Drive, Docs, Sheets, Slides, YouTube, and Google Maps". The FAQ for that also includes this confusing detail: What Gemini model does Gemini Spark run on? Gemini Spark runs on Gemini 3.5 Flash and Antigravity. The antigravity.google website currently lists Antigravity as a desktop app, a CLI agent tool (written in Go), the Antigravity SDK (an open source Python wrapper around a bundled closed source Go binary), and the original Antigravity IDE (a VS Code fork). I guess Gemini Spark, the user-facing hosted agent product, might be running on that Go binary, b...

### 4. Prompts are technical debt too

Source: seangoedecke.com RSS feed | Published: Wed, 20 May 2026 00:00:00 GMT

Link: <https://seangoedecke.com/prompts-are-technical-debt-too/>

It?s common and correct to say that ?all code is technical debt?. Adding code is a necessary evil for developing new features: you almost always have to do it, but each line of code adds to the complexity and maintenance burden of the system. All future changes to the system have to work with the existing code, or at least avoid breaking it. Once systems accumulate enough code, they become impossible for a single person to understand: instead of reading the code and understanding what it does, you must rely on guesses, theories and heuristics 1 . Sensible engineers write as little code as possible. They write a lot of prompts, though! Many large projects now have a set of codebase-specific prompt files: AGENTS.md, CLAUDE.md, those same files in sub-directories, and skills . If you?re building a program that uses AI 2 , you?ll have separate prompts for capabilities and for each tool , as well as a whole set of system prompts . Prompts are important. Minor tweaks to a LLM?s prompt can unlock significant

performance improvements. If the same model feels different across Codex, Cursor, OpenCode, and Copilot, it's almost certainly due to subtle differences in prompting. AI companies sp...

#### 5. WSJ: ?Google Unveils New Gemini AI Agent for Personal Tasks?

Source: Daring Fireball | Published: 2026-05-21T01:05:14Z

Link: <https://www.wsj.com/tech/ai/google-unveils-new-gemini-ai-agent-for-personal-tasks-b8093197?st=BFmPev>

Katherine Blunt and Rolfe Winkler, reporting for The Wall Street Journal from Google I/O (gift link): Google is supercharging its Gemini artificial-intelligence model to become more competitive in the era of agentic AI. The company has started rolling out what it calls Gemini Spark, a personal agent it says is capable of navigating a user's digital life and acting on his or her behalf. The agent will work across many of Google's products and run on the company's cloud infrastructure. [...] The company has been testing Spark with a limited number of users and plans to make it available next week to those who pay for AI Ultra, a new subscription tier that costs \$100 a month. A different top-level takeaway than the NYT's , which in turn was different from Bloomberg's . Ben Thompson, in a subscriber-only update at Stratechery , sums it up: Indeed, if you wanted a positive spin on Google's plethora of announcements, it's that the company is clearly fully committed to putting AI into anything and everything; if you want to put a negative spin, well, it's the exact same thing. One of the enduring critiques of Google is that the company is unfocused and unmanageable, which, to the extent...

#### 6. NYT: ?Powered by A.I., Google Changes Its Search Box for the First Time in 25 Years?

Source: Daring Fireball | Published: 2026-05-20T21:20:23Z

Link: [https://www.nytimes.com/2026/05/19/business/google-search-bar-ai-gemini.html?unlocked\\_article\\_code=1.jlA.95yh.ptfBUHf-rBtB&smid=url-share](https://www.nytimes.com/2026/05/19/business/google-search-bar-ai-gemini.html?unlocked_article_code=1.jlA.95yh.ptfBUHf-rBtB&smid=url-share)

Tripp Mickle, Kate Conger, and Brian X. Chen, opening The New York Times's report on yesterday's Google I/O keynote (gift link): For 25 years, Google's iconic search box was a long, slender bar where people typed in keywords like ?World Cup.? But over the past three years, artificial intelligence allowed people to type in longer, more complex questions like ?Who are the top 24 teams in the World Cup and what chance does the United States have of advancing?? On Tuesday, Google said the A.I. shift had inspired it to overhaul the dimensions of its search bar for the first time since 2001. The box is getting bigger and more interactive so that people can ask even longer questions and upload photographs and videos into queries. Interesting to me that this is the Times's biggest takeaway. But it speaks to how unchanged the google.com homepage has been since its earliest days . In addition, people can ask follow-up questions with a chatbot on Google's main search page. The company will also offer digital assistants, known as agents, to automate searches so that someone who may be apartment hunting can be notified of a new listing without opening a real estate site like Zillow. Odd, to me...

#### 7. ?You Do Not Need Fancy Equipment, You Do Not Need a Degree, to Make Money and to Do This as Your Job?

Source: Daring Fireball | Published: 2026-05-20T20:53:22Z

Link: <https://www.tiktok.com/@brye.shhh/video/7641047549758934285>

22-year-old pop singer-songwriter Brye , on TikTok: ?Lemons?, my biggest song ever, that went like super viral during quarantine back in 2020, was actually produced, if you can believe it, in GarageBand on my school iPad. My high school gave us all iPads and I produced ?Lemons? on there. I used to just like make beats on GarageBand in high school. I wrote musicals for my school with GarageBand on my iPad. And then I made that little demo for ?Lemons?, recorded it ... on my iPad ... with my horrible little plug-in mic, posted it to spite a guy who was being horrible to me, and it blew up. I love this, to say, how crazy is it that a song that could be on Sirius XM radio ? streamed a hundred million times, literally charted on the global top like viral 50 or whatever ? it was literally made on GarageBand. You do not need fancy equipment, you do not need a degree, to make money and to do this as your job. Obviously it's good to learn. It's fun to upgrade. But if you are working on a budget, GarageBand's free on any Apple

device. If Brye's story isn't exactly what Steve Jobs was talking about when he introduced GarageBand in 2004 and GarageBand for iPad in 2011 , well, I don't know wha...

8. What is the history of the ERROR\_ARENA\_TRASHED error code?

Source: The Old New Thing | Published: Tue, 19 May 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260519-00/?p=112339>

The storage control blocks were destroyed. The post What is the history of the ERROR\_ARENA\_TRASHED error code? appeared first on The Old New Thing .

9. "No way to prevent this" say users of only language where this regularly happens

Source: Xe Iaso's blog | Published: Wed, 20 May 2026 00:00:00 GMT

Link: <https://xeiaso.net/shitposts/no-way-to-prevent-this/CVE-2026-45584/>

In the hours following the release of CVE-2026-45584 for the project Microsoft Windows , site reliability workers and systems administrators scrambled to desperately rebuild and patch all their systems to fix a memory safety vulnerability resulting in arbitrary code execution inside the virus scanner Windows Defender. This is due to the affected components being written in C++, the only programming language where these vulnerabilities regularly happen. "This was a terrible tragedy, but sometimes these things just happen and there's nothing anyone can do to stop them," said programmer Dr. Annabelle Connelly, echoing statements expressed by hundreds of thousands of programmers who use the only language where 90% of the world's memory safety vulnerabilities have occurred in the last 50 years, and whose projects are 20 times more likely to have security vulnerabilities. "It's a shame, but what can we do? There really isn't anything we can do to prevent memory safety vulnerabilities from happening if the programmer doesn't want to write their code in a robust manner." At press time, users of the only programming language in the world where these vulnerabilities regularly happen once or...

10. Notes about reading messages with the Python email packages

Source: Chris's Wiki :: blog | Published: 2026-05-20T03:02:04Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/python/EmailPackagesNotes>

I have a long standing personal program to display MIME formatted email messages in the terminal in a sensible way (it was mentioned in this old entry on my email tools and its comments ). For a long time this was a Python 2 program, using the Python 2 version of the email package . Recently, I moved this program to Python 3 as part of my sudden enthusiasm for Python 3 conversions , using the Python 3 version of email and its sub-packages. In the process I have wound up with some notes and opinions on practical use of the Python 3 email packages. (The Python 2 version of email had its own quirks and oddities, but I worked all of those out that hard way years ago, have mostly forgotten them since, and they're not interesting any more now that the era of Python 2 is over.) The Python 3 email documentation will tell you that the modern interface for email messages is email.message.EmailMessage . The older email.message.Message is (theoretically) only there for Python 3.2 compatibility and you should ignore its methods and use only the EmailMessage methods. This is not entirely the case. If you look behind the curtain, you'll discover that many of the EmailMessage APIs for reading mes...

11. Read Cindy Cohn's new book, Privacy's Defender: My Thirty-Year Fight Against Digital Surveillance

Source: micahflee | Published: Wed, 20 May 2026 22:36:13 GMT

Link: <https://micahflee.com/read-cindy-cohns-new-book-privacys-defender-my-thirty-year-fight-against-digital-surveillance/>

I just finished reading Privacy's Defender: My Thirty-Year Fight Against Digital Surveillance , by my friend Cindy Cohn. It's excellent and you should buy a copy. Cindy is the Executive Director of the Electronic Frontier Foundation, and this memoir/legal history goes over three major legal battles

12. Nobody is destined for greatness.

Source: Westenberg. | Published: Thu, 21 May 2026 01:31:20 GMT

Link: <https://www.joanwestenberg.com/nobody-is-destined-for-greatness/>

Demosthenes lost his first appearance before the Athenian assembly. His voice came out thin and failed him mid-sentence, and the crowd laughed him off the platform. Plutarch tells us he walked home with his cloak pulled over his face, certain his public life had ended before it started. What he

### 13. [RSS Club] Let's meet up AFK

Source: Terence Eden's Blog | Published: Wed, 20 May 2026 11:34:04 +0000

Link: <https://shkspr.mobi/blog/2026/05/rss-club-lets-meet-up-afk/>

Shhhh! This post is only available to RSS subscribers like you ? My wife and I are preparing for a big Interrail journey through Europe. Whenever we go on holiday, we like to meet up with friendly locals to have a drink and chat. We did this on our last journey and it was great. So, if you're a member of RSS club and fancy showing some tourists a cool bar, awesome restaurant (with vegan op?

### 14. Assumptions weaken properties

Source: Computer Things | Published: Wed, 20 May 2026 15:13:16 +0000

Link: <https://buttondown.com/hillelwayne/archive/assumptions-weaken-properties/>

In some tests are stronger than others , I defined  $\text{STRONG} \Rightarrow \text{WEAK}$  to mean "any system passing test  $\text{STRONG}$  is also guaranteed to pass  $\text{WEAK}$ ". This uses the logical implication operator, defined as  $P \Rightarrow Q = !P \vee (P \wedge Q)$  . Implication may be the most overworked operator in logic. Among other things, it's also used in formal specification, where  $\text{Spec} \Rightarrow \text{Prop}$  means "any system satisfying  $\text{Spec}$  has property  $\text{Prop}$  " and  $\text{ASSUME} \Rightarrow \text{Spec}$  means "The assumption  $\text{ASSUME}$  must hold in order for the system to satisfy  $\text{Spec}$  ." Now let's mush these all together and do some math. To start, "the system has property  $\text{Prop}$  " is the same as "the system passes the test that checks  $\text{Prop}$  ", so test strength is also property strength. Now let "  $\text{ASSUME} \Rightarrow \text{Prop}$  " mean "the system passes  $\text{Prop}$  assuming  $\text{ASSUME}$  is true." In classic logic, if  $P$  is true, then obviously  $!Q \vee P$  is true. Further, that is equivalent (just draw the truth table!) to  $!Q \vee (P \wedge Q)$  . So for any propositions  $P$  and  $Q$  ,  $P \Rightarrow (Q \Rightarrow P)$  . In other words,  $\text{Prop} \Rightarrow (\text{ASSUME} \Rightarrow \text{Prop})$  . In other other words, "the system passes  $\text{Prop}$  " is a stronger property than "the system passes  $\text{Prop}$  whenever our assumptions hold." In other other other words, any assum...

### 15. What will better AI mean?

Source: the singularity is nearer | Published: 2026-05-20T00:00:00-07:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/05/20/what-will-better-mean.html>

I thought about posting this paper but rebranding it as the Claude Mythos technical report. As far as I can tell, there's no secret tricks the US frontier labs have, and that basically describes how Mythos was trained. What's in that paper just works, and for verifiable domains, it's only a matter of fixing bugs and scaling up. That's why Anthropic is so desperate for regulatory capture, AI has no moat. AI (and any form of search) has this property where you spend exponentially more money to get linear returns. So for a bit we'll live in an era where AI can in theory solve very hard problems, but it's very expensive to do so. The Internet has been fully mined, and it yielded 20T good tokens. For a Chinchilla optimal model, that's only 1T weights (1e26 training run if dense). 500 GB gets you all of human knowledge in a simple to query archive. For comparison, Wikipedia is 24 GB with mediocre compression. Technology proceeds in terms of S-curves, and AI has gone through a few of them already. I know I'm quite late to this, but I'm feeling optimistic that scaling will mostly stop yielding results. GPT 5.5 is to a point where it's really hard for me to stump it with any problem. What...

### 16. The Fourth Era of Computing

Source: Daniel De Laney | Published: Tue, 19 May 2026 00:00:00 +0000

Link: <https://danieldelaney.net/fourth-era/>

The polarity of hard and easy in computing is inverting. Raw configuration files, exposed APIs, terminal-first applications; these have all become easy. They are the hooks agents want, and we don't have to think about them. Polished GUIs are now becoming the hard path by comparison,

because we have to use them ourselves. This change has sparked the fourth era in the evolution of computing. By describing it and how we got here, I will overturn most of the things I've ever said about user interface design. I've always advocated for simple GUIs that normal people can use, but that argument is losing its potency. Ugly exposed wiring is exactly what makes something the right choice now. Era I: Terminal In the first era of computing, everything was hard. You memorized a specialized language and you typed it at a prompt. If you hadn't read the fine manual, the computer was entirely opaque to you. ' prompt that waits for a command and offers no hint about what to type" class="bordered" loading="lazy" /> Era II: Early GUI In the second era, we defended ourselves from the inhumanity of the CLI with the GUI. Functionality was behind discoverable buttons rather than typed commands. This was p...

17. Corona: Hoe en waar krijgen we het? Niemand die het weet.

Source: Bert Hubert's writings | Published: Mon, 27 Apr 2020 08:49:38 +0200

Link: <https://berthub.eu/articles/posts/corona-hoe-en-wat/>

Ik blog al maanden over SARS-CoV-2 aka COVID-19 aka 'het coronavirus', waarbij ik grotendeels dingen uitleg en verwijst naar positieve ontwikkelingen. Ook zijn er stukjes over medische experimenten, over ons immuunsysteem en recent iets korts over het vermoedelijke belang van goede ventilatie om COVID-19 besmettingen tegen te gaan. Ik houd van de wetenschap en ben trots op het mooie onderzoek wat hier in internationale samenwerking verricht wordt. Maar nu dan toch een zeurstuk: waarom weten we niet waar en hoe mensen in Nederland besmet raken?

18. Corona Update 21 April

Source: Bert Hubert's writings | Published: Tue, 21 Apr 2020 14:30:40 +0200

Link: <https://berthub.eu/articles/posts/corona-update-21-april/>

Ja, daar zijn we weer. Bij het beïndigen van de serie dagelijkse Corona updates had ik aangekondigd nog terug te komen als er weer wat te melden was. Bij deze een aflevering 'hoe gaat het nu met?', met daarin nieuws over hoe het verder gaat met eerder gemelde ontwikkelingen. Daarnaast twee stukjes op veler verzoek: hoe gaat het met het testen en kan de economie al weer open. Om te beginnen, op 25 maart zei ik 'De samenvatting is dat als iedereen zich aan de regels houdt (BLIJF THUIS, HOUD AFSTAND) we het waarschijnlijk net binnen de capaciteit gaan houden in Nederland?', en dat bleek dus te kloppen.

19. Tracing App Backend Development

Source: Bert Hubert's writings | Published: Tue, 14 Apr 2020 20:00:22 +0200

Link: <https://berthub.eu/articles/posts/tracing-app-backend/>

bert@hubertnet.nl / @bert\_hu\_bert / LinkedIn Very quick post. The Dutch government just launched a tender for 'digital technologies that can help with Corona'. The tender includes words on privacy sensitive contact tracing, it explicitly refers to PEPP-PT, the Pan-European Privacy-Preserving Proximity Tracing project. Responses are due on Tuesday the 14th of April, 12:00 (noon). The question/answer form asks if you can do a demo on the 18th of April and have something that could be rolled out on the 28th.