

1. Notes on Pope Leo XIV's encyclical on AI

Source: Simon Willison's Weblog | Published: 2026-05-25T23:58:17+00:00

Link: <https://simonwillison.net/2026/May/25/encyclical-on-ai/#atom-everything>

Dropped this morning by the Vatican: Magnifica Humanitas of His Holiness Pope Leo XIV on Safeguarding the Human Person in the Time of Artificial Intelligence . This is a very interesting document. It's some of the clearest writing I've seen on the ethics of integrating AI into modern society. Pope Leo XIV chose the name Leo in honor of Pope Leo XIII, who is known for his 1891 *Rerum novarum* encyclical on "Rights and Duties of Capital and Labor". This story on Vatican News further clarifies the significance of that decision: Meeting with the College of Cardinals for their first formal encounter after his election, Pope Leo XIV explained part of the reason for the choice of his papal name. "There are different reasons for this," he said, before going on to explain that he chose the name Leo "mainly because Pope Leo XIII, in his historic encyclical *Rerum novarum* addressed the social question in the context of the first great industrial revolution." "In our own day," he continued, "the Church offers to everyone the treasury of her social teaching in response to another industrial revolution and to developments in the field of artificial intelligence that pose new challenges for the def...

2. Quoting Armin Ronacher

Source: Simon Willison's Weblog | Published: 2026-05-24T18:46:53+00:00

Link: <https://simonwillison.net/2026/May/24/armin-ronacher/#atom-everything>

The most frustrating failure mode right now is that people submit issues that are not in their own voice. They contain an observed problem somewhere, but it has been thrown into a clanker and the clanker reworded it and made a huge mess of it. Typically, it was prompted so badly that the conclusions produced are more often than not inaccurate but always full of confidence. The result is complete guesswork on root causes, fake-minimal repros, suggested implementation strategies, analogies to adjacent but often the wrong code, and long lists of error classes that might or might not matter. [...] So at least personally, I increasingly want issue reports to be condensed to what the human actually observed: I ran this command. I expected this to happen. This happened instead. Here is the exact error or log. ? Armin Ronacher , on slop issues filed against Pi Tags: ai , github-issues , llms , ai-ethics , open-source , coding-agents , generative-ai , armin-ronacher , pi , slop

3. Mad House ? Usborne Creepy Computer Games

Source: Simon Willison's Weblog | Published: 2026-05-24T17:14:11+00:00

Link: <https://simonwillison.net/2026/May/24/usborne-mad-house/#atom-everything>

Tool: Mad House ? Usborne Creepy Computer Games Via Hacker News I learned that UK publisher Usborne published free PDFs of their 1980s Computer Books , some of which I remember working through on my Commodore 64 as a child. These were so great! Beautifully illustrated books with fun projects made up of code you could type into your own machine. I remember playing "Mad House" typed in from the 1983 book "Creepy Computer Games", so I fed that PDF into Claude and had it build an interactive version of that game in JavaScript and HTML: Build a vanilla JS artifact that exactly recreates the game Mad House from this book, make sure it's mobile friendly and has a suitable retro aesthetic Credit the book title and link to <https://usborne.com/us/books/computer-and-coding-books> " and " Tags: computer-history , games , tools

4. [Sponsor] exe.dev

Source: Daring Fireball | Published: 2026-05-25T23:57:24Z

Link: <https://exe.dev/?df>

A cloud for the agent era. Use exe.dev to get a pool of VMs with SSH, root, and web auth by default. Secrets injected at the network edge stay out of the LLM's hands. Persistent servers, internal tools, vibe coding, disposable devboxes, whatever. It's just a computer. ?

5. Awarding Jay Haynes His Being Right Points for Predicting Apple Hitting \$3 Trillion in Market Cap

Source: Daring Fireball | Published: 2026-05-25T20:32:29Z

Link: <https://daringfireball.net/linked/2014/01/29/haynes-aapl>

Here's a fun one. Back in 2014 I linked to a post by Jay Haynes in which he projected that with a very reasonable level of annual growth, Apple ought to reach a \$3 trillion market cap within 10 years. At the time of his writing, Apple's market cap was "just" \$450 billion, and no company had hit the \$1 trillion market. So projecting a \$3 trillion valuation in 10 years was a bold prediction. Apple hit \$3 trillion in just 8 years. Haynes's original blog went belly-up, alas, but he republished the piece on Medium, with a bit of additional commentary up front, in 2016. Re-reading Haynes's piece today, it holds up extremely well, including his case that the iPhone and iPad are almost textbook examples of Clayton Christensen's disruption theory (yet Christensen himself got it wrong). (Thanks to Nathan Peretic, longtime DF reader and owner of a perfect personal homepage, for prompting me to revisit this and award Haynes his well-earned Being Right Points.) ?

6. Thieves Are Texting Threats to Victims of iPhone Theft in London

Source: Daring Fireball | Published: 2026-05-25T19:23:32Z

Link: https://www.nytimes.com/2026/05/23/world/europe/phone-theft-threats-london.html?unlocked_article_code=1.IFA.OUt7.VJ_FoDpINr0L

Lizzie Dearden and Amelia Nierenberg, reporting for The New York Times (gift link): The crime Alex Pikula reported to the police was one they had heard before: An e-bike rider had zoomed past as Mr. Pikula left a theater in London's West End, ripping his phone from his hands. It was frustrating, Mr. Pikula thought, but that was that. He was wrong. His mother soon started receiving strange texts, claiming to have her son's emails and bank information. Then she received a video of a man brandishing a gun. Then came threats of sexual assault and death. "I know who you are and where you live," read one, full of obscenities and typos. "I've killed or [sic] far less than a phone before," it went on. "We will see if you value your life over this phone." All of the messages wanted her to do one thing: unlink her son's Apple ID from his stolen phone. The story only mentions the word iPhone twice, but phone appears over 30 times. "Apple ID" appears four times. There's zero mention of Android or Google. It's just implicitly assumed that the only phones worth stealing or threatening victims about are iPhones. The story makes no mention of Apple's Stolen Device Protection, which Apple recen...

7. Netherlands Seizes 800 Servers, Arrests 2 for Aiding Cyberattacks

Source: Krebs on Security | Published: Mon, 25 May 2026 13:21:49 +0000

Link: <https://krebsonsecurity.com/2026/05/netherlands-seizes-800-servers-arrests-2-for-aiding-cyberattacks/>

Authorities in the Netherlands have arrested the co-owners of two related Internet hosting companies for operating IT infrastructure used by Russia to carry out cyberattacks, influence operations and disinformation campaigns inside the European Union. The two men were the focus of a 2025 KrebsOnSecurity story about how their hosting companies had assumed control over the technical infrastructure of Stark Industries Solutions, an Internet service provider sanctioned last year by the EU as a frequent staging ground for cyber mischief from Russia's intelligence agencies.

8. Anti-robot techniques can be nice but the problem is, they're not static

Source: Chris's Wiki :: blog | Published: 2026-05-25T17:14:15Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/web/RobotsVsLowAttentionWebsites>

I've recently come up with what I expect would be a quite good anti-robot, anti-crawler tactic, which I will give the snappy label and summary of "robots don't POST". Simply require a HTTP cookie to see your web pages and then if visitors don't have the cookie, put up an interstitial page with a HTML form that requires them to POST it to get the cookie. All the form need is a "click me to get your entrance cookie", because right now, few or no robots or crawlers will

make that HTTP POST request; they only do HTTP GETs. To distract bad crawlers you might need some other links on the interstitial page, optionally going to content tarpits. (If you're going to do this in practice you'll want to exempt syndication feed requests and perhaps requests from bingbot, Googlebot, and so on. Although maybe not Googlebot any more .) The obvious problem with this technique is that if people start doing it in any quantity, the "robots don't POST" thing won't last. Bad crawlers will start hitting POST endpoints for forms that just have a "click me" button, and then POST endpoints for forms that have an "I am human" tick box to mark or a field to fill in or whatever the elaboration people come with...

9. GitHub Actions security in Python packages

Source: Andrew Nesbitt | Published: 2026-05-25T10:00:00+00:00

Link: <https://nesbitt.io/2026/05/25/github-actions-security-in-python-packages.html>

This is a written version of a talk I gave at PyCon US 2026 in Long Beach. Slides (PDF) , scripts, and datasets are at github.com/andrew/pycon . Of the roughly 864,000 packages PyPI lists, about 387,000 declare a repository URL on GitHub, mapping to 343,000 distinct repositories once you collapse the monorepos. 152,000 of those have something in `.github/workflows/` , and for practical purposes open source Python has one CI system: Travis CI, the previous default, accounts for 11% of the same population and stopped offering free open source builds in 2023, with everything else below 2%. Around 56,000 of those repositories reference `pypa/gh-action-pypi-publish` somewhere in a workflow file, which is to say they tag a commit, a runner spins up, a wheel gets built, and twine upload puts it on PyPI. About 22% of them have migrated to trusted publishing , where PyPI accepts a short-lived OIDC token minted by the workflow instead of a stored API key, and the publisher configuration on PyPI's side names the repository, the workflow file, and optionally a deployment environment. The other 44,000 or so still have a `PYPI_API_TOKEN` sitting in repository secrets. I think trusted publishing is on...

10. PHP - simple way to send HTTP headers before a script ends

Source: Terence Eden's Blog | Published: Mon, 25 May 2026 11:34:38 +0000

Link:

<https://shkspr.mobi/blog/2026/05/php-simple-way-to-send-http-headers-before-a-script-ends/>

Suppose you want PHP to keep processing after it has sent back an HTTP response. Normally, this doesn't work: `<?php header( "Location: https://example.com/" ); // Long operation. sleep(10); die();` Try it yourself. You'll have to wait 10 seconds before you get back `< HTTP/2 302 < location: https://example.com/` There are some complex ways to fix this - they usually involve?

11. Pluralistic: No honor among (ad-tech) thieves (25 May 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 25 May 2026 08:21:45 +0000

Link: <https://pluralistic.net/2026/05/25/lying-spies/>

Today's links No honor among (ad-tech) thieves: Including "and" and "the." Hey look at this: Delights to delectate. Object permanence: Budweiser nunchuks; GOP vote-suppressor voted illegally; Airbnb enshittifies; Oculus enshittifies; Nintendo copyfrauds its fans; Meritocracy to eugenics pipeline; Ultima Online crisis management; SNES cartridge urinal; JJ Abrams x Axanar, "Sex Criminals"; Beating school filters for fun; Orphan works; Japanese ATM heist; How the Sacklers rigged the game. Upcoming appearances: London, Kansas City, LA, Menlo Park, Toronto, NYC, Edinburgh. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. No honor among (ad-tech) thieves (permalink) It shouldn't come as a surprise to learn that a company that uses dishonest tactics to spy on you for profit will also use dishonest tactics to sell the resulting surveillance data. The only reason this wouldn't be obvious is if you've fallen into the trap of thinking "if you're not paying for the product, you're the product." Companies that cheat when the opportunity arises will cheat everyone: customers...

12. Games Are the Art Form of Our Time

Source: Dr Paris Buttfield-Addison | Published: Sun, 24 May 2026 00:00:00 +0000

Link: <https://hey.paris/posts/ncp-submission/>

Submissions for Australia's next National Cultural Policy closed in May 2026. I wrote one, on behalf of Tasmanian Game Makers , because I want games taken seriously as what they are: the art form most of us actually play. Read the full submission on Tas Game Makers Here are the highlights. Games are the art form of our time. In 2025, 82% of Australians played them, the average player was 35, and women now make up just over half of all players. On the government's own data, digital games development has been the fastest-growing of all cultural and creative activities in the country over the past decade. Games are also one of our most export-heavy creative industries: 93% of the Australian sector's revenue comes from overseas.

13. Leo wouldn't stand still

Source: Dr Paris Buttfield-Addison | Published: Thu, 21 May 2026 00:00:00 +0000

Link: <https://hey.paris/posts/leo-sprite-alignment/>

We're working on an adventure game, Leonardo's Moon Ship . The player character, Leo, has the usual pile of animations: idle left and right, walk in four directions, a couple of stair anims. Since implementing the basics of the character, Leo has had this slightly cursed behaviour where he'd grow about 3% taller the moment he started walking, then shrink back when he stopped. His feet would also lift off the ground by up to 27 pixels depending on which way he was facing. Not enough to be obviously broken in a screenshot. Exactly enough to make the game feel wrong.

14. Welcoming the Bhutanese Government to Have I Been Pwned

Source: Troy Hunt | Published: Mon, 25 May 2026 22:52:15 GMT

Link: <https://www.troyhunt.com/welcoming-the-bhutanese-government-to-have-i-been-pwned/>

Today, we welcome the 45th government onboarded to Have I Been Pwned's free gov service: Bhutan. The Bhutan Computer Incident Response Team, BtCIRT, now has access to monitor Bhutanese government domains against the data in HIBP. As Bhutan's national CIRT, BtCIRT is responsible for consuming threat

15. FediMeteo, timezones, and the art of not breaking what already works

Source: IT Notes | Published: Mon, 25 May 2026 09:14:00 +0000

Link: <https://it-notes.dragas.net/2026/05/25/fedimeteo-timezones-and-the-art-of-not-breaking-wh-at-already-works/>

I have already written about how FediMeteo was born , and about how HAProxy helps reduce the number of requests that reach snac . Seen from the outside, FediMeteo almost seems still. There is a static homepage, regenerated every hour. There are the city pages, with their forecasts. There are RSS feeds waiting to be fetched, JSON objects waiting to be requested, Fediverse instances refreshing data, subscribing, unsubscribing, retrieving profiles, and reading notes. That is the visible part. Behind it, however, FediMeteo is much more than a homepage, a few ActivityPub accounts, and a well-behaved reverse proxy. It is a chain of small pieces, in proper Unix style, each trying to do one thing and do it as well as possible. That chain, although almost invisible from the outside, was not born already tidy. It changed, was rewritten, adapted to new countries, timezones, ambiguous city names, external service limits, and also to my own mistakes. Some mistakes were small. Others were much less so. Because FediMeteo is a human project and, as such, imperfect. Imperfect in the way humans are imperfect, which today almost seems unfashionable. I like that. The first version of the bot was almo...

16. Distributing LLM inference in DwarfStar

Source: antirez.com | Published: Mon, 25 May 2026 16:54:59 +0200

Link: <http://antirez.com/news/167>

High end NVIDIA cards, and the server and power needed to run them, cost a lot of money, especially if you plan to reach enough VRAM to run massive models. The alternative, so far, has been Apple hardware, or the DGX Spark that, even if severely limited because of memory

bandwidth, still allows to run LLMs prompt processing (prefill) fast enough. The Mac Studio provided up to 512GB unified memory, a solution with modest memory bandwidth (but much better than the Spark) and compute at a price that was, after all, given the current situation, relatively fair. For instance, with DwarfStar the Mac Studio M3 Ultra 512GB can run DeepSeek v4 PRO at 150 t/s prefill and ~10-13 t/s decoding, not great but at a level that is usable for certain use cases. Even 2-bit quantized, DeepSeek v4 PRO resists very well, like Flash at the same quantization (today I made PRO write a C compiler, I'll publish the video soon). I would not consider a trivial fact to run a frontier model at home, with a ~12k total spending. One could expect this to get better and better, but the situation at the horizon appears cloudy. There is almost zero hope that NVIDIA setups will get less expensive, and even a small com...

17. Solving the board game Quoridor

Source: Grant Slatton's Blog | Published: Mon, 25 May 2026 20:01:38 +0000

Link: <https://grantslatton.com/solving-quoridor>

Algorithmic and optimization playtime

18. The Galmon GNSS Monitoring Project

Source: Bert Hubert's writings | Published: Wed, 08 Jan 2020 13:53:08 +0100

Link: <https://berthub.eu/articles/posts/galmon-project/>

Welcome to the Galmon project! Our goals are to monitor the major Global Navigation Satellite Systems, including GPS, GLONASS, Galileo and BeiDou, but also Space Base Augmentation Systems like WAAS, EGNOS and GAGAN. Galmon is an open source & open data project with a community of over 30 station operators running more than 50 receivers. Some history of the project may be found in the second part of this blog post, and on this page on the big Galileo 2019 outage.

19. A 2019 update on PowerDNS, Galileo, DoH and other things

Source: Bert Hubert's writings | Published: Wed, 01 Jan 2020 15:37:52 +0100

Link: <https://berthub.eu/articles/posts/update-2019-powerdns-galileo-ripe-doh/>

Last year I wrote a post listing what I had been doing in 2018, because it turned out it was a lot, so much that it was useful to summarise it somewhat. This year there is less to report, but people have wondered what I've been up to. And I've learned that if you leave room for speculation about your intentions, wrong answers tend to come up. This (professional) year for me mostly consisted of two things: DNS over HTTPS and Galileo.

20. What Is Life?

Source: Bert Hubert's writings | Published: Mon, 30 Dec 2019 14:44:26 +0100

Link: <https://berthub.eu/articles/posts/what-is-life/>

What is life? This question keeps many people awake at night and has led to rafts of definitions, some involving features (procreation, metabolism and so forth), some involving chemistry, entropy or energy flux and some are of a more philosophical bent. A BBC article noted there are over 100 definitions of life and it claims all are wrong. And indeed it does appear to be quite a challenge - every hard and fast rule is violated somewhere in biology.