

1. How we contain Claude across products

Source: Simon Willison's Weblog | Published: 2026-05-30T21:36:24+00:00

Link: <https://simonwillison.net/2026/May/30/how-we-contain-claude/#atom-everything>

How we contain Claude across products A complaint I often have about sandboxing products is that they are rarely thoroughly documented, and in the absence of detailed documentation it's hard to know how much I can trust them. Anthropic just published a fantastic overview of how their various sandbox techniques work across Claude.ai, Claude Code, and Cowork. We constrain where and how an agent can act with process sandboxes, VMs, filesystem boundaries, and egress controls. The goal is to set a hard boundary on what an agent can reach. For example, if credentials never enter the sandbox, they can't be exfiltrated, regardless of whether the cause is a user, a model finding a "creative" path, or an attacker. Claude.ai uses gVisor. Claude Code, run locally, uses Seatbelt on macOS and Bubblewrap on Linux. Claude Cowork runs a full VM (Apple's Virtualization framework on macOS, HCS on Windows). There's a lot in here, including some interesting stories of risks they missed such as the api.anthropic.com/v1/files exfiltration vector covered here previously. This reminded me it's time I took another look at Anthropic's open source srt (Anthropic Sandbox Runtime) tool - it's mature enough...

2. Running Python ASGI apps in the browser via Pyodide + a service worker

Source: Simon Willison's Weblog | Published: 2026-05-30T21:02:16+00:00

Link: <https://simonwillison.net/2026/May/30/pyodide-asgi-browser/#atom-everything>

Research: Running Python ASGI apps in the browser via Pyodide + a service worker Datasette Lite is my version of Datasette that runs entirely in the browser using Pyodide in WebAssembly. When I first built it four years ago I used Web Workers and code that intercepts navigation operations and fetches the generated HTML by running the Python app. This worked, but had the disadvantage that any JavaScript in tags would not be executed - breaking some Datasette functionality and a whole lot of Datasette plugins. This morning I set Claude Opus 4.8 the task (in Claude Code for web) of figuring out how to run Python ASGI apps in Pyodide using Service Workers instead, and it seems to work! Here's a basic ASGI FastCGI demo and here's a demo that runs Datasette 1.0a31. I'm still getting my head around exactly how it works, but once I've done that I plan to upgrade Datasette Lite itself. Tags: javascript, python, datasette, asgi, service-workers, pyodide, datasette-lite, claude-code

3. I Am Retiring from Tech to Live Offline

Source: Simon Willison's Weblog | Published: 2026-05-30T19:39:08+00:00

Link: <https://simonwillison.net/2026/May/30/retiring-from-tech-to-live-offline/#atom-everything>

I Am Retiring from Tech to Live Offline I've seen a lot of posts on forums from people threatening to quit their careers over AI. This is not one of those: Chad Whitacre is taking concrete steps, starting with this typewritten, scanned letter I'm retiring from tech. Well, "retiring" is euphemistic. I'm stepping away from tech, and that includes Open Source. [...] AI was the last straw. Have you heard of that island off India where the indigenous population kills any outsiders fool-hardy enough to land? They are doing the rest of us a favor by preserving a way of life we may need again someday, or at the very least should not want to see completely extinguished. A reminder. Never forget your roots. Here in Pennsylvania we have the Amish performing a similar function. Significantly less hostile, though still set apart, they bear witness to what was normal for all of us a couple short centuries ago: horse and buggy, wood stoves and lanterns. My intent is to be AI Amish, which means Internet Amish. Not 1780, but 1980. Neo-Amish. I'm fine driving a car and flipping a lightswitch, by which I mean that they don't make me into something I hate, which AI and [struck through: social media]...

4. Meta Is Launching Instagram, Facebook, and WhatsApp Subscriptions for "Fun Features"

Source: Daring Fireball | Published: 2026-05-30T15:34:02Z

Link: <https://techcrunch.com/2026/05/27/meta-officially-launches-instagram-facebook-and-whatsapp>

p-subscriptions-with-more-to-come-including-ai-plans/

Sarah Perez, reporting for TechCrunch: Meta is doubling down on its subscription offerings. On Wednesday, the social networking giant announced it's now rolling out its consumer subscription plans globally for its flagship apps, Instagram, Facebook, and WhatsApp, and beginning tests of new subscriptions for businesses, creators, and Meta AI users. For a few dollars per month, consumers subscribing to Instagram Plus (\$3.99/mo), Facebook Plus (\$3.99/mo), or WhatsApp Plus (\$2.99/mo) will gain access to extra features, like profile customization, super reactions, and story insights, among other things. In an announcement, Meta's head of product, Naomi Gleit, noted that "more fun features" will be added in the future. My first question about this was whether Meta would be using IAP on iOS and Android. On the one hand, Zuckerberg really resents Meta's subservient position to Apple and Google in the mobile ecosystem - that's what drove him to make a big wrong bet on the "metaverse" as the Next Big Thing. But on the other hand, what else are they going to do? Most people only use Meta's platforms via the phone apps, and if they're going to allow subscriptions via the apps, they have to p...

5. Daniel Jalkut on AI

Source: Daring Fireball | Published: 2026-05-30T15:27:20Z

Link: <https://mastodon.social/@danielpunkass/116639318125898071>

Daniel Jalkut, on Mastodon (cross-posted to Bluesky and Threads): My take on AI is, essentially, everybody who's against it is too against it and everybody who's for it is too for it. I concur with this take completely. (Sidenote: The different reply threads on the three networks speak loudly to the cultural and algorithmic differences between them. Good lord has Meta steered Threads into "make people argumentative" engagement.) ?

6. Yours Truly on TBN Yesterday

Source: Daring Fireball | Published: 2026-05-30T12:28:10Z

Link: <https://www.youtube.com/live/sQVwLUxFdMY?t=1997>

Fun show, good questions I thought. ?

7. Sharing the result of a single Windows Runtime IAsyncOperation among multiple coroutines, part 3

Source: The Old New Thing | Published: Fri, 29 May 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260529-00/?p=112368>

A variation where we try only once. The post Sharing the result of a single Windows Runtime IAsyncOperation among multiple coroutines, part 3 appeared first on The Old New Thing .

8. Microcode inside the Intel 8087 floating-point chip: register exchange

Source: Ken Shirriff's blog | Published: 2026-05-30T10:09:35.056-07:00

Link: <http://www.righto.com/2026/05/microcode-inside-intel-8087-floating.html>

In 1980, Intel introduced the 8087 floating-point chip, a co-processor that made floating-point operations up to 100 times faster. This chip was highly influential, and today most processors use the floating-point standard introduced by the 8087. The 8087 uses complicated algorithms to accurately compute functions such as square roots, tangents, and exponentials. These algorithms are implemented inside the chip in low-level code called microcode. I'm part of a group, the Opcode Collective, that is reverse-engineering this microcode. In this post, I take a close look at the microcode for one of the 8087's instructions: FXCH - and explain how the microcode works. The FXCH (Floating-point Exchange) instruction exchanges two floating-point registers. You might expect this instruction to be trivial, but there's more going on than you might expect; the microcode uses 14 micro-instructions to implement the exchange instruction. The Intel 8087 chip is packaged in a 40-pin DIP (dual in-line package). To explore the microcode, I opened up an 8087 chip and created a high-resolution image with a microscope. The large microcode ROM occupies a central position, holding the micro-instructions tha...

9. Spot checking polynomial identities

Source: John D. Cook | Published: Sat, 30 May 2026 21:06:49 +0000

Link: <https://www.johndcook.com/blog/2026/05/30/schwartz-zippel/>

If a polynomial identity holds at a few random points, it's very likely true. We'll make this statement more precise, but first let's look at some applications. You may want to test an identity that naturally presents itself as a statement that two polynomials are equal. Or you might use something like the binomial coefficient [?] The post Spot checking polynomial identities first appeared on John D. Cook .

10. Our unusual system of "web home directories" for people

Source: Chris's Wiki :: blog | Published: 2026-05-30T23:51:13Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/OurWebHomeDirectories>

One of the things we operate for the research side of the department is an old fashioned general purpose web server, where everyone has a home page area of their own in the traditional '/~/' style (cf). This web server has been there for a very long time, and one of the decisions that was made very early on was that for security reasons, the web server would not NFS mount people's regular home directories from our file servers . The traditional Apache way to do '/~/' home pages is to have some location under your home directory that's exposed as your web home page area ; the traditional name for this is ' public_html '. One alternative is to relocate this to a separate directory tree, but this directory tree is flat, which makes it awkward to have different pools of disk space for different people (which is absolutely required for us). Since we didn't want to use people's regular home directories for security reasons and we couldn't put everyone in one directory, we did the obvious hack: people have a different, special home directory on the web server. These home directories are in special 'webdir' filesystems on our file servers , and these webdir filesystems are the only NF...

11. Our servers mostly don't seem to have high peak power usage

Source: Chris's Wiki :: blog | Published: 2026-05-30T03:35:58Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/tech/ServersLowMaxPowerConsumption>

I wrote a bit ago about how our servers seem to have surprisingly low power consumption , where I looked at IPMI based power consumption information to look at their current, typical power usage and found that several of them were sitting at lower power usage than my desktops. I'm still interested in typical or average power usage for reasons beyond the scope of this entry, but now I'm also interested in maximum power usage that we've observed. One reason to be interested in maximum power usage is that servers are often given relatively high capacity power supplies. A 600 watt power supply is on the small side for even a 1U server, and my impression is that 800 watt and even 1200 watt PSUs are reasonably normal in basic servers. On the one hand, the vendors building servers don't know what they're going to be used for, so they're likely to be conservative. On the other hand, that's a lot of spare power for a system that is typically using, say, 24 watts. A PSU that is that far under its rating may not be all that efficient (although the one 600 watt server PSU I looked at had a 'platinum' rating). Obviously a server's PSU has to support not just its typical power draw but also its...

12. Pluralistic: Carneyism without Carney (30 May 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 30 May 2026 09:31:41 +0000

Link: <https://pluralistic.net/2026/05/30/rupture/>

Today's links Carneyism without Carney: Eh? Hey look at this: Delights to delectate. Object permanence: Replacing pharma patents with bounties; USTR v cheap leukemia meds; Plutocrats x wealth segregation; Anonymous Analytics; Scott Walker sells off donors; Anonymization v metadata; Probably; Amazon warehouse workers are the future of Amazon coders; Warcraft eggs; Brainwashing school; People who don't know The Onion is satire; "Company Town"; America is a scam. Upcoming appearances: London, Kansas City, LA, Menlo Park, Toronto, NYC, Edinburgh, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Carneyism without Carney (permalink) The "Third Way" in liberal politics involves saying things that working people love, but doing things that sociopathic plutocrats love. It works ?right up

until voters notice that you're not doing the things. That realisation breeds cynicism and fury and paves the way for fascist strongmen. It's really ugly, and no one does it uglier than Canada's Liberal Party. Remember that time Prime Minister Justin Trudeau marched with Gre...

13. Notes from May 2026

Source: Evan Hahn (dot com) | Published: Sat, 30 May 2026 00:00:00 +0000

Link: <https://evanhahn.com/notes-from-may-2026/>

My blog turned 16 this month! I did nothing to celebrate, but made some little tools and clicked some links about tech ethics. Things from me this month I published four little tools this month: ZIP Shrinker , a web app that shrinks ZIP files with higher compression ratios A command line tool to do (completely offline) translation Open Link in Unloaded Tab , a Firefox extension to open links without loading them png-cmp , a command line tool to compare PNG pixel data I also did some work on Helmet, my open source project: After over a year of quiet maintenance, I released version 8.2.0 with some small new features and documentation updates. In a step toward dropping GitHub, I moved the docs from a GitHub URL to helmet.js.org . And like every month, I wrote a few articles at Zelda Dungeon . I don't feel I wrote anything special this month, but my colleagues put together a feature about Zelda and mental health which was very affecting! Links to click on ?The vast majority of tech workers, at least those who I have encountered in my many years of reporting, are not vampiric Silicon Valley tech bro caricatures [? They] both like working with tech and ultimately want to see it serve th...

14. Arp 114:

Source: Maurycy's Blog | Published: Sat, 30 May 2026 00:00:00 +0000

Link: <https://maurycyz.com/astro/arp114/>

North is left. (mirrored). 26x18 arcmin (0.3 arcsec/pixel) dt {color: var(--em-color)} Total exposure time: 393 * 30 seconds = 3.2 hours. (across 2 nights) Exposure used in stack: 241 * 30 seconds = 2 hours. Telescope: C9.25 (230mm, f/10, fl=2300mm) Camera: IMX533 (16mm diagonal, square, color) Processing: (no software sharpening used) Calibration (dark + flats) Stacking (average w/ rejection) White balance and background subtraction Asinh stretch Light denoising Tone curve Crop This one has two entries in the Atlas of Peculiar Galaxies : Once as Arp 25 under "one heavy arm" and again as Arp 114 under "ellipticals close to and perturbing spirals". The spiral structure of NGC 2276 is somewhat unusual looking, although it's not an obvious case of tidal interaction. Hubble data suggests it's actually moving quite fast (950 km/sec) through space, which causes the lob-sided arms as it interacts with sparse gas clouds. If this is correct, the galaxy is not actually head-on as seen from earth: It tilts toward us on the eastern (top) side and away on the western (bottom) edge. This would require the galaxy to be stretched by the gravity of NGC 2300, so it's still an interacting galaxy, b...

15. Reading List 05/30/26

Source: Construction Physics | Published: Sat, 30 May 2026 12:02:52 GMT

Link: <https://www.construction-physics.com/p/reading-list-053026>

A California chemical leak, weapons-grade plutonium for nuclear reactor startups, a startup that will clean your house to get robot training data, Blue Origin's rocket explosion, and more.

16. This Week in Package Management: 30 May 2026

Source: Andrew Nesbitt | Published: 2026-05-30T10:00:00+00:00

Link: <https://nesbitt.io/2026/05/30/this-week-in-package-management.html>

Back for a second week, built from the package manager OPML feed collection and whatever I've posted or boosted on Mastodon . Security npm invalidated every granular access token with write access that bypassed 2FA following another Shai-Hulud-pattern attack, so CI pipelines that publish with one need to mint a new token. npm 11.16.0 ships phase one of the allowScripts install-script policy , an opt-in allowlist in package.json naming which dependencies may run lifecycle scripts; in this phase scripts outside the list still run but trigger a warning. pnpm 10.34.0 and 11.4.0 land the same security set on both maintained lines: a tarball-integrity

mismatch is now a hard failure instead of quietly re-resolving and overwriting the locked hash, unscoped `_authToken` is bound to the registry from the same config source so a workspace `.npmrc` can't redirect a credential set in `~/.npmrc`, git-resolution commit values must be a 40-char SHA to block `--upload-pack` injection from a hostile lockfile, and patch files can't reference paths outside the patched package. 10.34.1 followed by rejecting lockfile entries whose resolution: block has no integrity field at all, closing a path where a PR that...

17. On first looking into JAX

Source: Giles' blog | Published: Sat, 30 May 2026 18:45:00 +0000

Link: <https://www.gilesthomas.com/2026/05/on-first-looking-into-jax>

Much have I travell'd in the realms of gold, And many goodly states and kingdoms seen; Round many western islands have I been Which bards in fealty to Apollo hold. Oft of one wide expanse had I been told That deep-brow'd Homer ruled as his demesne; Yet did I never breathe its pure serene Till I heard Chapman speak out loud and bold: Then felt I like some watcher of the skies When a new planet swims into his ken; Or like stout Cortez when with eagle eyes He star'd at the Pacific -- and all his men Look'd at each other with a wild surmise -- Silent, upon a peak in Darien. John Keats, On First Looking into Chapman's Homer I've been working with PyTorch quite a lot for the last couple of years, and feel like I've come to a reasonably solid understanding of how it all fits together. Working through Sebastian Raschka's book "Build a Large Language Model (from Scratch)", training my own LLMs locally and in the cloud, rebuilding Andrej Karpathy's 2015-vintage RNNs -- over time, it all adds up! But, of course, there are other frameworks, and one I kept hearing about was JAX. While it's less dominant than PyTorch, it has a reputation for a certain cleanliness, a certain purity. And hav...

18. Galileo Notes: Ephemeris, Anomalies, I/NAV, SFRBX

Source: Bert Hubert's writings | Published: Tue, 30 Jul 2019 15:35:10 +0200

Link: <https://berthub.eu/articles/posts/galileo-notes/>

I've recently been spending quite some time on the EU's Galileo Navigation Satellite System. In this post you'll find some rough notes on things I've found out, both about how to receive data & how to understand the Galileo ephemeris parameters. Feedback is very welcome on bert@hubertnet.nl or [@bert_hu_bert](https://twitter.com/bert_hu_bert)! I'm no Galileo specialist, so I am sure this page could be better. Please also let me know if this page has been useful to you.

19. On the Pan-Genome

Source: Bert Hubert's writings | Published: Fri, 12 Jul 2019 09:10:21 +0200

Link: <https://berthub.eu/articles/posts/on-the-pan-genome/>

In this post I want to shine a light on an important development within the world of DNA which shows how much DNA and the world of computing are now intertwined. It turns out we share important problems, like in this case: file formats. Note: I expect to be posting more DNA material as part of my efforts to write a book on this fascinating subject. If you enjoy reading about DNA, you might like the links/presentations/videos found on this page.

20. From gnuplot to Matplotlib & Pandas

Source: Bert Hubert's writings | Published: Mon, 01 Apr 2019 09:59:00 +0100

Link: <https://berthub.eu/articles/posts/from-gnuplot-to-matplotlib-pandas/>

I've been using gnuplot since.. like forever. It is one of my best friends in plotting data and discovering what is going on. But for all its greatness, you do tend to run into a wall - once you step outside the things gnuplot is good at, suddenly large heaps of awk, sort, unique and odd shell scripts are required to get to the next level. This is no criticism of gnuplot - it is great for what it is for.