RSS Digest
Generated at: 2026-06-06 09:43:51

1. OpenAI Help: Lockdown Mode
Source: Simon Willison’s Weblog | Published: 2026-06-05T23:56:40+00:00
Link: https://simonwillison.net/2026/Jun/5/openai-help-lockdown-mode/#atom-everything
OpenAI Help: Lockdown Mode OpenAI first teased this in February , but now it’s live and "rolling out to eligible personal accounts, including Free, Go, Plus, and Pro, and self-serve ChatGPT Business accounts": Lockdown Mode is designed to help prevent the final stage of data exfiltration from a prompt injection attack by limiting outbound network requests that could transfer sensitive data to an attacker. Lockdown Mode does not prevent prompt injections from appearing in the content ChatGPT processes. For example, a prompt injection could appear in cached web content or in an uploaded file, and could still affect the behavior or accuracy of a response. This looks really good to me. The Lethal Trifecta occurs when an LLM system has access to all three of access to private data, exposure to untrusted content and a way to steal data and transmit it back to the attacker. The only way to solve the trifecta is to cut off one of the three legs, and by far the easiest leg to restrict without making your LLM systems far less useful is the exfiltration vectors to steal data. It looks to me like lockdown mode directly attacks that leg, using mechanisms that are deterministic and, crucially,...

2. Quoting Andreas Kling
Source: Simon Willison’s Weblog | Published: 2026-06-05T11:10:05+00:00
Link: https://simonwillison.net/2026/Jun/5/andreas-kling/#atom-everything
We will no longer accept public pull requests. [...] A substantial patch used to imply substantial effort, and that effort was a reasonable proxy for good faith. That assumption no longer holds. [...] Whether code was typed by hand is beside the point. What matters is who is responsible for it once it enters the browser. Ladybird is becoming a browser for real users. The people introducing changes to it must be the people who decide those changes belong in the project, and who will answer for the consequences. ? Andreas Kling , Changing How We Develop Ladybird Tags: ladybird , ai-ethics , open-source , generative-ai , ai , andreas-kling , llms

3. Nieman Journalism Lab: Twitter/X Punishes Accounts That Post Links
Source: Daring Fireball | Published: 2026-06-05T20:46:56Z
Link: https://www.niemanlab.org/2026/04/do-links-hurt-news-publishers-on-twitter-our-analysis-suggests-yes/
Laura Hazard Owen, writing for Nieman Journalism Lab back in April: I used Claude to help me scrape the 200 most recent tweets from 18 large publishers? X accounts and track the engagement (likes + comments + retweets) on each. Six of those publishers have paywalls: Bloomberg , CNN , Forbes , The New York Times , The Wall Street Journal , and The Washington Post . Nine don?t: Al Jazeera English , AP , BBC , Breitbart News , CBS News , Daily Wire , Fox News , NBC News , and Reuters . The last three accounts I looked at ? Leading Report , unusual_whales, and Globe Eye News ? are not news publishers, but aggregate breaking news in tweets without links. (Here, for example, is an example of a Leading Report tweet : ?BREAKING: Iran has halted direct talks with the US, per WSJ.? They?re sometimes referred to as engagement-maxing accounts. These charts make it pretty clear that links in tweets hurt engagement. The connection was so apparent in my analysis that a graph including all 18 publishers is almost unreadable: The traditional, link-loving publishers are clustered in the bottom left corner (lots of links, little engagement) in a nearly indistinguishable mass of bubbles, no matter ho...

4. Elon Musk?s X Is a Freak Show
Source: Daring Fireball | Published: 2026-06-05T20:24:58Z
Link: https://www.natesilver.net/p/social-media-has-become-a-freak-show
Nate Silver, back in April, under the headline ?Social Media Is Turning Into a Freak Show?, where by ?social media? he mostly discusses Twitter/X: But what does that remaining traffic consist of? I recently came across a bubble chart depicting the Twitter accounts that had

received the most ?engagement? in February 2026. It was depressing: most of the top accounts were extremely low-quality and highly partisan. I hadn?t even heard of many of them and only follow a handful of the top accounts. So I tracked down the original data myself and, with help from Claude, made my own improved version of the chart. Here, voilà, are the Twitter accounts with the most engagement so far in 2026: It?s not hard to notice that Twitter has become extremely right-leaning. But I?d argue there?s an equally important trend: the top accounts are of incredibly low quality. Elon, with the algorithmic boost he built in for himself, is at the eye of the storm, of course. But ?Catturd? literally gets far more engagement than the New York Times, for instance. There?s a common argument from proponents of the Musk-era X that the only problem is that left-leaning people have abandoned the platform. That the X a...

5. Checking in on Perplexity
Source: Daring Fireball | Published: 2026-06-05T15:26:29Z
Link: https://daringfireball.net/linked/2025/08/05/regarding-those-rumors-of-apple-pursuing-an-acquisition-of-perplexity
Yours truly, last August: I can?t see why Apple would want to get involved with a company like this though. Gurman?s report makes it sound like his sources are inside Apple, but man, this ?Apple + Perplexity? thing feels more like something Perplexity would be seeding than one that Apple executives would be leaking. Perplexity is still occasionally in the news (often not in good ways), but it seems to me they?ve slipped into the ?afterthought? tier of AI startups ? which is exactly why they started leaning into clownish stunts last year . Everyone who previously suggested Apple should ? or even might ? buy them has gone silent. ?

6. Why all the PRs?
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Fri, 05 Jun 2026 23:00:00 GMT
Link: https://idiallo.com/blog/why-all-the-prs
It’s a signal. That’s why we get AI-generated PRs. We told everyone, in order to get your resume taken seriously, you need to show your work. When I was getting started in my career, that meant having your own website that you contribute to regularly. So I did that. I built websites, I maintained them. I kept maintaining them even after I got the jobs because that’s how I actually honed my web programming skills. Where else was I going to try new frameworks, a new JavaScript paradigm, or try out Ruby on rails? I got the job, and I advised other developers to follow the same path. But then github became mainstream. Rather than just show a finished website, you could actually share the code that runs your project. Share a link to your github project and companies can review your code and directly gauge your experience. But even better, you can show your contribution to open source projects. Not just any projects. Popular projects. The github stars became a metric people look for. A signal that can be used to quickly assign a value to a candidate. But that?s the story told from the outside. I don?t think the github profile link was ever important, unless it was significantly good. Em...

7. Now that your newsletter is AI-generated, I’ve Unsubscribed
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Wed, 03 Jun 2026 21:10:00 GMT
Link: https://idiallo.com/blog/unsubscribed-from-ai-generated-newsletters
I’ve remained subscribed to some newsletters for over 20 years. The authors managed to keep my attention all that time. But then, one day, they decided to switch to an AI-generated newsletter without making any announcement. After a couple of weeks of blue high-tech image thumbnails, I simply hit unsubscribe. Here’s what happened: a person earned my trust. He maintained that trust for all those years. But then he thought the best way to improve was to take himself out of the equation. If you’re just going to present me with prompt-generated content, I hate to break it to you but I have access to ChatGPT, and I can do that myself. The reason the human voice matters to me is because there’s real experience behind the words. The oldest newsletter in my inbox is from when I was just 12 years old. It was from a French writer I used to read. After a decade of following him, the emails stopped coming. I was only reminded

a few years later, when the emails started coming back. I didn’t jump on it immediately. I didn’t even remember who it was. But when I read one at random, the words were different, the tone was nostalgic, and the name was unfamiliar. I dug deeper and found that the autho...

8. The web is changing, and we are not going back
Source: Software and Tech stories from an Insider - iDiallo.com | Published: Mon, 01 Jun 2026 19:34:37 GMT
Link: https://idiallo.com/blog/web-is-changing-we-are-not-going-back
Whenever I saw someone type a natural language query into Google, it made me cringe. "It’s not a person," I would say. "Type like you’re talking to a machine." This was especially true for programmers and it was before AI took over everything. Instead of "how do I write a function that reads a file?", I would suggest they use specific keywords, something that sounded more like machine language than conversation. "js function to read csv file" or "css gradient background property example." This got you better results. Even though Google was a sophisticated search engine, it was still doing a kind of keyword matching under the hood. But not anymore. You don’t get any advantage from writing in "machine language." Google understands natural language just as well. In fact, even better. How is it that in 2026, I Google things less than ever? It’s not that I know everything now. It’s more that I don’t want to call the friend who always talks too much. If the height of the Eiffel Tower ever comes up in conversation, I’ll type "eiffel tower wiki" and click through to Wikipedia. I don’t want to have a conversation about it. Googling something these days feels like Google is trying to join m...

9. I tested every IP KVM in my Homelab
Source: Jeff Geerling | Published: Fri, 05 Jun 2026 09:00:00 -0500
Link: https://www.jeffgeerling.com/blog/2026/i-tested-every-ip-kvm/
Since the PiKVM came out in 2017, there’s been an explosion of IP KVMs. I’ve tested almost every one . But what are they good for? You can use Remote Desktop, Screen Sharing, or VNC to remote control a computer from anywhere on a LAN. And if you don’t have a private VPN, you could use RealVNC , Raspberry Pi Connect , or wire up Tailscale or Pangolin for fully remote access. Those solutions are great, and so is SSH if you don’t need a full desktop.

10. Mr. Bessel?s eponymous functions
Source: John D. Cook | Published: Fri, 05 Jun 2026 12:04:28 +0000
Link: https://www.johndcook.com/blog/2026/06/05/mr-bessels-eponymous-functions/
Yesterday I wrote a post showing that the trapezoid rule evaluates the integral very efficiently. But how do we know what the exact integral is for comparison? If you ask Mathematica, it will tell you the integral equals ?2? J1(1) where J1 is a Bessel function. This may seem like rabbit out of a hat, [?] The post Mr. Bessel?s eponymous functions first appeared on John D. Cook .

11. It seems unlikely that I’ll have basic ARM servers to deal with
Source: Chris’s Wiki :: blog | Published: 2026-06-05T23:52:32Z
Link: https://utcc.utoronto.ca/~cks/space/blog/tech/ArmBasicServersSeemUnlikely
At work, we have a lot of basic servers. These basic servers are 1U rack servers, with unexciting amounts of disk bays, RAM, and CPU performance, because we buy them to do straightforward jobs (for example, much of our mail system hardware is basic servers). Right now, these are all x86 machines, but like a lot of people I’d love it if the x86 architecture had real competition (even if returning to the days of a multi-architecture Unix environment wouldn’t necessarily be fun). However, these days I don’t think we’re likely to see basic 1U ARM servers that we’re interested in, and why that is comes down to modern hardware and modern basic server power usage. At this point, the largest hardware difference between a reasonably performing basic x86 server and a reasonably performing basic ARM server will be the CPU (and surrounding chipset). Everyone uses the same RAM, the same PCIe busses, and quite possibly the same supporting chips for things like Ethernet and so on. And everyone is going into the same

rack form factor 1U boxes with very similar power supplies and physical sizes. My impression is that at the basic end, a comparably performing ARM server CPU is going to cost about...

12. The Emacs packages that I use (as of June 2026)
Source: Chris’s Wiki :: blog | Published: 2026-06-04T19:10:06Z
Link: https://utcc.utoronto.ca/~cks/space/blog/programming/EmacsPackages-2026-06
My Emacs configuration seems to have more or less settled down again after a flurry of changes, so it’s time to update my previous list of Emacs packages that I use , so that I can come back to this entry later and see how things have changed over time. A bunch of things haven’t changed since last time so I’m going to put the unchanged stuff at the bottom. Currently I’m using Emacs 30.2 everywhere so some of the things I’m mentioning here are now built in, which is why I’m not restricting this to third party packages. As before I’m going to exclude dependencies that are automatically installed by the Emacs package system (since I don’t use them, I just have them around in the background). In no particular order: In text-mode, I use Flyspell (although there’s a gotcha there). I’ve turned off Flyspell’s default binding to M-TAB, because I find Flyspell basically useless for actually correcting words. Almost all of what I use Flyspell for is simply marking words it thinks are misspelled or mistakes so I can spot them (and then fix them by hand). I now use Eglot as my LSP server environment , which is built in to GNU Emacs. I use LSP servers for Go , Python , shell programming, and...

13. Install-script allowlists
Source: Andrew Nesbitt | Published: 2026-06-05T12:00:00+00:00
Link: https://nesbitt.io/2026/06/05/install-script-allowlists.html
In most package managers a dependency?s install-time code runs by default the moment you install it: an npm postinstall, a Setuptools setup.py , a CPAN Makefile.PL , an RPM scriptlet, a Conda post-link, a Debian postinst . A handful require explicit per-package opt-in before any of that code runs, usually called an allowlist or a trusted-dependencies list depending on the tool. Per-package opt-in lists name which dependencies may run their install code: npm, pnpm, Bun, Deno, and Composer plugins all work this way. Global sandboxes (opam, Swift Package Manager, Nix, Guix, Portage) take a different shape, executing everything but constraining what that execution can reach. Identity and signature verification (RubyGems trust policies, Gradle dependency verification, NuGet trustedSigners, apt-secure) gates which artifacts get installed in the first place by who signed them, with no bearing on what their code subsequently does. An npm postinstall, a setup.py, a Makefile.PL or an RPM scriptlet fires during fetch or unpack. A Cargo build.rs or a Zig build.zig runs when the project is compiled, which on a fresh build is functionally the next step but is structurally distinct. JVM build fi...

14. JAX backends and devices
Source: Giles’ blog | Published: Fri, 05 Jun 2026 19:30:00 +0000
Link: https://www.gilesthomas.com/2026/06/jax-backends-and-devices
There’s nothing like writing your own code with a framework to clarify how things fit together! Continuing with my port of my PyTorch LLM code to JAX , I wanted to load up a large dataset: the 10,248,871,837 16-bit unsigned integers in the train split of gpjt/fineweb-gpt2-tokens . That’s just over 19GiB of data. from safetensors.flax import load_file ... full_dataset = load_file (dataset_dir / f "train.safetensors")["tokens"] When I ran that, I got a CUDA out-of-memory error: jax.errors.JaxRuntimeError: RESOURCE_EXHAUSTED: Out of memory while trying to allocate 19.09GiB. That makes sense! The allocation it was trying to do is exactly the size of the data I was trying to load. I have an RTX 3090 with 24 GiB, but some is already used up by the OS, various apps, and a model that the code creates earlier on. But in PyTorch land, I was used to things being loaded into RAM by default, and only moved over to the GPU when I asked it to do that. JAX was clearly loading to the GPU by default. How could I stop it from doing that for this case? The load into the GPU was happening inside Safetensors, in code I couldn’t directly control. Understanding how to do it helped me understand a lit...

15. In pursuit of desirable difficulties

Source: Westenberg. | Published: Sat, 06 Jun 2026 01:14:14 GMT
Link: https://www.joanwestenberg.com/in-pursuit-of-desirable-difficulties/
The psychologist Robert Bjork called them desirable difficulties. Learning sticks better when practice is made harder, rather than easier. Students who have to struggle to retrieve an answer remember it longer - and clearer - than students who are handed the same answer with minimal effort on their part. The

16. AI-indecision is a recursive trap. Don’t get stuck.
Source: Westenberg. | Published: Fri, 05 Jun 2026 04:06:59 GMT
Link: https://www.joanwestenberg.com/ai-indecision-is-a-recursive-trap-dont-get-stuck/
Jean Buridan was a 14th-century French philosopher and logician who twice served as rector of the University of Paris. His subject was the will, and he made an austere claim: the will follows the intellect. Show a rational creature the greater good and it’ll pick the greater good.

17. Premium: The Hater’s Guide To The AI Bubble 3.0
Source: Ed Zitron’s Where’s Your Ed At | Published: Fri, 05 Jun 2026 15:57:04 GMT
Link: https://www.wheresyoured.at/premium-the-haters-guide-to-the-ai-bubble-3-0/
Last year I wrote one of my favourite pieces ever ? The Hater?s Guide To The AI Bubble ? and followed it up with The Hater?s Guide To The AI Bubble Volume 2 several months later. Sadly, I?ve realized ?volume? is a

18. Pluralistic: Refining humanity (05 Jun 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 05 Jun 2026 20:49:39 +0000
Link: https://pluralistic.net/2026/06/05/defining-humanity/
Today’s links Refining humanity: What our technology is shows us what we’re not. Hey look at this: Delights to delectate. Object permanence: GNU Radio; France v "follow us on Twitter"; Aaronsw vindicated; Capitalism’s crooked refs. Upcoming appearances: Kansas City, LA, Menlo Park, Toronto, NYC, Edinburgh, South Bend. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. Refining humanity (permalink) One of the best ways to evaluate your own understanding of a subject is to attempt to explain it to someone else. Through explaining things, we discover how much of the "totally obvious" world is actually full of ambiguity, mystery and contradiction. There’s a great bit in Rowan Atkinson’s historical sitcom Blackadder that illustrates this principle. In "Ink and Incapability" Blackadder and friends have accidentally burned the only copy of Samuel Johnson’s original dictionary of the English language. To cover up their mistake, they decide that they will recreate the dictionary themselves. However, they founder on the first word they try to define, "A": Blackadder: Let’s...

19. Aggressive caching for a Mastodon reverse proxy: what to cache, what to never cache, and why content negotiation will eventually betray you
Source: IT Notes | Published: Fri, 05 Jun 2026 08:44:00 +0000
Link: https://it-notes.dragas.net/2026/06/05/aggressive_caching_for_a_mastodon_reverse_proxy/
I have written before about putting a cache in front of snac , and more recently about the HAProxy layer in front of FediMeteo . The general idea is always the same: the reverse proxy should absorb the repetitive, public work that has no business reaching the application server. This post is the same idea applied to a much louder neighbour: a Mastodon instance. The instance is mastodon.bsd.cafe , the proxy is nginx on FreeBSD, and the configuration below is what I am currently running in production. Mastodon is heavier than snac in every direction. It has Puma and Sidekiq behind it, more endpoints, more streaming, more federation patterns, and one specific characteristic that complicates everything: it serves multiple representations on the same URLs. The same path returns HTML to a browser, ActivityPub JSON to a remote instance, and sometimes plain JSON to an API client. If the proxy treats the URL as one thing, sooner or

later it will return the wrong thing to the wrong client. Most of the work below comes from that single observation. If I had to summarize this whole post in a single sentence, it would be this: Mastodon is not a website. It is a website, an API, and an Activity...

20. The Giant's Cup
Source: Herman's blog | Published: 2026-06-05T08:31:20.644159+00:00
Link: https://herman.bearblog.dev/the-giants-cup/
I recently completed my first long trail race, set in the Southern Drakensberg mountains with Emma, my siblings and their partners. It spanned 2 days with 30km on the first, and 15km on the second, winding through valleys and around mountains. It was spectacular. I love coming to the Drakensberg, which is a beautiful and unique mountain range, and I'm glad to have the excuse to be here. While I'm generally an active person, I've never been a runner. My first proper run (more than 3km) was in January this year after my brother twisted my (and Emma's) arm into signing up for this race. Since this was not a trivial distance, it meant I couldn't just wing it on my base fitness, and so Emma and I started running 2-3 times a week in preparation. We had 15 weeks until the event and settled on a fairly simple training plan: Run 5km each Wednesday, usually on the Sea Point Prominade, while increasing the speed gradually each week (it took me a few weeks to get to the initial 5km though) Longer trail runs on the weekend, generally on Table Mountain or in the Cape Winelands, increasing distance by about 1km each week Hop on the treadmill at a reasonably fast pace for about 20 minutes after g...

21. Spectre & Meltdown: tapping into the CPU's subconscious thoughts
Source: Bert Hubert's writings | Published: Sat, 06 Jan 2018 10:34:05 +0100
Link: https://berthub.eu/articles/posts/spectre-meltdown/
Comments are very welcome on bert@hubertnet.nl or @bert_hu_bert. Update: several constructive remarks have been used to improve this text. Thanks! In this post I will attempt to fully explain the Spectre and Meltdown vulnerabilities in an accessible way. I decided to write it up after I realised it took me more than a day to figure it out, even though I?ve been doing security related stuff on CPUs for 20 years.

22. Kolmo
Source: Bert Hubert's writings | Published: Tue, 24 Oct 2017 10:24:57 +0200
Link: https://berthub.eu/articles/posts/kolmo/
Welcome to Kolmo There are three avenues to learn about Kolmo (GitHub). There are two Youtube videos, one from the most excellent NLNOG 2017, the other from the equally excellent UKNOF 38. These videos explain a lot of the the history and the ?why? behind Kolmo. There is also the kolmo.org website, which is ?self-hosted? by the ws Kolmo-powered webserver. Finally, there is this post, in which I focus on what Kolmo actually does, without spending much time on the ?why? or the history.

23. Some brief notes on making ethernet cables
Source: Bert Hubert's writings | Published: Sat, 30 Sep 2017 20:07:40 +0200
Link: https://berthub.eu/articles/posts/ethernet-cables/
Some brief notes on making ethernet cables With thanks to the most excellent NLNOG for a lot of mostly correct advice! Friends will tell you ?crimping? ethernet cables is easy. Most of your friends either have overly rosy memories, or have practiced on 200 cables before telling you it is now easy. It is not easy. This page is for if you have a few extra euros/dollars/coins to spare, and don?t want to spend rare hours of concentrated effort to crimp an ethernet cable using unfamiliar equipment and supplies.

24. The back cover of C++: The Programming Language also raises questions not answered by the front cover
Source: The Old New Thing | Published: Fri, 05 Jun 2026 14:00:01 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20260605-01/?p=112391
Not doing the reading. The post The back cover of C++: The Programming Language also raises

questions not answered by the front cover appeared first on The Old New Thing .

25. Rotation revisited: Avoiding having to calculate the gcd when doing cycle decomposition
Source: The Old New Thing | Published: Fri, 05 Jun 2026 14:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20260605-00/?p=112389
Math is hard. Let’s go counting! The post Rotation revisited: Avoiding having to calculate the gcd when doing cycle decomposition appeared first on The Old New Thing .

26. Rotation revisited: Cycle decomposition in clang?s libcxx
Source: The Old New Thing | Published: Thu, 04 Jun 2026 14:00:00 +0000
Link: https://devblogs.microsoft.com/oldnewthing/20260604-00/?p=112384
Rotating in the minimum number of steps by performing cycle decomposition. The post Rotation revisited: Cycle decomposition in clang?s libcxx appeared first on The Old New Thing .