

1. Working with product managers

Source: seangoedecke.com RSS feed | Published: Mon, 08 Jun 2026 00:00:00 GMT

Link: <https://seangoedecke.com/working-with-product-managers/>

The relationship engineers have with product management is more dysfunctional than with any other part of the company. There's no shared culture or language like there is with other engineers, and the rules of "who gets to tell who what to do" aren't as clear-cut as they are with managers. Engineers don't have a lot in common with legal, or design, or sales, but they also don't need to interact much with those roles. In my experience, engineers are communicating with product managers almost every single day. Against the "product mommy" The worst version of the product/engineering relationship goes something like this: Engineers are technically competent but are too autistic to be fully trusted. They need a kind-but-stern parental figure who knows how to communicate to other stakeholders in the organization (for instance, by being comfortable using the word "stakeholders"), and how to keep engineers from going off in the wrong direction. This entire gross dynamic is neatly captured by the popular term "product mommy" 1 . I really, really don't like that term, or this entire dynamic in general. Almost none of my relationships with my product managers have been anything like this, th...

2. Siri AI at WWDC 2026

Source: Simon Willison's Weblog | Published: 2026-06-08T23:58:04+00:00

Link: <https://simonwillison.net/2026/Jun/8/wwdc/#atom-everything>

Given how badly burned anyone who took Apple's 2024 WWDC Apple Intelligence announcements at face value was, I'm holding to a strict "I'll believe it when I see it" policy for everything they announced today . The new Siri AI features do at least look feasible with today's technology, especially since Apple are licensing a custom Gemini-derived model that they can run on their own Private Cloud Compute . It sounds like they'll be taking advantage of vision LLMs to extract information from the user's screen, which neatly sidesteps the need for every existing application to ship custom code in order to integrate with Apple Intelligence. Vision LLMs were a much less mature category in June 2024. The new Core AI library looks like a good step in enabling developers to finally take full advantage of Apple's hardware for running their own models. It integrates with Meta's open source PyTorch ecosystem, using these Core AI PyTorch extensions : Core AI PyTorch Extensions (coreai-torch) is a Python package that bridges PyTorch and Core AI. You can use it to bring up an existing PyTorch model ? exported as a torch.export.ExportedProgram ? into a Core AI AIProgram ready to run on Apple har...

3. From the Annals of People Having Knowledge of the Matter, Siri AI Extensions Edition

Source: Daring Fireball | Published: 2026-06-09T01:37:27Z

Link: <https://www.bloomberg.com/news/articles/2026-03-26/apple-plans-to-open-up-siri-to-rival-a-i-assistants-beyond-chatgpt-in-ios-27>

Mark Gurman, reporting (?) for Bloomberg two short months ago: Apple Inc. plans to open Siri to outside artificial intelligence assistants, a major move aimed at bolstering the iPhone as an AI platform. The company is preparing to make the change as part of a Siri overhaul in its upcoming iOS 27 operating system update, according to people with knowledge of the matter. The assistant can already tap into ChatGPT through a partnership with OpenAI, but Apple will now allow competing services to do the same. The company is developing new tools to allow AI chatbot apps installed via the App Store to integrate with the Siri assistant, said the people, who asked not to be identified because the plans haven't been announced. The chatbots will also work with an upcoming Siri app and other features in the Apple Intelligence platform. That means, for instance, if users have Alphabet Inc.'s Google Gemini or Anthropic PBC's Claude installed, they'd be able to send queries to those services from within the Siri voice assistant, just like they have been able to with ChatGPT since Apple Intelligence launched in 2024. Maybe Apple ran out of time today, and will announce this tomorrow? Maybe they f...

4. ?Microsoft and OpenAI Broke Up ? Now They?re Ready to Fight?

Source: Daring Fireball | Published: 2026-06-04T20:25:27Z

Link: https://www.theverge.com/ai-artificial-intelligence/942242/microsoft-build-ai-agents-open-ai-competition?view_token=eyJhbGciOiJIUzI1NiJ9.eyJpZCI6IjdiRHFjMlJadmgiLCJwIjoilL2FpLWFyZGlmaWNpYWwtaW50ZWxsaWdlbmNILzk0MjI0Mi9taWNyb3NvZnQtYnVpbGQtYWktYWdlbnRzLW9wZW5haS1jb21wZXRpZGlvbilsImV4cCI6MTc4MTAzNjQ2OSwiaWF0IjoxNzgwNjA0NDY5fQ.jp0KO9OVCO-fGkk1Utt0NIEn97JWAl8zs0zhjf2V2MQ

Hayden Field and Tom Warren, writing for The Verge (gift link): This year?s Build had the vibe of a freshly single divorc e posting a thirst trap on Instagram. ?It?s always fun to be at developer conferences in times of great change,? Microsoft CEO Satya Nadella said onstage Tuesday, adding that events like this are about ?coming to grips with the new opportunity.? AI chief Mustafa Suleyman, in an interview with The Verge, put it even more bluntly. ?The goal is to prove that we can become one of the top four labs in the world,? Suleyman said. ?There?s three labs that matter, Google DeepMind, OpenAI, and Anthropic. We are not one of them at the moment, and that?s always been my intention. It?s why I came here. I want to build the very best frontier models in the world, fully multimodal, and in order to do that, we have to prove that we can do everything that we need to from the ground up, and we?re not just going to take from others.? Refreshingly blunt. But hasn?t that been Microsoft?s plan for Bing since it was announced in 2009? I mean I guess you can say that Bing is one of the top four search engines in the world. Maybe you can even say it?s one of the top two. But it?s irrele...

5. Lingon and Lingon Pro 10

Source: Daring Fireball | Published: 2026-06-04T18:43:47Z

Link: <https://www.peterborgapps.com/lingon/>

Peter Borg: Lingon makes scheduling apps, scripts, shortcuts, and commands feel simple. Create a task in minutes, run it on a schedule, and stay in control. Lingon helps you run whatever you want whenever you want without living in Terminal. Schedule apps, scripts, shortcuts, and commands with a clear, friendly UI. Run tasks at specific times, on intervals or at login. Optional notifications make it easy to keep control. Two separate apps. Lingon is the simpler Mac App Store version and free to use, while Lingon Pro is the advanced one-time purchase with extra power. Lingon and Lingon Pro are great apps. I?ve been meaning to recommend them for a while. Back in 2023 I wrote about a problem I was having with Maestral , the incredible ?works like Dropbox in the old days? open-source Dropbox client, where Maestral would just silently crash once in a while and I wouldn?t notice for a while. Then I would notice, manually re-launch Maestral, and have to wait while Maestral synced. Or, worse, I?d put a podcast recording in a shared folder and walk away from my computer, and my editor would never get the file because Maestral wasn?t running. My write-up described how I solved the problem w...

6. ppclp.ai announces 100x Productivity Gains

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Mon, 08 Jun 2026 19:50:00 GMT

Link: <https://idiallo.com/blog/100x-productivity-gain>

ppclp.ai, North America's third-largest AI-native manufacturer of premium wire-formed office fasteners, formerly known as Paper Clip Company, announced a landmark 100x improvement in its proprietary Organizational Productivity Index (OPI?), cementing what leadership is calling "a new era of operational excellence" and "a little bit of a miracle." The breakthrough follows an 18-month company-wide initiative called Project Streamline, during which all 340 employees completed mandatory efficiency training, adopted a new Jira-based workflow system powered by Rovo, and attended a two-day offsite in Scottsdale where a consultant named Derek asked everyone to "think about how they think about work." "I used to open fourteen browser tabs and just stare at them. Now I open fourteen browser tabs, and AI agents are looking at them for me. I've never felt more in control." ? Sandra K., Senior Clip Assembly Coordinator, 11 years at ppclp.ai Central to the transformation is Rovo, Atlassian's AI, which the ppclp.ai Jira Center of Excellence team has deeply integrated into the company's ticket lifecycle. Rovo now autonomously opens tickets when it detects workflow friction, assigns them to the ap...

7. Should we care any more about Googlebot crawling our sites?

Source: Chris's Wiki :: blog | Published: 2026-06-08T02:04:21Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/web/DoYouNeedGooglebotCrawling>

One piece of technology news of the time interval is that Google is no longer going to be providing Internet search, it's going to be providing 'answers' . On the Fediverse, I had a reaction to that : Given Google's apparent change to what Google Search will be, I'm wondering if I should even allow Googlebot to crawl my techblog. If they're going to AI slop up anything they tell people after a search, I'd rather have my writing excluded entirely rather than be garbled. (I have no idea how much traffic I get from people using Google Search, and if Google is going to do this I feel like I should punish them by pushing people away from them, even a little bit.) To echo what Paul Cantrell said on the Fediverse , we tolerated and even embraced Googlebot crawling our websites as part of a social bargain. Allowing Googlebot was a big part of how people found our sites and our work, which is to say through searches. That social bargain has been fading as Google put more and more things inline, but at least they were still providing links and directly showing our words (I know, usually). But now Google has said out loud that Googlebot is just the front end crawling ingester to an LLM syste...

8. How many consecutive hyphens can you have in a domain name?

Source: Terence Eden's Blog | Published: Mon, 08 Jun 2026 11:34:59 +0000

Link:

<https://shkspr.mobi/blog/2026/06/how-many-consecutive-hyphens-can-you-have-in-a-domain-name/>

A seemingly simple question which sent me down into the murky depths of standards. How many consecutive hyphens can you have in a domain name? It probably isn't sensible to name your online presence a-----hyphen.com - but is there anything technically stopping you? Table of ContentsHistoryTLD RestrictionsAnomaliesSo What? History Let's do some history! This is 1978's "HOST NAMES?"

9. Package Manager Patents

Source: Andrew Nesbitt | Published: 2026-06-08T10:00:00+00:00

Link: <https://nesbitt.io/2026/06/08/package-manager-patents.html>

Patents and applications relevant to package manager design, grouped by area. Mostly US filings, found through Google Patents searches on the obvious terms. Each entry lists the assignee, filing and grant dates, and current status, followed by a short summary of the core claim and a prior-art note where open-source predecessors are well-documented. Manifests and dependency resolution US6381742B2 - Software package management . Microsoft. Filed June 1998, granted 2002, expired 2018. Claims a distribution unit, a manifest file, and a code store data structure, with dependency resolution at install time and shared-component tracking at uninstall. Prior art: CPAN manifests (1995), dpkg control files (1995), RPM (1997), FreeBSD ports (1993). US7222341B2 - Method and system for processing software dependencies in management of software packages . Microsoft. Filed February 2002, granted 2007, expired 2019. Continuation of US6381742B2, sharing its June 1998 priority date. Claims the install-time loop: check installed, identify missing dependencies, fetch from specified sources, extract, register, update the code store. Prior art: as for US6381742B2. US9348582B2 - Systems and methods for s...

10. AI Is Slowing Down

Source: Ed Zitron's Where's Your Ed At | Published: Mon, 08 Jun 2026 15:37:39 GMT

Link: <https://www.wheresyoured.at/ai-is-slowing-down/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA , Anthropic and

11. Coding Is Designing

Source: Jim Nielsen's Blog | Published: Sun, 07 Jun 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/code-is-design/>

Code isn't just a way to implement a design, it's a way to find one. With an interface, you have to use it, feel it, interact with it, and poke at it to see the relationships between things. Change X, see Y react. If it doesn't feel right, tweak it. Change X again, now Y reacts differently. Better. Keep tweaking? this here, that there, until the relationships of all the disparate elements fall into place as a single whole. Balanced. Design is how it works? and code is the tool to specify how it works. Reply via: Email • Mastodon • Bluesky

12. Planescape: Torment, Part 2: ?to the Desktop

Source: The Digital Antiquarian | Published: Mon, 08 Jun 2026 15:36:41 +0000

Link: <https://www.filfre.net/2026/06/planescape-torment-part-2-to-the-desktop/>

This article tells part of the story of Dungeons & Dragons on the tabletop and on computers. Usually if you choose the longest dialog option, that's the best option. ? Chris Avellone Quite some years ago now, I briefly interviewed Brian Fargo about Interplay's 1988 adaptation of the William Gibson novel Neuromancer. He was plainly busy [?]

13. The sample efficiency black hole

Source: Dwarkesh Podcast | Published: Mon, 08 Jun 2026 18:09:37 GMT

Link: <https://www.dwarkesh.com/p/the-sample-efficiency-black-hole>

"We see these AIs as a galaxy glittering with capabilities, but at their center, invisible to the naked eye, holding all the constellations together, is an unimaginably massive black hole of data."

14. Hacking for Defense @ Stanford 2026 ? Lessons Learned Presentations

Source: Steve Blank | Published: Mon, 08 Jun 2026 13:00:15 +0000

Link:

<https://steveblank.com/2026/06/08/g-for-defense-stanford-2026-lessons-learned-presentations/>

We just wrapped up our Hacking for Defense class at Stanford. This was the 11th year we've taught Hacking for Defense, and the impact of asymmetric warfare, (drones, off-the-shelf technologies, etc.) disruptive technologies (AI, commercial access to space) and a startup friendly DoW acquisition system ? make it feel like a much different class than [?]

15. xAI is looking more like a datacentre REIT than a frontier lab

Source: Martin Alderson | Published: Mon, 08 Jun 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/xais-new-rental-business/?utm_source=rss&utm_medium=rss&utm_campaign=feed

xAI is renting huge amounts of GPU capacity to Anthropic and Google. Financial engineering ahead of the SpaceX IPO, a real compute shortage, or a genuine datacentre advantage? Probably all three.

16. De gietijzeren pan en big tech

Source: Bert Hubert's writings | Published: Mon, 08 Jun 2026 13:00:00 +0200

Link: <https://berthub.eu/articles/posts/de-gietijzeren-pan-en-big-tech/>

Al een jaar of tien koken we thuis op gietijzeren pannen. Inmiddels voelt koken op een anti-aanbakpan voor mij als koken op plastic. De voordelen zijn verder evident. Ik heb in die tien jaar tijd nog geen pan hoeven vervangen, bijvoorbeeld. Teflon- en vergelijkbare pannen verliezen na verloop van tijd altijd een deel van hun anti-aanbak laag, en ik heb het idee dat je dan behoorlijk wat van die laag hebt opgegeten tegen die tijd.

17. How About Some Actual Innovation

Source: Bert Hubert's writings | Published: Sun, 17 Apr 2016 20:45:41 +0200

Link: <https://berthub.eu/articles/posts/how-about-some-actual-innovation/>

A meandering walk through innovation that also reviews a book This article is part of a series on (European) innovation and capabilities. I care deeply about innovation. It is literally where the future comes from, but it is a curious thing. Innovation and its more powerful

partner, invention proceed at a snail's pace. In a world where change is the only constant is a lie. It took this world 20 years to invent a decent can opener.

18. How Many Hours for Multithreading the Server?

Source: Bert Hubert's writings | Published: Tue, 02 Sep 2014 20:25:58 +0200

Link: <https://berthub.eu/articles/posts/how-many-hours-for-multithreading/>

How many hours for multithreading the server? Or: dealing with overly detailed project planning

Many developers, me included, dread that moment. Someone sits down with you and wants to know how many hours each step of the project will take. And of course the thing is, if you are doing something that has been done many times before, you might be able to provide a detailed estimate. People that build houses work like this, but they still often get it wrong.