

1. Statement on the US government directive to suspend access to Fable 5 and Mythos 5

Source: Simon Willison's Weblog | Published: 2026-06-13T01:01:50+00:00

Link: <https://simonwillison.net/2026/Jun/13/us-government-directive-to-suspend-access/#atom-everything>

Statement on the US government directive to suspend access to Fable 5 and Mythos 5 Well this is nuts : The US government, citing national security authorities, has issued an export control directive to suspend all access to Fable 5 and Mythos 5 by any foreign national, whether inside or outside the United States, including foreign national Anthropic employees. The net effect of this order is that we must abruptly disable Fable 5 and Mythos 5 for all our customers to ensure compliance. Access to all other Anthropic models will not be affected. We received the directive from the government today at 5:21pm (ET). The letter did not provide specific details of its national security concern. Our understanding is that the government believes it has become aware of a method of bypassing, or "jailbreaking" Fable 5. We reviewed a demonstration of this specific technique being used to identify a small number of previously known, minor vulnerabilities. These vulnerabilities all appear relatively simple, and we have found that other publicly-available models are able to discover them as well without requiring a bypass. [...] To date, the government has only given us verbal evidence of a potent...

2. OpenAI WebRTC Audio Session, now with document context

Source: Simon Willison's Weblog | Published: 2026-06-12T23:53:04+00:00

Link: <https://simonwillison.net/2026/Jun/12/openai-webrtc/#atom-everything>

OpenAI WebRTC Audio Session, now with document context I built the first version of this tool in December 2024 to try out the then-new OpenAI WebRTC API for interacting with their realtime audio models. Last month OpenAI introduced a brand new model to that API called GPT?Realtime?2 , which they promoted as "our first voice model with GPT?5?class reasoning" - with a Sep 30, 2024 knowledge cut-off. I've been waiting for that model to show up in the ChatGPT iPhone app but it still hasn't, so I revisited my old playground. You can now pick the better model, and you can also paste in a big chunk of document context so you can have as audio conversation in your browser about whatever information you think would be useful to explore in a conversational way. Tags: audio , tools , ai , openai , generative-ai , llms , multi-modal-output , webrtc

3. Quoting Andrew Singleton

Source: Simon Willison's Weblog | Published: 2026-06-12T18:09:21+00:00

Link: <https://simonwillison.net/2026/Jun/12/andrew-singleton/#atom-everything>

Jenny owns a crematorium. John?s propane company gives her a \$20 billion investment in return for 5 percent of her operation. Jenny throws \$10 billion into the incinerator, then pays John \$10 billion to buy propane to burn that money to ashes. John reports that his AI investments have generated \$10 billion in revenue this quarter and that he owns 5 percent of a \$100 billion business. A reporter from Forbes is assigned to profile John and Jenny, and over the course of his research, he becomes embroiled in a passionate but confusing three-way love affair with them, which eventually turns into a polyamorous common-law marriage. His profile is glowing, but light on financial details. ? Andrew Singleton , AI Economics for Dummies Tags: ai

4. ? The Talk Show: Live From WWDC 2026

Source: Daring Fireball | Published: 2026-06-12T23:36:15Z

Link: https://daringfireball.net/2026/06/the_talk_show_live_from_wwdc_2026

Recorded in front of a live audience at The California Theatre in San Jose on Tuesday 9 June 2026, special guests Joanna Stern and Nilay Patel join John Gruber to discuss Apple?s announcements at WWDC 2026.

5. The WWDC 2026 Keynote and State of the Union on YouTube

Source: Daring Fireball | Published: 2026-06-12T17:28:19Z

Link: <https://www.youtube.com/watch?v=hF8swzNR1-o>

Apple's Developer app lets you download local copies of every session, including the State of the Union, except the keynote. Why this is I don't know. But if you want a local copy, you can grab it from YouTube. Speaking of the State of the Union, the full version runs just over an hour, but Apple cut together a 4.5-minute recap . If you haven't watched the full thing you should at least watch that recap. ?

6. The European Commission Response to Siri AI and the DMA

Source: Daring Fireball | Published: 2026-06-12T17:02:45Z

Link: https://www.linkedin.com/posts/thomas-regnier-24a05810b_what-is-the-true-story-behind-apples-decision-activity-7470439874664280064-TuEt

Thomas Regnier, spokesperson for the European Commission, in a statement posted to LinkedIn (with edited video, if you'd like to watch him read parts aloud): What is the true story behind Apple's decision not to roll out ?Siri AI? in the EU? This decision is Apple's and Apple's only. Because absolutely nothing in the DMA prohibits Apple from rolling out new features in the EU. Yes, the European Commission and Apple had a few contacts on ?Siri AI?. But instead of offering a compliant solution, Apple asked to be exempted from its interoperability obligations under the DMA ? and this for 18 months. That's not an option. EU rules are non negotiable. And it would mean that no AI agent other than ?Siri AI? could be chosen by EU consumers. Apple, like any other gatekeeper, cannot close the market. The DMA is very clear about that. Our developers have the right to compete. And our consumers the right to choose. Those who want to keep using Apple products in their current form can of course do it. But for those who want to use another AI agent, the DMA will give them the possibility to do so. Why this was posted to LinkedIn and not on the EC's own press website is as inexplicable as Regnie...

7. You can finally power on a Mac remotely

Source: Jeff Geerling | Published: Fri, 12 Jun 2026 09:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/power-on-your-mac-remotely/>

Apple FINALLY lets you turn on your Mac remotely , without having to press the power button. In the media, articles suggest it's a reaction to Mac mini power button complaints . While I agree the M4 mini's power button is in a really dumb spot, that's not why I care about this feature. The two bigger use cases for me have been a pain for years:

8. Understanding the rationale behind a rule when trying to circumvent it

Source: The Old New Thing | Published: Thu, 11 Jun 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260611-00/?p=112415>

I mean, technically I didn't do it. The post Understanding the rationale behind a rule when trying to circumvent it appeared first on The Old New Thing .

9. I can never fully embrace LLMs for code

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Fri, 12 Jun 2026 12:00:00 GMT

Link: <https://idiallo.com/blog/i-can-never-embrace-llms-to-write-code>

My younger sister graduated with a CompSci degree a few years ago. I've been behind her, motivating her and demystifying the world of programming from the very beginning. There was a piece of advice I repeated everyday, trying to make her understand how to operate. The problem was, she was trying to read and understand every line of code in a function before using it. I thought it was non-sense. Someone, much smarter than us has created that function, it's part of a vetted library, it has been tested already. All you have to do is use it. "After all, you don't need to understand how an internal combustion engine works, yet you feel safe driving your car, don't you?" Now, I find myself right at that inflection point. When I use an LLM to generate code, whether it is to define a single function or to create a long running job, I find this need to understand it. I cannot commit code that I don't understand. I posted about how I spent 10 hours reworking what the AI had created in 12 minutes . I didn't do so because I

didn't like the style of the code, or the naming convention. I did it because the code didn't work. As simple as that. Every time I generate code and trust it to be worki...

10. Why are cached input tokens cheaper with AI services?

Source: Xe Iaso's blog | Published: Fri, 12 Jun 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/why-llm-cached-token-cheaper/>

TL;DR: the GPU doesn't have to math as hard

11. Joint Guidance on Vulnerability Naming and Disclosure

Source: Andrew Nesbitt | Published: 2026-06-12T10:00:00+00:00

Link: <https://nesbitt.io/2026/06/12/joint-guidance-on-vulnerability-naming-and-disclosure.html>

FOR IMMEDIATE RELEASE Contact: naming@vna.example Subject: Vulnerability Naming Authority Announces Naming Process and Domain Allocation Embargo: None The Vulnerability Naming Authority (VNA), in coordination with the CVE Numbering Authority consortium and the National Telecommunications and Information Administration, has published a unified process for the assignment, registration, and disclosure of named vulnerabilities. The process introduces a controlled vocabulary, a centralised approvals registry, and a top-level domain, .vuln , allocated for use exclusively in disclosure communications. The process applies to any vulnerability disclosed publicly by an entity operating within the United States.

Vulnerabilities assigned only a CVE identifier remain out of scope. The Naming Process A named vulnerability is defined as a vulnerability that the discoverer intends to refer to by name in disclosure materials, including but not limited to: the discoverer's blog, the discoverer's employer's blog, the discoverer's employer's marketing department's blog, a conference programme, a podcast episode title, and any subsequent press coverage. Each named vulnerability is described by a struc...

12. Gadget Review: TP Link EH210 Ethernet Splitter (USB-C) ?????

Source: Terence Eden's Blog | Published: Fri, 12 Jun 2026 11:34:18 +0000

Link: <https://shkspr.mobi/blog/2026/06/gadget-review-tp-link-eh210-ethernet-splitter-usb-c/>

When I ran Ethernet around our house, I thought I was being clever. A CAT6 cable for every room - lush! Some of my rooms have lots of devices, so they get a nice big Ethernet switch with lots of ports and blinking lights. But most of my rooms don't have that many devices. Our gym had only an Internet connected TV so that I could watch Quibi while exercising. Recently we added a Kodi box so that?

13. Running a modern email system requires non-trivial staff time

Source: Chris's Wiki :: blog | Published: 2026-06-13T00:27:44Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/MailSystemsPeopleCost>

In a comment on my entry on universities, email, and the issues of running things in house , I mentioned that our departmental email system has a non-trivial cost in both hardware and staff time. For the hardware costs I can easily count servers , but things are more fuzzy for the staff time side of things so I'm not going to try to come up with a number (especially because our email system itself is somewhat complex because we have a number of unusual features, such as our simple mailing lists). The bigger issue is that the basic time cost to maintain your email system is in some sense an illusion, or at least an inadequate measure. The reality of modern email is that to run a modern email system you need to know about a lot of stuff, and worse the stuff that you need to know about keeps evolving. You can't set up a mail system and then walk away from it apart from software and hardware upgrades; instead, you have to keep on top of a perpetually changing and evolving landscape of anti-spam systems and especially what you need to keep your outgoing email being delivered. SPF, DKIM, DMARC, DMARC alignment, and so on are merely today's names; in a few years there will be more, diff...

14. Premium: The Silicon Valley Bubble (Part 1)

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 12 Jun 2026 17:16:51 GMT

Link: <https://www.wheresyoured.at/premium-the-silicon-valley-bubble-part-1/>

Friends, I believe we're approaching the end of this era. Both OpenAI and Anthropic have filed the paperwork to go public, starting a race for exit liquidity for two companies that burn billions of dollars a year and have no path to profitability. Both of these companies are

15. AI will be massively deflationary

Source: the singularity is nearer | Published: 2026-06-11T00:00:00-07:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/06/11/ai-will-be-deflationary.html>

The funny thing about Anthropic haters is that they still mostly believe Anthropic's marketing. They think Claude is a recursively self-improving silicon God, and that we are all a couple refusals away from falling into the perpetual underclass. This gives them way more power than they deserve. Of course they believe they should dictate morality to you and control the future light cone of the universe and everything they do is justified for the mission and so on and so forth. And yea, the model is pretty good, so you worry that if the model is good that they will get their way about other things. A lot of people in history have had dumb ideas about ruling the world, doesn't mean they can actually do it. They are a product of the rationalist cult (which surprise surprise, thinks the world is going to end soon but advocates for polyamory in the mean time). There's some short-term dangers from them, but long-term very little. Remember when SBF wanted to regulatory capture all of crypto? And similar to Claude, FTX was a successful product! But the totalizing nature of the ideology contains its downfall. The tech to make models is way more commodity than previous generations of techno...

16. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 12 Jun 2026 16:31:11 +0000

Link: <https://www.filfre.net/2026/06/this-week-on-the-analog-antiquarian/>

Opus 2: Henry VI, Part 1

17. Pluralistic: Google's new remote attestation scheme is every bit as terrible as its old remote attestation scheme (12 Jun 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 12 Jun 2026 20:45:52 +0000

Link: <https://pluralistic.net/2026/06/12/compelled-speech/>

Today's links Google's new remote attestation scheme is every bit as terrible as its old remote attestation scheme: Not even a QR code can produce a kissable pig. Hey look at this: Delights to delectate. Object permanence: Arrested at Toronto G20; Rule by rentiers; Wrong about the First Amendment; Mounties x Stingrays; (EU) Privacy without monopoly. Upcoming appearances: LA, Menlo Park, Toronto, NYC, Philadelphia, Chicago, Edinburgh, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Google's new remote attestation scheme is every bit as terrible as its old remote attestation scheme (permalink) Long before "agentic AI," we had the idea that software would act as your agent on the internet. That's why the old-fashioned technical term for a browser is a "user agent." Your browser acts on your behalf to retrieve information and then show it to you, in the format you choose. It's your agent: <https://pluralistic.net/2024/05/07/treacherous-computing/#rewilding-the-internet> This is a powerful and profound idea. It is because browsers are our "agents" that we...

18. I Am Not a Reverse Centaur

Source: Miguel Grinberg's Blog | Published: Fri, 12 Jun 2026 09:20:39 GMT

Link: <https://blog.miguelgrinberg.com/post/i-am-not-a-reverse-centaur>

About a year ago I wrote on this blog about how coding with LLMs would not work for me, even if there were no ethical or environmental concerns preventing me to use them. I'm not going to repeat the arguments I made that time because my views on the subject haven't changed. What has changed, however, is that the number of contributions I receive on my open source projects has gone up, and nearly all are now made with LLMs. The other day I had a very depressing thought regarding this. All these people who submit drive-by pull requests to my projects are pushing me to spend more and more of my time reviewing and merging code that was extruded by

machines. Cory Doctorow refers to people that perform this function as reverse centaurs . He calls these "frail and vulnerable people being puppeteered by uncaring, relentless machines." Ouch! Am I a reverse centaur now? Is my new purpose as a seasoned software engineer and open source developer to spend my days reviewing LLM code, in spite of having decided that I do not need nor want this technology myself? As you can guess from the title, I'm never going to become a reverse centaur. Let me tell you how I resist the forces that want me t...