

1. AI GPUs probably live longer than three years

Source: seangoedecke.com RSS feed | Published: Mon, 15 Jun 2026 00:00:00 GMT

Link: <https://seangoedecke.com/ai-gpus-live-longer-than-three-years/>

People who think current AI use is unsustainable often rely on the claim that inference GPUs only last "three years at the most" under load 1 . The idea here is that once the AI bubble money drains away, current infrastructure will rapidly become obsolete, and there won't be enough money floating around to buy a whole slate of brand-new GPUs. Inference costs would thus rapidly become way too expensive for current AI products to make any financial sense. Where does this "three years at the most" claim come from? Is it plausible? Sourcing the quote The original Tom's Hardware article quotes this tweet from Tech Fund, an anonymous former PM and tech investor, who quotes an anonymous "GenAI principal architect" at Google as saying "if you have a high utilization rate, then constant high utilization rate for a year or two, I think the lifespan will be three years at most?". This screenshot looks like it was from an interview. What interview? I scrolled back to October 2024 on Tech Fund's Twitter feed and saw a bunch of similarly-formatted screenshots , some of which were cited as coming from Tegus . Tegus is apparently a company with a business model of reaching out to insiders (in this...

2. Why AI hasn't replaced software engineers, and won't

Source: Simon Willison's Weblog | Published: 2026-06-14T23:54:11+00:00

Link:

<https://simonwillison.net/2026/Jun/14/why-ai-hasnt-replaced-software-engineers/#atom-everything>

Why AI hasn't replaced software engineers, and won't Arvind Narayanan and Sayash Kappor take on the question of AI job losses through the lens of a profession that is uniquely suited to AI disruption - software engineering. In this essay, we argue that there is enough evidence to reject the narrative that once AI capabilities reach a certain threshold, it will cause mass layoffs. Given that this is true even in a sector with very few regulatory barriers, most other professions are likely to be even more cushioned. The first good news is that the data still doesn't support the idea that AI is causing mass unemployment. In March 2025, New York became the first U.S. state to add an AI disclosure checkbox to WARN Act filings. In the full first year, more than 160 companies filed WARN notices. Not a single one checked the AI box AI speeds up the typing-code-into-a-computer phase, but it turns out software engineering is about a whole lot more than that: If writing code isn't the bottleneck, what is? The task-breakdown surveys point at things like meetings or debugging. This just leads to more questions: what are developers doing in those meetings and why can't it be done by AI? Won't d...

3. Anthropic Walks Back Policy That Could Have "Sabotaged" AI Researchers Using Claude

Source: Simon Willison's Weblog | Published: 2026-06-11T03:45:49+00:00

Link: <https://simonwillison.net/2026/Jun/11/anthropic-walks-back-policy/#atom-everything>

Anthropic Walks Back Policy That Could Have "Sabotaged" AI Researchers Using Claude Big scoop for Maxwell Zeff at Wired: "We're changing Fable 5's safeguards for frontier LLM development to make them visible." Anthropic said in a statement to WIRED. "We made the wrong tradeoff and we apologize for not getting the balance right." There's been a huge outcry about Anthropic's policy, tucked away in their system card , that Claude Fable/Mythos would identify "requests targeting frontier LLM development" and "limit effectiveness" without notifying the user. It's good news that they're dropping the invisible aspect of this. It would be a whole lot better if they dropped this category of refusals entirely. Update : More details from @ClaudeDevs on Twitter : We're rolling out changes to make Fable 5's safeguards for frontier LLM development visible. Starting this week, flagged requests will visibly fall back to Opus 4.8?the same as our safeguards for cyber and bio. You will see this every time it happens. On the API, any flagged requests will return a reason for their refusal (coming to server-side fallback in the next few days). We wanted to deploy Fable 5 to our users quickly and safely...

4. Setting a custom price for a model in AgentsView

Source: Simon Willison's Weblog | Published: 2026-06-09T21:35:31+00:00

Link: <https://simonwillison.net/2026/Jun/9/agentsview-custom-model-price/#atom-everything>

TIL: Setting a custom price for a model in AgentsView I've been really enjoying AgentsView by Wes McKinney as a tool for exploring my token usage across different coding agents running on my laptop. Claude Fable 5 came out today and wasn't yet included in the pricing database AgentsView uses. I used Fable to reverse-engineer AgentsView and figured out this recipe for setting custom prices. Here's my Claude Fable 5 usage for today so far, plotted by AgentsView as a treemap across my different local projects: Tags: ai , generative-ai , llms , llm-pricing , claude-mythos

5. Apple's WWDC Announcement of the New Apple Intelligence System

Source: Daring Fireball | Published: 2026-06-09T16:50:33Z

Link: <https://www.apple.com/newsroom/2026/06/apple-intelligence-brings-powerful-ai-capabilities-into-everyday-experiences/>

Apple Newsroom: These new capabilities are powered by the next generation of Apple Foundation Models, custom-built in collaboration with Google and its Gemini models for deeply integrated Apple Intelligence experiences. These latest models run on device and on servers using Private Cloud Compute. Every facet of the new Apple Intelligence architecture is built privacy-first, from the latest Apple Foundation Models to the core operating system technologies that integrate these models deep into Apple's platforms. Apple Intelligence uses on-device processing and Private Cloud Compute to help protect users' privacy. Private Cloud Compute gives users access to frontier-level intelligence, while extending the privacy and security of iPhone into the cloud. What's confusing about this Apple-Google partnership is that Google pretty much calls all things AI 'Gemini'. The models are 'Gemini', the assistant is 'Gemini', and the feature integrations are 'Gemini'. So Apple is taking pains to emphasize that they're building atop the Gemini models , not the Gemini assistant . One way to think about it is this. Let's say you're a Google Gemini app user. That's the assistant. Now you start using the...

6. [Sponsor] WorkOS Launches auth.md - an Open Protocol for Agent Registration

Source: Daring Fireball | Published: 2026-06-09T04:23:05Z

Link: https://youtu.be/Dqp_b8GHLXU?t=1074

Sign-up forms were built for humans in browsers, so how do AI agents programmatically register with services? Enter auth.md. By exposing a single, machine-readable Markdown file at your service root, AI agents can dynamically discover your OAuth Protected Resource Metadata, parse required scopes, and authenticate seamlessly. With native support in WorkOS AuthKit, you can now implement this protocol out of the box, giving AI tools a standardized, secure way to log into your application. Read the auth.md docs .

7. Remember When Chrome Went Bad on MacOS?

Source: Daring Fireball | Published: 2026-06-04T18:10:19Z

Link: <https://chromeisbad.com/>

Loren Brichter , back in 2020: Short story: Google Chrome installs an updater called Keystone on your computer, which is bizarrely correlated to massive unexplained CPU usage in WindowServer (a system process) [1] , and made my whole computer slow even when Chrome wasn't running . Deleting Chrome and Keystone made my computer way, way faster, all the time . Long story: I noticed my brand new 16" MacBook Pro started acting sluggishly doing even trivial things like scrolling. Activity Monitor showed nothing from Google using the CPU, but WindowServer was taking ~80%, which is abnormally high (it should use Doing all the normal things (quitting apps, logging out other users, restarting, zapping PRAM/SMC, etc) did nothing, then I remembered I had installed Chrome a while back to test a website. I deleted Chrome, and noticed Keystone while deleting some of Chrome's other preferences and caches. I deleted everything from Google I could find, restarted the computer, and it was like night-and-day. Everything was instantly and noticeably faster, and WindowServer CPU was well under 10% again. Not all Mac users, but many, found that just having Chrome installed slowed down their Macs

dramati...

8. How can I schedule work on a thread pool with low latency?

Source: The Old New Thing | Published: Fri, 12 Jun 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260612-00/?p=112417>

The thread pool is designed for throughput, not latency. The post How can I schedule work on a thread pool with low latency? appeared first on The Old New Thing .

9. Linux distribution packaging and third party 'package' systems

Source: Chris's Wiki :: blog | Published: 2026-06-14T03:22:27Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/DistrosAndOtherPackageManagers>

If you look at it from the right angle, the existence of third party package systems that sit on top of Linux distributions is rather odd. After all, these distributions already have packaging systems, yet here people are, ignoring them and writing new ones. A while back on the Fediverse, I said something on this general topic: I have feelings and some of the feelings are that everyone screwed this up, for actually natural reasons. I don't think any distro solved the 'how do we let people easily build and deploy software on us' problem, and so programmers do what programmers do and built themselves another layer of indirection to solve their immediate problem. (I mean, ship an entire container running a web server to do what you could do with a CGI or some PHP files or ... this is my face, etc.) (There are two sides of this, the system side (what I was talking about) and the user side .) On the one hand this is perfectly natural for Linux distributions to do. They built their package system to manage their own components, with all of the features that are important for that , not to make it easy for programmers to create a package for something that was the contents of '/opt/' an...

10. Did Frank Sinatra really think "Something" was a Lennon/McCartney song?

Source: Terence Eden's Blog | Published: Sun, 14 Jun 2026 11:34:16 +0000

Link: <https://shkspr.mobi/blog/2026/06/did-frank-sinatra-really-think-something-was-a-lennon-mc-cartney-song/>

Read enough articles about The Beatles and you'll repeatedly hit the claim that Frank Sinatra frequently introduced his cover of George Harrison's "Something" as his "favourite Lennon & McCartney number." Much like the misquote about Ringo not being the best drummer in The Beatles, I think this might be one of those semi-apocryphal lines which has taken on a life of its own. Here's what Paul?

11. Plugins case study: Pluggy

Source: Eli Bendersky's website | Published: 2026-06-13T20:21:00-07:00

Link: <https://eli.thegreenplace.net/2026/plugins-case-study-pluggy/>

Recently I came upon Pluggy , a Python library for developing plugin systems. It was originally developed as part of the pytest project - known for its rich plugin ecosystem - and later extracted into a standalone library. You're supposed to reach out for Pluggy if you want to add a plugin system ?

12. from hookswitch to grave

Source: computers are bad | Published: 14 Jun 2026 00:00:00 UT

Link: <https://computer.rip/2026-06-14-hookswitch-to-grave.html>

Through decades of consolidation, reorganization, and divestiture, AT&T left a famously complicated corporate history. One of the greatest enterprises in American history, arguably the greatest enterprise, AT&T has often rivaled the federal government in the size of its budget and workforce. One of the reasons, as we well know today, was monopolization and its close relative vertical integration. AT&T was the telephone system, or at least aspired to be, and for decades the meaning of "Universal Service" was that the service was designed, built, and operated by AT&T?universally. While AT&T's tangled origins are fertile ground for the historian, they also obscure many of the early stories of telephone history. Much of the work of the early independent telephone industry has been lost in the voluminous achievements of

AT&T. Even very basic facts become obscure. For example, who invented the telephone? Well, we all know the answer: Alexander Graham Bell. We have mostly forgotten that, at the time, this was a hotly contested question. One of the most prominent alternate claimants to the title was a man named Elisha Gray, today immortalized as the "Gray" in electrical distributor "Grayb...