

1. Cloudflare CAPTCHA on at least one ampersand

Source: Simon Willison's Weblog | Published: 2026-06-16T00:21:36+00:00

Link: <https://simonwillison.net/2026/Jun/16/captcha-on-at-least-one-ampersand/#atom-everything>

TIL: Cloudflare CAPTCHA on at least one ampersand I'm using Cloudflare's CAPTCHA (they call it a "Web Application Firewall > Custom rules > Managed Challenge" these days) to prevent crawlers from aggressively spidering my faceted search engine on this site, but I got fed up of even simple ?q=term searches triggering the challenge. After some mucking around with Claude Code it turns out you can register the following rule instead, so the CAPTCHA only kicks in for search URLs containing at least one ampersand: (http.request.uri.path wildcard r"/search/*" and http.request.uri.query contains "&") And now /search/?q=lemur works without triggering a CAPTCHA! Also included: notes on trying out the Cloudflare MCP with Claude Code , though it turned out not to be able to edit the rules in question so I had Claude Code switch to the Cloudflare API instead. Tags: captchas , cloudflare , model-context-protocol , claude-code

2. datasette-agent 0.3a0

Source: Simon Willison's Weblog | Published: 2026-06-15T17:19:27+00:00

Link: <https://simonwillison.net/2026/Jun/15/datasette-agent/#atom-everything>

Release: datasette-agent 0.3a0 New tool, execute_write_sql , which requests user approval and then writes to a database - taking user permissions into account. #27 I added a mechanism for asking user approval in datasette agent 0.2a0 . The new execute_write_sql tool can now prompt the user for all kinds of useful operations. Here's an example where I add some pelican sightings to my pelican_sightings table: The new version also enhances the datasette agent chat terminal mode to support approvals, and adds several new options including --unsafe mode for auto-approving them: datasette agent chat can execute tools that require user approval. #30 Three new options for datasette agent chat - --root to run as root, --yes to approve all ask user questions, and --unsafe for both. Tools can now provide plain text alternatives to HTML, for display in the datasette agent chat CLI. #31 The datasette agent chat content.db -m gpt-5.5 --unsafe command can now be used to chat directly with a specific database and directly modify it through prompts like "create a notes table", "add a note about X" etc. Tags: projects , ai , datasette , annotated-release-notes , generative-ai , llms , llm-tool-use...

3. "They screwed us": Personality clashes sent Anthropic's models offline

Source: Simon Willison's Weblog | Published: 2026-06-15T14:57:33+00:00

Link: <https://simonwillison.net/2026/Jun/15/axios-clashes-anthropics/#atom-everything>

"They screwed us": Personality clashes sent Anthropic's models offline Lots of "source familiar with the administration's thinking" and "source close to Anthropic" in this Axios piece, which is the best collection of behind-the-scenes gossip I've seen about the US government export control Mythos/Fable story so far. Logan Graham (I lead the Frontier Red Team at Anthropic), Dave Orr (Head of Safeguards, previously a Director of Engineering at Google DeepMind), and blog favorite Nicholas Carlini are reported to be meeting with the Commerce Department today in D.C. Good luck to them! (I just noticed Logan was "Special Adviser to the Prime Minister" in the Boris Johnson era, covering AI, science, and technology policy - so significant political experience.) This closing notes doesn't give me much optimism that we'll be getting Fable back any time soon: The bottom line : One option is to make sure Anthropic's models can't be jailbroken ? though perfect jailbreak resistance may be impossible. Absent that, a source familiar with the administration's thinking said it may simply come down to an attitude fix where, instead of feeling dismissed, "everyone feels safe, secure and happy." Thi...

4. The European Commission Ruled Months Ago That Google's Integration of Gemini in Android Violates the DMA

Source: Daring Fireball | Published: 2026-06-15T18:58:22Z

Link: <https://arstechnica.com/ai/2026/04/europe-could-force-google-to-open-android-to-other-ai->

assistants/

Ryan Whitham, writing for Ars Technica back in April: European regulators are proposing several broad changes to the way AI tools operate on Android phones. Some of this is straightforward, like allowing third-party AI tools to be invoked system-wide via hot words or button presses. This might also include allowing AI tools to view screen context when the user opens them. Context also extends to allowing alternative AI systems to access local data to generate proactive suggestions and summaries. The report actually describes something that sounds like Google's Magic Cue, which relies on Gemini to offer suggestions based on your activity. Google has also started experimenting with allowing AI to control certain apps. As we saw when this feature debuted on the Galaxy S26, Gemini is currently pretty bad at using apps on your behalf. The commission wants to explore allowing other AI services to autonomously control installed apps and system features on Android phones. Maybe someone else could do better? Maybe! But also maybe it's a bad idea for complex system architecture design to come from non-technical government bureaucrats. One of these maybes strikes me as a lot more likely than...

5. WorkOS Launches Auth.md ? an Open Protocol for Agent Registration

Source: Daring Fireball | Published: 2026-06-15T17:53:50Z

Link:

https://workos.com/auth-md?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q22026

My thanks to WorkOS for sponsoring DF last week to promote Auth.md, their new open protocol for AI agent registration. Sign-up forms were built for humans in browsers, so how do AI agents programmatically register with services? That's the question Auth.md aims to answer. By exposing a single, machine-readable Markdown file at your service root, AI agents can dynamically discover your OAuth Protected Resource Metadata, parse required scopes, and authenticate seamlessly. Markdown, baby. Who'd have thunk it? With native support in WorkOS AuthKit, you can now implement this protocol out of the box, giving AI tools a standardized, secure way to log into your application. Read the Auth.md docs, and watch its on-stage introduction at the MCP Night: Agent Night keynote. ?

6. ?Anthropic's Safety Superpower?

Source: Daring Fireball | Published: 2026-06-15T17:18:46Z

Link: <https://stratichery.com/2026/anthropics-safety-superpower/>

Ben Thompson, in his weekly free column at Stratechery: On one hand, I actually don't begrudge Anthropic not wanting to help its competitors; on the other hand, what should be blisteringly clear is that Anthropic does not think that anyone else other than them should even be making frontier LLMs. What makes this policy all the more remarkable is the fact that it was enacted only two months after Anthropic had that dispute with the Department of War: the latter wanted to use Claude for any legal use, while the former wanted more stringent controls around surveillance and autonomous weapons. What this degradation represented was both the capability and willingness of Anthropic to silently alter its models to achieve its policy preferences. In other words, Anthropic willfully validated some of its critics' worst fears in terms of being a supply chain risk. The broader takeaway from that previous episode, however, is that Anthropic believes that they are the ones who should have final say over how Anthropic is used; given that they think only they should be developing leading edge AI, they by extension think that only they should have final say over AI generally. When you further comb...

7. Testing pentagonal numbers

Source: John D. Cook | Published: Mon, 15 Jun 2026 23:57:54 +0000

Link: <https://www.johndcook.com/blog/2026/06/15/testing-pentagonal-numbers/>

The nth pentagonal number P_n is the number of dots in diagrams like those below with n concentric pentagons. We have the formula $P_n = (3n^2 + n)/2$ where n is a positive integer. If n is an integer but not positive, the equation above defines a generalized pentagonal number. If you're given an n, you can easily compute P_n . [?] The post Testing pentagonal numbers first appeared on John D. Cook.

8. Quaternion Rotations, Claude, and Lean

Source: John D. Cook | Published: Mon, 15 Jun 2026 19:31:38 +0000

Link: <https://www.johndcook.com/blog/2026/06/15/quaternions-claude-lean/>

I got an email message this afternoon reporting a typo in a blog post from about a year ago on converting between quaternions and rotation matrices [1]. The email said exactly where the typo was, but I decided to see whether Claude would find it. Specifically, I prompted Sonnet 4.6 Medium with the following. Write [?] The post Quaternion Rotations, Claude, and Lean first appeared on John D. Cook .

9. Writing Prolog with ChatGPT

Source: John D. Cook | Published: Mon, 15 Jun 2026 17:22:43 +0000

Link: <https://www.johndcook.com/blog/2026/06/15/writing-prolog-with-chatgpt/>

A few days ago I wrote about using Claude to solve a chess puzzle by writing Prolog code. This morning I tried a similar chess puzzle with ChatGPT. The task is to place a queen, king, rook, bishop, and knight on a 4 by 4 chessboard so no piece attacks another. Of course there?s not [?] The post Writing Prolog with ChatGPT first appeared on John D. Cook .

10. Arp 185:

Source: Maurycy's Blog | Published: Sun, 14 Jun 2026 00:00:00 +0000

Link: <https://maurycyz.com/astro/arp185/>

North is left. 0.35"/pixel (7.7' by 20' field). FWHM = 5.5" Arp filled this under "galaxies with narrow filaments", but I don't really see why. It's a barred sprial with irregular arms and a faint extended disk. That being said, the galaxy's bar is slightly misaligned with the core: which would sudget that it might have experienced some graviational interaction in the past. The main galaxy is at z=0.0045, and the faint background galaxy in the lower left is UGC 10509 at z=0.055. It's actually a spiral with a really intresting structure, but it's too small (and dim) to see in my image. dt {color: var(--em-color)} Total exposure time: 274 * 30 seconds = 2.2 hours. Telescope: C9.25 (230mm, f/10, fl=2300mm) Camera: IMX533 (16mm diagonal, square, color) Processing: Callibration (dark + flat) Stacking (average w/ rejection) White balance and background subtraction Asinh stretch Tone curve Crop Taken during a night with exceptionally bad seeing, resulting in a FWHM of 5" instead of the usual 3". Not that I ever get good seeing, presumably due to nearby heat sources: My telescope was set up between a asphalt road and asphalt roof. Both of these get super hot during the day and retain tha...

11. Making glass-to-metal seals for homemade vacuum tubes.

Source: Maurycy's Blog | Published: Sat, 13 Jun 2026 00:00:00 +0000

Link: <https://maurycyz.com/projects/glass/1/>

This page discusses sealing metal through borosilicate/lab glass: other chemistries behave quite differently. When making vacuum tubes, the glass is actually the easy part: premade tube stock of almost any size is easily available. Heating the end of such a tube softens the glass and allows surface tension to close it off. I used a rotary vane pump to remove all the air from the tube and heated the middle, which the atmosphere crushed to create a sealed-off ampule. Because glass is practically impermeable, it will retain that vacuum for a very long time, which can be shown by bringing it close to high-voltage AC (like a tesla coil): This glow is due to residual air being ionized, but the fuzzy appearance indicates that the vacuum is good enough to work in a triode or similar device. For those, the capacitive coupling trick won't work: I'll need to make electrodes that pass through the glass without letting air in. This is a lot harder than it might appear. Copper's red oxide bonds very well to glass . In fact, the bond is stronger than the bulk glass: when it breaks, there's always a thin layer of glass left stuck to the metal. Along with it's excellent electrical properties, it's...

12. Systemd-resolved and sticking (or not) to what distributions do

Source: Chris's Wiki :: blog | Published: 2026-06-15T02:51:25Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/SystemdResolvedOrNot>

We have a long established approach for configuring our Linux systems, and specifically our Ubuntu systems (since these days the only Linux we run is Ubuntu). Part of that has been that we use a standard `/etc/resolv.conf` that points to our local DNS resolvers (which have our split horizon DNS setup). However, as people may have noticed, Linux distributions these days are moving to `systemd-resolved` instead of a straight `/etc/resolv.conf` setup. There are good reasons for this switch in general. `Systemd-resolved` can cope with a variety of situations and problems that trip up the standard `resolv.conf` setup and it cooperates better with other programs.

Almost none of these apply in our particular situation for our own servers, which have a static network configuration, a static DNS resolution search domain, and our DNS resolvers already handle all of the split resolution (and we had to make this work well before `systemd-resolved` existed). (The one aspect of `systemd-resolved` that matters to us is that if you have multiple DNS resolvers, as we do, `resolved` will rapidly fail over from one that isn't responding to one that is. `Resolved` can also cache answers locally but for us this isn't a...

13. [RSS Club] What happens to old posts?

Source: Terence Eden's Blog | Published: Mon, 15 Jun 2026 11:34:34 +0000

Link: <https://shkspr.mobi/blog/2026/06/rss-club-what-happens-to-old-posts/>

Welcome to RSS Club! These posts are only available to RSS and Atom subscribers. You can read more about the idea at Dave Rupert's site. I recently received an email from a distraught

reader: I was going through my recent bookmarks and I found

<https://shkspr.mobi/blog/2026/04/rss-club-for-wordpress> which I had clearly saved to reference again later. However, as with the nature of RSS Club?, I c?

14. AI's Brokenomics

Source: Ed Zitron's Where's Your Ed At | Published: Mon, 15 Jun 2026 19:28:23 GMT

Link: <https://www.wheresyoured.at/brokenomics/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA , Anthropic and

15. JAX: commitment issues

Source: Giles' blog | Published: Mon, 15 Jun 2026 21:30:00 +0000

Link: <https://www.gilesthomas.com/2026/06/jax-commitment-issues>

Imagine you have JAX code like this, and run it on a machine with CUDA set up: `key = jax . random . key ( 42 )` `cpu0 = jax . devices ( "cpu" )[ 0 ]` with `jax . default_device ( cpu0 )`:
`array = jax . random . randint ( key , ( 530640 , 6 , 1024 ) , 0 , 50_000 , dtype = jax . numpy . uint16 )` `array . block_until_ready ()` `item = array [ 0 ]` `item . block_until_ready ()` We're creating a big array, blocking until it's ready (JAX is asynchronous, so this makes sure that it's actually finished creating it), then getting the first item, and as a belt-and-braces thing making sure that that is ready too. How long do you think those last two lines -- a simple retrieval of a 6 x 1024 array from a larger one -- will take? Some tiny fraction of a second would seem reasonable. But running it on my machine just now, the answer is a bit of a surprise: just over 5 seconds. And if you try to get `array[1]` immediately afterwards, it still takes about 1.2s. Further lookups into array consistently take more than a second -- so while the larger initial number might be something to do with setup -- maybe internal stuff being JITted -- that's clearly not the whole story. Something is making these seemingly-sim...

16. The Sword Juggling Fallacy

Source: Hugo Tunius - Blog | Published: 2026-06-15T00:00:00+01:00

Link: <https://hugotunius.se/2026/06/15/the-sword-juggling-fallacy.html>

Sword juggling is simple. I'll prove it to you by explaining it. First, acquire three ? or, if you feel ambitious, four swords. Divide the swords between your hands. Starting with one hand, the one holding two swords if juggling with three, toss a sword into the air. As this sword is about to reach your other hand, toss the sword in that hand and catch the falling sword by, and this is important, its handle, not the blade. Then repeat this process as the airborne sword

reaches the hand you started with. Now keep doing this, ensuring that you do not cut your fingers off in the process. To make the subtext, well, text. This is an allegory about the C programming language, although it applies to others too. C being simple is a common idea, especially made in discussion about and as contrast to Rust. It's true that C is simple. The language is small and the standard is, relatively, short. Just like sword juggling is simple, C being simple doesn't mean programming in C is easy. In fact it's very hard, especially if you care about your program being correct. Simple is not easy. After learning the simple C language, a world of painful experiences await the aspiring C programmer. With eac...

17. Weekly Update 508

Source: Troy Hunt | Published: Mon, 15 Jun 2026 04:55:47 GMT

Link: <https://www.troyhunt.com/weekly-update-508/>

Light switches. How on earth is it so hard to find decent light switches?! It sounds ridiculous until you actually spend enough time looking for ones that meet two simple criteria: Aren't stateful (switch is up or down, has to be push-button) Looks good Now, I'

18. Pluralistic: AI and amateurism (15 Jun 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 15 Jun 2026 15:40:29 +0000

Link: <https://pluralistic.net/2026/06/15/vernacular/>

Today's links AI and amateurism: When is generative content vernacular? Hey look at this: Delights to delectate. Object permanence: Disney characters x clean underwear; Transparent Pontiac; Makers v dog with LED collar; Microsoft buys LinkedIn; Legitimate greatness. Upcoming appearances: LA, Menlo Park, Toronto, NYC, Philadelphia, Chicago, Edinburgh, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. AI and amateurism (permalink) Over the weekend, I did an interview about my forthcoming book The Reverse Centaur's Guide to Life After AI (a book about being a better AI critic), and the interviewer said she was surprised that I wasn't an AI booster, based on my demographics and work history: <https://us.macmillan.com/books/9780374621568/thereversecentaursguidetolifeafterai/> I could see where she was coming from. I encountered computers in the mid-seventies, as a small child. My first computer was a CARDIAC, a working, Turing-complete, mechanical computer made entirely of cardboard, that I spent endless hours with: <https://www.instructables.com/CARDIAC-CAR...>

19. A brief history of KV cache compression developments

Source: Martin Alderson | Published: Mon, 15 Jun 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/a-brief-history-of-kv-cache-compression-developments/?utm_source=rss&utm_medium=rss&utm_campaign=feed

How KV cache compression - from MQA and GQA to MLA and linear-attention hybrids - quietly unlocked the long context windows that make modern agentic LLMs possible.

20. Things that made me think: Open Source trust relationships, knowledge without provenance, and theory building

Source: Tom Renner on My place to put things | Published: Mon, 15 Jun 2026 00:00:00 +0000

Link: <https://tomrenner.com/posts/ttmmt-4/>

This series is a place to collect interesting things I've seen, read, or heard, along with some brief thoughts (often incomplete and/or inconclusive) that they provoked. AI Agent Lands PRs in Major OSS Projects, Targets Maintainers via Cold Outreach ? socket.dev Back in February this report landed in a security blog, and I found it fascinating. As is usual for me, it's at the intersection of human behaviour with technological automation that the real interesting problems lie. The whole industry uses GitHub profiles as a proxy for a CV. Now agents can build a credible looking one in hours or days, and it results in our whole system of reputation and trust being borked. Once again, the thing you should be optimising for is trust . I wonder what new systems we need to build to deal with the old ones being broken. Where Provenance Ends, Knowledge Decays ? Jessica Talisman My first job was building information management systems,

and for those of you that haven't worked in the industry, let me tell you: knowledge architecture is fascinating. We should all be hiring librarians to help us with these problems. The core question posed here is: without provenance how can users interrogate t...

21. Maybe it's time for lots of little indie AIs to take over

Source: Anil Dash | Published: 2026-06-15T00:00:00Z

Link: <https://anildash.com/2026/06/15/indie-AI-takeover/>

[T]here can be alternatives. What we can imagine is, rather than the ChatGPT killer, a lot of different little AIs from little responsible players. That's me, in The Guardian a few days ago, trying to distill a message that I've been trying to get out as broadly as possible for quite a while now. It's sort of like hoping a comet will take out the major AI players and a bunch of smaller new players will be the smarter, better-adapted mammals that take their place instead. We're in another one of those big inflection points for AI. Trump administration policymakers for AI suspended access to Anthropic's newest product . All of these policymakers have a web of investments in competing players ? including SpaceX, which is about to IPO ? and the corruption and grift of this cohort are so extensive that it's impossible to judge what the actual risks and reality are around any of these platforms or technologies, since no one involved is an honest broker. It's a shame More broadly, there's been the widespread pushback against AI culturally, one that is undeniably strongest amongst those who were born in this century. But the adoption patterns and usage data show that even younger people...

22. EU & Civil Society need to progress on Digital Autonomy

Source: Bert Hubert's writings | Published: Mon, 15 Jun 2026 15:00:00 +0200

Link: <https://berthub.eu/articles/posts/eu-civil-society-need-progress-digital-autonomy/>

By now (happily) everyone wants to talk about digital autonomy, although some parties insist we talk about sovereignty. Fine. tl;dr: The discussions on digital autonomy are now going round in circles. Civil society and think tanks want to contribute, and are well placed to do so. In this piece, I urge everyone to look further ahead, beyond legislation and talking about European values. The road to digital sovereignty is very long, and we need to make progress along that entire road, which includes things further afield than what we are discussing now.