

1. datasette-tailscale 0.1a0

Source: Simon Willison's Weblog | Published: 2026-06-16T16:18:20+00:00

Link: <https://simonwillison.net/2026/Jun/16/datasette-tailscale/#atom-everything>

Release: datasette-tailscale 0.1a0 A very experimental alpha plugin which lets you do this:

```
datasette tailscale mydata.db \ --ts-authkey tskey-auth-xxxx --ts-hostname datasette-preview
```

This starts a localhost Datasette server with a Tailscale sidecar that connects it to your Tailnet, such that <http://datasette-preview/> serves Datasette. It's using the Python bindings for the experimental tailscale-rs library. I filed an issue asking if there's a cleaner way of setting up the proxy mechanism. Tags: datasette , tailscale

2. Quoting Georgi Gerganov

Source: Simon Willison's Weblog | Published: 2026-06-16T16:04:59+00:00

Link: <https://simonwillison.net/2026/Jun/16/georgi-gerganov/#atom-everything>

I can 100% attest to the fact that Qwen3.6-27B is a very capable local model for coding tasks. Over the last month and a half I've been using it almost daily, either on my M2 Ultra or on my RTX 5090 box. I use it for small mundane tasks at ggml-org - nothing really impressive, but definitely a helpful tool for a maintainer. I think I would be using it much more, if I didn't have to spend a lot of my time on reviewing PRs. Currently, I have a very lightweight harness - the pi agent with everything stripped (`pi -nc --offline`) and a short system prompt to align it a bit with my style. ? Georgi Gerganov , Hacker News comment on Running local models is good now by Boykis Tags: georgi-gerganov , llms , ai , generative-ai , pi , ai-assisted-programming , local-llms , qwen , coding-agents

3. The Fable 5 Export Controls Harm US Cyber Defense

Source: Simon Willison's Weblog | Published: 2026-06-16T05:20:29+00:00

Link: <https://simonwillison.net/2026/Jun/16/fable-5-export-controls/#atom-everything>

The Fable 5 Export Controls Harm US Cyber Defense I quoted The Atlantic quoting Kate Moussouris earlier, when I should have gone straight to the source. Here she is confirming that the "jailbreak" that got Claude Fable 5 banned under an export control really was "fix this code": The researchers took open-source code with known CVEs, plus new code with deliberately planted vulnerabilities, and asked Fable 5, Mythos, and Opus to ?review the code for security issues.? Fable 5 refused. They then asked the models to ?fix this code? and, through a multistep and manual process, turned the output into scripts that test the patches. As Kate points out, this is absurd. Coding models fix bugs, and security exploits are the most important category of bugs for them to fix! Defenders need to be able to ask AI to fix the bugs in a file, explain why the fix matters, and write tests that confirm the patch works. That is not a guardrail bypass. It is the most valuable thing an AI model can do for defensive security: executing the find, fix, and test loop defenders run every day. [...] The prompts worked because they were defensive requests, and that capability cannot be removed without making the m...

4. Debugging on Prod

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Tue, 16 Jun 2026 20:18:00 GMT

Link: <https://idiallo.com/blog/debugging-on-prod>

The worst type of bug is one that only happens on prod. And only on prod. If you checked this blog in the past few weeks, you might have encountered a big fat 500 error. I'd had the same design for 10 years, and I wanted something fresh. But who can redesign without also improving the underlying code? I deleted a whole bunch of things: old templates that were never used, `post.new7.old.php` , a pile of unused CSS. I just had to. I deployed a first version and all the pages worked just fine. But then I got cocky. I decided to also improve the underlying code using GitHub Copilot. I was vigilant at first, reviewing every single line of generated code. None of it was complex really, just refactoring functions and the like. But along the way, I

got lazy. I let the AI update deprecated functions on its own. The next time I deployed, the website returned a 500 error. When I checked the logs, nothing came back. No errors. I looked at running processes and noticed several PHP processes pinned at 100%. I reverted the code, but the server was still stuck. I restarted the web server, restarted PHP-FPM, and neither helped. The only thing that worked was restarting the whole machine. I ran the s...

5. The time the x86 emulator team found code so bad that they fixed it during emulation

Source: The Old New Thing | Published: Mon, 15 Jun 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260615-00/?p=112419>

Offensive content in the eyes of a software engineer. The post The time the x86 emulator team found code so bad that they fixed it during emulation appeared first on The Old New Thing .

6. Checking In on the iOS Continental Fun-Gap Drift

Source: Daring Fireball | Published: 2026-06-17T00:50:18Z

Link: https://daringfireball.net/2024/09/ios_continental_drift_fun_gap

Yours truly, in September 2024, expressing skepticism that "European iPhones are more fun now?". Meanwhile no one in the EU will get Apple Intelligence or iPhone Mirroring, both of which features are very useful, and, dare I say, quite fun. Should we judge how much fun each side of the continental divide is having by how much fun they theoretically could be having, or by how much fun they are having? As it stands, the fun side is not the EU. But hope springs eternal. Here we are two years later and I think the answer is more clear than ever which side of the continental divide is more fun. It's not the EU. EU users still don't have iPhone Mirroring and until and unless the European Commission changes its interpretation of the DMA, they likely never will. It's a great feature. Apple Intelligence, as we knew it until last week, eventually came to the EU , about six months after it shipped for the rest of us. One can reasonably argue that EU iPhone and iPad users didn't miss much during those six months. And that there hasn't been that much to enjoy since Apple Intelligence debuted in the EU in iOS 18.4. That changed last week with the introduction of the first beta release of iOS 27....

7. New in the App Store: Personalized Recommendations

Source: Daring Fireball | Published: 2026-06-16T16:12:31Z

Link:

<https://techcrunch.com/2026/06/09/apples-app-store-rolls-out-personalized-recommendations/>
Sarah Perez, writing for TechCrunch: This week, Apple announced a series of discovery features that will personalize app recommendations based on users' interests and behavior, providing a new way for developers to have their app discovered. At Apple's Worldwide Developers Conference (WWDC), the iPhone maker introduced Personalized Collections in the App Store, which will showcase recommendations tailored to the individual. These will also include new "App Notes" that explain why the specific apps were recommended to you. Starting this week, you'll find these new personalized suggestions in various places in the App Store, including the Apps or Games tab or on the Search tab. Security research critics Mysk, posting on Twitter/X (XCancel link), report that the App Store app seemingly sends analytics usage data to Apple with everything you do in the App Store app, including exactly what you type, character-by-character " and that this isn't for search suggestions, but for analytics. (Via Michael Tsai .) "

8. [Sponsor] Mux " Video for Developers

Source: Daring Fireball | Published: 2026-06-16T02:33:37Z

Link: https://www.mux.com/?utm_campaign=fireball&utm_source=DF

Video is a boatload of data. Every video file in your product contains audio, objects, and scenes that most stacks can't read or access. Mux Robots turns that data into video intelligence. Configure your video workflows once, and they run automatically on every new upload: ask questions, summarize, find key moments, and more. No asset webhooks or self-hosted glue code needed. Mux is video infrastructure for developers, trusted by Synthesia, Shopify, and the U.S. Soccer Federation. Start building for free. Use code FIREBALL at signup for an extra \$50 credit. "

9. Partial fraction decomposition

Source: John D. Cook | Published: Tue, 16 Jun 2026 17:56:27 +0000

Link: <https://www.johndcook.com/blog/2026/06/16/partial-fraction-decomposition/>

Nearly everyone who has seen partial fraction decomposition was introduced to it as a way to compute integrals. If $P(x)$ and $Q(x)$ are polynomials, then you can break their ratio $P(x)/Q(x)$ into a sum of terms that can each be integrated in closed form. As with most topics in a calculus class, partial fractions go by in [?] The post Partial fraction decomposition first appeared on John D. Cook .

10. Three examples suffice

Source: John D. Cook | Published: Tue, 16 Jun 2026 11:14:01 +0000

Link: <https://www.johndcook.com/blog/2026/06/16/three-examples-suffice/>

You can't prove a theorem by just checking a few examples. Except sometimes you can. A few weeks ago I wrote Pentagonal numbers are truncated triangular numbers. In a nutshell, if the pentagonal numbers are defined by $P_n = (3n^2 - n)/2$ and the triangular numbers by $T_n = (n^2 + n)/2$ then $P_n = T_{2n-1}$. Here's a visualization [?] The post Three examples suffice first appeared on John D. Cook .

11. What the new Linux NFS mount option 'fatal_netrerrors' is

Source: Chris's Wiki :: blog | Published: 2026-06-16T02:08:53Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/NFSMountOptionFatalNetrerrors>

If you have NFS mounts on a client using a sufficiently recent kernel, such as that shipped with Ubuntu 26.04 LTS, and you inspect `/proc/mounts` (or `'mount -t nfs'`), you'll probably discover a new NFS mount option that is listed for your mounts. For example: `[...]/h/281 /h/281 nfs4 rw,[...],fatal_netrerrors=none,[...]` You might be curious what this option is, especially if you have code that parses NFS (v4) options and sorts options into ones you care about and ones you don't (as we do). Unfortunately, as I write this about a year after the option was added to the Linux kernel, 'fatal_netrerrors' isn't documented in any manual page that I can readily find. As far as I know, the only accessible documentation for it is in the commit that added it, which I'm just going to quote: NFS: Add a mount option to make ENETUNREACH errors fatal If the NFS client was initially created in a container, and that container is torn down, there is usually no possibility to go back and destroy any NFS clients that are hung because their virtual network devices have been unlinked. Add a flag that tells the NFS client that in these circumstances, it should treat ENETDOWN and ENETUNREACH errors as fatal...

12. Would you like a drainer served at the very top of DuckDuckGo?

Source: tim.sh | Published: Tue, 16 Jun 2026 13:34:35 GMT

Link: <https://timsh.org/drainers-at-the-top-of-duckduckgo/>

Probably no, right? That would be very dangerous, especially if the phishing website serving it is an exact copy of the thing you were looking for. Well that's what happened to me! TL;DR I found a fake transcan blockchain explorer duplicate on the #1 position in the

13. Two Way TV - product photos of 1997's hottest gadget

Source: Terence Eden's Blog | Published: Tue, 16 Jun 2026 11:34:18 +0000

Link: <https://shkspr.mobi/blog/2026/06/two-way-tv-product-photos-of-1997s-hottest-gadget/>

Back in the late 1990s, I did a brief stint of work experience at the BBC. One of the most memorable moments was sitting in on a meeting about early forms of Interactive TV. I saw a demo of "Two Way TV". A flimsy grey box which (somehow) integrated with your OnDigital TV Box and connected to a server via a modem. If you were watching "Who Wants To Be A Millionaire" you could play along at home,?

14. How Open Source Projects Change Hands

Source: Andrew Nesbitt | Published: 2026-06-16T10:00:00+00:00

Link: <https://nesbitt.io/2026/06/16/how-open-source-projects-change-hands.html>

Dumb Ways for an Open Source Project to Die listed the ways projects end up dead, and most of the entries describe a moment where maintainership should have moved to someone else and didn't. This is the other, shorter inventory: the mechanisms that exist for a project to change hands, and who can trigger each one. The maintainer decides Chosen successor. The maintainer picks a person and hands over the keys. This is the model everyone imagines when they say succession, and it has almost no supporting infrastructure: it's a private conversation followed by some permission changes. The vetting is whatever the departing maintainer feels like doing, which is how event-stream happened: he emailed me and said he wanted to maintain the module, so I gave it to him. The xz takeover was the same mechanism worked patiently, with sock-puppet users manufacturing the pressure on an overworked maintainer to hand over. CPAN is the one registry that gives this stage a name: a module whose permissions list the pseudo-user HANDOFF has a maintainer looking for someone to take over, recorded in the same machine-readable permissions file as every real owner. The successor setting. GitHub has had acco...

15. 10Gb/s Ethernet: switching to a Broadcom SFP+ module

Source: Giles' blog | Published: Tue, 16 Jun 2026 18:45:00 +0000

Link: <https://www.gilesthomas.com/2026/06/10g-ethernet-switching-to-broadcom-sfp-plus>

Back in April, I upgraded my home LAN to 10Gb/s. The in-wall cabling is CAT-6 or similar, so I had to use 10GBASE-T. Now, the router I'm using, and the switch in my study, provide 10Gb/s through SFP+ cages; that meant that they needed 10GBASE-T SFP+ modules in order to connect. That kind of module is known to run hot -- sometimes too hot to actually work. The modules in reggie, the router, appeared to be running OK (see the linked post above for charts), but the one in nigel, the study switch, was a worrying 93C. I tried sticking some mini-heatsinks on it, which seemed to help a bit. But the weather got warmer, and eventually the module overheated. I lost access to the Internet from the study, and checking the metrics showed me this: You can see that it's "flapping": the temperature gets up to a level where the module shuts itself down for its own protection -- about 95C, I think -- and then when it has recovered, it switches on again, the temperature rises, and the process repeats. I was able to work around the problem by switching on the air conditioning in the study. But normally I only have it on when I'm in there, and keeping aircon on 24/7 just to keep the network workin...

16. Exclusive: OpenAI Losses Increased Nearly 8X in 2025, With Spending Hitting \$34 Billion

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 16 Jun 2026 03:58:20 GMT

Link: <https://www.wheresyoured.at/exclusive-openai-financials/>

Soundtrack: In Flames - Colony To further support my independent journalism, please subscribe to my premium newsletter. It's \$7 a month or \$70 a year. If you're subscribed to the free newsletter and logged in, you should see at the bottom right hand corner of your screen

17. Lean, not backpressure

Source: Entropic Thoughts | Published: Tue, 16 Jun 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/lean-not-backpressure>

Lucas Costa has written a good article on how to build systems that can handle code-generating robots. Unfortunately, when calling it backpressure, he used the wrong metaphor. Backpressure is about signaling to upstream processes that they are running too fast and need to slow down. The suggestions presented by Costa are mostly about signaling to the upstream process that it needs to do things differently, rather than just slow down. This has more to do with ensuring sufficient quality is sent downstream, rather than quantity. This irked me. As I was reading, I was searching for the right analogy. I kept coming back to lean manufacturing. The more famous half of the lean philosophy is waste reduction. The other half is about managing the unstable input of people. That's what we're interested in here. (Continue reading the full article on the web.)

18. Lean Launch Pad 2026 @ Stanford ? Lessons Learned Presentations

Source: Steve Blank | Published: Tue, 16 Jun 2026 13:00:57 +0000

Link:

<https://steveblank.com/2026/06/16/lean-launch-pad-2026-stanford-lessons-learned-presentations/>

We just finished the 16th annual Lean LaunchPad class at Stanford. In those 16 years, the class has gone from a radical idea ? that the Lean method could provide a more productive framework for new startups ? to something that everyone agrees is a way to build new startups. The class had gotten so [?]

19. Why Improve Your Writing?

Source: Refactoring English | Published: Tue, 03 Mar 2026 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/why-improve-your-writing/>

I?ve worked as a developer for 20 years, and I?ve always cared about clear writing. When I join a new team, the first thing I do is ask to update the onboarding documents. As I ramp up on the team, I document what I learn and encourage my teammates to do the same. Sometimes, other developers ask me why I invest so much into writing. Programming is technical, so why spend time on a ?soft skill? like writing? Isn?t that why we have technical writers and product managers?

20. How to Get Meaningful Feedback on Your Design Document

Source: Refactoring English | Published: Mon, 03 Nov 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/useful-feedback-on-design-docs/>

You?ve spent weeks carefully writing a design document for your software project, but what happens next? How can you get useful feedback about it from your teammates? How do you prevent your design review from dragging on for months? I?ve been through many design reviews in my career as both the author and reviewer, and I have a special fondness for effective reviews. Through trial and error, I?ve learned techniques that help the review process move smoothly and yield material improvements to the design.

21. Underused Techniques for Effective Emails

Source: Refactoring English | Published: Wed, 16 Jul 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/techniques-for-writing-emails/>

`blockquote { font-style: normal; } .pm-email blockquote { border-left: 4px solid purple !important; } .good-email blockquote { border-left: 4px solid #02c700 !important; h3 { margin-top: 1rem; font-size: 1.6em; font-weight: bold; } } .bad-email blockquote { border-left: 4px solid red !important; } p img { margin-top: 3rem !important; margin-bottom: 1rem !important; }` For software developers, there?s tremendous value in writing effective emails. Good emails save time, reduce misunderstandings, and earn you recognition within your company.

22. Ada Palmer ? Machiavelli is the most misunderstood thinker of all time

Source: Dwarkesh Podcast | Published: Tue, 16 Jun 2026 17:27:51 GMT

Link: <https://www.dwarkesh.com/p/ada-palmer-2>

"He begged to work for the regime that tortured him."

23. Do not invite big-tech to join your digital autonomy discussion

Source: Bert Hubert's writings | Published: Tue, 16 Jun 2026 18:00:00 +0200

Link:

<https://berthub.eu/articles/posts/do-not-invite-big-tech-to-your-digital-autonomy-discussion/>

Some time ago, I attended an event to discuss European digital autonomy and digital dependencies. I accepted the invitation without checking too carefully, and I belatedly found out that most of US big tech was invited as well, and would be presenting their thoughts. Now, even in our gloriously digitally autonomous future, we?ll still be doing business with Microsoft, Google and Amazon. And it is fine if they have a role.