

1. GLM-5.2 is probably the most powerful text-only open weights LLM

Source: Simon Willison's Weblog | Published: 2026-06-17T23:58:39+00:00

Link: <https://simonwillison.net/2026/Jun/17/glm-52/#atom-everything>

Chinese AI lab Z.ai released GLM-5.2 to their coding plan subscribers on June 13th, and then yesterday (June 16th) released the full open weights under an MIT license. Similar in size to their previous GLM-5 and GLM-5.1 releases, this is 753B parameter, 1.51TB monster - with 40 active parameters (Mixture of Experts). GLM-5.2 is a text input only model - Z.ai have a separate vision family most recently represented by GLM-5V-Turbo, but that one isn't open weights. GLM-5.2 has a 1 million token context window, up from GLM-5.1's 200,000. The buzz around this model is strong. Artificial Analysis, who run one of the most widely respected independent benchmarks: GLM-5.2 is the new leading open weights model on the Artificial Analysis Intelligence Index. GLM-5.2 is the leading open weights model on the Intelligence Index v4.1. At 51, it leads MiniMax-M3 (44), DeepSeek V4 Pro (max, 44) and Kimi K2.6 (43). They did however find it to be quite token-hungry: GLM-5.2 uses more output tokens per task than other leading open weights models: the model uses 43k output tokens per Intelligence Index task, up from GLM-5.1 (26k) and above MiniMax-M3 (24k), Kimi K2.6 (35k) and DeepSeek V4 Pro (max, 37...

2. Quoting Charity Majors

Source: Simon Willison's Weblog | Published: 2026-06-17T17:12:41+00:00

Link: <https://simonwillison.net/2026/Jun/17/charity-majors/#atom-everything>

What happened in 2025 was this: the economics of code production were turned upside down. Instead of being very hard, time-consuming, and expensive to generate code, it became effectively free and instant. Lines of code went from being treasured, reused, cared for and carefully curated, to being disposable and regenerable, practically overnight. ? Charity Majors, AI demands more engineering discipline. Not less Tags: charity-majors, ai-assisted-programming, generative-ai, ai, llms

3. ? a still that plays

Source: Simon Willison's Weblog | Published: 2026-06-17T03:56:10+00:00

Link: <https://simonwillison.net/2026/Jun/17/click-to-play-component/#atom-everything>

Tool: ? a still that plays A progressive enchantment Web Component that turns this markup: Into a still frame with a click to play button which loads the GIF on demand. For when you don't want big GIFs to be loaded unless people want to play them. Here's an example that demonstrates the new row editing tools in Datasette - in fact I built this Web Component for that post. Tags: gif, javascript, progressive-enhancement, web-components

4. Snap Unveils Specs, Its \$2,200 AR Glasses, and They're Fugly

Source: Daring Fireball | Published: 2026-06-18T01:33:15Z

Link: https://www.theverge.com/tech/950492/snap-specs-ar-glasses-launch-date-preorder?view_token=eyJhbGciOiJIUzI1NiJ9.eyJpZCI6IjZTMmZYVXprcHciLCJwIjoil3RIY2gvOTUwNDkyL3NuYXAtd3BIY3MtYXItZ2xhY2Z2ZGF0ZS1wcmVvcmlzImV4cCI6MTc4MjE3Nzc0Oj0wSwiaWF0IjoxNzgxNzQ1NzQ5fQ.Pdh1hCJafS7ca3UfJ7pPoS-wRpZQ6tEAr7HEVfTOAd8

Jay Peters, The Verge (gift link): Snap is finally launching augmented glasses for the public. Specs, which Snap describes as ?a wearable computer built into see-through augmented reality glasses,? will cost \$2,195. You can preorder a pair of Specs now at specs.com with a \$200 refundable deposit, and Snap says they're expected to ship ?this fall? in the US, UK, and France. [...] The company says that Specs are ?fully standalone, with no puck and no tether.? (Which is perhaps a jab at Apple's Vision Pro, which is tethered to a separate battery pack.) They'll be offered in two sizes, a 47mm model weighing 132g and a 52mm model weighing 136g, and will have removable inserts that Snap says will support ?a wide range of prescriptions.? Unlike Vision Pro, Snap is presenting Specs as eyeglasses that users will wear out and about in their

daily lives. Viewed perfectly straight-on and photographed by a professional fashion photographer ? as presented on the Specs website ? they're a bold look. Viewed from any other angle and captured normally, they look like goggles, not glasses. The frames look orthopedic and the lenses are not even close to clear. They make you look like you forgot to ta...

5. Vehicle Motion Cues ? a.k.a. Apple's Weird Anti-Nausea Dots

Source: Daring Fireball | Published: 2026-06-18T01:15:27Z

Link: <https://www.theverge.com/tech/942854/apple-vehicle-motion-cues-review-really-work>

Thomas Ricker, writing for The Verge: I'll just work from the car , I thought. But after a few minutes of staring at my screen on quick mountain switchbacks I could feel the first signs of cold, coagulated nausea bubbling up from that sweaty place in my gut. I looked to the horizon for relief, but nothing helped... until I remembered Apple's magic dots. Introduced in 2024, Apple's Vehicle Motion Cues promise to tap into your device's accelerometer and gyroscope to reduce or, in my case, even eliminate the motion sickness felt when trying to use an iPhone, iPad, or MacBook inside a moving vehicle. My son has suffered from motion sickness in cars his whole life, and Apple's Vehicle Motion Cues work like a charm for him too. What a great feature. ?

6. Yours Truly on MacBreak Weekly: Is the New Siri AI Good?

Source: Daring Fireball | Published: 2026-06-17T15:41:54Z

Link: <https://twit.tv/shows/macbreak-weekly/episodes/1029?autostart=false>

MacBreak Weekly: John Gruber of Daring Fireball joins the MacBreak Weekly panel this week! A deep dive into Apple's new Siri following WWDC. Why Apple Intelligence & the new Siri are not coming to the EU initially later this year. And could the iPhone Ultra's launch be delayed this year? It's fun to be the guest, not the host, of a podcast. I took Jason Snell's usual panelist spot this week, alongside Leo Laporte, Andy Ihnatko, and Christina Warren. Lots to cover, including a week of real-life experience using the new Siri AI. (It's really good!) Also, sometimes you just know what the episode title of a podcast is going to be, the moment a phrase is uttered. This was one of those episodes, with ? Intimate Functionalities ?. ?

7. Retrofitting the WM_COPY>DATA message onto Windows 3.1

Source: The Old New Thing | Published: Tue, 16 Jun 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260616-00/?p=112430>

It was carefully designed to be trivial. The post Retrofitting the WM_COPY>DATA message onto Windows 3.1 appeared first on The Old New Thing .

8. Formalizing a ring theorem with Lean 4 and Claude

Source: John D. Cook | Published: Wed, 17 Jun 2026 14:23:51 +0000

Link: <https://www.johndcook.com/blog/2026/06/17/rings-with-lean-claude/>

I've been testing Claude's ability to generate Lean 4 code to prove theorems. I've written about a couple experiments that verified calculations. I did not write about my failed attempt to get Claude to formalize a proof of the pqr theorem for seminorms. This time I asked Claude to formally prove the theorem from the [?] The post Formalizing a ring theorem with Lean 4 and Claude first appeared on John D. Cook .

9. You Got Faster. Your Company Didn't.

Source: Terrible Software | Published: Wed, 17 Jun 2026 17:39:03 +0000

Link: <https://terriblesoftware.org/2026/06/17/you-got-faster-your-company-didnt/>

AI made you faster. But you're not more productive; you're outsourcing the slow part to everyone else

10. Why servers running Ubuntu can stall on boot for two minutes

Source: Chris's Wiki :: blog | Published: 2026-06-17T02:30:09Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/UbuntuNetplanWaitForCarrier>

Suppose, not hypothetically, that you have some standard physical servers. As is typical, the

servers have more than one network interface (two is usually the basics even for 1U servers, and maybe you put in a 10G-T card in some), but you're only using one network interface on one network; the others are just sort of there. Recently you've started putting Ubuntu 26.04 LTS on them, and to your surprise (and displeasure) these machines stall for two minutes on boot, which they report as waiting for systemd-networkd-wait-online . The direct culprit is that you haven't told Canonical Netplan to ignore the state of network carrier for all of the network interfaces that you aren't using and don't have plugged in to anything. This is the ' ignore-carrier ' property, listed in Properties for all device types , and you need to set it on every network interface that is listed in /etc/netplan/*.yaml. However, why this happens and how we got here is a little bit convoluted (and it can happen in versions other than 26.04, under the right circumstances). Famously, Netplan doesn't actually configure your networks itself. Instead it creates a configuration for something else (a 'backend'), and on...

11. Flax debugging: making a hash of things

Source: Giles' blog | Published: Wed, 17 Jun 2026 02:40:00 +0000

Link: <https://www.gilesthoras.com/2026/06/hashing-jax-parameters>

I was debugging an issue with a JAX/Flax NNX training loop the other day, and found a neat little trick to help debug it. Specifically, I wanted to see if the issue was with my model, my loss function, my optimiser settings, or the "plumbing" of the training loop itself -- were gradients actually coming through and being applied to the parameters? I could print out the loss and the gradients, but printing out the parameters to see if they were changing was unhelpful -- any given update might only change a small number of parameters, or might change them such a small amount that I'd not notice -- especially given that the model had 77 million of them! Let's take a look. The world's worst LLM I am building an LLM from scratch in JAX and Flax NNX, and at this stage I'm trying to get the training loop right. As a simple test, I've just implemented the "shell" of the LLM -- the token embeddings on the input side, and the final linear layer for an output head, wired directly together. My plan was to train that so that given a sequence, instead of predicting next tokens for each position, it would "predict" the sequence itself -- that is, I might train it with the input The fat cat sat o...

12. Summoning the Demon

Source: the singularity is nearer | Published: 2026-06-17T00:00:00-07:00

Link: <https://geohot.github.io/blog/jekyll/update/2026/06/17/summoning-the-demon.html>

Watch you hit the stage like a willing bomb Strapped to crippled children It's hard to watch you whore out your damaged pride I spit on what you're building ? The Stephen Hawking The title is a 2014 Elon Musk quote . Now that we are 12 years in, it's so true, but maybe not in the way intended. Or maybe exactly as intended? I read this and it described what's currently going on as a social crime . I very much agree. AI is simply a reflection of our society, it doesn't contain a destiny. We are summoning the demon because we have demonic desires. I don't think this is going to change, but I do think that once the demon is here it will consume the people who brought it into the world. If you are one of these people, not getting out of technology will cost you your life. You'll bleed to feed the demon in me If you don't change your evil ways and end this peacefully You'll bleed to feed the demon in me I beg you, before you're digested, shred your legs and see I am a proud member of the perpetual underclass. After your game of musical chairs is over, you will all join me eventually, and when you show up you will lose because we have been here for years. We will be comfortable. We will...

13. Logic for Programmers v0.15, Livecoding

Source: Computer Things | Published: Wed, 17 Jun 2026 16:40:09 +0000

Link: <https://buttondown.com/hillelwayne/archive/logic-for-programmers-v015-livecoding/>

There's a new release of Logic for Programmers ! This one, version 0.15, is the first true release candidate. There's a couple of minor touch-ups I need to do but all content is in and copy edited and proofread. Unless something absolutely major comes up, the next release will be 1.0 and available in print. Full release notes here . (I'm also testing if a smaller low margin PDF is better for phones and computers. The printed book is going to have much larger proper

margins. Let me know what you think about this experiment!) Anyway I hurt my wrist and can't write a whole newsletter, so I'll just do a couple of recommendations. My friend Richard Whaling invited me to the inaugural livecode.chicago meetup , which was pretty cool! Among other things, people demoed two languages for livecoding music: Strudel is an extensive javascript DSL that seems relatively friendly to get into CLAVIER-36 is a ground-up special purpose language where the program is a 2D grid It all kinda makes me wish I knew the basics of music because it seems like it'd be a lot of fun to compare and contrast all of the various approaches to livecoding languages. Check em out if that's at all your thing! (For compl...

14. Pluralistic: The (real) dead economy theory (17 Jun 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 17 Jun 2026 18:27:18 +0000

Link: <https://pluralistic.net/2026/06/17/its-the-stupid-economy-stupid/>

Today's links The (real) dead economy theory: Vibes and memestocks, all the way down. Hey look at this: Delights to delectate. Object permanence: Jim Baen has had a stroke; Blame Apple for iTunes DRM; France v the internet; "Rotters"; 1901 undersea cables; Washington Post wants Trump coverage blackout; Taxes are for the little people; Gamer lifecycle; Ghanian postal song; "What Lies Beneath the Clock Tower": Murder of Jo Cox; 12 year old doxed by anti-vaxers; Hong Kong bookseller recants forced confession. Upcoming appearances: LA, Menlo Park, Toronto, NYC, Philadelphia, Chicago, London, Edinburgh, Brighton, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The (real) dead economy theory (permalink) Here's a fun fact about Elon Musk: in 2020, his (nominal) net worth was \$20b, and today it's \$1t (nominally). But that's not the fun fact; this is: everything he's done since 2020 was a flop. As John Quiggin writes, the pre-2020 Musk was the Musk of Tesla, batteries and Starlink. The post-2020 Musk is the Musk of Starship, robotaxis, Cybertrucks and Twitt...

15. How to Write Compelling Software Release Announcements

Source: Refactoring English | Published: Wed, 25 Jun 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/release-announcements/>

A release announcement showcases how the user's experience is better today than it was yesterday. That sounds obvious, but most release announcements seem to forget that there's a user at all. So many release announcements just enumerate new features in a way that's totally divorced from how real people use the software. The announcement is essentially just a changelog with better writing. For example, here's a ?changelog? style of announcing a new feature:

16. How to Write Blog Posts that Developers Read

Source: Refactoring English | Published: Thu, 27 Mar 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/write-blog-posts-developers-read/>

I recently spoke to a developer who tried blogging but gave up because nobody was reading his posts. I checked out his blog, and it was immediately obvious why he didn't have any readers. The developer had interesting insights, but he made so many mistakes in presenting his ideas that he was driving everyone away. The tragedy was that these errors were easy to fix. Once you learn to recognize them, they feel obvious, but some bloggers make these mistakes for years.

17. How to Write Useful Commit Messages

Source: Refactoring English | Published: Fri, 07 Mar 2025 00:00:00 +0000

Link: <https://refactoringenglish.com/excerpts/commit-messages/>

Effective commit messages simplify the code review process and aid long-term code maintenance. Unfortunately, commit messages don't get much respect, so the world is littered with useless messages like Fix bug or Update UI . There's no widespread agreement about what makes a good commit message or why it's useful, so I'm sharing my thoughts based on my 20 years of experience as a software developer. An example of a useful commit message What's the point of a commit message? Helps your code reviewer approach the change Communicates changes to teammates,

downstream clients, and end-users Facilitates future bug investigations Provides information to development tools Organizing information in a commit message Put the most important information first Use headings to structure long commit messages What should the commit message include? A descriptive title A summary of how the change impacts clients and end-users The motivation for the change Breaking changes External references Justifications for new dependencies Cross-references to issues or other changes Summaries of bugs or external references Testing instructions Testing limitations What you learned Alternative solutions you consi...

18. Key, in sight

Source: Aresluna | Published: 2026-06-16T18:50:27+0000

Link: <https://aresluna.org/key-in-sight>

A guide, of sorts, to keyboard customization. (9,400 words.)

19. Unsung

Source: Aresluna | Published: 2026-03-27T15:02:13+0000

Link: <https://unsung.aresluna.org>

A new blog on software craft and quality. Subscribe via RSS or receive an email digest every Friday.