

1. Saying the obvious thing

Source: seangoedecke.com RSS feed | Published: Sat, 27 Jun 2026 00:00:00 GMT

Link: <https://seangoedecke.com/saying-the-obvious-thing/>

Stating the obvious is surprisingly useful . Most of your knowledge lives below the threshold of conscious awareness, so it's possible for a piece of writing to remind you of what you already know. It's common to know you don't like something without being quite sure why, and reading an obvious statement (such as "accuracy matters , even when you agree with the broad strokes") can help clarify why you find certain things distasteful. Sometimes you can see some obvious truth that nobody seems to be talking about, and reading it in someone else's words can prompt an "oh god, I'm not crazy?" moment of catharsis. For many junior engineers, it's almost a rite of passage to notice that some percentage of software engineers do virtually no work . Since nobody talks about it (how would you even bring it up in the workplace?), they often feel like they're losing their minds: surely this state of affairs wouldn't be allowed to continue, so they must be completely misreading the situation. But in fact it's true. Stating the obvious is hard. It can even be dangerous: sometimes there's a good reason nobody says the obvious thing. But I think the bigger reason it's hard is for the same reason th...

2. All Chinese Models Will Be Illegal in 3... 2... 1...

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Sat, 27 Jun 2026 03:32:38 GMT

Link: <https://idiallo.com/blog/all-chinese-models-will-be-illegal>

The Washington Post reported that the US government will decide who can use state-of-the-art LLMs . After the ban of Fable and the limitations coming to ChatGPT 5.6, what's next? My bet is Chinese models. For all of Anthropic's doomsaying and propping up of their secret model Mythos, several open-weight models have proven capable of similar feats, and at a fraction of the cost. DeepSeek rocked the AI world in December 2024 with their initial release, nearly sending shockwaves through American stock markets. Last year, I looked into getting a BYD electric car. At the price they were selling for, I figured that even with a 100% tariff slapped on top, it would still be a bargain. Then I discovered that not only is there a steep import tariff, you simply cannot register the car in the United States. The car itself is illegal. According to reviews from people who actually own one, it's a fantastic vehicle that would outcompete most cars on the US market. Because of that, the US simply banned it. So what does this mean for large language models? If we're now told that state-of-the-art LLMs are too dangerous for the general public, what happens to Chinese models that are equally powerful...

3. Incident Report: CVE-2026-LGTM

Source: Simon Willison's Weblog | Published: 2026-06-26T17:58:54+00:00

Link: <https://simonwillison.net/2026/Jun/26/incident-report/#atom-everything>

Incident Report: CVE-2026-LGTM Spectacular hypothetical incident report by Andrew Nesbitt. Day 2, 16:00 UTC --- Two AI review agents from competing vendors, both attached to a downstream pull request bumping foxhole-lz4 , enter a disagreement loop over whether the package is malicious. After 340 comments and \$41,255 in inference spend, Finance revokes both API keys; one vendor's marketing team, cc'd on the cost anomaly alert, issues a press release citing "a 430% YoY increase in adversarial multi-agent security reasoning." The stock opens up 6%. Tags: security , ai , prompt-injection , generative-ai , llms , supply-chain , ai-security-research , andrew-nesbitt

4. Quoting OpenAI

Source: Simon Willison's Weblog | Published: 2026-06-26T17:10:43+00:00

Link: <https://simonwillison.net/2026/Jun/26/openai/#atom-everything>

We're beginning a limited preview of the GPT-5.6 series: Sol, our flagship model; Terra, a balanced model for everyday work; and Luna, a fast and affordable model. Terra has competitive

performance to GPT-5.5 while being 2x cheaper and Luna brings strong capability at our lowest cost. [...] We believe in broad access, and we plan to make GPT-5.6 Sol, Terra, and Luna generally available in the coming weeks. As part of our ongoing engagement with the U.S. government, we previewed our plans and the models' capabilities ahead of today's launch. At their request, we are starting with a limited preview for a small group of trusted partners whose participation has been shared with the government, before releasing more broadly. [...] GPT-5.6 is priced per 1M tokens across three model sizes: Sol is \$5 input / \$30 output; Terra is \$2.50 input / \$15 output; and Luna is \$1 input / \$6 output. GPT-5.6 also introduces more predictable prompt caching, including support for explicit cache breakpoints and a 30-minute minimum cache life. For GPT-5.6 and later models, cache writes are billed at 1.25x the model's uncached input rate, while cache reads continue to receive the 90% cached-input discount....

5. ? Bernie Sanders: Ideologue and Economic Ignoramus

Source: Daring Fireball | Published: 2026-06-27T23:35:00Z

Link: https://daringfireball.net/2026/06/bernie_sanders_ideologue

Sanders's tweet is better punctuated and capitalized, but it's the same argument as Trump's. Zero economic sense, 100 percent ideological wishful thinking.

6. Micron Executive Sumit Sadana Tells Tim Cook to Stop Hitting Himself

Source: Daring Fireball | Published: 2026-06-27T23:30:00Z

Link: https://www.wsj.com/tech/apple-raises-prices-on-macs-ipads-by-200-or-more-on-some-models-a7463f99?st=B1aQCP&reflink=desktopwebshare_permalink

From the bottom of Rolfe Winkler's report for The Wall Street Journal Thursday, on Apple's unprecedented price increases (gift link): Apple's price hikes arrived the day after Micron Technology, the big American maker of memory and storage, reported blowout quarterly earnings, touting gross profit margins that topped 80%. Shares jumped 16% after the close and appeared likely to power a Thursday rally among semiconductor stocks. [...] In an interview Wednesday night, Micron Chief Business Officer Sumit Sadana said the company couldn't make investments during the memory market's last downturn, when Micron's gross profits went negative, in part because certain customers took advantage to pay rock-bottom prices. "We told a couple of the customers who were being very aggressive with pricing at that time that this is not constructive," he said, without naming Apple, adding that low prices discouraged capital investments. "A lot of the industry investments got shut down in 2023 because of really poor pricing and really poor margins." I overlooked this segment when I read (and linked to) Winkler's report Thursday. It really does seem clear that Sadana is blaming Apple for not cutting Mi...

7. Apple Faced Bipartisan Opposition When It Last Lobbied to Buy Chinese RAM in 2022

Source: Daring Fireball | Published: 2026-06-27T22:39:51Z

Link: <https://www.warner.senate.gov/newsroom/press-releases/warner-rubio-urge-dni-to-review-risk-chinese-chipmaker-ymtc-presents-to-national-security/>

From a September 2022 letter to then-Director of National Intelligence Avril Haines, co-signed by Marco Rubio (then a Republican senator from Florida, currently secretary of state) and Mark Warner (Democratic senator from Virginia): We write to convey our extreme concern about the possibility that Apple Inc. will soon procure 3D NAND memory chips from the People's Republic of China (PRC) state-owned manufacturer Yangtze Memory Technologies Co. (YMTC). Such a decision would introduce significant privacy and security vulnerabilities to the global digital supply chain that Apple helps shape given YMTC's extensive, but often opaque, ties to the Chinese Communist Party (CCP) and concerning PRC-backed entities. In addition, we write to convey that any decision to partner with YMTC, no matter the intended market of the product offerings developed by such a partnership, would affirm and reward the PRC's distortive and unfair trade practices, which undermine U.S. companies globally by creating significant advantages to Chinese firms at the expense of foreign competitors. Last year, the Biden Administration described YMTC as China's "national champion memory chip producer," which supports t...

8. Brace expansion tree

Source: John D. Cook | Published: Sun, 28 Jun 2026 00:33:33 +0000

Link: <https://www.johndcook.com/blog/2026/06/27/brace-expansion-tree/>

Here's a crazy bash one-liner I found via an article by Peter Krumins: `echo`

`{w,t,}h{e{n{,ce{,forth}},re{,in,fore,with{,al}}},ither,at}` This prints 30 English words: when, whence, whenceforth, where, wherein, wherefore, wherewith, wherewithal, whither, what, then, thence, thenceforth, there, therein, therefore, therewith, therewithal, thither, that, hen, hence, henceforth, here, herein, herefore, herewith, herewithal, hither, hat This post will explain how [?] The post Brace expansion tree first appeared on John D. Cook .

9. When will the decimals in a/b repeat?

Source: John D. Cook | Published: Sat, 27 Jun 2026 18:12:03 +0000

Link: <https://www.johndcook.com/blog/2026/06/27/decimal-period/>

The previous post looked at how many digits are in the reduced fraction for the n th harmonic number. I was curious about how long the cycle of digits in a harmonic number might be. I wrote about the period length for the digits of fractions almost a decade ago. This post includes code so I can [?] The post When will the decimals in a/b repeat? first appeared on John D. Cook .

10. Height of harmonic numbers

Source: John D. Cook | Published: Sat, 27 Jun 2026 12:51:42 +0000

Link: <https://www.johndcook.com/blog/2026/06/27/height-of-harmonic-numbers/>

The previous post looked at writing the harmonic numbers as reduced fractions and estimating the number of digits in the numerator and denominator based on asymptotics. This is a follow up post with plots. We'll choose our base b to be 2. And we'll look at the total number of bits in both the numerator and [?] The post Height of harmonic numbers first appeared on John D. Cook .

11. How your Ubuntu 26.04 server boots with a software RAID array root filesystem

Source: Chris's Wiki :: blog | Published: 2026-06-27T03:29:54Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/linux/Ubuntu2604ServerSoftwareRaidBoot>

One answer to how your Ubuntu 26.04 server boots when its root filesystem is on a software RAID array is that it just does and you don't need to think about it. Unfortunately this wasn't the case in the pre-beta version . That's been fixed since (contrary to what I thought until now), but that raises the question of what changed between 24.04 LTS and 26.04 LTS to cause and then fix this. At some point between 24.04 and 26.04, Canonical switched from booting servers using `initramfs-tools` (and an initial ramdisk built with them) to booting them with `Dracut` . Initial ramdisks built with `initramfs-tools` will automatically assemble all software RAID arrays they see, or possibly anything they see in their embedded `/etc/mdadm/mdadm.conf` file (it's not clear to me right now). `Initramfses` using `Dracut` will normally only assemble software RAID arrays that are explicitly specified with `rd.md.uuid=` arguments , regardless of what's in their `/etc/mdadm.conf`. The 26.04 pre-beta server installer didn't add any such `rd.md.uuid` arguments to the kernel command line, causing the obvious problems. (At the time our workaround was to create an `/etc/default/grub.d` file that explicitly added the right `r...`

12. Arp 38: not that peculiar

Source: Maurycy's Blog | Published: Sat, 27 Jun 2026 00:00:00 +0000

Link: <https://maurycyz.com/astro/arp38/>

Data is from 2026-06-20. North is right. $0.53''/\text{pixel}$ ($27' \times 27'$ field). $\text{FWHM}=4'$ A nice, if rather small ($\sim 100''$) edge on spiral galaxy. Arp filed this one under "low surface brightness companions", but the companion doesn't seem to exist. There's a bright spot to the north (right) of the core , but hubble data shows it's just a bright HII region and star cluster.

This part of the sky isn't covered by SDSS, so I don't have widefield reference data. Total exposure time: $327 * 30 \text{ seconds} = 2.7 \text{ hours}$. Telescope: C9.25 (230mm, f/10, fl=2300mm) + 0.63 Starizona reducer Camera: IMX533 (16mm diagonal, square, color) Processing: Callibration (dark + flat) Stacking (average w/ outlier rejection) White balance and background subtraction (no

gradient removal) Asinh stretch (color preserving) Again with the walking noise! I'll try to check my calibration frames and restack this one. Related :

https://ned.ipac.caltech.edu/level5/Arp/Figures/big_arp38.jpeg : Arp's image.
[/astro/arp38/arp38.fits.gz](#) : raw stacks

13. Reading List 06/27/26

Source: Construction Physics | Published: Sat, 27 Jun 2026 12:02:02 GMT

Link: <https://www.construction-physics.com/p/reading-list-062726>

Trump refuses to sign a housing bill, the high cost of US-made doors, slow trucking, why we stopped making new land, and more.

14. This Week in Package Management: 27 June 2026

Source: Andrew Nesbitt | Published: 2026-06-27T10:00:00+00:00

Link: <https://nesbitt.io/2026/06/27/this-week-in-package-management.html>

Week six of the roundup, built from the package manager OPML feed collection and whatever I've posted or boosted on Mastodon . Releases Spack 1.2.0 makes the rewritten parallel installer the default, adds concretization groups and concretization caching, and ships SBOM generation alongside experimental build sandboxing and a spack isolate command. pnpm 11.9 computes a tarball's integrity from the downloaded file when a registry generates tarballs on demand and cannot supply a checksum in its metadata, storing it in the lockfile so later installs can verify it. The release also adds pnpm sbom --exclude-peers and speeds up audits on cyclic lockfiles. pixi 0.71.0 makes the conda-to-PyPI name mapping configurable, so you can supplement the generated mapping with your own entries rather than overwriting the whole file, which helps people on corporate networks hosting their own mappings. uv 0.11.24 makes project environments relocatable under preview and adds CPython 3.15.0b3. RubyGems 4.0.15 reduces peak memory when loading the full index and during bundle install , and Bundler 4.0.15 resolves Git LFS files in git sources from the real remote and implements a make jobserver. Dependabot...

15. Hazy Memory

Source: Tedium: The Dull Side of the Internet. | Published: 2026-06-27T14:09:29Z

Link: <https://feed.tedium.co/link/15204/17369108/apple-micron-ram-shortage-vertical-integration>

Who's to blame for the memory crisis that turned Macs and Steam Boxes into unobtainium this week? The memory-makers have a convenient answer.

16. The curious case of the disappearing Polish S

Source: Aresluna | Published: 2026-06-27T21:20:04+0000

Link: <https://aresluna.org/the-curious-case-of-the-disappearing-polish-s>

One keyboard bug three decades in the making. (New version of an essay originally posted in February 2015. 1,800 words.)

17. Pluralistic: Zuckerberg's increasingly bizarre war on whistleblowers (27 Jun 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 27 Jun 2026 11:14:05+0000

Link: <https://pluralistic.net/2026/06/27/zuckerstreisand-2/>

Today's links Zuckerberg's increasingly bizarre war on whistleblowers: Under no circumstances should you rush out and read the book that prompted Mark Zuckerberg to demand \$111m and eternal auctorial silence. Hey look at this: Delights to delectate. Object permanence: Flame warriors; Cryptography and casinos; TSA v dying 95 year old woman's adult diaper; Neoliberalism and Brexit; Beyond solutionism; How Thiel cheated with his Roth; Inequality's stabilizer; Palm Pilot school; Gillmor on PR flacks; "How I Edited an Agricultural Paper; Conservative judge chokes liberal judge; Hollywoodnomics; Rubber fingertips v fingerprint readers; Snowden's telepresence robot; "Shrill"; Moral hazard, "Three Rocks." Upcoming appearances: London, Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Zuckerberg's increasingly bizarre war on

whistleblowers ([permalink](#)) More than a decade ago, a group of young, internet-connected Belarusian dissidents launched a series of increasingly high-stakes, increasingly surreal confrontatio...