

1. The AI Compass

Source: Simon Willison's Weblog | Published: 2026-06-30T17:39:23+00:00

Link: <https://simonwillison.net/2026/Jun/30/the-ai-compass/#atom-everything>

The AI Compass This political compass style quiz by bambamramfan is pretty neat - answer 29 questions about AI and AI ethics to see which of the 30 archetypes you best fit. I'm impressed that my answers on my first time through the quiz categorized me as "The Garage Tinkerer", patron saint myself! It's implemented as a single page React app using the trick to avoid the necessary build step. Here's the code . Via @erisianrite.com Tags: ai , generative-ai , llms , ai-ethics

2. Have your agent record video demos of its work with shot-scraper video

Source: Simon Willison's Weblog | Published: 2026-06-30T16:54:26+00:00

Link: <https://simonwillison.net/2026/Jun/30/shot-scraper-video/#atom-everything>

shot-scraper video is a new command introduced in today's shot-scraper 1.10 release which accepts a storyboard.yml file defining a routine to run against a web application and uses Playwright to record a video of that routine. I've written before about the importance of having coding agents produce demos of their work; this is my latest attempt at enabling them to do that. Here's an example video created using shot-scraper video , exercising a still in development feature adding the ability to create new tables in Datasette from pasted CSV, TSV or JSON data: That video was created by running this command : shot-scraper video datasette-bulk-insert-storyboard.yml \ --auth datasette-demo-auth.json --mp4 (That --auth JSON file contains a cookie , as described here in the documentation.) Here's the datasette-bulk-insert-storyboard.yml file: output : /tmp/datasette-bulk-insert-demo.webm server : - uv - --directory - /Users/simon/Dropbox/dev/datasette - run - datasette - -p - 6419 - --root - --secret - " 1 " - /tmp/demo.db url : http://127.0.0.1:6419/demo/tasks viewport : width : 1280 height : 720 cursor : true wait_for : ' button[data-table-action="insert-row"] ' javascript : | () => {...

3. shot-scraper 1.10

Source: Simon Willison's Weblog | Published: 2026-06-30T15:10:14+00:00

Link: <https://simonwillison.net/2026/Jun/30/shot-scraper/#atom-everything>

Release: shot-scraper 1.10 The big new feature is shot-scraper video storyboard.yml , described in detail in Have your agent record video demos of its work with shot-scraper video . Tags: shot-scraper

4. ? A Tale of Two Modems

Source: Daring Fireball | Published: 2026-07-01T23:26:45Z

Link: https://daringfireball.net/2026/07/a_tale_of_two_modems

Cellular download speed and reception is nearly a solved problem for my needs. Battery life is not.

5. PlayStation Plus and Xbox Game Pass Subscriptions

Source: Daring Fireball | Published: 2026-07-01T19:55:26Z

Link: <https://daringfireball.net/linked/2026/07/01/valve-on-subsidizing-hardware>

Following up on my earlier post on Valve's righteous objection to selling game console hardware at a loss, I should have noted that PlayStation Plus starts at \$11/month (and goes up to \$20/month) and Xbox Game Pass starts at \$10/month (and goes up to \$23/month). One draw of these subscriptions is access to a library of game titles ? but another one is that you need one of these subscriptions to play online multiplayer games. Not every game demands online access but many (most?) do. There are very few serious PlayStation and Xbox gamers who don't pay for a subscription, and within a few years those subscriptions cost more than the (subsidized) hardware. It's not just about licensing fees for game titles players purchase anymore. Valve

didn't make any hay over this point, but should have. Because Steam Deck and Steam Machine are fundamentally more like PCs, all you need to play online multiplayer games is a free Steam account. ?

6. Valve Explains Why It Doesn't Subsidize Its Hardware Platforms

Source: Daring Fireball | Published: 2026-07-01T18:04:29Z

Link: <https://www.theverge.com/games/952004/valve-steam-machine-price-not-subsidizing>

Valve, in a statement to The Verge, explaining why it doesn't sell its handheld Steam Deck or new Steam Machine gaming devices at a loss (gift link): While this might seem like an easy solution, it doesn't align with our beliefs about how healthy ecosystems are built. If there's anything we're religious about at Valve, it's our belief that open systems are better in the long run, for ourselves and customers. The openness of the PC ecosystem in particular has enabled it to be the primary driver of hardware and software innovation, because anyone with an idea for a way to do something better was able to take a shot at it. When companies sell their hardware under cost for competitive advantage, or buy exclusive content for it, they're doing that to build a more closed system, one where you don't get to choose what software you want to use. We don't want that for PC hardware, and we don't think you should want it either. You shouldn't feel like you have to buy Valve hardware; you should be able to view it as just one option alongside all the devices for playing games, and select the one that makes sense for you. This means you get to decide which device fits your personal tradeoffs ar...

7. It rather involved being on the other side of this airtight hatchway: Changing administrative settings

Source: The Old New Thing | Published: Wed, 01 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260701-00/?p=112498>

Unlocking the door from the inside. The post It rather involved being on the other side of this airtight hatchway: Changing administrative settings appeared first on The Old New Thing .

8. 2026 mid-year link clearance

Source: The Old New Thing | Published: Tue, 30 Jun 2026 14:00:01 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260630-01/?p=112494>

Made it to another midpoint. The post 2026 mid-year link clearance appeared first on The Old New Thing .

9. Text AI watermarks will always be trivial to remove

Source: seangoedecke.com RSS feed | Published: Thu, 02 Jul 2026 00:00:00 GMT

Link: <https://seangoedecke.com/text-ai-watermarks/>

The European Union AI Act will begin to be enforceable in August 2026, one month from now 1 . One of the biggest new requirements is Article 50 , which requires all AI outputs to be ?detectable as artificially generated?. In other words, if LLM providers want to do business in the EU, they will have to apply a watermark to their outputs 2 : some hidden signature that can be used to identify AI content. LLM text watermarking is a fascinating problem. Like the best engineering problems, it is theoretically hard to solve perfectly, but has multiple partial solutions: for instance, Google's SynthID , and (as I'll argue) some quiet Unicode trickery from OpenAI and Anthropic. It will be interesting to see how the AI labs navigate these tradeoffs before the end of the year. Why text watermarking is hard I wrote about AI watermarking at the end of last year in AI detection tools cannot prove that text is AI-generated . It's easy to watermark an image, because digital images contain lots of noise that the human eye can't really see. For instance, you could apply a watermark like ?these twenty pixels in these exact spots will always share a color?. Text is much, much harder. Unlike images,...

10. I'm only interested in "native" installation systems

Source: Chris's Wiki :: blog | Published: 2026-07-02T01:34:27Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/NativeInstallersOnly>

There are a variety of ways to automatically or semi-automatically install systems, especially Linux systems and especially over the network. People have built a whole raft of them over the years, often with relatively impressive capabilities. A number of them have significant levels of automation and control, including for things we might want like sophisticated automatic disk setup . Despite that, we're interested in approximately none of them. As a practical matter, we only want to use the standard, native install systems for whatever we're running, which in this case is the Ubuntu server installer. There are two reasons for this. The obvious reason is that only native installers and install methods are officially supported by the Linux distribution (or whatever other free Unix we're using), and they're also the ones that are most used in practice. The native installers aren't always perfect, but other people work (often quite a lot) on making them work well, we can find plenty of people using them, and so on. If we use another installer, at the very least we're without support from the distribution. We're also likely to be off the beaten path, which means fewer people running...

11. Our long path from IPMI remote installs to network installs

Source: Chris's Wiki :: blog | Published: 2026-07-01T02:54:42Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/IPMIToNetworkInstalls>

Once upon what's now a long while ago, we had a bunch of SunFire X2100 and X2200 1U servers (for example, they were used in our first generation ZFS fileserver). One of the reasons that I loved these servers back in the days was that for free, they came with a full IPMI / BMC setup that supported both KVM over IP and virtual media. I installed any number of servers this way from the comfort of my office, using the KVM over IP as if I was at the console (because I effectively was) and the IPMI virtual media to feed a local ISO image to the server I was installing. In time those servers went the way of all servers (which is to say, into the e-waste dumpster, although it took a while) and none of our later servers gave me that '(re)install from the office' experience with their limited BMCs. For years, we lived with doing in person physical installs and reinstalls of our servers (well, I lived with that, my co-workers don't care as much). Recently we wound up with a bunch of servers in another building and a need to reinstall them all with Ubuntu 26.04. This pushed me into learning about UEFI network boot , working out how to network boot our customized Ubuntu ISO image , and then...

12. The CRA is not about open source

Source: Andrew Nesbitt | Published: 2026-07-01T10:00:00+00:00

Link: <https://nesbitt.io/2026/07/01/the-cra-is-not-about-open-source.html>

At FOSDEM in February and again at UN Open Source Week last week, the Cyber Resilience Act was the answer on offer whenever anyone asked what governments are doing about open source security, and the foundations and corporate advocates presenting it framed it as good news for open source. It is the largest piece of software legislation the EU has passed, the open source community spent two years lobbying over its text, and its main obligations come into force in December 2027. It is also not about open source. Where open source appears in the text, it appears as an exemption, a component risk, a compliance contact, or a paperwork trail. The CRA is a product-safety law in the CE-mark family, the same regime that already covers toys, radio equipment and machinery: rules a manufacturer must meet before selling a product in the EU, with a conformity mark at the end. The CRA extends that to anything with software in it. In the previous post I described this class of regulation as requiring haulage firms to certify which bridges their lorries crossed while employing no bridge inspectors, and the CRA is the regulation most often held up as the counterexample. Scope The things the CRA def...

13. Bring Back Crappy Forums

Source: Tedium: The Dull Side of the Internet. | Published: 2026-07-01T22:26:07Z

Link: <https://feed.tedium.co/link/15204/17371410/online-web-forums-retrospective>

Web forums were rough around the edges and faded in relevance as seemingly better options emerged. But what if we had stuck with them?

14. explaining myself through stories

Source: the website of jyn | Published: 2026-07-01T00:00:00+00:00

Link: <https://jyn.dev/explaining-myself-through-stories/>

When most people think ?river? I bet they think ?personal identity.?

15. The Winning Essays for the Big Questions About AI

Source: Dwarkesh Podcast | Published: Wed, 01 Jul 2026 22:13:17 GMT

Link: <https://www.dwarkesh.com/p/blog-prize-winners>

Abolishing pandemics/ Getting out of the way of AI automation/ Learning from Honk Kong MTR's business model

16. Summary of reading: April - June 2026

Source: Eli Bendersky's website | Published: 2026-06-30T18:52:00-07:00

Link: <https://eli.thegreenplace.net/2026/summary-of-reading-april-june-2026/>

"The Nuremberg Trial" by John Tusa and Ann Tusa - a detailed, meticulously researched account of the Nuremberg Trials. There's not a whole lot of side questing in this book - it's all focused on the trials themselves. Interesting read overall, though somewhat dry and academic. "Things Become Other Things: A Walking ?

17. Clickhouse is winning the Observability Wars

Source: matduggan.com | Published: Wed, 01 Jul 2026 13:14:03 GMT

Link: <https://matduggan.com/clickhouse-is-winning-the-observability-wars/>

For roughly the last ten years, a meaningful percentage of my working hours have been spent thinking about observability. If you're not familiar with the term, "observability" is what we call it now that "monitoring" doesn't sound expensive enough. The actual work

18. Pluralistic: Technocarcinization (01 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 01 Jul 2026 15:14:12 +0000

Link: <https://pluralistic.net/2026/07/01/ontogeny/>

Today's links Technocarcinization: Enshittification is the great leveler. Hey look at this: Delights to delectate. Object permanence: Grampa's backyard Disneyland; Elizabeth Warren on monopolies; Spotify v Apple (antitrust edn); Exxon lobbyist confesses; "When the Sparrow Falls." Upcoming appearances: London, Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Technocarcinization (permalink) "Carcinization" is a curious biological phenomenon: given enough time, across many environments, many species will evolve into crabs. The body-type of a crab, with its low center of gravity, sideways gait (useful for evading predators), ease of concealment and protected organs is suitable to many different environments: <https://en.wikipedia.org/wiki/Carcinisation> Lately, I've watched the American Big Tech platforms as they underwent their own form of technocarcinization, which is when every tech company turns into Facebook. For a long time, it seemed to me that you could make sense of the tech platforms by placing the...

19. Bulkdatasets AIVD en MIVD: de schaduw geheime dienst

Source: Bert Hubert's writings | Published: Wed, 01 Jul 2026 10:40:00 +0200

Link: <https://berthub.eu/articles/posts/de-schaduwgeheimediens/>

Uit een vandaag verschenen onderzoeksrapport (PDF) blijkt dat de AIVD en MIVD slordig en soms onrechtmatig omgaan met bulkdatasets. Bulkdatasets zijn grote bestanden vol data over vaak miljoenen willekeurige mensen, data die op een of andere manier in het bezit van de diensten gekomen is. Uit het rapport: ?De diensten kunnen bulkdatasets bijvoorbeeld verkrijgen van informanten, andere overheidsorganisaties, buitenlandse inlichtingendiensten, of tegen betaling van commerciële partijen. De diensten kunnen ook bijzondere bevoegdheden inzetten om bulkdatasets te verwerven.

