

1. April Report From Ookla: ?A Return to mmWave 5G?

Source: Daring Fireball | Published: 2026-07-02T22:30:40Z

Link: <https://www.ookla.com/articles/a-return-to-mmwave-5g>

Mike Dano, in a long (too long, I say) report for Ookla (makers of the nifty Speedtest app): Further, few other countries in the world followed in the mmWave footsteps of the U.S., with international spectrum regulators instead putting a focus on releasing mid-band spectrum for 5G. However, mmWave networks haven?t disappeared. New drive test data from Ookla?s RootMetrics, coupled with crowdsourced information from Ookla?s Speedtest Insights, shows the ongoing growth of mmWave 5G networks in the U.S., as well as the remarkable performance characteristics of those systems. Across all of RootMetrics? testing in the second half of 2025, in both urban (metro) and rural (state) areas, mmWave showed up in 2.2% of Verizon?s samples. For AT&T, that figure was 0.2%. For T-Mobile, that figure was almost 0% (and as a result, this report will mainly focus on Verizon and AT&T). Verizon?s mmWave connections showed up in 75 markets in the first half of 2024 (out of a total of 125 markets), a figure that rose to 91 in the second half of 2025. That?s almost triple the number of markets where RootMetrics recorded AT&T mmWave systems in the second half of 2025. 5G mmWave from T-Mobile, meanwhile, o...

2. Introducing the Safari MCP Server for Web Developers

Source: Daring Fireball | Published: 2026-07-02T21:55:31Z

Link: <https://webkit.org/blog/18136/introducing-the-safari-mcp-server-for-web-developers/>

Saron Yitbarek, writing on the WebKit blog, with a nice post-WWDC surprise: In Safari Technology Preview 247 , we?re introducing the Safari MCP server ? a Model Context Protocol server for web developers that makes your web development and debugging workflow faster and more powerful. We know agents are increasingly integral to the coding process and the Safari MCP server gives your agent the ability to know how your code actually renders in the browser by connecting it to a Safari browser window. Any MCP-compatible client can connect to the Safari MCP server. By connecting your agent to a Safari browser window, your agent can emulate what your users experience, giving it the information it needs to debug more autonomously, like access to the DOM, network requests, screenshots, and console output. MCP is Anthropic?s open protocol , so it was designed for Claude, but all sorts of other tools use it too ? Gemini CLI, Cursor, GitHub Copilot, and more. MCP really is open, not ?open?. ?

3. EveryMac Turns 30

Source: Daring Fireball | Published: 2026-07-02T21:31:04Z

Link: <https://everymac.com/whatsnew/>

EveryMac: Thirty years is a long time ? and a great deal has changed since then ? but what has not changed is that EveryMac.com has been there to provide you with detailed info on every Mac from the original 128k to the current line. Thank you very much for your support through the years. Daring Fireball turns 24 next month, which doesn?t sound that much younger than 30. But the way things work (in my mind at least) is that sites that are still around but were established years prior to my starting Daring Fireball are the real ?old guard?. I still feel like DF is a newcomer next to a site like EveryMac. 1996 for chrissakes. Steve Jobs wasn?t even back at Apple yet. What a great run it?s been and continues to be for EveryMac. Michael Tsai (whose eponymous blog is the same age as Daring Fireball) asked EveryMac proprietor Brock Kyle how it started . The final line of his answer: ?I miss the ethos of those days.? ?

4. llm-coding-agent 0.1a0

Source: Simon Willison?s Weblog | Published: 2026-07-02T19:33:12+00:00

Link: <https://simonwillison.net/2026/Jul/2/llm-coding-agent/#atom-everything>

Release: llm-coding-agent 0.1a0 Another Fable 5 experiment. Now that my LLM library has evolved into more of an agent framework it?s time to see what a simple coding agent would look like built on it. I started a new Python library using my python-lib-template-repository GitHub

template repository, then ran these two prompts (here's the Claude Code for web transcript):
Write a spec.md for this project - it will depend on the latest ?llm? alpha from PyPI and implement a Claude code style coding agent complete with tools for reading and editing files and executing commands Then: Commit the spec, then build it using red/green TDD in a series of sensible commits (each with passing tests and updated docs) - occasionally manually test it using the OpenAI API key in your environment Here's the spec , the resulting README file , and the sequence of commits . I've shipped a slop-alpha to PyPI, so you can run the new agent like this: `uvx --prerelease=allow --with llm-coding-agent llm code` It's pretty good for a first attempt! Here's the (Fable-authored) README , which lists recipes like `llm code --yolo` and `llm code --allow "pytest*" --allow "git diff*"` . It also presents a Python API based...

5. Using DSPy to evaluate and improve Datasette Agent's SQL system prompts

Source: Simon Willison's Weblog | Published: 2026-07-02T18:25:00+00:00

Link: <https://simonwillison.net/2026/Jul/2/dspy-datasette-agent-prompts/#atom-everything>

Research: Using DSPy to evaluate and improve Datasette Agent's SQL system prompts One of this morning's AIE keynotes covered dspy , which reminded me I've been meaning to see if it could help me improve the system prompt used by Datasette Agent - so I fired off an asynchronous research task in Claude Code for web using Claude Fable 5: Pip install the latest Datasette alpha and datasette-agent and dspy - then figure out how to use dspy to evaluate and improve the main system prompts used by Datasette Agent for the feature where it can execute read only SQL queries to answer user questions about data. Fable chose to test using GPT 4.1 mini and nano, and identified several promising looking directions for improvements. I particularly like this one: The schema listing gives only table names; the "don't call describe_table if you already have the information" advice caused column-name guessing (page_count, o.order_id, first_name) and error-retry loops in baseline traces. Either include column names in the prompt's schema listing or soften that advice. Tags: ai , datasette , generative-ai , llms , evals , dspy , datasette-agent , claude-mythos-fable

6. Understand to participate

Source: Simon Willison's Weblog | Published: 2026-07-02T17:07:14+00:00

Link: <https://simonwillison.net/2026/Jul/2/understand-to-participate/#atom-everything>

I saw Geoffrey Litt speak at AIE yesterday, and one framing he used particularly resonated with me: Understand to participate Geoffrey was talking about the challenge of collaborating with coding agents as they construct increasingly large and sophisticated changes, and the need to avoid taking on cognitive debt as your understanding drifts from how the code actually works. His argument is that you need to understand the code to a depth that enables you to participate further with the model: You can learn what the agent is doing to make sure you can be an active participant in the creative process. [...] You need a rich set of concepts in your mind to think creatively and fluently about how to move something forward. If you're lacking that fluency, your ability to participate in the project is meaningfully limited. The AIE talks are all recorded - all 300+ of them! - and should be trickling out over the next three weeks. Geoffrey's is one that I recommend catching on YouTube. Geoffrey also published a thread version of his talk on Twitter. Tags: geoffrey-litt , coding-agents , cognitive-debt , generative-ai , ai , llms

7. The case of the thread executing from an unloaded third-party DLL

Source: The Old New Thing | Published: Thu, 02 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260702-00/?p=112500>

Oops, I didn't realize that I was still doing that. The post The case of the thread executing from an unloaded third-party DLL appeared first on The Old New Thing .

8. FBI Seizes NetNut Proxy Platform, Popa Botnet

Source: Krebs on Security | Published: Thu, 02 Jul 2026 19:27:33 +0000

Link: <https://krebsonsecurity.com/2026/07/fbi-seizes-netnut-proxy-platform-popa-botnet/>

The Federal Bureau of Investigation (FBI) said today it worked with industry partners to seize

hundreds of domains associated with NetNut, a sprawling residential proxy service operated by the publicly-traded Israeli company Alarum Technologies [NASDAQ: ALAR]. The action comes roughly two weeks after KrebsOnSecurity published findings from multiple security firms connecting NetNut to the Popa botnet, a collection of at least two million devices that have been compromised by malicious software with little or no consent from victims.

9. Discovering rsync's -W option and our use for it

Source: Chris's Wiki :: blog | Published: 2026-07-03T02:42:01Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/RsyncWOptionUse>

Suppose, not hypothetically, that you use rsync to push an encrypted backup file from the machine it's created on to a fileserver, where it will be backed up by your regular backup system. Because this encrypted backup file is backed up every day, you only need one copy of it on the filesystem, so you use and reuse a fixed name for the file. In other words, we're using rsync somewhat as if it was scp, but with better control over what remote files can be written and (not) read. When you push (or pull) a file over rsync, rsync normally attempts to optimize what gets transferred by looking for common blocks in the file (how big a 'block' is depends on the file size, or you can fix it with the '--block-size' option, as covered in rsync(1)). This is a nice potential bandwidth saving, but it creates CPU and IO load on both ends as each of them checks their version of the file. In our specific case, we know that there aren't going to be common blocks; since the whole file is encrypted, it's basically random noise. Recently this backup process had some IO load problems, and today in the process of working on this I discovered rsync's '-W' option (also known as '--whole-file')....

10. This blog is written in en-GB

Source: Terence Eden's Blog | Published: Thu, 02 Jul 2026 11:34:59 +0000

Link: <https://shkspr.mobi/blog/2026/07/this-blog-is-written-in-en-gb/>

Someone left a comment on my blog recently asking if I'd mind making my language more inclusive. They didn't get some of the cultural references I'd used and suggested it would be easier if I used tropes which were more globally known. Here's the thing. No. All my blog posts start with a simple declaration: There's a reason for that. It is more than the?

11. This Page Left Intentionally Blank

Source: Jim Nielsen's Blog | Published: Thu, 02 Jul 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/intentionally-blank/>

I was popping off about negation being an act of creativity, when Blake Watson introduce me to the idea of the 'This Page Intentionally Left Blank'-Project (Internet Archive): In former times printed manuals had some blank pages, usually with the remark 'this page intentionally left blank'. In most cases there had been technical reasons for that. Today almost all blank pages disappeared [?] [this project] tries to introduce these blank pages to the Web again [?] to offer internet wanderers a place of quietness and simplicity on the overcrowded World Wide Web Ahead of its time. In our age of generative AI, a blank page is a deliberate act! So I went ahead made my own. Go ahead and crawl that bots. I don't use robots.txt, but I'm thinking of making one specifically to say 'Make sure you don't miss this page Botty Bot.' Reply via: Email

- Mastodon
- Bluesky

12. Understanding is the new bottleneck

Source: Geoffrey Litt | Published: 2026-07-02T10:00:00+00:00

Link: <https://geoffreylitt.com/2026/07/02/understanding-is-the-new-bottleneck.html>

July 2026 Understanding is the new bottleneck This is a written version of a talk I gave at the AI Engineer conference in July 2026, also shared as a tweet thread. Hot take: I think it's still important to understand the code that our agents...

13. Pluralistic: The difference between "today's task" and "accretive work" (02 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Thu, 02 Jul 2026 07:54:42 +0000

Link: <https://pluralistic.net/2026/07/02/canonization/>

Today's links The difference between "today's task" and "accretive work": Sometimes, "I got it working" is fine, but sometimes it isn't. Hey look at this: Delights to delectate. Object permanence: Series of tubes; Paralyzed teen beaten bloody by TSA; "Ultra unreal" Chinese lit; London property prices v Brexit; Biden v surprise billing; "Feeding Ghosts." Upcoming appearances: London, Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The difference between "today's task" and "accretive work" (permalink) One thing I've learned about paradoxes: often the answer to the riddle of "how can this one thing have such a contradictory set of features and effects?" is "it's not one thing, it's two things*." That's the idea that set me on the path to writing about "reverse centaurs" and AI. I was hearing from experienced programmers whom I knew to be reliable narrators of their own experience who described how AI was letting them write the best code of their lives; and from equally experienced and reliable coders who...

14. Digitale Autonomie 2.0: en nu echt

Source: Bert Hubert's writings | Published: Thu, 02 Jul 2026 13:37:00 +0200

Link: <https://berthub.eu/articles/posts/digitale-autonomie-2-0-surf-privacy-security/>

Afgelopen 25 juni deed ik het openingspraatje van de Surf Privacy en Security Conferentie. Nou heb ik vaker over digitale autonomie gesproken, maar deze keer heb ik het nadrukkelijk meer over wat er nu moet gebeuren. Of zoals de titel zegt ?en nu echt?. Ondanks dat ik nu letterlijk meer dan 50 praatjes over dit onderwerp heb gegeven had ik nog niet eerder de moeite genomen een goed transcript te maken.