

1. Better Models: Worse Tools

Source: Simon Willison's Weblog | Published: 2026-07-04T22:53:52+00:00

Link: <https://simonwillison.net/2026/Jul/4/better-models-worse-tools/#atom-everything>

Better Models: Worse Tools Armin reports on a weird problem he ran into while hacking on Pi: The short version is that newer Claude models sometimes call Pi's edit tool with extra, invented fields in the nested edits[] array. And not Haiku or some small model: Opus 4.8. The edit itself is usually correct but the arguments do not match the schema as the model invents made-up keys and Pi thus rejects the tool call and asks to try again. That alone is not too surprising as models emit malformed tool calls sometimes. Particularly small ones. What surprised me is that this is getting worse with newer Anthropic models as both Opus 4.8 and Sonnet 5 show it but none of the older models. In other words, the SOTA models of the family are worse at this specific tool schema than their older siblings. Armin theorizes that this is because more recent Anthropic models have been specifically trained (presumably via Reinforcement Learning) to better use the edit tools that are baked into Claude Code. This has the unfortunate effect that other coding harnesses, such as Pi, may find that their own custom edit tools are more likely to be used incorrectly. Claude's edit tool uses search and replace

....

2. June 2026 newsletter

Source: Simon Willison's Weblog | Published: 2026-07-03T14:50:50+00:00

Link: <https://simonwillison.net/2026/Jul/3/june-newsletter/#atom-everything>

The June edition of my sponsors-only monthly newsletter is out. If you are a sponsor (or if you start a sponsorship now) you can access it here . This month: Claude Fable 5, GPT-5.6, and US export restrictions GLM-5.2 is the new best open weights model Tokenmaxxing is so over Datasette Apps sqlite-utils and shot-scraper and Datasette Miscellaneous WASM projects Other model releases What I'm using Here's a copy of the May newsletter as a preview of what you'll get. Pay \$10/month to stay a month ahead of the free copy! Tags: newsletter

3. ?A Perfect Reflection of Trump?s Washington?

Source: Daring Fireball | Published: 2026-07-02T19:14:31Z

Link: <https://politicalwire.com/2026/06/19/a-perfect-reflection-of-trumps-washington/>

Taegan Goddard, two weeks ago at Political Wire: The Lincoln Memorial Reflecting Pool has become an almost too-perfect metaphor for Donald Trump?s presidency. He promised a quick, cheap fix. Instead, taxpayers got a no-bid project that ballooned to more than \$14 million, delivered a freshly painted pool in ?American Flag Blue,? and then promptly watched it turn green with algae as the new coating began to peel just days after it was supposedly finished. That is Trumpism in miniature: a grand declaration, a flashy cosmetic overhaul, a politically connected contractor, and an immediate failure blamed on someone else. Washington was supposed to be made beautiful again. Instead, the Reflecting Pool now looks like a murky swamp ? a fitting reflection of a capital overwhelmed by corruption and chaos. Like I just wrote , it?s all kayfabe . It was never about actually improving the Reflecting Pool. All that mattered was that Trump said he would. The only real part was the taxpayer money funneled to a Trump crony. Just like how in pro wrestling, they charge fans real money for tickets. That?s real too. Just like in pro wrestling, Trump has even explained it all away by pinning the blame on...

4. Claude Fable and Kayfabe

Source: Daring Fireball | Published: 2026-07-02T18:29:01Z

Link: <https://www.anthropic.com/news/redeploying-fable-5>

Anthropic: On Friday, June 12, the US government applied export controls to our newest models, Claude Fable 5 and Claude Mythos 5. This required us to restrict access to foreign nationals, whether inside or outside the United States. Because the order took effect immediately and we had no reliable way to verify nationality in real-time, we suspended access to both models for

all users. As of today, June 30, the export controls on Fable 5 and Mythos 5 have been lifted . That's a link to this tweet on Twitter/X from Commerce Secretary Howard Lutnick: Over the past two weeks, we have worked closely with Anthropic to analyze and approve Fable 5 to ensure alignment across the US Government and strengthen America's leadership in AI. I don't think a goddamn thing happened in the last two weeks, and if anything did happen, it sure as shit wasn't anything that Lutnick understood. There should be real government oversight regulating frontier AI, but this is just pantomime performative nonsense. The entire Trump 2.0 term (and much of the 1.0 term before it) can be summarized with a single word: kayfabe , the tacit agreement between professional wrestlers and their fans to pretend that overt...

5. Why Is Meta Destroying Its Engineering Organization??

Source: Daring Fireball | Published: 2026-07-02T17:09:42Z

Link: <https://newsletter.pragmaticengineer.com/p/why-is-meta-destroying-its-engineering>

Gergely Orosz, writing at The Pragmatic Engineer (which, sadly, is a Substack blog): The biggest problem: people stop caring about real work and focus on performative work. Let's check the four ingredients that Meta's leadership has decided to introduce to their workplace:

Tracking the keyboards and mouse clicks of all engineers, where legally possible
Reassign a good chunk of engineers to full-time data labeling
Let staff know that 10% of them will be laid off
Have a culture where devs optimize for any and all metrics measured during PSC
Measure token usage as part of PSC
Shake this mix up well, and what do you get? Two things: Everyone overuses AI to boost their personal stats. An engineering workforce that pretends to work with as much AI, and as little human input, as possible. It's a strange incentive where an outage caused by a failure to review code properly is not grounds for dismissal, but writing code by hand instead of having an AI agent write it could cost you your job. Every longer-tenured engineer is seeking a new job, or at least considering it. Those who have been around at Meta longer term have seen enough. PSC is 'Performance Summary Cycle', Meta's stringent...

6. C2PA only works if everything is signed

Source: seangoedecke.com RSS feed | Published: Mon, 06 Jul 2026 00:00:00 GMT

Link: <https://seangoedecke.com/c2pa-only-works-if-everything-is-signed/>

The European Union AI Act is Europe's attempt to comprehensively regulate AI usage. A big part of that is the requirement that AI-generated content be identifiable: either tagged with a watermark or with what the Act calls 'digitally signed metadata' 1 . Since all this becomes enforceable in a month, it's worth figuring out if it makes any sense. I recently discussed AI watermarking at length in Text AI watermarks will always be trivial to remove . What about digitally signed metadata? The most well-known implementation of digitally signed metadata is C2PA Content Credentials, which incorrectly 2 claims to be the technology that the AI Act gives as an example of how to do signed metadata properly. The idea here is that every single image file should contain unspoofable authorship metadata . Here's my position on it: C2PA broadly makes sense and is a good idea It is pointless to use C2PA for AI-generated images only It will take many years for C2PA to be adopted across all images Because C2PA makes such great safety theater, we're going to see a lot of hue and cry about it long before it becomes useful Lots to unpack. Let's start by considering images, since that's the easiest case...

7. Sometimes it actually is the network: a war story

Source: Chris's Wiki :: blog | Published: 2026-07-06T02:17:29Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/sysadmin/SometimesItsTheNetwork>

We've recently been having mysterious problems getting some of our backups to perform well . Also I recently wrote about how we'd wound up with a web server that frequently saturated its outgoing 1G interface with traffic (and it was a feature that it didn't have a faster network link). These two things turn out to not be as unrelated as we'd like, and there's a story or two there. The problematic backups aren't done through our usual Amanda-based backup system; instead, a particular central machine pulls /var/log and other relevant files from FreeBSD and OpenBSD hosts (either through rsync or direct SSH access) and then writes them over NFS to one of our NFS filesystems (and then they get backed up by the Amanda backups). The problems

manifested as terrible NFS performance; the central machine's load average would go to 50 or 70, everything NFS related on it would be slow, and so on. While we like NFS in general we've had quite a share of problems with it (cf , also), so we immediately assumed that this was yet another instance of an NFS problem. First we reduced the IO load from rsync s in various ways (eg), then we spent quite a while digging at various NFS metrics, both f...

8. Making a Shuffle Button

Source: Jim Nielsen's Blog | Published: Sun, 05 Jul 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/notes-shuffle/>

I made some updates to my notes blog , including a change to how my ?Shuffle? feature worked. Figured I'd blog about it. Shuffle? On a Blog? At the time of this writing, I have 974 ?notes? that I've published. For fun, I have a ?shuffle? button that digs up a random note from the past. I like to press it from time to time and re-encounter some insight from the past. It's like going through an old album, pulling out a random photo, and thinking, ?Oh yeah, I remember this! Good times.? Like old photos, there's also the occasional ?that didn't age so well?. But I find it fun to randomly dig up old insights from others and continue to be inspired. How Did It Work Before? Since my site is built and hosted as static files without a runtime server, this feature required JavaScript to work. Every page had a snippet like this: `button id = "js-shuffle" > Shuffle button > script > // All 974 note IDs injected by my SSG const notelds = ['id-1', 'id-2', id-3', 'id-4', '...']; document.querySelector("#js-shuffle").addEventListener('click', () => { // randomly grab an item in 'notelds' const randomId = '...'; window.location.href = '/n/${randomId}/' }) script > Essentially: inject every note ID...`

9. Travel notes: PLDI Boulder

Source: Max Bernstein's Blog | Published: Sun, 05 Jul 2026 00:00:00 +0000

Link: https://bernsteinbear.com/blog/travel-notes-pldi-boulder/?utm_source=rss

I had another excellent PLDI this past June. It was my fourth 1 . I continued to meet new people and learn new things! Overall: I got to meet a lot of new people, which was exciting. I had some good chats about research. I asked a question at a talk! I got to show Aaron and Jacob PLDI and see them enjoy it. I missed hanging out with CF Bolz-Tereick and Chris Fallin, the usual suspects at conferences I attend. I'm looking forward to next year. This post is more about the conference than the town of Boulder (unlike the last PLDI post about Seoul) because I didn't do much Boulder exploring. Sunday I got in late on Sunday. Then I had to take a long car ride from Denver airport to Boulder. I don't think I had ever flown into Denver with intent to go to Boulder before so it was a bit of a surprise. Jacob offered to have a late dinner with me so we had a tasty meal at Gaia Masala and Burger. Shout out to Harry, our server. Monday Monday was a workshop day. I signed up for EGRAPHS and mostly stayed in that workshop. People kept throwing around the term ?Knuth-Bendix?, as they have been for several years, and only in one of these workshop talks did someone explain it in a way that made any...