

1. The Special Value Pi 4 was extremely short-lived

Source: Jeff Geerling | Published: Wed, 08 Jul 2026 09:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/special-value-pi-4-extremely-short-lived/>

The 'Special Value' Pi 4 pictured above is probably the rarest Raspberry Pi I own?even rarer than my blue special edition Pi . A Raspberry Pi reseller briefly listed a special 'value edition' Pi 4 . But the product page 404's now. While it was up, my curiosity got the better of me, and now I have two 'value' Pi 4s. What makes them a 'value'? They're only certified to run at 1.25 GHz (retail Pi 4s run at 1.8 GHz, and can usually be overclocked).

2. Rewriting Bun in Rust

Source: Simon Willison's Weblog | Published: 2026-07-08T23:57:21+00:00

Link: <https://simonwillison.net/2026/Jul/8/rewriting-bun-in-rust/#atom-everything>

Rewriting Bun in Rust Jarred Sumner has been promising this blog post (since May 9th) about his Zig to Rust rewrite of Bun for significantly longer than it took him to finish the rewrite. Honestly, it was worth the wait. This is a detailed description of an extremely sophisticated piece of agentic engineering, featuring dynamic workflows, trial runs, adversarial review and all sorts of other interesting tricks. Jarred spends the first half of the post praising Zig for getting Bun this far. Then we get to a core idea in the piece, emphasis mine: Our bugfix list felt bad and I was tired of going to sleep worrying about crashes in Bun. I don't blame Zig for that - other users of Zig don't have the bugs we had, and mixing GC with manually-managed memory is an uncommon enough thing for software to need that no language really designs for it. We wouldn't have gotten this far if not for Zig, and I'll always be grateful. Until very recently, programming language choice was a one-way decision for a project like Bun. Everyone knows you should never stop the world and rewrite a large piece of software from the ground up. Joel Spolsky highlighted that in Things You Should Never Do, Part I b...

3. Introducing GPT?Live

Source: Simon Willison's Weblog | Published: 2026-07-08T23:20:48+00:00

Link: <https://simonwillison.net/2026/Jul/8/introducing-gptlive/#atom-everything>

Introducing GPT?Live OpenAI finally upgraded the model used by ChatGPT voice mode! I've had preview access for a few weeks in the iPhone app, and the new model is very impressive. It also has the ability to spin off harder tasks to GPT-5.5: For questions that require web search, deeper reasoning, or more complex work, it delegates to our latest frontier model behind the scenes and brings the result back into the conversation when it's ready. While it works, GPT?Live can keep talking with you and maintain the flow of conversation. At launch, GPT?Live will use GPT?5.5 in the background. As we release new frontier models, we'll continuously update the model used by GPT?Live. The previous voice mode in the ChatGPT app was based on a GPT-4o era model, with a knowledge cut-off some time in 2024. I had mostly stopped using voice mode because the age and relative weakness of the model greatly limited how useful it was as a brainstorming partner. During the preview period I encountered a pretty obscure bug: the model was interrupting me to laugh at things I said, which weren't even intended as jokes! It felt rude and condescending - I reported it to OpenAI and as far as I can tell they mad...

4. Quoting Kenton Varda

Source: Simon Willison's Weblog | Published: 2026-07-08T20:03:34+00:00

Link: <https://simonwillison.net/2026/Jul/8/kenton-vara/#atom-everything>

I just declared a moratorium against AI-written change descriptions (e.g. PR and commit messages, also issues/tickets) from my team. AI was writing change descriptions that were worse than useless to me as I tried to review PRs: outlining details of the code that could easily be seen by looking at the code, but omitting the higher-level framing needed to understand broadly what the code is doing. ? Kenton Varda Tags: kenton-vara , ai-assisted-programming , generative-ai , ai , llms

5. Felons, Fraudsters Flog Offensive Cybersecurity Startup

Source: Krebs on Security | Published: Wed, 08 Jul 2026 12:31:39 +0000

Link:

<https://krebsonsecurity.com/2026/07/felons-fraudsters-flog-offensive-cybersecurity-startup/>

A cybersecurity startup dangling millions of dollars to acquire zero-day security vulnerabilities in popular software is run by a pair of far-right conspiracy theorists and convicted felons whose most recent ventures included fake intelligence companies and a now-defunct AI-based lobbying platform they operated under assumed names.

6. And Then the Billionaire Paid Off \$550 Million of Our Debts

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Wed, 08 Jul 2026 08:00:00 GMT

Link: <https://idiallo.com/blog/billionaire-paid-off-550-million-dollars-of-our-debts>

Imagine being worth \$2 billion. Would you give away \$550 million? That's a quarter of your wealth, more money than I could spend in several lifetimes. Yet that's how much Evan Spiegel, the CEO of Snapchat, and his wife have donated to a charity in California. Specifically, they donated to Undue Medical Debt, an organization that buys Californians' medical debts and expunges them. A noble act. I'm not one to tell you that billionaires shouldn't exist, or what is or isn't fair in a system I don't control. But one thing I've come to see over the years is that public good deeds are rarely what they seem. Whether it's a feel-good story on the news, a TV show pimping your car, or another turning your house into a mansion, there's an underlying truth that often gets obscured by the appearance of a good deed. Bill Gates, who was once the richest man in the world, pledged to leave almost all his money to charity. Over the years this story has been repeated, and I'm the last person to tell you how to feel about it. For me, it's a good thing to help a charitable organization that's trying to help others. Last year, Bill Gates renewed this pledge, stating that he would give away around \$200 b...

7. ? What's Good for the iOS Goose Is Often Not Good for the MacOS Gander

Source: Daring Fireball | Published: 2026-07-09T01:45:08Z

Link: https://daringfireball.net/2026/07/whats_good_for_the_ios_goose_is_often_not_good_for_the_macos_gander

App icons in MacOS are not mere buttons. You can drag them, move them, and drop things on them. You click them to select, and double-click to launch. They are richer objects that deserve a richer visual vocabulary.

8. ?PARRY Encounters the DOCTOR? ? Chatbot on Chatbot Action Circa 1973

Source: Daring Fireball | Published: 2026-07-08T22:38:06Z

Link: <https://www.rfc-editor.org/info/rfc439/>

Back in the primordial days of AI, PARRY was an ELIZA-style chatbot created by psychiatrist Kenneth Colby to simulate the words of a paranoid schizophrenic. Someone had the genius idea to connect it to ELIZA (a.k.a. ?DOCTOR?). Vint Cerf published the transcript as an RFC in 1973, and I laughed my ass off reading it. A taste: Do you know anything about bookies? What makes you think I know anything about bookies? A bookie didn't pay me off once. What does that suggest to you? I beat up the bookie who did not pay off, and I thought he might use his friends in the underworld to get even with me. What else would you like to discuss? What do you know about racketeers? Thanks to Bob Peterson for the link. Vint Cerf, by the way, retired from Google last week at the spry age of 83. ?

9. Mac Apps Can Escape From Squircle Jail If They're Not in the Mac App Store

Source: Daring Fireball | Published: 2026-07-08T21:03:21Z

Link: <https://tyler.io/2026/07/05/escape-from-squircle-jail/>

Tyler Hall: We all know about macOS Tahoe's terrible app icons and how 3rd party developers have been confined to squiracle jail . If you're lucky enough to distribute an app outside the Mac App Store, you can break free of squiracle jail using NSDockTilePlugIn . It's not strictly

the intended use-case of that API. And it's not allowed in the Mac App Store, either. But it can solve the problem. So today's release of Iris adds three additional app icons to choose from in the app's Special Preferences Settings pane. And since they use the NSDockTilePlugIn API, the custom icon remains even when you quit the app. Iris is far from the only app using this or other techniques to sort-of escape squircle jail. All the apps I'm aware of doing something like this are Mac-assed apps. It's a good litmus test. ?

10. The other kind of control flow guard check: The combined validate and call

Source: The Old New Thing | Published: Wed, 08 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260708-00/?p=112510>

A two-in-one package. The post The other kind of control flow guard check: The combined validate and call appeared first on The Old New Thing .

11. BMCs and a surprising USB network device on your server

Source: Chris's Wiki :: blog | Published: 2026-07-09T01:50:12Z

Link: <https://utcc.utoronto.ca/~cks/space/blog/tech/BMCsAndUSBEthernetDevices>

Suppose, not entirely hypothetically, that you're installing a server with two network ports and during the (Linux) installation, a third network device shows up with a funny name like 'enp1s0f4u1u2c2', which is a USB Ethernet device (despite you not having any such thing plugged in to the server's USB ports). To your further surprise, your server installer can even lease a DHCP IP on this interface, say "169.254.3.1". Congratulations, your server has a BMC, and this BMC probably speaks Redfish, which is sort of the modern, cloud influenced version of IPMI. One of the things you'd like to do on a server with a BMC is have the server (the 'host') talk directly to the BMC, for example to get sensor information that only the BMC has or to configure the BMC. In the world of IPMI, people had to put together special methods to talk to the BMC, which required kernel drivers, extracting information from SMBIOS, and so on. This isn't the greatest, and is also not at all like how you talk to the management agent in cloud virtual machines, where you generally talk to the management agent by making HTTP requests to a special IP address. IPMI is in part a network protocol, but talking to...

12. Arp 90

Source: Maurycy's Blog | Published: Wed, 08 Jul 2026 00:00:00 +0000

Link: <https://maurycyz.com/astro/arp90/>

North is up (mirrored, exact). 0.53"/px [15' x 9.2' field]. FWHM = 3.9" Arp 90 is the two interacting galaxies in the east. The west of the image has a wierd distorted galaxy (?), with a similar redshift to the interacting pair. This means it could be a tidal fragment, or a gravitationally bound irregular galaxy... I really don't know. Total exposure time: 314 * 30 seconds = 2.6 hours. Used in stack: 276 * 30 seconds = 2.3 hours. Telescope: C9.25 (230mm, f/10, fl=2300mm) + 0.63 Starizona reducer Camera: IMX533 (16mm diagonal, square, color) Processing: Callibration (dark + flat) Stacking (average w/ outlier rejection) White balance and background subtraction (no gradient removal) Asinh stretch (color preserving) Rotation + crop Tone curve Taken during one night, with typical (ML-free) processing... except I rotated the image before cropping to align it to north. This should make it easier to cross reference. Related : https://ned.ipac.caltech.edu/level5/Arp/Figures/big_arp90.jpeg : Arp's image. /astro/arp90/stack.fits.fz : Raw stacks.

13. The console wars have been lost

Source: Xe Iaso's blog | Published: Thu, 09 Jul 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/console-wars-lost/>

Valve wins by doing absolutely nothing

14. Agents are monads (but not that kind)

Source: Xe Iaso's blog | Published: Wed, 08 Jul 2026 00:00:00 GMT

Link: <https://xeiaso.net/blog/2026/hyle-pneuma/>

I managed to write this without using the word endofunctor once. Wait, shit, I just did. Uhhh,

oops!

15. Family Feud: Mac-assed Mac App Edition

Source: Jim Nielsen's Blog | Published: Wed, 08 Jul 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/mac-assed-family-feud/>

?We asked 100 people: What are the top three companies on earth best positioned to make a world-class Mac-assed Mac app ?? Buzz! ?Apple!? Survey says: Yes! Apple at the number one spot. Makes sense. Who better to make the very definition of a great Mac app than the people who make the Mac? No brainer, I suppose. Granted, they've had some misses , but nobody bats 1000. Ok, let's keep going. ?We asked 100 people: What are the top three companies on earth best positioned to make a world-class Mac-assed Mac app?? Buzz! ?Anthropic!? Survey says: Wow, that's odd huh? You'd think Anthropic would be right there at number two. Not only do they have billions of dollars, but they also develop, maintain, and control the super intelligence we'll all soon be subservient too, right? Surely if anyone (besides Apple) is well positioned to make a world-class Mac app, it would have to be Anthropic ? right? And yet, here we are with Claude Desktop as an Electron app . Ok, let's keep going. ?We asked 100 people: What are the top three companies on earth best positioned to make a world-class Mac-assed Mac app?? Buzz! ?Adobe!? Maybe not . Buzz! ?Google!? Not so much . I'm sorry, but that's three strikes...

16. Writing an LLM from scratch, part 34b -- from bigrams to GPT-2, one component at a time (in JAX)

Source: Giles' blog | Published: Wed, 08 Jul 2026 18:45:00 +0000

Link: <https://www.gilesthomas.com/2026/07/llm-from-scratch-34b-building-and-training-gpt-2-small-in-jax>

This post is the capstone of the most long-running series on my blog . In December 2024 (!), I started reading Sebastian Raschka 's book " Build a Large Language Model (from Scratch) ", and worked through it carefully. Being who I am, despite trying to apply a strict "no side quests" policy, I found myself zooming off and digging into all kinds of things. It's time to wrap it up. I had decided that the endpoint would be to build and train an LLM from scratch just using my notes -- no reference to the book, no reference to the model code I'd written when following the book. After an X/Twitter poll, I decided to use JAX for that, just to make sure that I really was building it from scratch and not regurgitating bits of PyTorch code like a bad coding LLM spitting out half-digested lumps of Stack Overflow. In my last post , I showed how I built a JAX training script that mirrored what I had built for the original PyTorch version of the model. To test it as I went along, I used it to train a really dumb "LLM", which instead of trying to predict the next token for every token in an input sequence, instead predicted the input -- that is, if you fed it The fat cat sat on the mat It would...

17. poppy the training box, part 1: the beginnings

Source: Giles' blog | Published: Thu, 09 Jul 2026 01:00:00 +0000

Link: <https://www.gilesthomas.com/2026/07/poppy-the-training-box-1-the-beginnings>

For a while I've been planning to put together a separate machine for local LLM training. Until now, I've been using my desktop PC, perry . I have an RTX 3090 installed, and can get useful training runs done (most recently, a 163M-parameter GPT-2 small style LLM in JAX), but there are a couple of problems. perry is my daily driver. If he's doing a training run, then everything is just a little bit sluggish as CPU and GPU alike are busy. Although I don't play games often, it's annoying to have the option ruled out for days at a time. While the GPU is busy with a training run, I can't do other experiments in parallel -- for example, to scope out what the next step might be. And relatedly to all of those: the two-day limit to the training runs I've been doing is something I set because that's the maximum amount of time I'm willing to have perry tied up. It would be really interesting to try longer training runs! I also have longer-term plans; a multi-GPU box would be interesting to put together -- not just to have more power locally, but so that I could test larger-scale cloud multi-GPU training runs before starting to pay for expensive machines. US\$15.92 an hour to rent a machine i...

18. Weekly Update 511: Live from my Riad in Marrakech

Source: Troy Hunt | Published: Wed, 08 Jul 2026 13:54:46 GMT

Link: <https://www.troyhunt.com/weekly-update-511/>

How's this for a location?! I mean, last week was nice with Scott in Mallorca, but Marrakech is, well, wow ? Anyway, about those data breaches... This week I'm talking about the futility of attempting to remove piss from a pool , yet here we are, with

19. The Decline of Deviance 2

Source: Experimental History | Published: Wed, 08 Jul 2026 16:10:54 GMT

Link: <https://www.experimental-history.com/p/the-decline-of-deviance-2>

2 Decline, 2 Deviance