

1. QuadRF can spot drones and see WiFi through my wall

Source: Jeff Geerling | Published: Fri, 10 Jul 2026 09:00:00 -0500

Link:

<https://www.jeffgeerling.com/blog/2026/quadrf-can-spot-drones-and-see-wifi-through-my-wall/>

The QuadRF (pictured above) a phased-array radio built around a Raspberry Pi 5 and an FPGA board with picosecond-level timing. It does advanced signal processing and beamforming. It can see WiFi through walls and track drones in flight. If the open source community can come up with something like this, just imagine what governments are capable of. When you plug a computer into a network, tools like Wireshark can show all the hidden traffic you might not even know is there. WiFi packets are the same, but those travel through the air, allowing snooping without physical access.

2. John Ternus Calls Sam Altman

Source: Daring Fireball | Published: 2026-07-11T03:45:00Z

Link: <https://www.youtube.com/watch?v=CIASuxd8jQY>

?Yeah, who?s this?? ?You know who this is.? ?Yes I do, yes I do. I sent a guy to deliver the package ... he didn?t call. Is everything alright?? ?Tell you what. Forget the money.? ?

3. ?No Interest?

Source: Daring Fireball | Published: 2026-07-11T03:41:00Z

Link: <https://x.com/drewpusateri/status/2075708238650089981>

Drew Pusateri, director of communications at OpenAI, on Twitter/X (or XCancel): Our statement in response to this suit: We have no interest in other companies? trade secrets. We remain focused on building innovative technology that empowers people everywhere. Let?s say I think you stole my wallet. I approach a police officer and tell him my suspicion and describe the evidence that makes me think you have my wallet in your pocket. ?Is that his wallet in your pocket?? the cop asks. ?I have no interest in other people?s wallets? would be a rather curious answer. ?

4. Ice Cold

Source: Daring Fireball | Published: 2026-07-11T03:14:28Z

Link: <https://www.threads.com/@alexheath/post/Daol2jaEioX>

Alex Heath, on Threads: At WWDC, Apple execs I met with were ice-cold when I asked about their OpenAI partnership. Now we know why: Apple just sued OpenAI for allegedly stealing trade secrets related to consumer hardware (Apple and OAI senior leaders are in Sun Valley this week. Yikes!) I noticed the same ice-cold reaction to my questions about ChatGPT and Siri. (In fact, I think Heath and I even talked about it together when we bumped into each other at the end of the day on Monday during WWDC week.) At the time I took it as Apple execs not wanting to distract from the fact that Siri AI with Apple Intelligence, with no third-party plug-in, was truly competitive. But in hindsight their coolness, I think, was about this. There was no response like, say, ? We think ChatGPT is great and we?re happy to keep it available in iOS 27 as an option for our customers who love that experience, but the new Siri AI truly stands on its own. ? That?s a typical Apple answer. But what they actually said was just more like ? The ChatGPT extension remains available. ? I didn?t think much of it at the time but now it stands out. ?

5. Quoting Nilay Patel

Source: Simon Willison?s Weblog | Published: 2026-07-10T17:05:26+00:00

Link: <https://simonwillison.net/2026/Jul/10/nilay-patel/#atom-everything>

The reality is to make augmented reality glasses, you need to put a camera next to your eyes that is continuously recording everything you see and processing that to put information over it. There is not another way around it. And there?s certainly not a chip that can fit in the

stem of a glasses that is both powerful enough and power miserly enough to do that in real time. You have to send that data to a cloud. You gotta do it. [...] Or you can build something the size of a Vision Pro with a battery pack that lives somewhere else. Those are the current choices in this world. And it means if you want to build the product that everyone thinks is the next thing, you are going to have to invade people's privacy. And maybe you shouldn't. Like, there's an incredible argument for, nope, you shouldn't do that. Nope, the trade-offs required to make this product are so high at a societal level that we should stop it. ? Nilay Patel , The Vergecast Tags: ai-ethics , augmented-reality , nilay-patel , privacy , ai

6. llm-meta-ai 0.1

Source: Simon Willison's Weblog | Published: 2026-07-09T16:12:20+00:00

Link: <https://simonwillison.net/2026/Jul/9/llm-meta-ai/#atom-everything>

Release: llm-meta-ai 0.1 Let's LLM run prompts against the new muse-spark-1.1 model. Tags: llm , meta

7. llm 0.31.1

Source: Simon Willison's Weblog | Published: 2026-07-09T16:06:15+00:00

Link: <https://simonwillison.net/2026/Jul/9/llm/#atom-everything>

Release: llm 0.31.1 Fix for a bug with OpenAI Chat Completion endpoints where a tool call with empty arguments could result in a JSON error from some providers. #1521 This bug came up when I was testing llm-meta-ai . Tags: llm

8. The case of the mysterious changes to integers when there shouldn't have been any code generation effect

Source: The Old New Thing | Published: Fri, 10 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260710-00/?p=112514>

Decoding where those integer came from. The post The case of the mysterious changes to integers when there shouldn't have been any code generation effect appeared first on The Old New Thing .

9. Game Review: Lovers In A Dangerous Spacetime ?????

Source: Terence Eden's Blog | Published: Fri, 10 Jul 2026 11:34:18 +0000

Link: <https://shkspr.mobi/blog/2026/07/game-review-lovers-in-a-dangerous-spacetime/>

My new year's resolution is to play more video games with my wife. Specifically co-operative games. I hate playing competitively; it's rubbish to achieve victory at the expense of someone else. So I asked for recommendations and picked the cheapest things which looked reasonable. Several people recommended Lovers In A Dangerous Spacetime. It's a neat little game which is just short enough to?

10. Package Management as Org Chart

Source: Andrew Nesbitt | Published: 2026-07-10T10:00:00+00:00

Link: <https://nesbitt.io/2026/07/10/package-management-as-org-chart.html>

Conway's Law says organisations produce systems that copy their own communication structure. Dependency management tooling is part of the system. A resolution strategy is an opinion about how disagreements get settled; a manifest format records who is allowed to depend on whom. Monorepo, single version policy : every package in the tree must agree on one version of each dependency; upgrading anything means upgrading everyone at once. The org chart has been forbidden from appearing in the code. It works when there's a standing migration workforce whose job is dragging every consumer along. Monorepo with workspaces : one repository, but each package keeps its own manifest and can pin its own versions. One strong platform team, and leadership that believes coordination was solved at the 2019 offsite. Git submodules : a dependency is referenced by exact commit SHA in a separate repository, and updating it is a manual two-step that never gets automated. Two teams agreed to collaborate and recorded, to the commit hash, exactly how little they meant it. Bazel : every dependency edge between targets is declared explicitly in a BUILD file; nothing is ambient or inferred. Adopted when the o...

11. Building intuition about LLM parameter counts

Source: Giles' blog | Published: Fri, 10 Jul 2026 21:30:00 +0000

Link: <https://www.gilesthomas.com/2026/07/llm-parameter-counts>

When I was building my GPT-2 implementation in JAX, I started with just token embeddings for the input, and a separate output head (as I was not using weight tying). It wasn't an LLM -- no Transformer blocks, no attention, no feed-forward networks. I was somewhat surprised when I noticed that even that stripped-down model had 77 million parameters with the "small" settings I was using to train -- specifically, an embedding dimension of 768. However, I realised I shouldn't be -- with a vocab size of 50,257, each of those components is essentially a $768 \times 50,257$ matrix, and that is indeed over 38 million numbers. But the finished LLM at the end of the project was only 163 million parameters -- that meant that the input and output components alone were almost half of it. That felt like a surprisingly large percentage. I had a similar shock when I was first looking into the feed-forward network, and realised that it had roughly twice as many parameters as the attention layers. When we learn about the internals of LLMs, a lot of the focus is on the attention mechanism. This makes sense -- it's the hardest part to get your head around. The rest of the setup, at least for simple GPT...

12. Premium: The Hater's Guide To The Memory Crisis

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 10 Jul 2026 17:51:30 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-the-memory-crisis/>

Hi premium readers! I'll be taking a week off of the premium next week? July 17? to have some well-earned rest. This will mark only the second time I've missed a premium piece since I started this newsletter in June 2025, and I

13. Adam Brown? A deep but accessible introduction to general relativity

Source: Dwarkesh Podcast | Published: Fri, 10 Jul 2026 16:23:56 GMT

Link: <https://www.dwarkesh.com/p/adam-brown-gr>

And why black holes are the ultimate power plants.

14. Dot product: Component vs. Geometric definition

Source: Eli Bendersky's website | Published: 2026-07-10T18:42:00-07:00

Link: <https://eli.thegreenplace.net/2026/dot-product-component-vs-geometric-definition/>

The goal of this post is to answer a simple question: why are the following two definitions of the vector dot product in Euclidean space [1] equivalent for vectors $\vec{a}$ and $\vec{b}$:

Component definition: $\vec{a} \cdot \vec{b} = \sum_{i=1}^n a_i b_i$ Geometric definition: $\vec{a} \cdot \vec{b} = |\vec{a}| |\vec{b}| \cos \theta$

15. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 10 Jul 2026 16:20:37 +0000

Link: <https://www.filfre.net/2026/07/this-week-on-the-analog-antiquarian/>

Opus 4: Henry VI, Part 3

16. Pluralistic: "Rights for robots" and the AI slavery fantasy (10 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 10 Jul 2026 09:52:06 +0000

Link: <https://pluralistic.net/2026/07/10/posthuman-as-in-no-humans/>

Today's links "Rights for robots" and the AI slavery fantasy: When we were robots in Egypt? Hey look at this: Delights to delectate. Object permanence: Awkward questions for the entertainment industry; ISP v record label death penalty; Coming out on Splash Mountain; Negative Swiss bond yields; Are We Having Fun Yet? Upcoming appearances: London, Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em.

Colophon: All the rest. "Rights for robots" and the AI slavery fantasy (permalink) While the AI bubble is primarily a material phenomenon (driven by the calculation that bosses are easy marks for a sales pitch that sees them replacing workers with software), there is an inescapable ideological component to it: the desire for a world without people in it:

<https://pluralistic.net/2026/05/13/vibe-governance/#k-hole> If you'd like an essay-formatted version of this thread to read or share, here's a link to it on pluralistic.net, my surveillance-free, ad-free, tracker-free blog:
<https://pluralistic.net/2026/07/10/posthuman-as-in-no-humans/#he...>