

1. Using uvx in GitHub Actions in a cache-friendly way

Source: Simon Willison's Weblog | Published: 2026-07-14T00:56:20+00:00

Link: <https://simonwillison.net/2026/Jul/14/uvx-github-actions-cache/#atom-everything>

TIL: Using uvx in GitHub Actions in a cache-friendly way I finally found a cache-friendly recipe for using uvx tool-name in GitHub Actions workflows that I like. The trick is setting a UV_EXCLUDE_NEWER: "2026-07-12" environment variable at the start of the workflow and then using that as part of the GitHub Actions cache key. This means any uvx tool-name commands will resolve to the most recent version as-of that date, and you can bust the cache and upgrade the tools by bumping the date in the future. My goal here is to use Python tools in GitHub Actions without every run of the workflow hitting PyPI to download a fresh copy of the tool and its dependencies. Tags: packaging , pypi , python , github-actions , uv

2. DOOMQL

Source: Simon Willison's Weblog | Published: 2026-07-13T22:34:41+00:00

Link: <https://simonwillison.net/2026/Jul/13/doomql/#atom-everything>

DOOMQL Peter Gostev built this using GPT-5.6 Sol. This is a lot of fun: DOOMQL started with a deliberately unreasonable question: what if SQLite were the game engine, not merely the place where a game stores data? The result is a small, original Doom-like game in which SQL owns movement, collision, enemies, combat, progression and every RGB pixel on screen. It's implemented as a Python terminal script - I tried it out like this: `cd /tmp git clone https://github.com/petergpt/doomql cd doomql uv run host/doomql.py` Here's the huge SQL query that implements a full ray tracer in SQLite using a recursive CTE. Running the above script creates a `/tmp/doomql/.doomql/doomql.sqlite` SQLite database, which you can explore using Datasette like this: `uvx --prerelease=allow --with datasette-apps datasette \ /tmp/doomql/.doomql/doomql.sqlite \ -p 4444 --root --secret 1 --internal internal.db` The `--with datasette-apps` option installs the new Datasette Apps plugin, which supports creating custom HTML+JavaScript apps that can run SQL queries directly within the Datasette interface. I created a new app, pasted the copy-paste prompt into Claude chat (Fable 5) and told it : Build an app that displays th...

3. datasette code-frequency chart on GitHub

Source: Simon Willison's Weblog | Published: 2026-07-13T21:45:27+00:00

Link: <https://simonwillison.net/2026/Jul/13/datasette-code-frequency/#atom-everything>

datasette code-frequency chart on GitHub Out of curiosity I decided to see if I could find a useful illustration of the impact of coding agents and Opus 4.5 class models on my own output. The best I've found so far is this GitHub chart of frequency of code changes to my Datasette open source project: The big spike in activity at the end aligns with Opus 4.8, GPT-5.5, Fable 5 and GPT-5.6 Sol. Tags: github , ai , datasette , generative-ai , llms , ai-assisted-programming , coding-agents

4. Remember Musk's Suit Alleging a Conspiracy Between Apple and OpenAI?

Source: Daring Fireball | Published: 2026-07-13T15:24:53Z

Link: <https://arstechnica.com/tech-policy/2025/08/elon-musk-sues-apple-openai-to-block-exclusive-iphone-chatgpt-integration/>

Ashley Belanger, reporting for Ars Technica back in August 2025: After a public outburst over Grok's App Store rankings , on Monday, Elon Musk followed through on his threat to sue Apple and OpenAI. At first, Musk appeared fixated on ChatGPT consistently topping Apple's ?Must Have? app list ? which Grok has never made ? claiming Apple seemed to preference OpenAI, an Apple partner, over all chatbot rivals. But Musk's filing shows that the X and xAI owner isn't just trying to push for more Grok downloads on iPhones ? he's concerned that Apple and OpenAI have teamed up to completely dash his ?everything app? dreams , which was the reason he bought Twitter . For what it's worth, I just looked at the App Store editorially-curated ?Popular

iPhone Apps? list , and ChatGPT is in the top spot. Given the lawsuit Apple filed against OpenAI last week, I'd say this is pretty good anecdotal data that these editorial decisions in the App Store aren't driven by favoritism. The top 10 downloads list for free iPhone apps currently looks like this (for me here in the U.S. App Store): Paramount+ Netflix Game Controller ChatGPT Kalshi: Trade the World Cup Peacock TV: Stream TV & Movies Netshort - Popular Dr...

5. ?Every Frame Perfect?

Source: Daring Fireball | Published: 2026-07-12T19:48:09Z

Link: <https://tonsky.me/blog/every-frame-perfect/>

Nikita ?Tonsky? Prokopov: The rule of thumb is: If I take a screenshot of your app at any moment, you should be able to explain what I see. Why care about every frame? It builds trust. Users can't see the code, so UI is the only way for them to judge the quality of the app. If UI looks good, that means developers had time to polish it, which means that they probably spent a comparable amount of time to iron out the code. It's a heuristic, but a reasonable one. Now, what does it mean in practice? I can think of a few things: No white flashes between screens. No partially loaded content. No relayout while content loads. Internally consistent. If one part of the UI says ?1 update available?, another part should not say ?Checking for updates...? Precise animations. Animations often end up being forgotten. A UI might look great in both start and end states but very janky in between. ?Every frame perfect? is a great mantra for UI craftsmanship. If you care about every frame, that discipline will be palpable, even though almost no users will ever examine your animations and transitions frame-by-frame, and most will happen too quickly to see in real time. If you cut corners on interstia...

6. TwoMillionKit: Use Private Cloud Compute in MacOS 27 Foundation Models Without an Entitlement

Source: Daring Fireball | Published: 2026-07-12T18:04:05Z

Link: <https://github.com/insidegui/TwoMillionKit>

Guilherme Rambo: Apple ships the fm command-line tool in macOS 27, which can be used to run inference with the local system model or Private Cloud Compute from Terminal or scripts. You know what else can run command-line tools? Mac apps! ? I decided to spend some of my Codex tokens and take GPT 5.6 Sol for a spin. I asked it to create this Swift package. All it does is provide a LanguageModel implementation that uses the fm command-line tool under the hood, meaning that any Mac app can use the Private Cloud Compute model without requiring a special entitlement from Apple. The main limitation is that this will not work for sandboxed Mac apps, so any Mac app distributed via the Mac App Store won't be able to use it. But for developers of Mac apps distributed outside the Mac App Store, this provides a simple and entitlement-free way to use Private Cloud Compute in their apps. Use sparingly and at your own risk. This is a workaround for Apple's current limitation that only grants access to Private Cloud Compute to ?developers in the App Store Small Business Program with fewer than two million first time App Store downloads?. Hence Rambo's clever name for the framework. ?

7. Lessons Learned from CISA's Recent GitHub Leak

Source: Krebs on Security | Published: Mon, 13 Jul 2026 15:03:28 +0000

Link: <https://krebsonsecurity.com/2026/07/lessons-learned-from-cisas-recent-github-leak/>

The Cybersecurity and Infrastructure Security Agency (CISA) has issued a postmortem on a data leak in which a contractor published dozens of internal CISA credentials -- including AWS Govcloud keys -- in a public GitHub repository for almost six months before being notified by KrebsOnSecurity. Experts say the gaps identified in the agency's initial response provide important lessons that all security teams should absorb.

8. Why don't we just make the entire stack out of guard pages?

Source: The Old New Thing | Published: Mon, 13 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260713-00/?p=112528>

Guard pages all the way down? The post Why don't we just make the entire stack out of guard pages? appeared first on The Old New Thing .

9. [RSS Club] Half a million steps is about 10 marathons

Source: Terence Eden's Blog | Published: Mon, 13 Jul 2026 11:34:09 +0000

Link: <https://shkspr.mobi/blog/2026/07/rss-club-half-a-million-steps-is-about-10-marathons/>

Shhh! RSS Club posts are only available to feed subscribers. Keep the secret! I'm not a big fan of the "Quantified Self" movement, but since buying a ridiculously cheap smartwatch, I've been interested to see just how many steps I take per day. During our recent Interrail holiday, we travelled about 6,400 Km by train over 49 days. But how much did we cover by foot? My watch pairs with the Open ?

10. Mandatory Update: A Short Story

Source: micahflee | Published: Mon, 13 Jul 2026 15:33:31 GMT

Link: <https://micahflee.com/mandatory-update-a-short-story/>

Guess what? Lately I've been nerding out on creative writing. I even joined a local fiction writers group where we meet up at coffee shops. Like, in person! Anyway, I submitted a short story to the DEF CON 34 Creative Writing Short Story Contest . My story is called

11. Pluralistic: Why aren't AI companies competing directly with their customers? (13 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 13 Jul 2026 08:15:24 +0000

Link: <https://pluralistic.net/2026/07/13/go-meta-meta/>

Today's links Why aren't AI companies competing directly with their customers? Why sell picks and shovels if you've already struck gold? Hey look at this: Delights to delectate. Object permanence: Pro-logging Lorax; Anti-DRM picture book; Whose cops did America train? Day on a Device; Theresa May x Fresh Prince; Ratfucking Corbyn supporters; Infosec v W3C DRM; Remix Public Enemy; Hungary's Cold War cartoons; America is unpaving its roads; Olympic overruns are world champions; Vivendi lobbyist is now UN copyright boss; Boris's racism; Facebook's insider stalkers; Semantic drift, ethical drift. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em.

Colophon: All the rest. Why aren't AI companies competing directly with their customers?

(permalink) "I often wonder what the Vintners buy/One half so precious as the Goods they sell"

-The RubÆiyÆt of Omar KhayyÆm I first encountered that quote from someone extolling the virtues of bookstores, and it stuck with me, because for most of my childhood, every bookstore vis...

12. Winners and losers in the coming AI margin collapse (part 2)

Source: Martin Alderson | Published: Sun, 12 Jul 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/the-upcoming-ai-margin-collapse-part-2-winners-and-loser-s/?utm_source=rss&utm_medium=rss&utm_campaign=feed

Part two: as AI inference margins collapse, who actually captures the value? The hardware supply chain and consumers win, model inference commoditises, and the frontier labs' escape routes are managed agents and staying ahead.

13. Control the ideas, not the code

Source: antirez.com | Published: Mon, 13 Jul 2026 13:39:52 +0200

Link: <http://antirez.com/news/169>

Look at the past history of this blog. There are many blog posts about programming with AI, a few of them date back to January 2024 (like this: <https://antirez.com/news/140>). I'm a relatively well regarded programmer, after all. I don't have the need to still be in the ?loop? as a old man that seeks for relevance, I recently rejoined Redis, and now I also am developing a new open source software for local LLM inference that received a good welcome in the community. Why I keep doing this, of saying what people don't want to hear? Why I keep announcing how future programming will be by default? Because I feel the urge of lowering the impact for people less prepared to the change than me, often younger than me, and that,

unlikely me, didn't see many of those things coming (In 2022 I published, before ChatGPT existed, a book preannouncing many things that now happened and other things that I believe *will* happen, so I feel like I can say this without sounding egocentric). So mine is a trick. People feel more and more programming is completely modified by AI and don't know what they should do, if they can really start coding in a completely different way, without looking much at the c...