

1. Reverse-engineering is cheap now

Source: Simon Willison's Weblog | Published: 2026-07-20T19:24:05+00:00

Link: <https://simonwillison.net/2026/Jul/20/cheap-reverse-engineering/#atom-everything>

I keep hearing anecdotes from people who used coding agents to reverse-engineer and automate devices in their homes. I think this is an interesting illustration of the impact of the reduced cost of writing code. Prior to agents, it was entirely possible to reverse-engineer home devices. The problem was the ROI - was it really worth all of that effort? More importantly, any experienced programmer knows that undocumented, unstable APIs like that may well change or break in the future. Is that initial work worth the effort if you're committing yourself to a frustrating cycle of maintenance in the future? Coding agents change that equation entirely. The effort to get a simple automation working has dropped, as has the cost of trying and failing to get it to work. Since the code is so cheap, the idea of having to maintain it in the future - or throw it away and start again - carries way less psychological baggage. Tags: reverse-engineering , coding-agents , ai-assisted-programming , generative-ai , ai , llms

2. Who's Afraid of Chinese Models?

Source: Simon Willison's Weblog | Published: 2026-07-20T17:09:19+00:00

Link: <https://simonwillison.net/2026/Jul/20/afraid-of-chinese-models/#atom-everything>

Who's Afraid of Chinese Models? Interesting proposal from Ben Thompson that both addresses the hypocrisy of labs outlawing distillation against their models despite training on unlicensed data, and could help US open models compete more effectively with their Chinese counterparts: The U.S. should pass a law that (1) makes explicit that collecting data for training models is fair use, and (2) bars terms of service that forbid distillation, for U.S. companies at a minimum. Stopping distillation ? which is literally just querying the API ? is nearly impossible; the U.S. should go the other way and lean into a new copyright policy that both indemnifies the labs and also guarantees that what they learned fuels further innovation for everyone else. Ben also theorizes that Alibaba's decision to release Qwen 3.8 Max as open weights - a reversal from their decision not to release Qwen 3.7 Max in May - may have been influenced by a recent speech by Xi Jinping, who said: We should seize this rare, historic opportunity to encourage open source, openness, collaboration and sharing. And on the subject of Qwen 3.8 Max - a new 2.4T parameter model (nearly as large as the 2.8T Kimi K3) - here's a...

3. Inkling: Our open-weights model

Source: Simon Willison's Weblog | Published: 2026-07-16T15:35:25+00:00

Link: <https://simonwillison.net/2026/Jul/16/inkling/#atom-everything>

Inkling: Our open-weights model Mira Murati's Thinking Machines Lab just released their first open-weights model. Inkling is "a Mixture-of-Experts transformer with 975B total parameters, 41B active" - an Apache-2.0 licensed multimodal model trained on 45 trillion tokens of text, images, audio and video. They're also promising Inkling-Small, a 276B (12B active) model, but that's still being tested and the weights will be released "once that work is complete". The model card is much shorter than I've come to expect from US AI labs. It links to even shorter Training Data Documentation with almost nothing of interest in it - it's best summarized by these two paragraphs: The datasets Thinking Machines Lab uses to develop its AI services includes content that is in the public domain as well as content that may be subject to intellectual property protection. Thinking Machines Lab's services were developed using publicly available content obtained from the open internet and publicly accessible data repositories. Certain datasets were also obtained from third parties. By Thinking Machines' own admission, this is not a frontier model. It's instead intended as a strong base model for fine-tu...

4. [Sponsor] WorkOS MCP: Manage Your Auth Platform From Any AI Agent

Source: Daring Fireball | Published: 2026-07-20T23:13:48Z

Link: https://workos.com/blog/management-mcp-server?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q32026

Debugging SSO, managing users, adjusting auth policies, configuring branding: every configuration task has lived behind a UI that only a human can drive. The WorkOS MCP server gives agents the same access as your dashboard login. Hundreds of operations, discoverable at runtime. Connect in one command via OAuth, with scoped tokens instead of a master API key. Pass a screenshot of your marketing site and ask your agent to match the login page. If a human had to do it before, an agent can do it now. Connect your agent ? ?

5. ?Who?s Afraid of Chinese Models??

Source: Daring Fireball | Published: 2026-07-20T16:27:36Z

Link: <https://stratechery.com/2026/whos-afraid-of-chinese-models/>

Ben Thompson, on the hype regarding Kimi K3, at Stratechery: This is a point that bears repeating: because U.S. open weight model makers must follow the frontier labs' terms of service, they (1) are worse than Chinese alternatives and (2) end up distilling the distillation, just with a detour through Chinese labs. Wouldn't it be better if western open weight model makers could go to the source? To that end, here's an even more interesting question around distillation: why exactly is it bad? After all, what are large language models but the distillation of all of the knowledge on the open Internet, scraped by the frontier labs and distilled into the models that are themselves being distilled? Who is exactly being wronged here? In fact, this paradox is the solution. I believe that open weight models are good for innovation (and, per the above, I think that labs on the frontier will be fine), but it's a problem to be dependent on China. The U.S. should pass a law that (1) makes explicit that collecting data for training models is fair use, and (2) bars terms of service that forbid distillation, for U.S. companies at a minimum. Stopping distillation ? which is literally just querying...

6. MG Siegler: ?OpenAI Makes ChatGPT ChatGPT Again?

Source: Daring Fireball | Published: 2026-07-17T19:48:53Z

Link: <https://spyglass.org/chatgpt-brings-back-chatgpt/>

MG Siegler, writing at Spyglass: In that light, we can see the dilemma. But there were also probably about a dozen better ways to do this roll-out ? as highlighted by how fast they're fixing these things ? and they just missed the mark. While they undoubtedly knew there would be some backlash, they probably didn't realize there would be this much . Again. All that points to perhaps a disturbing trend where OpenAI doesn't really understand their user base. Which you almost can understand given how they clearly stumbled into ChatGPT in the first place. Still, here they are with those billion users. A problem that all of their competitors would love to have. But still a problem if you want to fundamentally change what you are as a product . But those are bigger, existential issues. All I wanted was chat back in the ChatGPT app. Front and center. And I got it. It's not perfect. It's still bloated. But at least it's usable again now and not a confusing mess of ideas out of the box. MG is pretty scathing, but I still think he's taking it too easy on OpenAI for what a colossal fuck-up this remains. There was nothing wrong with the way things were, with two apps ? the simple ChatGPT and t...

7. Making an agile version of a Windows Runtime delegate in C++/WinRT, part 1

Source: The Old New Thing | Published: Mon, 20 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260720-00/?p=112545>

The easy case is easy. The post Making an agile version of a Windows Runtime delegate in C++/WinRT, part 1 appeared first on The Old New Thing .

8. Expensive Is Just a Brand Now

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Mon, 20 Jul 2026 23:00:00 GMT

Link: <https://idiallo.com/blog/expensive-is-just-branding>

You've probably heard this one before. It might have come from a parent, a teacher, or a friend who bought an expensive item and wanted to justify it. It goes something like this: A poor man buys a cheap pair of boots for a few dollars. They fall apart within a year, so he buys another pair. And another. Over five years, he spends far more than the man who bought one expensive, well-made pair up front. A pair that's still going strong to this day. And the lesson was simple: buying cheap is expensive. Buy quality, and you save money in the long run. Now swap the boots for a jacket, a phone, a kitchen knife, whatever you like. In my experience, that used to be true. I have old clothing still hanging in my closet, and they still hold together. Price and quality were correlated. If a company was going to charge you more, it needed to justify that premium with better materials, better construction, better everything. The market punished brands that didn't. But something changed along the way. This is just my own personal observation, but I'm sure you can relate: I pay too much for things that don't last. When I was growing up, we separated the real from the knockoff with a simple test....

9. Volume to Area ratio for Regular Solids

Source: John D. Cook | Published: Mon, 20 Jul 2026 14:35:21 +0000

Link: <https://www.johndcook.com/blog/2026/07/20/volume-area-regular-solids/>

The volume of a sphere of radius r is $V = \frac{4}{3}\pi r^3$ and the surface area is $A = 4\pi r^2$ and so the ratio of volume to area is $V / A = r / 3$. Surprisingly, the same ratio holds for all regular solids if r is the radius of the largest sphere [?] The post Volume to Area ratio for Regular Solids first appeared on John D. Cook .

10. Solving a chess puzzle with Grok 4.5

Source: John D. Cook | Published: Mon, 20 Jul 2026 14:33:31 +0000

Link: <https://www.johndcook.com/blog/2026/07/20/grok-chess/>

I've written several posts about using Claude or ChatGPT to generate Prolog or Lean code to solve a chess puzzle. I didn't think Grok would be up to the task, though I didn't try it. I've heard good things about Grok 4.5, so I gave it a shot. It did great. Here's the problem, a [?] The post Solving a chess puzzle with Grok 4.5 first appeared on John D. Cook .

11. Memory Safety's Hardest Problem

Source: matklad | Published: 2026-07-20T00:00:00+00:00

Link: <https://matklad.github.io/2026/07/20/memory-safety-hardest-problem.html>

Uplifting a lobsters comment for easier reference.

12. Public Transport - Don't Make Me Think!

Source: Terence Eden's Blog | Published: Mon, 20 Jul 2026 11:34:23 +0000

Link: <https://shkspr.mobi/blog/2026/07/public-transport-dont-make-me-think/>

In the last year, I've been through over a dozen cities and used public transport in all of them. It is wild just how confusing and complex buying a ticket can be. While some cities obviously take a user-centred approach to ticketing, others appear to take a Kafkaesque delight in a bureaucratic maze of zones, apps, and intricate restrictions. One of the seminal texts in computer systems design?

13. Why I Stopped "Creating Content"

Source: Refactoring English | Published: Mon, 20 Jul 2026 00:00:00 +0000

Link: <https://refactoringenglish.com/blog/why-i-stopped-creating-content/>

The first time I saw Michelangelo's David in person, I wept. Never in my life have I witnessed a more beautiful piece of content. I wasn't the only one who felt that way. Museum patrons all around me were looking up from their phones to admire David as well. Despite it being legacy content, David was so scrollstoppingly captivating that it generated stronger engagement than TikTok and Instagram combined, despite those platforms offering fresher, more personally relevant content.

14. I added a blogroll

Source: matduggan.com | Published: Mon, 20 Jul 2026 11:29:27 GMT

Link: <https://matduggan.com/i-added-a-blogroll/>

I realized that it might be nice if you happen to stumble on this website if I had a way to recommend other websites you might enjoy. As it turns out this is a "blogroll", a concept I have never heard of before today but whatever. It'

15. OpenTK: nu ook met Internetconsultaties & notificaties

Source: Bert Hubert's writings | Published: Mon, 20 Jul 2026 13:37:00 +0200

Link: <https://berthub.eu/articles/posts/opentk-internetconsultaties/>

Sinds 2024 run ik OpenTK, met als doel zo goed mogelijk te laten zien wat er in de Tweede Kamer allemaal gebeurt. En, om daar ook nuttige email-notificaties over te versturen. OpenTK is begonnen omdat ik (vanuit de Internetwereld) ineens geconfronteerd werd met een wet die over DNS (mijn hobby) ging en over een paar dagen in stemming zou komen. Dat was heel erg schrikken en ik dacht, dat kan ons nooit meer gebeuren.