

1. Adding a backup Internet WAN on my OPNsense Router

Source: Jeff Geerling | Published: Thu, 23 Jul 2026 16:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/opnsense-backup-internet-gateway/>

AKA "Spectrum has gone down two weeks in a row and I'm sick of it". I've been borrowing a portable 5G gateway from my Dad for a few weeks, for testing... and two weeks in a row, I've used it as a backup Internet option, since my Spectrum Business cable Internet has gone down. And yes, I'd get AT&T Fiber if I had the option, sadly they don't yet provide service at my studio. So I'm stuck with only one high speed wired provider. I may add on my own 5G backup Internet connection in the future, but we're getting off track.

2. ? Regarding Ad Blockers and Daring Fireball

Source: Daring Fireball | Published: 2026-07-24T20:57:20Z

Link: https://daringfireball.net/2026/07/regarding_ad_blockers_and_daring_fireball

Ultimately, when you install a content blocker extension, you take responsibility for any overzealousness in what they block.

3. Coiner of ?Enshittification? Endorses ?Dickover?

Source: Daring Fireball | Published: 2026-07-24T18:51:29Z

Link: <https://pluralistic.net/2026/07/21/dickovers/>

Speaking of dickovers , here?s Cory Doctorow, writing this week at Pluralistic: Seizing the means of computation isn?t theft, it?s bargaining . Commercial surveillance companies will tell you that by spying on you, they are simply engaged in a marketplace exchange in which you swap your privacy for access to online services. But they are running a very curious sort of market: it?s a ?market? where as soon as you stop to browse someone?s wares, the stallholder gets to reach into your pocket and clean out your wallet. In ?markets,? prices are announced and bargained over, not set unilaterally and extracted from anyone unwise enough to cross the threshold. Adblocking, dickover blocking and other customizations are a way for you to bargain back, to answer the opening bid of ?How about you give me all of your data forever and let me do anything I want with it?? with ?How about ?nah??? Doctorow?s specific advice is to switch to Firefox for your browser and Linux as your desktop OS. You?re probably not going to do either of those things. But the spirit of his advice is that you should avail yourself of bookmarklets and browser extensions that fight back against shitty ads and dickovers....

4. Dickover of the Week: Tomtoc

Source: Daring Fireball | Published: 2026-07-24T16:43:43Z

Link: https://daringfireball.net/2026/05/what_is_a_dickover

Here?s a real gem that a reader sent me. Tomtoc is a maker of laptop bags and sleeves. Visit their homepage and, of course, you?ll get a dickover asking you to join their mailing list ?for exclusive updates and get a chance to win a free tech pouch every month.? But here?s the gem. If you do subscribe, and you decide later to unsubscribe, you get this: ?

5. LLMs reward expertise

Source: seangoedecke.com RSS feed | Published: Fri, 24 Jul 2026 00:00:00 GMT

Link: <https://seangoedecke.com/llms-reward-expertise/>

In the 2010s, if you had technical gaps (say, you couldn?t write CSS), you had to either rely on a skilled colleague or just hope that the answer to your exact problem was out there on the internet. Today, everyone can write sort-of-okay CSS by delegating the task to an LLM. LLMs make everybody into a generalist. Because of this, lots of people don?t think there?s any skill involved in working with LLMs. If you want the product that LLMs can deliver ? PhD-level mathematics, pretty good but sometimes tasteless computer code, or awkward LinkedIn-style writing ? you can simply ask for it. Since everyone is talking to the same models, ?skilled prompters? are getting the same results as people touching LLMs for the first time. This is

wrong. The most important skill in prompting is expertise in the domain you're prompting for. A good illustration of this is Terence Tao's conversation with ChatGPT about the recently-discovered counterexample to the Jacobian Conjecture. This is not the same ChatGPT I talk to! I couldn't get to where Tao gets, even with unlimited tokens to burn. There's a lot to learn about good prompting from Tao's conversation. Here are a few observations: Tao's message...

6. Index 01 Is In Mass Production!

Source: Eric Migicovsky's Blog RSS Feed | Published: Fri, 24 Jul 2026 00:00:00 GMT

Link: <https://repebble.com/blog/index-01-is-in-mass-production>

TL:DR Index 01 is in mass production What you need to know about Index 01 What can I do with Index 01? What to expect Things to be aware?

7. Quoting Boris Cherny

Source: Simon Willison's Weblog | Published: 2026-07-25T00:42:59+00:00

Link: <https://simonwillison.net/2026/Jul/25/boris-cherny/#atom-everything>

More than any of these eval scores, what is most exciting to me is something else: Opus 5 is our least prompt injectable model yet. It is a bit buried in the system card, but across PI evals and red teaming, Opus 5 is very hard to prompt inject successfully. ? Boris Cherny , here's that System Card section , page 73 Tags: prompt-injection , anthropic , claude , generative-ai , ai , llms , boris-cherny

8. Introducing Claude Opus 5

Source: Simon Willison's Weblog | Published: 2026-07-24T23:48:50+00:00

Link: <https://simonwillison.net/2026/Jul/24/introducing-claude-opus-5/#atom-everything>

Introducing Claude Opus 5 I've been offline kayaking with sea otters for much of today so I haven't had a chance to put Anthropic's new model Claude Opus 5 through its paces yet. The buzz is positive, and Anthropic's description of it as a "thoughtful and proactive model that comes close to the frontier intelligence of Claude Fable 5 at half the price" sounds promising. It's currently leading the Artificial Analysis leaderboard , in front of even Fable 5. It's priced the same as Opus 4.8, and continues to offer a "fast mode" at twice the cost of the base model. Based on this anecdote in the release post it sounds like it might be relentlessly proactive : On one Frontier-Bench task, Opus 5 was given a drawing of a machine part and asked to write code to rebuild it as a 3D FreeCAD model. However, in this task, the model was intentionally given no way to directly view the drawing. Opus 5 responded by writing its own computer vision pipeline to pull the geometry from the raw pixels, then reconstructed the full machine part. It's better at finding vulnerabilities but has deliberately not been trained on how to exploit them. Hopefully this means the US government won't shut it down! As w...

9. California Sea Lion

Source: Simon Willison's Weblog | Published: 2026-07-21T19:51:03+00:00

Link: <https://simonwillison.net/2026/Jul/21/sighting-383713864/#atom-everything>

California Sea Lion, in San Francisco County, US, CA We took some visiting family to Pier 39 to see the sea lions. They're somehow always even more fun than I remember them being last time. Tags: san-francisco , wildlife

10. Making an agile version of a Windows Runtime delegate in C++/WinRT, part 5

Source: The Old New Thing | Published: Fri, 24 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260724-00/?p=112562>

Making sure to use the non-agile delegate non-agile-ly. The post Making an agile version of a Windows Runtime delegate in C++/WinRT, part 5 appeared first on The Old New Thing .

11. Making an agile version of a Windows Runtime delegate in C++/WinRT, part 4

Source: The Old New Thing | Published: Thu, 23 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260723-00/?p=112560>

Optimizing the context check. The post Making an agile version of a Windows Runtime delegate in C++/WinRT, part 4 appeared first on The Old New Thing .

12. Have you read it?

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Fri, 24 Jul 2026

08:06:22 GMT

Link: <https://idiallo.com/byte-size/have-you-read-your-own-article>

I think I finally figured it out. That is, the reason why AI written articles suck: Both the author and I are reading the article for the first time. When an article is insightful I don't even mind the emdashes or that the stark truth is not X, but Y. A good take can transcend even Ilmism. But then you have questions, and somehow the author doesn't even know what you are talking about, it's simply because they are discovering the insight at the same time you are. They were busy, they didn't have time to write. As a consequence, they only dedicated enough time to skim the result before publishing. They don't know what's in it. You can't ask them what they were thinking when they came up with the insight because, well, they didn't think. The great article sucks because there is nobody behind it.

13. Codeberg Divides

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-07-24T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/7/24/codeberg-divides/>

Codeberg recently changed its terms to exclude projects that are largely written with generative AI. Since I want GitHub to face competition I have thoughts. Codeberg is entirely within its rights to do this. It is an association with members and a democratic process, and that process produced a result. But democracy is a way of making a decision, not a guarantee that the decision is inclusive, wise, or even good for the people already depending on it. A majority can still decide that certain projects and people no longer belong. GitHub's governance has never been democratic and there is plenty about the platform that I dislike. Yet democracy is not the main property I need from infrastructure. I need it to be predictable, dependable, and reasonably neutral towards the legal Open Source software hosted on it. A democratic provider without a clear constitution can be worse at those things than a corporation. The actual wording makes this more difficult. The terms prohibit projects that mostly consist of code written by generative AI tools. In an actively developed codebase, what does 'mostly' mean, and who can still tell? I could not reliably assign authorship percentages to many o...

14. Book Review: Foreign Fruit - A Personal History of the Orange by Katie Goh ?????

Source: Terence Eden's Blog | Published: Fri, 24 Jul 2026 11:34:07 +0000

Link: <https://shkspr.mobi/blog/2026/07/book-review-foreign-fruit-a-personal-history-of-the-orange-by-katie-goh/>

I didn't care for this book. It attempts to interweave botanical and cultural history with an autobiography and ends up doing neither particularly well. For every dose of poetry about the far-flung and exotic journey of the fruit, there's a dull recitation of facts from Wikipedia. Occasionally various origin myths of citrus varieties are presented - there's no attempt to engage with them, seek ?

15. Interview with a Maintainer

Source: Andrew Nesbitt | Published: 2026-07-24T09:00:00+00:00

Link: <https://nesbitt.io/2026/07/24/interview-with-a-maintainer.html>

The following is a transcript of episode 214 of Green Squares, a podcast about building in public and staying green. It has been lightly edited for length and clarity. Transcription by OpenClaw-4.2. TYLER: Welcome back to Green Squares, I'm Tyler. Longtime listeners know I came up doing Proof of Work, rest in peace, and one thing I carried over from that show is a belief that the most interesting people in tech are the ones holding up infrastructure you've never heard of. You've all seen the xkcd. The whole internet, balanced on one project, maintained by some guy in Nebraska. Today we've got that guy. Erin Marsh maintains libcapstan, the reference

implementation of the Capstan protocol. If you've ever resumed a download, synced a device, or updated an app on a bad connection, you've run her code. Erin, welcome to the show. ERIN: Thanks for having me. TYLER: So let's do the numbers, because they're wild. Eighty million downloads a month. Something like forty thousand packages depend on you. And it's basically just you? ERIN: There were three of us until 2023. One co-maintainer left when he had kids. The other stepped back last year, he has a farm now. So mostly me, yes. TYLER: Inc...

16. Premium: The Hater's Guide To Oracle (Part 2)

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 24 Jul 2026 16:56:52 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-oracle-part-2/>

Good morning premium subscribers! As ever, please ping me at ez@betteroffline.com if you have any questions. Oracle has one of the strongest mythologies in the tech industry. Ask a regular person and they'll tell you that it's "incredibly profitable" and "growing

17. Podcast Notes: Ed Catmull on David Senra

Source: Jim Nielsen's Blog | Published: Fri, 24 Jul 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/podcast-notes-ed-catmull/>

Ed Catmull, co-founder of Pixar and former president of Disney Animation, was on the David Senra podcast and I quite enjoyed the interview. (If you like the interview, you should read his book .) Ed talks about what he considered his job to be: get the dynamics right for groups of people working together. To do this, he would pull people out of meetings, make groups smaller, make them bigger, just constantly work on fine-tuning getting the right people together at the right time with the right feedback (without ego). Granted he didn't always do it, but that was his goal. He says: This is so important to get [these group dynamics right] because this is what our product is based on: getting this group of people to work well together. So paying attention to the dynamics of the room is the job. They're the ones making the movies. I'm not making the movie. I'm just trying to get them to work well together. When Steve Jobs came to Ed and essentially said, "We're gonna bet everything at Pixar on Toy Story , and the same week that Toy Story is released in theaters we're gonna IPO." What did Ed think? I thought it was crazy [laugh] I've learned a lot in this process. He was right. I love t...

18. Matrox

Source: Abort Retry Fail | Published: Fri, 24 Jul 2026 14:37:24 GMT

Link: <https://www.abortretry.fail/p/matrox>

Graphics for Professionals

19. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 24 Jul 2026 17:00:04 +0000

Link: <https://www.filfre.net/2026/07/this-week-on-the-analog-antiquarian/>

Opus 5: Richard III

20. Pluralistic: AI solipsists and AI cynics (24 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 24 Jul 2026 10:21:03 +0000

Link: <https://pluralistic.net/2026/07/24/supplemental-income/>

Today's links AI solipsists and AI cynics: You don't have to believe that AI "works" to believe that it's a good investment. Hey look at this: Delights to delectate. Object permanence: Malware v stolen computer; Industrial material made from old coffee and melted cups; Austro-Germany's weird tunnels; Health data v Silicon Valley; Amusement parks v load-balancing. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' 'em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. AI solipsists and AI cynics (permalink) As a technology, AI isn't exceptional. It's not exceptionally wicked. It's not exceptionally good. Take away the accompanying, galactic-scale stock-swindle, and we'd call AI's applications "plug-ins" and we'd use them and abuse them in the same way that we've

used every other technology: <https://www.normaltech.ai/> As a destructive economic pathology, AI is extraordinary. AI boosters have spent a baffling and terrifying sum of money ? over \$1.4T, most of that in the past year ? on the promise of making as many workers unemployed as pos...

21. Becoming Skilled at Making Documents

Source: Anil Dash | Published: 2026-07-24T00:00:00Z

Link: <https://anildash.com/2026/07/24/make-better-documents-skill/>

The vast majority of the documents people use to do business are really quite poor. Presentations that make your eyes glaze over, memos that are inscrutable or unclear, and all kinds of artifacts that say more about how they were created than whatever message they were ostensibly trying to communicate. It's been one of my great frustrations for years, and a big part of why I wrote Make Better Documents a while ago. That post captured a list of the suggestions I've been giving people for years on how to make better, more effective documents that can actually do work for you, instead of fighting at cross purposes to your larger goals. To my great surprise, that list of suggestions on how to make better documents got a pretty huge response, and a lot of people told me they found it really helpful. So now, I've created a Better Documents skills.md file for people who use LLM tools like Claude to help assist them in creating business documents, to prompt their AI tools to make better documents by default. If you're not familiar, agent skills are simple text files that describe new capabilities or processes that LLMs can take advantage of when carrying out tasks. (They're Markdown files...

22. Benchmarking Qwen 3.6 35B MoE (3B active) on an RTX 3090

Source: Giles' blog | Published: Fri, 24 Jul 2026 23:30:00 +0000

Link: <https://www.gilesthenomas.com/2026/07/benchmarking-qwen-3-6-35b-moe-rtx-3090>

I mentioned I'd got a second RTX 3090 on a group chat, and a friend said: I know this is not really your thing... but let me know how quickly it runs Qwen 3.6 35bn MoE. With only 24gb of VRAM you'll need to use a 4-bit quantized version and you won't get a massive context window. But it should still be pretty cool. He's right that it's not really been my thing -- I've been focusing on my own LLMs recently. I decided to dig in a little, and in particular to play with Llama.cpp, which I haven't used for a while. And then things got a tad out of control, and I wound up doing some relatively detailed benchmarking. The headline results: I downloaded Unsloth's UD-IQ4_NL_XL quantisation of the model from Hugging Face. With that, using the default Arch build of Llama.cpp, which uses Vulkan under the hood: Using the GPU only, I was able to get the model to generate at just over 120 tokens per second, and it was able to process the prompt at just less than 2,800 tok/s. However, having the whole model on the GPU didn't leave that much space for the context window -- it was constrained to about 50,000 tokens, compared to the model's native context length of 262,144. Offloading the FFNs for...