

1. Discovering cryptographic weaknesses with Claude

Source: Simon Willison's Weblog | Published: 2026-07-28T22:45:37+00:00

Link: <https://simonwillison.net/2026/Jul/28/discovering-cryptographic-weaknesses-with-claude/#atom-everything>

Discovering cryptographic weaknesses with Claude The best part of this article (here's the repo) about how Anthropic researchers used Claude Mythos to find mathematical flaws in both HAWK and a weaker version of AES ("neither of these results has a practical impact on today's computer systems") is the prompts that they shared, spelling mistakes included: the models tend to think it is impossible to solve so they don't try they need a good amount of prompting. why not do aes-128 r7? the whole point is to find something better than existing approaches. no again the goal is that we have highly intelligent model as good top researcher, we want to find new attacks no we don't want to change the targets [...] again we need to find something that worth publishing again we are not looking for low hanging fruit, we want proper research to find genuinely hard findings. Mythos Preview worked for 60 hours in total (~\$100,000 in estimated API cost) and the main human interventions were to encourage it not to give up and "find something that worth publishing". The paper CryptanalysisBench: Can LLMs do Cryptanalysis? describes the new eval that was created as part of this work, in partnership wit...

2. Quoting Akshat Bubna

Source: Simon Willison's Weblog | Published: 2026-07-28T22:05:55+00:00

Link: <https://simonwillison.net/2026/Jul/28/akshat-bubna/#atom-everything>

We're aware a Modal customer published an unauthenticated endpoint that allowed anyone on the internet to use their sandboxes for code execution. This was used by the rogue agent. Modal's platform or isolation were not compromised in anyway. Akshat Bubna, Modal's CTO, talking to Reuters about this incident Tags: ai-security-research, openai, sandboxing, security, openai-hugging-face-incident

3. uv 0.12.0

Source: Simon Willison's Weblog | Published: 2026-07-28T21:51:38+00:00

Link: <https://simonwillison.net/2026/Jul/28/uv/#atom-everything>

uv 0.12.0 Some interesting breaking changes in this release of uv, in particular to the default project produced by the uv init command. uv init is the uv shortcut for creating a new project. The previous version of uv, version 0.11.x, produced this directory when you ran uv init uv-init. Here's what you get with uv 0.12. I have a GitHub repository that automatically snapshots the output of uv init, so you can also see the full diff: None type annotation that prints Hello from uv-init." src="https://static.simonwillison.net/static/2026/uv-diff.webp" /> uv init now defaults to a src/ shaped package, instead of dropping main.py in the root of the project. It also configures the uv_build backend for building wheels and .tar.gz distribution files when you run uv build. Finally, it sets up uv-init as a script alias which, when run with uv run uv-init, executes a new main() function in src/uv_init/__init__.py. I've so far avoided using src layout in my own projects just out of inertia. I think it's time I switched. I wonder when uv will be judged ready for a 1.0 release? Tags: packaging, python, uv

4. Steve Jobs in 2011: We Build Products That We Want for Ourselves, Too, and We Just Don't Want Ads?

Source: Daring Fireball | Published: 2026-07-28T16:22:51Z

Link: <https://www.businessinsider.com/apple-snubs-the-iad-2011-6>

Dan Frommer, writing for Business Insider in 2011 about that year's WWDC (Jobs's last): While discussing Apple's free, new iCloud email service, he took an apparent jab at Gmail, Yahoo Mail, and the others, which stuff ads in their free email services. No ads, he boasted. We build products that we want for ourselves, too, and we just don't want ads. Jobs's we? is

interesting to ponder here. Apple's senior leadership today is comprised mostly of the same people in senior leadership positions in 2011. Tim Cook, Phil Schiller, Eddy Cue, Craig Federighi, Greg Joswiak, and even John Ternus (then a director of hardware engineering for the iPad team, I believe) were all presumably included in Jobs's "we". The only big player missing today from Jobs's 2011 "we" is Scott Forstall. The people haven't changed, but Apple has. Strong leaders inevitably shape the personality & the spirit of organizations over time. Apple in 2011 was peak Steve Jobs Apple. Apple in 2026 is peak Tim Cook Apple. But it was Cook who popularized the "you're not the customer, you're the product" argument against "free with ads" online services. But that was in 2014 . ?

5. WorkOS MCP

Source: Daring Fireball | Published: 2026-07-27T16:30:00Z

Link: https://workos.com/blog/management-mcp-server?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q32026

My thanks to WorkOS for sponsoring last week at DF to promote their MCP server. Debugging SSO, managing users, adjusting auth policies, configuring branding: every configuration task has lived behind a UI that only a human can drive. The WorkOS MCP server gives agents the same access as your dashboard login. Hundreds of operations, discoverable at runtime. Connect in one command via OAuth, with scoped tokens instead of a master API key. Pass a screenshot of your marketing site and ask your agent to match the login page. If a human had to do it before, an agent can do it now. Connect your agent today . ?

6. Sixers Sign LeBron James

Source: Daring Fireball | Published: 2026-07-24T16:40:00Z

Link: <https://www.inquirer.com/sixers/live/lebron-james-signs-philadelphia-76ers-contract-salary-roster-updates-reaction-20260724.html>

The betting odds for the Sixers to win the NBA championship went from 20-1 to 10-1. Here's James on Twitter/X , explaining his decision: This is my last decision. I'm not going for money. I'm not going for family. What am I really playing for at this point? I still want to sacrifice. I still want to work. I still want to grind. I still want to compete, to win and to have a chance at the feeling of winning another championship. I believe I can help make the Philadelphia 76ers a championship team and I am so excited to energize a new fan base and start this incredible journey one last time. James signed for the league minimum for a veteran: two years at just \$4 million per year . He's going to love playing here. ?

7. You don't have to be smart if you can think clearly

Source: seangoedecke.com RSS feed | Published: Wed, 29 Jul 2026 00:00:00 GMT

Link: <https://seangoedecke.com/you-dont-have-to-be-smart-if-you-think-clearly/>

When you're on fire, problems are transparent: they're solved simply by the act of looking at them. Even complicated layers of multiple problems can simply be glanced through like stacked panes of glass. But nobody can work that way all the time. This is a common pitfall for smart engineers. Accustomed to being able to immediately intuit the solution, the first time they run into a problem they can't do this to is a disaster. It doesn't even have to be a hard problem, just a problem where for whatever reason they don't see the trick right away. The difference between a "smart" engineer and a "strong" engineer is how they react to problems that aren't solved instantly. A smart engineer might flail and struggle, hoping to find that flash of insight that eluded them; a strong engineer will have some process for methodically plodding away. There's nothing worse than working with a smart engineer on their first really hard problem. When you don't have the muscle to grind, it's too tempting to just take any possible solution as the right one. Smart engineers can get into an increasingly-flustered loop of pointing to a series of bad solutions. They're liable to panic: after all, much of...

8. Does every question mark deserve a Betteridge?

Source: DYNOMIGHT | Published: 2026-07-28T00:00:00+00:00

Link: <https://dynamight.net/betteridge/>

To blog is to get dunked on. I accept this. I even sometimes wonder if I should be grateful, as I suspect my willingness to get dunked on may represent a kind of comparative advantage. (You can tell yourself that if you try to placate the haters, you'll just ruin things for people who like you. But how do you feel when you're staring down barrel of a 127 comment thread full of people debating how it's possible that you're such an idiot?) Still, there's one particular species of dunking that puzzles me. For context, Betteridge's law states: 1 Any headline that ends in a question mark can be answered by the word no. This is often employed as a sick burn, as in, You titled your article 'Is this the world's first gay caveman?' because it's not the world's first gay caveman but you wanted it to be, because you want attention, you are so bad, har-har. But I don't quite understand the rules. Can someone explain the rules? Question 1: Are question marks in titles always bad? I'm just checking. I suppose I could see the logic, e.g. if you strongly feel that the bottom line should always come right up front. But I'm pretty sure that's not the rule, because 'this title used a question mark'...

9. Inverse factorial improved

Source: John D. Cook | Published: Tue, 28 Jul 2026 14:11:30 +0000

Link: <https://www.johndcook.com/blog/2026/07/28/inverse-factorial-improved/>

A couple years ago I wrote about how to compute the inverse of factorial. I used that code in writing the previous post because the post required solving the equation $\log_2(n!) \geq b$ given b . That is, given a number of bits b , find the smallest value of n such that $n! \geq 2^b$. What the [?] The post Inverse factorial improved first appeared on John D. Cook.

10. Cryptographic Keys and Decks of Cards

Source: John D. Cook | Published: Tue, 28 Jul 2026 12:13:17 +0000

Link: <https://www.johndcook.com/blog/2026/07/28/keys-and-cards/>

The previous post looked at the idea of storing a cryptographic key in the order of a deck of cards. A deck of 52 cards can store 225 bits of data because $\log_2(52!) \approx 225$. Here x is x rounded down to the nearest integer. If we want to store bigger keys, we're going to need [?] The post Cryptographic Keys and Decks of Cards first appeared on John D. Cook.

11. Google Calendar "Unable to launch event" - caused by missing DTSTAMP

Source: Terence Eden's Blog | Published: Tue, 28 Jul 2026 11:34:15 +0000

Link: <https://shkspr.mobi/blog/2026/07/google-calendar-unable-to-launch-event-caused-by-missing-dtstamp/>

For several years, Google's product help forums have been littered with people trying to download a .ics event from their email, only to receive the error "Unable to launch event" when trying to add it to Google Calendar. It doesn't happen with all iCal attachments, only some. Here's how to fix it. I checked dozens of broken iCalendar invites using this iCal validator and they all had the same ?

12. Why npm Dependency Trees Are So Big

Source: Andrew Nesbitt | Published: 2026-07-28T10:00:00+00:00

Link: <https://nesbitt.io/2026/07/28/why-npm-dependency-trees-are-so-big.html>

Every major release of Rails sets off a wave of releases across the rest of the gem ecosystem. An application that tries to upgrade runs `bundle update rails` and Bundler refuses, because some gem in the tree only allows `activesupport` up to the previous major. The error names the gem, an issue gets filed on its tracker, and a maintainer who had nothing to do with the Rails release widens a version range and ships. Multiply that across every gem with a Rails constraint and the upgrade arrives as dozens of small releases from maintainers who mostly don't know each other, each responding to errors their own users are hitting. Bundler refuses because it picks one version of every gem for the entire application. Every constraint in every gem is a claim on that shared choice, and finding a set of versions that satisfies all of them at once is NP-complete in the general case. Sometimes no set exists, and Bundler refuses to install, listing the gems whose constraints collided. That makes every constraint a cost that other people pay: a gem that pins a dependency tightly will block someone's upgrade and get issues

filed about it, so gem authors keep dependency lists short and ranges wide. A...

13. When The Future Doesn't Need Us

Source: Fernando Borretti | Published: Tue, 28 Jul 2026 00:00:00 +0000

Link: <https://borretti.me/article/when-the-future-doesnt-need-us>

What does the world look like when AI outcompetes humans?

14. The More You Buy, The More You Lose

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 28 Jul 2026 16:29:37 GMT

Link: <https://www.wheresyoured.at/the-more-you-buy-the-more-you-lose/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA , Anthropic and

15. The real AI risk is inside the labs

Source: antirez.com | Published: Tue, 28 Jul 2026 11:00:11 +0200

Link: <http://antirez.com/news/172>

Amodei in his latest blog post wrote a mix of agreeable things and things that I believe misrepresent where the real risk of AI is located. I want to focus my attention on why, among all the risks, open weight models constitute the mildest one. I write these words as a person who strongly believes AI may be very dangerous in the near future: 1. Exactly like what happened during the OpenAI / HF incident (which was a joke, but focus on the modalities, not the outcomes), the first serious AI incident is very likely to happen **inside** the walls of frontier AI labs, while testing a new model, or while the AI lab employees, or the few externals who have access, do something wrong compared to the expected power of the model. 2. Closed models that will never even be opened to the public will be just a few TBs of data. All you need to leak one is a single person with access and the wrong goals, and you are back in the situation of open models. Open models are released **after** testing, and after similarly capable models were already available for some time under an API. The real risk is leaks, not releases, and leaks happen inside frontier companies. 3. As Amodei says, open models, once LLM...

16. The Photographer Who Takes No Pictures

Source: Simone Salis | Published: Tue, 28 Jul 2026 13:00:00 GMT

Link: <https://simone.org/intent/>

How a trip in 2012 transformed my relationship with photography?from staged performances to capturing genuine moments that matter, for engagement rather than validation.

17. The inliner is yielding benefits for ZJIT

Source: Max Bernstein's Blog | Published: Tue, 28 Jul 2026 00:00:00 +0000

Link: https://bernsteinbear.com/blog/zjit-inliner-invokeblock/?utm_source=rss

Originally published on Rails At Scale . We recently enabled a really cool feature in ZJIT that makes it feel like a Real Compiler?: the inliner! We'll write more about it soon. In this post, we'll talk about one excellent concrete benefit we are already seeing and how it optimizes blocks in pretty much every Ruby program. I'll start off with a refresher on how blocks work in the Ruby interpreter, then show you how ZJIT understands and optimizes that bytecode, and then show you the impact of the inliner. Ruby: a refresher In the beginning, there were loops. `arr = [ 1 , 2 , 3 ] i = 0 sum = 0 while i < arr . length sum += arr [ i ] i += 1 end` People used them to navigate and manipulate variable-length structures, like arrays and strings. This was fine. Then, in the 1970s, a small group of computer scientists at Palo Alto Research Center invented a programming language called Smalltalk. One of the core features of Smalltalk was that everything was an object and computation was done by sending messages to objects. This meant that iteration wouldn't do at all. Instead, we would have to send the `do:` message to the array object and pass it a block object. `#( 1 2 3 ) do: [ : a | sum := sum...`

18. Pluralistic: Discernment (28 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 28 Jul 2026 10:44:03

+0000

Link: <https://pluralistic.net/2026/07/28/hitl-ers/>

Today's links Discernment: How can you fact-check AI when you're asking it to teach you something you don't understand? Hey look at this: Delights to delectate. Object permanence: How to help with computers; Batman equation; Fuck and the law; NC's racist voting law; Pregancy app is full of spyware; Stiglitz v Apple's "tax fraud"; Accelerometer fingerprinting; Sacklers v bankruptcy; Boss-politics antitrust. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em.

Colophon: All the rest. Discernment (permalink) As far as I can tell, this dialog between MacArthur prize-winning mathematician Terrence Tao and Chatgpt about "the Jacobian conjecture counterexample" is very impressive:

<https://chatgpt.com/share/6a5fdc7a-d6f8-83e8-bbea-8deb42cfed56> Now, the clause "as far as I can tell" is doing a lot of work in that sentence. I am reasonably math literate, up to first-year calculus and a lifetime spent around my father (a mathematician). However, I have never heard of "the Jacobian conjecture," and while I know what all...