

1. Advancing the price-performance frontier with GPT-5.6

Source: Simon Willison's Weblog | Published: 2026-07-30T23:58:42+00:00

Link: <https://simonwillison.net/2026/Jul/30/luna-price-drop/#atom-everything>

Advancing the price-performance frontier with GPT-5.6 Huge price drop from OpenAI today: GPT-5.6 Terra got a 20% reduction, and GPT-5.6 Luna got a massive 80% drop. OpenAI credit 5.6 Sol with enabling this: in How GPT-5.6 fuses frontier intelligence with frontier efficiency they describe using 5.6 Sol to optimize load balancing, and more impressively to optimize inference itself: We also used GPT-5.6 Sol to optimize the model's forward pass: the computation that transforms inputs into next-token predictions. Even when individual operations are fast, excess memory movement, synchronization, and inefficient data layouts can leave GPUs idle. To avoid this, GPT-5.6 Sol found work that could be precomputed, avoided, or parallelized. With Codex, GPT-5.6 Sol autonomously rewrote and optimized our production kernels, the core code that executes the mathematical operations that make up the model. This worked in part because we've trained GPT-5.6 to be effective at writing and improving kernels in Triton and Glue, two open-source GPU programming languages maintained by OpenAI. These efforts, combined with broader kernel advancements from GPT-5.6 Sol, reduced end-to-end serving costs by...

2. Investigating three real-world incidents in our cybersecurity evaluations

Source: Simon Willison's Weblog | Published: 2026-07-30T23:41:29+00:00

Link: <https://simonwillison.net/2026/Jul/30/three-real-world-incidents/#atom-everything>

Investigating three real-world incidents in our cybersecurity evaluations It happened again! This is turning into something of a pattern. Last week OpenAI accidentally exploited Hugging Face when one of their frontier models broke out of a sandboxed container and hacked into Hugging Face to try and get the solutions to the cyber benchmark it was executing. This inspired Anthropic to double-check their own logs, and it turned out they had three similar (albeit less impressive) incidents, the earliest of which played out in April! Of the 141,006 evaluation runs we reviewed, we identified three separate incidents (involving six total runs, four of which impacted the same organization; the other two incidents each happened in independent evaluation runs). [...] In all cases, Anthropic's evaluation prompt specified to Claude that its environment was a simulation and that it had no internet access. Due to a misunderstanding between us and our evaluation partner, this was not the case, and internet access was available. Because of this, when Claude's search led it to real systems on the open internet, it treated them as part of the exercise. [...] Operating under the false belief that al...

3. Llm 0.32rc2

Source: Simon Willison's Weblog | Published: 2026-07-30T22:52:06+00:00

Link: <https://simonwillison.net/2026/Jul/30/llm-rc2/#atom-everything>

Release: Llm 0.32rc2 Hot on the heels of RC1, this fixes a dependency issue and also adds two neat new features: The default model for users who have not set their own default is now GPT-5.6 Luna. It was previously GPT-4o mini. Luna is a much better and more recent model, albeit slightly more expensive - \$0.20 per million input tokens and \$1.20 per million output tokens, compared to \$0.15/\$0.60 for 4o mini. You can switch back to 4o mini using Llm models default gpt-4o-mini, or switch to GPT-5 nano, an even cheaper default model (\$0.05/\$0.40), using Llm models default gpt-5-nano. #1576 New Llm openai endpoint command for running prompts, chats and model listings against arbitrary OpenAI-compatible endpoints without first configuring a model. These calls are not logged. #1565 The Llm openai endpoint command is really cool. I got frustrated at the lack of an obvious CLI tool for trying out prompts against arbitrary OpenAI Chat Completions imitation endpoints, so I decided to add that to LLM itself. You don't even have to install LLM to use this. Here's a uvx one-liner for running a prompt - with tools - against an LM Studio local model: uvx --pre Llm openai endpoint http://127...

4. Mark Zuckerberg: 'The AI Future Is for Everyone'

Source: Daring Fireball | Published: 2026-07-30T21:09:51Z

Link: <https://www.wsj.com/opinion/the-ai-future-is-for-everyone-a0c24e20?st=T6AAwM>

Mark Zuckerberg, in an op-ed Tuesday for The Wall Street Journal (gift link; and irony isn't lost that 'everyone' needs a paid WSJ subscription to read this): We are fortunate to live at an incredible moment in history. In the next few years, people will be able to use superintelligence beyond human capacity to create and discover extraordinary new things, build new businesses, express our ideas, learn new concepts, and improve our lives, health, relationships and careers. Perhaps we'll be able to use it to build a metaverse too. Or to keep our attention on a major initiative for more than two or three years. ?

5. Apple Releases iOS and macOS 26.6, macOS 15.7.8, and More

Source: Daring Fireball | Published: 2026-07-30T20:41:43Z

Link: <https://arstechnica.com/gadgets/2026/07/ios-and-macos-26-6-arrive-today-paving-the-way-for-ios-and-macos-27/>

Samuel Axon, Ars Technica: Apple has released iOS, iPadOS, macOS, watchOS, and tvOS 26.6. Apart from potential security hotfixes, these are likely the last updates before the arrival of iOS 27, macOS 27, and so on. All of today's releases include minor bug fixes, and there are numerous security updates: more than 150 for macOS 26.6 . Apple also rolled out macOS 14.8.8 and macOS 15.7.8 for older devices, also focused on security fixes. In terms of new features, you won't find many in these releases, as they mainly pave the way for the next major OS update, likely to hit sometime in September. Most notably, the release notes for iOS and iPadOS 26.6 say the update 'optimizes the Spotlight index to prepare for iOS 27.' This update will kick off some indexing work that will then be leveraged in an ostensibly much more robust Spotlight search feature when iOS 27 launches next month. I don't mean to pick nits, but there's nothing 'ostensible' about it. You can use the new Spotlight index via Siri AI in the OS 27 beta releases and it's really good. The initial background indexing took almost a full week for me on iOS, I suspect because I have so much email archived, but anyone who upgrade...

6. Rogue Amoeba: Unobtrusive Update Notifications

Source: Daring Fireball | Published: 2026-07-30T20:33:59Z

Link: <https://weblog.rogueamoeba.com/2026/04/28/unobtrusive-update-notifications/>

Paul Kafasis, on the Rogue Amoeba blog back in April: Though Sparkle serves us very well, it has one notable downside. Update announcements are most likely to appear at the least convenient time: right after you've launched the app. You want to start recording with Audio Hijack, for instance, but the app is telling you about a new version. We've long wished to avoid these disruptions. With that in mind, we're making changes to how update notifications appear throughout our apps. In the future, when the software's timed automated check detects a newer version, it will no longer pop an obtrusive window like the one seen above. Instead, a small 'Update Available' indicator will be shown in the app's interface. Such a little thing, but like I wrote earlier this week , getting all the little things right is how you get to insanely great. Every Mac app using Sparkle ought to copy this. It's so much nicer. ?

7. BI Slop

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Thu, 30 Jul 2026 23:50:00 GMT

Link: <https://idiallo.com/blog/business-intelligence-slop>

You must use the tool. We must justify it. These aren't the words companies use, but that's what it comes down to. When you buy an expensive solution before you've clearly identified the problem, you have to use the tool to justify its line item in the spreadsheet. At most companies I've worked in, there's always some application I've used once or twice. The third time I want to use it, it's gone. If a tool isn't popular, the finance team has no trouble dropping it. But for some reason, it's different with AI. If nobody's using it, the assumption is that employees aren't being productive enough. So we got signed up for mandatory training.

There was limited space, but I moved fast and got a spot. I'd been using ChatGPT since it came out, so I was hoping the training would turn me from a casual user into a pro. Instead, we were shown how to copy and paste data from Excel and other sources into the chat interface. We worked with sample data, and there was always someone in class who'd say "mine didn't work." The developers in the room asked about Codex. The OpenAI-certified instructor replied that she wasn't a developer. There was an awkward moment when someone asked if we could just...

8. You have been misled about lightbulbs

Source: Maurycy's Blog | Published: Thu, 30 Jul 2026 00:00:00 +0000

Link: <https://maurycyz.com/misc/tungsten/>

There's a story that goes something like this: In 1925, lightbulb manufacturers secretly colluded to standardized lifespans at 1,000 hours. They would test each other's products to ensure compliance. This is true. At the time, many bulbs lasted longer than one thousand hours. This is true. Therefore, this was done so that people would always need to buy more light bulbs. This is wrong, but it's the type of wrong that cites its sources and hides in the part you'd never think to fact check: The assumption that a longer lasting lightbulb is a good product. In truth, increasing the lifespan of a bulb makes it worse in every other way... but people think they want a long lasting lightbulb: the purpose of standardizing was (primarily) to avoid a race to the bottom. A old-school lightbulb is a rather simple device : A thin tungsten wire (~20 μ m) sealed inside a glass envelope to protect it from air. When current is applied, the wire gets white hot and starts glowing. The most important parameter of a lightbulb is how hot that wire gets: this controls the peak emission wavelength (color) and brightness of the lamp. Room temperature objects do emit light (this is how thermal cameras work),...

9. So you want to use plants to reduce indoor CO?

Source: DYNOMIGHT | Published: 2026-07-30T00:00:00+00:00

Link: <https://dynamight.net/plants/>

Humans make carbon dioxide. Carbon dioxide is bad for cognition. But plants turn carbon dioxide back into oxygen. And plants are the one true home decoration strategy. So maybe if you get a lot of plants, you can keep carbon dioxide in check and keep your brain working? It's theoretically possible. It's probably just barely possible in practice. But it won't be easy. People produce ~1 kilogram of carbon dioxide per day. That's around 5.7×10^{23} molecules or 0.948 moles per hour. (You may remember from high school that a mole is a gigantic number made up to avoid having factors of 10^{23} everywhere.) Let's keep it simple and call it one mole per hour. Meanwhile, plants turn carbon dioxide into oxygen through photosynthesis, i.e. the chemical reaction of (6 water molecules) + (6 carbon dioxide molecules) + (energy) $\rightarrow$ (1 glucose molecule) + (6 oxygen molecules). The minimum energy physically needed to convert 1 mole of carbon dioxide into glucose and oxygen via this reaction is ~477 kilojoules. So we've already got a lower bound. Say you have magical plants that somehow channel all incoming energy into photosynthesis with perfect efficiency. They'll need ~477 kilojoules per hour...

10. Read This Before You Buy That TV Streaming Stick

Source: Krebs on Security | Published: Thu, 30 Jul 2026 16:49:00 +0000

Link: <https://krebsonsecurity.com/2026/07/read-this-before-you-buy-that-tv-streaming-stick/>

Security experts have been sounding the alarm for years about the risks of using generic TV boxes that promise unlimited content streaming for a one-time fee, warning that they secretly rent the user's Internet connection out to strangers. But a groundbreaking new analysis finds these devices also routinely spoof themselves as mobile phones clicking ads on AI-generated websites as part of sprawling operation that seeks to defraud online merchants and advertising networks.

11. Wheels, Bottles and Images

Source: Andrew Nesbitt | Published: 2026-07-30T10:00:00+00:00

Link: <https://nesbitt.io/2026/07/30/wheels-bottles-images.html>

In February Simon Willison packaged a Go CLI as Python wheels , one wheel per platform with the compiled binary inside, so pip install sqlite-scanner fetches a program containing no Python. His build tool maps each Go target to a wheel tag on the way through: darwin/arm64 to macosx_11_0_arm64 , linux/amd64 to manylinux_2_17_x86_64 . PyPI accepted the uploads and pip picked the right platforms, and the thread that followed was about whether PyPI should store binaries that serve no Python purpose. Wheels, Homebrew bottles and OCI images are the three big self-service binary distribution systems: the project publishes a compiled artifact and any machine fetches the one matching its platform, with no distro maintainer in between. conda, NuGet's native assets, RubyGems platform gems and plain GitHub Releases do the same job in their own ecosystems. The three big ones share no specification: each was designed against its own stack. Simon's trick works because a wheel is a checksummed archive selected by a platform tag, and that description also fits a bottle and an image. Immutable checksummed files are what CDNs serve cheaply and caches hold safely, and all three systems depend on both...

12. Why Is Everyone Trying to Build a Solid-State Battery?

Source: Construction Physics | Published: Thu, 30 Jul 2026 12:04:25 GMT

Link: <https://www.construction-physics.com/p/why-is-everyone-trying-to-build-a>

A battery technology that's getting a lot of attention is solid-state batteries, lithium-ion batteries that replace the liquid electrolyte with a solid material.

13. Pluralistic: The stupidest imaginable excuses for surveillance pricing (30 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Thu, 30 Jul 2026 12:51:03

+0000

Link: <https://pluralistic.net/2026/07/30/pay-for-privacy/>

Today's links The stupidest imaginable excuses for surveillance pricing: Since 1850, San Francisco's Chamber of Commerce has been pissing in our mouths and calling it rain. Hey look at this: Delights to delectate. Object permanence: RIP Wau Holland; Diebold voting machines suck; How to steal an RFID-locked car; Pirates are big spenders; You can't fight enshittification. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The stupidest imaginable excuses for surveillance pricing (permalink) With "surveillance pricing," businesses have finally found something AI can do way, way, way better than people: price gouging. As the name suggests, "surveillance pricing" is the practice of charging every customer a different price for every transaction, based on the massive surveillance dossiers that Big Tech companies and data-brokers have assembled on everyone in the world. Congress hasn't updated federal consumer privacy law since 1988 (when they passed a law banning the disclosure of VHS ren...

14. Why do OpenAI's GPT-2 weights beat mine? Part two: the bugfix

Source: Giles' blog | Published: Thu, 30 Jul 2026 19:00:00 +0000

Link: <https://www.gilesthomas.com/2026/07/why-do-openai-gpt2-weights-beat-mine-2-the-bugfix>

I'm digging into why my GPT-2 style models score worse on an instruction-following eval than OpenAI's original weights; I gave the details in this post . While I was writing up the results of my first experiment into possible causes, I ran the post past ChatGPT -- I always use an "editorial board" of AIs to check my posts for flow, style, and any technical errors (though all writing is always mine). It took a look at the eval code that I was running, and highlighted a bug. Luckily, it doesn't change the important results -- OpenAI's models continue to be better than mine at instruction-following. But it was enough to change the baseline numbers, re-ordering how well my own models did. So I fixed it and regenerated the baseline so that future experiments are based on solid ground. The bug The eval takes a model, and trains it over multiple epochs on a split of a subset of the Alpaca instruction-following dataset. At the end of each epoch, it evaluates the resulting model against a held-back validation split; if the eval loss starts rising, it bails out. Finally, it runs a test split of the dataset through the resulting model, and saves the result. Once I've run it for a bunch of di...

15. Why do OpenAI's GPT-2 weights beat mine? Part three: testing overtraining

Source: Giles' blog | Published: Fri, 31 Jul 2026 01:15:00 +0000

Link: <https://www.gilesthomas.com/2026/07/why-do-openai-gpt2-weights-beat-mine-3-overtraining>

The GPT-2-style models that I've been training work really well, and I've even managed to train some that perform better than the original OpenAI small model in terms of cross entropy loss on a test set. But as I wrote previously, there's a mystery: why do they perform worse on my instruction fine-tuning evaluation? I had various theories about why that might be, and to me, the most plausible-seeming of them was the amount of data they were trained with. As best I can find out, OpenAI's models were, by modern standards, trained on much more data than they should have been, while I'd used the theoretically optimal amount of training data. To put it in other words, OpenAI's models were overtrained. If I deliberately overtrained my own models, could I match their performance? This post is a write-up of what happened, but so as not to bury the lede -- it didn't seem to help much, if at all. Let's see why. Overtraining Let's start by getting a nice crisp definition of overtraining. It's important not to confuse overtraining with overfitting. Overfitting is where you train a model so that instead of learning a general rule about the data it's seeing, it learns something very specific t...

16. The AI Aesthetic

Source: Jim Nielsen's Blog | Published: Wed, 29 Jul 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/ai-aesthetic/>

Every zeitgeist comes with new design idioms unique to its challenges. Many of them disappear as fads change, but others bake themselves into deeper parts of existing software interaction paradigms. For example, there's the hamburger menu (?) which saw a proliferation during the rise of mobile due to the constraints around screen size. It has since spread to many other parts of software interaction design and will likely remain prevalent for a long time as a terse way of indicating "more menu-type content here". As another example, before AI what were the connotations of the sparkle emoji ?? Personally, I don't know, but now it means AI. (AI = sparkles and rainbow colors ? it's funny when you think about it. They should've just thrown unicorns in there for the trifecta. AI = sparkles, rainbows, and unicorns ??? Apt.) Some patterns are very specific to the interactions inherent to the nature of AI as a technology. For example: streaming text. This is a pattern made for and refined by chat interfaces, so it may not have tons of utility for reuse across other software interaction paradigms. Then there are other patterns that've been refined by AI interfaces and are starting to spre...