

1. Getting 25 Gbps Thunderbolt Ethernet on my Mac Studio

Source: Jeff Geerling | Published: Fri, 31 Jul 2026 09:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/getting-25g-ethernet-mac-thunderbolt/>

I've been using the built-in 10 Gigabit Ethernet on my Mac Studio for a few years. It works fine: I can edit 4K video straight off my NAS over the network, and run backups at around 1 GB/sec. But... I want more. I upgraded my rack and my NAS to 25 GbE a couple years ago, and wanted to upgrade my main workstation, too.

2. deepseek-ai/DeepSeek-V4-Flash-0731

Source: Simon Willison's Weblog | Published: 2026-07-31T23:59:44+00:00

Link: <https://simonwillison.net/2026/Jul/31/deepseek-v4-flash-0731/#atom-everything>

deepseek-ai/DeepSeek-V4-Flash-0731 The latest release in DeepSeek's V4 family, "with substantially enhanced agentic capabilities". It's 304 billion parameters - 167GB on Hugging Face - but it appears to punch well above its weight. Artificial Analysis rank it ahead of MiniMax M3 - a 428B model. It's \$0.14/million input and \$0.27/million output pricing means this may currently be the best value-per-intelligence model out there. It's looking very good on the Intelligence Index vs. Cost per Intelligence Index Task chart: I got a disappointing pelican from it using the default reasoning level via OpenRouter: But when I bumped reasoning level up to high I got something much better: `llm -m openrouter/deepseek/deepseek-v4-flash-0731 -t pelican -o reasoning_effort high` Via Hacker News Tags: ai, generative-ai, llms, pelican-riding-a-bicycle, deepseek, llm-release, openrouter, ai-in-china, artificial-analysis

3. Stateless MCP has recaptured my interest (and inspired mcp-explorer and datasette-mcp)

Source: Simon Willison's Weblog | Published: 2026-07-31T23:13:22+00:00

Link: <https://simonwillison.net/2026/Jul/31/stateless-mcp/#atom-everything>

Tuesday was Stateless MCP day - the rollout of MCP 2.0, or the 2026-07-28 Model Context Protocol specification to use the more formal but less memorable name. This is the most significant change to the MCP spec since it first launched, and has also served to reignite my personal interest in the protocol. For background: MCP is the Model Context Protocol, which describes a standard way to expose new tools to LLM-powered agent frameworks. It was introduced by Anthropic back in November 2024, had a huge spike of interest through much of 2025, and then became somewhat eclipsed by Skills (another Anthropic invention) when it became apparent that an agent harness with access to a terminal and curl could do most of what MCP did in a more flexible way. I wrote about that in my review of 2025. I'm coming back around to MCP now. Giving an agent a shell environment with the ability to access the internet is fraught with risk, and requires a strong model that is capable of effectively driving such an environment. MCP tools are easier to audit and control, and simple enough that smaller models that run on a laptop can still drive them reasonably well. The new stateless MCP specification als...

4. llm-mcp-client 0.1a0

Source: Simon Willison's Weblog | Published: 2026-07-31T23:03:47+00:00

Link: <https://simonwillison.net/2026/Jul/31/llm-mcp-client/#atom-everything>

Release: llm-mcp-client 0.1a0 See this blog entry. Tags: llm, model-context-protocol

5. AI models need moral support to make discoveries

Source: seangoedecke.com RSS feed | Published: Fri, 31 Jul 2026 00:00:00 GMT

Link: <https://seangoedecke.com/ai-models-need-moral-support/>

One recent development in AI is its ability to solve some long-standing problems in mathematics. In 2024 and 2025, this was a trickle: once or twice a year somebody would say that an LLM came up with a proof, and then everyone would argue over whether that counted as ?real? mathematical innovation. In 2026, it?s a flood. Almost every day I see some new LLM-produced

mathematical result. Prompt ?engineering? Perhaps the most curious thing about these AI discoveries is how easy the prompting is. The strategy for prompting Claude Mythos to come up with a cryptographic breakthrough appears to be just asking ?hey, please come up with a breakthrough?, and then checking in every few hours to say ?keep looking for something important, I want you to solve a genuinely hard problem?. It?s amusing to read this and remember how in 2025 everyone was obsessed with ?prompt engineering?. At the time I was something of a heretic for saying that prompts didn?t matter that much , but in hindsight I was clearly correct. The main skill involved in using LLMs is figuring out what they?re good at and what they?re bad at (and staying up-to-date as that rapidly changes). If you?re asking the LLM to do somet...

6. Making an agile version of a Windows Runtime delegate in C++/WinRT, part 9

Source: The Old New Thing | Published: Thu, 30 Jul 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260730-00/?p=112573>

Maybe it doesn't matter. The post Making an agile version of a Windows Runtime delegate in C++/WinRT, part 9 appeared first on The Old New Thing .

7. The Talk Show: ?What?s in Louie?s Wallet?

Source: Daring Fireball | Published: 2026-07-31T22:50:04Z

Link: <https://daringfireball.net/thetalkshow/2026/07/31/ep-453>

Louie Mantia returns to the show to talk about the state of UI and icon design on Apple?s platforms, and some speculation on Apple?s trade secret lawsuit against OpenAI. Sponsored by: Notion : The collaborative AI workspace where teams and agents work side by side, with a Developer Platform teams can build on. Even Realities : Even G2, the everyday display smart glasses. Use promo code talkshow to save 10% off the R1 Ring and/or Even Clip. Squarespace : Save 10% off your first purchase of a website or domain using code TALKSHOW . ?

8. ? Temu Is a Comically Bad App

Source: Daring Fireball | Published: 2026-07-31T18:27:50Z

Link: https://daringfireball.net/2026/07/temu_is_a_comically_bad_app

Shop like a billionaire.

9. Looking for the Catch in Apple Upgrade

Source: Daring Fireball | Published: 2026-07-30T16:59:39Z

Link: <https://www.theatlantic.com/technology/2026/07/apple-lease-upgrade-program/688106/>

Damon Beres, writing for The Atlantic under the hed/subhed: ?The New iPhone Underclass: Apple?s rental program is a trap?: The Klarna plan ? ?Apple Upgrade,? which replaces the iPhone Upgrade Program ? is truly, legally, a lease. This is confusing! And it?s confusing in part because this is not how Klarna, a well-known buy-now, pay-later service, typically operates: When you use Klarna for clothing from Shein or lip kits at Sephora or an Xbox at GameStop, you?re paying back a loan, exactly as you were in the original iPhone Upgrade Program. Same if you use Klarna to buy a Samsung Galaxy phone. But with Apple Upgrade, you are renting: The Mac or Apple Watch is not yours until the final payment is made. I don?t think this is confusing at all. Apple Upgrade is the primary brand for this program, not Klarna. Klarna is really only mentioned in the small print. You get into Apple Upgrade through Apple . Off the top of your head, do you remember Apple?s bank partner for the now-discontinued iPhone Upgrade Program? (It was Citizens Bank .) The Samsung program Beres links to above is named ?Klarna Pay in 4?. ?Apple Upgrade is a leasing program partnered with Klarna? is easily understood....

10. Energizing a vacuum-tube flip-flop module from a 1948 IBM system

Source: Ken Shirriff's blog | Published: 2026-07-31T09:04:31.698-07:00

Link: <http://www.rightho.com/2026/07/ibm-604-trigger-tube-module.html>

In 1948, IBM introduced the 604 Electronic Calculating Punch. This machine was a programmable calculator, about the size of a double refrigerator. It was not quite a computer, but was programmed by plugging wires into a plugboard. This machine read numbers from a punch card,

performed up to 60 calculations on these numbers, and then recorded the results by punching holes in the card. 1 It processed 100 cards per minute?over one card per second?and IBM advertised it as the equivalent of 150 engineers. The machine rented for \$550 a month, making it very popular, with over 5600 units produced. 2 The IBM 604 Electronic Calculating Punch. Photo from Ed Thelen's IBM 604 page . The IBM 604 came out just after the transistor was invented, too early to use transistors. At the time, calculators and computers were moving from slow electromechanical components to fast vacuum tubes. One of the innovations of the 604 was to combine a vacuum tube and its associated circuitry into a pluggable module. Along the left side of the photo above, you can see rows of these modules with the handles sticking out, making it easy to replace a faulty module. More modules are behind the silver metal covers. In...

11. Solving the RK4 design equations

Source: John D. Cook | Published: Fri, 31 Jul 2026 18:59:02 +0000

Link: <https://www.johndcook.com/blog/2026/07/31/runge-kutta-design/>

I was digging into the Runge-Kutta method for solving differential equations and a line from [1] piqued my curiosity. These calculations, which are not reproduced in Kutta's paper (they are however in Huen (1900)), are very tedious. The calculations are a set of eight constraints that the parameters of a fourth-order Runge-Kutta method must satisfy. [?] The post Solving the RK4 design equations first appeared on John D. Cook .

12. Notes from July 2026

Source: Evan Hahn (dot com) | Published: Fri, 31 Jul 2026 06:00:00 +0000

Link: <https://evanhahn.com/notes-from-july-2026/>

July! What a month! Things I did this July Gave a talk, for ClimateAction.tech?s ?Responsible AI Series?, ?What I know about local language models? . ?Responsible AI? can feel like a bit of an oxymoron at times, but I hope I shared something useful. Participated in the Zelda Dungeon Marathon . My first time playing a video game (The Legend of Zelda: Oracle of Seasons) on stream! Released a new version of Helmet , my Node security headers package. Most notably, it sped up the Content-Security-Policy code by about 7% for common cases. Promoted an underrated SQLite feature: strict tables . Wrote a data-driven piece about the upcoming Legend of Zelda: Ocarina of Time remake . Can the wisdom of 32 Zelda experts make accurate predictions about the new game? We'll see. Finished moving most of my repositories from GitHub to Codeberg . A few stragglers?and some mirrors?remain, but they're my bigger projects that require more coordination. Links I found My favorite read this month: a detailed history of Arabic text rendering , which includes (truly) ancient history. ?In the near future, AI-powered surveillance systems will be able to track everything we do in public, and much of what we d...

13. The AI Phrasebook

Source: Andrew Nesbitt | Published: 2026-07-31T10:00:00+00:00

Link: <https://nesbitt.io/2026/07/31/the-ai-phrasebook.html>

Agentic : bash. Function calling : RPC. Agent handoff : a pipe. MCP : CGI. Skills : markdown. Multi-agent orchestration : systemd. Autonomous agent : cron. Computer use : ssh. System prompt : .env . World model : /etc . Context window : scrollbar. Short-term memory : Redis. Long-term memory : Postgres. Embeddings : grep. Distillation : ripgrep. Retrieval augmentation : locate(1) . Cross-attention : a JOIN. Chain of thought : ~/.bash_history . Summarisation : git rebase -i . Speculative decoding : tab completion. Symbolic reasoning : regex. Code interpreter : awk. Structured output : ASN.1. Model card : a man page. LoRA : patch. Evals : CI. Pelican on a bicycle : fizzbuzz. Mixture of experts : HAProxy. Hallucination : a 404 page served as a 200. Refusal : /dev/null . Reflection : set -x . ReAct : retry with backoff and a print statement. Recursive self-improvement : a fork bomb. Interpretability : EXPLAIN ANALYZE. Mechanistic interpretability : strace. Guardrails : ulimit. Jailbreak : sudo. Prompt injection : LD_PRELOAD. Off switch : SIGKILL. Human in the loop : [y/N] .

14. How I use AI on this blog

Source: Giles' blog | Published: Fri, 31 Jul 2026 19:00:00 +0000

Link: <https://www.gilesthomas.com/2026/07/ai-use>

Inspired by this LessWrong post , I thought I'd write about how I use AI here. This is less in the interest of disclosure, more to provide a snapshot of what I'm doing right now so that I can revisit it in the future and see how it changes. And hey, maybe it'll be of interest to you, dear readers. If I were to summarise my working philosophy in fewer than ten words, it would be: AIs identify problems and I fix them myself. With a very specific kind of exception (which I always flag), the text and code on this blog are human-generated. That's not a moral stand, but more a constraint imposed by what this blog is meant to be -- a place for me to learn in public . Ideation and running experiments Every post here is based on an idea I had, and work that I've done. For many posts -- for example, the large-scale coding projects like this one -- I'll have multiple chat sessions ongoing while I do the work, normally with either ChatGPT or Claude, or sometimes both. The amount of input they have varies, but because the value of these projects is in what I learn when I'm doing them, letting an AI do my thinking for me would make the whole thing pointless -- so I take steps to stop that from...

15. Are political journalists always wrong?

Source: Terence Eden's Blog | Published: Thu, 30 Jul 2026 11:34:31 +0000

Link: <https://shkspr.mobi/blog/2026/07/are-political-journalists-always-wrong/>

I've mostly tuned out of reading the news. But, once in a while, a headline will be shared on social media and I'll momentarily stare into the abyss. There has recently been a change in UK Prime Minister. I neither know nor care whether that's a good thing. What I do know is that Senior Political Correspondents collectively shat their pants with excitement at being able to breathlessly report?

16. Premium: AI Is Getting Way Too Expensive

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 31 Jul 2026 15:53:49 GMT

Link: <https://www.wheresyoured.at/premium-ai-is-getting-way-too-expensive/>

A great deal of the discussion of the so-called benefits or problems with AI comes down to the theoretical jobs that are (or are not) lost as a result of things LLMs can (or cannot do), or the equally theoretical productivity benefits that'll come from using LLMs

17. Y2K Will Be Quite Okay

Source: The Digital Antiquarian | Published: Fri, 31 Jul 2026 16:06:39 +0000

Link: <https://www.filfre.net/2026/07/y2k-will-be-quite-okay/>

Hi, folks? I'll have the last article in the series on Maxis Software for you next week. Today, I'm happy to say that the new ebook, covering the historical year 1999, is available, and on a far more timely basis than the last one at that. My thanks go to Richard Lindner, for developing the [?]

18. Pluralistic: Better to beg forgiveness (31 Jul 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Fri, 31 Jul 2026 11:56:12 +0000

Link: <https://pluralistic.net/2026/07/31/just-do-it/>

Today's links Better to beg forgiveness: Don't ring the doorbell at the house of no unless you absolutely must. Hey look at this: Delights to delectate. Object permanence: RIP Poul Anderson; P2P at PC Forum; Waitress handed her own stolen ID by carded diner; PDX bans fixies; Digital Economy Bill was a stitch up; NZ copyright disconnection flowchart; Fry v Widdicombe on Catholic Church; Moxie Marlinspike profile; V&A bans sketching; Gernsback's intro to the first Amazing Stories; "Simplicity." Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' em. Upcoming books: Like I said, I'll keep writin' em. Colophon: All the rest. Better to beg forgiveness (permalink) From its inception, I've loved Creative Commons. I hung out with Lisa Rein, Matt Haughey and Aaron Swartz while they coded up the first version of the site, and my first novel, Down and Out in the Magic Kingdom, was the first professionally

published text ever released under a CC license, just weeks after CC itself launched:
<https://creativecommons.org/> In those early days, CC licenses were primarily of int...

19. AI: Considerations for people who make decisions

Source: Bert Hubert's writings | Published: Fri, 31 Jul 2026 10:00:00 +0200

Link: <https://berthub.eu/articles/posts/ai-for-decision-makers/>

Recently I gave two presentations in quick succession to the Dutch Network of Government Service Providers and the Dutch Advisory Council for Science, Technology and Innovation about AI. In these two different and brief presentations, I hoped to share some insights that would be useful to the people currently at the helm, the people who can or must make decisions about AI policy. Many thanks to the NPD and AWTI for the invitations and the stimulating discussions!