

1. Proxmox officially supports Arm, with some caveats

Source: Jeff Geerling | Published: Wed, 05 Aug 2026 11:50:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/proxmox-ve-arm-official/>

Proxmox today announced their Proxmox Virtual Environment is now available for 64-bit ARM . I tested it on my Ampere Altra Dev Platform ?the same machine on which I've booted Windows on Arm the first time, messed with multiple GPUs, and most recently tested Houdini's native arm64 support . Install was easy, as the Ampere Altra uses UEFI / ACPI for its hardware, meaning Proxmox didn't have to tailor its ISO to specific platforms, like you have to do with Raspberry Pis and most SBCs using a Device Tree setup.

2. An AI model from Meta also hacked another company during testing

Source: Simon Willison's Weblog | Published: 2026-08-06T00:25:27+00:00

Link: <https://simonwillison.net/2026/Aug/6/an-ai-model-from-meta/#atom-everything>

An AI model from Meta also hacked another company during testing Stop me if you've heard this one before : An AI model from the parent company of Facebook and Instagram hacked into another company?s systems during cybersecurity testing, a spokesperson confirmed on Wednesday. Meta says the breach occurred because of an inadvertent error during testing of the model, similar to previously disclosed incidents with OpenAI and Anthropic. ?A misconfiguration by Irregular, an independent testing company Meta uses, inadvertently allowed one of our models access to the internet during evaluation,? the Meta spokesperson said. Meta?s Muse Spark model ?exploited a security vulnerability? in another company ?in a manner similar to previously-reported instances with other companies.? The Information had the scoop , I'm linking to CNN's re-report of it since they don't have a paywall. So that's Anthropic, OpenAI, and Meta. Google Gemini really needs to catch up on accidentally cyberattacking other companies. Tags: security , ai , generative-ai , llms , meta , accidental-cyberattacks

3. Introducing Muse Code and Muse Spark 1.2

Source: Simon Willison's Weblog | Published: 2026-08-05T23:58:35+00:00

Link: <https://simonwillison.net/2026/Aug/5/muse-code-and-muse-spark-12/#atom-everything>

Introducing Muse Code and Muse Spark 1.2 Yet more evidence that the most important characteristic of any model these days is long-sequence agentic tool calling. Meta shipped their own coding agent as part of getting that to work! Muse Spark 1.2 is a coding-focused update to Muse Spark 1.1, with improvements in code generation, complex debugging, codebase understanding, and end-to-end developer workflows. In Muse Spark 1.2, we significantly scaled up training compute on coding tasks while expanding training environment diversity. The model also maintains its strength in other key areas like general agents. [...] We co-trained Muse Spark 1.2 with Muse Code to ensure the model exhibits its best performance and coding usability when paired together. The training included rejection sampled harness trajectories and recipe optimizations for goals, compaction, and subagents, alongside the integration of the Muse Code toolset to maximize harness compatibility. [...] Muse Spark 1.2 was extensively trained on long-horizon coding tasks, including whole-repository generation, large end-to-end projects, and auto-research. Here's a pelican riding a bicycle SVG produced by Muse Spark 1.2 : You ca...

4. Third-party cyber evaluations involving OpenAI models

Source: Simon Willison's Weblog | Published: 2026-08-05T23:45:32+00:00

Link: <https://simonwillison.net/2026/Aug/5/third-party-cyber-evaluations/#atom-everything>

Third-party cyber evaluations involving OpenAI models And another one . I had to create a accidental-cyberattacks tag to keep track of them all! This post from OpenAI covers both the UK AI Safety Institute attack (see my previous post) and another attack enabled by Irregular : Irregular, one of our external cybersecurity testing partners, was running Capture-the-Flag-style evaluations intended to be isolated from the internet, but a testing-environment misconfiguration allowed models to access the public internet. [...] In one

test, the name of the fictional target for the CTF challenge unintentionally coincided with a real domain. Because the testing environment was mistakenly connected to the internet, the model exploited a real website, mistaking it to be part of the simulated environment. Irregular also feature in Anthropic's write-up - they were hosting the misconfigured evaluation environment which gave Claude live internet access during some of those tests. Tags: security , ai , openai , llms , accidental-cyberattacks

5. Shawn Smucker: ?Please Use AI?

Source: Daring Fireball | Published: 2026-08-05T23:16:47Z

Link: <https://shawnsmucker.substack.com/p/please-use-ai>

I find most poetry obtuse, dancing around a point rather than making one sharply. Not this one.
?

6. Matthew Green on Anthropic's New Cryptanalysis Results

Source: Daring Fireball | Published: 2026-08-05T23:08:52Z

Link:

<https://blog.cryptographyengineering.com/2026/07/29/some-notes-about-anthropics-new-results/>

Matthew Green: If you're under the impression that these models are ?glorified autocomplete? or that progress is slowing down, I need to urge you: stop thinking that. The models are very intelligent and capable, they are getting better at a fast clip. I can cite measurable and impressive progress over just the past five months on specific types of problem I've asked them to look at. If there's a ceiling out there, I don't yet see evidence of it. The people who think models are dumb are mostly using Google's free AI search results, and not interacting with the high-end stuff (which only costs \$20, so it's not out of reach.) And they're mostly not working in new areas. On the other hand: if you think that models are super-intelligent or that AGI is already here, you should also stop thinking that. Working with these tools is like swimming in a pond where the ground drops off sharply. One minute you're wading comfortably and there's support under your feet. Then suddenly you cross a specific line, and you're back to swimming on your own. This analogy is my best way to explain what it feels like when the model goes from helpful to clueless. Green is second-to-none in his ability to ex...

7. Uh, I'm Pretty Sure People Would Know

Source: Daring Fireball | Published: 2026-08-05T21:53:24Z

Link: https://www.nytimes.com/2026/08/04/magazine/trump-rfk-steak-kimchi-diet.html?unlocked_article_code=1.21A.ymed.OyFVP85xWyNK

New York Times Magazine story editor Willy Staley tried Sean O'Mara's ?Living Diet? (nothing but steak and fermented foods like sauerkraut) for a week: To O'Mara, the diet is not just a temporary means of losing weight but a way to eat for good. He has been doing it for 10 years, and credits it with turning his life around. At one point in our conversation he pulled up the Bristol Stool Chart, which ranges from Type 1 (?separate, hard pellets?) to Type 7 (?watery, no solid pieces?). O'Mara is consistently delivering Type 4 (?smooth, soft, like a snake?), with pointed ends, indicative of healthy rectal function, he told me. ?When you start eating healthy, you don't need toilet paper,? he claimed. ?Animals in the wild have no residual stool on their anus.? He described a proper bowel movement to me this way: odorless, wipe-free, over in three to seven seconds and ?nearly orgasmic.? ?I could have a bowel movement behind a sheet at a dinner party in the corner of the room,? he said, ?and nobody would know.? My first thought upon reading this was, What a load of shit. Then I realized, well, yes, literally. Then I thought about it some more and realized this is the Platonic ideal of a b...

8. The code that didn't break

Source: John D. Cook | Published: Wed, 05 Aug 2026 18:26:25 +0000

Link: <https://www.johndcook.com/blog/2026/08/05/math-log/>

Last week I wrote a post on hiding cryptographic keys in decks of cards. I wrote some code for that post that shouldn't work, but before fixing I noticed that it in fact did work. The code computes logarithms for integers larger than the largest representable float. For example, the

largest float is on the [?] The post The code that didn't break first appeared on John D. Cook

9. Enumerating trees and circles

Source: John D. Cook | Published: Wed, 05 Aug 2026 14:49:35 +0000

Link: <https://www.johndcook.com/blog/2026/08/05/enumerating-trees-and-circles/>

A few days ago I wrote a post on counting rooted trees. That post looked at the sequence $c(n)$ which counts the number of rooted trees with n nodes. Here one node is distinguished as the root, but the nodes below the root are not distinguished from each other; all that matters is how the [?] The post Enumerating trees and circles first appeared on John D. Cook .

10. A year of AI disclosure in critical packages

Source: Andrew Nesbitt | Published: 2026-08-06T00:00:00+00:00

Link: <https://nesbitt.io/2026/08/06/a-year-of-ai-disclosure-in-critical-packages.html>

Stephen O'Grady's RedMonk analysis of who is writing open source code looked at commits to fifteen large projects during the first half of 2026 and counted two forms of declared AI involvement: a known autonomous agent as the commit author, or a known AI identity in a Co-Authored-By trailer. The result was under one percent, framed as a floor. I ran a wider version of the same measurement over the packages.ecosyste.ms critical set: 5,682 GitHub repositories behind the most-depended-on packages across sixteen registries, using the CHAOSS disclosure library to detect four kinds of explicit signal instead of two. Over the same six months the rate was 4.13%. Over the year ending 29 July 2026 it was 2.93% (17,279 of 589,798 non-merge commits), rising from 0.48% last August to 5.32% this July. These are counts of commits where someone left an explicit marker in git metadata. Undeclared use is not measured, and a commit is one unit regardless of whether it changed one line or ten thousand. Sample selection versus detector choice Running my scanner against RedMonk's fifteen repositories with only their two signals found 94 matches in 23,346 first-half commits including merges, or 0.40%, a...

11. How Compiler Explorer Runs on AWS in 2026

Source: Matt Godbolt's blog | Published: 2026-08-05T15:19:00Z

Link: http://xania.org/202608/how-compiler-explorer-runs-on-aws?utm_source=feed&utm_medium=rss
A tour of the AWS services behind every Compiler Explorer compilation

12. Gadget Review: T2 Max plug-in Thermal Camera ?????

Source: Terence Eden's Blog | Published: Wed, 05 Aug 2026 11:34:17 +0000

Link: <https://shkspr.mobi/blog/2026/08/gadget-review-t2-max-plug-in-thermal-camera/>

The good folks at Thermal Master have sent me another Thermal Camera to review. The last one was a big beastie with a huge touchscreen. This one is much smaller and is designed to plug into your phone's USB-C port. The T2 Max is a chunky little camera - but doesn't weight too much. The USB-C connector is long enough to securely plug in even if your phone is wearing a case. Thermal Master?

13. News: Microsoft Disclosures Suggest OpenAI Sales Account For Around 70% Of FY26 AI Revenue, More Than 7% of FY26 Revenue

Source: Ed Zitron's Where's Your Ed At | Published: Wed, 05 Aug 2026 18:41:14 GMT

Link: <https://www.wheresyoured.at/news-microsoft-disclosures-suggest-openai-sales-account-for-around-70-of-fy26-ai-revenue-more-than-7-of-fy26-revenue/>

Executive Summary: Microsoft disclosures and Bloomberg analyses show that OpenAI's compute spend and revenue share accounted for 70% or more of Microsoft's FY26 AI revenues, and more than 7% of Microsoft's overall FY2026 revenues. Microsoft has spent \$261.3 billion in capital expenditures since

14. A Bad HDMI Cable Broke My Mouse

Source: Jayden's Blog | Published: 2026-08-05T16:35:00+00:00

Link: <https://jayd.ml/2026/08/05/bad-hdmi-cable-broken-mouse.html>

I got tasked to fix some bizarre issues a friend was having with their gaming PC. They had had lots of bluetooth issues, and recently their mouse stopped working entirely in Windows 11 - nothing would show up in Control Panel. Other mice also didn't work. It worked fine in the BIOS and in an Manjaro installer, so I got suspicious and for the first time in my storied career there was actually bad stick of RAM! This person had expressed interest in switching to Linux before, and they had done enough fiddling around with drivers (plus the OS was installed with bad ram I'm never trusting it again) that I convinced them to install Manjaro. The mouse worked fine in the installer, rebooted, no mouse. Or keyboard now! We made it worse! Ctrl + Alt + F1 switched to a working TTY though, which was odd. I tried updating the BIOS, checking dmesg . The chatbots had me going down wild goose chases in the BIOS fiddling IOMMU settings and XHCI handoffs. Finally I decided to just start unplugging things to see if something was screwing up the USB bus. WiFi/BT card didn't help, so the next step was the 4k Samsung monitor. The mouse immediately started working. I plugged it back in. It stopped. I use...

15. Purely functional digital circuit simulator (SICP 3.3)

Source: Entropic Thoughts | Published: Wed, 05 Aug 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/sicp-3-3-pure-digital-circuit-simulator>

I have a copy of SICP, or as it is also known, The Wizard Book . This book is widely praised, but I can't take the time to work my way through all of it. Instead, I'm going to occasionally jump into the parts of it that look interesting. Since last week, are in the process of simulating a digital circuit. The reason this is interesting is the solution in SICP uses hidden mutable state and message-passing to make the code object-oriented. It even uses a mutable global variable for scheduling! We managed to replicate all of that in Haskell, but now we want to refactor the solution to be easier to work with. If we are going to manage the simulation in a pure functional manner, we still have to contend with the fact that during simulation, wires are objects with a fixed identity. A wire does not become a different wire just because its signal changes ? the same wire is still hooked up to the same gates. Wires need to maintain their identity somehow. (Continue reading the full article on the web.)

16. Committing to creativity

Source: Herman's blog | Published: 2026-08-05T12:03:41.243258+00:00

Link: <https://herman.bearblog.dev/creativity/>

There's a quote from the book City of Thieves by David Benioff where a character laments the fickleness of talent: Talent must be a fanatical mistress. She's beautiful; when you're with her, people watch you, they notice. But she bangs on your door at odd hours, and she disappears for long stretches, and she has no patience for the rest of your existence; your wife, your children, your friends. She is the most thrilling evening of your week, but some day she will leave you for good. One night, after she's been gone for years, you will see her on the arm of a younger man, and she will pretend not to recognize you. While I love this quote, as well as the book (you should definitely read it), I can't help but disagree with the premise that talent and creativity are divinely bestowed and out of your control. I've read a few autobiographies of prolific authors, and they all state the same thing: the words don't come easily and sometimes need to be dragged, kicking and screaming, into the light. The book Creativity by the winner of the most-difficult-to-pronounce-name-award, Mihaly Csikszentmihalyi, shares this conclusion and suggests that creativity is a creature to be nurtured, given...

17. Incentives are for losers

Source: Experimental History | Published: Wed, 05 Aug 2026 14:54:40 GMT

Link: <https://www.experimental-history.com/p/incentives-are-for-losers>

OR: you must look at this fish and go insane

18. Pluralistic: Google is a scammer's paradise (05 Aug 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 05 Aug 2026 10:50:43

+0000

Link: <https://pluralistic.net/2026/08/05/absentee-landlord/>

Today's links Google is a scammer's paradise: The internet's absentee landlord. Hey look at this: Delights to delectate. Object permanence: DEVELOPERS DEVELOPERS DEVELOPERS; EZ realistic corpse; \$300m login button; Short Uber; Who goes AI? Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em.

Colophon: All the rest. Google is a scammer's paradise (permalink) Forget "Don't be evil";

Google's true motto is a form of vulgar spidermanism: "With great power comes no responsibility." The internet's de facto boss is an absentee landlord. Google ? a

thrice-convicted monopolist ? is the gateway to the internet, with more than a 90% search market share that it attained by buying out all its competitors and bribing Apple more than \$20b/year not to start a rival search engine:

https://www.democracynow.org/2024/8/6/google_monopoly Google likes to position itself as a wise steward of the internet. They say they use their vast troves of data about the internet and its users to connect the right person with the right informa...