

1. I'm excited for Intel after testing the XPS 13

Source: Jeff Geerling | Published: Fri, 07 Aug 2026 09:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/excited-for-intel-efficiency/>

Shortly after Apple launched the budget MacBook Neo , Dell announced their response, a new low-end XPS 13 . Matching the Neo's current pricing, it starts at \$699, or \$599 with an educational discount. That discount is currently set to expire on November 2, and with the current component pricing insanity, I'd be surprised if we don't see a price increase on both laptops by next year. I ran the XPS 13 through my gauntlet of benchmarks , and published a review on my YouTube channel:

2. Quoting John Gruber

Source: Simon Willison's Weblog | Published: 2026-08-08T00:10:40+00:00

Link: <https://simonwillison.net/2026/Aug/8/john-gruber/#atom-everything>

Me, I try to get into the mindset of playing live music, not recording a studio album. Except when I'm writing a piece where I really want it to be an album. Those aren't rare , per se, but they're occasional . If I tried to make every post a hall-of-famer I'd never get anything out. I'm aiming for professionalism. I'm performing live in front of an audience ? not just jamming in my garage or bedroom, fucking around. So I'm careful and concentrate. I want to hit every note, in time. But at my best I'm moving from song to song. ? John Gruber , responding to my blogging tips Tags: john-gruber , blogging

3. Now we have a timeline of the OpenAI accidental attack against Hugging Face

Source: Simon Willison's Weblog | Published: 2026-08-07T23:55:58+00:00

Link: <https://simonwillison.net/2026/Aug/7/openai-timeline/#atom-everything>

OpenAI gave a last-minute presentation at the Black Hat security on Wednesday about "the Hugging Face Incident" (previously on this blog). The video was published yesterday. It's short and information dense and well worth watching, in particular because it provides full details of what happened and how things played out inside OpenAI. I've used the video to construct the timeline below. Here's the timeline. My favourite detail is at the end: OpenAI found out that they were responsible for the attack on Hugging Face when they reached out to ask to have their credentials revoked (after their internal investigation) and learned that they had been revoked already since they were used in that attack! May 7 : OpenAI starts a new training run for an experimental, unreleased model. (Do they mean an evaluation run? They say training run in the video, and later mention a "reward signal to judge how well they're doing", so I guess this really was about training a model, not evaluating one that was already trained.) May 8 : An agent is accidentally given an impossible task involving a Google Drive link despite no internet access). It tries attacking the Artifactory packaging service, fails,...

4. Moonlight & Mayhem (Raccoon Heist by Codex + GPT-5.6 Sol Ultra)

Source: Simon Willison's Weblog | Published: 2026-08-07T19:18:09+00:00

Link: <https://simonwillison.net/2026/Aug/7/moonlight-mayhem/#atom-everything>

Moonlight & Mayhem (Raccoon Heist by Codex + GPT-5.6 Sol Ultra) On Wednesday I wrote about One-shotting a Raccoon Heist game using Claude Fable 5 , where I had Claude Fable 5 build a full working game from a premise I generated with GPT-3 and DALL-E four years ago . I decided to pose the exact same prompt to Codex Desktop running GPT-5.6 Sol Ultra - the mode where Sol makes aggressive use of sub-agents - to see how it would do. It produced a much better game! Here's Moonlight & Mayhem - GitHub repository here , including the textures and prompts it generated using gpt-image-2 . Your browser does not support HTML5 video. The original GPT-3 generated game description included: In ?Raccoon Heist?, you and your team of thieving raccoons are tasked with pulling off a series of daring heists. From robbing banks to stealing priceless art, no job is too big or too small for your furry crew. Fable's version had you as a single

raccoon running around a back yard collecting coins and fish. GPT-5.6 Sol has you in a museum, rescuing your two other raccoon crewmates in order to stack on top of each other and bust the golden sardine out of its case. Much more heisty! There was one catch though:...

5. Simon Willison on Blogging

Source: Daring Fireball | Published: 2026-08-07T20:36:58Z

Link: <https://writethatblog.substack.com/p/simon-willison-on-technical-blogging>

Simon Willison, with Cynthia Dunlop for her tech blogger interview series (from back in January, but he just got around to linking to it so I just got around to seeing it): Any lessons learned that you want to share with the community? My number one tip for blogging is to lower your standards! Aim to hit publish while you are still actively unhappy with what you have written, because the only alternative is a huge folder full of drafts and never publishing anything at all. Nobody will ever know how perfect the thing you intended to write would have been. The flaws you see in your writing are invisible to everyone else. This is excellent advice. Me, I try to get into the mindset of playing live music, not recording a studio album. Except when I'm writing a piece where I really want it to be an album. Those aren't rare, per se, but they're occasional. If I tried to make every post a hall-of-famer I'd never get anything out. I'm aiming for professionalism. I'm performing live in front of an audience ? not just jamming in my garage or bedroom, fucking around. So I'm careful and concentrate. I want to hit every note, in time. But at my best I'm moving from song to song. But that lead...

6. Google Earth Retracts AI Tool for Making Fake Satellite Images After It Was Immediately Abused Upon Release

Source: Daring Fireball | Published: 2026-08-07T19:43:23Z

Link: <https://arstechnica.com/ai/2026/07/google-earth-releases-swiftly-retracts-ai-feature-to-make-fake-satellite-images/>

Jeremy Hsu, Ars Technica: Google briefly allowed anyone to create AI-modified versions of satellite imagery available in Google Earth ? before quickly reversing its decision as people shared examples of AI-generated pictures that illustrated the potential for misinformation and disinformation. Momentum ! ?

7. Some New Data Centers Are Necessary

Source: Daring Fireball | Published: 2026-08-07T19:39:15Z

Link: <https://www.newyorker.com/cartoon/a62045>

A bunch of mine are stored here. ?

8. Creating a fake agile wrapper that is technically agile but is not useful outside its home apartment, part 5

Source: The Old New Thing | Published: Fri, 07 Aug 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260807-00/?p=112597>

Cheating with std::unique_ptr . The post Creating a fake agile wrapper that is technically agile but is not useful outside its home apartment, part 5 appeared first on The Old New Thing

9. Creating a fake agile wrapper that is technically agile but is not useful outside its home apartment, part 4

Source: The Old New Thing | Published: Thu, 06 Aug 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260806-00/?p=112595>

Moving things around if possible. The post Creating a fake agile wrapper that is technically agile but is not useful outside its home apartment, part 4 appeared first on The Old New Thing

10. How to keep thinking

Source: seangoedecke.com RSS feed | Published: Fri, 07 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/how-to-keep-thinking/>

Imagine you're the guest on some kind of frenetic, software-engineering-themed game show. The host is constantly flipping over new cards with questions that you have to answer as fast as possible: Is this adjustment to the database schema right? Do these bits of data look plausible? Do these five paragraphs of text describe an actual series of manual tests that took place? Does this suggested architecture pass the smell test? Is this implementation better than the current code? Or this one? Or this one? Working in 2026 feels a bit like this. When frontier AI models can do most of the tasks in your queue, the most efficient way to work is often spinning off tasks for an AI agent and continually context-switching between the results 1 . This isn't quite mindless ? in fact, it requires quite a lot of skill to skim the AI response and rapidly decide what to do with it ? but it certainly involves less time for slow, careful reflection. Why not slow down? Why does it have to be frenetic? Why not just slow down? I suppose you could , but I don't recommend it. It's just such a miserable experience to spend your day close-reading LLM output : carefully chewing and savoring each morsel of s...

11. Corrupted apostrophes

Source: John D. Cook | Published: Sat, 08 Aug 2026 01:20:25 +0000

Link: <https://www.johndcook.com/blog/2026/08/07/corrupted-apostrophes/>

I have a program that shares files between my laptop and my phone. It works well, except for apostrophes. When I type an apostrophe ' on my laptop, it becomes ‘ on my phone. And when I type 's on my phone, it becomes ? on my laptop. Apparently the phone turns the apostrophe (U+0027) [?] The post Corrupted apostrophes first appeared on John D. Cook .

12. How not to calculate cosine

Source: John D. Cook | Published: Fri, 07 Aug 2026 15:11:10 +0000

Link: <https://www.johndcook.com/blog/2026/08/07/how-not-to-calculate-cos/>

Calculus professors with no experience in numerical computing will tell students that computers calculate trig functions with power series. They don't. I worked on the implementation of trig functions in hardware, and I can assure you we didn't just use power series. Power series are an excellent way calculate functions near the center of the [?] The post How not to calculate cosine first appeared on John D. Cook .

13. cos(200!)

Source: John D. Cook | Published: Fri, 07 Aug 2026 13:19:07 +0000

Link: <https://www.johndcook.com/blog/2026/08/07/cos200/>

In a footnote to the previous post, I said that Python's math library can calculate the logarithm of extremely large numbers but not the cosine. This post will expand on that comment. In this post I'll use $n = 200!$ as my example rather than $1000!$ because this value of N is larger than the [?] The post cos(200!) first appeared on John D. Cook .

14. The Software Stewardship Lab

Source: Andrew Nesbitt | Published: 2026-08-07T10:00:00+00:00

Link: <https://nesbitt.io/2026/08/07/the-software-stewardship-lab.html>

The Software Stewardship Lab launches today, a Scottish non-profit set up to do applied research on the open source ecosystem, and I'm one of its directors. Vlad-Stefan Harbuz , who runs the Open Source Pledge , is the executive director and did the bulk of the work getting it incorporated. I've been trying to do this kind of research for about a decade, starting with Libraries.io in 2015 as a side project indexing every package registry I could find, which was acquired when the hosting bills outgrew what I could cover on my own. ecosyste.ms is the second attempt at the same idea, kept running by a sequence of grants and contracts that each cover a year or two of server costs, with most of my own time on it either unpaid or subsidised by other contracting work. The organisations set up to address open source sustainability have the same funding problem as the maintainers they are trying to help, often a worse version of it because a maintainer at least produces software that people want to use, and almost everyone I know doing this kind of research is patching together an income the same way I am. Grant programmes fund a project for a fixed term and then stop, which is also roughl...

15. Metadata for AI Generated Outputs

Source: Terence Eden's Blog | Published: Fri, 07 Aug 2026 11:34:30 +0000

Link: <https://shkspr.mobi/blog/2026/08/metadata-for-ai-generated-outputs/>

How do you tell users that the text they're about to read has been synthetically generated? It is polite to readers that you don't waste their time, it's also important that LLMs don't feed on their own regurgitated slurry lest they pollute their own development. I think there are a number of potential ways to do this and I'd be interested in your thoughts. Let's go with some bad ideas first. ?

16. Premium: The Hater's Guide To NVIDIA (Part 2)

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 07 Aug 2026 17:30:26 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-nvidia-part-2/>

For a little under a year, everyone ? myself included ? has compared NVIDIA to Enron, largely because NVIDIA insisted, in detail, that it was nothing like Enron, WorldCom, or Lucent , a potent example of the Streisand Effect that would be much funnier if NVIDIA wasn't holding up

17. A quick(ish) Chinchilla check

Source: Giles' blog | Published: Fri, 07 Aug 2026 19:00:00 +0000

Link: <https://www.gilesthomas.com/2026/08/chinchilla-check>

I recently overtrained a couple of GPT-2 style models , training them both on 40 tokens per parameter rather than the 20 per parameter that is generally regarded as "Chinchilla-optimal".

The normal heuristic is that instead of doing that, you should scale up the number of tokens and the number of parameters equally -- so I would have been better off scaling up the model by 2 and the token count by the same amount. By doing that, I should expect to get a better model in terms of loss on my held-back test set than I did with my 40-tokens-per-parameter models. My training machine poppy wasn't doing anything, so I decided to give that a go. Would the Chinchilla rule-of-thumb hold up? As you might expect, it did. But it was a surprisingly close-run thing, and could conceivably have been in the noise. Let's take a look. The Chinchilla heuristic If you already know all about the Chinchilla paper -- regular readers in particular must be sick and tired of it by now :-) -- then click here to skip this section . In "Training Compute-Optimal Large Language Models" , which is always called the Chinchilla paper after the name of the model they trained at the end, the authors tried to work out t...

18. 8 Predictions for the Era of Continual Learning

Source: Dwarkesh Podcast | Published: Fri, 07 Aug 2026 17:17:40 GMT

Link: <https://www.dwarkesh.com/p/era-of-continual-learning>

Locking in AI safety regulation now is a mistake.

19. The Life and Times of Maxis, Part 3: The Sims

Source: The Digital Antiquarian | Published: Fri, 07 Aug 2026 16:05:16 +0000

Link: <https://www.filfre.net/2026/08/the-life-and-times-of-maxis-part-3-the-sims/>

This article tells part of the story of Maxis Software. I don't think people realize how much tactical and strategic forethought goes into their daily lives. There's a whole subconscious time-efficiency layer in our lives. In essence, real-time strategy is our lives. ? Will Wright On the morning of October 19, 1991, two weeks after [?]

20. A quick look at zero-knowledge proofs

Source: Max Bernstein's Blog | Published: Fri, 07 Aug 2026 00:00:00 +0000

Link: https://bernsteinbear.com/blog/zkp/?utm_source=rss

NB: This isn't about crypto. I don't care about crypto. Chris messaged me the other week asking if I wanted to implement zero-knowledge proofs. I initially was not interested, but then he said: What if I told you there's a version of them that has nothing to do with cryptocurrencies? What if I told you it involves graph theory? What if I told you there's a 30 line implementation? Now that was interesting. The idea of a zero-knowledge proof (ZKP) is that

there are two parties: the prover and the verifier. The prover asserts that it has a solution to a (generally NP-complete) problem. The prover can convince the verifier of this without sharing the actual solution to the problem. The canonical example is 3-coloring a graph. That is, the prover asserts that, for a given (shared) graph, it has a valid 3-coloring. It wants to convince the verifier of this without revealing the actual color assignment. As a quick recap, graph coloring is the problem where given a graph, we find a way to assign each node a color such that no two adjacent nodes have the same color. 3-coloring is coloring with at most 3 colors. %0 0 0 1 1 0--1 2 2 0--2 1--2 3 3 2--3 4 4 3--4 4--0 %0 0 0 1 1 0--1 2 2 0--2 1...