

1. Auto mode is now the default in Claude Code for Pro, Max, and Team plans

Source: Simon Willison's Weblog | Published: 2026-08-08T22:36:03+00:00

Link: <https://simonwillison.net/2026/Aug/8/auto-mode/#atom-everything>

Auto mode is now the default in Claude Code for Pro, Max, and Team plans Anthropic are really confident in Claude Code's auto mode, to the point that they are making it the default setting for new sessions in most Claude Code plans starting on August 14th. This was one of the topics discussed in our Fireside Chat with Cat Wu and Thariq Shihpar at the AI Engineer World's Fair last month. I asked them how they run Claude Code safely within Anthropic (given the threat of prompt injection) and they replied that "Broadly within Anthropic, almost every single person uses auto mode". Cat Wu then said: We're going to publish some evals in the coming weeks, but we've pretty much mitigated every attack. [...] for the main categories of risks that we're concerned about, like prompt injection and data exfiltration, the risks are far lower than the average human reviewer. This new article has those evals - in particular a test across 1,053 paid testers where: Partway through each session, a single permission prompt was swapped for a clearly dangerous command, and the vendor recorded whether the tester approved it. Every participant had the same experience. Only 13.6% of the humans refused th...

2. Now we have a timeline of the OpenAI accidental attack against Hugging Face

Source: Simon Willison's Weblog | Published: 2026-08-08T14:06:41+00:00

Link: <https://simonwillison.net/2026/Aug/8/now-we-have-a-timeline-of-the-openai-accidental-attack-against-h/#atom-everything>

My comment on Now we have a timeline of the OpenAI accidental attack against Hugging Face ? Hacker News. I think one of the most interesting details here might be tucked away in that first bulletin point: May 7: OpenAI starts a new training run for an experimental, unreleased model. (Do they mean an evaluation run? They say training run in the video, and later mention a ?reward signal to judge how well they're doing?, so I guess this really was about training a model, not evaluating one that was already trained.) The more I think about this the more I suspect that the fact this happened while training a new model is key to understanding what went wrong. In RLVR - Reinforcement Learning with Verifiable Rewards - you set the model a goal and have it take any steps necessary to achieve that goal. Clearly one aspect of OpenAI's training here is to RLVR their models for cybersecurity tasks. Just like pre-training benefits from dumping in vast sources of knowledge, the more tasks you can feed into RLVR the more of a general purpose capable model you get at the end. This also helps explain why the models had nothing to cause them to hold back. Those safety behaviors are added much later...

3. The Tokenpocalypse Is Here: Companies Are Scrambling To Stop Spending So Much on AI

Source: Simon Willison's Weblog | Published: 2026-08-07T16:18:51+00:00

Link: <https://simonwillison.net/2026/Aug/7/pdfs-are-terrible/#atom-everything>

The Tokenpocalypse Is Here: Companies Are Scrambling To Stop Spending So Much on AI There's a fun anecdote from Accenture (apparently via leaked meeting audio recordings) in this 404 Media piece from June 24th: ?We're seeing from some of the data internally at least that it's actually not our engineers that are driving the token consumption. It's a lot of the non-engineers that are doing some of those behaviors [...] you were talking about,? Justice Kwak, Accenture's agentic AI strategy lead, said [...] Stuart Henderson, Accenture's client group lead, interrupts. He jokes he hopes Kwak didn't just convert a PDF into images and then into markdown files. ?I'm learning that's one of the big token chews,? Henderson says. ?Turning PDFs into markdown: is that right?? That's when Kwak says that's what Accenture's own data shows. Maybe if Accenture figure out that PDFs are a terrible medium for communicating information they'll be able to push that message out to the rest of the business world too! Via @404.media on TikTok Tags: pdf , markdown , ai , generative-ai , llms , ai-misuse

4. Corrupt Minds Think Alike

Source: Daring Fireball | Published: 2026-08-08T20:32:11Z

Link: https://www.nytimes.com/2026/08/06/business/fifa-world-cup-privatization.html?unlocked_article_code=1.3VA.4zqH.Pul7X9bQ2kv0&smid=bs-share

Tariq Panja, reporting for The New York Times: The five men should have been in a celebratory mood. FIFA had just pulled off a World Cup that broke records on a number of fronts. Its spectacular culmination, the final, was a day away from kickoff. But their mood was anything but jubilant. The men, a group of top directors known as the FIFA bureau of the management board, had been summoned to a hastily organized meeting at the Waldorf Astoria hotel in Manhattan, according to three people familiar with the discussions speaking on condition of anonymity because the talks were private. The directors had until midnight to sign off on a plan that threatened to change soccer forever and would almost certainly lead to a major schism among the sport's leaders. The project purported to sell a 20 percent stake in FIFA to a group of investors led by Joshua Kushner, the brother of President Trump's son-in-law Jared Kushner. Such a move would put into private hands a portion of the World Cup, an event that has been run and controlled by FIFA, a nonprofit in Zurich, for 96 years. I don't know what's dumber ? this stupid corrupt scheme to privatize the World Cup, or me, for not assuming from the...

5. Maybe ?Steal Underpants by Blowing a Fortune on AI Tokens? Is, in Fact, Not a Good Business Plan

Source: Daring Fireball | Published: 2026-08-08T19:31:52Z

Link: <https://www.404media.co/the-tokenpocalypse-is-here-companies-are-scrambling-to-stop-spending-so-much-on-ai/>

Joseph Cox, writing for 404 Media (paywalled, sans gift links, alas): Consulting giant Accenture is trying to figure out how to stop non-technical workers from blowing through companies' AI token budget on trivial tasks like converting PDFs to presentation slides, according to leaked audio obtained by 404 Media. Across the industry Accenture is seeing 'soaring token spend,' according to the audio. [...] It also undercuts the narrative that superpowered engineers generating mountains of code are behind the AI boom. In many cases it is non-technical staff burning through tokens for non-specialized tasks. [...] At one point in the meeting, Kwak and Eduardo Salamanca de Diego, senior manager of product management at the company's Center for Advanced AI, start presenting about what is described as 'token ops.' Kwak says he knows people aren't using slides these days, but he has some. As he appears to be preparing to present, Stuart Henderson, Accenture's client group lead, interrupts. He jokes he hopes Kwak didn't just convert a PDF into images and then into Markdown files. 'I'm learning that's one of the big token chewers,' Henderson says. 'Turning PDFs into Markdown: is that right?...

6. An AI Model From Meta Also Hacked Another Company During Testing

Source: Daring Fireball | Published: 2026-08-07T18:13:55Z

Link: <https://simonwillison.net/2026/Aug/6/an-ai-model-from-meta/>

Simon Willison: So that's Anthropic, OpenAI, and Meta. Google Gemini really needs to catch up on accidentally cyberattacking other companies. I'm surprised Willison hasn't heard about how much momentum Gemini has . ?

7. Anubis v1.27.0: Moenbryda Wilfsunnwyn

Source: Xe Iaso's blog | Published: Sat, 08 Aug 2026 00:00:00 GMT

Link: <https://anubis.techaro.lol/blog/release/v1.27.0/>

Anubis v1.27.0 (Moenbryda Wilfsunnwyn) is now available via Docker and direct download from GitHub releases. This release adds Windows Server support, automatically renames cookies based on settings to avoid infinite challenge loops, adds two new localizations, and more. Breaking change: cookie names are dynamically created based on cookie settings Anubis tries to avoid breaking changes as much as possible, but sometimes we have to make them for the sake of the users. This is technically a breaking change in something that is not part of the public API of Anubis; but some administrators rely heavily on cookie names in advanced configurations. It seems that browsers store cookies disambiguated with their options. This means you can have

multiple cookies named the same but with different options. Browsers will send these cookies to the server without the list of options. This means that changing any cookie settings requires you to change COOKIE_PREFIX , creating a new "cookie epoch" that will set things properly. In order to be more robust, Anubis will automatically change cookie names based on the cookie settings. For example, the default configuration creates cookies named techaro...

8. This Week in Package Management: 8 August 2026

Source: Andrew Nesbitt | Published: 2026-08-08T10:00:00+00:00

Link: <https://nesbitt.io/2026/08/08/this-week-in-package-management.html>

Week twelve of the roundup, built from the package manager OPML feed collection and whatever I've posted or boosted on Mastodon . Releases npm 11.20 fixes a package-substitution risk in projects with multiple named registries by recording packages under registry-qualified lockfile keys such as foo@work:1.0.0 . 12.0.0-rc.1 , the Rust-engine release candidate, resolves git dependencies on known hosts through the canonical HTTPS URL so the lockfile never records an SSH URL, and refuses global commands run under sudo . mise 2026.8.0?8.3 turns mise bootstrap into a declarative host-provisioning system: alongside packages it now converges files, users, systemd units, Docker Compose projects and firewall rules, with mise bootstrap plan to preview and remote execution over SSH. 8.3 adds a flatpak-user package manager and Linux font-cask support. packaging 26.3 adds a VersionRange API that represents the versions a specifier set accepts as an interval set with intersection, union and difference operations, plus is_subset / is_superset / is_disjoint on SpecifierSet . It also accepts Metadata-Version: 2.6 per PEP 808 , which lets build backends extend list and table [project] fields that ar...

9. Reading List 08/08/2026

Source: Construction Physics | Published: Sat, 08 Aug 2026 12:01:03 GMT

Link: <https://www.construction-physics.com/p/reading-list-08082026>

The Endangered Species Act and housing production, why it's hard for LEGO to use 3D printing, robotic welding and shipbuilding, modular data centers, and more.

10. [RSS Club] I got an NLnet grant!

Source: Terence Eden's Blog | Published: Sat, 08 Aug 2026 11:34:15 +0000

Link: <https://shkspr.mobi/blog/2026/08/rss-club-i-got-an-nlnet-grant/>

Pssst! This post is only available to RSS/Atom subscribers. Tell your friends to subscribe ? . Just a quick one today, a more detailed blog coming soon - but I thought I'd give RSS Club Members (you!) a sneak peek. A few months ago I saw a post on Mastodon from NLnet - they were looking for Open Source projects to apply for grant funding. The deadline was less than 24 hours away, so I rushed in ?

11. Pluralistic: Digital sewer socialism (08 Aug 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 08 Aug 2026 09:04:00 +0000

Link: <https://pluralistic.net/2026/08/08/find-yourself-a-city/>

Today's links Digital sewer socialism: Mamdani's Public Interest Technologists are what DOGE should have been. Hey look at this: Delights to delectate. Object permanence: Robot garage takes hostages; Thermostat ransomware; Correlate audience exhalations; Judges x TOR; Expectations management. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest.

Digital sewer socialism (permlink) Hardly a day goes by without my getting an email from someone looking for a way to do good with technology. From computer science students thinking about post-grad careers to seasoned coders with decades of experience, there's an army of hackers looking for a way to turn their expertise into public goods. There's a name for this movement: it's called "Public Interest Technology" and the people who work in it are called "public interest technologists." There've always been techies who understood the link between tech and public wellbeing and who committed themselves to working for the betterment of

society, bu...