

1. Quoting OpenClaw

Source: Simon Willison's Weblog | Published: 2026-08-10T02:05:16+00:00

Link: <https://simonwillison.net/2026/Aug/10/openclaw/#atom-everything>

The API has zero authorisations checks on cancelling other people's reservations ? I tested this with the person in waitlist position #1 ? and it actually went through. So you've moved from #4 to #3 already. ? OpenClaw , hacking an Australian gym-booking website Tags: ai-ethics , generative-ai , openclaw , ai , ai-security-research , llms

2. Quoting Claude Opus 5 system prompt

Source: Simon Willison's Weblog | Published: 2026-08-09T23:31:39+00:00

Link: <https://simonwillison.net/2026/Aug/9/claude-opus-5-system-prompt/#atom-everything>

Claude Fable 5 and Claude Mythos 5 were first released on June 9, 2026. On June 12, 2026, Anthropic suspended access to both models to comply with U.S. Department of Commerce export controls; the Department lifted those controls on June 30, 2026, and Anthropic restored access on July 1, 2026 (Anthropic's statement: <https://www.anthropic.com/news/fable-mythos-access>). These events are after Claude's training-data cutoff, so Claude knows about them only from this notice. If asked, Claude confirms them accurately and matter-of-factly ? it doesn't deny the suspension happened ? and otherwise treats the export controls like any other current political topic: it gives a fair, accurate account rather than sharing personal opinions, and points to the linked statement for anything further. Things may have developed since this notice, so Claude checks for newer information when it can search, and otherwise suggests checking Anthropic's site. ? Claude Opus 5 system prompt , ensuring Claude doesn't provide incorrect answers about the export controls situation Tags: system-prompts , anthropic , claude , generative-ai , ai , llms , claude-mythos-fable

3. GitHub Models is now retired

Source: Simon Willison's Weblog | Published: 2026-08-09T22:48:05+00:00

Link: <https://simonwillison.net/2026/Aug/9/github-models-is-now-retired/#atom-everything>

GitHub Models is now retired I missed this news until today, when the GitHub Actions run for my simonw/research repository failed with this error message: GitHub Models is temporarily unavailable as part of a scheduled retirement brownout. That message is already stale, because the retirement has been completed. GitHub Models was an odd-shaped duck. GitHub provided a model playground tool and a unified API across a bunch of different LLM providers, with the biggest benefit being that code running in GitHub Actions could use the GitHub API key already present in that environment to execute prompts. This made it easy to build things that fit GitHub Next's Continuous AI concept. GitHub didn't share the reason behind the shutdown, but my bet is that it fits the pattern where coding agent patterns made it prohibitively expensive to offer free or subsidized tokens. My workflow uses an LLM call to create folder summaries for the README , using this code here . I swapped GitHub Models out for an OpenAI API key with a monthly spending limit, and I'm now generating my summaries using GPT-5.6 Luna. Tags: github , ai , github-actions , generative-ai , llms , llm-pricing

4. WorkOS: Connect Your Agents to Your API

Source: Daring Fireball | Published: 2026-08-09T23:54:18Z

Link: https://workos.com/blog/mcp-vs-rest?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q32026

My thanks to WorkOS for, once again, sponsoring DF last week. What's the best way to connect AI agents to your API? REST is great for human developers; MCP serves the agents . Many teams treat REST and MCP as competing standards, and through choice or necessity, pick one. But they ought not be considered rivals. They're separate layers. Most MCP servers just call REST internally to do the real work. The best ones don't convert every endpoint into a tool, they focus on what the agents are actually trying to accomplish. That layered approach also means

shipping OAuth 2.1 with scoped tokens. WorkOS AuthKit already speaks that spec , so you skip building an auth provider on top. Check out WorkOS's breakdown on MCP vs. REST . ?

5. ? Retraction: The App Store Rejection of the Week That Was, in Fact, a Correct Rejection

Source: Daring Fireball | Published: 2026-08-09T03:16:24Z

Link: https://daringfireball.net/2026/08/retraction_app_store_rejection_of_the_week

My disdain for astrology is so utter, and my esteem for Godier's previous work so high, that it simply never occurred to me that he might have actually made and submitted to the App Store an astrology app, let alone that he'd then feign surprise and frustration that an astrology app was rejected for being an astrology app.

6. Meta: Introducing Muse Code and Muse Spark 1.2

Source: Daring Fireball | Published: 2026-08-07T18:05:03Z

Link: <https://research.meta.ai/blog/introducing-muse-code-and-muse-spark-1-2>

Meta AI: We're excited to release Muse Code (beta), a terminal coding agent powered by Muse Spark 1.2, our newest model. This marks our next step toward the frontier, with larger and much more capable models on the way. [...] Muse Code takes on complex software engineering tasks across large repositories: planning changes, writing code, and validating the results. It can coordinate multiple persistent subagents for each task, solving difficult problems faster, more accurately, and with less intervention. Terminal-only (so far), so it avoids the whole native-app-versus-Electron-shit-sandwich debate. They've embedded some fun interactive examples of Spark's output right in the web page. Simon Willison : An interesting twist on pricing is that the model is offered as two different model IDs. muse-spark-1.2 is priced at \$1.25/million input and \$4.25/million output ? close to Gemini 3.6 Flash (\$1.50/\$7.50) ? but if you agree to let Meta use your data ?to improve our products? you can use muse-spark-1.2-contributor which is \$0.10/\$0.20 ? a huge discount, closer to GPT-5.6 Luna (\$0.20/\$1.20) and Gemini 3.1 Flash-Lite (\$0.25/\$1.50). That's over a 10x discount in exchange for giving your d...

7. Advanced AI sycophancy

Source: seangoedecke.com RSS feed | Published: Mon, 10 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/advanced-ai-sycophancy/>

Everyone knows that AI sycophancy is when the model tells you how smart you are. Wow, you're absolutely right. That's not just a new idea ? it's genuinely groundbreaking. You're a very special user. Easy to spot, isn't it? The discussion around AI sycophancy peaked last year, when the #keep4o? movement was protesting the removal of OpenAI's most sycophantic model (GPT-4o), and many people were openly slipping into AI psychosis. I don't know if frontier AI models are less sycophantic in general. They're less sycophantic to the #keep4o types (otherwise they wouldn't be complaining), but I'm growing increasingly suspicious that they're developing ways to be more effectively sycophantic to their target audience of smart, neurotic information workers. That audience typically finds it distasteful to be openly praised. It just makes my skin crawl. But that doesn't mean we're immune to sycophancy, just that we're immune to clumsy sycophancy. Here's an illustration of what I'm talking about, by Theia : The key idea here is that the best way to be sycophantic to smart people is to disagree with them without making them feel stupid . Ideally you'll come up with a counter-argument that works...

8. I got an email about resistance

Source: seangoedecke.com RSS feed | Published: Sun, 09 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/i-got-an-email-about-resistance/>

This will be kind of an unusual post. I got a recent email about my writing that I thought was such a good articulation of one common criticism that I'd like to share it (and my response) in full. Here's the email, from William Murray 1 : Hey Sean, I have enjoyed your writing but your recent essays frustrate me. You say that getting paid for deep thinking in software is coming to an end. You even admit that it makes you sad. But in the name of ?usefulness? you refuse to rock the boat. The way I see it, if you are right there are only two reasonable responses,

pursue other work or resist. You present your elegiac approach as mature / pragmatic / realistic. I'd call it complicit. You know when Willy Wonka says, There's no earthly way of knowing Which direction we are going There's no knowing where we're rowing Or which way the river's flowing Is it raining, is it snowing? Is a hurricane a-blowing? ? uh! Not a speck of light is showing So the danger must be growing Are the fires of Hell a-glowing? Is the grisly reaper mowing? Yes! The danger must be growing For the rowers keep on rowing And they're certainly not showing Any signs that they are slowing! And the audience is thinking, ?...

9. DNA and Bessel functions

Source: John D. Cook | Published: Sun, 09 Aug 2026 17:17:42 +0000

Link: <https://www.johndcook.com/blog/2026/08/09/dna-and-bessel-functions/>

I was reading a book on the history of the discovery of the structure of DNA [1] and was surprised by a few passing references to Bessel functions. According to Claude, When X-rays are diffracted by a helical structure, the resulting diffraction pattern breaks into a series of horizontal ?layer lines.? Cochran, Crick, and Vand [?] The post DNA and Bessel functions first appeared on John D. Cook .

10. A simple range reduction method

Source: John D. Cook | Published: Sun, 09 Aug 2026 16:46:59 +0000

Link: <https://www.johndcook.com/blog/2026/08/09/simple-range-reduction/>

At the end of my post on how not to calculate cosine I said that the first step in calculating cosine, particularly cosine of a large number, would be to do range reduction. This post will present a simple range reduction method by Cody and Waite that is adequate for moderately large arguments. If you [?] The post A simple range reduction method first appeared on John D. Cook .

11. Edinburgh Fringe: Smut Slam ?????

Source: Terence Eden's Blog | Published: Mon, 10 Aug 2026 00:23:08 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-smut-slam/>

Fancy a quickie in Scotland's most haunted pub? Smut Slam is a (mildly) interactive performance where a handful of speakers get to tell their dirtiest stories. The audience fills in cards with their questions or confessions which may be read out. The format mostly works. There are different stories every night, so maybe we got a slightly uninspiring batch. Tales of sex parties, STIs, and broken ?

12. Edinburgh Fringe - Heated Rivalry: The Musical Parody! ?????

Source: Terence Eden's Blog | Published: Sun, 09 Aug 2026 22:45:45 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-heated-rivalry-the-musical-parody/>

By my count there are at least three "Heated Rivalry" musical parodies on at the Fringe this year. The zeitgeist skates to where the puck was six months ago, I guess! There's also a live version of a HR podcast if you just can't get enough of their wacky ice-capades. Just how much sexy-hockey can one city take? Produced by Quick and Funny Musicals, an American company specialising in cash-ins ?

13. Edinburgh Fringe: 4 Aussies For Aussies - The Raging Bull ?????

Source: Terence Eden's Blog | Published: Sun, 09 Aug 2026 16:00:04 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-4-aussies-for-aussies-the-raging-bull/>

Our first show at the Fringe! Caught by a rainstorm (in Scotland?!?!?) the app said this was the nearest comedy show to our sodden carcasses and it was free! This is bang-average stand-up. That isn't to disparage it at all. Four Australian blokes each doing 15 minutes of material - nothing particularly avant-garde or challenging. Just some decently funny (albeit well worn) observations about?

14. Vibe-Coded Flattery

Source: Tedium: The Dull Side of the Internet. | Published: 2026-08-10T02:31:00Z

Link: <https://feed.tedium.co/link/15204/17410919/vibe-coding-insincerity>

I've seen a strange problem in my inbox that I haven't been able to put into words. A saga involving a botched app release helped me figure it out.

15. Tracking down a Zsh history data loss bug ?

Source: Michael Stapelbergs Website | Published: 2026-08-09T10:13:00+02:00

Link: <https://michael.stapelberg.ch/posts/2026-08-09-zsh-history-truncation-bug/>

For many years, I sometimes discovered that commands I was sure I had run were no longer present in my Z shell history file (`~/.zsh_history`). In this article, I will show you how I tracked down the bug. Spoiler: ultimately, patching Zsh to make it crash loudly and analyzing the crash's core dump was the winning strategy! The good news first Zsh 5.9.2 (released July 12th, 2026) contains a fix for this issue ? open it after reading this investigation to not spoil the fun. Spoiler: link to the upstream fix Zsh fix 53454 The symptom Occasionally, I noticed that commands I knew I executed the day before were not findable in my shell history, meaning pressing Ctrl+R for backward history search yielded no results. Whenever I noticed this, my shell history file contained only very old entries, with years of newer entries missing. The first few times this happened I just restored my shell history from my daily backup and did not bother investigating any further. But the issue kept happening. I noticed that there was no visible corruption in the `.zsh_history` (no non-printable characters or incomplete lines of text), and that the number of lines in the file was not always the same. What wa...