

1. Introducing Muse Glimmer

Source: Simon Willison's Weblog | Published: 2026-08-10T23:56:03+00:00

Link: <https://simonwillison.net/2026/Aug/10/introducing-muse-glimmer/#atom-everything>

Introducing Muse Glimmer Meta are back in the open weights game! Muse Glimmer is a brand new 30B model under a clean Apache 2.0 license (a step up from the janky Llama licenses of old). They claim to have optimized it for exactly the kind of things I'm looking for in a local model: End-to-end Agentic Task Completion. Muse Glimmer achieves strong success rates on full-task benchmarks including DeepSearch QA, MCP-Atlas, ?-Bench and SWE-Bench, which measure its ability to work within scaffolds, write and debug code, and resolve multi-turn requests from start to finish. Reliable Tool Use. The model handles a wide range of function calls, invoking tools with precise schemas throughout extended workflows. Multi-Step Reasoning. Muse Glimmer chains reasoning over long horizons, sustaining coherent plans across complex, extended workflows. [...] Here's a pelican which I generated using LM Studio's 18.16 GB version of the model : I also tried it out with my llm-coding-agent plugin, running against a fresh checkout of Datasette with the prompt: how does auth work? Here's the response , at the end of a long transcript showing all of the tool calls it made to explore the codebase. I ran this u...

2. SQLite compressed text-history prototypes

Source: Simon Willison's Weblog | Published: 2026-08-09T22:05:00+00:00

Link: <https://simonwillison.net/2026/Aug/9/sqlite-text-history-prototype/#atom-everything>

Research: SQLite compressed text-history prototypes I'm perennially interested in options for storing revision histories in relational databases. While out on a dog walk I had a new idea: how about taking the full text of every prior version in a big JSON array of strings and then applying zlib or zstd compression to the whole thing? Surely that would compress really well due to all of the repeated strings. The new GPT?Live voice mode in the ChatGPT iPhone app has got really good, so I discussed the prototype with that. You still can't share URLs to voice conversations, but here's what I said copied from the transcript as a proper stream of consciousness: I have an interesting idea for a scheme for saving all previous versions of a piece of text that's constantly edited in a SQLite database um column in as efficient a way as possible. Okay, so I built these kinds of systems in the past, and it's always difficult to come up with a efficient way to do this. Like the easiest way is you have a row for every previous copy of the previous previous value of the string. But if it's a long document Like 20 kilobytes of data, that means that every single edit adds another 20 kilobytes of dat...

3. datasette-auth-tokens 0.4a13

Source: Simon Willison's Weblog | Published: 2026-08-06T23:07:47+00:00

Link: <https://simonwillison.net/2026/Aug/6/datasette-auth-tokens/#atom-everything>

Release: datasette-auth-tokens 0.4a13 Upgraded for compatibility with 'sqlite-utils 4. Tags: datasette

4. [Sponsor] Drata

Source: Daring Fireball | Published: 2026-08-10T22:25:00Z

Link: <https://drata.com/daring>

Leverage autonomous AI agents to automate compliance, manage internal and third-party risk, and continuously prove your security posture. ?

5. What's New in iOS 27 Beta 5

Source: Daring Fireball | Published: 2026-08-10T21:54:25Z

Link: <https://9to5mac.com/2026/08/10/heres-whats-new-with-ios-27-beta-5/>

Zac Hall at 9to5Mac has a good rundown of new features and changes in beta 5. Some app icons got tweaks (the Siri app in particular looks much cooler now), and a bunch of the existing American and British Siri speaking voices now have the Pace and Expressivity adjustment

sliders. Also interesting: the OS 27 releases were stuck at beta 4 for 22 days. That's a week longer than usual. Will Hains (my Kotoba friend) maintains a wonderful single-page website tracking iOS version release date history . For these summer ?next big version? comparisons, it's best to filter by ?*.0? releases. Do that at Hains's website and you can see a fairly clear pattern: betas 1 through 4 typically last about two weeks each; then starting with beta 5, Apple switches to a roughly weekly schedule until the end of the beta cycle. I wouldn't read anything into that other than to expect beta 6 next week and beta 7 the week after that. ?

6. ?Friday Night Baseball? Will Start Broadcasting Games Live, in Apple Immersive, on Vision Pro

Source: Daring Fireball | Published: 2026-08-10T20:47:43Z

Link: <https://www.apple.com/newsroom/2026/08/apple-mlb-announce-september-friday-night-baseball-schedule/>

Apple Newsroom: For the first time, baseball fans can take in all the action of ?Friday Night Baseball? live in Apple Immersive on Apple Vision Pro, leveraging 3D video recorded in 8K with a 180-degree field of view. On August 28, viewers with Vision Pro will experience the iconic Red Sox vs. Yankees rivalry from unique camera angles that place viewers right in the game, immersed in the drama of each at-bat, the celebrations in the dugout after a big hit, close plays at the outfield wall, and more. The experience features bespoke commentary from Paul Severino (play-by-play), Xavier Scruggs (analyst), and Michelle Margaux (sideline reporter); dedicated replays; and an immersive graphics package that includes team lineups and stats floating in the viewer's space. During breaks in the action, the broadcast remains live, immersing viewers in the stadium experience from all angles as ambisonic microphones capture sounds in Spatial Audio. Hot damn, they even picked the best rivalry in sports to kick it off. ?

7. No, local models will not win

Source: seangoedecke.com RSS feed | Published: Tue, 11 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/local-models-will-not-win/>

Every time a new open-weight AI model is released, people say that local models are the future. Why spend billions of dollars building out datacenters when everyone will just be able to run AI models on their laptops or phones? I think this idea is doomed. No matter how strong open-weight models get, most inference will always happen in AI datacenters. Local models are too weak to be widely used Local models are never going to be as powerful . I think this point should be obvious: all of the current frontier models (closed and open-weights) are far too big to run on anything but a full GPU cluster in a datacenter. Of course, smaller models are getting more intelligent over time. In a year you might be able to run something about as strong as GPT-5.6-Sol on your laptop. But by then, you'll think of GPT-5.6-Sol as too weak to be useful. Many people deny this last point, but it's true: almost everyone's revealed preference is to use the strongest available model in their price range . If AI progress had stalled at GPT-4, I think we could have built some very powerful tools around it, but who'd use GPT-4 today? As LLMs have gotten more capable, our expectations around them have grown:...

8. Failing after Success

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Mon, 10 Aug 2026 07:00:00 GMT

Link: <https://idiallo.com/blog/failing-after-success>

What are you supposed to do after you've created the most successful product? Not just successful, but the most useful. In an ideal world, you would sail into the sunset and enjoy your riches. In this real world, you have to come up with another, even more successful product, or you're a failure. That's why it feels like Stack Overflow has failed. I remember being in the middle of debugging my code when I copied an error and pasted it into Google. Right there in the search results, below the clearly marked sponsored link, was a link to Stack Overflow, where a user had previously faced the exact issue I was facing. I clicked on it and got "Error 503 Service Unavailable." Two seconds later, I was on Reddit, checking if SO was

down for everyone else. Then Hacker News. The memes started flowing shortly after. Developers were touching grass for the first time in years. Productivity was down 100%. If there were any way to measure the success of a company, that would be it. Stack Overflow went from receiving hundreds of thousands of programming questions a day to less than a thousand today. Stack Overflow never positioned itself as just another Q&A forum. It wasn't ExpertsExchange. The...

9. When the Aliens Finally Came to Visit

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Mon, 10 Aug 2026 04:26:49 GMT

Link: <https://idiallo.com/byte-size/when-the-aliens-came-to-visit>

The aliens arrived, and their first instinct was to check our networks. They accessed one of our most important websites. A government website. Unfortunately, they were met with a reCAPTCHA. Are you a robot? They are not robots. But after fifteen minutes of squinting and trying to select all the fire hydrants and traffic lights, they gave up. Defeated by a checkbox, they returned to the stars.

10. Learning from historical mistakes

Source: John D. Cook | Published: Mon, 10 Aug 2026 14:34:21 +0000

Link: <https://www.johndcook.com/blog/2026/08/10/learning-from-historical-mistakes/>

The following extraordinary paragraph comes from Knuth's TAOCP Volume 4A, right before the last set of exercises. Many of the exercises below ask a modern reader to find and/or to correct errors in the literature of bygone days. The point is not to gloat over how smart we are in the 21st century; the point [?] The post Learning from historical mistakes first appeared on John D. Cook .

11. Inverse differential equations

Source: John D. Cook | Published: Mon, 10 Aug 2026 14:32:42 +0000

Link: <https://www.johndcook.com/blog/2026/08/10/inverse-differential-equations/>

In science and engineering classes, you might describe a system using Newton's laws and end up with a differential equation. You then solve the differential equation, analytically or numerically, to see how the solutions behave. You might also do the opposite, especially in a mathematics class: look at what differential equation a set of functions [?] The post Inverse differential equations first appeared on John D. Cook .

12. Edinburgh Fringe: 15 Minutes of Shame ?????

Source: Terence Eden's Blog | Published: Mon, 10 Aug 2026 20:43:10 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-15-minutes-of-shame/>

Sex work is work, so we're told. There's no difference between picking up a night job stacking shelves and a night job jacking-off on screen. Right? Perhaps a couple need a little bit of extra cash and the siren song of OnlyFans is right there. Let's fire up the webcam and release ourselves from financial drudgery! Then reality hits. The script by Sean Stillmaker sounds like it was ripped?

13. Edinburgh Fringe - Kat Ronson: Millennial Girl ?????

Source: Terence Eden's Blog | Published: Mon, 10 Aug 2026 17:45:36 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-kat-ronson-millennial-girl/>

OMG! If you remember A/S/Ling your way through MSN to fill in your MySpace top 8, then you'll love this show. You're also over 30 and, therefore, old. Ew. The stage is littered with a Nokia 3310, Bliss Magazine, and Charlie Red. Katie (15/F/UK) talks us through the complexities of her nascent love life, what shag band colours mean, and which quotes sum up millennial angst. All the while?

14. Edinburgh Fringe - Michael Brunström: William Tell vs the Algorithm ?????

Source: Terence Eden's Blog | Published: Mon, 10 Aug 2026 15:57:27 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-michael-brunstrom-william-tell-vs-the-a>

Igorithm/

The terror of AI has come to comedy. There are a dozen shows this year about the threat that technology poses to the way we live, laugh, and love. In this case, AI stands for Austrian Imperialism and the fight is for Switzerland's independence. Well, sort of. It gets a bit lost in an endless stream of rambling diversions. There's a fine line between surrealist comedy and utter nonsense. If you?

15. A License to Act

Source: Jim Nielsen's Blog | Published: Sun, 09 Aug 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/license-to-act/>

In the past, search engines and the open web granted access to other people's knowledge but the task fell to you to synthesize an understanding from it and build a capability to act. Now with LLMs, it's possible to build things you don't fully understand which then require a perpetual license in order to have the capacity to maintain and modify. LLMs don't just do things for you, they do things instead of you ? and the doing is a huge part of the point. The struggle to understand, to synthesize, to explain, it all has a byproduct: you! Your own capacity to act in the world. But now you can rent that capacity. Every time you reach for the LLM instead of struggling through it yourself, you're borrowing against someone else's capacity that you're not developing in yourself. It's as if I used to be a car mechanic , fixing my own vehicles or looking to other mechanics for advice or a demonstration on approaching my problems. That process produced a working car ? the output I wanted ? but it also produced a capability within myself to fix the next one that came along. But now, for expedience and speed, I drop all my cars off the LLM dealer for fixes and maintenance. Fast and efficient ?...

16. Pluralistic: The bureaucratic AI arms-race is mutually assured destruction (10 Aug 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 10 Aug 2026 07:01:46 +0000

Link: <https://pluralistic.net/2026/08/10/deep-state-wopr/>

Today's links The bureaucratic AI arms-race is mutually assured destruction: The only way to win is not to play. Hey look at this: Delights to delectate. Object permanence: Seymour Cray's tunnels; War on moisture; \$5 wrench cryptanalysis; Nauru files; People's Ride; German transit upholstery fashion; Monopolies v small business; Linkedin will put you in ads; NZ Parliament kicks itself off the internet; Adblock Plus v Facebook adblock-block; Cracking 100m VWs for \$40; Trump as defective machine learning; Canada sucks at internet law; Goodhart's Law of AI. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The bureaucratic AI arms-race is mutually assured destruction (permalink) According to an Economist editorial, "AI is breaking the British state" by making it too easy to file complaints, demands and appeals, which will "drown the state" with "demands as well-crafted as a first-class lawyer's": <https://archive.is/VTrj9> Let's pause a moment to appreciate the Economist's touching cred...

17. UPS finally refunded the overcharged tariff

Source: Old Vintage Computing Research | Published: 2026-08-09T20:36:22.854-07:00

Link: <https://oldvcr.blogspot.com/2026/08/ups-finally-refunded-overcharged-tariff.html>

Last November UPS mistakenly charged me a \$684 tariff on \$355 worth of computer parts shipped from Europe. It was shortly afterwards reduced to \$51.30, but I had prepaid the tariff to avoid warehouse fees, so UPS owed me \$632.70. They finally issued the refund in July after repeatedly E-mailing their accounting office in India approximately monthly. (Phone number? Ha ha ha.) I don't know if these overcharges are still happening, but I've basically stopped importing anything from outside the United States unless they pay the tariff for you, or until there is more clarity. I'd also like to get that \$51.30 back ? especially since they were Apple parts, and one of the parts even said " MADE IN USA ." Until then, any new Australian computer exploration is staying local to the in laws' in regional New South Wales, and it will be awhile

before we get any new Japanese systems to play with.

18. Watch out for cache read costs

Source: Martin Alderson | Published: Mon, 10 Aug 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/watch-out-for-cache-read-costs/?utm_source=rss&utm_medium=rss&utm_campaign=feed

For long agentic sessions cache reads are most of the bill, not input or output tokens. KV caches have shrunk dramatically while cache read pricing has barely moved.