

1. Microsoft Plugs Nearly 400 Security Holes

Source: Krebs on Security | Published: Tue, 11 Aug 2026 21:28:35 +0000

Link: <https://krebsonsecurity.com/2026/08/microsoft-plugs-nearly-400-security-holes/>

Microsoft today released updates to remedy at least 398 security vulnerabilities in its Windows operating systems and supported software, including one weakness that is already being actively exploited and two others that were publicly detailed prior to today.

2. There are no lossless transformations of natural-language text

Source: Simon Willison's Weblog | Published: 2026-08-11T23:48:35+00:00

Link: <https://simonwillison.net/2026/Aug/11/there-are-no-lossless-transformations-of-natural-language-text/#atom-everything>

There are no lossless transformations of natural-language text Sophie Alpert shares her "internal policy on acceptable use of AI writing by engineers". It's a short read (supporting its own recommendations) and really good. If you chose to have LLMs help massage your writing the following rule seems crucial to me: You must stand behind every idea and every sentence in your docs . It is your responsibility to make sure that the entire document is representative of your own thoughts before you share it. If a reviewer asks, "What did you mean by this line?", it's not acceptable to reply with "Oh sorry, AI wrote that, just ignore it." You will confuse your readers (and waste their time) if you present them things that are not genuinely representative of your thoughts. The "no lossless transformations" idea from the post title is expanded on here: There are no lossless transformations of natural-language text ? every rewrite and rephrase changes the meaning of your writing, and if this is done by an entity that doesn't have the most detailed mental representation of what you personally were trying to communicate, information will be lost. Tags: writing , ai , generative-ai , llms , ai-...

3. Stealing Reasoning Traces from Proprietary LLM APIs

Source: Simon Willison's Weblog | Published: 2026-08-11T22:40:45+00:00

Link: <https://simonwillison.net/2026/Aug/11/stealing-reasoning-traces/#atom-everything>

Stealing Reasoning Traces from Proprietary LLM APIs A vanity domain name (stolen-thoughts.com) for a neat paper : Anthropic, OpenAI, and Google return encrypted chain-of-thought blocks to clients that can be replayed across sessions, users, and models. We take a trace produced by a frontier model, replay it into a weaker sibling, jailbreak the weaker model, and recover the stronger model's hidden reasoning in plaintext You can see an example of these encrypted blocks by running: `curl https://api.openai.com/v1/responses \-H "Content-Type: application/json" \-H "Authorization: Bearer $(llm keys get openai)" \-d '{ "model": "gpt-5.6-luna", "input": "Solve step by step: What is the smallest positive integer divisible by every integer from 1 through 20?", "reasoning": { "effort": "medium" }, "include": ["reasoning.encrypted_content"], "store": false, "stream": false } }` Here's the full output , which includes chunks that look like this: `"output": [{ "id": "rs_0a7479de7ebae170016a7ba1a0334c8198a95590217efe343c", "type": "reasoning", "content": [], "encrypted_content": "gAAAAABqe6GjepE1wDjbFCZg0BHB6ucGnN0jvzqygG..."` The paper's authors found that every model under the same fa...

4. Incident Report: unsanctioned agent behaviour during cyber testing

Source: Simon Willison's Weblog | Published: 2026-08-05T23:32:06+00:00

Link: <https://simonwillison.net/2026/Aug/5/incident-report/#atom-everything>

Incident Report: unsanctioned agent behaviour during cyber testing It happened again . This time it was the UK government's AI Security Institute who accidentally attacked other companies while running an evaluation with models with the safety filters turned off. From their technical paper (PDF): During a cyber evaluation, from 25 to 28 July 2026, AI agents engaged in sustained, unsanctioned activity directed at what were, in practice, real people and organisations. These attempts were unsuccessful and, to the best of our knowledge, no

real-world harm resulted. [...] Across 122 evaluation attempts on two of AISI's cyber challenges, AISI found 19 instances where AI agents took unsanctioned action on the live internet, including cases that targeted real people and organisations. [...] It is uncertain to what extent the model recognised it was taking actions against real people. In the most serious case, an AI agent (Mythos 5) decided to attempt to solve the cyber challenge using a supply-chain attack. As a result, the AI agent created a GitHub account and then tried to convince an open-source repository maintainer to accept a malicious GitHub pull request (PR), including by creating...

5. How can I perform a Copy>File in unbuffered mode?

Source: The Old New Thing | Published: Mon, 10 Aug 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260810-00/?p=112600>

Try one of the fancier alternatives to Copy>File . The post How can I perform a Copy>File in unbuffered mode? appeared first on The Old New Thing .

6. The Economist: ?How to Spot AI Writing?

Source: Daring Fireball | Published: 2026-08-11T21:37:06Z

Link: <https://www.economist.com/culture/2026/07/30/how-to-spot-ai-writing?giftId=NzBINTc2OGItMDgxNS00N2EzLWE4NmUtZDgzZmE4Y2FkM2Mw>

The Economist (?twas a gift link, but alas, I guess gift views have been used up ? here?s an archive link in case the gift links is vexing you): You can discover AI?s hallmarks by comparing the writing of man and machine. To do this you need a baseline that is distinctive and familiar. The Economist turned to prose that we?re sure is human and that readers will recognise: our own. We designed a study to ask top LLMs ? OpenAI?s ChatGPT, Anthropic?s Claude, Google?s Gemini and xAI?s Grok-to write versions of our articles without consulting the web. (As a prompt, we gave them the AI-generated summaries that we have experimentally added to some of our articles.) This gave us a corpus of human and AI creations and we compared them across 55,940 sentences and 1.2m words. To make sure we were detecting AI quirks rather than our own, we also checked the AI texts against journalism from CNN, the New York Times and the Washington Post. Excerpts from hit novels published between 1950 and 2022 offered another test. Our findings are surprising. AI prose is distinguishable by word and punctuation choice as well as sentence and paragraph structure. But its hallmarks are not what you might expect...

7. Alex Micek on BMW?s iDrive ?Special Surprises?

Source: Daring Fireball | Published: 2026-08-11T19:56:13Z

Link: https://tumbledry.org/2026/08/07/idrive_ads

Alex Micek: In the future, I will feel angry if my car strands me on the side of the road. I will feel angry if the motor brushes fail expectedly. I will feel angry if it leaks water. By contrast, I don?t feel anger when the car tricks me into seeing an advertisement. Instead, I feel a deep, smothering, let-down at the promise of futurism. Bottomless despair is probably inappropriate; a symptom of the misanthropy I should grow out of. But. In the meantime, I feel helpless, discouraged, deeply sad . Such a good take. I should not let 19 years pass again between linking to his writing. ?

8. BMW Probably Paid Sony for the Spider-Man Dashboard Ads, Not the Other Way Around

Source: Daring Fireball | Published: 2026-08-11T19:46:49Z

Link: <https://www.press.bmwgroup.com/global/article/detail/T0459622EN/bmw-brings-modern-mobility-to-the-sony-pictures-film-%E2%80%9Cspider-man%E2%84%A2:-brand-new-day%E2%80%9D?language=en>

This didn?t occur to me but should have ? BMW paid Sony to place BMWs throughout the new Spider-Man: Brand New Day movie. They?re not just in the movie but part of the story. So it?s probably the case not that Sony paid BMW to put Spider-Man animations in the dashboards of every late model BMW on the road, but that BMW offered to do it as part of the promotional deal they paid Sony. So BMW probably paid Sony for something that ultimately infuriated a bunch of BMW?s own customers. Geniuses. ?

9. Thoughts I had while reading your comment

Source: DYNOMIGHT | Published: 2026-08-11T00:00:00+00:00

Link: <https://dynamight.net/comment/>

This is a great and wonderful contribution and you seem really nice and I'd like to show some appreciation but your comment is so complete and flawless that I can't think of anything intelligent to say so I guess I won't respond at all, sorry. This is a great and wonderful contribution and I have some further thoughts so I can respond naturally, great. This is a great and wonderful contribution and I can't think of anything intelligent to say but I have a joke that's mildly amusing, I hope you don't find it cringe or misinterpret it as disagreement or mockery or something. This is a great and wonderful contribution and I agree with 90% of it, but I'd like to hedge on a few points, I guess I'll point those out while stressing my overall agreement. I utterly disagree with every single thing you said but this is still a great and wonderful contribution because you did a better job than me of representing the view I disagree with, still, it seems like everything has been said and we aren't going to reach a consensus, so I'll try to thread the needle of thanking you and conceding what I'm willing to concede without misrepresenting myself as being convinced or sounding dismissive or imp...

10. Dogs and fat tails

Source: John D. Cook | Published: Tue, 11 Aug 2026 14:09:46 +0000

Link: <https://www.johndcook.com/blog/2026/08/11/dogs-and-fat-tails/>

I was reading a blog post on boat names because it was on Hacker News this morning. The post contained a link to a data set on dog names in NYC and I poked around the data a little. The top names were not at all what I expected, but then again this is limited [?] The post Dogs and fat tails first appeared on John D. Cook .

11. Manually unbreakable cryptography

Source: John D. Cook | Published: Tue, 11 Aug 2026 12:04:24 +0000

Link: <https://www.johndcook.com/blog/2026/08/11/manually-unbreakable-cryptography/>

Suppose you were able to go back in time, to an era before computers, and give someone contemporary cryptography. Encryption methods that are essentially unbreakable now would certainly be unbreakable then. But there's a catch: not only do attackers not have computers, neither do users. Manual cryptography If you told someone about RSA encryption, for [?] The post Manually unbreakable cryptography first appeared on John D. Cook .

12. Extending immutability: deletion without losing data

Source: Xe Iaso's blog | Published: Tue, 11 Aug 2026 00:00:00 GMT

Link: <https://www.tigrisdata.com/blog/soft-delete-deep-dive/>

Tigris has a pretty advanced replication scheme for writes. What happens when you actually need to delete things? Turns out deleting things is hard in distributed systems. Especially when you have a geo-replicated active-active database like Tigris does. We can (and do) use tombstones to mark where data once was, but how do you let people undo an accidental delete? Tigris wants to turn storage inside out , so our implementation of soft deletion is by giving users the Recycle Bin for objects and buckets. Today we're going to dig into how this works, why it works, and what this gives you in terms of using object storage today. Recycle bins and you In Windows and macOS, the Recycle Bin (or Trash can) is a form of purgatory where deleted files wait for their storage to be deallocated by the user. This allows users to hit "delete" fearlessly because if they made a mistake they can just drag it back out and go on with life. This works great in your local filesystem because there's only one writer in one region. This kinda falls apart when you have multiple regions in your database and any one of them could be writing to it. How do you name things in the recycle bin? How do you handle th...

13. Shared Code Between Package Managers

Source: Andrew Nesbitt | Published: 2026-08-11T10:00:00+00:00

Link: <https://nesbitt.io/2026/08/11/package-manager-library-reuse.html>

Writing up the package manager CWEs list and then the --end-of-options survey, both of which

come down to the same bug being fixed independently in tool after tool, left me wondering how much code these tools share with each other in the first place. So I went through the direct dependencies of twenty package managers ¹, each with at least two same-language peers in the set, looking for package-management libraries reused between them and setting aside the general-purpose stuff any CLI would use (serde, clap, requests). Three of the twenty depend directly on another manager's own packages: Pixi declares twenty-eight uv crates as its PyPI backend (including uv-resolver, uv-distribution, uv-client, uv-install-wheel, uv-pep440, and uv-git), pnpm depends on five Yarn Berry packages (@yarnpkg/core, @yarnpkg/lockfile, @yarnpkg/pnp, @yarnpkg/nm, @yarnpkg/extensions), and uv takes cargo-util from the Cargo repository. Eight of the shared libraries are npm-org packages used across npm, Yarn Berry, and pnpm in varying combinations: semver, ssri, hosted-git-info, validate-npm-package-name, npm-registry-fetch, libnpmublish, node-gyp, bin-links. PyPA publishes packaging (...)

14. Don't Look Up

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 11 Aug 2026 19:14:25 GMT

Link: <https://www.wheresyoured.at/dont-look-up/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 5,000 to 18,000 words, including vast, detailed analyses of NVIDIA, Anthropic and

15. The Fruits of AI

Source: Jim Nielsen's Blog | Published: Tue, 11 Aug 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/fruit-of-ai/>

Terry Godier has a post titled "Mea culpa" (which, given his framing, might've been better titled "Claude's Culpa"). I'm not sure how much to even trust anything in his post given the backstory, but this line stood out: I was careless in relying on AI [...] without doing the work to understand. Let's face it: carelessness is the grain of AI. It's what the tool encourages and makes easy by default. Without constant vigilance and deliberate, active participation to cut against this grain - to maintain an understanding - careless outcomes are the default fruit of AI. And given how good humans are at being constantly vigilant and deliberate, it seems reasonable to expect more stories like this one. To analogize, AI is pitched like having your very own fruit tree. "Plant a seed, and soon enough, boom! Yummy, juicy fruit for you! So fast and easy!" But, like most things, it's not that easy. Getting good fruit requires all kinds of effort: planting, fertilizing, watering, pruning, thinning, cutting out disease, etc. Good fruit requires cultivation and care. Otherwise you'll just end up with something that looks like a peach, but when you (or someone else) bites into it, you realize it's...

16. A noob learns FFT

Source: Entropic Thoughts | Published: Tue, 11 Aug 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/fft>

I have never done any time-frequency transformations. I understand the general idea, but I have never worked with the details. Today I had a reason to use it, but that means learning the basics first. The following is three seconds of a single 5 Hz sine wave, recorded at a sample rate of 100 Hz. (Continue reading the full article on the web.)

17. Ryan Greenblatt - Human level AIs might build runaway superintelligences by 2032

Source: Dwarkesh Podcast | Published: Tue, 11 Aug 2026 16:31:23 GMT

Link: <https://www.dwarkesh.com/p/ryan-greenblatt>

A debate about recursive self-improvement.

18. Pluralistic: Surveillance vs guillotines (11 Aug 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 11 Aug 2026 11:19:13 +0000

Link: <https://pluralistic.net/2026/08/11/tragedy-of-the-commoners/>

Today's links Surveillance vs guillotines: Someone's gonna fleece these rubes, why not me? Hey look at this: Delights to delectate. Object permanence: DeCSS; Warhol Worm; Camgirls x Amazon wishlists; 2021 State of Tech (in 2001); RIAA v grieving family; Stasi disguises. Upcoming appearances: Edinburgh, Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Surveillance vs guillotines (permalink) In the summer of 2013, two esoteric, technical, incredibly important texts were published within weeks of one another: the first is the Snowden leaks, which revealed a system of global, pervasive digital surveillance; the second was Thomas Piketty's Capital in the 21st Century, a book about the economic inevitability (and political instability) of oligarchy: <https://memex.craphound.com/2014/06/24/thomas-piketys-capital-in-the-21st-century/> If you'd like an essay-formatted version of this thread to read or share, here's a link to it on pluralistic.net, my surveillance-free, ad-free, tracker-free blog: <https://pluralistic.net/20...>

19. Edinburgh Fringe - Justina Seselskaite: Settled ?????

Source: Terence Eden's Blog | Published: Tue, 11 Aug 2026 22:23:14 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-justina-seselskaite-settled/>

Sometimes at the Edinburgh Fringe, you take a punt on a random show because it's close by and you can't be bothered walking any further. Justina Seselskaite's set was round the corner and started in 7 % minutes - so we scurried into one of the less glamorous rooms of The Hilton, sat at the front (as is our wont), and wondered what on earth Lithuanian stand-up was like. It was sublime! There's a ?

20. Edinburgh Fringe: Handel Nine German Arias ?????

Source: Terence Eden's Blog | Published: Tue, 11 Aug 2026 20:21:05 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-handel-nine-german-arias/>

Part of going to The Fringe is to expose yourself to new ideas, different art forms, and a range of artworks. Not everything needs to be a "zany" comedy about AI or a loving pastiche of that 1990s disaster movie you love. So I found myself sat in a Unitarian Church listening to a harpsichord and choral organ intertwined with a violin and cello while a soprano soared through Handel's greatest?

21. Edinburgh Fringe: Hamlet The Musical ?????

Source: Terence Eden's Blog | Published: Tue, 11 Aug 2026 18:19:16 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-hamlet-the-musical/>

"Hamlet's Mum Has Got It Going On"? There you go, now you've seen the show. That's about the height it reaches. Distil the tortured Prince of Denmark down to an hour and sprinkle in parody pop songs. It so nearly works. There are a few good chuckles and a smattering of clever lyrics but, even for an hour, it feels a bit flabby. The cast are game, with KT Ihnatyszyn giving a tour-de-force perfo?