

1. Markdown SVG upgrades

Source: Simon Willison's Weblog | Published: 2026-08-16T23:59:37+00:00

Link: <https://simonwillison.net/2026/Aug/16/markdown-svg-upgrades/>

I started building my markdown-svg-renderer tool in May , but I've since added enough features to it that it's worth talking about here again. It's evolved into my ideal tool for sharing Markdown transcripts that include SVG documents. Given my proclivity for drawing pelicans riding bicycles this is a problem that I needed to solve! The tool is very simple. Navigate to markdown-svg-renderer in your browser and paste in some Markdown to see it rendered... or save that Markdown to a CORS-friendly URL or a GitHub Gist and paste in a URL to that document. The URL option will give you a bookmarkable page, for example <https://tools.simonwillison.net/markdown-svg-renderer?url=https%3A%2F%2Fgist.github.com%2Fsimonw%2F6f9e48293be5c916652d29f0dc0b0657> - which bakes in the URL to this Gist . If you visit the Gist you'll see raw SVG: In the rendered tool that looks like this instead: As you can see, that SVG block in the Markdown has been transformed into a rendered SVG (in this case animated) plus several tabs. The tabs are the really fun bit. The PNG and JPEG tabs render that SVG to those image formats in the browser and lets you copy or download them - useful for sharing on platforms that...

2. Qwen 3.8 27B is excellent, but it defaults to wildly overthinking things

Source: Simon Willison's Weblog | Published: 2026-08-16T22:00:39+00:00

Link: <https://simonwillison.net/2026/Aug/16/qwen-38-27b/>

Friday's big release was Qwen 3.8 27B , an Apache 2 licensed 27B parameter vision-capable LLM from Alibaba's Qwen research lab. I've been looking forward to this one: 27B is an excellent size for running a model on a reasonably specced laptop, and its predecessor Qwen 3.6 27B was impressive. Qwen's self-reported benchmarks for this model are eye-opening. They show a boost from both Qwen 3.6 27B and the closed-weight Qwen 3.7-Plus, which was one of Qwen's strongest models of any size as recently as May this year . It will be interesting to hear what independent benchmarks have to say about the model. I've been running the model on two different machines: my 128GB M5 Max MacBook Pro, and an NVIDIA DGX Spark . On both machines I'm running LM Studio and their 17GB Q4_K_M quantized build . I also tried using llama-server directly on the Spark. The default of extra high results in spectacular over-thinking Qwen's documentation describes the model as defaulting to xhigh for the reasoning effort, and the LM Studio GGUF I've been trying preserves that default: Qwen3.8 comes with official support for reasoning_effort , which can be used to adjust reasoning depth and control cost: xhigh (def...

3. Quoting Dario Amodei

Source: Simon Willison's Weblog | Published: 2026-08-16T15:05:36+00:00

Link: <https://simonwillison.net/2026/Aug/16/dario-amodei/>

I do agree that the public has a negative view of AI (and that this is a big problem), but I don't think it is primarily caused by me or any other AI leader warning about AI's risks. I think it is fundamentally a crisis of trust. I think that ordinary people don't trust companies, governments, or the tech industry and always suspect that we are cooking up some new way to screw them over. The causes of this go back decades and AI is just the latest iteration of it. I don't think that a glitzy marketing campaign with a positive spin (which some have advocated that Anthropic do) is the way to win back that trust ? at this point, saying that AI will cure cancer is more a cliché than it is inspiring, and most people think it is deceptive. The thing that will work is actually curing cancer . I think by far the most accurate criticism of AI companies including Anthropic is that we haven't yet delivered on our big promises to benefit the world. That is totally on us, and I think it's the criticism you should be making, instead of all this stuff about messaging and marketing. ? Dario Amodei Tags: anthropic , ai , ai-backlash

4. The Information Profiles Cami Clark ? Dario Amodei's Wife, Ivanka Trump's Friend,

One-Time Would-Be Pornographer, and Anthropic's First Lady?

Source: Daring Fireball | Published: 2026-08-16T19:30:00Z

Link:

<https://www.theinformation.com/articles/anthropics-first-lady-took-winding-road-top?rc=jfy0lk>

Cory Weinberg, Jemima McEvoy, Jessica E. Lessin, and Stephanie Palazzolo, writing for the paywalled-without-gift-links The Information: As Amodeli has hopscotched the globe to preach about the potential and risks of AI from New Delhi to Davos to Sun Valley Clark has almost always been near his side. Several people who know the couple describe Clark as Amodeli's emotional ballast, someone he has sought counsel from during the turbulence of Anthropic's growth and clashes with Washington over the future of AI. Clark may well be the most consequential person within Anthropic's orbit who doesn't have a formal role at the startup. You might even call her the first lady of Anthropic. [...] And before Female Algorithm Technologies, Clark tried her hand at building a pornographic film company, Eddice, that planned to make movies with high-production values and female protagonists. She also hoped it would have online shopping capabilities. In 2011, she tried to raise money for Eddice from Jeffrey Epstein, who'd pleaded guilty to charges of soliciting a minor for prostitution three years earlier, according to emails made public as part of the Epstein files released by the Department of Jus...

5. Anthropic's Watermark? Text Adulteration in Claude Is a Perversion of Writing

Source: Daring Fireball | Published: 2026-08-16T19:26:33Z

Link: https://daringfireball.net/2026/08/anthropics_watermark_text_adulteration_in_claude_is_a_perversion_of_writing

It's unacceptable for a tool to sacrifice an iota of clarity, coherence, meaning, quality, etc. for the purpose of embedding hidden clues within the text to suggest its provenance. The idea that anything other than *my* needs should factor into the generation of text *for me* is patently offensive.

6. Anthropic's Weak Watermarks Appease a Weak Law?

Source: Daring Fireball | Published: 2026-08-16T17:30:00Z

Link: https://blog.j11y.io/2026-08-12_Anthropics-weak-watermarks-appease-a-weak-law/

James Padolsey, on the Claude-text-watermarking-to-comply-with-an-EU-regulation imbroglio: The same thought that led to this law could have applied to calculators at the time of their inception, had their outputs revealed themselves through artefacts. Thankfully, a sum borne of the brain is treated no differently from one produced by a calculator. Likewise with spellcheckers. To make assistance suspect only once the tool becomes capable enough to compose a whole sentence is not a principled boundary. It is a moral premium placed on difficulty itself. Anthropic has nevertheless chosen a blanket, model-level implementation that appears broader than the law's minimum requirement. That may be convenient compliance engineering, but it discards distinctions the law expressly attempted to preserve. The result is a signal broad enough to implicate harmless and assistive use, yet fragile enough to be removed by a motivated person through substantial recomposition. It risks concentrating suspicion on ordinary and assistive users while remaining weakest against deliberate deception. Padolsey is the creator of Declaude, a delightfully simple web app that allows you to Paste in AI-flavored t...

7. AI text watermarking is not a big deal

Source: seangoedecke.com RSS feed | Published: Sun, 16 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/ai-text-watermarking-is-not-a-big-deal/>

People are pretty unhappy about Anthropic's recent announcement that they're planning to include a hidden watermark in Claude model outputs. Will this lead to a mass exodus from Anthropic models? Will the introduction of watermarking be a meaningful change for users? No. AI text watermarking is not a big deal. It doesn't make the text worse, it doesn't make AI outputs more detectable in practice, it doesn't violate user privacy, and everyone's going to be doing it by 2027 regardless. Watermarked text is not lower-quality There is no meaningful difference in quality between watermarked and unwatermarked text. I wrote about this more here

, but the two popular ways to do it ? Google?s SynthID-Text and Meta?s TextSeal ? are completely transparent to the user. They work by replacing the pseudo-random logit sampler with a different pseudo-random logit sampler. Suppose you were gambling on coin flips with your friends, and instead of flipping a coin you decided to do this: Check the current time since midnight in seconds Count that many words forward in the Encyclopaedia Britannica Count whether the word you land on has an even or odd number of letters 1 That would still be random enough...

8. Proportion of 1s in a Hadamard matrix

Source: John D. Cook | Published: Mon, 17 Aug 2026 01:16:39 +0000

Link: <https://www.johndcook.com/blog/2026/08/16/proportion-of-1s-in-a-hadamard-matrix/>

The first post in the recent series of posts on Hadamard matrices describes a way of constructing new Hadamard matrices from two other Hadamard matrices by taking their Kronecker product. Starting with a Hadamard matrix H_0 and a Hadamard matrix G , you can construct a sequence of Hadamard matrices by $H_{n+1} = G \otimes H_n$ for $[?]$ The post Proportion of 1s in a Hadamard matrix first appeared on John D. Cook .

9. Thoughts on visiting the Edinburgh Fringe as a newbie

Source: Terence Eden?s Blog | Published: Sun, 16 Aug 2026 11:34:21 +0000

Link: <https://shkspr.mobi/blog/2026/08/thoughts-on-the-edinburgh-fringe-as-a-newbie/>

This was my first ever time at The Fringe?. Here are some scattered thoughts looking at the good, the bad, and the annoying aspects of this madly extravagant exhibition of talent. We saw 23 shows together, and one separate show each. There were more hits than misses, but the frustration of wasting your time on a bad show is rather heavy when there are literally hundreds of other shows on offer. P?

10. Book Review: Coding Literacy by Annette Vee ?????

Source: Terence Eden?s Blog | Published: Thu, 13 Aug 2026 11:34:09 +0000

Link: <https://shkspr.mobi/blog/2026/08/book-review-coding-literacy-by-annette-vee/>

This is a disappointing and frustrating book - not least because it so nearly hits the mark. The subtitle is "How Computer Programming Is Changing Writing" - but it never gets close to discussing that topic, let alone deeply examining it. Instead, it is a book mostly focussed on the history of what we mean by "literacy" and, by contrast, what it meant to be "illiterate". As per the standard?

11. Edinburgh Fringe - Heated Rivalry: The Unauthorized Musical Parody ?????

Source: Terence Eden?s Blog | Published: Wed, 12 Aug 2026 21:22:09 +0000

Link: <https://shkspr.mobi/blog/2026/08/edinburgh-fringe-heated-rivalry-the-unauthorized-musical-parody/>

There are at least four different Heated Rivalry shows at the Fringe this year. Earlier this week we saw Edinburgh Fringe - Heated Rivalry: The Musical Parody!. This is different because it is the unauthorized musical parody. Oh, and it is in the much larger Udderbelly venue. Some people spend the fringe watching every possible interpretation of Timon of Athens - we decided to go for doubles of ?

12. How I think about reducing AI costs

Source: Martin Alderson | Published: Sun, 16 Aug 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/how-i-think-about-reducing-ai-costs/?utm_source=rss&utm_medium=rss&utm_campaign=feed

A practical framework for reducing AI costs - auditing token spend, killing legacy models, moving workloads to open weights providers, and fixing the agent and tool inefficiencies that quietly burn tokens.