RSS Digest
Generated at: 2026-08-20 10:17:14

1. smolmachines / smolvm as a sandbox for untrusted Python & JavaScript
Source: Simon Willison’s Weblog | Published: 2026-08-19T23:16:00+00:00
Link: https://simonwillison.net/2026/Aug/19/smolmachines-untrusted-sandbox/
Research: smolmachines / smolvm as a sandbox for untrusted Python & JavaScript I tasked Claude Fable 5 running in Claude Code for web with the following research task: Put https://smolmachines.com through its paces as a fast secure sandbox. Explore what it would take to use this to run untrusted Python and JavaScript code in a way that is limited in what RAM and CPU time it can take up (protection against "while true") with no network access and filesystem access only to designated files Goal is to be able to use this to execute user-provided tasks for things like data transformations It quickly ran into a problem: the Claude Code for web environment can’t run smol machines . Quoting the notes it wrote : This Claude Code container: Linux 6.18.5-fc-v20 (itself a Firecracker guest), 4 vCPU, 15GB RAM. No /dev/kvm, no vmx/svm CPU flags ? no nested virt. smolvm machine run fails as expected: "kvm not available". Plan B: GitHub Actions ubuntu runners DO expose /dev/kvm ? run the real test battery via a temporary workflow on this branch, collect logs, remove workflow in final commit. And Plan B is what it did , installing smolvm and running these tests directly in a GitHub Actions runner...

2. Quoting Jeremy Morrell
Source: Simon Willison’s Weblog | Published: 2026-08-19T22:56:31+00:00
Link: https://simonwillison.net/2026/Aug/19/jeremy-morrell/
My hypothesis is that there is a new opportunity for Extensible Software on the web . LLMs radically lower the cost of authoring extensions, and modern sandbox primitives lower the deployment cost and provide good security boundaries. We can build our app as a solid, accountable core, and allow users to safely extend it in many directions by having LLMs fill in the missing pieces. We can give our users super powers. ? Jeremy Morrell , Extensible Software in the age of LLMs Tags: sandboxing , llms , ai , generative-ai

3. Conceptual integrity and counting lines of code
Source: Simon Willison’s Weblog | Published: 2026-08-19T22:46:07+00:00
Link: https://simonwillison.net/2026/Aug/19/conceptual-integrity-and-counting-lines-of-code/
Last week I recorded an episode of the Talking Postgres podcast with Claire Giordano on the subject of "How AI is changing software development". We had a really great conversation. Here are a couple of my highlights from a lightly edited transcript (prompt to Claude: "very minor edits to remove disfluencies"). This is the latest version of an argument I’ve been trying to build about why sometimes it does make sense to talk about lines of code as an indicator of productivity with coding agents, at 35:01 : A lot of people will tell you it makes no sense to measure productivity in lines of code. I?d actually disagree, because there?s a hard limit. In the before-times, a software engineer could produce a few hundred lines of production-ready code per day ? and 200 lines of working, debugged, production-level code is an incredibly good day. Most days you?d produce 50 or 60. If agents let you produce a thousand lines of debugged code, that really is a very meaningful improvement ? as long as the code is the same quality: maintainable, tested, all of that. You can get to that point with agents, but it takes a huge amount of skill and knowledge and experience. That?s what senior engineer...

4. ?Dickover? Makes It Into The Guardian
Source: Daring Fireball | Published: 2026-08-18T16:03:32Z
Link: https://www.theguardian.com/technology/2026/aug/18/dickovers-baggravation-botiquette-18-new-words-tech-hellscape
Stuart Heritage, writing for The Guardian, under the splendid headline ?Dickovers, Baggravation and Botiquette: 18 New Words to Describe Our Tech Hellscape?: The ?dickover? is the scourge of the modern age. Coined in May by John Gruber, a tech writer, it is defined as ?a modal panel,

popover, or curtain presented by a website or app, deliberately obscuring its own content to frustrate the user?. Essentially, it is the thing (or more often, things) you have to click away in order to read the thing you wanted to read. It is the ?accept cookies? box, the ?allow notifications?? box, the ?sign in with Google? box, the pop-up video box, etc. We shouldn?t fool ourselves into thinking that dickovers are the only terrible thing about technology in 2026. Until now, though, we have lacked the vocabulary to describe them, so here are some suggestions. This was a delight to see, but it would have been nicer if it had included a link. I remain quite fond of my self-exemplifying definition . Amongst Heritage?s suggested coinages, I like baggravation (maddening self-checkout retail kiosks) and Schrödinger?s Wi-Fi (public Wi-Fi, especially on trains and planes, that simultaneously does and does no...

5. OpenAI Pot Complains That Google Kettle Is Black
Source: Daring Fireball | Published: 2026-08-18T14:54:28Z
Link: https://x.com/thsottiaux/status/2083373529081291076?s=12
Thibault Sottiaux, the OpenAI genius in charge of Codex and the new ChatGPT Homer Simpson car, on Twitter/X: I don?t come often to GMail [sic] (OpenAI is a Slack company), but I swear, every time I do open it, there is a new button somewhere around the top right corner. I mean, he?s right about Google?s zero-taste just-keep-cramming-shit-in approach to Gmail, but it?s almost hard to believe he could post this without any self-awareness that it?s the exact same approach his Codex team has applied to ChatGPT. Gmail debuted in 2004 homely but wonderfully simple (which I?d argue was Google?s original, but now long-lost UI brand). It was just your email. Now it?s too complex to explain. ChatGPT?s original incarnation was attractive-in-today?s-minimalist-fad-way and wonderfully simple. It was just AI chat. Now it?s too complex to explain. (ChatGPT?s original incarnation still exists, in the form of the hard-to-find-if-you-don?t-already-have-it-installed Coke ChatGPT Classic , just like how Gmail, after Google started needlessly complicating it, used to have a ?classic? view that many users preferred but Google soon took away .) ?

6. Nature Is Healing: MacOS 27 Golden Gate Beta 6 Adds Redesigned Traffic Light Window Controls
Source: Daring Fireball | Published: 2026-08-18T12:32:38Z
Link: https://9to5mac.com/2026/08/17/macos-27-golden-gate-beta-6-features-redesigned-traffic-light-window-controls/
Zac Hall, 9to5Mac: macOS 27 Golden Gate beta 6 introduces redesigned traffic light window controls. The new look resembles older Mac OS X systems with more detail in each button. Nice. But these changes can?t improve fast enough for me ? I am impatient. The MacOS user interface had been in a slow, steady, sad, no-fun decline for years, and then, a year ago, the bottom fell out with the disaster of 26 Tahoe. But all the changes in 27 Golden Gate are in the right direction. Golden Gate is not a mixed bag. Designers with taste and talent are clearly back at the helm . And I love that they?re still pushing UI and icon changes in mid-August, entering the home stretch of the .0 beta cycle. ?

7. Hands-on with Raspberry Pi’s CM5 Programming Jig
Source: Jeff Geerling | Published: Wed, 19 Aug 2026 02:13:00 -0500
Link: https://www.jeffgeerling.com/blog/2026/cm5-programming-jig/
In the before-times, when Raspberry Pi CM5s were (relatively) affordable, I built a number of Pi clusters (example), and one of the most annoying parts of the build was flashing Raspberry Pi OS to all the Pis. One, two, or even three Pis isn’t a big deal, but once you hit 4+, the process of plugging the Compute Module into a carrier board, plugging that into a computer, managing Raspberry Pi Imager, and trying to match up details like a hostname, MAC address, and the physical Pi itself, gets annoying.

8. Why did the Microsoft Entertainment Pack for Windows have a special sticker announcing that it also had Tetris?
Source: The Old New Thing | Published: Tue, 18 Aug 2026 14:00:00 +0000

Link: https://devblogs.microsoft.com/oldnewthing/20260818-00/?p=112621
A contingency plan. The post Why did the Microsoft Entertainment Pack for Windows have a special sticker announcing that it also had Tetris? appeared first on The Old New Thing .

9. Good writing is obvious, not original
Source: seangoedecke.com RSS feed | Published: Tue, 14 Jul 2026 00:00:00 GMT
Link: https://seangoedecke.com/good-writing-is-obvious-not-original/
When you write, you should try to say things that are obviously true, and spend very little time worrying about whether you?re being original. Ironically, this is the best way to do truly original writing. Every important idea has been discussed already. I learned this in grad school for philosophy 1 , but it?s true across the board. Almost every field 2 has generations of very smart people who?ve spent their whole lives thinking about the biggest problems. If you try to only write about ideas that are truly brand-new, you will restrict yourself to writing about ideas that weren?t important enough to be considered before : in other words, trivia and ephemera. What happens if you focus on the obvious instead? Writing about things that are obviously true is surprisingly difficult. Writing about anything is difficult, because reality has a surprising amount of detail . Still, writing about obvious things is especially difficult, because they?re almost too obvious to see . It?s trivial for me to come up with non-obvious things I believe (for instance, I don?t like database foreign keys), but it?s hard for me to articulate the parts of what I believe that are truly fundamental. I thin...

10. What Is Reasoning
Source: Armin Ronacher’s Thoughts and Writings | Published: 2026-08-19T00:00:00+00:00
Link: https://lucumr.pocoo.org/2026/8/19/what-is-reasoning/
A few weeks ago a paper was shared that showed how to extract reasoning traces from closed-weight models. Together with online discussions about tricking models into leaking them, it made me investigate it more out of curiosity. Twitter seems full of half-truths and confusion about how this works, so perhaps this helps some to understand what is happening. Hiding Traces Reasoning traces are usually hidden from us. We have lamented this , but mostly have to accept it. Open-weight models thankfully reveal them, and from their behavior you can see that their traces can be long and confusing. This is probably a good reason to separate them from what is normally shown to users. At minimum, UIs need to detect them. The industry has done a good job at making reasoning traces sound special and exotic, but they really are just text: the model is trained to emit its thinking into a scratchpad as part of its response, before its final answer. GPT-OSS?s Harmony response format makes this easy to see: analysis I need to work this out ... assistant final The answer is ... The markers are special tokens, but the reasoning between them uses ?the same text? as the final answer (just that GPT chain...

11. Anubis continues to expose new ways people configure webservers
Source: Xe Iaso’s blog | Published: Wed, 19 Aug 2026 00:00:00 GMT
Link: https://xeiaso.net/notes/2026/anubis-csp-worker-hell/
TL;DR: if admins turn off browser features, those features won’t work in confusing ways that are annoying to debug

12. Asymmetric Agents
Source: Terence Eden?s Blog | Published: Wed, 19 Aug 2026 11:34:52 +0000
Link: https://shkspr.mobi/blog/2026/08/asymmetric-agents/
One of the promises of the AI-filled future is that we’ll all have highly capable autonomous servants "agents" to work for us. Even back in the 1990s, it was a common trope in future-gazing to insist that a person’s personal slave agent would negotiate on their behalf. Ask your meek digital pal agent to find you a restuarant in town, tonight, for a hot date, catering to your dietary preferences, ?

13. Why do OpenAI’s GPT-2 weights beat mine? Part two: IFT dropout
Source: Giles’ blog | Published: Fri, 14 Aug 2026 19:00:00 +0000

Link: https://www.gilesthomas.com/2026/08/why-do-openai-gpt2-weights-beat-mine-2-ift-dropout
Wait, dropout! https://arxiv.org/pdf/2101.03961 -- p11 -- Regularizing large sparse models. Our paper considers the common NLP approach of pre-training on a large corpus followed by fine-tuning on smaller downstream tasks such as summarization or question answering. One issue that naturally arises is overfitting since many fine-tuning tasks have very few examples. During fine-tuning of standard Transformers, Raffel et al. (2019) use dropout (Srivastava et al., 2014) at each layer to prevent overfitting. Our Switch Transformers have significantly more parameters than the FLOP matched dense baseline, which can lead to more severe overfitting on these smaller downstream tasks. Could that be it?! Model conf has dropout in it, so... Drop rate was zero in the GPT-2 weights JSON files I had! Nothing in the paper but Raschka and also https://huggingface.co/docs/transformers/v4.48.2/en/model_doc/gpt2#transformers.GPT2Config When I did the extended/two-epoch runs, I wound up doing many more epochs for the IFT. Matching pre-training OK, firstly let’s do a run where we match the models’ pre-train dropout. Test loss Base dropout IFT epochs IFT score IFT rank OpenAI weights: medium 3.231442 Yes...

14. Use the built-in GELU, don’t roll your own!
Source: Giles’ blog | Published: Thu, 20 Aug 2026 02:00:00 +0000
Link: https://www.gilesthomas.com/2026/08/built-in-gelu
Unsurprisingly, PyTorch’s own built-in GELU function is faster than the hand-rolled one I’ve been using to date. But I was surprised at how much faster using it made things when training my models. I discovered this accidentally just now while working on something unrelated, but am logging the details here for anyone else that might find it useful. The headline numbers: the same code, training the same model on the same data, ran at about: 21,000 tokens per second using the hand-rolled GELU from Sebastian Raschka ’s book " Build a Large Language Model (from Scratch) ". 25,000 tokens per second using PyTorch’s built-in GELU with no arguments. 25,000 tokens per second using the built-in GELU with approximate="tanh" , which uses the same maths as Raschka’s version under the hood. That’s a 20% increase in throughput for both of the built-in versions -- definitely nothing to be sneezed at. And what is particularly interesting is that there aren’t that many GELUs going on -- it’s a GPT-2 small-style model, with 12 layers. So that’s 12 GELUs handling tensors shaped (batch_size, seq_len, 4 * d_emb) , which is (6, 1024, 3072) for my training setup. Given that the rest of the model is doing...

15. Pluralistic: The ordinariness of evil (19 Aug 2026)
Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 19 Aug 2026 07:36:27 +0000
Link: https://pluralistic.net/2026/08/19/banaility/
Today’s links The ordinariness of evil: Stop selling AI (bad). Hey look at this: Delights to delectate. Object permanence: Dot-crash bumwad; AP v fair use; Nym Wars; What’s ctrl-F? Pentagon lost $6.5T in a year; Voter ID is voter suppression; Become unoptimizable. Upcoming appearances: Sydney, Melbourne, Brighton, London, South Bend. Recent appearances: Where I’ve been. Latest books: You keep readin’ em, I’ll keep writin’ ’em. Upcoming books: Like I said, I’ll keep writin’ ’em. Colophon: All the rest. The ordinariness of evil (permalink) Maybe it seems weird that AI bosses won’t stop publicly rending their garments about the terrible potential of their products, from the jobspocalypse that will ensue when AI can do our jobs better than us, to the impending moment when the word-guessing programs learn too many words, wake up and turn us all into paperclips. It seems weird that they won’t stop fretting about this terrible potential ? until you realize that every public pronouncement about this terrible potential is also a public boast about its potential, period. That’s very, very important, because all AI really has is potential. Actual, existing AI is useful at the margins, if you...