

1. Anthropic's best AI model struggles to attract users as cheaper tools thrive

Source: Simon Willison's Weblog | Published: 2026-08-23T20:24:52+00:00

Link: <https://simonwillison.net/2026/Aug/23/anthropics-best-ai-model-struggles-to-attract-users-as-cheaper-t/>

Anthropic's best AI model struggles to attract users as cheaper tools thrive A few interesting numbers in this FT story gathered from "people with knowledge of the matter": Anthropic's "annualized revenue" for July is up to \$65bn - it was \$47bn in May, and I collected more historic numbers here . Anthropic expect Q3 to be profitable according to the same model they used to declare Q2 profitable. "It also told investors that it had 6,000 customers that spend \$100,000 annually or more." As for OpenAI, "annualised revenue has jumped 35 per cent in the quarter to date and is now over \$40bn, with the launch of GPT 5.6 in July jolting the company's performance after a sluggish start to the year". This article also introduced me to the Ramp AI index , which uses billing data from 70,000 Ramp credit card using companies to estimate model adoption. Here's Ramp's breakdown of Anthropic model spend for July 2026, which looks reasonable given that Opus 5 was only released on July 24th, and supports the idea that Fable's cost has made it a less popular model: Opus 4.8: 28.0% Sonnet 4.6: 8.3% Fable 5: 8.0% Opus 4.6: 6.9% Sonnet 5: 3.6% Opus 5: 3.5% Opus 4.7: 1.7% Sonnet 4.5: 1.3% Haiku 4.5: 1.0...

2. Quoting Drew Breunig

Source: Simon Willison's Weblog | Published: 2026-08-23T19:55:30+00:00

Link: <https://simonwillison.net/2026/Aug/23/drew-breunig/>

Prior to Fable, it felt silly to waste too much time improving your coding harness or context strategies. A new model would arrive at the same price (or cheaper!) and paper over most of your problems. But then Fable landed. It was (and still is!) incredible . But the cost was so high and Opus was good enough (as was 5.6, K3, and even GLM) for most of the code we needed. So we started to think about what work went where. ? Drew Breunig , Fable & The End of the Free Lunch Tags: drew-breunig , anthropic , claude , llm-pricing , ai , llms , generative-ai , claude-mythos-fable

3. Quoting Matt Webb

Source: Simon Willison's Weblog | Published: 2026-08-21T15:06:26+00:00

Link: <https://simonwillison.net/2026/Aug/21/matt-webb/>

After I released version 1.0, I figured I would have to do the rotations myself. So I sat down with ChatGPT and I didn't get it to write the code, but I got it to educate me. With a patient, interactive tutor, I was able to finally do what I hadn't by reading books and asking mathematician friends ? I learnt how to use quaternions just enough to make the app work. So learning doesn't stop just because I outsource a bunch of thinking to AI. It pushes me to learn more. I like that as an outcome. ? Matt Webb , Galactic Compass 2: now with new augmented reality mode Tags: matt-webb , generative-ai , chatgpt , education , ai , llms

4. Cooper Sharp Proves That American Cheese Can Be Great Cheese

Source: Daring Fireball | Published: 2026-08-23T17:36:28Z

Link: <https://sixcolors.com/member/2026/08/this-week-in-apple-lets-fight/>

John Moltz, in his weekly members-only column for Six Colors (no gift links, but Moltz's column is worth the subscription in and of itself): This week Anthropic announced that in order to identify AI from human-generated content it would be watermarking all text it generates by subtly altering the randomness at which its models picked the next word. While this might seem a perfectly cromulent means of accomplishing a worthy goal, it really rubbed John Gruber the wrong way, as if someone suggested putting autocomplete Swiss on his Philly cheesesteak AI. Cooper Sharp cheese is the way to go. The best new cheesesteak joint in town ? Uncle Gus's ? only offers Cooper Sharp, in fact. Surely a fair number of the hundreds and hundreds of bees that have nested in Gruber's extremely fancy bonnet are a result of the fact that this all

ostensibly in response to a mandate from the EU. But, as James Thomson has suggested, Anthropic might actually want to watermark so that it doesn't scoop up its own output to train its models, thereby degrading the results. In fact, the company would probably love to be able to figure out other companies' watermarks and have them not be able to figure out its wa...

5. Bluesky's Active User Base Is Shrinking, and Remains Tiny Compared to Threads and Twitter/X

Source: Daring Fireball | Published: 2026-08-21T13:36:58Z

Link: <https://techcrunch.com/2026/08/11/blueskys-active-user-base-is-shrinking-as-its-focus-expands-beyond-the-app/>

Sarah Perez, writing at TechCrunch last week: Decentralized social network Bluesky was one of the bigger beneficiaries of the exodus from Elon Musk's X in November 2024, following the U.S. elections. But now, nearly two years later, the social network and would-be X competitor is struggling to hold on to its momentum. According to data from digital intelligence provider Similarweb, Bluesky's mobile app had 10.4 million monthly active users worldwide in June 2026, down 27.2% year-over-year. In addition, mobile daily active users continued to decline, falling 25.6% year-over-year in July to around 3 million. As always, take all third-party analytics and usage data about app usage with an enormous grain of salt. Third-party data companies can only estimate this data, they can't know it. It's like comparing polling data to election results. But I think it's obviously true that Threads and Twitter/X both have hundreds of millions of active users, and Bluesky has like maybe 10 million. Even if Bluesky has three times more actual users than Similarweb's data suggests, that would still suggest that Bluesky's user base is an entire order of magnitude smaller than Threads or Twitter/X. A...

6. Bluesky Is Full of Anti-AI Zealots

Source: Daring Fireball | Published: 2026-08-21T12:52:47Z

Link: <https://bsky.app/profile/masnick.com/post/3mtk7cuvbok2x>

Mike Masnick, in a thread on Bluesky: Multiple people I know have told me that they love the idea of Bluesky, and want it to succeed, but have abandoned it for X because they use agentic tools in their work and find them incredibly useful, and feel that any mention of their usage here leads to hate and ridicule. I wouldn't bring this up if it was just one person, but it was multiple people since the beginning of August who mentioned in passing something to the effect of either 'I'd spend more time on Bluesky if people weren't so anti-AI' or 'As much as I hate to be there all the useful AI talk is on X.' [...] I'm not telling anyone what to do or what tech to use or what to like or not like. I'm just relaying that the reputation this place now has is one that a large segment of the population who are NOT fascists, saying they don't feel comfortable here because a main area of their interest is considered worthy of being mocked and attacked. And even as I've explained that I've been able to have positive conversations about productive AI usage here with many of you, the reputation of the community here as insisting their own lived experience can't possibly be real has driven them aw...

7. What exactly is modified about a modified Bessel function?

Source: John D. Cook | Published: Sun, 23 Aug 2026 18:40:57 +0000

Link: <https://www.johndcook.com/blog/2026/08/23/modified-bessel-function/>

Special functions often have arcane names that not very helpful without some context. The previous post goes into some reasons for this. This post will expand on a point at the end of the post about 'modified' functions. Things are given their names for reasons. Discovering those reasons may help you understand their motivation and [?] The post 'What exactly is modified about a modified Bessel function?' first appeared on John D. Cook.

8. Why special function terminology is arcane

Source: John D. Cook | Published: Sun, 23 Aug 2026 17:56:03 +0000

Link: <https://www.johndcook.com/blog/2026/08/23/arcane-terminology/>

Special functions are special because they're useful. They can also be shrouded in arcane terminology. These two facts are related. The more widely useful a function is, the more likely

it is that the function will be discovered independently multiple times. Independent discoveries lead to varying definitions and notations. For example, there are two widely [?] The post Why special function terminology is arcane first appeared on John D. Cook .

9. Death to px, long live ch!

Source: Terence Eden's Blog | Published: Sun, 23 Aug 2026 11:34:52 +0000

Link: <https://shkspr.mobi/blog/2026/08/death-to-px-long-live-ch/>

Pixels are a lie. Even if you think you're drawing something with "pixel perfect" accuracy, your monitor is lying to you. There is no grid of platonically perfect squares. In CSS, pixels are a double lie: Note that 1px doesn't necessarily equal one physical device pixel. On HD displays, it may span multiple physical pixels. Similarly, 1cm in CSS often doesn't correspond to one hundredth of SI?

10. Fixing an eMachines EL1200 BIOS bug with Claude

Source: Downtown Doug Brown | Published: Sun, 23 Aug 2026 19:28:10 +0000

Link:

<https://www.downtowndougbrown.com/2026/08/fixing-an-emachines-el1200-bios-bug-with-claude/>

I'm no stranger to pushing hardware past its limits. For example, seven years ago, I tracked down an issue that prevented 16 GB of RAM being used in a motherboard that only supported 8. It ended up being a one-line GRUB hack to fix one of the ACPI tables that was mistakenly overlapping a PCI [?]

11. Welcoming the Sri Lankan Government to Have I Been Pwned

Source: Troy Hunt | Published: Sun, 23 Aug 2026 06:39:45 GMT

Link: <https://www.troyhunt.com/welcoming-the-sri-lankan-government-to-have-i-been-pwned/>

Today, we welcome the 48th government onboarded to Have I Been Pwned's free gov service: Sri Lanka. Sri Lanka CERT now has access to monitor Sri Lankan government domains against the data in HIBP, helping identify exposed government accounts and respond when they appear in new data breaches.

12. The summer of open weights

Source: Martin Alderson | Published: Sun, 23 Aug 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/the-summer-of-open-weights/?utm_source=rss&utm_medium=rss&utm_campaign=feed

Winter 2025 was the tipping point for coding agents. Summer 2026 is shaping up to be the same for open weights - and compute supply, not model quality, is now setting the price.