

1. EVE Online: The Move to Python 3 Begins!

Source: Simon Willison's Weblog | Published: 2026-08-25T22:59:30+00:00

Link: <https://simonwillison.net/2026/Aug/25/eve-online-move-to-python-3/>

EVE Online: The Move to Python 3 Begins! EVE Online has been one of the most interesting case studies in Python at scale for over twenty years now. They've been running on Stackless Python since their launch in 2003, and their last major upgrade was 16 years ago, to Stackless Python 2.7 in 2010 . Their upgrade to Python 3 will start using the futurize script against 2.4 million lines of code, followed by careful manual review of the ~20,000 places where Python 2 and 3 behavior differ - for example `1 / 2` is 0 in Python 2 but is 0.5 in Python 3. There's nothing in this announcement about how they plan to replace Stackless, but at their conference last year they presented Scheduling in Carbon: Leaving Stackless Python Behind describing how they replaced Stackless in the Carbon engine for their more recent game EVE Frontier, using their (now open source) carbonengine/scheduler library. Via Lobster.rs Tags: eve-online , migrations , python , python3 , stackless

2. Debugging Ubiquiti's 5G Backup on AT&T

Source: Jeff Geerling | Published: Tue, 25 Aug 2026 14:18:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/unifi-u5g-backup-debugging/>

For a mobile 'mini homelab' project I'm working on, I wanted to use 5G Internet as a primary 'on-to-go' connection, but still have the ability to plug in another Internet connection when I put the mobile rack I'm building in a fixed location. I'll have more on that project soon, but I figured this project was a good way to check out Ubiquiti's solutions, especially considering many of their smaller gear fits nicely within the dimensions of a mini rack (the 3U DeskPi RackMate TT is pictured above).

3. Foot Guns for Sale

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Tue, 25 Aug 2026 07:15:00 GMT

Link: <https://idiallo.com/blog/foot-gun-for-sale>

I don't think it's going to work out the way everybody thinks it will. The current narrative, at least the one pushed by the companies selling the shovels, is that AI will become centralized. Anthropic and OpenAI will offer safe, vetted AI. Developers will become mere prompt-engineers, submitting requests to these benevolent gatekeepers. They have tamed the dragon. We will benefit only if we become tenants to their API-driven fiefdoms, paying by the token for the privilege of renting intelligence. ?They didn?t care that they?d seen it work in practice because they already knew it couldn?t work in theory? - Clay Shirky I complain about AI frequently on this blog, but not because I don?t think it is useful. I use it quite frequently on my day to day. But what I hate is hype and fake narratives. In fact, I believe that the opposite of their narrative will come true. Technology will continue to improve whether Moore?s Law becomes a relic of the past or not. It?s been called dead, yet CPUs and GPUs are becoming faster than ever. I do not believe for one second that developers with pre-LLM experience will end up on the losing side. And if technology continues to improve, then we won?t n...

4. XCancel, the Twitter/X Mirror, Shuts Down After Cease and Desist From the Fine Folks at X Corp

Source: Daring Fireball | Published: 2026-08-25T22:44:18Z

Link: <https://xcancel.com/>

XCancel: On Monday 24th August at 8PM EST, we received a letter from X Corp. asking to cease and desist the service XCancel. The service XCancel is stopped until further notice. We are seeking legal advice and won?t share more details for now. Thank you for the trust you have put in these two years of XCancel. I didn?t want to jinx them, but this obvious threat hanging over

their head was the main reason I never linked only to XCancel rather than directly to tweets on x.com. This sucks not just because some people don't want to visit x.com, but because XCancel was a faster-loading website with no bullshit. ?

5. Dolly Parton Dies at 80

Source: Daring Fireball | Published: 2026-08-25T22:39:12Z

Link: <https://www.nytimes.com/2026/08/25/arts/music/dolly-parton-dead.html>

Bill Friskics-Warren for The New York Times: Her death was announced in a video posted on social media by her nephew Bryan Seaver. No cause was provided. Concerns about Ms. Parton's health were raised last year as she postponed and later canceled a Las Vegas residency, and intensified to the point that she took to social media in October to announce, "I ain't dead yet." This May, she said that she had responded well to treatment after her immune and digestive systems had, as she put it, "got all out of whack" in recent years. From her hardscrabble beginnings to glittering stardom, Ms. Parton's life was the stuff of make-believe. With her voluptuous figure, sparkly stage costumes and blond bouffant wigs, she was likened by the writer Chet Flippo to the Big Rock Candy Mountain, the Depression-era symbol of abundance. Her songs spoke to millions, confronting privation and loss even as they radiated possibility and hope. She was, quintessentially, a uniter of people. In 2019, The New York Times ran a story with the headline "Is There Anything We Can All Agree On? Yes: Dolly Parton." Hillary Clinton, on Threads : Dolly Parton did all the good she could, for all the people she could, fo...

6. Update Regarding the Base Prices of the M5 Max Mac Studio

Source: Daring Fireball | Published: 2026-08-25T22:17:40Z

Link:

https://daringfireball.net/2026/08/configurations_and_pricing_for_new_mac_minis_and_mac_studios

I made a mistake in my table earlier today regarding the base price of the M5 Max Mac Studio , but my mistake was because of a rather sneaky design in Apple's purchasing flow. Apple's web interface looks like this: Both of the available M5 Max SoCs have 18-core CPUs, but the entry model (that starts at \$2,500) has a 32-core GPU, and the upgraded chip has a 40-core GPU. That upgrade is labeled "+ \$300", which makes it look as though the starting price for the 18/40-core model is \$2,800. That's the price I put in the original version of my chart. But if you select that option, you'll notice that the actual starting price jumps from \$2,500 to \$3,100 ? a \$600 difference, not \$300. The reason is that the \$2500 18/32-core version only comes with one option for RAM: 32 GB. The 18/40-core chip has three tiers for RAM: 48, 64, and 128 GB. So when you switch from the 18/32-core chip to the 18/40-core chip, you don't just get 8 extra GPU cores (which Apple is saying costs \$300), but the base RAM jumps from 32 to 48 GB, and that costs you an additional \$300. Apple does show a blue notification stripe when you change chips here, which states "Based on your new selection, other parts of your co..."

7. Theatre Review: Cats at Regent's Park Open Air Theatre ?????

Source: Terence Eden's Blog | Published: Tue, 25 Aug 2026 11:34:22 +0000

Link: <https://shkspr.mobi/blog/2026/08/theatre-review-cats-at-regents-park-open-air-theatre/>

Cats is so silly! In the olden days, after a wild animal was slaughtered, the tribe's shaman would wear the skin and re-enact the hunt. As he became one with the beast, the people slowly understood the ways of nature. Cats is a similarly spiritual experience wherein we watch androgynous sylphs gyrate across the stage and believe (if only for a moment) that the human has become feline, all in?

8. Uranium from Píbram (Czechia)

Source: Maurycy's Blog | Published: Tue, 25 Aug 2026 00:00:00 +0000

Link: <https://maurycyz.com/misc/rocks3/>

All images are links to high quality versions. RC-102: 1.7 kCPS [43 uSv/h]. AB+: 55 CPS, 1960 CPS Uraninite nodules in calcite (?). Unlike the Colorado plateau rocks, the veins are nice and clean (although not chemically pure) which results in a shiny surface. RC-102: 1.1 kCPS [37

uSv/h] Chunk of Uraninite growing on calcite. It has nice botryoidal formations of massive Uraninite along with some areas where the shape of the underlying calcite crystal is visible. A few of the "bubbles" are broken, showing the individual crystals that grew on the matrix. RC-102: 300 kCPS [14 uSv/h]. AB+: 25 CPS, 704 CPS Ugly Uraninite chip. Might glue this inside a cloud chamber or something: it's small, crumbly and weathered. RC-102: 3.2 kCPS [112 uSv/h] Uraninite coating "granite". RC-102: 1.3 kCPS [32 uSv/h]. AB+: 26 CPS, 1200 CPS Granite with Uraninite veins running through it. A nice crystal of calcite is visible in the third photo. Uraninite veins running through calcite (?). This might look cool if cut and polished. RC-102: 732 CPS [15 uSv/h]. AB+: 1.75 CPS, 240 CPS This piece has a (somewhat metamorphic) sedimentary host rock, which is very different from the other pieces. Geology notes : Th...

9. Hardening the Override Flag

Source: Andrew Nesbitt | Published: 2026-08-25T10:00:00+00:00

Link: <https://nesbitt.io/2026/08/25/hardening-the-override-flag.html>

I've been reading through the recent run of incidents where package-manager infrastructure was the attack surface, as one does on a Sunday afternoon. One question I kept coming back to was whether the tools have any defences at the flag and config level. Every package manager has an override that turns off a check: `--allow-unauthenticated` , `--break-system-packages` , `--ignore-scripts` , an env var that points the resolver at a different registry. If a compromised install script can pass those as easily as a person at a keyboard can, the check does very little. So I went looking at how command-line tools more generally handle their dangerous overrides, `sudo` and `curl` and `rm` and the rest, and the catalogue got long enough to write up on its own. Names and channels GNU `rm` goes back and re-reads `argv` after option parsing has finished. `getopt_long` accepts any unambiguous prefix of a long option, which would make `--no-p` a valid spelling of `--no-preserve-root` , so `rm` checks the literal string and rejects anything shorter with `?you may not abbreviate the ?no-preserve-root option?. --preserve-root` has been the default since 2006 , and the override requires the full flag every time, on top of...

10. Adding diagrams to my static site generator with D2

Source: Giles' blog | Published: Tue, 25 Aug 2026 23:00:00 +0000

Link: <https://www.gilesthomas.com/2026/08/adding-d2>

A lot of the time when I've been writing posts for this blog, I've felt that a diagram would really help. But they're a pain to produce well, and I think I underuse them as a result. I wanted to fix that, and wound up adding D2 support to my static site generator. I think it works pretty well! In the past, I've tried drawing my own diagrams in LibreOffice and exporting as SVG, but my complete lack of artistic skill doesn't help: Asking an AI to do it for me helped in simple cases: ...but with something less standard (there must be a million neural network diagrams in their training sets) it can be really fiddly to get something right. I did some investigations into the various diagram-generating tools out there, and decided to give D2 a go. It has a simple language for specifying what your diagram should show, and the output is pretty nice: Here's the source for that diagram: `style.fill: transparent tokens: Tokens tokens -> llm.token-embeddings llm: "" { token-embeddings: Token embeddings position-embeddings: Position embeddings plus: "+" { shape: circle width: 36 height: 36 style.font-size: 24 style.fill: transparent } token-embeddings -> plus position-embeddings -> plus input-em...`

11. The AI Hater's Manifesto

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 25 Aug 2026 16:57:38 GMT

Link: <https://www.wheresyoured.at/the-ai-haters-manifesto/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year , \$18 a quarter , or \$7 a month , and in return you get a weekly newsletter that's usually anywhere from 10,000 to 18,000 words, including vast, detailed analyses of

12. A curmudgeon tries a language server

Source: Entropic Thoughts | Published: Tue, 25 Aug 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/curmudgeon-tries-language-server>

I write code roughly the same way I did ten years ago: I make edits in my text editor, I switch over to a terminal window and I run a command to compile and execute the code. I look at the result, then switch back to my editor. If I am unsure of what happens, I either add tracer prints to the code and restart the program, or I restart the program in a debugger. If I hit an exception that crashes my program, I fix the problem and then restart it. I'm envious of what Lisp programmers do. Lisp development typically happens by typing code directly into the REPL to add, remove, and replace parts of the live, running system. A Lisp programmer does not need to switch over to something else because they are already inside their program process. They never compile and execute the code because the code is already running and they edit it by hot-swapping code. They never restart in a debugger because they are already inside their program process and can inspect anything they want. They don't have to restart the program on exceptions, because the condition system allows resuming the crashing code from anywhere in the stack after the system has been patched with a fix. (Continue reading th...

13. Have You Heard the Good News About Microlighter?

Source: Jim Nielsen's Blog | Published: Tue, 25 Aug 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/shipping-microlighter/>

Dave Rupert wrote about shipping microlighter : a tool for handling syntax highlighting using the CSS Custom Highlights API . I saw his post the day he released it, and I had an implementation PR up for my blog by end of day. Then, like I do with so many things, I let it sit there. This is the period where my subconscious takes over. It does the work of, "How do I actually feel about that? Do I want to merge it? Do I have any regrets about what I did?? If I still want to merge it after a few days, that's usually a good sign that I'll be happy with the work. (Sometimes after a few days I say, "What the hell was I thinking?? and then it's easy to simply close the PR with zero regrets.) Well it's a few days later and I still feel good about it, so time to ship! My PR for this is pretty straightforward: Remove highlight.js dependency (and related plumbing) On paths that 1) match my post pages (i.e. /YYYY/:slug), and 2) have code on them, pull microlighter deps from a CDN and run it. Done. Granted, there are trade-offs to this approach. I get it. Dave's explainer for this tool on The ShopTalk Show vibed with me because I've been in his shoes many times: "Whoops, somehow syntax highlig...

14. The Negative Option

Source: Tedium: The Dull Side of the Internet. | Published: 2026-08-25T02:45:00Z

Link: <https://feed.tedium.co/link/15204/17427939/columbia-house-history-retrospective>

In the 1980s and 1990s, Columbia House could do no wrong?it was a great way to get 12 CDs for a penny. But its business model kind of set the stage for a lot of bad e-commerce.

15. Dylan Patel ? Anthropic & OpenAI will have most of the world's compute by 2028

Source: Dwarkesh Podcast | Published: Tue, 25 Aug 2026 15:32:57 GMT

Link: <https://www.dwarkesh.com/p/dylan-patel-3>

"Every force is screeching towards centralization."

16. A Cautionary Tale About Data Breach Claims, Verification and Carhartt

Source: Troy Hunt | Published: Tue, 25 Aug 2026 21:51:08 GMT

Link:

<https://www.troyhunt.com/a-cautionary-tale-about-data-breach-claims-verification-and-carhartt/>

You're not going to believe this, but turns out you can't always take criminals at their word.

Actually, I'll walk that back a bit as it may not even be the cybercrime guys who got this wrong, but it all starts here: "Cyber

17. Actually good advice

Source: Experimental History | Published: Tue, 25 Aug 2026 14:28:03 GMT

Link: <https://www.experimental-history.com/p/actually-good-advice>

don't tell someone they look like Beaker from The Muppets, etc.

18. Extending life perception

Source: Herman's blog | Published: 2026-08-25T11:59:44.765013+00:00

Link: <https://herman.bearblog.dev/extending-life-perception/>

I write about journaling a lot, since it's honestly my most important practice. In my most recent post on the subject I made the case that active recall is an effective way to remember things, whether that be ideas, thoughts, or events. I'd like to expand on that post and make the case that journaling extends life perception. To wax philosophical for a moment: long life perception and long lifespan are, when it comes down to it, effectively the same thing. If we live in the infinite now, then the only thing separating an actual long life from the perception of a long life is perception. This intuitively makes sense, and you've likely experienced this yourself. Years that are filled with novelty tend to feel longer and are more memorable than years of consistently similar days. It's partly why time feels so much longer as a child, since all experiences are fairly novel and there's explicit progression through the different phases of school, which stands in stark contrast to, say, a few years of working the same job and rarely taking vacation or trying anything new. Those few years can easily condense into a handful of memories, and so the perception of those years is fairly short...

19. Forgejo hack: How to set a starting issue and pull request number

Source: Miguel Grinberg's Blog | Published: Tue, 25 Aug 2026 13:38:15 GMT

Link: <https://blog.miguelgrinberg.com/post/forgejo-hack-how-to-set-a-starting-issue-and-pull-request-number>

I'm currently working on migrating my open source projects from GitHub to a self-hosted Forgejo instance. As part of this effort I often end up looking through the Forgejo source code to figure out if there are hidden ways to configure certain things to my liking when I can't do it on the administration UI. I thought I'd start putting my discoveries in writing here, in case they can help others. So here goes the first one. One of the aspects of the migration that is tricky is how to transition issues and pull requests. What makes the most sense to me is to only use Forgejo to track issues and pull requests going forward, leaving all the issues and pull requests created up to the migration point on GitHub. Of course whether this is a good or bad idea is debatable, but considering all the options I have decided that this is the solution that is going to inflict the least pain on me. The one problem with this approach is that I would end up having duplicate issue numbers, because Forgejo would start creating issues and pull requests all the way back from #1, and all those low numbers have been used on the GitHub side. So I wanted to hack my Forgejo instance so that issues start from...