

1. Just a rumour of a bug is enough to find a security exploit these days

Source: Simon Willison's Weblog | Published: 2026-08-28T22:12:02+00:00

Link: <https://simonwillison.net/2026/Aug/28/just-a-rumour-of-a-bug/>

Just a rumour of a bug is enough to find a security exploit these days Anil Madhavapeddy is a professor of computer science at Cambridge and a core maintainer of the OCaml compiler. In this somewhat alarming post he reports that security issues in OCaml projects are seeing evidence of attempted exploits within minutes of patches being shared for discussion: This normally takes a few days and a release within a week or two is reasonable. Within about ten minutes (!) this website was fielding probes for percent-encoded traversal sequences, indicating that automated watchers are keeping an eye on public repositories. Modern coding agents have become so effective at finding flaws that the slightest hint at a new bug can be enough information for them to find it, something Anil has been able to demonstrate using his own agents, switching to DeepSeek V4 Pro? when Claude Fable refused the task. Anil points out that this rate of discovery appears incompatible with existing open source embargo practices for new issues. If an issue can become an exploit this fast, we need to figure out new processes for keeping our communities safe. rclone maintainer Nick Craig-Wood confirms in the Hacker N...

2. Building a mini Homelab that fits in my carry-on

Source: Jeff Geerling | Published: Fri, 28 Aug 2026 10:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/mini-homelab-network-fits-in-carry-on/>

I'm traveling to Chicago for VCF Midwest next month. I'll be demoing NTP time history on vintage Macs, with my own GPS-derived NTP service hosted on an Xserve G5, synced via NTP or a strange AppleTalk timing extension from the 1990s . So I built a little 'portable homelab' (pictured above) that supports 1-10 Gbps networking, can run off a small battery for at least an hour, switches between multiple WANs (so I can get my own 5G Internet connection, in case I need it), and gives me 12 wired Ethernet connections.

3. On forcing all derived classes to implement a specific non-virtual method, part 2

Source: The Old New Thing | Published: Fri, 28 Aug 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260828-00/?p=112654>

Explicitly denying that you implement the method. The post On forcing all derived classes to implement a specific non-virtual method, part 2 appeared first on The Old New Thing .

4. Making the unnecessary easier

Source: John D. Cook | Published: Fri, 28 Aug 2026 14:35:36 +0000

Link: <https://www.johndcook.com/blog/2026/08/28/making-the-unnecessary-easier/>

I watched a few videos this morning, looking for ideas of what I could use AI to do. In one video, someone had an agent monitor tech news sites every 30 minutes to notify him of a variety of developments. No doubt that?s less effort than visiting a bunch of sites every half hour, but [?] The post Making the unnecessary easier first appeared on John D. Cook .

5. Premium: The Hater's Guide To Circular Financing (Part One)

Source: Ed Zitron's Where's Your Ed At | Published: Fri, 28 Aug 2026 15:55:36 GMT

Link: <https://www.wheresyoured.at/premium-the-haters-guide-to-circular-financing-part-one/>

[NVIDIA Company Meeting, the present day, YMCA playing] JENSEN HUANG : We love NVIDIA, don?t we folks? We?re the biggest, most-beautiful semiconductor company, we make the biggest, hottest GPUs for Clammy Sammy and Wario Amodei ?s huge, beautiful AI labs, but they can?t

6. This Week on The Analog Antiquarian

Source: The Digital Antiquarian | Published: Fri, 28 Aug 2026 16:14:15 +0000

Link: <https://www.filfre.net/2026/08/this-week-on-the-analog-antiquarian/>

Rehabilitating Richard

7. What GLM-5.3 Flash running on Chinese hardware actually means

Source: Martin Alderson | Published: Sat, 29 Aug 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/glm-5-3-flash-chinese-hardware/?utm_source=rss&utm_medium=rss&utm_campaign=feed

Z.AI ran a full model release on domestically manufactured Chinese chips. Impressive, but EUV lithography is a hard wall and the gap with Western inference hardware is more likely to widen than close.