

1. You have to beat the models at something

Source: seangoedecke.com RSS feed | Published: Sun, 30 Aug 2026 00:00:00 GMT

Link: <https://seangoedecke.com/you-have-to-beat-the-models-at-something/>

In 2025, I wrote that software engineers ought to be assessed by ?value over replacement? : not how much money they made for their company, but how much they would have made compared to the average engineer in their position. I?ve always found it vaguely silly when engineers put ?built a product that made \$X? on their resumes, when they just did the JIRA tickets that came across their desk. Today, value over replacement is even more important. A replacement-level engineer in the 2010s was fine : maybe not worth promoting, but still worth paying , because writing code had a high fixed cost. Now writing code costs a hundred bucks a month . What are you doing that GPT-5.6-Sol or Claude Opus 5 wouldn?t do in your position? Why is it worth paying an extra two or three orders of magnitude for? This is a scary thought. But you?re not doing yourself any favors by pretending that LLMs can?t actually write code and it?s all just a scam, or that LLM-written code is inherently so bad as to cause companies using it to collapse next year. We are not going to wake up in 2027 to find that the AI craze is over and everyone is writing code by hand again. You ought to put some serious thought into w...

2. Introducing Hy4 Preview

Source: Simon Willison's Weblog | Published: 2026-08-29T23:53:13+00:00

Link: <https://simonwillison.net/2026/Aug/29/hy4/>

Introducing Hy4 Preview New open weight text input (no vision) LLM from Chinese company Tencent today: 770B total parameters, 49B active parameters, 1M token context window, 1.56TB on Hugging Face . This is a big size increase from their previous Hy3 in July, which was 295B, 21B active, 256,000 context, 598GB. I recently started using model chat templates to better understand their capabilities. Here's Hy4's chat_template.jinja on Hugging Face, which includes this section: {% - if not reasoning_effort is defined %} {% - set reasoning_effort = 'high' %} {% - elif reasoning_effort not in ['high' , 'no_think'] %} {% - if reasoning_effort is none %} {{- raise_exception('reasoning_effort error : None, should be no_think/high') }} {% - else %} {{- raise_exception('reasoning_effort error : ' + reasoning_effort + ', should be no_think/high') }} {% - endif %} {% - endif %} So it looks like there are just two reasoning effort levels: "high" (the default) and "no_think" (reason by disabled). I tried my "Generate an SVG of a pelican riding a bicycle" prompt with the default high reasoning via OpenRouter and got this : Quoting the reasoning trace: [...] Let's maybe add a helmet? It could imp...

3. ? Thoughts and Observations on Apple?s First Immersive MLB Broadcast, a Yankees 1-0 Win Over the Red Sox

Source: Daring Fireball | Published: 2026-08-29T17:53:34Z

Link:

https://daringfireball.net/2026/08/thoughts_and_observations_apple_immersive_mlb_broadcast

It felt like something new and different. Something at least as different from watching on TV as watching on TV is different from merely listening on radio. It?s profound.

4. Apple Announces Price Increase for Apple TV and Apple One Subscriptions

Source: Daring Fireball | Published: 2026-08-28T21:52:37Z

Link: <https://9to5mac.com/2026/08/28/apple-announces-price-increase-for-apple-tv-and-apple-one-subscriptions/>

Chance Miller, 9to5Mac: Here are the full details on the price increases: Apple TV (monthly): \$14.99 (up from \$12.99) Apple TV (annual): \$119 (up from \$99) Apple One Individual: \$21.95 (up from \$19.95) Prices for other Apple One plans are unchanged. When Apple increased the price of Apple Music last month, it also raised the Apple One Family and Apple One Premier plans, but not the Individual plan. The higher prices go into effect today for new subscribers. Apple says existing subscribers will be notified of the price increase about a month before they are

charged the new price. Seems poetic that Tim Cook's last action as CEO is a goose to services revenue. ?

5. ?I?m the Guy Who Destroys Antique Books After We Scan Them Into Our Company's Insatiable AI Platform?

Source: Daring Fireball | Published: 2026-08-28T21:35:22Z

Link: <https://www.mcsweeney.net/articles/im-the-guy-who-destroys-antique-books-after-we-scan-them-into-our-companys-insatiable-ai-platform>

Jack Loftus, writing at McSweeney's: People assume scanning is the exciting part. And, sure, it is impressive. They shear off the spine, separate the pages, and run everything through these enormous machines so the computer can absorb every sentence ever written and use it to create some fairly compelling LinkedIn posts. But then, boy oh boy, the wheeled cart carrying those decimated remains comes to me. That is where the human touch enters the process, assuming you count industrial shredders and a giant burn pit out back as an extension of the hand, which I certainly do. (Via Michael Sippey .) ?

6. Technical note: transfer files over an ethernet patch cable

Source: Maurycy's Blog | Published: Sat, 29 Aug 2026 00:00:00 +0000

Link: <https://maurycyz.com/misc/etherfiles/>

It's possible to just connect two computers together with Ethernet and do a bit of IP configuration: # On sender ... ip address add dev eth0 fd42:dead:beef::1/48 ip link set dev eth0 up # On receiver ... ip address add dev eth0 fd42:dead:beef::2/48 ip link set dev eth0 up After a few seconds, pings should work: # On receiver ... / # ping fd42:dead:beef::1 64 bytes from fd42:dead:beef::1: icmp_seq=1 ttl=64 time=0.649 ms 64 bytes from fd42:dead:beef::1: icmp_seq=2 ttl=64 time=0.376 ms 64 bytes from fd42:dead:beef::1: icmp_seq=3 ttl=64 time=0.414 ms 64 bytes from fd42:dead:beef::1: icmp_seq=4 ttl=64 time=0.340 ms ... and so should this: # On receiver ... socat - TCP6-LISTEN: 1234 | dd status = progress > big_file.tar.gz # On sender ... socat - 'TCP6-CONNECT: [fd42:dead:beef::2] : 1234 ' big_file.tar.gz These commands assume you are using Linux, but this trick works everywhere. A nothing-special patch cable and ethernet jack should have no problem hitting ~900 Mbits/second, which is 6.7 GB per minute. Fancy network cards will allow speeds orders of magnitude faster, but it's already much faster than USB flash or cloud storage. Other options ... for transferring a file larger than 10 G...

7. A simple "copy this code" button in JavaScript

Source: Terence Eden's Blog | Published: Sat, 29 Aug 2026 11:34:49 +0000

Link: <https://shkspr.mobi/blog/2026/08/a-simple-copy-this-code-button-in-javascript/>

Next to all the code samples on this blog is a little "copy" button. That makes it easier to grab any of the code I've shared. The HTML and JS is delightfully simple: ? The navigator.clipboard.writeText needs a user interaction to w?

8. This Week in Package Management: 29 August 2026

Source: Andrew Nesbitt | Published: 2026-08-29T10:00:00+00:00

Link: <https://nesbitt.io/2026/08/29/this-week-in-package-management.html>

Week fifteen of the roundup, built from the package manager OPML feed collection and whatever I've posted or boosted on Mastodon . pnpm 12 pnpm 12.0 is the stable release of the Rust rewrite. The commands, flags, settings and lockfile format from pnpm 11 carry over unchanged; a companion what's different post lists the seven behavioural changes. Git dependencies on GitHub, GitLab and Bitbucket now resolve through the host's canonical HTTPS URL regardless of how they were specified, so lockfiles record only the canonical HTTPS URL for those hosts. Dependency cycles are broken canonically during peer resolution, so the lockfile is a pure function of the dependency graph. Cycle-heavy workspaces resolve peers 2-3x faster with roughly 25% less memory. Globally installed tools resolve to the version pinned in the current project, and pnpm can provision npm, Yarn and Bun as managed tools with a pnx command for one-off runs. Unrecognised keys in pnpm-workspace.yaml are now errors. On the 11.x branch, 11.23 reworks the registries setting so an Artifactory or GitLab registry can keep its tarball URLs out of

pnpm-lock.yaml , and 11.24 restores pnpm approve-builds --global and groups recursiv...

9. The Server Called Paranoia: Defend Autistici/Inventati

Source: micahflee | Published: Sat, 29 Aug 2026 19:40:15 GMT

Link: <https://micahflee.com/the-server-called-paranoia-defend-autistici-inventati/>

For twenty-five years, an Italian hacker collective has built communications infrastructure designed to survive censorship, surveillance and police raids. On August 26, the United States designated it a terrorist organization. On September 25, the wind-down period ends.

10. Reading List 08/29/26

Source: Construction Physics | Published: Sat, 29 Aug 2026 12:03:07 GMT

Link: <https://www.construction-physics.com/p/reading-list-082926>

Sinking London homes, US magnesium manufacturing, the World Humanoid Robot games, catastrophic flooding in Nepal, and more.

11. The Rise and Fall of Agent Civilizations

Source: Dwarkesh Podcast | Published: Sat, 29 Aug 2026 22:47:53 GMT

Link: <https://www.dwarkesh.com/p/openai-huggingface>

The whole OpenAI/Hugging Face story in plain English