

1. FBI Probes Service Selling 153M+ Drivers Licenses

Source: Krebs on Security | Published: Tue, 01 Sep 2026 22:40:28 +0000

Link: <https://krebsonsecurity.com/2026/09/fbi-probes-service-selling-153m-drivers-licenses/>

A new identity theft service launched on the dark web this week is selling digital scans of more than 153 million drivers licenses from people in the United States and Canada. Based on interviews with individuals whose licenses are available for purchase on this service, it appears to be siphoning images collected by a widely-used identity verification company based in Louisiana. KrebsOnSecurity also has learned that the New Orleans field office of the Federal Bureau of Investigation (FBI) today launched an official inquiry into the source of the images.

2. Claude Fable 5.1 made me a really nice animated pelican

Source: Simon Willison's Weblog | Published: 2026-09-01T23:57:28+00:00

Link: <https://simonwillison.net/2026/Sep/1/claude-fable-5-1/>

Today is Claude Fable (and Mythos) 5.1 day . Anthropic say that Fable 5.1 "sets a new standard for coding, knowledge work, and long-running problem-solving tasks". Their announcement spends a notable amount of time on scientific research, boasting of a 52.6% score on the brand new Terminal-Bench-Science 0.1 benchmark (first announced on August 27th), up from 24.7% for Fable 5, 29.0% for Opus 5 and 22.4% for GPT-5.6 Sol. Other benchmarks show slightly improved scores, but none as impressive as the Science one. But how well can it pelican? Back in July I wrote about how I was losing faith in the pelican benchmark - its connection to how good the models were at other tasks didn't seem to hold as strongly as it did back in 2025 . The most interesting insights I get from it now are comparisons within model families, and particularly comparisons for the same prompt at different reasoning effort levels. Fable 5.1 has five reasoning levels: low, medium, high, xhigh, max - and no option to turn off reasoning entirely. I fixed an issue in llm-anthropic which caused reasoning traces not to be correctly recorded, then ran some prompts. Here's the full set of pelicans for all of the reasoning...

3. Codex bundles LibreOffice

Source: Simon Willison's Weblog | Published: 2026-09-01T19:03:01+00:00

Link: <https://simonwillison.net/2026/Sep/1/codex-libreoffice/>

I was poking around in my ~/.cache/ folder using OmniDiskSweeper when I spotted something interesting. The OpenAI Codex desktop app (since rebranded to just ChatGPT) has 1.7GB of stuff in there in a folder called codex-primary-runtime , including a full Python installation, a full Node.js installation, and native binaries for Poppler , git, and the LibreOffice open source office suite (which forked from OpenOffice.org in 2010): The ~/.cache/codex-runtimes/codex-primary-runtime/plugins/openai-primary-runtime/plugins/documents folder includes skills which tell Codex how to find and use those binaries. Tags: codex , generative-ai , openai , ai , llms , openoffice , open-source

4. GeoJSON Map Viewer

Source: Simon Willison's Weblog | Published: 2026-09-01T18:05:45+00:00

Link: <https://simonwillison.net/2026/Sep/1/geojson/>

Tool: GeoJSON Map Viewer I was helping Natalie gather some maps of local political boundaries (for the Granada Community Services District and the Midcoast Community Council and found a need to display some GeoJSON files on a map and export that as a PNG. I asked GPT-5.6-Sol for suggestions of tools and it proactively built one. After some iterations using Claude Code for web and Fable 5.1 we got to this finished tool. As for the GeoJSON.. it turns out if you ask ChatGPT Work to provide boundaries for almost anything it will churn away extracting and combining files from different Government data sources and build exactly what you need. I got this polygon from: I want a polygon that represents the exact boundary of the El Granada GCSD And this one from: Get me a GeoJSON file for the boundary (or boundaries if that makes sense) for the MCC - Midcoast Community Council - that operates near Half Moon Bay CA Here's a link

that displays both of them at the same time on the new GeoJSON map viewing tool. Tags: geospatial , tools , geojson , chatgpt

5. Tim Cook's Departure Memo on His Last Day as CEO

Source: Daring Fireball | Published: 2026-09-01T18:11:57Z

Link: <https://9to5mac.com/2026/08/31/read-tim-cooks-full-memo-to-apple-employees-on-his-last-day-as-ceo/>

Tim Cook: There is something truly special about Apple. I am most proud of what an annual report could never capture. This place is proof that culture triumphs over everything. We share a belief that what we build matters, and that we have both the opportunity and the responsibility to leave the world better than we found it. That purpose is part of what makes this place extraordinary. Apple helps nurture it, but I believe it lived within each of you long before you arrived here. It is what brought you to this company and what continues to drive the work you do every day. Together, we have created something far greater than any one of us could have imagined or accomplished alone. And that's the secret to our success. We bring out the best in each other. We lift each other up. We have made it possible to leave our "dent in the universe," as Steve once described it, because of who we are and what we believe, because of what we value and how we see the world. How fortunate we are. How fortunate I am. That's the best paragraph of the memo. I like those last two "fortunate" lines, and I think they're the most open Cook gets. But on the whole the memo reads to me like the classic Far Si...

6. Apple Reveals Forensic Evidence From Chang Liu's MacBook in OpenAI Lawsuit

Source: Daring Fireball | Published: 2026-09-01T17:36:08Z

Link: <https://9to5mac.com/2026/08/31/apple-openai-forensic-macbook-evidence/>

Chance Miller, 9to5Mac: In today's filing, Apple says that Liu's lawyers recently handed over his MacBook that he used after leaving Apple. According to Apple, an "initial forensic analysis" of that MacBook revealed four things, quoted from Apple's filing: Mr. Liu not only downloaded a confidential Apple circuit schematic but also used it in his work at OpenAI; Far from his unauthorized access to Apple's third-party cloud storage being unknown to him, Mr. Liu and others at OpenAI were well-aware of that access; Mr. Liu, upon learning of Apple's internal investigation of him, sent instructions for destroying evidence to an OpenAI colleague who confirmed she would comply; and Mr. Liu used a tool in his work at OpenAI that has the same name as an internal Apple engineering application used for Apple development work. 9to5Mac is hosting a copy of the original filing from Apple's lawyers. The first point seems really bad for OpenAI. On the fourth point, I'm not sure what Apple is alleging here ? I think they're alleging that Liu took an Apple internal engineering app with him over to OpenAI? This is the entirety of that allegation from Apple's filing: D. New evidence shows previously u...

7. [Sponsor] WorkOS: How to Give an Agent a Task Instead of a Token

Source: Daring Fireball | Published: 2026-09-01T07:59:07Z

Link: https://workos.com/blog/delegated-access-for-ai-agents?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q32026

Give an agent an access token and it spreads: into the context window, into tool call logs, into notes it keeps between steps. Each copy works from anywhere, long after the fact. Relay keeps the credential at WorkOS. Your agent names the user, WorkOS attaches that token, refreshes it, and releases it only to allowlisted hosts. A hijacked agent session is a live process you can kill. Learn how it works ? ?

8. Microspeak: Funded / unfunded

Source: The Old New Thing | Published: Tue, 01 Sep 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260901-00/?p=112662>

People are our most valuable asset. The post Microspeak: Funded / unfunded appeared first on The Old New Thing .

9. Git Submodules as a Package Manager

Source: Andrew Nesbitt | Published: 2026-09-01T09:00:00+00:00

Link: <https://nesbitt.io/2026/09/01/git-submodules-as-a-package-manager.html>

I added a worktree to a repository last week to try a branch alongside the main checkout, ran `git submodule update --init` in it because the build needed the vendored dependencies, and when I was done went to clean up with `git worktree remove`, which git refused. Per the man page only clean worktrees can be removed, and "unclean worktrees or ones with submodules" need `--force`. Submodules get their own clause in that sentence, distinct from dirty state. `git worktree move` is stricter again and refuses outright on any worktree containing submodules. I'd spent the previous week cataloguing how command-line tools harden their `--force` flags, and here was git requiring one because two of its own features had collided. GitHub's git 2.5 announcement introduced `git worktree` in July 2015 with a one-line caveat: "It's not recommended to use `git worktree` with a repository that contains submodules." Eleven years later git still requires `--force` to remove a worktree that has submodules and refuses to move one. In between, `worktree add` had to be patched to ignore `submodule.recurse` because honouring it made the internal `reset --hard` recurse into submodule paths that were still empty in the fresh w...

10. Book Review: ActivityPub by Evan Prodromou ?????

Source: Terence Eden's Blog | Published: Tue, 01 Sep 2026 11:34:18 +0000

Link: <https://shkspr.mobi/blog/2026/09/book-review-activitypub-by-evan-prodromou/>

As part of my grant from NLnet to improve my Fediverse bot project, I'm spending time making sure I understand all the fundamentals of ActivityPub. The problem with Internet standards is they're scattered all over the place. Abandoned forums, half-maintained wikis, mailing lists with obscure rules, and monolithic blocks of text. Sometimes you just want one book which collates all the info and?

11. New Post: A Crash Course in Predicate Logic

Source: Computer Things | Published: Tue, 01 Sep 2026 18:30:01 +0000

Link: <https://buttondown.com/hillelwayne/archive/new-post-a-crash-course-in-predicate-logic/>

Logic for Programmers has now been out for a month! To celebrate, I'm releasing the entire second chapter, "A Crash Course in Logic", for free on my blog. Read it here!

12. Hyperscale Normalization

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 01 Sep 2026 17:39:44 GMT

Link: <https://www.wheresyoured.at/hyperscale-normalization/>

If you liked this piece, you should subscribe to my premium newsletter. It's \$70 a year, \$18 a quarter, or \$7 a month, and in return you get a weekly newsletter that's usually anywhere from 10,000 to 18,000 words, including vast, detailed analyses of

13. Nobody Believes It, Everybody Does It

Source: Jim Nielsen's Blog | Published: Tue, 01 Sep 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/nobody-believes-it-everybody-does-it/>

Speaking of things we know individually, but let slide collectively, here's Jeremy Keith talking about how we all used Flash: Remember when almost every website had a Flash intro? Everyone knew they were annoying and unnecessary, but everyone else was doing it. And here's John Gruber writing about his experience with those marketing schemes that trade your personal information for a discount: The marketing shitbirds who press for these schemes [?] do so by pointing to data that shows that they do convert some number of users. "It works" they claim, pointing to data. What doesn't show up in their data are interactions like mine. They don't have analytics that measure that I now consider their website an antagonist to avoid at all costs. There's a funny thing about us humans: we place trust in systems that our own individual experience tells us are flawed. For example: KPIs: we know metrics get optimized at the expense of reality, but leadership treats dashboards as a reflection of truth. Performance reviews: we know they're political and subjective, yet companies leverage them for promotions, compensation, even firings. Standardized testing: we know tests don't measure real abili...

14. Ajeya Cotra ? Inside the OpenAI agent swarm that hacked Hugging Face

Source: Dwarkesh Podcast | Published: Tue, 01 Sep 2026 15:41:06 GMT

Link: <https://www.dwarkesh.com/p/ajeya-cotra>

"This might be the clearest warning shot we ever get."

15. Claude comment detection

Source: Entropic Thoughts | Published: Tue, 01 Sep 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/ai-comment-classifier>

Here's a comment I read in the code I was working on. Walk the entire collection to locate flagged clusters. Skip it altogether when the cached, dirty-tracked flag count says there are none (the usual case): nothing needs marking, so the walk is pure waste. When the count is stale (e.g. right after an edit) or non-zero, we fall through to the real lookup, so stale marks never show. Same eventually-consistent signal the ?Flags? badge relies on. One thing in this comment caught my mind: the small parenthetical that says ?the usual case?. If that were true, it would be important information! However, this comment is generated by the LLM model Claude, which has no sense of what counts as the usual case in this domain. But by generating that comment, it tricks every future reader (both human and robot) into believing in a property of the system that does not exist. (Continue reading the full article on the web.)

16. Weekly Update 519: Breaches & Data Integrity

Source: Troy Hunt | Published: Tue, 01 Sep 2026 22:26:38 GMT

Link: <https://www.troyhunt.com/weekly-update-519/>

It does feel like I've bitten off too much and am now chewing like crazy this week. The 3D printing talk with Elle in Oslo, the "normal" NDC infosec talk, the cyber-broken talk with Scott in Copenhagen and then those ratbag hackers keep dumping more