

1. Research acceleration: The view inside OpenAI

Source: Simon Willison's Weblog | Published: 2026-09-06T23:57:40+00:00

Link: <https://simonwillison.net/2026/Sep/6/research-acceleration-the-view-inside-openai/>

Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I think it's their new AGI. Both this piece and the new essay An Alien Mind (by Chief Scientist Jakub Pachocki) talk about it, and this one doesn't even bother to expand the acronym. Included are details on how OpenAI's own research team are using coding agents. Like pretty much everyone else 2026 has been the year that agentic engineering really took off at OpenAI, best illustrated by this chart: I'm intrigued at what caused that significant acceleration in AI spend per researcher in late July - my best guess is that's when internal employees gained access to the model later released as GPT-6 Astra. Tags: ai , openai , generative-ai , chatgpt , llms , coding-agents , november-2025-inflection , recursive-self-improvement

2. The purpose of DNS is to spread scams

Source: Simon Willison's Weblog | Published: 2026-09-06T14:40:07+00:00

Link: <https://simonwillison.net/2026/Sep/6/the-purpose-of-dns-is-to-spread-scams/>

The purpose of DNS is to spread scams Terence Eden shares some daunting statistics in support of his take that "the Domain Name System's purpose seems to be a vector for criminals to run scams on people at a terrifyingly high rate". On this Interisle report (via Andrew Campling), Terence says: It says 85 million new registrations of gTLDs were made in 2025. Of those 8.5 million were added to blocklists by May 2025. It reckons that a 10% abuse rate is the likely floor for these numbers and it's probably closer to 20%. One in five newly registered domains with a gTLD are scams. That's a bloody crisis. I had no idea. Apparently ICANN have been discussing this problem for years. Tags: dns , scams , terence-eden

3. There's No Limit to How Bad Code Can Get

Source: Simon Willison's Weblog | Published: 2026-09-06T09:08:06+00:00

Link: <https://simonwillison.net/2026/Sep/6/theres-no-limit-to-how-bad-code-can-get/>

My comment on There's No Limit to How Bad Code Can Get ? Lobste.rs. [In reply to a comment about burning it down to start from scratch when technical debt becomes overwhelming] In my experience it's so rare for that to work. You announce the old thing is irrecoverably drowning in tech debt. You spin up a team to rewrite it from scratch. Work begins. Meanwhile the old thing remains a moving target: it's running the core business, so changes are still necessary. The developers working on it know that it's going to be made obsolete by the new thing soon, so they don't have any incentive to go beyond the smallest effort possible to add the new features. Technical debt continues to mount. Meanwhile, the team working on the new thing are ambitious and probably a little naive. They start out at a great pace - it's greenfield after all - but as time progresses it becomes apparent that nobody fully understands the behavior and scope of the thing they are replacing. If it was well documented and tested it wouldn't need to be replaced, after all... After months (or even years) without delivering value, the pressure is on to "ship it", so the new system is launched to handle a subset of what...

4. Meta AI Has a Native Mac App Now, and It Seems Decent

Source: Daring Fireball | Published: 2026-09-06T20:44:27Z

Link:

<https://9to5mac.com/2026/08/19/meta-ai-is-now-available-as-a-more-capable-desktop-app-for-mac/>

Zac Hall, writing for 9to5Mac on August 19: The new Meta AI desktop app is version 1.0 beta at launch and weighs just 16MB once installed. Based on initial inspection, it runs natively on Apple silicon Macs with macOS 15 or later, using an AppKit and SwiftUI shell with WebKit for richer chat content. In other words, this is not an Electron app or a repackaged iPad release. It includes Mac-specific hooks, making the desktop experience more robust than the Meta AI

mobile app. Quick Invoke uses Option-Space to place a compact Meta AI composer over whatever you are doing. A separate dictation feature lets you hold a shortcut, speak, and have the resulting words typed into any app, from Mail and documents to code editors. Meta AI can also attach another Mac window to a conversation. With Screen Recording and Accessibility permission, the app reads the window's visible text and captures a screenshot for the next question. This is context gathering rather than computer control for now. Passes the Settings window sniff test too ? it actually has a Settings window. Meta, institutionally, seems more keen on native Mac apps than they are native Windows apps . ?

5. WorkOS: How to Give an Agent a Task Instead of a Token

Source: Daring Fireball | Published: 2026-09-06T19:06:27Z

Link: https://workos.com/blog/delegated-access-for-ai-agents?utm_source=daringfireball&utm_medium=newsletter&utm_campaign=q32026

My thanks to WorkOS for once again sponsoring Daring Fireball. Give an agent an access token and it spreads: into the context window, into tool call logs, into notes it keeps between steps. Each copy works from anywhere, long after the fact. Relay keeps the credential at WorkOS. Your agent names the user, WorkOS attaches that token, refreshes it, and releases it only to allowlisted hosts. A hijacked agent session is a live process you can kill. Learn how it works at WorkOS's blog. ?

6. Matt Haughey: ?The Car Industry A/B Tested Selling a Car With and Without CarPlay and the Results Are Not Shocking?

Source: Daring Fireball | Published: 2026-09-06T19:04:12Z

Link: <https://a.wholelottanothing.org/the-car-industry-a-b-tested-selling-the-same-car-with-and-without-carplay-and-the-results-are-not-shocking/>

Matt Haughey: One case in point is the Honda Prologue EV. It's a partnership between Honda and Chevrolet, where the Prologue is basically a Chevy Blazer EV with Honda badges, slight visual differences in the sheetmetal and interior, but all the underpinnings are straight up Chevy. Even though the keys say Honda on a Prologue, it's a General Motors key design with a Honda logo over it. But one big difference the Prologue has from the Blazer it is based on is this: it includes Apple CarPlay and Android Auto integration. [...] For the years these cars were sold together, the Honda outsold the Blazer by 43%, 73%, and 166% in each time period. Today in 2026, people are buying a Honda Prologue over one and a half times more often than the Chevy Blazer EV the Prologue is based on. It's not a perfect test. It's reasonable to argue that Honda has a better brand and reputation than Chevrolet, and is just plain better at selling cars. But the trendline over the last three years is pretty stark. ?

7. It took a year to ship WebAssembly in Anubis

Source: Xe Iaso's blog | Published: Sun, 06 Sep 2026 00:00:00 GMT

Link: <https://anubis.techaro.lol/blog/2026/anubis-wasm/>

After a year of work, hundreds of commits, 5 generations of pull requests, dozens of tests, rewriting part of Anubis in Rust, the first compiler bug of my career, and at least three times making my tower run out of ram I think I have finally done it. The next version of Anubis will ship with WebAssembly-based proof of work checks that admins can enable in their thresholds or bot rules: thresholds : - name : moderate - suspicion expression : all : - weight > = 10 - weight action : CHALLENGE challenge : algorithm : argon2id difficulty : 6 This makes Anubis challenges use a memory-hard proof of work function (argon2id) instead of just a CPU hard one. It also means that the "hey Claude vibeslop me a CUDA Anubis solver" route is on its way to being fundamentally dead. Overall, this project taught me a lot of things about how WebAssembly works in practice, what the rough edges are when operating on the bleeding edge like I am, and led to the first genuine compiler bug of my career. Today I want to take you through the journey and process behind adding these WebAssembly based proof of work checks and all the problems that came up along the way. Why use WebAssembly in the first place? T...

8. Conflict resolution is ?fun?

Source: Xe Iaso's blog | Published: Tue, 01 Sep 2026 00:00:00 GMT

Link: <https://www.tigrisdata.com/blog/conflict-resolution-is-fun/>

One of the hardest problems in distributed systems is conflict resolution, or the same basic problem as merge conflicts in Git. Git merge conflicts happen when your branch differs from upstream in a way that Git can't easily work its way around so it exposes both sides of the changes to humans and has the human (or their agent) figure out which side is "correct?". Distributed systems don't really have this same flow as the scale of changes is often impossible for any human or team of humans to keep up with. As a result, we really want to have our own business logic define which side of a conflict wins. FoundationDB doesn't let us do that out of the box, so we had to make our own layer. The neat part about being able to do this is that this lets us control the replication behaviour of buckets based on user needs. The three main ways it differs are with single-region buckets, multi/dual-region buckets, and global buckets. I'm going to cover the replication differences in the order of complexity. Single-region buckets One of the easiest conflict resolution methods we have is a single-region bucket, which prevents any need for it. In this mode all actions are reverse proxied to the...

9. The purpose of DNS is to spread scams

Source: Terence Eden's Blog | Published: Sun, 06 Sep 2026 11:34:20 +0000

Link: <https://shkspr.mobi/blog/2026/09/the-purpose-of-dns-is-to-spread-scams/>

I imagine everyone here has received an unsolicited message telling them that their tax is overdue and that they urgently need to visit Genuine-Tax-Payment-Website.fart or that a parcel is delayed at customs and you can pay a small sum for its release at Almost-The-Right-Acronym.ak You know it is a scam. Most people just mark as spam and move on with their day. But a significant number of people ?

10. Pluralistic: How corporate America built a better Roach Motel (07 Sep 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 07 Sep 2026 05:29:25 +0000

Link: <https://pluralistic.net/2026/09/06/hotels-california/>

Today's links How corporate America built a better Roach Motel: Hostages beat customers every day. Hey look at this: Delights to delectate. Object permanence: 3-hole punches; Warner Bros v Warner Bros; Privacy wars get worse; "Weapons of Math Destruction"; HK elections go to pro-democracy reformers; Stock buybacks are swindles; NZ v Open source; MSFT patches DRM faster than other bugs; Wikipedia's worst arguments; Pirates x Iceland. Upcoming appearances: Brighton, London, Budapest, Edmonton, South Bend, Hudson, Calgary, Winnipeg, Vancouver, Victoria, Ottawa. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. How corporate America built a better Roach Motel (permalink) "If economists wished to study the horse, they wouldn't go and look at horses. They'd sit in their studies and say to themselves, 'What would I do if I were a horse?'" -Ely Devons Half a century ago, a group of lavishly financed economists from the University of Chicago (the "neoliberals") convinced governments all over the world to completely upend the way they treated monopolies. Up until then, th...

11. Weekly Update 520: The Unscripted Edition

Source: Troy Hunt | Published: Sun, 06 Sep 2026 23:31:36 GMT

Link: <https://www.troyhunt.com/weekly-update-520/>

I've started playing around with YouTube's "create video thumbnail", which hopefully will give me back a bit of time in my day (it used to be a manual job in Photoshop) and be a bit more interesting. And on that note, the imagery it&

12. Have the frontier labs mixed up AI safety and security?

Source: Martin Alderson | Published: Sun, 06 Sep 2026 00:00:00 GMT

Link: https://martinalderson.com/posts/ai-safety-vs-security/?utm_source=rss&utm_medium=rss&utm_campaign=feed

The recent agent sandbox escapes at OpenAI and Anthropic look less like a technical failure and

more like a philosophical one - treating security controls as though they only need to work most of the time.

13. Debian Code Search: Fast TurboPFor with Go SIMD

Source: Michael Stapelbergs Website | Published: 2026-09-06T09:00:00+02:00

Link: <https://michael.stapelberg.ch/posts/2026-09-06-dcs-fast-turbopfor-go-simd/>

This August, I accomplished what I wanted for many years: I deleted the last cgo dependency in Debian Code Search! This was made possible by Go's recently introduced SIMD support, because now we can implement the TurboPFor integer compression format as efficiently ? more efficiently, in fact, by using the newer AVX512 instruction set! ? as the reference implementation. Background: Why does DCS need a fast Integer Codec? Debian Code Search (DCS) is a search engine that allows searching all the Open Source source code within Debian, with either literal search expressions or regular expression search queries. A search engine uses an inverted index: a map from term to documents containing the term. Each document is typically represented most efficiently by using an id, so the index consists of many lists of document ids. When searching, it is important to quickly decode these lists to answer the search query. However, there is a point of diminishing returns where the decoding speed, even though it can still be measurably improved quite a bit, no longer influences the overall query duration. From 2012 (its inception) to 2019, Debian Code Search used to use a small index format, and que...