

1. Generating running routes with GPT-6 Astra and ChatGPT Work

Source: Simon Willison's Weblog | Published: 2026-09-12T23:56:42+00:00

Link: <https://simonwillison.net/2026/Sep/12/astra-running-routes/>

Here's a neat thing I had ChatGPT Work with GPT-6 Astra (Max) do this morning: I live at . Figure out 5K and 10K running routes from me that loop from my house. Use OSM data. It worked for 27 minutes and produced exactly what I'd asked for, as both an embedded visualization and downloadable GPX file and GeoJSON files. Here's that 5K route: When I asked it how it had created the route, it replied: I used Nominatim to locate the address and Overpass to download local OpenStreetMap roads and trails , then calculated the loops locally. Frustratingly, the actual code it ran and exact details of what it did weren't visible to me in the ChatGPT UI. I see this lack of transparency is an anti-feature. By the time I thought to ask for a copy of the Python code it had used, ChatGPT was unable to provide it. This appears to be because the thread had been compacted. I think any LLM system that uses compaction needs to both preserve the pre-compacted text and make that text available via agent tool calls, to protect against this kind of problem. As for displaying the map to me, that used the visualize skill . It created a file called /workspace/el-granada-5k-share.html to embed directly into th...

2. California Brown Pelican

Source: Simon Willison's Weblog | Published: 2026-09-12T21:16:09+00:00

Link: <https://simonwillison.net/2026/Sep/12/sighting-399708714/>

California Brown Pelican, in San Mateo County, CA, US The Pacifica Pier shut down at the start of June after a crack in the concrete walkway made access to the pier unsafe. It has since been entirely taken over by pelicans! Tags: wildlife

3. Quoting Paul Ford

Source: Simon Willison's Weblog | Published: 2026-09-12T18:00:21+00:00

Link: <https://simonwillison.net/2026/Sep/12/paul-ford/>

For a while, I must admit, it looked as if software developer roles like mine were done for. How could we fight against tireless robots? But our industry is slowly realizing that making truly cutting-edge software still requires humans to think and work together, to maximize their skill sets and to practice their respective crafts. A.I. can write very good software, but it also makes it easy to do someone else's job badly, which is part of why all those projects fail. Now that everyone can code, it's become clearer why many shouldn't. ? Paul Ford , A.I. Was Supposed to Give Us New Killer Apps. What Happened? Tags: paul-ford , generative-ai , deep-blue , ai , llms

4. Unread 5.0

Source: Daring Fireball | Published: 2026-09-12T19:27:37Z

Link: <https://www.goldenhillsoftware.com/2026/09/unread-50/>

John Brayton: Unread 5.0 is available now from the App Store . This update adds these improvements and more: Improvements around hero images, article list thumbnails, and widget thumbnails Syncing with FreshRSS and Miniflux The ability to quickly mark old articles read I'm a NetNewsWire man personally, but I've long admired Unread as a different, very good take on a feed reader. Just a gorgeous app. One of the very best things about RSS is that feed readers take such different shapes and approaches. ?

5. Post Peek ? Litterbox-Inspired Tweet Viewing Extension for Chrome

Source: Daring Fireball | Published: 2026-09-12T18:05:01Z

Link: <https://github.com/tsvb/post-peek>

Tim VanBenschoten: Post Peek is an independent, open-source project inspired by Litterbox , the Safari extension by Zhenyi Tan (And a Dinosaur). ?

6. Where Are the AI-Generated Killer Apps?

Source: Daring Fireball | Published: 2026-09-12T16:23:57Z

Link: https://www.nytimes.com/2026/09/12/opinion/ai-software-coding-apps.html?unlocked_article_code=1.AIE.C9EE.H2j5a1eHFUbQ&smid=nytcore-ios-share

Paul Ford, writing for The New York Times (gift link): For a while, I must admit, it looked as if software developer roles like mine were done for. How could we fight against tireless robots? But our industry is slowly realizing that making truly cutting-edge software still requires humans to think and work together, to maximize their skill sets and to practice their respective crafts. A.I. can write very good software, but it also makes it easy to do someone else's job badly, which is part of why all those projects fail. Now that everyone can code, it's become clearer why many shouldn't. Ford, linking to his own column on Bluesky : I just want one killer app before we get a killer robot. ?

7. AI is breaking our proxies for expertise

Source: seangoedecke.com RSS feed | Published: Sun, 13 Sep 2026 00:00:00 GMT

Link: <https://seangoedecke.com/ai-is-breaking-our-proxies-for-expertise/>

Mathematicians are broadly not anti-AI. They're more culturally open to using AI as a tool than, say, artists or writers ¹ . However, now that more and more genuinely prestigious problems have fallen to AI, that might be changing. Almost five thousand mathematicians (including twenty-five Fields medalists) have signed a declaration called A Severe Misalignment of AI in Mathematics . The core argument goes something like this: In recent months, the success of AI in solving major mathematical problems has made headlines even outside mathematical circles. But solving problems is only a tool and proxy for achieving the primary goal of conceptual understanding and insight. Forgetting this in the world of AI may turn the tool against the primary goal. Indeed, the mass production at faster and faster pace of ?true/false? statements could destroy fertile ground instead of breathing life into new ideas. A lot of people online have interpreted this as the expected complaint from any field that gets automated: translators did it, artists and programmers have been doing it, and now it's the turn of the mathematicians. I think this is too dismissive. Understanding the concrete problem mathemati...

8. Microcode in Intel's 8087 floating-point chip: the scale instruction

Source: Ken Shirriff's blog | Published: 2026-09-12T08:45:16.101-07:00

Link: <http://www.righto.com/2026/09/8087-microcode-reverse-engineering-fscale.html>

In the 1970s, floating-point arithmetic was a mess. Computer manufacturers had a dozen incompatible arithmetic standards. Moreover, floating-point systems were designed around hardware simplicity rather than mathematical rigor, leading to problems with numerical stability. This changed when Intel introduced the 8087 floating-point coprocessor chip in 1980, designed to be as accurate as possible, even in the corner cases. The 8087 became popular because it could be installed in the IBM PC, making floating-point operations up to 100 times faster in applications ranging from spreadsheets to CAD. But more importantly, the 8087 became the floating-point standard used by most computers today. The 8087 implemented its instructions in complex low-level code called microcode. I'm part of a group, the Opcode Collective, that is reverse-engineering this microcode, and I've recently made some progress. In this post, I examine the microcode for one of the 8087's instructions? FSCALE ?and describe how this microcode works. The FSCALE (Floating-point Scale) instruction provides a quick way to scale a number by a power of two, much faster than a multiplication. I figured that FSCALE was a simple,...

9. P(doom)

Source: Armin Ronacher's Thoughts and Writings | Published: 2026-09-12T00:00:00+00:00

Link: <https://lucumr.pocoo.org/2026/9/12/pdoom/>

This week some flavor of ?AI is going to kill us all? went viral. In particular one where an employee put his personal probability of that happening above 10%. Which made me go to the Wikipedia page of P(doom) and I realized that Dario Amodei's apparent probability of something

bad happening seems to be between 10-25%. And well, Dario then wrote about pacing the frontier . And Sam read it and wants to pace too . And well, so does Musk . I encourage you strongly to read the post, because I think it's a good one. And yet, when I read the post I could not help but feel in strong opposition to it, despite the fact that I think I'm on the same page with regard to all observations and, to a large degree, the concerns. I thought it might be interesting to write down my present-day thoughts on this, even if for no other reason than for myself to look back at it a year or two from now. What Is Doom? What I really appreciate about Dario's post is that he lays out a scenario that is not a huge stretch but also one that describes a clear, unfortunate outcome we should fight: persistent botnets and other forms of nuisance. And well, we don't have to look very far to see the issues left and rig...

10. Everyone should slow down AI development except for me

Source: Xe Iaso's blog | Published: Sat, 12 Sep 2026 00:00:00 GMT

Link: <https://xeiaso.net/notes/2026/everyone-slowdown-but-me/>

The power of science is staggering! Every day we're making new advancements in the field of generative artificial intelligence via judicious application of large language model inference technology. However, as we approach the next critical threshold in capability, we must look forward to ensure that we're not inadvertently creating a bad user experience in the form of mass societal collapse due to our technology replacing organic contributions in the workplace. As such, I am calling for the AI industry globally to pause all frontier model research and development. This will let Techaro's Lygma AGI lab catch up so we can dominate the world with our Intelliga series of models (where if you pay we remove the subliminal advertising that says being a catgirl is an ideal outcome). We want to give people the ability to have cat ears and we believe that AGI is the only way to do it. Our plan is to invent artificial general intelligence and then ask it to figure out how to give people cat ears. I think that this is a flawless plan that has absolutely no downsides for anyone involved, and if we do this together, bro, we can absolutely make sure that AGI development happens at a pace that i...

11. ActivityPub - How to send an updated user profile to Mastodon and the Fediverse

Source: Terence Eden's Blog | Published: Sat, 12 Sep 2026 11:34:02 +0000

Link: <https://shkspr.mobi/blog/2026/09/activitypub-how-to-send-an-updated-user-profile-to-mastodon-and-the-fediverse/>

Let's suppose you've updated the description of your ActivityPub account from "World's Number 1 Taylor Swift Fan" to "This account is now a Nickleback Truther". How do you let the rest of the Fediverse know that you've changed your allegiance? By default, most Mastodon instances won't periodically poll your account information just to see if you've updated it. So how does the information get?

12. Reading List ? 09/12/2026

Source: Construction Physics | Published: Sat, 12 Sep 2026 12:32:46 GMT

Link: <https://www.construction-physics.com/p/reading-list-09122026>

City permitting speeds, more efficient ground source heat pump drilling, an Amazon Air crash in Miami, cooling paint, and more.

13. This Week in Package Management: 12 September 2026

Source: Andrew Nesbitt | Published: 2026-09-12T10:00:00+00:00

Link: <https://nesbitt.io/2026/09/12/this-week-in-package-management.html>

Week seventeen of the roundup, built from the package manager OPML feed collection and whatever I've posted or boosted on Mastodon . Releases pnpm 12.4 can manage npm, Cargo, and Python dependencies in the same workspace: enable Cargo and Python in pnpm-workspace.yaml and one pnpm install handles all three, each from its own manifest, backed by pnpm's shared store, authentication, and offline cache. It also adds a pnpm pipeline command for CI-style task graphs and binaries for six more architectures. 12.4.1 fixes installs on filesystems that refuse hard links or copy-on-write clones. On the 11.x line, 11.26 backports workspace: ranges in catalogs, the trust-policy flags on remove and update , and pnpm change check for CI.

RubyGems and Bundler 4.1.0.beta1 adds content-addressable gem support, ML-DSA post-quantum signatures for the signed-gem workflow, an opt-in OS credential store for gem and bundler credentials, and a --cooldown flag on gem install , update and outdated . Pixi 0.80.0 adds experimental conda-script support: a comment header in a script file declares channels, dependencies, and an entrypoint, and Pixi solves and caches the environment then runs the script without a...

14. A Dick Smith VZ200 without the Dick Smith (but with a serial port)

Source: Old Vintage Computing Research | Published: 2026-09-12T19:17:52.821-07:00

Link: <https://oldvcr.blogspot.com/2026/09/a-dick-smith-vz200-without-dick-smith.html>

Australians! They walk among us ! Do not be deceived by those charming faces! They might be your parent! They might sleep in the same bed as you at night! A half-breed might write the very blog you read! For the preservation of our precious bodily fluids , we must remove the scurrilous larrikin influence of Australia upon our American home computers, and I know exactly where to start! Why allow Dick Smith's name (even though he'd already sold Dick Smith Electronics' controlling interest to Woolies by then but stop ruining my intro) to corrupt this, um, rubber-keyed diminutive beige home computer when we can return it to its prior, pristine, plasticky state as it once rolled out from a glorious Asian factory ? Let us restore it to its purest virginal roots! Let this American, possibly sold in Canada, original of a Hong Kong-manufactured super-cheap home computer that was also sold in Europe but never mind that flourish once more! ... after, of course, we compare this seppo VTech VZ200 with the Dick Smith units and the bumper crop of crap American home computers in 1983, and fix the keyboard and video. Then , to avoid wrecking it with a botched logic board repair, we'll bolt on a US...

15. Pluralistic: LLMs are real, AI is fake (12 Sep 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Sat, 12 Sep 2026 11:14:10 +0000

Link: <https://pluralistic.net/2026/09/12/god-in-the-box/>

Today's links LLMs are real, AI is fake: No, it didn't "go rogue." Hey look at this: Delights to delectate. Object permanence: Why 9/11 Means We Must Support My Politics; Blogs x 9/11; Jimmy Wales v Britannica's EiC; Hollywood astroturfs Australia; Agents v queer YA; Soft landings for dirty cops; Leaked Stingray manual. Upcoming appearances: Budapest, Edmonton, South Bend, Hudson, Calgary, Winnipeg, Vancouver, Victoria, Ottawa. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. LLMs are real, AI is fake (permalink) Once you understand the corporate culture of AI "hyperscalers" consists primarily of everyone cooking their brains by locking themselves in the bathroom, holding flashlights under their chins, and saying "Aaaaaaaaay Eyeeeeeee" until they wet themselves in terror, a lot of things snap into focus: <https://pluralistic.net/2023/06/04/yyyyyy-eyeeeeee/> It explains how a company can simultaneously be staffing up an enterprise sales division while also constantly freaking out at the thought that its product has "a 10% chance of ending humanity": <https://www.l...>