

1. Data Broker Radaris Loses Domains in Privacy Fight

Source: Krebs on Security | Published: Wed, 16 Sep 2026 18:14:22 +0000

Link: <https://krebsonsecurity.com/2026/09/data-broker-radaris-loses-domains-in-privacy-fight/>

The consumer data broker Radaris.com has long had a reputation for ignoring requests to remove personal information from its vast empire of people-search services online. That reputation caught up with the company recently in a lawsuit alleging Radaris violated a New Jersey privacy law that provides for hefty fines against data brokers that publish personal information on state law enforcement officials. In the face of repeated stonewalling and prevarication by attorneys for Radaris, the judge in the case ordered that radaris.com and more than a dozen other data broker domains be transferred to the plaintiffs.

2. datasette 1.0a40

Source: Simon Willison's Weblog | Published: 2026-09-16T23:51:43+00:00

Link: <https://simonwillison.net/2026/Sep/16/datasette/>

Release: datasette 1.0a40 Same security fix as 0.65.5 , plus some neat new features and bug fixes: Plugins can now launch and manage background tasks using the new `datasette.add_background_task()` method. Thanks, Alex Garcia . I've migrated Datasette to `httpx2` for features like the internal `datasette.client.get()` method. A whole lot of bug fixes , many of them stemming from a recent effort to triage issues for a 1.0 stable release. Tags: security , datasette

3. datasette 0.65.5

Source: Simon Willison's Weblog | Published: 2026-09-16T23:51:08+00:00

Link: <https://simonwillison.net/2026/Sep/16/datasette-2/>

Release: datasette 0.65.5 Security fix for an issue where a trailing newline in a requested table name could bypass table permissions and expose private rows, reported by dpfkdemtp in GHSA-h547-rmjf-5m2m . Tags: security , datasette

4. Claude Cowork and chat are now one Claude

Source: Simon Willison's Weblog | Published: 2026-09-16T18:09:49+00:00

Link: <https://simonwillison.net/2026/Sep/16/one-claude/>

Claude Cowork and chat are now one Claude In hopefully good news for anyone who, like me, was increasingly confused at Cowork v.s. Claude v.s. Claude Code: Starting today, Claude Cowork and chat are merging into one Claude. Bring a quick question, or hand over a report due at noon, and Claude takes it from there, even after you've closed your laptop. [...] This is rolling out to Pro and Max plans first, in the Claude app on web, desktop, and mobile over the coming weeks to existing and new users on these plans. I guess this means Claude is becoming a general agent in its own right. Echoes of OpenAI renaming their Codex desktop app to ChatGPT a few weeks ago. On the one hand, this saves me some work, in that I was planning to finally figure out the boundaries between Cowork and regular Claude and write a follow-up to my piece on Understanding ChatGPT Work . I have a hunch that figuring out what this actually means in terms of features and surfaces is still going to take quite a bit of work. Via Hacker News Tags: ai , generative-ai , llms , anthropic , claude , general-agents

5. Jev means structured output is interesting again

Source: seangoedecke.com RSS feed | Published: Wed, 16 Sep 2026 00:00:00 GMT

Link: <https://seangoedecke.com/jev-means-structured-output-is-interesting-again/>

I don't write blog posts about new models. That's Simon Willison's beat, and he's very good at it. But I want to write about Jev , which is a different kind of AI model: a "System One" model. As it turns out, it's not that different from an ordinary LLM with structured output, but the interface it uses is very cool and I hope it becomes more widespread. How Jev is different from LLMs Ordinary LLMs take in some human-language prompt and produce some

human-language output. They do so autoregressively : first they produce one token, then the next, then the next, and so on. User: Who invented the sandwich? LLM: The sandwich was invented by the Earl of Sandwich. This makes them extremely flexible, since they can do literally anything a computer can do. But it also makes them slow and weird. Slow, because they have to run a whole new generation pass per-token, and weird, because the space of human language is so broad that you can get really odd behavior from a model trained on it. Jev takes a human-language prompt, but it does not produce human-language output. It only produces structured output. User: { state: "What color is the sky?", choices: ["blue", "red", "yellow"] } Jev: { a...

6. Why do Microsoft job levels start in the high 50's instead of starting at a sane number like 1?

Source: The Old New Thing | Published: Tue, 15 Sep 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260915-00/?p=112700>

It used to start at 1, but that was the old system. The post Why do Microsoft job levels start in the high 50's instead of starting at a sane number like 1? appeared first on The Old New Thing .

7. Xcode 27.2 Is Out, but 27.1 Is Not, Which Means Developers Still Can't Get Started on Duo-Adapted Apps

Source: Daring Fireball | Published: 2026-09-16T22:40:41Z

Link: https://developer.apple.com/documentation/xcode-release-notes/xcode-27_2-release-notes

A note flagged 'Important' atop the Xcode 27.2 release notes: The iOS SDK and simulator support for iPhone Duo will be available later this month in an upcoming release of Xcode 27.1.

Adapting apps for the Duo is not just a simple 'recompile with the new SDK'. It's different.

For just about all apps it's going to be a significant amount of work, and for some apps it's

going to be a lot of work. But without Xcode 27.1, developers cannot get started. Developers

wanted this SDK a week ago. And when day one preorder buyers start using their Duos on October

23, they're going to wish Apple had shipped this SDK last week too. ?

8. Is That a Duo in Ternus's Pocket or Was He Just Happy That Apple TV Shows Won 28 Emmys?

Source: Daring Fireball | Published: 2026-09-16T22:21:30Z

Link: <https://x.com/DEADLINE/status/2099645816361332775>

Note the reaction from the Deadline reporter. ?

9. Woz Launches Merch Store

Source: Daring Fireball | Published: 2026-09-16T21:38:27Z

Link: <https://x.com/stevewoz/status/2100074363605397658>

Steve Wozniak: I've decided it's time to have a little more fun on X this year! ?? I get to

speak at some amazing events, meet fascinating people, hear great stories, and occasionally

find myself in places I never expected to be. So I figured? why not share some of those moments

here? He announced this on X, you will note, not 'social media'. ?

10. You can run git on object storage if you re-make packfiles

Source: Xe Iaso's blog | Published: Tue, 15 Sep 2026 00:00:00 GMT

Link: <https://www.tigrisdata.com/blog/objgit-packfiles/>

It sure seems that a bunch of companies are trying to ship a git product of some kind as of

late. Wonder why that is. Either way, I'm building a Git server backed by object storage as an

open-source project . It sounded simple enough to start: Git looks like a filesystem, so let's

use a filesystem as a translation layer on top of object storage to make Git speak object

storage. This model worked? ok, I guess? But it didn't work for real-world size repositories,

so I needed a different approach. Git stores everything in Objects , so why not store those as

objects in Tigris? Turns out Git packfiles and how they intersected with my (admittedly

somewhat terrible) filesystem shim were the main reason why it was slow. I ended up having to

invent my own packfile format with a columnar store that's object storage native. This is the

fruit of all of my performance analysis, metrics annotations, and more Texas-style distributed systems work than you can make your k8s cluster shake sticks at. This approach worked surprisingly well for production-sized repositories, so I'm sticking with this new Packfile format for now. It seems the least obtrusive change to make Git Objects feel like object...

11. Converting between cosine similarity and concentration ratio

Source: John D. Cook | Published: Wed, 16 Sep 2026 16:05:11 +0000

Link: <https://www.johndcook.com/blog/2026/09/16/concentration-ratio/>

I've written three posts on cosine similarity lately. The first looked at interpreting cosine similarity. The second looked at an approximation related to the first. The third looked at how ranking according to cosine similarity works better than cosine similarity itself. Normalized word vectors are points on a high dimensional sphere, and geometry in high [?] The post Converting between cosine similarity and concentration ratio first appeared on John D. Cook .

12. Coffee + milk ? latte

Source: John D. Cook | Published: Wed, 16 Sep 2026 15:06:58 +0000

Link: <https://www.johndcook.com/blog/2026/09/16/coffee-milk-latte/>

Yesterday I wrote about the canonical example of how vector embeddings of words add: ?king? ? ?man? + ?woman? ? ?queen? This should be interpreted as saying that the word vector for king, minus the word vector for man, plus the word vector for woman, is in some sense close to the word vector for queen. This post will [?] The post Coffee + milk ? latte first appeared on John D. Cook .

13. Fibonacci product

Source: John D. Cook | Published: Wed, 16 Sep 2026 12:04:09 +0000

Link: <https://www.johndcook.com/blog/2026/09/16/fibonacci-product/>

The product of four consecutive Fibonacci numbers equals the product of two consecutive integers. For example, $3 \times 5 \times 8 \times 13 = 39 \times 40$. I ran across this theorem in a note [1] that says ?The product of any four consecutive Fibonacci numbers is twice a triangular number.? Since triangular numbers have [?] The post Fibonacci product first appeared on John D. Cook .

14. How to get a DOI for your blog posts

Source: Terence Eden?s Blog | Published: Wed, 16 Sep 2026 11:34:28 +0000

Link: <https://shkspr.mobi/blog/2026/09/how-to-get-a-doi-for-your-blog-posts/>

Each new post on this blog now has a Digital Object Identifier. This post looks at the how and the why of getting one, whether it is useful, and any issues you might experience if you go down this path. Table of ContentsBackgroundGetting a DOI the easy wayLet's Go Rogue!Automatic Submission of New ContentManual Submission of Old ContentGetting the DOIGenerating your own DOIMaking the DOI?

15. Flock cameras are riddled with security vulnerabilities and hard-coded credentials

Source: micahflee | Published: Wed, 16 Sep 2026 20:48:47 GMT

Link: <https://micahflee.com/flock-cameras-are-riddled-with-security-vulnerabilities-and-hard-coded-credentials/>

This morning, DDoSecrets published an exciting new dataset: Filesystem images of the partitions from an in-use Flock ALPR camera. 404 Media and Wired published a joint investigation into it. I downloaded the dataset and am now thoroughly nerd-sniped. Hackers from a collective called stegan0gram collected the data. ?

16. The LLMs yearn for the spines

Source: Computer Things | Published: Wed, 16 Sep 2026 19:58:30 +0000

Link: <https://buttondown.com/hillelwayne/archive/the-llms-yearn-for-the-spines/>

Earlier this year I worked on a couple of AI-generated TLA+ specifications, and one thing I've noticed was that they all used the word "spine" somewhere. Then I saw some non-TLA+ public projects use "spine" too and got curious if it was a new LLMism (LLisM?). If it is, we should

see it appear in a discontinuous jump in the number of code projects that use "spine". The proper way to thoroughly explore this would be via the GitHub archive data dump but that'd take forever and I have a full time job now. Instead, I'm going to use GitHub Search . We can't filter code samples by date but can filter pull requests, which seems like an acceptable enough proxy. Here's a graph of the number of public PRs each year with "spine" in the title: I stopped 2026 at 2026-09-01, meaning the first nine months of this year have seen 20 times more "spine" PRs as all of 2025. Now it could be the case that that GitHub grew, like, a lot in the last year. Maybe the 20x increase is due there being 20x as many pull requests. We can account for this by comparing it against all public PRs : 1 If "spine" wasn't an LLM, we would expect only a 1.5xish increase from 2025, not a 20x increase. One other possible c...

17. Pluralistic: How an AI moratorium can save AI bosses (16 Sep 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Wed, 16 Sep 2026 07:52:58

+0000

Link: <https://pluralistic.net/2026/09/16/beggar-thy-neighbor/>

Today's links How an AI moratorium can save AI bosses: If you can't impose switching costs, just eliminate the competition. Hey look at this: Delights to delectate. Object permanence: Flash Worms; This Film is Not Yet Rated; Libdem copyright sabotage; Religion worth more than Big Tech; Geographic tubemap; Selective censorship resistance. Upcoming appearances: Budapest, Edmonton, Boston, South Bend, Hudson, Calgary, Winnipeg, Vancouver, Victoria, Ottawa. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. How an AI moratorium can save AI bosses (permalink) There's lots of reasons to believe the "hyperscaler" model of AI can never be profitable, and not just because of its gigantic expenditures and negative unit economics (the companies lose money with every new customer and every new use, and they lose more money with each generation of their products): <https://pluralistic.net/2025/09/27/econopocalypse/#subprime-intelligence> The industry strenuously denies this, of course. They insist that they are only days away from turning their balance sheets right side up. All...