

1. Raspberry Pi locks down Pi 5 RAM upgrades in firmware

Source: Jeff Geerling | Published: Mon, 21 Sep 2026 12:00:00 -0500

Link: <https://www.jeffgeerling.com/blog/2026/raspberry-pi-ram-lockdown/>

Raspberry Pi added a feature in their firmware that restricts users from swapping RAM chips , sometimes even RAM chips of the same capacity from other Raspberry Pi boards. Apparently there were people buying cheap 1 or 2 gig Pi 5s, soldering in some cheap and sometimes unreliable 8 gig RAM chips, and selling them as new 8 gig Pi 5s . Now I totally understand where Raspberry Pi's coming from. They don't wanna support Pis being sold like this, because flaky RAM and dodgy resellers are a huge problem. Especially with how profitable can be in the current AI bubble.

2. [Sponsor] Mux: Turn Your Video Into Context

Source: Daring Fireball | Published: 2026-09-21T23:04:15Z

Link: https://www.mux.com/?utm_campaign=fireball&utm_source=DF

Video isn't just something to stream; it's structured data you build with. Mux Robots, hosted AI workflows, turns video into context. Generate chapters, find key moments, translate audio, and more, through one API call, with no model hosting to maintain. Each workflow is evaluated against real video, not generic benchmarks. Configure the workflows once, and every new upload runs automatically. Mux is video infrastructure trusted by Patreon, Substack, and Perplexity. Start building for free. Use code FIREBALL for an extra \$50 credit. ?

3. Lex Friedman Brings Back Strategy

Source: Daring Fireball | Published: 2026-09-21T22:52:44Z

Link: <https://llexontech.org/strategy-is-back-and-im-the-developer>

I first recommended Strategy back in May 2009 , including a link to Friedman's review at Macworld . I described it thus: Strategy is a beautiful \$2 Risk-like strategy game for the iPhone. Lots of fun and the gameplay has a great feel. Still lots of fun and still has a great feel. It now costs \$5 , but, if you're an existing user ? even if you only paid \$2 back when I told you to 17 years ago ? it's a free update. And it's a great story: the original Strategy game finally gave up the compatibility ghost with iOS 27 . Friedman made his own version of the game, including multiplayer support, but before releasing it on his own under a new name, he reached out to the original developers and it turns out everyone was happy to make a deal. Sometimes all you need to do is ask. (First thing you should do after installing the new version is go into settings and turn on ?Classic Hex Edges?. The default style ? Lex edges? ? doesn't look right.) ?

4. GM Confirms They're Still Smoking Crack

Source: Daring Fireball | Published: 2026-09-21T21:52:50Z

Link: <https://x.com/JoannaStern/status/2102105565195288859>

Joanna Stern, on X: GM has confirmed it is NOT rolling phone projection to GM EVs. Instead this announcement is specific to JUST their new gas powered light-duty pickups, which have always supported Apple CarPlay and Android Auto. ?

5. It's Not Just True, It's False

Source: Software and Tech stories from an Insider - iDiallo.com | Published: Tue, 22 Sep 2026 01:49:51 GMT

Link: <https://idiallo.com/byte-size/its-not-just-true-its-false>

We've given in to using LLMs for writing at work. From emails and Slack comments to full-on documentation, everything is AI-generated. However, I've realized it only ever saves me time when I'm being lazy. That is, it's great at converting meetings into transcripts, summarizing those transcripts into action items, and keeping information accessible. But it's pretty bad because the result is ridden with invisible errors. But what counts as an error if you trust

the LLM and never meticulously read the output? And how much time are you really saving if you end up spending that time proofreading? I went back to a meeting transcript after a developer struggled to turn my AI-generated summary into a coherent ticket. Looking through the source transcript, I noticed that at one point it recorded the word "embroidery." I mean, we work for a large telecom company, the chances of me using that word are slim to none. Nevertheless, the AI summarized the conversation and included "embroidery" as an action item. In the words of Claude himself: It's not just true, it's false. I chuckled, but it made me think about all the BI slop we generate daily in the name of using AI technology. It feels help...

6. What's the highest legal FILETIME? Is it safe to use?

Source: The Old New Thing | Published: Mon, 21 Sep 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260921-00/?p=112711>

Being wary of how people use it and how it interacts with other types. The post What's the highest legal FILETIME ? Is it safe to use? appeared first on The Old New Thing .

7. Comparing exception behavior of magic statics, std::call_once, and std::async

Source: The Old New Thing | Published: Fri, 18 Sep 2026 14:00:00 +0000

Link: <https://devblogs.microsoft.com/oldnewthing/20260918-00/?p=112709>

Do you try again or give up? The post Comparing exception behavior of magic statics, std::call_once , and std:: async appeared first on The Old New Thing .

8. Jev introduces a new shape of LLM - System One, aka Decision Models

Source: Simon Willison's Weblog | Published: 2026-09-21T23:09:20+00:00

Link: <https://simonwillison.net/2026/Sep/21/jev/>

Last week TypeSafe AI unveiled Jev , their first example of a new category of model that they are calling "System One models" (I'm with Maggie Appleton, I think "decision models" is a better name for these). Jev is an interesting variant on the usual LLM format: it still accepts text inputs, but instead of text output it returns floating point numbers corresponding to categories, yes/no questions, ratings, and associated confidence scores. TypeSafe describe Jev like this: Think of Jev as a frontier-intelligence function call: unstructured state in, typed probabilistic decisions out. It's also very fast, and really cheap . Regular LLMs are priced in terms of input and output tokens, with output generally charged at significantly higher rates. Jev charges only for input - output is free - and the input price of their first model is \$0.042 per million tokens - cheaper even than OpenAI's GPT-5 Nano (\$0.05/million). Jev lets you ask questions about text or semi-structured data. You compose a "state" object containing a string, array of strings, or set of name-value pairs - this might describe an article, or a customer, or any other kind of record. You then send that to their API with o...

9. Cloudflare Python Workers are now generally available

Source: Simon Willison's Weblog | Published: 2026-09-21T22:25:44+00:00

Link: <https://simonwillison.net/2026/Sep/21/cloudflare-python-worker/>

Cloudflare Python Workers are now generally available After a two year preview, Cloudflare's support for running Python code in their server-side Workers platform is now stable: "Python is now a first-class, fully supported language on the Cloudflare Developer Platform". A neat thing about this is how it works. Cloudflare are running Python compiled to WebAssembly via Pyodide in their V8-based workerd runtime. This comes with some limitations, documented here - most notably both multiprocessing and threading are non-functional in the WebAssembly VM. One particularly interesting detail of this is the local development environment story - their pywrangler development tool (confusingly packaged as workers-py on PyPI) runs a full local simulation of their stack, including executing code with Pyodide in WebAssembly in V8 in a 123MB workerd binary, which for me ended up in node_modules/@cloudflare/workerd-darwin-arm64/bin/workerd . Python Workers represent a significant investment in the wider Python ecosystem by Cloudflare. The release announcement is credited to Gyeongjae Choi, Dominik Picheta, and Hood Chatham - Gyeongjae and Hood are both Pyodide core maintainers. Via Hacker News Ta...

10. MCP was always a bad idea?

Source: Simon Willison's Weblog | Published: 2026-09-20T20:24:41+00:00

Link: <https://simonwillison.net/2026/Sep/20/hn-49779718/>

My comment on MCP was always a bad idea? ? Hacker News. This article entirely misses the value that MCP brings today. Sure, there's almost no reason to use MCPs if you are running a full-blown terminal agent (Claude Code, Codex, Meta Muse, OpenClaw etc) with unfettered internet access - just let it call APIs directly. If you want to operate something that's less YOLO than that, you'll find yourself wanting: Control over exactly which external services it can access A way to handle authentication that doesn't allow the agent to directly access API keys A sensible UI to allow users to connect and authenticate further services Strong audit logging for what's going on MCP makes all of that so much easier to provide. Thinking MCP is obsolete because full coding agents don't need it misses out on all of the other things we might want to build. Tags: hacker-news , model-context-protocol

11. Haversine law

Source: John D. Cook | Published: Mon, 21 Sep 2026 23:13:22 +0000

Link: <https://www.johndcook.com/blog/2026/09/21/haversine-law/>

Suppose you want to solve a triangle. You know two sides and the angle between them. Then you can solve for the third side using the law of cosines. Now suppose you want to solve a big triangle, a triangle on the surface of the earth so large that the curvature of the earth matters. You [?] The post Haversine law first appeared on John D. Cook .

12. Abusing CD-Rs, DVD-Rs, and BD-Rs for Science

Source: Jayden's Blog | Published: 2026-09-05T11:41:00+00:00

Link: <https://jayd.ml/2026/09/05/burned-media-torture-test.html>

I've been on an optical media kick lately, having a grand old time burning discs and whatnot. One bit of advice that consistently comes up: don't let your burned discs get exposed to sunlight; the dye layer is UV after all, and since the sun is a deadly laser it can damage your precious bits. Naturally, this advice needs to be tested to determine what actually happens if you disregard it. I had a friendly LLM shart out some scripts to generate max-size iso images and sha256sum them from the disc, and burned using k3b: Verbatim CD-R (at max speed, 40x) Verbatim CD-RW (at max speed, 40x) Verbatim DVD-R (at max speed, 16x) I also have a second DVD-R that long story short had 615 buffer underruns while burning. It passed verification, but it seemed like a poor test subject. I exposed it though because it would be interesting if using Burnfree made it degrade faster Verbatim BD-R (at max speed, 12x) The CD and DVD were burned with a brand new (old stock) Hitachi-LG data storage GHD0N , the Blu-ray with a brand new LG WH16NS60 . The blanks were also just purchased from Amazon, so hopefully they are fresh. I verified each burn with k3b as well as my slop-script, and placed them in a zipl...

13. Writing Rust code that's faster than state-of-the-art libraries by asking agents to make the code faster

Source: Max Woolf's Blog | Published: Mon, 21 Sep 2026 09:30:00 -0700

Link: <https://minimaxir.com/2026/09/agent-iteration/>

c?mon, try doing a breakthrough

14. Cursed Napster

Source: Tedium: The Dull Side of the Internet. | Published: 2026-09-21T14:38:49Z

Link: <https://feed.tedium.co/link/15204/17468568/napster-ai-pivot>

If you haven't looked at the Napster website lately, get ready for some major AI-related indigestion. The former music-sharing rebel is now peddling AI digital twins of teachers and serving the Saudi market.

15. A Great Story Capturing the Essence of Being Online in 2026

Source: Jim Nielsen's Blog | Published: Mon, 21 Sep 2026 19:00:00 GMT

Link: <https://blog.jim-nielsen.com/2026/online-in-2026/>

I made a note of it , but Bryan Cantrill wrote a post about the ex-Anthropoc employee who pegged the probability of AI killing humanity in the next decade at >10%. The article is a good one and you should read it, but what I want to capture here is a side story from Cantrill on the Oxide and Friends podcast where he talks about how proud he was of this phrase from his post: "These ghoulish claims strike brazenly at the hearth?". He notes that someone online commented on that phrase, saying it sounded like a mix of classic literature and a line from The Simpsons . Amused, Cantrill says he came up with the line all by himself, so he knew it wasn't a classic literature reference nor a line from The Simpsons . But just to be sure, he Google'd it: "These ghoulish claims strike brazenly at the hearth?" with "The Simpsons?" tagged on to the end. To his surprise, Google's AI told him that the line was, in fact, a line from The Simpsons . Here's Google's response (as reported by Cantrill): During the Springfield mayoral debate on Channel 6, Sideshow Bob dramatically delivers the line "these ghoulish claims strike brazenly at the hearth of Springfield?" while turning the debate against Mayor...

16. Pluralistic: The Claude Delusion (21 Sep 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Mon, 21 Sep 2026 11:24:31 +0000

Link: <https://pluralistic.net/2026/09/21/sunsetting/>

Today's links The Claude Delusion: What if we're the ones having hallucinations? Hey look at this: Delights to delectate. Object permanence: 9/11 v sex; 9/11 v flags; US customs v diplomatic Vegemite; CIA torture survivors speak; HP's inky timebomb; "The Raven" pop-up; Forex v Zimbabwe's internet; Shirky v "expert Wikipedia": McSweeney's v RIAA; Italy's internet disconnection law; Facehugger mask; "Goliath"; Netzpolitik v German surveillance; Ted Cruz v internet founders; Epipencil; Framework laptops; "The Actual Star;" Pilots tell passengers to tackle suspected terrorists; \$2000 rotary phone; French DRM activists demand arrest; US child rearing costs up 40%; Calyx beats telcos; Wells Fargo whistleblower retaliation; David Graeber's "Mutual Aid"; Still censorship (even if it's not a First Amendment violation). Upcoming appearances: Berkeley, Edmonton, Boston, South Bend, Hudson, Calgary, Winnipeg, Paris, Vancouver, Victoria, Ottawa, Kilkenny, Montreal. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. The Claude Delusion (permalink) One of the less remarked-upon...