

1. SF October 14th: A Birds of a Feather Session on Agentic Engineering

Source: Simon Willison's Weblog | Published: 2026-09-23T02:53:19+00:00

Link: <https://simonwillison.net/2026/Sep/23/bof-agentic-engineering/>

SF October 14th: A Birds of a Feather Session on Agentic Engineering I'm hosting an evening event with Jesse Vincent in San Francisco on Wednesday 14th October for people who are building weird and interesting things with and on top of coding agents. Think of it as an agentic show-and-tell: ?Compare notes with other builders and experimenters on things you're trying, what you're learning, and what you haven't figured out yet. We're especially interested in work you haven't discussed publicly, odd experiments, or unfinished projects that don't have an obvious market. ?Expect one flowing conversation with an informal show-and-tell. Sharing something you're working on is encouraged but no presentation is required. This isn't about product pitches, it's about much earlier explorations than that. This agentic AI stuff is weird! Let's celebrate and lean into that weirdness. Tags: events , ai , generative-ai , llms , coding-agents , jesse-vincent , agentic-engineering

2. Claude Opus 5.5, GPT-6 Sol, GPT-6 Luna, and a new price war

Source: Simon Willison's Weblog | Published: 2026-09-22T23:46:41+00:00

Link: <https://simonwillison.net/2026/Sep/22/opus-and-sol-and-luna/>

Yesterday was Grok 4.7 (pelicans) and MiMo v2.6 Flash/Pro (more pelicans). Today Anthropic released Claude Opus 5.5 , and around an hour later OpenAI released GPT-6 Sol and GPT-6 Luna . It's going to take a while to get a good read on all of these new models, but here are my impressions so far. GPT-6 Sol and Luna are half the price of their GPT-5.6 equivalents GPT-5.6 Luna was already my favorite model for building applications against, because it combined excellent performance with being really cheap . Somehow GPT-6 Luna is half the price of that again - and GPT-6 Sol had a similar reduction compared to GPT-5.6 Sol. Here's what the pricing landscape looks like today: Model Input Cached input Output GPT-6 Luna \$0.10/M \$0.01/M \$0.50/M GPT-5.6 Luna \$0.20/M \$0.02/M \$1.20/M Grok 4.7 \$2/M \$0.50/M \$6/M GPT-6 Sol \$2/M \$0.20/M \$10/M GPT-5.6 Terra \$2/M \$0.20/M \$12/M Claude Opus 5.5 \$4/M \$0.20/M \$20/M GPT-5.6 Sol \$4/M \$0.40/M \$20/M Claude Fable 5.1 \$10/M \$0.25/M \$50/M GPT-6 Astra \$10/M \$1/M \$50/M Note that GPT-5.6 has a scheduled 25% price increase for November, so GPT-6 is half the price of the promotional pricing for those models. (With GPT-5.6 Terra priced the same as GPT-6 Sol, any...

3. Llm 0.36

Source: Simon Willison's Weblog | Published: 2026-09-22T18:48:37+00:00

Link: <https://simonwillison.net/2026/Sep/22/llm/>

Release: Llm 0.36 New OpenAI models: gpt-6-sol for GPT-6 Sol and gpt-6-luna for GPT-6 Luna . #1702 Model plugins can now declare supports_conversation = False for models that only accept single-turn prompts. LLM raises Llm.ConversationNotSupported when these models receive assistant or tool history, and llm chat rejects them before starting a session. See Models that do not support conversations . The first plugin to use this is Llm-typesafe . #1692 Reasoning traces in the Markdown output of Llm logs are now wrapped in tags. #1701 Plus bug fixes from five new contributors . Tags: openai , llm

4. Meta's Muse Logomark, Designed by Jessica Hische

Source: Daring Fireball | Published: 2026-09-22T22:51:09Z

Link: <https://thedieline.com/jessica-hische-metas-muse-and-the-ethical-landmines-of-design/>

Bill McCool, writing at Dieline: This past Wednesday, lettering artist, designer, and author Jessica Hische took to Threads to announce that she had worked alongside Meta's team to create the icon and logotype for their agentic AI, Muse. Acknowledging that people would be ?mad because it's AI and Meta,? Hische sought to explain at least why she had taken the gig in the first place. Obviously, she took some flak for working with Meta. While plenty of comments voiced support, just as many chastised her for even considering the project, which I can mostly

sum up as a general narrative thread of "How could you?" FWIW, it's pretty rich to criticize someone for working with Meta while you're very much using a Meta product to issue said criticism. But I digress. I like the Muse mark/icon. It's both an obvious "M" and yet at the same time, not an "M". It's friendly but not un-serious. And, as Meta chief AI officer Alexandr Wang noted today on X (not Threads), it doesn't look like a asshole. (Yet Meta AI's logo/icon is a little asshole-y.) ?

5. Meta's New Muse AI Agent Read Jason Aten's Messages Database

Source: Daring Fireball | Published: 2026-09-22T21:26:24Z

Link: <https://www.inc.com/jason-aten/metas-new-muse-ai-agent-read-my-private-messages-i-never-asked-it-to/91408202>

Jason Aten, writing for Inc: Then, yesterday, I was having a conversation with my Primary Technology podcast co-host, Stephen Robles, about the new iPhones. Moments later, I got a push notification from Muse suggesting that the conversation we were having would make for a good column and offered to put together research for me to write about. It even flagged a message from my editor about having a column ready for Monday. Not only had I not asked it to do that sort of thing, I never gave it permission to read my messages. In fact, I remember explicitly choosing not to let it have access to my messages, calendar, and other personal information. Stranger still was what happened when I asked Muse how it knew. Like I told Aten on Threads, Muse was right about one thing: it did give him a good idea for a column. Later on, Aten writes: I also want to be clear about something else: If your response to this is that, of course, Meta would do this ? it's Meta ? you are just letting them off the hook. You're normalizing behavior that should not be normal, and you should not let them off the hook. My response to this is a little more complex than that. At the top level, this is why I don't r...

6. Why Didn't Google Build Muse?

Source: Daring Fireball | Published: 2026-09-22T20:51:12Z

Link: <https://spyglass.org/why-didnt-google-build-muse/>

MG Siegler, writing at Spyglass: The number one response I've gotten to my thoughts on Meta's Muse personal assistant agent is some variation of "I will never trust Meta with my data." But the number two response is arguably more interesting (and related): "Why didn't Google build this?" The short answer to this is undoubtedly: they're working on it. But the longer answer is probably more interesting here too. Siegler's take is basically that Google lacks focus. They've got multiple AI agent products, none of which are hitting ? which as he points out has been their decades-long problem with messaging apps. Say what you want about Meta as a company and Mark Zuckerberg as a leader, but they (and he) have focus. I make fun of them (and him) for going all-in on "the metaverse" only to lose all interest in it just a few years later, but to play devil's advocate, that's how focus works. ?

7. Nathaniel Bowditch

Source: John D. Cook | Published: Tue, 22 Sep 2026 14:01:30 +0000

Link: <https://www.johndcook.com/blog/2026/09/22/nathaniel-bowditch/>

A couple days ago a friend told me about the book Carry On, Mr. Bowditch, a fictional account of the life of Nathaniel Bowditch (1773-1838). I've been listening to the book on Audible, and apparently it's only lightly fictionalized. Bowditch was a self-educated mathematician and astronomer, best known for his book The American Practical Navigator, [?] The post Nathaniel Bowditch first appeared on John D. Cook.

8. Are LLMs still surprisingly bad at some simple tasks?

Source: Terence Eden's Blog | Published: Tue, 22 Sep 2026 11:34:27 +0000

Link: <https://shkspr.mobi/blog/2026/09/are-llms-still-surprisingly-bad-at-some-simple-tasks/>

Last year I ran an experiment to test the ability of modern LLMs to correctly answer a relatively straightforward question. Every single one of them got it wrong. Some missed information, some made up false statements, none were right. Of course the fanbois variously claimed that I was holding it wrong, my prompts were shit, I should have chosen better

defaults, and - my favourite - that it?

9. Package Manager Threat Model, Revisited

Source: Andrew Nesbitt | Published: 2026-09-22T09:00:00+00:00

Link: <https://nesbitt.io/2026/09/22/package-manager-threat-model-revisited.html>

I've spent a fair chunk of the last four months pointing the package manager CWE list and the threat model questions at real package managers, in some cases with help. Ten targets, mostly clients, a couple of registries, a mix of tools I'd used and tools I'd only read about. I also re-ran the advisory search that produced the original CWE list, filtered to everything published since May, to see what everyone else had been finding in the same window. The two lists have now been merged into a maintained page at [/package-manager-threat-model/](#). Advisories 126 advisories against 23 package managers and registries between May and mid-September. About a quarter of everything ever filed against the same set of tools, which is mostly explained by three coordinated audits (one of pnpm producing twenty advisories over June and July, one of Homebrew, which I help maintain, producing thirteen, one of Flatpak producing ten in a single day) plus the OCI-registry implementations working through the same authorisation and DoS bugs the self-hosted language registries had five years earlier. Path traversal is still the top of the table by a distance, 33 advisories, and almost all of them come from...

10. Unfinished Work in Package Security

Source: Andrew Nesbitt | Published: 2026-09-22T09:00:00+00:00

Link: <https://nesbitt.io/2026/09/22/unfinished-work-in-package-security.html>

In May 2026, an Nx contributor installed a malicious dependency that had been published 77 minutes earlier. The repository had a seven-day cooldown configured, but its pinned pnpm version predated support for the setting. The client skipped the check and installed the package. The stolen credentials contributed to a malicious Nx Console extension release a week later. Package-manager defaults have improved even in the last few months, including release cooldowns and restrictions on installation scripts. Going back through what I've written this year, I keep running into work left to each consumer, in the gaps between controls that work on their own: checking that a configured policy takes effect across the tools they use, or finding someone who can publish a fix when a scanner turns up a vulnerability. Installation permissions Nx had already suffered a different compromise in August 2025, when an npm post-install script stole credentials and uploaded them to public GitHub repositories. The script ran with the developer's permissions, including access to unrelated credentials and the network, with no separate approval for either. pip's build-system interface can execute a backend...

11. Where're All The AI Chips?

Source: Ed Zitron's Where's Your Ed At | Published: Tue, 22 Sep 2026 14:35:46 GMT

Link: <https://www.wheresyoured.at/wherere-all-the-ai-chips/>

If you liked this piece, you should subscribe to my premium newsletter, and you can subscribe on the following links: \$70 a year, \$18 a quarter, or \$7 a month. In return you get a weekly premium newsletter including vast, detailed analyses of NVIDIA, Anthropic and OpenAI&

12. Cheaper LLM labelling

Source: Entropic Thoughts | Published: Tue, 22 Sep 2026 00:00:00 +0200

Link: <https://entropicthoughts.com/cheaper-llm-labeling>

I have a small project where I needed to label commits as either 'maintenance' or 'new development'. The obvious way to do it is with a cheap but relatively capable LLM, like GPT 5.6 Luna. I tested it on a small set of commits and manually verified its labelling, and it emitted the same label as I would have for the entire test set. That was good enough for me to roll out on a wider scale. If we have Simon Willison's llm CLI tool installed (and you should, it's great!), we can call it in a pipe from Perl, and read its response. My script had a loop that retried the request a few times, but without that bookkeeping, the code for this is simple enough. In[1]: sub classify { my (\$msg) = @_ ; my \$pid = open2(my \$output, my \$prompt, 'llm -m

```
openrouter/openai/gpt-5.6-luna' ); print $prompt prompt_template($msg); close $prompt; # Slurp the entire response and chomp off the # trailing newline. chomp(my $result = do { local $/; }); waitpid($pid, 0); return $result; } (Continue reading the full article on the web.)
```

13. Apple's Performa in-store demo software from 1993

Source: Downtown Doug Brown | Published: Tue, 22 Sep 2026 06:12:01 +0000

Link:

<https://www.downtowndougbrown.com/2026/09/apples-performa-in-store-demo-software-from-1993/>

Last year, I bought a Macintosh Performa 410 on eBay. I wasn't expecting it to be special in any way. The main reason I was interested in it was because I suspected it would have originally come with System 7.1P3, which was a weird Performa-specific OS version that nobody had preserved at the time. If [?]

14. tokens too cheap to meter

Source: the website of jyn | Published: 2026-09-16T00:00:00+00:00

Link: <https://jyn.dev/tokens-too-cheap-to-meter/>

tokens are going to be as cheap as electricity within the decade

15. Pluralistic: Bonta sold us out to Trump's oligarchs (22 Sep 2026)

Source: Pluralistic: Daily links from Cory Doctorow | Published: Tue, 22 Sep 2026 08:19:01 +0000

Link: <https://pluralistic.net/2026/09/22/happy-chudmas/>

Today's links Bonta sold us out to Trump's oligarchs: A "settlement" that promises less than nothing. Hey look at this: Delights to delectate. Object permanence: Warren on taxing the rich and class warfare; Psychology v reproducibility; Copyright troll sentenced for beating Uber driver; "Halloween Moon"; Enshittification of solar. Upcoming appearances: Berkeley, Edmonton, Boston, Brighton, South Bend, Hudson, Calgary, Winnipeg, Paris, OVancouver, Victoria, Ottawa, Kilkenny, Montreal. Recent appearances: Where I've been. Latest books: You keep readin' em, I'll keep writin' 'em. Upcoming books: Like I said, I'll keep writin' 'em. Colophon: All the rest. Bonta sold us out to Trump's oligarchs ([permalink](#)) Well, fuck. California Attorney General Rob Bonta just surrendered to the Trump-aligned Ellison billionaires who want to take over and destroy Warners, merging it with the chudded out husk they've made of Paramount, leaving these two colossal, corrupt, useless assholes to control Warners, Paramount and Tiktok: <https://www.reuters.com/legal/litigation/paramount-settles-with-california-other-states-clearing-major-hurdle-warner-bros-2026-09-21/> In announcing the settlement, Bonta's offi...

16. Weekly Update 522: Live From Oslo with Scott Helme

Source: Troy Hunt | Published: Wed, 23 Sep 2026 05:18:34 GMT

Link: <https://www.troyhunt.com/weekly-update-522/>

Heads up: the first 7 mins is a bit quiet until we worked out the external mic was misbehaving - sorry! But get through that and have a listen to Scott's experiences with how Report URI is identifying malware-infected machines within orgs, all due to CSP reporting. It&

17. On reading books

Source: Herman's blog | Published: 2026-09-22T13:27:56.346630+00:00

Link: <https://herman.bearblog.dev/on-reading-books/>

I grew up in a house stuffed full of books, with no television, and electricity for only an hour or two per day (while the generator was running). There were shelves full of old science fiction (Asimov and Heinlein), as well as books on history, some good old fantasy, and other smatterings collected from second-hand bookstores over the years. My mother read to us each night when we were children, starting with Dr Seuss, all the way to Harry Potter, where she did the spooky voices during the Voldemort scenes, which had me walking up and down the room with my fingers in my ears, humming to myself, until the scary parts were over. Naturally, this led to me reading a lot as a teenager and young adult. Up until I left home at 18, I would fall asleep each night reading by candlelight, which sounds very romantic and, in hindsight,

definitely contributed to good sleep growing up. At university, I read substantially less. I suddenly had a computer in my room, reliable electricity, and downloaded series to watch. I never stopped reading, but certainly felt my attention span erode as a result (I also think Family Guy circa 2012 rotted my brain irreparably). As you can probably tell by now, I...